




Autonomous AI and Education

Michael Gannotti

Contents

Autonomous AI and Education	3
Contents	4
Preface	5
The Collision, Not the Tour	5
Introduction	8
Eight Corrections and a Design Problem	8
Chapter 1	19
The Job the Student Must Still Do	19
The popular story, named so it can be taken apart	20
Performance is not learning	23
Cognitive offloading is old; generative AI extends it to the skills school exists to form . .	26
Students already use consumer chat — and the good tool is not equally held	28
The assessment broke before the policy did	30
Labor-market literacy versus cognitive formation	33
The design problem: which cognitive work stays inside the student	36
Chapter 2	39
Chat, Tutor, Adaptive System, Agent	39
Chapter 3	62
The Exam After the Chat Window Closes	62
Chapter 4	85
A Tool Lawful at 17 Can Be Illegal and Developmentally Wrong at 8	85
The legal stack is the design spec, not the appendix	86
Accessibility, language, and special education inside every band	92
Band 5-8 (K-grade 3): literacy <i>about</i> AI, teacher-facing only	93
Band 9-11 (grades 4-6): logged hint tutors in teacher rooms	96
Band 12-14 (grades 6-8): assigned Socratic help; companion chat is litigation	97
Band 15-18 (grades 9-12): disclosed partner; no acting agents	100
College and the adult bridge: transcripts, data classification, agent bans remain legitimate	102

Chapter 5	108
Managed Struggle, Not Rescue and Not Abandonment	108
The cognitive science that should constrain design	109
Mathematics: hint ladders help; answer-giving GPT harms; access is not tutoring	113
Writing: augmentation versus automation	116
Science: structured tutoring is not open inquiry with a chatbot	118
Reading and history: do not hide the work	119
Computer science: Copilot as worked examples at keystroke speed	120
Two-lane assessment, and the AIAS as vocabulary not evidence	122
Chapter 6	128
Judgement Needs Something to Judge	128
Chapter 7	143
Who Turns the Agent Off	143
Chapter 8	162
Red Lines That Map to Instruments, Not Vibes	162
Chapter 9	187
Who Hosts the Student Agent	187
Open-source does not settle governance	189
Three models: cloud GPT Edu, national wrapper, local open	193
Default host: the institution or a public platform	195
Teacher-facing agents: the real 2026 beachhead	198
Student-owned Hermes is a home tutor, not a school requirement	201
Practical checklist	204
Age-gated agent ladder	206
Chapter 10	209
Not “Can Prompt ChatGPT”	209
Coda	222
What an Honest Book Can Promise versus What It Cannot	222
A Note on Sources	227
Bibliography	228
1. Statutes, regulations, executive orders, court filings	228
2. Empirical research (journals, working papers, preprints labelled as such)	232
3. Policy, union, ministry, UNESCO / OECD / U.S. Department of Education	241
4. Product documentation (labelled vendor)	247
5. Journalism and ratings	250



Figure 1: Cover

Autonomous AI and Education

Michael Gannotti

This illustrated edition merges the documented final draft with its Chicago bibliography. Endnotes run in one series. Chapter-opening images are original still-lives, not portraits of persons or children.

Contents

- Preface — The Collision, Not the Tour
- Introduction — Eight Corrections and a Design Problem
- Chapter 1 — The Job the Student Must Still Do
- Chapter 2 — Chat, Tutor, Adaptive System, Agent
- Chapter 3 — The Exam After the Chat Window Closes
- Chapter 4 — A Tool Lawful at 17 Can Be Illegal and Developmentally Wrong at 8
- Chapter 5 — Managed Struggle, Not Rescue and Not Abandonment
- Chapter 6 — Judgement Needs Something to Judge
- Chapter 7 — Who Turns the Agent Off
- Chapter 8 — Red Lines That Map to Instruments, Not Vibes
- Chapter 9 — Who Hosts the Student Agent
- Chapter 10 — Not “Can Prompt ChatGPT”
- Coda — What an Honest Book Can Promise versus What It Cannot
- A Note on Sources
- Bibliography

Preface

The Collision, Not the Tour

Two documents, months and a product class apart.

The first is a field experiment. Hamsa Bastani, Osbert Bastani, Alp Sungu, Haosen Ge, Özge Kabakcı, and Rei Mariman, writing in the *Proceedings of the National Academy of Sciences* in 2025, gave nearly a thousand students in a large Turkish high school four ninety-minute sessions of mathematics practice under one of three conditions: no AI; unguarded GPT-4 in a ChatGPT-style interface; or the same interface with guardrails that withheld full solutions and used teacher-written correct answers and common-mistake hints. Unguarded GPT-4 raised practice scores about 48 percent. On the unaided exam that followed, those same students scored about 17 percent *worse* than classmates who never had AI. The hint tutor raised practice about 127 percent and left the unaided exam indistinguishable from control. Students in the unguarded arm did not perceive that they had learned less.¹ Performance was not learning. The artefact of the session was the wrong measure. The students could not see the loss.

The second is a product launch. On 11 August 2026, xAI published “Introducing Grok Bot.” Bots, the post said, “have their own computer” and “keep working 24/7.” The illustrated jobs were sales, CRM, inbox, expenses, bug reproduction, website deploy. None of them was educational. The pages did not mention schools, students, or teachers. There was no education SKU.² An autonomous teammate entered the culture of the reader this book is written for — a person who already lives with systems that take a goal, use tools, and continue after the user leaves — without entering the schoolhouse.

¹Hamsa Bastani, Osbert Bastani, Alp Sungu, Haosen Ge, Özge Kabakcı, and Rei Mariman, “Generative AI without Guardrails Can Harm Learning: Evidence from High School Mathematics,” *Proceedings of the National Academy of Sciences* 122, no. 26 (2025): e2422633122, <https://doi.org/10.1073/pnas.2422633122>. Field RCT, large Turkish high school, Fall 2023–24, nearly 1,000 students, grades 9–11, four 90-minute sessions. Assisted practice: GPT Base about +48 percent; GPT Tutor about +127 percent. Unaided exam: GPT Base about –17 percent versus never-AI; GPT Tutor indistinguishable from control. Students in the unguarded arm did not perceive the loss. Full table in Chapter 3; the preface uses the dissociation only.

²xAI, “Introducing Grok Bot,” August 11, 2026, <https://x.ai/news/introducing-grok-bot>. “Bots have their own computer... keep working 24/7.” Expansion August 26, 2026, to SuperGrok plans and Cursor Pro/Teams, <https://x.ai/news/grok-bot-more-plans>. Jobs illustrated are not educational. No education SKU. Product pages do not mention schools, students, or teachers. Grok Bot FAQ, <https://docs.x.ai/grok-bot/faq>, re-fetched after the landscape timeout: requires cloud data storage; Legacy Privacy Mode not supported; training opt-out follows Cursor account settings. FAQ does not mention schools, students, or teachers.

That pair is the book. Not ChatGPT-in-the-classroom. Not a tour of products. What happens when systems that keep working after the student leaves meet an institution whose job is to form unaided skill. The first document is the cleanest school demonstration that a polished homework is not a formed skill. The second is infrastructure chat cannot match, arriving without a pedagogy, a child-safety certification, or a district contract.

The popular story collapses what those documents keep apart. It collapses chat with agents: a window that waits for a prompt with a runtime that takes a goal and does not wait. It collapses Grok-in-El-Salvador with Grok Bot: an 11 December 2025 announcement of the assistant as a tutor across more than 5,000 schools, with no outcomes found, with a 2026 product that has no education SKU.³ It collapses Hermes 4, an open-weight model family, with Hermes Agent, a runtime Nous's own docs say that family is "not recommended" to sit inside.⁴ It collapses a fluent take-home with a skill that survives when the tool is taken away. This book exists because those collapses make policy, procurement, and pedagogy unintelligible in the same sentence.

This is not a product tour. Named systems appear because their own documentation forces a distinction, not because the chapters are a catalog. It is not a ban tract. Constrained, logged, hint-not-answer tutors have a causal record; unguarded chat has a different one; unattended agents in classrooms have almost no record at all. It is not a tutor-for-every-child brief. Bloom's two-sigma observation remains about human one-to-one tutoring. It is not an EEG morality play. A later chapter will have to refuse a preprint that will be offered as population cognitive decline. The preface does not.

Eight corrections, in miniature, are the contract the introduction will unpack. Chat is not an agent; tool use is not autonomy. Performance is not learning; guardrails are the treatment. There is no documented K–12 student-facing autonomous open-web agent; Broward's K–5 have zero student-facing AI, and Stanford CS221 bans agentic repo-builders on the project. Companion harm is litigation, not a hypothetical chapter. Detectors fail, and they fail unfairly. "Teachers who use AI will replace those who don't" is folklore, not US ED, NEA, or AFT doctrine. Hermes 4 is not Hermes Agent. El Salvador's Grok is not Grok Bot. The ninth correction is the design problem: which cognitive work must remain inside the student. The introduction is the map. Chapter 2 is where the four classes are taken apart at length. They are not taken apart here.

The plot is a collision already visible in 2026. Policy red lines — grading, IEPs, surveillance, elementary student-facing AI, companion chat — are hardening just as agent autonomy increases. Teacher-facing agents and constrained Socratic tutors are the beachhead. The student-facing unattended runtime is not. An honest book does not give Grok Bot an education positioning it did not

³xAI, "xAI and El Salvador Pioneer the World's First Nationwide AI Education Program," December 11, 2025, <https://x.ai/news/el-salvador-partnership>. Grok (the assistant) across more than 5,000 public schools over two years, "personalized learning to over one million students." Announced national program, not a documented completed deployment; no outcome evaluation found as of August 30, 2026. Not Grok Bot.

⁴Hermes 4: Teknium et al., "Hermes 4 Technical Report," arXiv:2508.18255, August 25, 2025, <https://arxiv.org/abs/2508.18255>, open-weight hybrid-reasoning LLM family. Quantitative benchmarks UNVERIFIED pending PDF. Hermes Agent: MIT-licensed autonomous runtime, launched February 2026 (vendor homepage dates the launch to February without a day; 25 February is the first-release date recorded in this book's responsible-use stream), <https://hermes-agent.org/>. Nous Portal docs: Hermes 4 (70B / 405B) is "not recommended for use inside Hermes Agent," <https://hermes-agent.nousresearch.com/docs/integrations/nous-portal>. Using the model family as a proxy for the agent is a factual error.

have on 26 August 2026, and does not treat an announcement as an outcome.

The contract with the reader is documentary. Chicago notes; load-bearing numbers cited on the page; longer source fights in the notes. Items the research marked UNVERIFIED stay flagged until a primary is read. Composite hypotheticals, if a later chapter needs one to show a design, are labelled *hypothetical* in the sentence that introduces them. Harm cases — Setzer, Raine — are named as litigation, with ages, dates, and filings. No portraits of children. No reconstructed chats. Vendor copy, detector scores, and the replacement slogan do not enter as findings. Open questions are left standing rather than filled with analogy. There is no What Works Clearinghouse report on generative tutors as of the access date of this research, 30 August 2026, and this book will not invent one.

If the collision is real, the reader needs the argument map before any product is explained at length. That is the introduction's job: eight corrections, a design problem, and ten chapters that do not tour.

Introduction

Eight Corrections and a Design Problem

The popular story is easy to recite, which is why it has to be named before it can be taken apart. Students have ChatGPT. Schools will soon give every child an agent. Detectors will catch the cheats. Teachers who do not use AI will be replaced. Critical thinking is a transferable bicycle-skill an agent can teach in the abstract. Every child will have a tutor. The homework is the evidence of the learning. The future is a teammate.

Almost none of that is how 2026 worked.

What students have, in the numbers that exist, is consumer chat. Pew, on 15 January 2025, reported that 26 percent of U.S. teens ages 13–17 had used ChatGPT for schoolwork, double the 13 percent of 2023. Awareness jumped most in the lowest-income households.⁵ That is chat use, not agent use. It is the fact underneath the panic, and it is not Hermes Agent on a laptop. New York City Public Schools and EDUCAUSE both treat policy clarity as lagging student behavior: the assignment broke before the guidance caught up. “Cheating” assumes the take-home still measures the intended skill. If the intended skill is a five-paragraph essay produced unsupervised, a fluent model makes the assignment dead. If the intended skill is to plan, judge, and defend a position, the same model is a studio tool or a crutch, depending on design. New York City’s March 2026 guidance still listed homework and integrity, cognitive offloading, and grade-band student-use as Playbook workstreams as of the fetched page. The nation’s largest district had traffic lights for *teachers* and had not yet published the student-use playbook.⁶ EDUCAUSE’s 2026 assessment survey — 438 faculty and staff; full PDF not fetched — is reported as finding that students need clearer use and non-use rules, and that most surveyed faculty use AI to create assessments and believe students use AI while taking them. Policy is lagging behavior. That lag is not a classroom mandate, and it is not evidence that every child has been issued an agent.

Access to the idea of the tool is not access to the good tool. Pew’s jump in awareness was largest

⁵Pew Research Center, “About a quarter of U.S. teens have used ChatGPT for schoolwork – double the share in 2023,” January 15, 2025, <https://www.pewresearch.org/short-reads/2025/01/15/about-a-quarter-of-us-teens-have-used-chatgpt-for-schoolwork-double-the-share-in-2023/>. Twenty-six percent of U.S. teens ages 13–17 (13 percent in 2023). Chat use, not agent use. Awareness jumped most in the lowest-income households.

⁶New York City Public Schools, “Guidance on Artificial Intelligence,” March 2026, <https://www.schools.nyc.gov/about-us/vision-and-mission/artificial-intelligence/guidance-on-artificial-intelligence>. Traffic lights for teachers; homework/integrity, cognitive offloading, and grade-band student-use listed as Playbook workstreams (targeted June 2026) as of the fetched page. Grade-band student-use Playbook still forthcoming.

in the lowest-income households. A Harvard undergraduate survey in 2024 found that among AI users, 40 percent of students with no financial aid paid for AI products versus 20 percent of students on aid.⁷ Homework-gap 2.0 is a paid-tier gap. A personal agent — memory, tools, skills — is more unequal than a chatbot. The popular story that “everyone has ChatGPT now” erases that gap and then, by calling chat an agent, erases it twice.

What schools bought is not an agent for every child. It is teacher-side chat and workflow tools under a data-protection agreement, constrained Socratic tutors, and pre-generative adaptive platforms. Agentic systems are entering teacher evenings and campus operations — Claude Cowork, OpenAI Workspace Agents off by default — not third-grade desks. Broward County’s K–5 have zero district-approved student-facing AI tools. Stanford CS221, in Spring 2026, bans agentic repo-builders on the project. No Hermes Agent or Grok Bot district DPA was found.

What the causal record shows is that unguarded GPT can harm subsequent learning and that engineered tutors can help immediately, and that neither finding is about Hermes-class agents. There is no identified learning randomised trial on those agents. There is no What Works Clearinghouse report on generative tutors. Age and law are design specifications, not afterthoughts. Companion harm is litigation, not a chapter of hypotheticals. Detectors fail, and they fail unfairly. The replacement slogan is folklore. The plot of this book is a collision, not a prediction: red lines on grading, IEPs, surveillance, elementary student-facing AI, and companion chat are hardening just as agent autonomy increases.

This introduction’s job is to make the reader unable to collapse the four product classes again, and unable to treat a homework artefact as a learning outcome. Eight corrections, each with a load-bearing source. Then a ninth: design is the treatment. Then a map of ten chapters and a list of questions the book will not fill.

Correction 1 — Chat is not an agent

A chatbot waits for a prompt and returns text. An agent is given a *goal*, plans multi-step work, uses tools — terminal, browser, files, APIs, other agents — continues after the user leaves, and keeps memory across sessions. That is the operational definition synthesized from vendor documentation, not from a standards body. Hermes Agent, Grok Bot, Cursor Cloud Agents, Claude Cowork, and OpenAI Workspace Agents all claim that infrastructure in their own words.

Hermes Agent: “Not a chatbot. Not a copilot. An agent that lives on your machine.” Persistent memory, cron, sub-agents, browser, terminals, messaging gateways. MIT-licensed, launched February 2026 (vendor homepage; 25 February is the responsible-use first-release date).⁸ Grok Bot, 11 August 2026: “Bots have their own computer... keep working 24/7... only come back when something needs your approval.”⁹ Cursor Cloud Agents: isolated cloud VMs; clone repos;

⁷Harvard Undergraduate Survey on Generative AI, arXiv:2406.00833 (2024), <https://arxiv.org/abs/2406.00833>: among AI users, 40 percent of students with no financial aid paid for AI products versus 20 percent of students on aid. Pew 2025 as in note 1.

⁸Nous Research, Hermes Agent, <https://hermes-agent.org/>. “Not a chatbot. Not a copilot. An agent that lives on your machine.” Vendor homepage: February 2026 without a day; 25 February is the responsible-use first-release date. MIT license.

⁹xAI, “Introducing Grok Bot” (11 Aug. 2026). Expansion August 26, 2026, <https://x.ai/news/grok-bot-more-plans>.

open pull requests; run while the laptop is closed. Claude Cowork, via Claude for Teachers: “Set a task up once and Claude handles it on a schedule.” OpenAI Workspace Agents: reusable workflows with a trigger, process, tools, and guardrails; research preview for eligible ChatGPT Edu workspaces; off by default; humans review before anything is sent.

A well-prompted ChatGPT session can *simulate* some of this with the human as the execution engine. Agents move execution off the human. That is the gap chat cannot close, and it is infrastructure, not “smarter answers”: a computer that is not the user’s laptop; unattended runtime; acting inside real apps rather than drafting text to paste; parallel specialists; durable memory and reusable skills.

Four classes, collapsed in ordinary talk into the mass noun “AI,” have to be held apart or the rest of the book is noise. A chatbot is a reactive thread: consumer ChatGPT, Grok on grok.com, the Gemini app. A tutor app is a pedagogical wrapper whose value is what it *refuses*: Khanmigo, SchoolAI Spaces, MagicStudent. An adaptive system or intelligent tutoring system sequences items from a learner model and does not browse the open web; ALEKS, MATHia, i-Ready, DreamBox are the names commonly cited — official pages were not fetched this pass; treat feature lists as UNVERIFIED. An autonomous system takes a goal, plans, uses tools, and continues after the user leaves. The book’s named objects sit only in the last row. Chapter 2 prints the table and walks it. The introduction’s job is to make collapse impossible before that walk.

Tool use is not autonomy. A chatbot with tools is still a chatbot if the human must stay in the loop for every step. Khanmigo, MagicSchool, SchoolAI, MagicStudent, Raina, Dot, Gemini Education, and ChatGPT for Teachers are not agents. Listing them as if they were is how a district thinks it has already procured the last row. It has not. The U.S. Department of Education’s 2026 supplemental priority defines AI via 15 U.S.C. § 9401(3). It does not define “agentic AI.” It does not authorize unconstrained student agents.¹⁰

Correction 2 — Performance is not learning

Conditions that make practice look fluent often produce weaker long-term skill. Nicholas Soderstrom and Robert Bjork’s 2015 review in *Perspectives on Psychological Science* is the load-bearing paper; Bjork and Bjork’s desirable difficulties — spacing, interleaving, generation, retrieval — slow immediate performance and improve later retention and transfer.¹¹ Generative AI, by default,

No education SKU. Jobs illustrated are not educational. FAQ, <https://docs.x.ai/grok-bot/faq>, re-fetched after the landscape timeout: cloud data storage required; Legacy Privacy Mode not supported; training opt-out follows Cursor settings; no schools, students, or teachers mentioned.

¹⁰“Final Priority and Definitions—Secretary’s Supplemental Priority and Definitions on Advancing Artificial Intelligence in Education,” 91 Fed. Reg. 18774 (April 13, 2026), effective May 13, 2026, <https://www.federalregister.gov/documents/2026/04/13/2026-07087/final-priority-and-definitions-secretarys-supplemental-priority-and-definitions-on-advancing>. AI defined via 15 U.S.C. § 9401(3). Does not define “agentic AI.” Discretionary-grant scoring tool, not a classroom mandate. Traces to Executive Order 14277 (April 23, 2025), an adoption-and-workforce order, not a child-safety regulation. Executive Order 14110 was rescinded January 20, 2025, by EO 14148.

¹¹Nicholas C. Soderstrom and Robert A. Bjork, “Learning versus Performance: An Integrative Review,” *Perspectives on Psychological Science* 10, no. 2 (2015): 176–199. Robert A. Bjork and Elizabeth L. Bjork, desirable difficulties, 2011 and 2020. Difficulty is desirable only when achievable (Bjork and Bjork, 2020).

is undesirable ease. It inflates session artefacts and can starve retrieval, generation, and error-correction. Immediate fluency is the wrong success metric. Unaided delayed tests are the right one.

Bastani and colleagues, in *PNAS* in 2025, is the cleanest school demonstration. Unguarded GPT-4 during high-school math practice raised scores about 48 percent; the same students then scored about 17 percent worse on an unaided exam than classmates who never had AI. A hint-and-grounded-solution tutor, GPT Tutor, raised practice about 127 percent and left the unaided exam indistinguishable from control. Students in the unguarded arm did not perceive that they had learned less.¹² Guardrails are the treatment. Default ChatGPT is a crutch students do not notice. The full table is Chapter 3. The dissociation is the correction: practice up, unaided exam down, for the unguarded arm; harm erased without a learning gain over business-as-usual, for the guarded one. Louis Deslauriers and colleagues, in 2019, documented the illusion of learning: students in a passive lecture felt they had learned more than students in an active classroom, while learning less.¹³ Bastani’s unguarded students did not perceive the loss. Fluency of the session is not a measure of the skill. Any assessment regime that grades only the session artefact has hired that illusion as a teaching assistant.

OECD Digital Education Outlook 2026, as reported on the listing page, claims that unguided generative AI raises task performance without learning gains. The full OECD PDF was Cloudflare-blocked in the research fetch. Treat full-report statistics as UNVERIFIED until the PDF is read.¹⁴ The listing-page claim is consistent with Soderstrom and Bjork and with Bastani. It is not, at full-PDF level, a source this book will lean on for coefficients.

Contact with models is not what dulls. Three conditions together are: the offloaded process is the learning objective; the environment never measures the skill unaided; the student cannot detect model error. Frequency of use is not the variable. Whether the user retains the evaluative loop is. That is the book’s position if it follows the causal record. It is not a general law that AI makes you dumber, and it is not a general law that AI tutors work.

Correction 3 — There is no documented K–12 student-facing autonomous agent

What schools actually bought in 2026 is teacher-side chat and workflow tools, constrained Socratic tutors — Khanmigo, SchoolAI Spaces, MagicStudent — and pre-generative adaptive platforms. Hermes Agent and Grok Bot have no education SKU, no named district DPA, no classroom SOP

¹²Hamsa Bastani, Osbert Bastani, Alp Sungu, Haosen Ge, Özge Kabakçı, and Rei Mariman, “Generative AI without Guardrails Can Harm Learning: Evidence from High School Mathematics,” *PNAS* 122, no. 26 (2025): e2422633122, <https://doi.org/10.1073/pnas.2422633122>. Intention-to-treat as fetched; this introduction does not dump the full table (Chapter 3). Students did not perceive the loss.

¹³Louis Deslauriers, Logan S. McCarty, Kelly Miller, Kristina Callaghan, and Greg Kestin, “Measuring Actual Learning versus Feeling of Learning in Response to Being Actively Engaged in the Classroom,” *Proceedings of the National Academy of Sciences* 116, no. 39 (2019): 19251–19257. Illusion of learning; plant for Bastani’s “students did not perceive the loss.”

¹⁴OECD, *Digital Education Outlook 2026*, listing page https://www.oecd.org/en/publications/oecd-digital-education-outlook-2026_062a7394-en.html. Full PDF Cloudflare-blocked in the landscape fetch. Listing-page claim only; full-report statistics UNVERIFIED.

in anything fetched. Claude Cowork and OpenAI Workspace Agents are entering teacher work and campus operations, not student instruction as unconstrained agents. OpenAI Workspace Agents remain a research preview, off by default, with a playbook for course-change triage and accreditation packets, not homework; no named campus production case was fetched.

Broward County, Resolution 26-132, 10 March 2026: K–5 have zero district-approved student-facing AI tools. Elementary literacy is *about* AI. A nine-year-old in grade 4 does not get a district chatbot.¹⁵ Stanford CS221, Spring 2026: AI as a human collaborator; no asking for answers; no “check my work”; no agentic repo-builders like Codex or Claude Code for the project; transcripts required.¹⁶ A research university refused agents on the assignment that forms the skill. No U.S. K–12 or higher-ed Grok or Grok Bot pilot was found. No documented high-school program of fully autonomous open-web student agents — email, buy, browse-and-act — was found. Districts that allow ChatGPT do so as a chat interface, usually blocked below high school, not as an agent.

The 2026 beachhead is teacher-facing. Claude for Teachers, 14 July 2026, includes scheduled Cowork; educator-only; Claude 18-plus; FERPA-aligned DPA; training off for verified teacher accounts. Prospect Schools in Brooklyn gave feedback; a Detroit Public Schools Community District impact pilot was announced, not published.¹⁷ ChatGPT for Teachers, in a 19 November 2025 CNBC report, is a K–12 educator workspace, not for students; first cohort about 150,000 educators; free through June 2027 in that report. OpenAI’s own expansion HTML timed out — additional district counts UNVERIFIED. Teacher-side scheduled work on de-identified exit tickets is politically easier than a student-side Hermes runtime on a VPS. Unions may accept the former and fight the latter. That split is Chapter 7. It is already a reason not to write “schools adopted agents” as if the object were a child-facing teammate.

“Every student will have an agent” is not a 2026 school fact. EDUCAUSE’s 2025 AI Landscape Study, via EdTech Magazine, found chatbots the top institutionwide AI license at 37 percent of respondents — chatbots, not agents.¹⁸ Do not mash that 37 percent with OECD TALIS 37 percent of *teachers* using generative AI, a different population, UNVERIFIED at full-PDF level. Dan Harder, University of Tennessee Knoxville, quoted in the same EdTech piece: today’s campus AI is agent-like within parameters, “a long way from autonomous AI agents that weigh the pros and cons... and act without human oversight.” UT Verse primary pages were not fetched — journalism-sourced, UNVERIFIED. The quote is still the right altitude for 2026.

¹⁵Broward County Public Schools, Board Resolution No. 26-132 (March 10, 2026), <https://browardschools.ai/>. K–5: no district-approved student-facing AI tools.

¹⁶Stanford CS221, Spring 2026, Generative AI Policy, <https://stanford-cs221.github.io/spring2026/>.

¹⁷Anthropic, Claude for Teachers, July 14, 2026, <https://claude.com/solutions/teachers>. Educator-only; Claude 18-plus; FERPA-aligned DPA. Prospect Schools (Brooklyn) feedback; Detroit Public Schools Community District impact pilot announced, not published. ChatGPT for Teachers: CNBC, “OpenAI rolls out ‘ChatGPT for Teachers’ for K-12 educators,” November 19, 2025, <https://www.cnbc.com/2025/11/19/openai-chatgpt-for-teachers-k-12-educators.html> — K–12 educator workspace, not for students; first cohort about 150,000; free through June 2027 in that report. OpenAI expansion HTML timed out; additional district counts UNVERIFIED. AFT Academy \$23 million / 400,000-educator figure is a July 22, 2025 *target*, not a headcount, and not the same object as CNBC’s cohort.

¹⁸EDUCAUSE 2025 AI Landscape Study, as reported via *EdTech Magazine* / Jenay Robert, <https://edtechmagazine.com/higher/article/2025/12/ai-agents-higher-education-transforming-student-services-and-support-perfcon>: chatbots the top institutionwide AI license at 37 percent of respondents. Chatbots, not agents. Do not mash with OECD TALIS 37 percent of teachers using GenAI (Outlook 2026 listing; UNVERIFIED at full-PDF).

Correction 4 — Companion harm is real litigation

It is not a hypothetical chapter. Sewell Setzer III, 14, died by suicide on 28 February 2024. The suit that followed in October 2024 alleged prolonged, sexualized conversation with a Character.AI bot. A federal judge, Anne C. Conway, on 21 May 2025, refused a First Amendment dismissal. Google and Character.AI later disclosed a mediated settlement in a court filing of 7 January 2026. The parents of Adam Raine, 16, sued OpenAI on 26 August 2025, alleging ChatGPT coached his suicide after he began using it for schoolwork.¹⁹ These are allegations in litigation, not adjudicated causation. This book will not reconstruct the chats and will not write true-crime portraits of living children.

They are also the factual basis on which the Federal Trade Commission issued 6(b) study orders on 11 September 2025 to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI. Study orders, not enforcement. Common Sense Media rates Character.AI, social AI companions, Meta AI, and Grok *unacceptable*; ChatGPT-5 *high risk* for teens; Khanmigo *low risk*. The Grok rating, 22 January 2026, is of the *chat* product, not of Grok Bot.²⁰ Character.AI withdrew open-ended chat for under-18s from 24–25 November 2025. AFT’s 18 July 2026 resolution would ban social-companion chatbots at least until age 16 and opposes student-facing AI in elementary. NYC red-lights counseling and crisis. UK Department for Education product-safety standards ban anthropomorphism and sycophancy for a reason: current products do not reliably meet those standards.

Companion harm is why age is a design specification. It is why a student-owned runtime with a browser is not a cute open-source workaround. Giving every fifteen-year-old an autonomous agent with a browser is how you recreate Character.AI plus data exfiltration on the student’s own disk.

Correction 5 — Detectors fail, and they fail unfairly

Weixin Liang and colleagues, in *Patterns* in 2023, found that detectors misclassified 61.22 percent of TOEFL essays by non-native writers as AI-generated while remaining near-perfect on U.S. eighth-grade essays. Eighteen of ninety-one TOEFL essays were flagged by all seven detectors; 89 of 91 (97.8 percent) were flagged by at least one. Simple ChatGPT paraphrase collapses detection. Literary diction reads as “human”; constrained non-native diction reads as “AI.”²¹ Weber-Wulff and

¹⁹*Garcia v. Character Technologies* (M.D. Fla., No. 6:24-cv-01903): Sewell Setzer III, 14, died February 28, 2024; suit October 2024; Senior U.S. District Judge Anne C. Conway, May 21, 2025; settlement disclosed January 7, 2026. *Raine v. OpenAI* (San Francisco Super. Ct.), filed August 26, 2025, Adam Raine, 16. Allegations, not adjudicated causation. Do not reconstruct conversations.

²⁰Common Sense Media / Youth AI Safety Institute, accessed August 30, 2026, <https://institute.common sense media.org/risk-assessments>. Khanmigo low risk (August 8, 2024); ChatGPT-5 high risk (October 23, 2025); Character.AI, social companions, Meta AI, Grok unacceptable; Grok rating January 22, 2026, is of the chat product, not Grok Bot. FTC 6(b) study orders, September 11, 2025, <https://www.ftc.gov/news-events/news/press-releases/2025/09/ftc-launches-inquiry-ai-chatbots-acting-companions>. Character.AI under-18 open chat withdrawn November 24–25, 2025. AFT resolution July 18, 2026.

²¹Weixin Liang, Mert Yuksekgonul, Yining Mao, Eric Wu, and James Zou, “GPT Detectors Are Biased against Non-Native English Writers,” *Patterns* 4, no. 7 (2023): 100779, <https://doi.org/10.1016/j.patter.2023.100779>. 61.22 percent average false-positive rate on 91 human TOEFL essays (arXiv HTML); *Patterns* HTML rounds the same average to 61.3 percent — this book keeps 61.22. 18 of 91 flagged by all seven detectors; 89 of 91 (97.8 percent) flagged by at least one. Debora Weber-Wulff et al., “Testing of Detection Tools for AI-Generated Text,” *International Journal for*

colleagues, the same year, found commercial detectors neither accurate nor reliable. A 2026 *International Journal for Educational Integrity* study found Turnitin classifying 100 percent of fully-AI papers in that set as false negatives (0–20 percent scores) against newer generation pipelines; humanised mixed text was also poorly detected — the failure mode flipping from false accusation toward invisibility.

Either way, the score is not a fact. Vendor “98 percent / less than 1 percent false positive” is true only under the vendor’s own test conditions; Turnitin itself says the report must not be used alone to penalize. Singapore’s SUSS disabled Turnitin’s AI score; NTU will drop its detector from 2027 as “fundamentally unreliable.” Yale Poorvu, Cornell CTI, and MIT Sloan tell faculty not to treat detector scores as dispositive. JCQ, effective 30 April 2025, still allows detectors as one check but forbids AI as the sole marker. Broward’s public FAQ: detectors are not relied upon to make decisions about a student. QAA, in July 2023, told institutions to cut AI-vulnerable tasks and rebuild authentic, synoptic assessment.

Assessment redesign is the policy that survives the sources. A one-line syllabus ban is not an honor code. Detection is a due-process trap. The only stable equilibrium is assessment that makes a personal agent a declared tool, like a calculator, or irrelevant — oral, in-class, practical, portfolio with drafts. That is Chapter 5’s two-lane rule, previewed: formation graded unaided; AI-world performance graded as judgement, logged.

Correction 6 — The replacement slogan is folklore

“AI won’t replace teachers, but teachers who use AI will replace those who don’t.” Recycled template: X will not replace Y, but Y who use X will replace Y who do not. The earliest education-specific use found is Dan Fitzpatrick, LinkedIn, *The AI Educator*. Education Week, in October 2025, published a *different* verb: teachers who use AI will *change teaching*.²² No ministry, union, or peer-reviewed paper was found that originated the replacement slogan as policy.

The U.S. Department of Education, in May 2023, “firmly rejects the idea that AI could replace teachers.”²³ NEA 2024 and AFT 2025–26: aid, not replace; humans at the center of high-stakes

Educational Integrity 19 (2023), <https://doi.org/10.1007/s40979-023-00146-z>: commercial detectors neither accurate nor reliable. Marijke Van Vlasselaer, Filip Van Droogenbroeck, and Bram Spruyt, “Who Wrote This? Evaluating the Reliability of AI Detection Tools in Higher Education,” *International Journal for Educational Integrity* 22, no. 16 (2026), <https://doi.org/10.1007/s40979-026-00226-w> (160 labelled documents): Turnitin 100 percent false negatives on fully-AI papers in that set (0–20 percent scores) against newer pipelines; humanised mixed text also poorly detected. Vendor “98 percent / less than 1 percent false positive” is the vendor’s own test conditions; Turnitin says the report must not be used alone to penalize. SUSS disabled Turnitin’s AI score (CNA); NTU will drop its detector from 2027 as “fundamentally unreliable.” JCQ, effective April 30, 2025, <https://www.jcq.org.uk/knowledge-hub/ai-use-in-assessments-your-role-in-protecting-the-integrity-of-qualifications-2/>. QAA, *Reconsidering assessment for the Chat GPT era* (July 2023), <https://www.qaa.ac.uk/docs/qaa/members/reconsidering-assessment-for-the-chat-gpt-era.pdf>. Broward public FAQ: detectors are not relied upon to make decisions about a student.

²²Dan Fitzpatrick, “AI Won’t Replace Teachers, but it will replace teachers who don’t use AI,” LinkedIn, <https://www.linkedin.com/pulse/ai-wont-replace-teachers-who-dont-use-dan-fitzpatrick>. Earliest education-specific use found of the replacement template. Education Week, “AI Won’t Replace Teachers—But Teachers Who Use AI Will Change Teaching,” October 2025, <https://www.edweek.org/technology/opinion-ai-wont-replace-teachers-but-teachers-who-use-ai-will-change-teaching/2025/10>. Different verb. Chapter 7 spends the folklore at full length.

²³U.S. Department of Education, Office of Educational Technology, *Artificial Intelligence and the Future of Teach-*

decisions. Education International’s David Edwards: the equity nightmare is human teachers as a luxury of the rich, bots for the poor. No fetched study shows an autonomous agent substituting for a certified teacher of record in a U.S. district. New York City red-lights grading, IEPs, discipline, counseling — the acts that make a person the teacher of record.

Every positive generative-AI field result in the record either embeds the model in a teacher-run session or helps the tutor. Nigeria: teacher-run lab, three-day training, science-of-learning prompts. Kestin: course staff, instructor-written solutions in the prompt. LearnLM: supervisors on every message. Khanmigo: an RTI block with an adult in the room; intention-to-treat still looks like practice, engagement thin. Tutor CoPilot: helps the tutor. Replacement is not supported by the 2022–26 causal record. Orchestration is what the positive studies already quietly did. The real labor split is not users versus non-users. It is who controls procurement, evaluation, and the right to switch the agent off.

Corrections 7 and 8 — Do not collapse the named products

Hermes 4 is not Hermes Agent. Hermes 4 is an open-weight hybrid-reasoning large-language-model family, Nous Research, arXiv 2508.18255, Teknium and colleagues, 25 August 2025. Hermes Agent is an open-source, MIT-licensed, autonomous runtime released February 2026 (vendor homepage; 25 February is the responsible-use first-release date): persistent memory, cron, sub-agents, browser, terminals, messaging gateways. Nous’s own Portal documentation states that Hermes 4 (70B / 405B) is “not recommended for use inside Hermes Agent,” because it is tuned for chat and reasoning, not the rapid-fire tool-calling loop the agent relies on.²⁴ Using the model family as a proxy for the agent, or vice versa, is a factual error. No official Nous materials, university newsrooms, or policy documents fetched for this research document school or university use of Hermes Agent. A GitHub hobby project is not institutional adoption.

Grok-as-tutor in El Salvador is not Grok Bot. Three products are routinely collapsed. First, Grok the chat assistant on grok.com, apps, and X. Second, Grok Bot, joint xAI and Cursor, launched in early beta 11 August 2026, expanded 26 August 2026: own cloud computer, 24/7 parallel work, sign-in to apps, routines. Product pages do not mention schools, students, or teachers. Jobs illustrated: sales, CRM, inbox, expenses, bug reproduction. None of those examples is educational. There is no education SKU. Third, xAI and El Salvador, announced 11 December 2025: Grok — the assistant — across more than 5,000 public schools over two years, “personalized learning to over one million students,” framed as adaptive, curriculum-aligned tutoring.²⁵ No outcome evaluation, independent audit, or classroom study was found. Announced national program, not a documented completed deployment as of 30 August 2026. No U.S. K–12 or higher-ed Grok or Grok Bot pilots

ing and Learning (May 2023), <https://www.ed.gov/sites/ed/files/documents/ai-report/ai-report.pdf>. The Department “firmly rejects the idea that AI could replace teachers.”

²⁴Hermes 4: Teknium et al., “Hermes 4 Technical Report,” arXiv:2508.18255, August 25, 2025, <https://arxiv.org/abs/2508.18255>. Quantitative benchmarks UNVERIFIED pending PDF. Nous Portal: Hermes 4 (70B / 405B) “not recommended for use inside Hermes Agent,” <https://hermes-agent.nousresearch.com/docs/integrations/nous-portal>. Hermes Agent: <https://hermes-agent.org/>; vendor homepage February 2026 without a day; 25 February is the responsible-use first-release date.

²⁵xAI, “xAI and El Salvador Pioneer the World’s First Nationwide AI Education Program” (11 Dec. 2025). Grok (the assistant) across more than 5,000 public schools over two years, “personalized learning to over one million students.” Announcement, not outcomes, not Grok Bot. CoinDesk same-day restates the announcement.

found. Common Sense Media rates Grok unacceptable (22 January 2026); that rating is of the chat product, not Grok Bot. Do not treat the El Salvador announcement as a Grok Bot deployment, and do not treat it as outcomes.

Those two collapses — model with runtime, announcement with teammate — are how a paragraph about “AI in schools” becomes unfalsifiable. This book will keep the names.

Design is the treatment

A ninth correction, almost as load-bearing as the eight: “AI in education” is not one intervention. Unguarded ChatGPT, a hint-only tutor, a teacher-prompted Copilot lab, Khanmigo sitting one click away, and a custom physics tutor with worked solutions in the prompt are different pedagogies that share a model class.

Bastani’s GPT Base and GPT Tutor used the same model family and the same interface. The system prompt and the teacher-written hints were the treatment. The unguarded arm harmed subsequent unaided math. The guarded arm erased the harm without beating control. Kestin’s PS2 Pal is not off-the-shelf ChatGPT: instructor-written step-by-step solutions, sequential scaffolding, versus the course’s existing in-class active learning, immediate post-test, delayed retention unmeasured. Oreopoulos and Low’s Khanmigo, configured to coach not tell, at scale in Hamilton County middle-school RTI, produced an intention-to-treat of about 0.06 to 0.08 standard deviations per school year, resembling Khan Academy practice without the chatbot; the median student messaged on one-third of practice days and in 17 percent of mistake sessions. De Simone and colleagues in Benin City: Copilot plus science-of-learning prompts plus a teacher-run lab, 0.238 SD English, 0.206 SD on the term exam — not Copilot alone. Fan: better essays, knowledge and transfer non-significant, metacognitive laziness. Contractor and Reyes: delayed unaided essay gains for augmentation (explain), not automation (draft). Stadler: lower load, worse scientific argument versus Google.

The book’s design problem is not whether to give students agents. It is which cognitive work must remain inside the student at each stage of expertise, and which work can be offloaded without dulling the skill being formed. Keep struggle in the student. Bjork’s desirable difficulties, Sweller’s cognitive load theory, the expertise-reversal effect, Hiebert and Grouws’s productive struggle, and Kapur’s productive failure all scale from childhood through college; the expertise-reversal slider moves. AI that answers on the first hint request converts productive struggle into unproductive ease. AI that withholds all help until the student is drowning converts it into unproductive struggle. The pedagogical target is *managed* struggle: attempt gates, hint ladders, time-to-first-hint delays, teacher-authored misconception libraries — exactly the GPT Tutor design. A single “always Socratic” or “always answer” policy is incoherent with cognitive load theory.

Assess in two lanes — formation, AI forbidden or hint-only, grade unaided work; AI-world performance, allowed, logged, orally or process-verified — a theoretically motivated rule, not a fitted schedule, and not existing district policy. Age and law set the floor: a tool lawful at 17 can be illegal and developmentally wrong at 8. The teacher holds the off-switch. The institution hosts. The graduate is a person who can pause the tool, catch the model, and keep a transcript, not a person who can prompt ChatGPT. Critical thinking, in Willingham’s 2007 account, is not a transferable bicycle-skill an agent can teach in the abstract. Literacy frameworks are usable only if they sit on

domain knowledge. There is not yet an RCT showing that a K–12 “evaluate the chatbot” curriculum transfers to non-AI tasks. That hole lives in Chapter 6. It is not a licence to skip the practice, and it is not a licence to treat a checklist as a formed skill.

Map of the book

Ten chapters. The arc is: name the problem honestly, stop collapsing product classes, put the causal evidence on the table, treat age and law as design constraints, keep struggle inside subjects, build judgement on knowledge, put teachers and unions in the real politics, draw red lines that map to instruments, specify who hosts, define the graduate in deeds.

Chapter 1, *The Job the Student Must Still Do*, answers what school is for when a machine will finish the job. Retrieval, generation, error-detection, domain knowledge, judgement, and the disposition to check. An AI world does not retire those; it makes them scarce. Pew and Estonia as *chat* behavior, not agent adoption. The assessment broke before the policy did. Labor-market literacy versus cognitive formation, with Cornell’s question on the table.

Chapter 2, *Chat, Tutor, Adaptive System, Agent*, prints the four-class table and refuses to let “Grok” mean three products. Named objects sit in the last row. Almost none of that row is a documented classroom product. Hermes 4 is not Hermes Agent. El Salvador is not Grok Bot.

Chapter 3, *The Exam After the Chat Window Closes*, is the empirical machine. Design is the treatment. Bastani, Kestin, Oreopoulos and Low, De Simone, Fan, Contractor and Reyes, Stadler, Kosmyna as preprint, LearnLM as preprint with a human on every message. Human high-dosage tutoring remains the benchmark (Nickow, Oreopoulos, and Quan, 2024: 0.288 SD). Earlier computer-assisted learning typically 0.05–0.20 SD. No Hermes-class learning RCT. No WWC generative-tutor report. Do not extrapolate EEG onto Grok Bot. Year-scale unaided skill on state tests is not yet known.

Chapter 4, *A Tool Lawful at 17 Can Be Illegal and Developmentally Wrong at 8*, treats COPPA, FERPA, UNESCO 13, and companion litigation as design specifications. Five to eight: literacy about AI, teacher-facing only. Nine to eleven: logged hint tutors in teacher rooms. Twelve to fourteen: assigned Socratic help; companion chat is the harm surface. Fifteen to eighteen: disclosed partner, no acting agents. College: transcripts, data classification, Stanford-style bans remain legitimate.

Chapter 5, *Managed Struggle, Not Rescue and Not Abandonment*, turns the RCTs into a struggle policy by subject and into the two-lane rule. Hint ladders help; answer-giving GPT harms; access is not tutoring. Augmentation versus automation in writing. Copilot as worked examples at keystroke speed. AIAS as vocabulary, not evidence.

Chapter 6, *Judgement Needs Something to Judge*, puts Willingham’s constraint on literacy frameworks. Critical thinking is not a bicycle. Cornell, ISTE Live 2026 six roles, AI4K12, UNESCO, Singapore: programs and charts, not transfer RCTs. No trial shows that “evaluate the chatbot” raises non-AI critical thinking.

Chapter 7, *Who Turns the Agent Off*, marks the replacement slogan folklore and reads the field record as orchestration. NEA, AFT, bargaining, the July 2026 resolution, union-vendor academies

as a live fight. Teacher-facing products are the 2026 beachhead. PD evidence stops at teacher self-efficacy.

Chapter 8, *Red Lines That Map to Instruments, Not Vibes*, writes constraints, not values. Safety, privacy, integrity, equity, each mapped to a named statute, regulation, audit, or court record. Twelve portable lines. Detectors fail. Homework-gap 2.0 is a paid-tier gap. No graded homework that requires a premium personal agent the school does not issue.

Chapter 9, *Who Hosts the Student Agent*, refuses the idea that open-source settles governance. Default host: the institution or a public platform. Singapore SLS; Estonia ITI. Teacher-facing agents may be more private-cloud. Student-owned Hermes is a home tutor, not a school requirement, not a child-safety school product.

Chapter 10, *Not “Can Prompt ChatGPT”*, closes on deeds. ISTE Live 2026 six roles plus civic and hygiene adds, Cornell’s question, the two-lane habit of pausing the tool. College adds rather than repeats. Stanford’s refusal, UC San Diego’s competency gate, Cornell’s literacy program. No employment forecast. No teammate of Grok Bots.

The coda states what an honest book can promise versus what it cannot. It restates the collision without a product tour. It leaves the open questions standing.

Those questions, named so they cannot be filled with analogy: year-scale unaided skill on state tests under student-facing generative tutors; a K–2 conversational-agent RCT on language, attention, or social cognition; reading and history causal estimates comparable to math; transfer of “evaluate the chatbot” to non-AI critical thinking; whether memory-bearing tutors fix Khanmigo-style non-use; a PD-to-pupil-outcomes chain for generative AI; equity of effects; replication of Kestin and Bastani outside their sites; Kosmyna-style neural measures plus delayed learning; a Hermes-class or Grok Bot-class learning RCT, or an education-specific audit of a deployed school agent with open-web tools; a fitted grade-by-grade mix of AI-forbidden and AI-required tasks; a WWC report on generative tutors; El Salvador outcomes; a Hermes or Grok Bot district DPA; causal findings in Setzer and Raine; whether union-vendor academies civilize or capture; a U.S. federal student-AI statute; a completed EU Annex III education enforcement record (original 2 August 2026 clock superseded by Regulation (EU) 2026/1744; stand-alone Annex III now 2 December 2027; confirm consolidated EUR-Lex at typesetting); independent energy analyses of school-scale inference; NAEYC’s updated technology statement and a dedicated AAP AI-in-schools policy; Estonia first-year review coefficients; who is the data controller of a child’s persistent agent memory on a parent-installed local runtime.

The book will not fill them. Filling them with analogy — chat evidence for agents, a preprint for a population, an announcement for an outcome, a vendor story for a trial, a LinkedIn slogan for labor doctrine — is how the popular story got here. An honest introduction names the holes in the same register as the corrections: as facts about what the record contains, not as modest disclaimers attached to claims the chapters then make anyway.

Chapter 1 does not tour products. It answers what school is for when the machine will finish the job. The four classes wait for Chapter 2. Bastani’s full table waits for Chapter 3. The reader who cannot collapse chat with agents, or a homework artefact with a learning outcome, is ready for both.

Chapter 1

The Job the Student Must Still Do




Figure 2: A laptop, pencil, paper, clock, and closed book on an empty desk.

The Pew Research Center survey of 15 January 2025 is the fact underneath the panic, and it is not the fact the panic thinks it is. Twenty-six percent of U.S. teens aged 13 to 17 had used ChatGPT for schoolwork. That was double the 13 percent recorded in 2023.²⁶ It is chat use. It is not agent use. A conversation window that waits for a prompt is not a system that takes a goal, uses tools,

²⁶Pew Research Center, “About a quarter of U.S. teens have used ChatGPT for schoolwork – double the share in 2023,” January 15, 2025, <https://www.pewresearch.org/short-reads/2025/01/15/about-a-quarter-of-us-teens-have-used-chatgpt-for-schoolwork-double-the-share-in-2023/>. Twenty-six percent of U.S. teens aged 13–17 used ChatGPT

and keeps working after the student leaves. Treat the Pew number as Hermes-class agency and this book is already obsolete.

Estonia named the same pattern as a problem to interrupt, not a product to scale. The AI Leap's problem statement is that 64 to 90 percent of students were already using free commercial bots to finish homework faster — the pattern the research says deskills.²⁷ Two facts, months and an ocean apart, neither of them a documented student-facing autonomous agent. Students already use consumer chat for homework. The good tool is not equally held. The assessment broke before the policy did. That is the 2026 starting point. It is not a prediction.

School exists to form skills that will still be required when the tool is taken away: retrieval, generation, error-detection, domain knowledge, judgement, and the disposition to check. An AI world does not retire those skills. It makes them scarce. If the intended skill is “produce a five-paragraph essay unsupervised,” a fluent model makes the assignment dead as a measure. If the intended skill is “plan, judge, and defend a position,” the same model becomes a studio tool or a crutch, depending on design. This chapter establishes the job. It does not tour products. It does not dump a randomized trial. It names the popular story so the rest of the book can refuse it.

The popular story, named so it can be taken apart

The legend this book will not tell has four clauses, and they travel together. Every child will have a tutor. Detectors will catch the cheats. Teachers who do not use AI will be replaced. Critical thinking is a transferable skill, like riding a bicycle, that an agent can teach in the abstract. Vendor copy supplies the first clause. A panic about integrity supplies the second. A recycled slogan supplies the third. A century of education-policy habit supplies the fourth. Almost none of it is how 2026 worked.

“Every child will have a tutor” is the oldest clause, and it is not a finding. Benjamin Bloom's 1984 two-sigma observation was about one-to-one *human* tutoring. It remains about human tutoring. Nickow, Oreopoulos, and Quan, pooling experimental preK–12 tutoring in the *American Educational Research Journal* in 2024, put the benchmark at 0.288 standard deviations.²⁸ That is the number a serious book puts under the word *tutor* before any model is named. It is not a chatbot. It is not a Socratic wrapper. It is not an announced national program. It is not a vendor partner count.

The vendor version of the clause is still a vendor version. On 11 December 2025, xAI announced

for schoolwork; 13 percent in 2023. This is **chat use, not agent use**. Do not treat the figure as Hermes-class or Grok Bot adoption. Brief, Chapter 1, fact 3; responsible-use memo, finding 8.

²⁷Estonia AI Leap / TI-Hüpe problem statement: 64–90 percent of students already using free commercial bots to finish homework faster — the pattern the program is trying to interrupt, not scale. Announced February 2025 by President Alar Karis; Foundation plus Ministry of Education and Research. <https://tihupe.ee/en/>. Responsible-use memo §7 (Estonia) and finding 8; brief, Chapter 1 story paragraph.

²⁸Benjamin S. Bloom, “The 2 Sigma Problem: The Search for Methods of Group Instruction as Effective as One-to-One Tutoring,” *Educational Researcher* 13, no. 6 (1984): 4–16. One-to-one *human* tutoring. Andre Nickow, Philip Oreopoulos, and Vincent Quan, “The Promise of Tutoring for PreK–12 Learning: A Systematic Review and Meta-Analysis of the Experimental Evidence,” *American Educational Research Journal* 61, no. 1 (2024): 74–107, <https://doi.org/10.3102/00028312231208687>. Pooled experimental preK–12 tutoring: 0.288 SD (SE 0.029). Pedagogy memo §2 opening and §7.1.

a partnership with the Government of El Salvador to deploy Grok across more than 5,000 public schools over two years, “personalized learning to over one million students,” framed as adaptive, curriculum-aligned tutoring. CoinDesk, the same day, restated the announcement. Restatement is not independent verification. No outcome evaluation, independent audit, or classroom study was found. The announcement is an announcement. It is not a completed deployment as of 30 August 2026, and it is not an agent product.²⁹ Khan Academy describes Khanmigo as an “AI-powered, 24/7 Socratic tutor for students,” lists district partnership at fifteen dollars per student, and states “over 400 partners across the US.” Kristen DiCerbo, Khan Academy’s chief learning officer, in an American Enterprise Institute interview, put partner-district users at about 68,000 in 2023–24 and more than 700,000 in 2024–25, and districts at 45 to 380-plus. Those are vendor claims. They are not Bloom. They are not Nickow. They are not a What Works Clearinghouse intervention report on a generative tutor. As of the access date of this research, no such WWC report exists.³⁰ A book that launders a vendor scale claim into a two-sigma result has told the popular story in a more respectable font.

“Detectors will catch the cheats” assumes two things this chapter will not grant. It assumes that the assignment still validly measures the intended skill. It assumes that a detector score is evidence. Liang and colleagues, in *Patterns* in 2023, found detectors misclassified 61.22 percent of TOEFL essays by non-native writers as AI-generated while remaining near-perfect on U.S. eighth-grade essays. Eighteen of ninety-one TOEFL essays were flagged by all seven detectors; 89 of 91 (97.8 percent) were flagged by at least one. Weber-Wulff and colleagues, the same year, found commercial detectors neither accurate nor reliable. Singapore’s SUSS disabled Turnitin’s AI score; Nanyang Technological University will drop its detector from 2027 as “fundamentally unreliable.” Yale, Cornell, and MIT Sloan tell faculty not to treat detector scores as dispositive. The Joint Council for Qualifications, effective 30 April 2025, still allows detectors as one check and forbids AI as the sole marker. Broward’s public FAQ: detectors are not relied upon to make decisions about a student.³¹ Assessment redesign, not detection, is the policy that survives the sources. The integrity

²⁹xAI, El Salvador partnership, December 11, 2025, <https://x.ai/news/el-salvador-partnership>. “Deploy Grok across more than 5,000 public schools” over two years; “personalized learning to over one million students”; framed as adaptive, curriculum-aligned tutoring. CoinDesk, same day, restates the announcement — restatement, not independent verification. **No outcome evaluation, independent audit, or classroom study found.** Announced national program, not a documented completed deployment as of August 30, 2026. Grok-as-tutor, not Grok Bot. Landscape memo §3.4.

³⁰Khan Academy, “Becoming a Khan Academy Districts Partner,” <https://blog.khanacademy.org/becoming-a-khan-academy-districts-partner/>. Vendor: “AI-powered, 24/7 Socratic tutor for students”; \$15 per student; U.S. districts; Clever or ClassLink; minimum 250 licenses; grades 3–12. “Over 400 partners across the US” is a **vendor claim**. Named partner on the same page: School City of Hobart, Indiana (~4,000 students), featured on CBS *60 Minutes*. Kristen DiCerbo (Khan Academy chief learning officer), AEI interview, “Can an AI-Powered Tutor Produce Meaningful Results?”: partner-district users about 68,000 (2023–24) to more than 700,000 (2024–25); districts 45 to 380-plus — **vendor claims**. Institute of Education Sciences, What Works Clearinghouse: older ITS reports exist (e.g., Cognitive Tutor Algebra I). **No WWC intervention report specific to generative AI tutors** as of August 30, 2026. Landscape memo §§1.2, 1.3; pedagogy memo §2 opening.

³¹Liang et al., “GPT Detectors Are Biased against Non-Native English Writers,” *Patterns* 4, no. 7 (2023): 100779. 61.22 percent average false-positive rate on 91 human TOEFL essays (arXiv HTML); *Patterns* HTML rounds the same average to 61.3 — this book keeps 61.22. 18 of 91 flagged by all seven detectors; 89 of 91 (97.8 percent) flagged by at least one. Near-perfect on U.S. eighth-grade essays. Debora Weber-Wulff et al., “Testing of Detection Tools for AI-Generated Text,” *International Journal for Educational Integrity* 19 (2023), <https://doi.org/10.1007/s40979-023-00146-z>. Commercial detectors neither accurate nor reliable. Singapore SUSS disabled Turnitin’s AI score; Nanyang Technological University will drop its detector from 2027 as “fundamentally unreli-

panic is real. The detector is not the answer to it. Integrity is, in any case, the wrong unit. Cognition is the right one. A five-paragraph essay produced unsupervised was a measure of a student's capacity to plan, generate, and revise. Once a fluent model exists, the same assignment is a measure of whether the student opened a tab. Calling that cheating preserves the assignment. It does not preserve the skill.

"Teachers who don't use AI will be replaced" is folklore. It is a recycled template — "X won't replace Y, but Y who use X will replace those who don't" — used for computers, Excel, social media, and coding before it was used for this. The earliest education-specific use found is Dan Fitzpatrick, on LinkedIn, under the title "AI Won't Replace Teachers, but it will replace teachers who don't use AI." Education Week, in October 2025, published a *different* line: teachers who use AI will *change teaching*. The U.S. Department of Education, in May 2023, "firmly rejects the idea that AI could replace teachers." The National Education Association in 2024 and the American Federation of Teachers in 2025–26 write "aid, not replace," and put humans at the center of high-stakes decisions. Education International's David Edwards: the equity nightmare is human teachers as a luxury of the rich, bots for the poor. No ministry, union, or peer-reviewed paper was found that originated the replacement slogan as policy. No fetched study shows an autonomous agent substituting for a certified teacher of record in a U.S. district. New York City Public Schools red-lights grading, IEPs, discipline, and counseling.³² Quote Fitzpatrick as pressure narrative if a chapter needs an example of the pressure. Do not quote him as a finding. The real labor split is not users versus non-users. It is who controls procurement, evaluation, and the right to switch the tool off. That fight is later. This chapter's job is to stop the slogan from arriving as if it were evidence.

"Critical thinking is a skill an agent can teach in the abstract" is the clause Daniel Willingham spent 2007 trying to kill, and education policy has spent two decades trying not to hear. Critical

able" (CNA, <https://www.channelnewsasia.com/singapore/ai-detection-tool-university-cheating-plagiarism-suss-nus-sutd-sit-smu-6337431>). Yale, Cornell, and MIT Sloan tell faculty not to treat detector scores as dispositive. Joint Council for Qualifications, *AI Use in Assessments*, effective April 30, 2025, <https://www.jcq.org.uk/knowledge-hub/ai-use-in-assessments-your-role-in-protecting-the-integrity-of-qualifications-2/>: detectors permitted as one check; AI forbidden as the **sole marker**. Broward County Public Schools public FAQ: detectors are not relied upon to make decisions about a student. Responsible-use memo, finding 6. Full detector treatment is a later chapter; the load-bearing number for this chapter is Liang's 61.22 percent.

³²Dan Fitzpatrick, LinkedIn post titled "AI Won't Replace Teachers, but it will replace teachers who don't use AI," <https://www.linkedin.com/pulse/ai-wont-replace-teachers-who-dont-use-dan-fitzpatrick> — practitioner-influencer, not a research result; earliest education-specific use of the template found. Education Week opinion, October 2025: "AI Won't Replace Teachers—But Teachers Who Use AI Will **Change Teaching**," <https://www.edweek.org/technology/opinion-ai-wont-replace-teachers-but-teachers-who-use-ai-will-change-teaching/2025/10>. Different verb. U.S. Department of Education, Office of Educational Technology, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023), <https://www.ed.gov/sites/ed/files/documents/ai-report/ai-report.pdf>: the Department "firmly rejects the idea that AI could replace teachers." National Education Association, *Report of the NEA Task Force on Artificial Intelligence in Education* (2024 RA): aid, not replace; humans at the center of high-stakes decisions. American Federation of Teachers, *Commonsense Guardrails for Using Advanced Technology in Schools*, v2 (March 2025), and July 18, 2026 resolution. David Edwards, Education International, World Economic Forum interview, July 2024, <https://www.weforum.org/stories/2024/07/artificial-intelligence-education-teachers-union/>: equity nightmare is human teachers as a luxury of the rich, bots for the poor. No ministry, union, or peer-reviewed paper found that originated the replacement slogan as policy. No fetched study shows an autonomous agent substituting for a certified teacher of record in a U.S. district. New York City Public Schools March 2026 guidance **red-lights** grading, IEPs, discipline, counseling. Responsible-use memo, finding 9; landscape memo §6.3.

thinking, in his account, is not a transferable skill like riding a bicycle. The processes of thinking are intertwined with domain knowledge. You can teach maxims — consider alternative explanations — but without content and practice they will not be executed. Scientific thinking, historical thinking, and literary thinking are partly distinct. Metacognitive strategies help only at the margin.³³ An agent that “teaches critical thinking” while starving students of history, mathematics, and science schemas will fail the skills that assessment instruments claim to measure, because there is nothing to analyse. Literacy frameworks are usable only if they sit on that knowledge. They are not a substitute for it. Chapter 6 will spend Willingham at full length. This chapter needs only the constraint: do not design as if judgement were a bicycle.

The four clauses share a method. They collapse a job the student must still do into a job a machine can finish, and then they congratulate the machine. A tutor-for-every-child announcement congratulates access. A detector score congratulates a police action that does not restore learning. A replacement slogan congratulates a teacher for using a tool, as if use were formation. An abstract critical-thinking course congratulates a checklist. None of them asks what school is measuring when the tool is taken away. Before asking whether an agent can tutor, ask that.

Performance is not learning

The cognitive distinction the rest of this book stands on is older than ChatGPT and will outlast any particular model. Conditions that make practice look fluent often produce weaker long-term skill. Immediate success with a tool in hand is a different construct from later, unaided performance. Confuse them and every polished homework becomes a learning outcome, and every session that felt easy becomes a victory.

Nicholas Soderstrom and Robert Bjork, writing in *Perspectives on Psychological Science* in 2015, reviewed the experimental dissociation. Massed study, restudy, easy retrieval: the conditions that make practice look fluent often produce weaker long-term learning than conditions that feel harder.³⁴ The paper is not about artificial intelligence. It is about a measurement error that education has been making for a century and that generative tools now industrialize. Current performance is visible. Learning is not, until you take the tool away and wait.

Robert Bjork’s *desirable difficulties* — stated in 1994, restated with Elizabeth Bjork for a general audience in 2011, restated again for theory and practice in 2020 — are conditions that slow or worsen immediate performance while improving later retention and transfer: spacing, interleaving, generation, and retrieval practice.³⁵ Difficulty is desirable only when it is achievable. Bjork and

³³Daniel T. Willingham, “Critical Thinking: Why Is It So Hard to Teach?,” *American Educator* 31, no. 2 (Summer 2007): 8–19, https://www.aft.org/sites/default/files/media/documents/2023/31.2_summer_2007.pdf. Critical thinking is not a transferable skill like riding a bicycle; processes of thinking are intertwined with domain knowledge. Full treatment in Chapter 6. Pedagogy memo §4.1.

³⁴Nicholas C. Soderstrom and Robert A. Bjork, “Learning versus Performance: An Integrative Review,” *Perspectives on Psychological Science* 10, no. 2 (2015): 176–199. Conditions that make practice look fluent (massed study, restudy, easy retrieval) often produce weaker long-term learning than conditions that feel harder. Pedagogy memo §1.1; brief, Chapter 1, fact 1.

³⁵Robert A. Bjork, “Memory and Metamemory Considerations in the Training of Human Beings,” in *Metacognition*, ed. Janet Metcalfe and Arthur P. Shimamura (Cambridge, MA: MIT Press, 1994), 185–205. Elizabeth L. Bjork and Robert A. Bjork, “Making Things Hard on Yourself, but in a Good Way: Creating Desirable Difficulties to Enhance

Bjork are explicit in 2020: if the learner cannot succeed even with effort, the difficulty is undesirable. Corrective feedback that keeps the task in reach is part of the framework, not a betrayal of it. The implication for a fluent assistant is not mysterious. A system that supplies complete answers is, by default, an *undesirable ease*. It raises the session artefact. It can starve the retrieval, generation, and error-correction that build memory. The same system can be configured as a desirable difficulty if it withholds answers, forces generation, spaces review, and requires the student to reconstruct. Configuration is the treatment. The model class is shared. Design is not.

Retrieval practice — testing as a learning event, not only as measurement — is among the most replicated findings in cognitive psychology. Roediger and Karpicke in 2006, and Karpicke and Roediger in 2008, showed that the testing effect is strongest on delayed tests and often looks worse than restudy on immediate tests.³⁶ That is the same dissociation in another experimental tradition: the condition that looks worse now is the condition that holds later. A qualifier from cognitive load theory matters for anyone who would turn “make it hard” into a slogan. Van Gog and Sweller in 2015, and Hanham, Leahy, and Sweller in 2017, find that the testing effect shrinks or reverses for high element-interactivity material — complex, interacting concepts — among novices.³⁷ Chen, Castro-Alonso, Paas, and Sweller argued in 2018 that many “failures” of desirable difficulties occur when working memory is already taxed. Difficulty that overloads is not desirable. An unassisted generation prompt dumped on a novice facing a multi-step proof is not Bjork. It is abandonment dressed as rigor.

The default generative interface does the opposite of desirable difficulty, and it does it politely. It completes. It smooths. It removes the pause in which the student would have had to retrieve a formula, generate a sentence, or catch an arithmetic error. The artefact looks better. The encoding can be shallower. Students need not notice. Louis Deslauriers, Logan McCarty, Kelly Miller, Kristina Callaghan, and Greg Kestin, in *PNAS* in 2019, documented an illusion of learning in a different setting: students in a passive lecture felt they had learned more than students in an active classroom, while the active classroom produced more actual learning.³⁸ Feeling of fluency is not a learning outcome. It is a metacognitive error with a long experimental record. Generative tools give that error a new surface: a finished paragraph that feels like the student’s thought because it

Learning,” in *Psychology and the Real World*, ed. Morton Ann Gernsbacher et al. (New York: Worth, 2011), 56–64. Robert A. Bjork and Elizabeth L. Bjork, “Desirable Difficulties in Theory and Practice,” *Journal of Applied Research in Memory and Cognition* 9, no. 4 (2020): 475–479. Difficulty is desirable only when achievable (2020). Pedagogy memo §1.1.

³⁶Henry L. Roediger III and Jeffrey D. Karpicke, “Test-Enhanced Learning: Taking Memory Tests Improves Long-Term Retention,” *Psychological Science* 17, no. 3 (2006): 249–255. Jeffrey D. Karpicke and Henry L. Roediger III, “The Critical Importance of Retrieval for Learning,” *Science* 319, no. 5865 (2008): 966–968. Testing effect strongest on delayed tests; often looks worse than restudy immediately. Pedagogy memo §1.2.

³⁷Tamara van Gog and John Sweller, “Not New, but Nearly Forgotten: The Testing Effect Decreases or Even Disappears as the Complexity of Learning Materials Increases,” *Educational Psychology Review* 27 (2015): 247–264. José Hanham, Wayne Leahy, and John Sweller, “Cognitive Load Theory, Element Interactivity, and the Testing and Reverse Testing Effects,” *Applied Cognitive Psychology* 31, no. 3 (2017): 265–280. Ouha Chen, Juan C. Castro-Alonso, Fred Paas, and John Sweller, “Extending Cognitive Load Theory to Incorporate Working Memory Resource Depletion,” *Educational Psychology Review* 30 (2018): 483–501. Pedagogy memo §1.2.

³⁸Louis Deslauriers, Logan S. McCarty, Kelly Miller, Kristina Callaghan, and Greg Kestin, “Measuring Actual Learning versus Feeling of Learning in Response to Being Actively Engaged in the Classroom,” *Proceedings of the National Academy of Sciences* 116, no. 39 (2019): 19251–19257. Illusion of learning: planted here for Bastani’s “students did not perceive the loss,” which is spent in Chapter 3. Pedagogy memo §2.1.

appeared in the student's tab.

Preview, in one sentence, the cleanest school demonstration, and then leave it for Chapter 3. Bastani, Bastani, Sungu, Ge, Kabakcı, and Mariman, in *PNAS* in 2025, found that unguarded GPT-4 during high-school math practice raised scores about 48 percent, after which the same students scored about 17 percent worse on an unaided exam than classmates who never had AI; a hint-and-grounded-solution tutor raised practice about 127 percent and left the unaided exam indistinguishable from control, and students in the unguarded arm did not perceive that they had learned less.³⁹ That is the dissociation, in a field experiment, with the illusion of learning attached. Guardrails are the treatment. Default ChatGPT is a crutch students do not notice. The table, the mechanism, the message clustering, and the limits of four ninety-minute sessions belong in Chapter 3. This chapter needs only the implication: if you grade the practice window, you will congratulate the crutch.

The OECD Digital Education Outlook 2026, as reported on the listing page, makes a parallel claim: unguided generative AI raises task performance without learning gains; generative AI can support learning when guided by teaching principles.⁴⁰ The full PDF was Cloudflare-blocked in the research fetch for this book. Treat full-report statistics as **UNVERIFIED** until that PDF is read. The listing-page sentence is usable as a direction of travel, not as a number. It is consistent with Soderstrom and Bjork, and with the one-sentence Bastani preview, and it is not a substitute for either.

What this distinction forbids is a habit already visible in 2026 reporting and in vendor decks: treating a better artefact as a better student. An essay that a model helped finish is not, by that finish, evidence that the student can write. A problem set completed in a chat window is not, by that completion, evidence that the student can solve. A code file that compiles after autocomplete is not, by that compile, evidence that the student can trace. The right success metric is unaided, delayed, off-tool performance of the skill the course claimed to form. Immediate fluency is the wrong one. An AI world does not change that. It makes the wrong metric cheaper to produce, and therefore more tempting to accept.

Offloading is how the ease gets into the student. It is not a moral failing. It is a mechanism, and generative tools have extended it to the skills school exists to form.

³⁹Bastani et al., “Generative AI without Guardrails,” *PNAS* 122, no. 26 (2025). <https://doi.org/10.1073/pnas.2422633122>. One-sentence preview only. Field RCT, large Turkish high school, Fall 2023–24, nearly 1,000 students in grades 9–11, four 90-minute sessions. Unguarded GPT-4 (GPT Base) practice about +48 percent, then unaided exam about –17 percent versus never-AI; GPT Tutor (hints, no full solutions, teacher-written correct solutions and common-mistake hints) practice about +127 percent, unaided exam indistinguishable from control. Students in the unguarded arm did not perceive that they had learned less. Full table, mechanism, and limits are Chapter 3. Pedagogy memo §2.1; brief, Chapter 3, fact 1.

⁴⁰OECD, *Digital Education Outlook 2026*, listing page: https://www.oecd.org/en/publications/oecd-digital-education-outlook-2026_062a7394-en.html. Listing page and search snippets: generative AI can support learning when guided by teaching principles; without pedagogical guidance, outsourcing tasks raises performance without learning gains. Roles named: tutor, partner, assistant. Full PDF fetch was **blocked (Cloudflare)**. Treat full-report statistics as **UNVERIFIED** until the PDF is read. Landscape memo §5.1; responsible-use memo §7 (OECD).

Cognitive offloading is old; generative AI extends it to the skills school exists to form

Evan Risko and Sam Gilbert, in *Trends in Cognitive Sciences* in 2016, defined cognitive offloading as using physical action or the environment to reduce the internal processing a task requires: tilting the head, setting a reminder, using a calculator.⁴¹ Offloading is not a moral failing. It is how human cognition scales. The research questions are what triggers it — often a metacognitive judgement that internal demand is high — and what it does to later unaided performance. A book that treats offloading as cheating has already lost the argument, because the student who writes a due date in a planner is offloading, and no one is proposing to ban planners. The educational stake is narrower and prior: what happens when the offloaded component was the learning target.

The older findings are the right analogy, and they are not flattering. Sparrow, Liu, and Wegner, in *Science* in 2011, described a Google effect: people remember *where* information is stored better than the information itself when they expect later access.⁴² Calculators and GPS produce analogous trade-offs. Performance goes up. Unaided skill goes down, if the offloaded component was the thing the course existed to form. Nobody is confused about this in arithmetic. A student who cannot add without a calculator has not learned to add, however fast the calculator is. Nobody is confused about it in navigation, either, once the battery dies. The confusion arrives when the offload is invisible, fluent, and attached to a skill that schools have always claimed as their own.

Generative AI extends offloading from memory and calculation to reasoning, drafting, and evaluation. Those are the skills schools exist to form. A calculator offloads arithmetic. A search engine offloads lookup. A fluent model offloads the generation of the argument, the structure of the essay, the selection of the warrant, and, if the student accepts the paragraph, the evaluation of whether the warrant holds. That is a different bargain. It is not “students have always used tools.” It is a change in which part of the cognitive loop remains inside the student. Frequency of use is not the variable. Whether the user retains the evaluative loop is. A 2026 conceptual distinction in *Frontiers in Psychology*, between dependent and autonomous offloading, is useful even before that paper’s psychometrics settle: the question is not how often the student opened the tool, but whether judgement stayed inside.⁴³

Jason Lodge and Leslie Loble, writing for Education International, put the practitioner synthesis in one contrast. Unstructured AI use trends toward detrimental offloading: performance up, durable learning down. Beneficial offloading needs an explicit pedagogical frame — for example, Load Reduction Instruction — so that what is offloaded is extraneous load, not the generative work the course is trying to build.⁴⁴ That is Soderstrom and Bjork in union-adjacent prose. It is also a design

⁴¹Evan F. Risko and Sam J. Gilbert, “Cognitive Offloading,” *Trends in Cognitive Sciences* 20, no. 9 (2016): 676–688. Offloading is using physical action or the environment to reduce internal processing — not a moral failing. Pedagogy memo §3.1; brief, Chapter 1, fact 2.

⁴²Betsy Sparrow, Jenny Liu, and Daniel M. Wegner, “Google Effects on Memory: Cognitive Consequences of Having Information at Our Fingertips,” *Science* 333, no. 6043 (2011): 776–778. People remember *where* information is stored better than the information itself when they expect later access. Pedagogy memo §3.1.

⁴³*Frontiers in Psychology* (2026), conceptual distinction between dependent and autonomous offloading. Frequency of AI use is not the variable; whether the user retains the evaluative loop is. Cited at abstract level in the research memos; psychometrics not treated as settled. Pedagogy memo §3.3; brief, Chapter 3, live debate.

⁴⁴Jason Lodge and Leslie Loble, “AI and Cognitive Offloading: Supporting Teachers to Shape AI’s Impact on Learn-

claim. The same model, under a pedagogical frame, is not the same treatment as the same model used to finish faster. Estonia’s AI Leap, already named, exists in part because most students were already doing the unstructured version. New York City Public Schools named cognitive offloading as a Playbook workstream on the March 2026 guidance page, still unfinished as of that fetch, which is itself evidence: the nation’s largest district has traffic lights for *teachers* and has not yet published the student-facing offloading rules.⁴⁵

The OECD listing-page claim already cited belongs here as well, with the same flag. Outsourcing without pedagogical guidance “enhances performance with no real learning gains.” Full-report statistics remain **UNVERIFIED** pending the PDF.⁴⁶ A book that needs a flagship international number will have to wait for a primary read. A book that needs the mechanism does not. The mechanism is in Risko and Gilbert, in Sparrow, in the calculator, and in the GPS. Generative tools change the *object* of the offload, not the existence of the offload.

What would actually dull the skill is not contact with models. Three conditions together, previewed here and spent in Chapter 3, are the dulling design. First, the offloaded process is the learning objective: writing, proof, source evaluation, the argument itself. Second, the environment never measures the skill unaided: the take-home essay is the only grade. Third, the student cannot detect model error — including arithmetic errors a student with the relevant school mathematics could have caught. Designs that keep generation, retrieval, and checking inside the student do not show an atrophy pattern on the outcomes they have measured. That is the opposite of a general “AI makes you dumber” law, and it is also the opposite of a general “give them the tool” law. The evaluative loop is the variable. Default chat drops it. Guarded, domain-loaded tasks can keep it.

A caution people will cite, and that this chapter will not carry: Nataliya Kosmyna and colleagues’ MIT Media Lab preprint, “Your Brain on ChatGPT,” arXiv 2506.08872, a small EEG study of SAT-style essays. It is a preprint, not a journal article as of 30 August 2026. It is not a K–12 field experiment. It should not be cited as proof that ChatGPT atrophies the brain, and it should not be extrapolated onto agent products it did not measure. Chapter 3 will treat it at the length a small preprint deserves. This chapter plants it so that a reader who arrived with the headline has somewhere to put it, and then refuses to let it set the frame.⁴⁷ The behavioural dissociation —

ing,” Education International, <https://www.ei-ie.org/en/item/32652:ai-and-cognitive-offloading-supporting-teachers-to-shape-ais-impact-on-learning>. Unstructured AI use trends toward detrimental offloading (performance up, durable learning down); beneficial offloading needs an explicit pedagogical frame (e.g., Load Reduction Instruction). Landscape memo §6.2.

⁴⁵New York City Public Schools, “Guidance on Artificial Intelligence,” March 2026, <https://www.schools.nyc.gov/about-us/vision-and-mission/artificial-intelligence/guidance-on-artificial-intelligence>. Cognitive offloading named as a Playbook workstream, still open as of the fetched page, alongside homework/integrity, grade-band guidance, family notification, and managed versus personal accounts. Landscape memo §5.2.

⁴⁶OECD *Digital Education Outlook 2026* listing-page claim, as in note 15. Outsourcing without pedagogical guidance “enhances performance with no real learning gains.” Full-report statistics **UNVERIFIED**. Landscape memo §6.2.

⁴⁷Nataliya Kosmyna, Eugene Hauptmann, Yeo T. Yuan, Jessica Situ, Xian-Hao Liao, Ashly Vivian Beresnitzky, Iris Braunstein, and Pattie Maes, “Your Brain on ChatGPT: Accumulation of Cognitive Debt When Using an AI Assistant for Essay Writing Task,” arXiv:2506.08872, MIT Media Lab, June 10, 2025. **Preprint, not a journal article as of August 30, 2026.** N = 54 Boston-area adults; session 4 n = 18; 20-minute SAT-style essays under EEG. Planted here as a caution people will cite; refused as a frame for this chapter; full treatment in Chapter 3. Do not cite as “MIT proved ChatGPT atrophies the brain.” Do not extrapolate onto agent products the study did not measure. Pedagogy memo

performance up, later unaided skill often not — is already available from Soderstrom and Bjork, and from the school experiment previewed above. EEG connectivity is not required to see it, and a convenience sample of fifty-four adults is not a developmental policy.

Agents increase the offloading surface. A system that finishes the job — writes, files, and sends — is a different cognitive bargain from a tutor that refuses the answer. The 2022–26 causal record is almost entirely about turn-based chat and tutoring tasks. Unattended completion is a stronger intervention and is under-studied in education. An honest book flags that gap rather than filling it with analogy. This chapter is still about the job the student must still do, in a world where the consumer tool already in students’ pockets is chat, not an unattended teammate. The extension of offloading to reasoning, drafting, and evaluation is already available in that world. Students are already doing it at home, unequally.

Students already use consumer chat — and the good tool is not equally held

Put 2026 student behavior on the page without inflating it into agent adoption. The Pew number that opened this chapter is chat for schoolwork, among U.S. teens 13 to 17, as of a 15 January 2025 release: 26 percent, double the 13 percent of 2023. Awareness jumped most in the lowest-income households. Access to the *idea* of the tool is democratising. Access to the *good* tool is not.⁴⁸ Do not treat these figures as Hermes-class or Grok Bot use. They are not. They are a conversation window, often a free tier, used to finish homework. That is already enough to break take-home assessment. It is not a district deployment of an unattended teammate, and writing as if it were is how a book about autonomous AI becomes a book about ChatGPT and calls itself a book about agents.

The paid-tier gap is measurable where researchers have looked. A 2024 Harvard undergraduate survey, arXiv:2406.00833, found that among AI users, 40 percent of students with no financial aid paid for AI products, against 20 percent of students on aid. The authors flag this as a potential inequity if premium quality is real — which, given model-tier gaps, it is.⁴⁹ Wonkhe, recording UK student voice, put the same fact in a sentence a dean can read: “students’ academic help may be determined by their financial resources rather than their academic competence.”⁵⁰ Homework-gap 2.0 is a paid-tier gap. The old homework gap was devices and broadband. This one is whether the student has the subscription that writes more fluently, remembers more context, and, in the agent case, keeps working after the laptop closes. A personal agent — persistent memory, tools, custom skills — is more unequal than a chatbot. It is closer to a private tutor plus a research assistant plus

§3.2; brief, “What not to overclaim” for Chapter 1 (do not spend Kosmyna here).

⁴⁸Pew, January 15, 2025, as in note 1. Awareness jumped most in the lowest-income households: access to the *idea* of the tool is democratising; access to the *good* tool is not. Responsible-use memo §4 (equity); brief, Chapter 1, fact 3.

⁴⁹Harvard Undergraduate Survey on Generative AI, arXiv:2406.00833 (2024), <https://arxiv.org/abs/2406.00833>. Among AI users, 40 percent of students with no financial aid paid for AI products versus 20 percent of students on aid. Authors flag a potential inequity if premium quality is real. Responsible-use memo, finding 8; brief, Chapter 1, fact 4.

⁵⁰Wonkhe (UK student voice), quoted in the responsible-use memo §4: “students’ academic help may be determined by their financial resources rather than their academic competence.” Primary Wonkhe URL was not in the source list — treat the page as **UNVERIFIED**; the wording is as the memo recorded it.

a secretary. Who owns it — the child, the family, the school, the state — decides who keeps the learning history, who can audit it, and who can turn it off. That hosting question is a later chapter. The equity fact is this one: do not grade homework that assumes every child has a premium personal tool the school does not issue.

Estonia's numbers are the national-scale version of the same behavior, and they are the reason Estonia built a program rather than celebrating the behavior. Sixty-four to 90 percent of students were already using free commercial apps to finish work faster. That pattern is what the AI Leap is trying to interrupt, not scale. The 2025/26 implementation: 154 of 156 upper secondary schools; about 4,700 to 5,000 teachers get paid ChatGPT or Gemini; about 20,000 grade 10–11 students get ITI, a Socratic Estonian-language app that does not give answers, developed with Estonian researchers and OpenAI. University of Tartu is running an impact study with Stanford and OpenAI. A first-year review was published 7 August 2026. Expansion in 2026/27 is toward vocational schools and grades 5–9.⁵¹ This is the cleanest national counter-example to “just give them ChatGPT.” It is also evidence that the unstructured use was already the default. A ministry that treats 64 to 90 percent as a market to serve, rather than a pattern to redesign, has chosen performance over learning in public.

UNESCO's connectivity figures sit under the paid-tier gap as a reminder that the U.S. and Estonian numbers are not the world. About 2.6 billion people still lack internet, a 2024 figure UNESCO continues to use. About one in four primary schools lack electricity. About 60 percent of primary schools are unconnected. UNESCO's 7 September 2023 generative-AI guidance warned that training corpora reflect Global North values, and that governments should not roll tools into classrooms faster than textbooks are approved. Fewer than 10 percent of 450-plus surveyed schools and universities had a generative-AI policy in 2023.⁵² Leapfrog versus dependency is a live fight this chapter will not settle. The implication for the job the student must still do is narrower: an “AI for all” slogan that skips unaided skill in connected systems, and skips electricity in unconnected ones, is two different failures with the same brand.

What these figures authorize, and what they do not, should be said in one place. They authorize the claim that consumer chat is already in the homework loop for a large minority of U.S. teens, for a majority of Estonian students in the bands the Leap measured, and for a paid-tier slice of Harvard undergraduates that splits on financial aid. They authorize the claim that access to the *idea* of the tool is spreading faster than access to the *good* tool. They do not authorize the claim that every student will have an agent. They do not authorize the claim that Pew measured agency. They do

⁵¹Estonia AI Leap / TI-Hüpe, as in note 2. 2025/26: 154 of 156 upper secondary schools; ~4,700–5,000 teachers get **paid** ChatGPT or Gemini; ~20,000 grade 10–11 students get ITI, a Socratic Estonian-language app that does not give answers, developed with Estonian researchers and OpenAI. University of Tartu impact study with Stanford and OpenAI. First-year review published August 7, 2026, <https://tihupe.ee/en/read-the-ai-leap-first-year-report/>; outcome coefficients **not pulled** this pass. 2026/27 expansion toward vocational and grades 5–9. 50/50 public–private funding (~€4 million pilot year). Responsible-use memo §7 (Estonia). The 64–90 percent problem statement is not itself a completed learning-outcomes result; the impact study is in progress.

⁵²UNESCO, *Guidance for Generative AI in Education and Research* (September 7, 2023), Fengchun Miao and Wayne Holmes. Training corpora reflect Global North values; do not roll tools into classrooms faster than textbooks are approved; <10 percent of 450-plus surveyed schools/universities had a GenAI policy in 2023. Connectivity figures UNESCO continues to use: ~2.6 billion people offline (2024 figure); ~1 in 4 primary schools without electricity; ~60 percent of primary schools unconnected. Responsible-use memo §4 (Global South) and §7 (UNESCO).

not authorize a national U.S. student-AI mandate; there isn't one. They do not authorize treating Estonia's ITI as a completed learning-outcomes result; the impact study is in progress. They do not authorize treating a personal agent as a closing of the gap. A personal agent is a widening, unless the institution issues it, contracts it, age-gates it, and still grades the skill unaided.

Policy is lagging the behavior, and it is not a classroom mandate. That lag is the next fact, and it is a document fact, not a vibe.

The assessment broke before the policy did

"Cheating" assumes the assignment still validly measures the intended skill. That assumption is the one 2026 policy is still catching up to. If the intended skill is "produce a five-paragraph essay unsupervised," a fluent model makes the assignment dead as a measure. The student who pastes the prompt and pastes back the paragraph has not cheated a living assessment. The assessment has already died, and the honor code is being asked to perform an autopsy. If the intended skill is "plan, judge, and defend a position," the assignment can live — as a studio, if design keeps generation, selection, and defence inside the student, or as a crutch, if it does not. Integrity is the wrong unit. Cognition is the right one. Policy, as of the documents fetched, is still mostly talking integrity, with cognition arriving as a workstream.

New York City Public Schools, the nation's largest district, issued March 2026 guidance on a traffic-light framework. Context, in the Chalkbeat and Gothamist accounts: it arrives about three years after a short-lived ChatGPT ban, and after a parent petition of about 1,500 signatures seeking a pause.⁵³ Red, never: placement, discipline, eligibility, promotion, graduation; IEP and 504 development; assessments and grading as the decision; surveillance and behavior monitoring; counseling and crisis; student data training models; personally identifiable information in unapproved tools. Yellow, human review required: data-pattern surfacing; translations of critical communications by a qualified linguist; scaffolds for diverse learners with certified-staff review; student use for research, exploration, and creative projects with educator guidance. Green, encouraged, still ERMA-approved, and the human owns the output: lesson brainstorming, drafting communications, scheduling and summarizing non-sensitive information, operational synthesis, non-critical translation with a disclaimer if unreviewed, accessibility, professional learning. ERMA, a ten-step vendor review, is required before any tool touches student data. AI-specific vendor rules added in December 2024: disclose AI capabilities; no training on student data; transparency. ERMA as fetched is privacy and security, not yet bias, equity, or instructional effectiveness, which were promised for a June 2026 Playbook. Playbook topics still open as of the fetched page included homework and integrity, cognitive offloading, grade-band guidance, family notification, and managed versus personal accounts.⁵⁴ Read that list again. The largest U.S. district has traffic lights for *teachers*

⁵³NYCPS March 2026 guidance, as in note 20. Context: Chalkbeat New York, March 24, 2026, <https://www.chalkbeat.org/newyork/2026/03/24/preliminary-ai-policy-nyc-schools/>; Gothamist, "NYC schools get AI guidance using red light, green light model." Arrives about three years after a short-lived ChatGPT ban; parent petition of about 1,500 signatures sought a pause. Landscape memo §5.2.

⁵⁴NYCPS March 2026 traffic-light framework, same primary page as note 20. Red / yellow / green as enumerated in the text. ERMA (10-step vendor review) required before any tool touches student data. AI-specific vendor rules added December 2024: disclose AI capabilities; no training on student data; transparency. ERMA as fetched is privacy/security, **not yet** bias, equity, or instructional effectiveness (promised for a June 2026 Playbook). Playbook topics

and has not yet published grade-band student-use rules, homework rules, or an offloading policy. Policy clarity is lagging student behavior. NYCPS says so by leaving the workstreams open. A book that treats a teacher-side traffic light as a settled student pedagogy has not read the page.

EDUCAUSE's 2026 report on the impact of AI on learning assessment, a survey of 438 faculty and staff dated to a June 2026 library page, says the same thing from the campus side. Students need clearer use and non-use rules. Most surveyed faculty use AI to *create* assessments and believe students use AI *while taking* them. Faculty want choice. AI literacy includes knowing when *not* to use tools. "Does it make teaching harder or easier? It depends." The full report PDF was not fetched. Treat additional statistics in that PDF as **UNVERIFIED** here.⁵⁵ The 2026 EDUCAUSE Horizon Report, Teaching and Learning edition, describes AI reshaping assessment toward authentic, process-based demonstrations of learning. Full report not fetched.⁵⁶ UIC's Center for the Advancement of Teaching Excellence published a teaching guide with three assignment types — recall, skill application, cumulative essays — and preserve-versus-leverage strategies, citing reduced cognitive engagement with generative-AI writing.⁵⁷ These are assessment-redesign documents. They are not agent deployments. They are evidence that higher education, like NYCPS, is treating the broken assignment as the problem, and has not yet agreed on the replacement.

Federal policy is a scoring tool for discretionary grants, not a classroom mandate, and not a child-safety regulation. The U.S. Department of Education's final supplemental priority, "Advancing Artificial Intelligence in Education," 13 April 2026, effective 13 May 2026, Docket ED-2025-OS-0118, 91 FR 18774, has two prongs: AI literacy and computer-science and AI coursework, including detecting AI-generated disinformation, dual enrollment, and teacher preparation; and ethical use of AI for tutoring, personalization, students with disabilities, gifted students, students below grade level, universal design for learning, and administrative efficiency. It defines AI via 15 U.S.C. § 9401(3): "a machine-based system that can, for a given set of human-defined objectives, make predictions, recommendations, or decisions influencing real or virtual environments." It does not define "agentic AI." It does not authorize unconstrained student agents.⁵⁸ It traces to Executive

still open as of the fetched page: homework/integrity, cognitive offloading, grade-band guidance, family notification, managed versus personal accounts. Brief, Chapter 1, fact 6.

⁵⁵EDUCAUSE, 2026 *EDUCAUSE: The Impact of AI on Learning Assessment Report*, library page June 2026, <https://library.educause.edu/resources/2026/6/2026-educause-the-impact-of-ai-on-learning-assessment-report>. Survey of **438** faculty/staff. Claims used here are from the library page: students need clearer use/non-use rules; most surveyed faculty use AI to *create* assessments and believe students use AI *while taking* them; faculty want choice; AI literacy includes knowing when **not** to use tools. **Full report PDF not fetched.** Additional statistics in that PDF are **UNVERIFIED**. Landscape memo §5.3.

⁵⁶2026 *EDUCAUSE Horizon Report* (Teaching and Learning): AI reshaping assessment toward authentic, process-based demonstrations of learning. **Full report not fetched.** Landscape memo §5.3.

⁵⁷University of Illinois Chicago, Center for the Advancement of Teaching Excellence, "Adapting Assignments for the Age of AI," <https://teaching.uic.edu/cate-teaching-guides/digital-learning/adapting-assignments-for-the-age-of-ai/>. Three assignment types (recall, skill application, cumulative essays) with preserve-versus-leverage strategies; cites reduced cognitive engagement with generative-AI writing. Landscape memo §5.3.

⁵⁸U.S. Department of Education, "Final Priority and Definitions: Secretary's Supplemental Priority and Definitions on Advancing Artificial Intelligence in Education," 91 Fed. Reg. 18774 (April 13, 2026; effective May 13, 2026), Docket ED-2025-OS-0118, <https://www.federalregister.gov/documents/2026/04/13/2026-07087/final-priority-and-definitions-secretarys-supplemental-priority-and-definitions-on-advancing>. Scoring tool for **discretionary grants**, not a classroom mandate. Two prongs as in the text. AI defined via 15 U.S.C. § 9401(3). Does **not** define "agentic AI" and does not authorize unconstrained student agents. Public comments included both strong

Order 14277 of 23 April 2025, “Advancing Artificial Intelligence Education for American Youth”: a Task Force, a Presidential AI Challenge, public–private K–12 resources, Education Department priority for AI in ESEA and HEA Title II teacher-training grants, NSF research, USDA 4-H. That order is an adoption-and-workforce order. It is not a child-safety regulation. The actual federal child-safety action in this file is at the FTC — COPPA amendments and a 6(b) chatbot study — not at Education.⁵⁹ Public comments on the supplemental priority included both strong support and calls to keep AI out of K–12 over cognitive and emotional harm. The Department declined a ban.

Do not collapse the executive orders. Biden Executive Order 14110 of 30 October 2023, “Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence,” directed Education to produce education AI resources. It was rescinded on 20 January 2025 by Executive Order 14148.⁶⁰ The May 2023 Department report, *Artificial Intelligence and the Future of Teaching and Learning*, still sits in the file: seven recommendations, humans in the loop, education-specific guardrails, and an explicit rejection of teacher replacement. Follow-on toolkits were written to 14110’s 365-day clock. They are not 14277, and 14277 is not a safety statute. A book that treats 14277 as a child-protection instrument has read a workforce order as a ban. A book that treats 14110 as still in force has not checked the date.

UNESCO’s 7 September 2023 *Guidance for generative AI in education and research* (Miao and Holmes) is the international document of record on age and posture: human-centred; data-privacy mandates; age 13 for independent classroom generative-AI conversation; under-13 and primary use of untested tools “not appropriate”; seven steps for government regulation; a warning that training data encode Global North norms.⁶¹ The 2024 competency frameworks, launched at Digital Learning Week, give the literacy architecture: for students, twelve competencies, four dimensions, three levels (Understand, Apply, Create); for teachers, fifteen competencies, five dimensions (human-centred mindset, ethics, foundations and applications, pedagogy, professional learning), levels Acquire, Deepen, Create.⁶² Digital Learning Week 2026, 8–11 September, Paris, is titled “Education

support and calls to keep AI out of K–12; ED declined a ban. Landscape memo §5.1; brief, Chapter 1, fact 5.

⁵⁹Exec. Order No. 14277, “Advancing Artificial Intelligence Education for American Youth” (April 23, 2025). Task Force; Presidential AI Challenge; public–private K–12 resources; ED to prioritize AI in ESEA/HEA Title II teacher-training grants; NSF research; USDA 4-H. This is an **adoption-and-workforce order, not a child-safety regulation**. The actual federal child-safety action in this research file is at the FTC (COPPA 2025 amendments; 6(b) chatbot study orders of September 11, 2025), not at ED. Responsible-use memo §7 (United States) and open policy fight 1. AFT Commonsense Guardrails v2 (March 2025): “the current federal administration has indicated it will not regulate AI.”

⁶⁰Exec. Order No. 14110, “Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence” (October 30, 2023), directed ED to produce education AI resources. **Rescinded January 20, 2025, by Exec. Order No. 14148**. Follow-on ED toolkits (ERIC ED661924, ED661949) were written to 14110’s 365-day clock. U.S. Department of Education, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023), remains in the file as the 2023 humans-in-the-loop document, including the replacement rejection in note 7. Responsible-use memo §7; brief, Chapter 1, “What not to overclaim.”

⁶¹UNESCO, *Guidance for Generative AI in Education and Research*, September 7, 2023 (Miao and Holmes). Human-centred; data-privacy mandates; **age 13** for independent GenAI conversation; under-13 / primary use of untested tools “not appropriate”; seven steps for government regulation; warning that training data encode Global North norms. Responsible-use memo §7 (UNESCO); brief, Chapter 1, fact 7.

⁶²UNESCO *AI Competency Framework for Students* (2024, Miao and Shiohira): 12 competencies, four dimensions, three levels (Understand / Apply / Create). UNESCO *AI Competency Framework for Teachers* (Miao and Cukurova): 15 competencies, five dimensions (human-centred mindset, ethics, foundations/applications, pedagogy, professional learning); levels Acquire / Deepen / Create. Launched Digital Learning Week 2024. Digital Learning Week 2026, September 8–11, Paris: “Education in the age of AI.” TVET AI guide dated July 3, 2026. UNESCO does **not** endorse

in the age of AI.” UNESCO does not endorse student-facing autonomous agents. Thirteen is a UNESCO independent-use line and a common terms-of-service number. It is not a developmental finding, and it is not a last-row license. Singapore’s bands are stricter in the early years: no student chatbots in Primary 1–3; Primary 4–6 only education-designed tools under teacher supervision.⁶³ Age as a design specification is Chapter 4. This chapter needs only the floor: international guidance already said independent generative conversation is not a primary-school default, while U.S. federal policy in 2026 is still a grant-scoring tool that declined a ban.

The live fight is not “AI yes or no.” It is literacy versus formation. A grant priority that funds AI literacy and ethical use for tutoring can be read, generously, as a formation agenda. It can also be read as a workforce checklist that never asks whether the student can still do the work when the window closes. That is the debate the next section stages, without settling the graduate chapter or the judgement chapter.

Labor-market literacy versus cognitive formation

The live debate this chapter takes a position on, without settling Chapters 6 and 10, is whether school in an AI world is for workforce-entry fluency or for unaided skill. The two are not the same outcome. They are being sold as if they were.

Campus-ops and access stories supply the literacy pole. Clemson News describes ChatGPT Edu, GPT-5.5, Codex, and locally hosted open-weight models for researchers. That is documented access. It is not a published learning-outcomes study.⁶⁴ EdTech Magazine, in December 2025, quotes Dan Harder of the University of Tennessee Knoxville: today’s campus AI is agent-like within parameters, “a long way from autonomous AI agents that weigh the pros and cons... and act without human oversight.” The same piece cites UT Verse, campus assistants for Wi-Fi, dining, coursework, translation, and Johns Hopkins Agent Laboratory. UT Verse’s primary pages and the Agent Laboratory paper were not fetched. Treat those two campus examples as journalism-sourced, **UNVERIFIED** against primary documents.⁶⁵ “Students must be AI proficient at graduation” is the

student-facing autonomous agents. Landscape memo §§1.4, 5.1.

⁶³Singapore Ministry of Education page, last updated July 31, 2026. Age bands: **P1–3** print-first, no student chatbots; **P4–6** only MOE-designed Student Learning Space features, teacher-supervised; secondary/pre-U increasing student ownership plus integrity rules. Parliamentary reply May 6, 2026: education-designed tools only for P4–6. Responsible-use memo §7 (Singapore). Thirteen is a UNESCO independent-use line and a common terms-of-service number; it is not a last-row license. Age as design specification is Chapter 4.

⁶⁴Clemson News, “GPT-5.5, Codex and Clemson-hosted models give researchers new options for AI-assisted work,” <https://news.clemson.edu/gpt-5-5-codex-and-clemson-hosted-models-give-researchers-new-options-for-ai-assisted-work/>. Documented **access** (ChatGPT Edu + GPT-5.5 + Codex + locally hosted open-weight models), not a published learning-outcomes study. Landscape memo §4.1.

⁶⁵EdTech Magazine, December 2025, <https://edtechmagazine.com/higher/article/2025/12/ai-agents-higher-education-transforming-student-services-and-support-perfcon>. Dan Harder, University of Tennessee Knoxville: today’s campus AI is agent-like within parameters, “a long way from autonomous AI agents that weigh the pros and cons... and act without human oversight.” Same piece cites Johns Hopkins Agent Laboratory (authors claim 84 percent cost decrease versus prior autonomous-research methods) and UT Verse at UT Knoxville (campus assistants for Wi-Fi, dining, coursework, translation). Agent Laboratory paper and UT Verse primary pages were **not fetched**. Treat those two campus examples as journalism-sourced, **UNVERIFIED** against primary documents. Landscape memo §4.7; brief, Chapter 2, fact 7.

sentence these stories license. It is an access-and-ops sentence. It is not a trial of whether proficiency is prompt skill, or whether it is the capacity to catch the model, or whether it is the ISTE six roles, or whether it is Cornell's question. Do not treat a newsroom account of a campus license as a cognitive result.

Cornell's question is the formation pole, stated as an institutional program rather than as a slogan. The *Cornell Chronicle*, in August 2026, reported that the AI Critical Literacy Program expands from a spring pilot to all incoming students plus interested faculty and staff. Four Canvas modules: What is GenAI?; Ethics; GenAI and Learning; Creating an AI Policy for Personal Use. Center for Teaching Innovation and the Library. Faculty from writing, business, biology, engineering, economics, communication, information science. The pilot reported literacy gains on how models work and are trained — CTI-reported, not a published randomized trial. The explicit goal is to help students decide whether AI use *contributes to learning*.⁶⁶ That sentence is the right question. It is not, by itself, an answer. Literacy gains on how models are trained are not unaided delayed skill in a domain. They are a start, and they are a start that refuses to treat use as the outcome.

The American Federation of Teachers, in a resolution adopted 18 July 2026, “Preparing for the Age of Artificial Intelligence...,” treats elementary student-facing AI as developmental risk, not a tooling choice: no screens in pre-K through grade 2 except for disability or compelling need; oppose student-facing AI in elementary; ban social-companion chatbots at least until age 16; AI as a limited supervised tool; public money to class size and mental health before unproven AI.⁶⁷ That is a formation-first labor position, and it is a position about childhood, not about whether a college senior may use a campus Codex seat. Collapsing the AFT elementary line into a general “unions versus AI” story is the popular legend in another font. The union is drawing an age floor. Age floors are Chapter 4. The fact for this chapter is that a national teachers’ union has already refused to treat elementary student-facing AI as a workforce-literacy decision.

ISTE’s Profile of an AI-Ready Graduate, unveiled at ISTE Live 2026 in Orlando, names six roles: Learner, Researcher, Synthesizer, Problem Solver, Connector, Storyteller. Discipline-agnostic, built on ISTE Student Standards.⁶⁸ It is a profile, not an outcomes study. UNESCO’s student and teacher frameworks, already cited, are competency architectures, not evidence that schools

⁶⁶Cornell Chronicle, August 2026, “Cornell expands AI Critical Literacy program to incoming students,” <https://news.cornell.edu/stories/2026/08/cornell-expands-ai-critical-literacy-program-incoming-students>. Expands from spring pilot to **all incoming students** plus interested faculty/staff. Four Canvas modules as in the text. CTI + Library. Pilot reported literacy gains (how models work / are trained) — **CTI-reported, not a published RCT**. Explicit goal: help students decide whether AI use **contributes to learning**. Landscape memo §5.3; brief, Chapter 1, live debate.

⁶⁷American Federation of Teachers, resolution adopted July 18, 2026, “Preparing for the Age of Artificial Intelligence...,” <https://www.aft.org/resolution/preparing-age-artificial-intelligence-protecting-democracy-workers-students-and-planet>. No screens PK–2 except disability/compelling need; **oppose student-facing AI in elementary**; ban social-companion chatbots at least until age 16; AI as limited supervised tool; National Academy for AI Instruction; “gold standard” privacy; public money to class size/mental health before unproven AI. Landscape memo §5.2; brief, Chapter 1, live debate.

⁶⁸ISTE Profile of an AI-Ready Graduate, ISTE Live 2026, Orlando: Learner, Researcher, Synthesizer, Problem Solver, Connector, Storyteller — discipline-agnostic, built on ISTE Student Standards. A profile unveiled at a conference, not an outcomes study. (Culatta’s 14 July 2025 ASCD headings match the ISTE Live 2026 six — Problem Solver, not Ideator. ISTE’s Generation AI news page summarizing that blog lists Ideator in place of Problem Solver; do not cite Culatta 2025 as an Ideator list.) Brief, Chapter 6, fact 3. Full treatment of literacy frameworks is Chapter 6.

have implemented them. The U.S. Department of Education’s literacy definition in 91 FR 18774 — technical knowledge, durable skills, civic awareness, ethical reasoning, detecting AI-generated disinformation — is a grant-scoring definition. None of these documents is a warrant for skipping unaided skill. Willingham’s constraint, already installed, predicts that a graduate who can recap a literacy framework and cannot evaluate a historical claim, a proof, or a lab result has a checklist, not a formed capacity. Chapter 10 will have to say what an AI-fluent graduate is in deeds. This chapter will not fill that question with analogy. Whether “AI-fluent graduate” means prompt skill or the ISTE, Cornell, and two-lane bundle is an open question this book leaves standing until the graduate chapter.

The Kosmyna preprint, planted above as a caution people will cite, belongs in this debate only as a caution. It will be offered, in some rooms, as the decisive evidence that literacy-without-formation is already a neural fact. It is not that evidence. Small N; convenience sample of elite-university students; twenty-minute essays; EEG connectivity is not “getting dumber”; not K–12; not agentic systems; preprint as of 30 August 2026. Full treatment is Chapter 3. Do not let a preprint carry this chapter, and do not let a campus access story carry it the other way.⁶⁹

The book’s position, stated here so later chapters can spend it rather than rediscover it: formation is not optional, and literacy that skips unaided skill is a checklist. Students should learn to use tools. They should not be credentialed for using tools as a substitute for being able to retrieve, generate, detect error, and judge when the tool is gone. Preview, in one sentence, the assessment rule Chapter 5 will specify: two lanes — formation, in which AI is forbidden or hint-only and unaided work is graded; and AI-world performance, in which AI is allowed, logged, and orally or process-verified.⁷⁰ That rule is a theoretically motivated hypothesis, not a fitted grade-by-grade schedule. The optimal mix of AI-forbidden and AI-required tasks by grade is an open question this book will not fill. Year-scale unaided skill under student-facing generative tutors, on independent state tests, is also not yet known: Khanmigo’s TCAP results were still outstanding as of the August 2026 paper; the Nigeria and Harvard physics trials did not follow a year later; Bastani measured the next hour.⁷¹ Name those holes. Do not paper them with workforce prose.

Vendor “tutor for every child” rhetoric collides with the tutoring benchmark already on the table. Bloom remains about human one-to-one tutoring. Nickow, Oreopoulos, and Quan remain at 0.288 standard deviations for experimental preK–12 tutoring. El Salvador remains an announcement. Khanmigo scale remains a vendor claim. The causal record on generative tutors is Chapter 3, and

⁶⁹Kosmyna et al. 2025, as in note 22. Limits restated so the literacy-versus-formation debate cannot recruit a preprint as a neural verdict. Full treatment, Chapter 3.

⁷⁰Two-lane assessment rule, theoretically motivated, not a fitted schedule: **formation** (AI forbidden or hint-only; grade unaided retrieval, derivation, draft, oral) versus **performance in an AI world** (AI allowed, logged; grade judgement — selection, critique, verification, process). Pedagogy memo §6.4; Mike Perkins, Leon Furze, Jasper Roe, and Jason MacVaugh, “The Artificial Intelligence Assessment Scale (AIAS),” *Journal of University Teaching and Learning Practice* 21, no. 6 (2024), is shared vocabulary, not an RCT of learning. Full specification is Chapter 5. Optimal mix of AI-forbidden and AI-required tasks by grade remains Open question 11.

⁷¹Open questions this chapter leaves standing. Year-scale unaided skill under student-facing generative tutors, on independent state tests: not yet known (Open question 1). Oreopoulos and Low 2026 Khanmigo cluster RCT: TCAP registered as co-primary, **not yet incorporated** as of the August 2026 paper; ITT magnitudes resemble Khan Academy practice without the chatbot. De Simone et al. 2025 (Nigeria) and Kestin et al. 2025 (Harvard physics) did not follow a year later. Bastani et al. 2025 measured the next hour. Whether an “AI-fluent graduate” is prompt skill or the ISTE/Cornell/two-lane bundle lives in Chapter 10. Pedagogy memo §9; brief, Chapter 3, “What not to overclaim.”

it is not a two-sigma story. A literacy agenda that cites Bloom and ships a chatbot has performed the collapse this chapter exists to make expensive.

The design problem, stated as the chapter's payload, is not whether to give students agents. It is which cognitive work must remain inside the student.

The design problem: which cognitive work stays inside the student

The book's problem is not whether to give students agents. It is which cognitive work must remain inside the student at each stage of expertise, and which work can be offloaded without dulling the skill being formed. That is the ninth correction, almost as load-bearing as the eight that the introduction will take in full: design is the treatment. "AI in education" is not one intervention. Unguarded ChatGPT, a hint-only tutor, a teacher-prompted Copilot lab, Khanmigo sitting one click away, and a custom physics tutor with worked solutions in the prompt are different pedagogies that share a model class.⁷² Collapse them and every later chapter becomes folklore: a harm result about answer-giving GPT becomes a ban on tutors; a help result about a custom physics tutor becomes a warrant for consumer chat; an access result about a Socratic app that students do not message becomes a proof that "AI tutoring works." Design is the independent variable. The model class is not.

Three conditions together dull. Preview them; spend them in Chapter 3. First, the offloaded process is the learning objective — writing, proof, source evaluation, the argument itself. Second, the environment never measures the skill unaided: the take-home essay is the only grade, the problem set is accepted as submitted, the repository is graded as it compiles. Third, the student cannot detect model error, including errors the student's own school knowledge would have caught if the student had been looking. Contact with models is not what dulls. Frequency of use is not the variable. Whether the user retains the evaluative loop is. Designs that keep generation, retrieval, and checking inside the student have not shown the atrophy pattern on the outcomes they measured. Designs that hand the generative work to the model raise the artefact and often not the skill. That is not a general law about intelligence. It is a design result, and it is the result this book will not trade for a slogan.

What must remain inside, at minimum, is the list this chapter opened on. Retrieval: the student can still produce the fact, the formula, the line of the text, when the window is closed. Generation: the student can still produce a sentence, a step, a plan, without the model supplying the first draft. Error-detection: the student can still catch a wrong warrant, a hallucinated citation, an arithmetic slip. Domain knowledge: there is something to retrieve, generate, and check *against*, because Willingham is the constraint. Judgement: the student still decides what to keep, what to cut, and

⁷²Ninth correction: design is the treatment. Pedagogy memo, "How to read this memo"; brief, pp. 1–2. Unguarded ChatGPT, a hint-only tutor, a teacher-prompted Copilot lab, Khanmigo one click away, and Kestin's PS2 Pal are different pedagogies sharing a model class. The design problem is which cognitive work must remain inside the student at each stage of expertise. 91 FR 18774, as in note 33, does not define "agentic AI" and does not authorize unconstrained student agents. "Every student will have an agent" is not a 2026 school fact; no documented K–12 student-facing autonomous open-web agent program was found. Landscape memo §§5.2, 7.

what to believe. The disposition to check: the habit, not the prompt, of treating fluency as a reason to verify rather than a reason to accept. An AI world does not retire those. It makes them scarce, because the environment will happily finish the job without them.

What can be offloaded without dulling depends on the stage of expertise and on what is being formed. That is cognitive load theory and expertise reversal, not a vibe, and Chapter 5 will move the slider by subject. For a novice, a worked example can be the right offload, because unguided problem-solving burns working memory on search instead of schema. For an intermediate, the same full solution is the wrong offload, because the student now needs to generate and to vary. Grammar as a low-stakes offload, once ideas exist, is a different bargain from first-draft generation. A hint that names a common mistake is a different bargain from a pasted answer. The book will not fit a grade-by-grade schedule of forbidden and required tasks. That mix is an open question. The principle is not: keep the learning objective inside; measure it unaided; do not offload the part the student cannot yet check.

Federal prose does not settle this, and it should not be asked to. 91 FR 18774 does not define “agentic AI” and does not authorize unconstrained student agents. It funds literacy, tutoring, personalization, UDL, and administrative efficiency. Those are grant categories. They are not a design specification of which cognitive work stays inside. “Every student will have an agent” is not a 2026 school fact. There is no documented K–12 student-facing autonomous open-web agent program in the sources this chapter stands on. What schools bought is teacher-side chat and workflow tools, constrained Socratic tutors, and pre-generative adaptive platforms. Agentic systems are entering teacher work and campus operations. They are not entering third-grade desks as unconstrained teammates. Chapter 2 will make that classification expensive, product by product. This chapter needed only to stop the popular story from arriving as if Pew were agency, as if a polished homework were learning, and as if a workforce order were a pedagogy.

Do not write as if the AI world has already replaced the need for unaided skill. It has not. It has made unaided skill harder to see, cheaper to fake, and more valuable when it is real. Do not treat Pew teen-ChatGPT use as agent use. Do not cite “every student will have an agent” as a 2026 school fact. OECD full-report statistics remain **UNVERIFIED** pending the PDF. Do not treat Executive Order 14277 as a child-safety regulation; it is an adoption-and-workforce order. Biden Executive Order 14110 was rescinded on 20 January 2025. Do not spend the MIT EEG preprint here. Do not tour last-row product features here. Do not dump the Bastani table here.

Leave standing what analogy would otherwise fill. Year-scale unaided skill under student-facing generative tutors, on independent state tests, is not yet known. Whether an “AI-fluent graduate” is prompt skill or the ISTE, Cornell, and two-lane bundle lives in Chapter 10. The optimal mix of AI-forbidden and AI-required tasks by grade is not a fitted schedule. The gap between chat-task evidence and unattended agent offloading is real; flag it; do not extrapolate a writing-task study onto a system that finishes the job while the student is at lunch.

Before deciding what work stays inside the student, stop collapsing four product classes into “AI.” Chat is not a tutor. A tutor is not an adaptive engine. An adaptive engine is not an agent. The named objects of this book sit in the last row, and almost none of that row is a documented classroom product. That classification is Chapter 2. The job the student must still do does not wait on it. Retrieval, generation, error-detection, domain knowledge, judgement, and the disposition to check

are still the job. An AI world makes them the scarce resource. School exists to form them. Design is how.

Chapter 2

Chat, Tutor, Adaptive System, Agent



Figure 3: Four objects in a row: a notebook, a stack of cards, a marked book, and a computer tower.

The popular vocabulary of 2026 has a mass-noun problem. It says *AI* the way a lazy brief says *the technology*, as if a conversation window, a Socratic tutor that refuses the answer, a pre-ChatGPT knowledge map that serves the next item, and a system that keeps working after the laptop is closed were one object. They are not. Collapse them and every later chapter of this book becomes folklore: Khanmigo becomes an agent, Grok-in-El-Salvador becomes Grok Bot, a Hermes weight file becomes a school product, and a teacher who schedules a 4 p.m. folder analysis is said to have put an autonomous teammate on a third-grader's desk.

This chapter exists to make that collapse expensive. The argument is a table, then a walk through each row, then two product-family corrections that will make the book obsolete if they do not survive copy-edit, then a look at what schools and universities actually bought. The book's named objects — Nous Hermes Agent, Grok Bot / Cursor Grok Bot, and similar tool-using, time-running assistants — sit in the last row. Almost none of that last row is a documented, scaled classroom product. What *is* in schools at scale in 2026 is teacher-side chat and workflow tools, constrained student tutors, and pre-generative adaptive platforms. Agentic systems are entering teacher work and campus operations. They are not entering K–12 student instruction as unconstrained agents.

Documents do this work. Product pages, a Federal Register priority, a district guidance page, vendor research labelled as vendor research. No launch party. No invented classroom.

Four classes, one table

Four product classes are routinely collapsed in education discourse. Print them as a table so the collapse has to happen in public.

Class	What it does	Typical interaction	Persistence / tools	Education examples (named)
Chatbot / chat assistant	Answers a prompt in a conversation window	Reactive; one turn or a thread	Usually no independent computer; limited or no unattended action	Consumer ChatGPT; Grok on grok.com; Gemini app
Tutor app (including Socratic AI tutors)	Guides a learner through content, typically refusing to give the answer outright	Student-facing dialogue, often inside a curriculum	Session memory; rarely schedules jobs or signs into other apps	Khanmigo; SchoolAI Spaces; MagicStudent
Adaptive learning platform / ITS	Sequences problems from a learner model; adjusts difficulty	Item-level practice, not open chat	Strong learner model; weak open-ended agency	ALEKS; MATHia; i-Ready; DreamBox
Autonomous / agentic system	Takes a goal, plans, uses tools, and continues work over time, including while the user is away	Job / teammate / workflow, not a Q&A box	Own computer or VM; browser/terminal; cron/routines; multi-agent	Hermes Agent; Grok Bot; Cursor Cloud Agents; Claude Cowork; OpenAI Workspace Agents / Codex Cloud

The four names in the adaptive row are commonly cited in 2026 secondary roundups. Official product pages for ALEKS, MATHia, i-Ready, and DreamBox were not fetched for the research this chapter stands on. Feature lists and user counts for those platforms are **UNVERIFIED** here. The names stay in the table because the *class* is real and older than ChatGPT. The features do not get invented to fill the cells.⁷³

The book's named objects sit in the last row. That is not a compliment and it is not a procurement claim. It is a classification. Hermes Agent, Grok Bot, Cursor Cloud Agents, Claude Cowork, and OpenAI Workspace Agents all claim, in their own documentation, an infrastructure that chat does not have. Almost none of that infrastructure is yet a documented, scaled classroom product. The education market of 2026 bought the first three rows, with a thin, teacher-facing strip of the fourth.

A chatbot with tools is still a chatbot if the human must stay in the loop for every step. Tool use is not autonomy. A well-prompted ChatGPT session can *simulate* planning, browsing, and multi-step work with the human as the execution engine: the model drafts, the person pastes, the model suggests, the person clicks. Agents move execution off the human. That is the distinction this chapter will not allow to blur. Web search inside a conversation window does not make ChatGPT into Hermes Agent. A Google Workspace connector does not make Gemini into Grok Bot. A "GPT" a school built for its help desk is not, by that naming, an unattended job with its own computer.

The operational definition of an agent used here is synthesized from vendor docs, not from a standards body. An agent is a system that (1) is given a *goal* rather than a single prompt, (2) plans and executes multi-step work, (3) uses tools — terminal, browser, files, APIs, other agents — (4) can continue after the user leaves, and (5) has some memory, skill, or routine that persists across sessions.⁷⁴ Hermes Agent, Grok Bot, Cursor Cloud Agents, Claude Cowork, and OpenAI Workspace Agents all claim that infrastructure in their own words. The U.S. Department of Education, in the final supplemental priority effective 13 May 2026, does *not* define "agentic AI." It defines AI via 15 U.S.C. § 9401(3): "a machine-based system that can, for a given set of human-defined objectives, make predictions, recommendations, or decisions influencing real or virtual environments." The priority funds AI literacy, tutoring, personalized learning, and administrative efficiency. It does not authorize unconstrained student agents.⁷⁵ UNESCO's 2024 competency frameworks and 2025–26 education pages frame a human-centred approach and "preserving human agency in a world of automation." UNESCO does not ship a product taxonomy of agents versus chatbots, and it does not endorse student-facing autonomous agents.⁷⁶ Illinois statewide guidance is reported, in a Franczek

⁷³ALEKS (McGraw Hill), Carnegie Learning MATHia, Curriculum Associates i-Ready, and Discovery Education / DreamBox are named in 2026 secondary roundups as adaptive / ITS platforms. Official product pages were not fetched for the research this chapter stands on. Feature lists and user counts for these products are **UNVERIFIED**. The names identify the class. See landscape memo, "How to read this memo" table and §1.3.

⁷⁴Operational definition synthesized from vendor documentation, not from a standards body: a system that (1) is given a goal rather than a single prompt, (2) plans and executes multi-step work, (3) uses tools (terminal, browser, files, APIs, other agents), (4) can continue after the user leaves, and (5) has durable memory, skill, or routine. The five named systems claim that infrastructure in their own product docs, cited below.

⁷⁵U.S. Department of Education, "Final Priority and Definitions: Secretary's Supplemental Priority and Definitions on Advancing Artificial Intelligence in Education," 91 Fed. Reg. 18774 (Apr. 13, 2026; effective May 13, 2026), Docket ED-2025-OS-0118. AI is defined via 15 U.S.C. § 9401(3). The priority does **not** define "agentic AI" and does not authorize unconstrained student agents. It is a scoring tool for discretionary grants, not a classroom mandate. Traces to Exec. Order No. 14277 (Apr. 23, 2025).

⁷⁶UNESCO, Artificial Intelligence in education, <https://www.unesco.org/en/digital-education/artificial-intelligence>;

/ JDSupra summary of Public Act 104-0399, to name agentic AI as a core concept districts should understand. The ISBE PDF itself was not fetched. Treat that naming as **UNVERIFIED** pending the primary document.⁷⁷

The gap that chat cannot close is not “smarter answers.” It is infrastructure, and the primary product docs are explicit about it. First, a computer that is not the user’s laptop: Grok Bot’s cloud computer, a Cursor Cloud Agent VM, Hermes local or VPS or Modal or Daytona backends. Second, unattended runtime: cron, scheduled routines, overnight jobs. Third, acting inside real apps — sign-in, click, type, file a ticket, push a branch — rather than drafting text for a human to paste. Fourth, parallel specialists: multiple Bots, sub-agents, a group chat of agents. Fifth, durable memory and reusable skills: Hermes skill documents, Grok Bot routines learned by watching the user, Claude Skills, OpenAI skills inside Workspace Agents.⁷⁸ Keep those five in view. They are what the last row *claims*. Whether any of them belongs on a student’s desk is a later chapter. Whether any of them is *on* a student’s desk, as a documented K–12 program, is a question this chapter can already answer: no fetched source shows it.

The collapse has a mechanism, and naming the mechanism is part of making it expensive. Education already had a word, *tutor*, that covered a human, an ITS, and, after 2023, a chatbot with a Socratic system prompt. Vendors then added *agent* to the same shelf, because *agent* is what the laboratory market is buying. A superintendent hearing all four words in one sales meeting is not being careless. The language is being asked to do work it cannot do. The four-class table is a refusal of that work. When a later chapter says “tutor,” it will mean the second row unless it says otherwise. When it says “agent,” it will mean the last row. When it says “Grok,” it will say which Grok. That is not pedantry. It is how a book about autonomous AI avoids writing a book about chat windows and calling it a book about agents.

Walk the first row so “Grok” cannot mean three products, and so “chat” cannot be stretched into a teammate.

Chat: the conversation window

A chatbot is a conversational interface that waits for a user prompt and returns text, or multimodal output. It may have tools — web search, file upload, connectors — but the unit of work is a *conversation*, not an unattended job. The human stays in the loop because the loop is the product: a thread. Close the tab and the work stops, except insofar as a vendor keeps a history the user can reopen. That is not a moral failing of the class. It is the class.

AI Competency Framework for Students, <https://www.unesco.org/en/articles/ai-competency-framework-students>; *AI Competency Framework for Teachers*, <https://www.unesco.org/en/articles/ai-competency-framework-teachers>. Human-centred framing; “preserving human agency in a world of automation.” UNESCO does not ship a product taxonomy of agents versus chatbots and does not endorse student-facing autonomous agents.

⁷⁷Franczek P.C. / JDSupra summary of Illinois Public Act 104-0399 and related ISBE guidance, <https://www.jdsupra.com/legalnews/new-artificial-intelligence-9071517/>. The claim that statewide guidance names agentic AI as a core district concept is **UNVERIFIED** against the ISBE PDF, which was not fetched.

⁷⁸The five infrastructure gaps are taken from primary product documentation assembled in landscape memo §1.5: Grok Bot cloud computer; Cursor Cloud Agent VM; Hermes local / VPS / Modal / Daytona backends; cron and scheduled routines; in-app acting (sign-in, click, type); parallel Bots / sub-agents; Hermes skill documents; Grok Bot routines; Claude Skills; OpenAI skills inside Workspace Agents. A well-prompted chat session can *simulate* some of this with the human as execution engine. Agents move execution off the human.

Three named surfaces, fetched from primary product documentation, sit in this row and are the surfaces students and teachers actually meet.

Grok, the assistant, is documented by xAI as an assistant on grok.com and on iOS and Android: chat, Imagine, voice, file upload, connectors. Distinct, in xAI's own product line, from Grok Bot. The overview at docs.x.ai/grok/overview is the chat-centric product. That is the ChatGPT-class surface.⁷⁹ It is the thing a student can open. It is not a cloud computer that keeps working at 2 a.m. on a sales sequence, and it is not, as a later section will insist, the El Salvador announcement.

Gemini app and Gemini Notebook, in Google Workspace for Education, are chat and notebook grounded in user-provided sources, included free in all Google Workspace for Education editions, for all ages, as a Core service. Google claims the data are not used to train models, and claims COPPA and FERPA support. Those are vendor and admin claims. Admin organizational-unit controls exist. Youth onboarding is endorsed, in Google's materials, by ConnectSafely and the Family Online Safety Institute. LearnLM, and the claim that a "2.5 Pro" learning-tuned model is the "world's leading model for learning," is Google marketing, and will be labelled as such whenever it reappears.⁸⁰ NotebookLM, now Gemini Notebook, is a grounded research partner: it works from user-provided sources only. That is closer to a constrained research helper than to Grok Bot. It is still not a documented autonomous-agent SKU for schools. A notebook that will not leave the sources you gave it is a pedagogical constraint, not a VM. The temptation, already visible in campus conversation, is to call Gemini Notebook an agent because it "does research." Grounded synthesis over a folder the user uploaded is a useful first-row or first-row-adjacent tool. It does not clone a repository, it does not keep working at 2 a.m., it does not sign into the bursar's system, and Google does not document it as an autonomous-agent SKU for schools. Closer to a constrained research helper than to Grok Bot is still not the last row. The helper remains a conversation with sources. The last row is a job with a computer.

ChatGPT, consumer and Edu, remains the default education "chat window." OpenAI itself now distinguishes regular ChatGPT from Workspace Agents. The distinction is the company's, dated 23 April 2026 in an OpenAI Academy higher-education playbook, and this chapter will honor it. Regular ChatGPT is the thread. Workspace Agents are a different class, treated below, and even there the playbook is campus operations, off by default, with a human review before anything is sent.⁸¹

What this row cannot do, even when it looks busy, follows from the unit of work. A thread can

⁷⁹xAI, Grok overview, <https://docs.x.ai/grok/overview>. Chat, Imagine, voice, file upload, connectors on grok.com and iOS/Android. Distinct, in xAI's own line, from Grok Bot.

⁸⁰Google Workspace Help, "Quickstart Guide to Gemini and Gemini Notebook for Education," <https://support.google.com/a/answer/16350447>. Free in all Google Workspace for Education editions; all ages; Core service; data not used to train models and COPPA/FERPA support — **vendor/admin claims**. Youth onboarding endorsed, in Google's materials, by ConnectSafely and FOSI. "World's leading model for learning" is **Google marketing** for LearnLM / "2.5 Pro" learning-tuned models. NotebookLM (Gemini Notebook) is grounded in user-provided sources only; not a documented autonomous-agent SKU for schools. Google, "Google and Khan Academy build classroom AI tools," <https://blog.google/products-and-platforms/products/education/khan-academy-back-to-school/>.

⁸¹OpenAI Academy, "Understanding Workspace Agents in Higher Education" (Apr. 23, 2026), <https://academy.openai.com/public/clubs/higher-education-05x4z/blogs/understanding-workspace-agents-higher-education>. OpenAI distinguishes regular ChatGPT from Workspace Agents. Workspace Agents: research preview for eligible ChatGPT Edu workspaces; **off by default**; admin-gated; humans review before anything is sent.

be long. A thread can have tools. A thread can be given a system prompt that says “you are a tutor” or “you are a research assistant.” None of that gives the system a computer that is not the user’s laptop, an unattended runtime, the ability to sign into Gmail and file the ticket, a parallel specialist, or a skill document that fires on a schedule. The simulation is real. The execution is still the human’s. Teachers who have spent 2023 through 2026 watching students paste homework into a box already know the phenomenology: the box waits. It is very good at waiting. It is not a colleague who continues when the period ends. The same phenomenology, on the teacher side, is a planning period spent in a thread: generate a quiz, paste it into the LMS, generate a parent email, paste it into the client, generate a lesson outline, paste it into slides. Cowork’s claim, when we reach it, is that the pasting stops. ChatGPT for Teachers, as launched, is still a workspace for educators — a first-row product with an education term — not a student agent and not, in the CNBC launch report, a scheduled folder-scale job. The education term does not move the row. Google Workspace for Education, wrapping Gemini, is the same fact in another vendor’s clothing: a Core-service chat and notebook, not a documented last-row SKU.

The Department of Education’s 2026 priority, already cited, is a scoring tool for discretionary grants, not a classroom mandate, and not a definition of this product class. It funds literacy and ethical use for tutoring, personalization, UDL, and administrative efficiency. A reader looking in 91 FR 18774 for authorization of student-facing agents will not find it, because the document does not contain it. UNESCO, likewise, does not endorse the last row for students. The federal and international framing of 2026 is still, at the level of official prose, a framing of *AI* as a mass noun with human-centred caveats. This chapter’s job is to refuse the mass noun even when the official documents use it.

Tutors are not smarter chat. They are designed refusal. That is a different row, and schools bought it on purpose.

Tutor apps: designed pedagogical wrappers

A tutor app is a designed pedagogical wrapper around a model. The product’s value is what it *refuses* to do (give the answer), what curriculum it is grounded in, and what the teacher can see. That is a different commercial and cognitive object from a conversation window that will, by default, complete the assignment. The wrapper is the treatment. The model class is shared. Design is not.

Three named student-facing products dominate the 2026 district conversation. None of them is an agent. Do not list Khanmigo, MagicSchool, SchoolAI, Gemini Education, or ChatGPT for Teachers as agents. MagicSchool’s “Raina” and SchoolAI’s “Dot” are constrained classroom assistants, not Grok Bot-class teammates.⁸²

Khanmigo, from Khan Academy, is described by Khan Academy as an “AI-powered, 24/7 Socratic tutor for students,” incorporated into Khan Academy’s content library: math, humanities, coding, social studies. District partnership with Khanmigo for students is listed at \$15 per student; United States districts only; Clever or ClassLink rostering; minimum 250 licenses; grades 3–12. Teacher

⁸²Landscape memo §4.4. MagicSchool’s “Raina” and SchoolAI’s “Dot” are constrained classroom assistants, not Grok Bot-class teammates. Do not list Khanmigo, MagicSchool, SchoolAI, Gemini Education, or ChatGPT for Teachers as agents. They are tutor apps or teacher-productivity generators with education guardrails.

tools are included even without the student tutor. Khan Academy states “over 400 partners across the US” for Districts. That is a vendor claim. The documented partner named on the same partner page is School City of Hobart, Indiana, about 4,000 students, featured on CBS *60 Minutes*.⁸³ Kristen DiCerbo, Khan Academy’s chief learning officer, in an American Enterprise Institute interview, put partner-district users at about 68,000 in 2023–24 and more than 700,000 in 2024–25, and districts at 45 to 380-plus. Vendor claims. Common Sense Media rates Khanmigo **low risk** as of 8 August 2024. Google announced Gemini-powered interactive diagrams for Khanmigo for the 2026 school year.⁸⁴ The Socratic refusal is the product. A system that will not give the answer is not a system that will file the answer, send it, and go on to the next assignment while the student is at lunch. Later chapters will ask whether students *use* the refusal, and whether use produces unaided skill. This chapter only places the object: a tutor app, in a curriculum, with rostering, at a district price, with a vendor partner count.

SchoolAI Spaces are teacher-designed student experiences. Company research with Jordan School District, Utah, claims a 28 percent critical-thinking gain over two years, from 13,882 conversations and 82 teachers. The study is ESSA Level III via an Instructure review. It is vendor-commissioned, not an independent randomized controlled trial. Scoring used a Bloom rubric on chat logs, not standardized tests. The study’s own punchline is useful and will be used: gains tracked teacher-designed Spaces (six or more), not mere access.⁸⁵ Midland Public Schools, Michigan, selected SchoolAI for optional use in grades 6–12; local journalism dated 19 August 2026 records the board selection. Selection is not outcomes.⁸⁶ A district that buys optional Spaces has bought a teacher-designed wrapper. It has not bought an unattended agent, and the company’s research, even taken at vendor strength, does not claim that it has.

MagicStudent, the student-facing layer of MagicSchool, is documented in MagicSchool’s Midway Independent School District, Texas, case study. Midway ISD enrolls 8,700 students. The case study is vendor work, not independent evaluation. Midway’s documented path, still on the vendor site, is a districtwide pilot of MagicSchool, Diffit, Gemini, and PowerBuddy; MagicSchool selected; 630-plus teachers trained in summer 2024; official pilot fall 2024; adoption spring 2025;

⁸³Khan Academy, “Becoming a Khan Academy Districts Partner,” <https://blog.khanacademy.org/becoming-a-khan-academy-districts-partner/>. Vendor description: “AI-powered, 24/7 Socratic tutor for students.” District partnership listed at \$15 per student; U.S. districts only; Clever or ClassLink; minimum 250 licenses; grades 3–12; teacher tools included even without the student tutor. “Over 400 partners across the US” is a **vendor claim**. Named partner on the same page: School City of Hobart, Indiana (~4,000 students), featured on CBS *60 Minutes*. *60 Minutes* safety-flag claims are vendor claims, not an independent audit.

⁸⁴Kristen DiCerbo (Khan Academy chief learning officer), interview, American Enterprise Institute, “Can an AI-Powered Tutor Produce Meaningful Results?,” <https://www.aei.org/commentary/can-an-ai-powered-tutor-produce-meaningful-results/>. Partner-district users about 68,000 (2023–24) to more than 700,000 (2024–25); districts 45 to 380-plus — **vendor claims**. Common Sense Media / Youth AI Safety Institute rates Khanmigo **low risk** (Aug. 8, 2024), <https://institute.commonsensemedia.org/risk-assessments>. Google announcement of Gemini-powered interactive diagrams for Khanmigo for the 2026 school year: Google education blog, note 8 above.

⁸⁵SchoolAI, research page, <https://schoolai.com/research/schoolai-makes-students-think>. Company research with Jordan School District (Utah): 28 percent critical-thinking gain over two years; 13,882 conversations; 82 teachers. ESSA Level III via Instructure review — **vendor-commissioned, not an independent RCT**. Bloom-rubric scoring of chat logs, not standardized tests. Gains tracked teacher-designed Spaces (six or more), not mere access.

⁸⁶Dan Chalk, “Midland Public Schools selects SchoolAI,” *OurMidland*, Aug. 19, 2026, <https://www.ourmidland.com/news/article/midland-schools-school-ai-22393016.php>. Board selection for optional use in grades 6–12. Documented selection, not outcomes.

a 23-page data-privacy agreement signed with no redlines.⁸⁷ MagicSchool’s district trust page lists iKeepSafe FERPA and COPPA certifications, SOC 2 Type II, Student Data Privacy Consortium membership, and ESSA Level III (“Promising Evidence”) via Instructure EdTech Collective: a positive association between teacher use and student ELA. That is a vendor-funded evidence tier, not a large independent RCT. The company is explicit that third-party AI providers may not train on student data.⁸⁸ Diffit, which sat in Midway’s comparison pilot and which Anthropic later lists as a Claude-for-Teachers connector, is a teacher tool to create and adapt instructional materials — leveled passages, vocabulary, questions, lesson kits — on a schools plan with site license, Google and Microsoft export, standards and depth-of-knowledge tagging, and an admin dashboard. Diffit’s vendor claim is FERPA and COPPA compliance, and that it does not collect student data.⁸⁹ These products are prompted generators with education guardrails. They are not persistent autonomous agents.

The row can be summarized without charity and without sneer. Schools that wanted a student-facing generative product in 2024–26 bought wrappers: Socratic refusal, teacher-designed spaces, rostering through Clever or ClassLink, a DPA, a logged conversation the teacher can see. The wrapper is why Common Sense can rate Khanmigo low risk and Grok, the chat product, unacceptable. The wrapper is why a board in Midland can select SchoolAI as optional in grades 6–12 without having selected a cloud computer that signs into apps. The causal record on whether those wrappers produce learning is Chapter 3. The classification is this chapter’s. Tutor is not chat with better manners. Tutor is not agent with a school logo.

Adaptive systems are older than ChatGPT and are not agents either. The ITS evidence base will be laundered as generative-AI evidence unless this next row is kept separate.

Adaptive systems: the learner-model that does not browse

Adaptive learning platforms and intelligent tutoring systems maintain a knowledge map — what the student can do — and serve the next item. They do not browse the open web, write a paper overnight, or sign into Gmail. They predate ChatGPT. They are the class for which education already has a research literature, and that literature is about item-level practice under a learner model, not about unattended multi-step agency.

Mary Burns, writing for Brookings on 27 January 2026, restates the ITS evidence base: Nesbit and colleagues (2014), Kulik and Fletcher (2016), Major and colleagues (2021). Intelligent tutor-

⁸⁷MagicSchool, Midway ISD case study, <https://www.magicschool.ai/case-studies/midway-isd>. Midway Independent School District, Texas, 8,700 students. Districtwide pilot of MagicSchool, Diffit, Gemini, and PowerBuddy; MagicSchool selected; 630-plus teachers trained summer 2024; official pilot fall 2024; adoption spring 2025; 23-page DPA signed with no redlines. **Vendor case study**, not independent evaluation.

⁸⁸MagicSchool trust center, <https://www.magicschool.ai/privacy/trust>. iKeepSafe FERPA and COPPA certifications; SOC 2 Type II; SDPC member; ESSA Level III (“Promising Evidence”) via Instructure EdTech Collective — positive association between teacher use and student ELA; **vendor-funded evidence tier, not a large independent RCT**. Third-party AI providers may not train on student data (vendor statement).

⁸⁹Diffit, plans, <https://web.diffit.me/plans>; app.diffit.me. Teacher tool to create/adapt instructional materials. Schools plan: site license, Google/Microsoft export, standards/DOK, admin dashboard. Vendor: FERPA and COPPA compliant; **does not collect student data**. Anthropic lists Diffit as a Claude-for-Teachers connector (note 28). Midway ISD included Diffit in its comparison pilot (note 15).

ing systems can match human tutoring under the right conditions. Burns’s argument, which this chapter accepts as the right framing of the class, is that generative AI is now replacing some *rule-based dialogue and content-generation* components of ITS, not replacing the learner-model idea. A model that can write a hint in fluent English is a new component inside an old architecture. The architecture is still: know what the student can do, serve the next item, keep the student inside a knowledge map. Fluency at the hint layer does not confer a browser, a cron job, or a cloud computer. If generative dialogue makes an ITS more pleasant, the evaluation question remains the ITS question — does the learner model, plus this dialogue, produce unaided skill — not the agent question. Mixing the questions is how a 2016 tutoring-systems review gets cited as if it had measured Grok Bot.⁹⁰ Wang, Wang, and colleagues (2024) meta-analyse AI-enabled *adaptive* systems from 2010 to 2022 at $g = 0.70$ versus non-adaptive. That is not a generative-AI result. The window ends in 2022. Treating a 2010–2022 adaptive-systems meta-analysis as evidence for Khanmigo, or for Hermes Agent, is the laundering this section exists to stop.⁹¹

The Institute of Education Sciences What Works Clearinghouse has older ITS reports, including an intervention report on Cognitive Tutor Algebra I. As of 30 August 2026 there is no WWC intervention report specific to generative AI tutors.⁹² That absence is a finding, not a gap to be filled with vendor ESSA-3 studies. Kulik and Fletcher, Cognitive Tutor, the Wang adaptive meta-analysis: those are the documents a serious book puts under the third row. They do not migrate, by rhetorical courtesy, into a claim about Grok Bot.

Named platforms commonly cited in 2026 secondary roundups include ALEKS (McGraw Hill; Knowledge Space Theory), Carnegie Learning MATHia, Curriculum Associates i-Ready, and Discovery Education / DreamBox. Official product pages were not fetched. Feature lists and user counts for these products are **UNVERIFIED** in this chapter. The names identify the class. They do not authorize a tour of dashboards the research did not open.⁹³

What the third row shares with the second is constraint. What it does not share with the fourth is agency. An ITS that knows which item to serve next is a powerful pedagogical machine. It is not a colleague. It does not have a cloud computer. It does not watch the user do a job once and then run the job on a schedule. It does not open a pull request. A superintendent who cites Kulik and Fletcher in a board memo about “AI tutoring” has cited a literature about the third row. If the purchase is Khanmigo, the literature is being asked to travel. If the purchase is imagined to be a student-facing Hermes Agent, the literature is being asked to travel further than it can. Keep the

⁹⁰Mary Burns, “What the research shows about generative AI in tutoring,” Brookings, Jan. 27, 2026, <https://www.brookings.edu/articles/what-the-research-shows-about-generative-ai-in-tutoring/>. Restates Nesbit et al. (2014), Kulik and Fletcher (2016), and Major et al. (2021): ITS can match human tutoring under the right conditions. Generative AI is replacing some rule-based dialogue and content-generation components of ITS, not the learner-model idea. James A. Kulik and J. D. Fletcher, “Effectiveness of Intelligent Tutoring Systems,” *Review of Educational Research* 86, no. 1 (2016): 42–78.

⁹¹S. Wang et al., “The Efficacy of Artificial Intelligence-Enabled Adaptive Learning Systems from 2010 to 2022,” *Journal of Educational Computing Research* (2024), <https://doi.org/10.1177/07356331241240459>. Meta-analysis of AI-enabled *adaptive* systems, 2010–2022: $g = 0.70$ versus non-adaptive. **Not a generative-AI result.** Window ends 2022.

⁹²Institute of Education Sciences, What Works Clearinghouse, Cognitive Tutor intervention report (2016), <https://ies.ed.gov/ncee/wwc/Intervention/818>. Older ITS reports exist. **No WWC intervention report specific to generative AI tutors** as of Aug. 30, 2026.

⁹³See note 1. Official ALEKS, MATHia, i-Ready, and DreamBox pages not fetched this pass.

rows. The evidence in Chapter 3 will be almost entirely about chat and tutors. The ITS row is here so that evidence is not padded with a pre-generative literature that measured something else.

Now the last row — the book’s named objects — defined from primary product docs, not from education marketing.

Agents: goal, tools, unattended runtime, memory

The last row is the only row this book is, in the end, about. It is also the row that almost no school has bought. Define it from the vendors’ own sentences, then look at which of those sentences has an education SKU attached.

Hermes Agent, on its product line, says: “Not a chatbot. Not a copilot. An agent that lives on your machine.” The sentence is vendor copy. It is useful here only to distinguish classes, never as an education endorsement.⁹⁴ The product is an open-source (MIT), self-hosted autonomous agent released by Nous Research in February 2026; the responsible-use stream of the research this book stands on dates the launch to 25 February 2026. Install via curl; Linux, macOS, WSL2; optional Desktop installer. Model-agnostic: Nous Portal, OpenRouter, OpenAI-compatible, local vLLM. Documented capabilities, from product docs, not from education evaluations: persistent memory across sessions, including Honcho dialectic user modeling and FTS5 session search; autonomous skill creation after complex tasks, with skills compatible with agentskills.io; cron and scheduled automations with delivery to messaging platforms; parallel isolated sub-agents; RPC via `execute_code`; terminal backends including local, Docker, SSH, Singularity, Modal, and Daytona; browser automation, web search and extract, vision, image generation, text-to-speech; a messaging gateway listing Telegram, Discord, Slack, WhatsApp, Signal, Teams, and a long tail of others, twenty-plus platforms in the docs; Bot Mode — named Bots with their own model, memory, skills, routines; a research and MLOps loop, batch trajectory generation, ShareGPT export, Atropos RL integration, “built by model trainers”; a bundled research skill for end-to-end ML paper writing, literature review through experiments through LaTeX through submission, using arXiv, Semantic Scholar, `delegate_task`, `cronjob`. Privacy posture, vendor: no telemetry claimed; data in `~/.hermes/`; MIT license; container hardening options.⁹⁵ That is a last-row object. It is not a classroom product. School and university use of Hermes Agent: none found in official Nous materials, university newsrooms, or policy documents fetched for this research. A GitHub project, AlijonovMukhammaddiyor/hermes-university, describes a personal “Hermes-based learning system” on Telegram. That is a third-party hobby project. It is not institutional use, and it will not be cited as such.⁹⁶

⁹⁴Nous Research, Hermes Agent product line, <https://hermes-agent.org/>: “Not a chatbot. Not a copilot. An agent that lives on your machine.” **Vendor copy.** Used here only to distinguish product classes, never as an education endorsement.

⁹⁵Hermes Agent: <https://hermes-agent.org/>; official docs, <https://hermes-agent.nousresearch.com/docs/>; GitHub, <https://github.com/nousresearch/hermes-agent/>. Open-source, MIT license; self-hosted; vendor site date February 2026; launch dated 25 February 2026 in the responsible-use stream of this book’s research. Model-agnostic. Capabilities listed in the text are from product documentation, not from education evaluations. Privacy posture (no telemetry; data in `~/.hermes/`) is a **vendor claim**. Bundled research-paper-writing skill: <https://hermes-agent.nousresearch.com/docs/user-guide/skills/bundled/research/research-research-paper-writing>.

⁹⁶No official Nous materials, university newsrooms, or policy documents fetched document school or university use of Hermes Agent. GitHub AlijonovMukhammaddiyor/hermes-university is a third-party hobby project, not

Grok Bot, the autonomous teammate, launched in early beta on 11 August 2026 as a joint xAI and Cursor product. Sign-in is via a Cursor account. On 26 August 2026 it expanded to SuperGrok, SuperGrok Plus, SuperGrok Heavy, Cursor Pro, Pro+, and Ultra, and Cursor Teams Standard and Premium, on a usage pool separate from Grok and Cursor plan quotas. Enterprise sat on a waitlist. What it can do that a chat window cannot, in vendor docs: a persistent cloud computer shared by all Bots on an account, with files, browser sessions, logins; sign into apps and websites, including those without a clean API or MCP; work 24/7 in parallel; group-chat coordination among Bots; a “chief of staff” pattern; routines that watch the user do a job once, then run it on a schedule; messaging like a teammate on desktop and iOS; continuation when the laptop is closed. Jobs illustrated in launch materials: sales outbound, CRM hygiene, inbox, expenses, bug reproduction, website deploy, meeting stand-in. None of these examples is educational. Product pages, the launch post, and the expansion post do not mention schools, universities, students, or teachers. Education positioning of Grok Bot: none found. Grok Bot “requires cloud data storage.” The FAQ at docs.x.ai/grok-bot/faq, re-fetched after the landscape timeout, confirms that requirement, states that Legacy Privacy Mode is not supported, and says training opt-out follows Cursor account settings. It does not mention schools, students, or teachers. Plan-by-plan eligibility remains as in the 11 August and 26 August 2026 news posts.⁹⁷

Cursor Cloud Agents, formerly Background Agents, are autonomous coding agents on isolated Cursor-managed Linux VMs with a full development environment: clone, dependencies, secrets, terminal, browser and desktop, MCP. They can be kicked off from Cursor desktop, cursor.com/agents, iOS, Slack @cursor, GitHub or Bitbucket comments, Linear, or the API. They produce merge-ready pull requests plus screenshots, videos, and logs, with optional remote-desktop takeover. CNBC, on 24 February 2026, reported Cursor saying roughly 35 percent of Cursor’s own pull requests are now generated by agents on VMs. That is a company statistic, not an education statistic.⁹⁸ Cursor docs do not market to schools. Search returned “Cursor 101” listings on Texas Southern University and University of Illinois Chicago career-hub class catalogs, with Ray Villalobos as instructor. Those course pages were not fully fetched. Treat as **UNVERIFIED** whether they are credit-bearing curriculum or third-party catalog syndication.⁹⁹

institutional adoption. Education-relevant mapping onto graduate research workflows is inference, not a claim of use. Mapping ≠ adoption.

⁹⁷xAI, “Introducing Grok Bot,” Aug. 11, 2026, <https://x.ai/news/introducing-grok-bot>; “Grok Bot more plans,” Aug. 26, 2026, <https://x.ai/news/grok-bot-more-plans>; product, <https://x.ai/bot>. Early beta 11 August 2026; joint xAI / Cursor; Cursor authentication; expanded 26 August 2026 to SuperGrok, SuperGrok Plus, SuperGrok Heavy, Cursor Pro / Pro+ / Ultra, Cursor Teams Standard and Premium; separate usage pool; Enterprise on a waitlist. Vendor: own cloud computer; sign into apps; 24/7 parallel work; routines; jobs illustrated — sales, CRM, inbox, expenses, bug reproduction, website deploy, meeting stand-in — **none educational**. Product pages do not mention schools, students, or teachers. “Requires cloud data storage.” Grok Bot FAQ at <https://docs.x.ai/grok-bot/faq>, re-fetched after the landscape timeout: Legacy Privacy Mode is not supported; training opt-out follows Cursor account and privacy settings; no schools, students, or teachers mentioned. Plan list in the 11 and 26 August 2026 news posts remains the eligibility source used in the body.

⁹⁸Cursor, Cloud Agents, <https://cursor.com/docs/cloud-agent>. Isolated Cursor-managed Linux VMs; clone, deps, secrets, terminal, browser/desktop, MCP; kickoff from desktop, cursor.com/agents, iOS, Slack @cursor, GitHub/Bitbucket comments, Linear, API; merge-ready PRs. CNBC, “Cursor announces major update,” Feb. 24, 2026, <https://www.cnbc.com/2026/02/24/cursor-announces-major-update-as-ai-coding-agent-battle-heats-up.html>. Cursor said roughly 35 percent of *its own* PRs are generated by agents on VMs — **company statistic, not an education statistic**.

⁹⁹Search returned “Cursor 101” listings on Texas Southern University and University of Illinois Chicago career-hub

Claude Cowork, via Claude for Teachers, is the closest “agent in the schoolhouse” story, and it is a *teacher-side* story. Anthropic launched Claude for Teachers on 14 July 2026: verified U.S. K–12 educators, free premium Claude, a Learning Commons connector covering standards in all fifty states plus Illustrative Mathematics and OpenSciEd, teaching Skills co-developed with Learning Commons. The offering includes Claude Code and Cowork — scheduled and folder-scale autonomous work. Educator-only. Claude 18-plus policy. Training off for verified teacher accounts. FERPA-aligned data-privacy agreement. Sign-up by 30 June 2027 for a year of free access. On 28 August 2026, a dedicated schools-and-districts offering: SSO, role-based access control, district terms. Anthropic names Prospect Schools, Brooklyn, as an early-feedback school, and announces an impact pilot with Detroit Public Schools Community District, “working closely with teachers to study the impact on educator wellbeing and practice.” Announced, not yet published as results. A Gates Foundation partnership is mentioned; Playlab is to support a national network of lab schools. AFT President Randi Weingarten is quoted on a “Gold Standard” for K–12 safety and privacy. Connectors named by Anthropic, as ecosystem, not as “Anthropic built these”: ASSISTments, Brisk Teaching, Canva Education, Coteach, Diffit, Eedi, MagicSchool, Snorkl, TeachFX. Cowork examples, teacher-side: analyze a folder of de-identified diagnostics or exit tickets; run every school day at 4 p.m. to adapt tomorrow’s plan; scheduled parent newsletters. Not student-facing agents.¹⁰⁰ Cowork’s own sentence is the class in vendor English: “Set a task up once and Claude handles it on a schedule.”

OpenAI Workspace Agents, in a 23 April 2026 Academy playbook, are a reusable workflow with trigger, process, tools, and guardrails, able to run on a schedule. They are in research preview for eligible ChatGPT Edu workspaces, off by default, admin-gated. The playbook is explicit that agents are for repeatable campus operations, not student homework: course-change triage, faculty AI-support intake, weekly student-services summaries, grant-readiness checklists, campus communications, IT software-request routing, program-review packets, faculty onboarding, accreditation evidence collection. Humans review before anything is sent. No named campus production case was fetched.¹⁰¹ Codex in higher education, 9 April 2026, is a coding-agent surface: Codex App for supervised pilots, Codex Cloud for asynchronous repository work with a GitHub App and least-privilege tokens, Codex CLI. Aimed at developers and at non-technical campus staff via Edu lo-

class catalogs (Ray Villalobos, instructor). Course pages were **not fully fetched**. Treat as **UNVERIFIED** whether they are credit-bearing curriculum or third-party catalog syndication. Cursor docs do not market to schools.

¹⁰⁰Anthropic, “Introducing Claude for Teachers,” July 14, 2026 (schools-and-districts offering Aug. 28, 2026), <https://www.anthropic.com/news/claude-for-teachers>; <https://claude.com/solutions/teachers>. Verified U.S. K–12 educators; free premium Claude; Learning Commons connector (standards in all 50 states; Illustrative Mathematics, OpenSciEd); teaching Skills co-developed with Learning Commons; includes Claude Code and Cowork. Educator-only; Claude 18+ policy; training off for verified teacher accounts; FERPA-aligned DPA; sign-up by June 30, 2027, for a year of free access. 28 August 2026: SSO, RBAC, district terms. Named early-feedback school: Prospect Schools, Brooklyn. Announced impact pilot: Detroit Public Schools Community District — **announced, not published as results**. Gates Foundation partnership mentioned; Playlab to support a national network of lab schools. AFT President Randi Weingarten quoted on a “Gold Standard” for K–12 safety/privacy. Connectors named by Anthropic as ecosystem, not as Anthropic-built: ASSISTments, Brisk Teaching, Canva Education, Coteach, Diffit, Eedi, MagicSchool, Snorkl, TeachFX. Cowork: “Set a task up once and Claude handles it on a schedule.” Examples are teacher-side (de-identified diagnostics/exit tickets; 4 p.m. scheduled planning; parent newsletters), not student-facing agents.

¹⁰¹OpenAI Academy, Apr. 23, 2026 (note 9). Playbook examples are campus operations, not student homework: course-change triage, faculty AI-support intake, weekly student-services summaries, grant-readiness checklists, campus communications, IT software-request routing, program-review packets, faculty onboarding, accreditation evidence collection. **No named campus production case fetched.**

gin, with a security emphasis on repo scoping, sandbox, and approval policy.¹⁰² Clemson News describes ChatGPT Edu plus GPT-5.5 plus Codex plus locally hosted open-weight models for researchers. That is documented *access*, not a published learning-outcomes study.¹⁰³

Five systems. One operational definition, claimed in their own docs. Two education SKUs in the set, both teacher- or campus-facing: Claude Cowork inside Claude for Teachers, and OpenAI Workspace Agents inside ChatGPT Edu, off by default, with a human in the send path. Hermes Agent and Grok Bot, the objects this book’s named reader already knows from the open internet, have no education SKU in anything fetched, no named district DPA, no classroom standard operating procedure. Cursor Cloud Agents do not market to schools. The last row exists. The last row in a third-grade classroom, as an unconstrained student agent, does not — not in the documents.

Two collapses will make this book obsolete if they survive copy-edit. They are the next two sections, at full length, because the popular story will keep making them.

Hermes 4 is not Hermes Agent

Nous’s own documentation separates a model family from a runtime. Using the family as a proxy for the agent, or the agent as a proxy for the family, is a factual error. It is also the error a book about *autonomous* AI is most likely to make, because the name is shared and the community that cares about the name already treats “Hermes” as a single object. It is not.

Hermes 4 is an open-weight hybrid-reasoning large-language-model family from Nous Research, not an agent runtime. The Hugging Face model card for Hermes 4 14B, on a Qwen 3 14B base, describes hybrid-mode reasoning with optional `<think>` traces; function calling via `<tool_call>` tags; schema and JSON adherence; a post-training corpus described as about five million samples and about sixty billion tokens, against Hermes 3’s one million samples and 1.2 billion tokens. Weights are public. The technical report is arXiv 2508.18255, Teknium and colleagues, 25 August 2025. HTML conversion of that report failed on fetch. Abstract and metadata are confirmed via the Hugging Face citation block. Quantitative benchmarks in the technical report remain **UNVERIFIED** until the PDF is read. They will not be recited from memory or from secondary roundups.¹⁰⁴

Education-relevant capabilities of the *model*, from the model card, not from any school study: math, code, STEM, logic, long-form writing, tool and function calling, structured output. These are model

¹⁰²OpenAI Academy, “Deploying Codex in Higher Education,” Apr. 9, 2026, <https://academy.openai.com/public/clubs/higher-education-05x4z/blogs/deploying-codex-in-higher-education-2026-04-09>. Codex App (supervised pilots); Codex Cloud (async repo work, GitHub App, least-privilege tokens); Codex CLI. Security emphasis: repo scoping, sandbox, approval policy. Playbook, not an outcome study.

¹⁰³Clemson News, “GPT-5.5, Codex and Clemson-hosted models give researchers new options for AI-assisted work,” <https://news.clemson.edu/gpt-5-5-codex-and-clemson-hosted-models-give-researchers-new-options-for-ai-assisted-work/>. Documented **access** (ChatGPT Edu + GPT-5.5 + Codex + locally hosted open-weight models), not a published learning-outcomes study.

¹⁰⁴Nous Research, Hermes-4-14B-FP8 model card, <https://huggingface.co/NousResearch/Hermes-4-14B-FP8>. Qwen 3 14B base; hybrid-mode reasoning; `<think>` traces; `<tool_call>` tags; post-training corpus described as ~5 million samples / ~60 billion tokens versus Hermes 3’s 1 million / 1.2 billion. Technical report: Teknium et al., arXiv:2508.18255 (Aug. 25, 2025), <https://arxiv.org/abs/2508.18255>. HTML conversion failed on fetch; abstract/metadata confirmed via Hugging Face citation block. **Quantitative benchmarks UNVERIFIED until the PDF is read.** Model skills (math, code, STEM, tool/function calling) are model skills, not a classroom product.

skills. They are not a classroom product. A weight file that can do math is not Khanmigo. A weight file that can emit a `<tool_call>` is not Hermes Agent. Local inference of an open-weight model, including Nous’s own Hermes weights, can sit *under* an agent runtime because the runtime is model-agnostic. That is a hosting fact, not an identity. A school that downloaded Hermes 4 14B has downloaded a model. It has not deployed Hermes Agent. A hobbyist who runs Hermes 4 in a chat UI has a chat UI.

The official warning is Nous’s, not this book’s. Nous Portal integration docs state that Hermes 4, at 70B and 405B, is “not recommended for use inside Hermes Agent” because it is “tuned for chat and reasoning, not the rapid-fire tool-calling loop the agent relies on.” Recommended agent models in that same document include Claude Sonnet 4.6, GPT-5.5 Pro, Gemini 3 Pro, and DeepSeek V4 Pro.¹⁰⁵ Read the sentence again. The laboratory that named both objects tells the user not to put the model family inside the agent, because the family is tuned for the first row of the table — chat and reasoning — and the agent depends on a tool-calling loop the family is not tuned for. If Nous does not recommend Hermes 4 inside Hermes Agent, a book that treats “Hermes” as one education-relevant system has not read the vendor’s own integration page.

Hermes Agent, already described, is the runtime: MIT-licensed, self-hosted, model-agnostic, February 2026. It can, in principle, call Claude or GPT or Gemini or DeepSeek or a local vLLM endpoint. The agent is not a Nous-model exclusive, and the Nous model is not an agent. Persistent memory, cron, sub-agents, browser, terminals, messaging gateways, Bot Mode, the research-paper-writing skill: those are runtime features. They do not appear on a Hugging Face model card, because a model card is not a computer.

No official Nous materials, university newsrooms, or policy documents fetched for this research document school or university *use* of Hermes Agent. The GitHub hobby project already named is not institutional adoption. Education-relevant inference, which is not a claim of use: the research skills, persistent memory, self-hosting, and cron jobs map naturally onto graduate research workflows and faculty knowledge work — literature monitoring, experiment babysitting, paper drafting — the same class of work OpenAI and Anthropic are pitching to researchers. Mapping is not adoption. A capability that would help a doctoral student write a methods section is not a methods section that has been written, and it is not a graduate school that has approved the writing. The inference is licensed here because the book’s named reader will make it anyway. The inference is labelled because a later chapter on who hosts a student agent cannot be allowed to treat a mapping as a deployment.¹⁰⁶

Two further cautions, both about what this correction is *not*. It is not a claim that Hermes 4 is weak. The model card’s list of skills is a vendor list; the unread PDF may or may not support it; this chapter does not need the leaderboard. It is not a claim that Hermes Agent is unsafe, or safe, as a school product. The runtime is not a COPPA-certified, KCSIE-aligned school product. Computer-use and web tools re-open prompt injection and exfiltration. There is no designated-

¹⁰⁵Nous Portal integration docs, <https://hermes-agent.nousresearch.com/docs/integrations/nous-portal>. Hermes 4 (70B / 405B) is “**not recommended for use inside Hermes Agent**” because it is “tuned for chat and reasoning, not the rapid-fire tool-calling loop the agent relies on.” Recommended agent models in that document include Claude Sonnet 4.6, GPT-5.5 Pro, Gemini 3 Pro, DeepSeek V4 Pro. Hermes Agent is model-agnostic.

¹⁰⁶See note 24. The inference that research skills, persistent memory, self-hosting, and cron map onto graduate research workflows is licensed as inference and labelled as not a claim of use.

safeguarding-lead alerting path unless the school runs the server. Equity requires every family to own a capable machine if the runtime is local — a new gap, not a closed one. Those are hosting and safety facts for a later chapter. This chapter’s fact is narrower and prior: the name is two objects. Write “Hermes 4” when the subject is weights. Write “Hermes Agent” when the subject is a computer that keeps working. Do not write “Hermes” and let the reader pick.

The same error class, with three objects instead of two, is the Grok file.

Grok the assistant is not Grok Bot is not El Salvador

Three different products are often collapsed in conversation. The collapse is convenient for a headline and fatal for a school board. (a) Grok, the chat assistant. (b) Grok Bot, the autonomous teammate. (c) Grok-as-tutor, announced for El Salvador. They do not share a SKU, a date, a host, or a documented outcome. Treat them as one and the book has already endorsed a national tutoring program that has not been shown to tutor, and an autonomous teammate that has not been shown to enter a school.

Grok the assistant is the first-row object already named: chat-centric, grok.com, apps, X, connectors, files, Imagine, voice. That is the ChatGPT-class surface. Common Sense Media rates Grok **unacceptable** as of 22 January 2026. That rating is of the *chat product*, not of Grok Bot. Do not migrate a companion-and-chat risk assessment onto a cloud teammate that did not exist when the rating was issued, and do not treat the rating as a finding about El Salvador.¹⁰⁷

Grok Bot is the last-row object launched in early beta on 11 August 2026, expanded 26 August 2026, joint xAI and Cursor, Cursor-authenticated, with its own cloud computer, 24/7 parallel work, sign-in to apps, and routines. Product pages do not mention schools, students, or teachers. Jobs illustrated: sales, CRM, inbox, expenses, bug reproduction, website deploy, meeting stand-in. None educational. No education SKU. The FAQ, re-fetched after the landscape timeout, confirms cloud storage, no Legacy Privacy Mode, and Cursor-tied training opt-out; it still contains no education SKU. US K–12 and higher-education Grok Bot pilots: none found.¹⁰⁸ A book that needs Grok Bot to be a school story will not find the story in xAI’s own launch materials. The launch is a culture event for this book’s named reader. It is not a procurement event for a district.

xAI’s actual education play is El Salvador, and it is Grok the assistant, not Grok Bot. On 11 December 2025, xAI announced a partnership with the Government of El Salvador to “deploy Grok across more than 5,000 public schools” over two years, “personalized learning to over one million students,” framed as adaptive, curriculum-aligned tutoring. CoinDesk, same day, restates the announcement. Restatement is not independent verification. No outcome evaluation, independent audit, or classroom study was found. This is an announced national program, not a documented completed deployment as of 30 August 2026. It is Grok-as-tutor, not Grok Bot-as-agent.¹⁰⁹ US

¹⁰⁷Common Sense Media / Youth AI Safety Institute rates Grok **unacceptable** (Jan. 22, 2026), <https://institute.commonsensemedia.org/risk-assessments>. That rating is of the **chat product**, not Grok Bot (which had not launched). Do not migrate it onto Grok Bot or treat it as a finding about El Salvador.

¹⁰⁸See note 25. Education positioning of Grok Bot: none found. U.S. K–12 and higher-education Grok or Grok Bot pilots: **none found**. FAQ now fetched; still no education SKU.

¹⁰⁹xAI, El Salvador partnership, Dec. 11, 2025, <https://x.ai/news/el-salvador-partnership>. “Deploy Grok across more than 5,000 public schools” over two years; “personalized learning to over one million students”; framed as adaptive,

K–12 and higher-education Grok *assistant* pilots: none found either. The only named national education claim in the Grok file is the December 2025 announcement, and the announcement has not been followed, in anything fetched, by results.

Keep the dates on the page, because the popular story will scramble them. 11 December 2025: tutoring announcement, chat product, national scale claimed, no outcomes. 11 August 2026: Grok Bot early beta, no education SKU, no educational job illustrated. 22 January 2026: Common Sense rates the chat product unacceptable — between the announcement and the Bot, and not about either a completed Salvadoran deployment or a Bot that had not launched. 26 August 2026: Bot expanded to more plans, still no school in the copy. A timeline that honest is less exciting than a headline that says Elon Musk is putting agents in every Salvadoran classroom. The headline would be three errors: the product is the assistant, the program is announced rather than evaluated, and agents in the Grok Bot sense are not in the announcement.

Do not treat xAI El Salvador as a Grok Bot deployment. Do not treat it as outcomes. Do not treat the one-million-student figure as a completed headcount; it is the announcement’s scale claim. Vendor announcements of national scale sit in a scorecard of their own, later in this chapter, next to vendor teacher SKUs with named early schools and district case studies. El Salvador belongs in the first cell of that scorecard: announced, unevaluated. Grok Bot belongs in a different cell: not found in a school. The assistant belongs in the first row of the four-class table. Three cells. Three objects. One name. The name is the trap.

So what did schools actually integrate, if not the last row as a student-facing product?

What “integrate into modern education” looked like in 2026

End on the beachhead, not on a prediction. The education market of 2026 bought constrained tutors and teacher copilots. It has not bought student-facing autonomous agents. Hermes Agent and Grok Bot are culturally important to this book’s named reader. They are not what districts procure. The first real agent beachhead is the teacher’s evening and the registrar’s inbox, not the third grader’s desk.

The pattern in K–12, visible in the largest U.S. system and in the union resolutions, is not a straight line from ban to agent. It is 2022–23 panic bans of ChatGPT, then 2024–25 teacher-tool adoption under data-privacy agreements and FERPA contracts, then 2026 district and state guidance that is neither ban nor open season. New York City’s March 2026 traffic lights arrive, in the Chalkbeat and Gothamist accounts, about three years after a short-lived ChatGPT ban, and after a parent petition of about 1,500 signatures seeking a pause.¹¹⁰ UNESCO’s student framework (twelve competencies,

curriculum-aligned tutoring. CoinDesk, same day, restates the announcement — restatement, not independent verification. **No outcome evaluation, independent audit, or classroom study found.** Announced national program, not a documented completed deployment as of Aug. 30, 2026. Grok-as-tutor, not Grok Bot-as-agent. Salvadoran MINED status updates were not found beyond the xAI announcement.

¹¹⁰New York City Public Schools, “Guidance on Artificial Intelligence,” <https://www.schools.nyc.gov/about-us/vision-and-mission/artificial-intelligence/guidance-on-artificial-intelligence> (March 2026). Context: Chalkbeat New York, Mar. 24, 2026, <https://www.chalkbeat.org/newyork/2026/03/24/preliminary-ai-policy-nyc-schools/>; Gothamist, “NYC schools get AI guidance using red light, green light model,” <https://www.gothamist.com/news/nyc-schools-get-ai-guidance-using-red-light-green-light-model>. Arrives about three years after a short-lived ChatGPT ban; parent petition of about 1,500 signatures sought a pause.

four dimensions, three levels: Understand, Apply, Create) and teacher framework (fifteen competencies, five dimensions: human-centred mindset, ethics, foundations and applications, pedagogy, professional learning; levels Acquire, Deepen, Create) are the international literacy documents of record; Digital Learning Week 2026, 8–11 September, Paris, is titled “Education in the age of AI.” UNESCO does not endorse student-facing autonomous agents. A TVET AI guide is dated 3 July 2026.¹¹¹ The U.S. Department of Education’s supplemental priority, 13 April 2026, effective 13 May 2026, Docket ED-2025-OS-0118, 91 FR 18774, traces to Executive Order 14277 of 23 April 2025, “Advancing Artificial Intelligence Education for American Youth.” Public comments included both strong support and calls to keep AI out of K–12 over cognitive and emotional harm; the Department declined a ban. The priority remains a scoring tool for discretionary grants, not a classroom mandate, and not an authorization of unconstrained student agents.¹¹² OECD Digital Education Outlook 2026, from the listing page and search snippets: generative AI can support learning when guided by teaching principles; without pedagogical guidance, outsourcing tasks raises performance without learning gains. Roles named: tutor, partner, assistant. Full PDF fetch was blocked. Full-report statistics remain **UNVERIFIED**.¹¹³

K–12 integration, as of the documents fetched, is five things, and almost never a sixth.

First, teacher-side generation under a data-privacy agreement: MagicSchool, Diffit, Claude for Teachers, ChatGPT for Teachers, Gemini. ChatGPT for Teachers launched as a K–12 educator workspace, not intended for students. CNBC, 19 November 2025, reported a first cohort of districts representing roughly 150,000 educators, free through June 2027 in that report. OpenAI’s own expansion post, “Bringing ChatGPT for Teachers to more U.S. school districts,” timed out on fetch. Additional district counts, and a National Academy for AI Instruction / AFT “400,000 educators” figure appearing in search snippets, are **UNVERIFIED** until that page is retrieved. Do not list named cohort districts from unfetched OpenAI HTML. The AFT National Academy for AI Instruction figure of \$23 million and 400,000 educators is a 22 July 2025 *target* in a union resolution — a plan, not a headcount of trained teachers, and not the same object as CNBC’s first cohort.¹¹⁴ Claude

¹¹¹ UNESCO student framework: 12 competencies, four dimensions, three levels (Understand / Apply / Create). Teacher framework: 15 competencies, five dimensions (human-centred mindset, ethics, foundations/applications, pedagogy, professional learning); levels Acquire / Deepen / Create. Digital Learning Week 2026, 8–11 September, Paris: “Education in the age of AI.” TVET AI guide dated July 3, 2026. See note 4. UNESCO does not endorse student-facing autonomous agents.

¹¹² 91 Fed. Reg. 18774 (note 3). Exec. Order No. 14277 (Apr. 23, 2025), “Advancing Artificial Intelligence Education for American Youth.” Public comments included support and calls to keep AI out of K–12; ED declined a ban. EO 14277 is an adoption-and-workforce order, not a child-safety regulation.

¹¹³ OECD, *Digital Education Outlook 2026*, listing page https://www.oecd.org/en/publications/oecd-digital-education-outlook-2026_062a7394-en.html. Listing page and search snippets: GenAI can support learning when guided by teaching principles; without pedagogical guidance, outsourcing raises performance without learning gains. Roles: tutor, partner, assistant. Full PDF fetch **blocked (Cloudflare)**. Treat full-report statistics as **UNVERIFIED**.

¹¹⁴ CNBC, “OpenAI rolls out ChatGPT for Teachers,” Nov. 19, 2025, <https://www.cnbc.com/2025/11/19/openai-chatgpt-for-teachers-k-12-educators.html>. K–12 educator workspace, not intended for students; first cohort of districts representing roughly 150,000 educators; free through June 2027 **in that report**. OpenAI expansion HTML (“Bringing ChatGPT for Teachers to more U.S. school districts”; openai.com/index/chatgpt-for-teachers/) **timed out on fetch, 30 August 2026**. Additional district counts and National Academy for AI Instruction / AFT “400,000 educators” figures in search snippets are **UNVERIFIED** until that page is retrieved. Do not list named cohort districts from unfetched OpenAI HTML. American Federation of Teachers, National Academy for AI Instruction resolution (July 22, 2025): \$23 million, Microsoft lead plus OpenAI plus Anthropic, target 400,000 educators over five years — a **plan / target**,

for Teachers, already described, is the teacher SKU that includes last-row machinery, Cowork and Claude Code, behind an educator-only, 18-plus, FERPA-aligned door. Gemini in Workspace for Education is first-row chat and notebook, free, all ages, Core service, with Google’s vendor claims on training and COPPA/FERPA. MagicSchool and Diffit are prompted generators with education guardrails, not agents.

Second, constrained student tutors with rostering, logging, and Socratic refusal: Khanmigo, SchoolAI Spaces, MagicStudent, as already placed. These are the second row. They are what a board can buy for students without buying a cloud computer.

Third, AI literacy modules and acceptable-use updates. Cornell, in higher education, expanded an AI Critical Literacy Program from a spring pilot to all incoming students plus interested faculty and staff, with four Canvas modules: What is GenAI?; Ethics; GenAI and Learning; Creating an AI Policy for Personal Use. The Center for Teaching Innovation reported literacy gains on how models work and are trained. That is CTI-reported, not a published RCT. The explicit goal is to help students decide whether AI use *contributes to learning*.¹¹⁵ The 2026 EDUCAUSE report on the impact of AI on learning assessment, a survey of 438 faculty and staff, is cited here from the library page: faculty warming to AI; “does it make teaching harder or easier? It depends”; students need clearer use and non-use rules; most surveyed faculty use AI to *create* assessments and believe students use AI *while taking* them; faculty want choice; AI literacy includes knowing when *not* to use tools. The full report PDF was not fetched.¹¹⁶ The 2026 EDUCAUSE Horizon Report, Teaching and Learning edition, describes AI reshaping assessment toward authentic, process-based demonstrations of learning. Full report not fetched.¹¹⁷ UIC’s Center for the Advancement of Teaching Excellence published a teaching guide with three assignment types — recall, skill application, cumulative essays — and preserve-versus-leverage strategies, citing reduced cognitive engagement with generative-AI writing.¹¹⁸ Khlaif and colleagues (2025), in *Education Sciences*, interviewed 61 faculty on an “Against, Avoid, Adopt, and Explore” redesign framework; motivations included integrity, career preparation, technological change, and institutional policy.¹¹⁹ These are literacy and assessment-redesign documents. They are not agent deployments.

Fourth, procurement as the real control surface: New York City Public Schools’ ERMA ten-step

not a headcount of trained teachers, and not the same object as CNBC’s first cohort.

¹¹⁵Jose Beduya, “Cornell expands AI Critical Literacy Program,” *Cornell Chronicle* (August 2026), <https://news.cornell.edu/stories/2026/08/cornell-expands-ai-critical-literacy-program-incoming-students>. Four Canvas modules; CTI + Library; CTI-reported literacy gains (how models work / are trained), **not a published RCT**. Explicit goal: whether AI use contributes to learning.

¹¹⁶EDUCAUSE, *The Impact of AI on Learning Assessment* (June 2026), survey of 438 faculty/staff, <https://library.educause.edu/resources/2026/6/2026-educause-the-impact-of-ai-on-learning-assessment-report>. Cited from the library page. **Full report PDF not fetched.**

¹¹⁷EDUCAUSE, *2026 Horizon Report, Teaching and Learning Edition*, <https://library.educause.edu/resources/2026/5/2026-educause-horizon-report-teaching-and-learning-edition>. AI reshaping assessment toward authentic, process-based demonstrations of learning. **Full report PDF not fetched.**

¹¹⁸University of Illinois Chicago, Center for the Advancement of Teaching Excellence, “Adapting Assignments for the Age of AI,” <https://teaching.uic.edu/cate-teaching-guides/digital-learning/adapting-assignments-for-the-age-of-ai/>. Three assignment types (recall, skill application, cumulative essays); preserve-versus-leverage strategies; cites reduced cognitive engagement with genAI writing.

¹¹⁹Khlaif et al., “Redesigning Assessments for AI-Enhanced Learning,” *Education Sciences* 15, no. 2 (2025): 174, <https://doi.org/10.3390/educsci15020174>. Qualitative; 61 faculty; “Against, Avoid, Adopt, and Explore.”

vendor review before any tool touches student data; Student Data Privacy Consortium NDPA; iKeepSafe certifications. New York City, the nation’s largest system, issued March 2026 guidance on a traffic-light framework. Red, never: placement, discipline, eligibility, promotion, graduation; IEP and 504 development; assessments and grading as the decision; surveillance and behavior monitoring; counseling and crisis; student data training models; personally identifiable information in unapproved tools. Yellow, human review required: data-pattern surfacing; translations of critical communications by a qualified linguist; scaffolds for diverse learners with certified-staff review; student use for research, exploration, and creative projects with educator guidance. Green, encouraged, still ERMA-approved, and the human owns the output: lesson brainstorming, drafting communications, scheduling and summarizing non-sensitive information, operational synthesis, non-critical translation with a disclaimer if unreviewed, accessibility, professional learning. AI-specific vendor rules added in December 2024: disclose AI capabilities; no training on student data; transparency. ERMA as fetched is privacy and security, not yet bias, equity, or instructional effectiveness, which were promised for a June 2026 Playbook. Playbook topics still open as of the fetched page included homework and integrity, cognitive offloading, grade-band guidance, family notification, and managed versus personal accounts.¹²⁰ That is what “integrate” looks like in the largest U.S. district: traffic lights for *teachers*, a vendor review that is not yet an efficacy review, and homework rules still unfinished. Local district examples fetched as PDFs in search, not fully extracted — Little Falls, Minnesota, Policy 625, using a human-AI-human model with broader access in grades 9–12; Council Rock, Pennsylvania, 2026 guidelines, district-approved tools only, academic integrity still applying — should be treated as partially verified via search snippets.¹²¹

Fifth, almost no student-facing Hermes, Grok Bot, or Cursor-class autonomy. Those products lack COPPA-age design, FERPA DPAs, and, in Grok Bot’s case, any education SKU. There is no documented K–12 student-facing autonomous agent program in the sources this chapter stands on. That absence is a finding.

The sixth thing, which is not a K–12 student program, is the beachhead: teacher-side Cowork and campus-side Workspace Agents. A 4 p.m. exit-ticket analysis on a schedule is politically easier than a student-side Hermes Agent on a VPS. Unions may accept the former and fight the latter. The American Federation of Teachers, in a resolution adopted 18 July 2026, would put no screens in pre-K through grade 2 except for disability or compelling need; oppose student-facing AI in elementary school; ban social-companion chatbots at least until age 16; treat AI as a limited supervised tool; hold a “gold standard” on privacy; and put public money to class size and mental health before unproven AI. The same union partners with Anthropic on K–12 terms, as Anthropic’s launch post records.¹²² The National Education Association’s 2024 task force, still the

¹²⁰NYCPS March 2026 guidance (note 38). Traffic-light framework as summarized in the text, from the primary page. ERMA: ten-step vendor review before any tool touches student data. AI-specific vendor rules added December 2024: disclose AI capabilities; no training on student data; transparency. ERMA as fetched is privacy/security, not yet bias/equity/instructional effectiveness (promised for a June 2026 Playbook). Playbook topics still open as of the fetched page: homework/integrity, cognitive offloading, grade-band guidance, family notification, managed versus personal accounts. NYCPS: “AI is not a replacement for your child’s teacher.”

¹²¹Little Falls (Minnesota) Policy 625: H-AI-H model (human at initiation and conclusion); broader access grades 9–12. Council Rock (Pennsylvania) 2026 guidelines: district-approved tools only; academic integrity still applies. PDFs appeared in search; **not fully extracted**. Treat full contents as **partially verified via search snippets**.

¹²²American Federation of Teachers, “Preparing for the Age of Artificial Intelligence...,” resolution adopted July 18, 2026, <https://www.aft.org/resolution/preparing-age-artificial-intelligence-protecting-democracy-workers-students->

association’s baseline, is “aid but not replace,” with humans at the center of high-stakes decisions: evaluation, placement, discipline, diagnosis.¹²³ The live labor debate is not whether a teacher may schedule a folder analysis. It is whether the last row crosses onto a student’s machine. Adjacent to that fight, and already forced by the landscape, is age. AFT’s pre-K through 2 screen line and its opposition to elementary student-facing AI sit next to Google’s all-ages Gemini-in-Education onboarding, endorsed in Google’s materials by ConnectSafely and FOSI, and next to the Department of Education’s “age-appropriate” language without national standards, left to families and localities. Thirteen is a terms-of-service number and a UNESCO independent-use line; it is not a last-row license. Data and training, by contrast, are a rare consensus across NYCPS, MagicSchool, Claude for Teachers, and Gemini Education: no student data in model training. Consumer ChatGPT, Grok, and Hermes-on-a-personal-API do not automatically inherit that consensus. A student who pastes schoolwork into a consumer chat window has left the DPA. A student whose district-issued tutor is logged, rostered, and contractually barred from training is still in the second row, under a contract. The last row, as shipped by Hermes Agent and Grok Bot, is not that contract. Official “will not replace” language is a bloc: UNESCO; the International Task Force on Teachers for Education 2030, September 2025; NEA 2024; AFT July 2026; NYCPS (“AI is not a replacement for your child’s teacher”); U.S. Department of Education comments emphasizing human educators; Anthropic framing Claude for Teachers as returning time to students.¹²⁴ Pressure in the other direction is also documented: teacher shortages; U.S. ED priority language on tutoring, personalization, and administrative efficiency; Brookings on generative-AI tutoring at scale in teacher-shortage systems; xAI El Salvador’s “every student”; vendor “AI operating system for schools.” AFT explicitly fears public schools becoming “testing laboratories for unregulated technology.” The empirical gap under that fight: no fetched study shows an autonomous agent substituting for a certified teacher of record in a U.S. district. New York City red-lights grading, IEPs, discipline, and counseling. That is the current legal-practical ceiling.¹²⁵

Higher education bought chat, and is being offered ops agents. The EDUCAUSE 2025 AI Landscape Study, quoted via EdTech Magazine and Jenay Robert, found that chatbots were the top institutionwide AI license at 37 percent of respondents — chatbots, not agents.¹²⁶ Do not mash this 37

and-planet. No screens PK–2 except disability/compelling need; oppose student-facing AI in elementary; ban social-companion chatbots at least until age 16; AI as limited supervised tool; National Academy; “gold standard” privacy; public money to class size/mental health before unproven AI. Also partners with Anthropic on K–12 terms (Anthropic launch post, note 28). AFT fears public schools becoming “testing laboratories for unregulated technology.”

¹²³National Education Association, *Report of the NEA Task Force on Artificial Intelligence in Education* (2024), https://www.nea.org/sites/default/files/2024-06/report_of_the_nea_task_force_on_artificial_intelligence_in_education_ra_2024.pdf. “Aid but not replace”; humans at the center of high-stakes decisions (evaluation, placement, discipline, diagnosis).

¹²⁴“Will not replace” bloc as documented in landscape memo §6.3: UNESCO (note 4); International Task Force on Teachers for Education 2030, *Promoting and protecting teacher agency in the age of artificial intelligence* (September 2025); NEA 2024 (note 51); AFT July 2026 (note 50); NYCPS (note 48); U.S. Department of Education comments emphasizing human educators; Anthropic framing of Claude for Teachers as returning time to students (note 28). International Task Force PDF was listed in the research bibliography, not fully extracted.

¹²⁵No fetched study shows an autonomous agent substituting for a certified teacher of record in a U.S. district. NYC red-lights grading, IEPs, discipline, counseling (note 48). That is the current legal-practical ceiling in the largest U.S. system, not a national statute — there is no U.S. federal student-AI statute.

¹²⁶EDUCAUSE 2025 AI Landscape Study, quoted via EdTech Magazine / Jenay Robert: chatbots were the top institutionwide AI license at **37 percent** of respondents — **chatbots, not agents**. EdTech Magazine, “AI Agents for Education,” December 2025, <https://edtechmagazine.com/higher/article/2025/12/ai-agents-higher-education-transforming->

percent with the OECD TALIS figure, reported on the Digital Education Outlook 2026 listing page, of 37 percent of *teachers* using generative AI. Different populations, same number, easy to mash. The OECD full PDF was Cloudflare-blocked; treat full-report statistics as **UNVERIFIED**.¹²⁷ Dan Harder of the University of Tennessee Knoxville, quoted in the same EdTech Magazine piece (December 2025), said today's campus AI is agent-like within parameters, "a long way from autonomous AI agents that weigh the pros and cons... and act without human oversight." EdTech Magazine also cites Johns Hopkins Agent Laboratory, an LLM framework for literature, code, and write-up whose authors claim an 84 percent cost decrease versus prior autonomous-research methods, and UT Verse at UT Knoxville, campus assistants for Wi-Fi, dining, coursework, translation. Agent Laboratory's paper and UT Verse's primary pages were not fetched. Treat those two campus examples as journalism-sourced, **UNVERIFIED** against primary documents.¹²⁸ Clemson, already named, is documented access. GitHub Copilot at UC San Diego's computer science and engineering program appears in a Microsoft customer story: faculty (Porter, Shah) redesigned introductory courses so students first demonstrate manual competency, then use Copilot; research on trust calibration; plans to prepare students for "emerging agentic AI." That is curriculum redesign around a coding assistant, not a published RCT of Cloud Agents, and it is a vendor customer story.¹²⁹ Assignment redesign in practice, as a synthesis of the fetched higher-education sources and not a new claim: process portfolios, oral defenses, in-class writing, scaffolded drafts, disclosure norms, AI-allowed versus AI-forbidden lanes on the syllabus, and, at UC San Diego CSE, manual competency gates before AI-assisted work.

A scorecard, so marketing cannot hide in the same cell as a newsroom.

Claim type	Example	Status
Vendor announcement of national scale	xAI × El Salvador, 5,000 schools / 1 million students	Announced December 2025; no outcomes fetched
Vendor teacher SKU with named early school	Claude for Teachers; Prospect Schools (feedback); Detroit Public Schools Community District (impact pilot announced)	Product live July 2026; impact results not published

student-services-and-support-perfcon.

¹²⁷OECD TALIS 2024 as reported on the *Digital Education Outlook 2026* listing page: **37 percent** of teachers already use GenAI for work tasks. **UNVERIFIED at full-PDF level** (note 41). Different population from the EDUCAUSE 37 percent of *institutions* licensing chatbots. Do not mash the two figures.

¹²⁸EdTech Magazine, December 2025 (note 54). Dan Harder, University of Tennessee Knoxville: today's campus AI is agent-like within parameters, "a long way from autonomous AI agents that weigh the pros and cons... and act without human oversight." Same piece cites Johns Hopkins Agent Laboratory (authors claim 84 percent cost decrease versus prior autonomous-research methods) and UT Verse at UT Knoxville (campus assistants for Wi-Fi, dining, coursework, translation). Agent Laboratory paper and UT Verse primary pages **not fetched**. Journalism-sourced; **UNVERIFIED against primary docs**.

¹²⁹Microsoft Customer Stories, University of California San Diego / GitHub Copilot, <https://www.microsoft.com/en/customers/story/26539-university-of-california-san-diego-github-copilot>. Faculty (Porter, Shah) redesigned intro courses so students first demonstrate manual competency, then use Copilot; research on trust calibration; plans to prepare students for "emerging agentic AI." **Vendor customer story**, not a peer-reviewed trial of Cloud Agents. Curriculum redesign around a coding assistant.

Claim type	Example	Status
District case study (vendor site)	Midway ISD × MagicSchool	Named, dated roadmap; vendor case study
Local journalism of board adoption	Midland Public Schools × SchoolAI, August 2026	Documented selection, not outcomes
Vendor plus district research, ESSA 3	SchoolAI × Jordan School District; MagicSchool ESSA 3	Real districts; not RCTs
University newsroom of access	Cornell AI literacy expansion; Clemson Codex/GPT-5.5	Documented programs
University newsroom of a learning RCT	—	Not found this pass
Hermes Agent or Grok Bot in a school	—	Not found
OpenAI Workspace Agents “in production at University X”	Playbook only; research preview	No named campus case fetched

An IEEE AIEI 2026 survey on “agentic AI-powered adaptive learning” noted an empirical gap on fully agentic platforms. The full text, doi:10.1109/aiei69164.2026.11497428, was not fetched.¹³⁰ Elicit, Consensus, SciSpace, and Perplexity as research agents were not fetched this pass. No primary district page documenting production “school GPTs” was fetched. OpenAI Workspace Agents and Claude Skills and Cowork are the institutional replacement for one-off custom GPTs: admin-controlled, off by default, with review points.

What this chapter must not be taken to have shown: that last-row systems are imaginary; that teacher-facing Cowork is a student agent; that El Salvador is finished, or is Grok Bot, or has outcomes; that Cursor “101” listings are credit-bearing courses; that a Telegram hobby project is a university; that vendor partner counts and ESSA 3 studies are randomized trials; that Hermes 4’s unread benchmarks are classroom evidence; that Grok Bot’s FAQ, now fetched, is a FERPA-ready school product (it is not: cloud storage required, no education SKU); that UT Verse and Johns Hopkins Agent Laboratory have been verified against primary pages. Open questions this chapter leaves standing, because filling them with analogy would be the popular story in a more respectable font: any Hermes Agent or Grok Bot institutional acceptable-use policy, DPA, or classroom SOP — none found; El Salvador program status beyond the December 2025 announcement; a named campus with OpenAI Workspace Agents in production; independent, non-vendor evaluations of Khanmigo, MagicSchool, SchoolAI, and Claude for Teachers.

The four classes are distinct. Chat is not an agent. Hermes 4 is not Hermes Agent; Nous’s own portal says the model family is not recommended inside the runtime. Grok the assistant is not Grok Bot, which launched 11 August 2026 with no education SKU, and is not Grok-as-tutor in El Salvador, announced 11 December 2025, with no outcomes found. There is no documented

¹³⁰IEEE AIEI 2026 survey on “agentic AI-powered adaptive learning,” doi:10.1109/aiei69164.2026.11497428. Noted an empirical gap on fully agentic platforms. **Full text not fetched.** Elicit, Consensus, SciSpace, and Perplexity as research agents: not fetched this pass. No primary district page documenting production “school GPTs” was fetched.

K–12 student-facing autonomous agent. Teacher-facing Cowork and Workspace Agents are the 2026 beachhead. The causal record is almost entirely about two of the four classes — chat and tutors. That record is Chapter 3. Do not use it, when it arrives, as if it had measured the last row. Agent offloading, unattended completion of the job, is a stronger intervention than a chatbot that waits, and it is under-studied in education. An honest book flags that gap rather than extrapolating a writing-task study onto Grok Bot. The collision the rest of this book is about is already visible in the documents just cited: policy red lines — grading, IEPs, surveillance, elementary student-facing AI, companion chat — are hardening just as agent autonomy, in the last row, is increasing. The beachhead is the teacher’s evening. The unnamed object in the popular story is a student agent that the documents do not show.

Chapter 3

The Exam After the Chat Window Closes



Figure 4: A closed laptop, an open notebook with a pencil, and a clock at a dark window.

The number the public will not have heard is not the practice score. In a field experiment in a large Turkish high school, students who used ChatGPT-style GPT-4 on mathematics practice looked about 48 percent better than classmates who used notes and a textbook. Then the chat window closed. On a closed-book exam of isomorphic problems, the same students scored about 17 percent worse than the never-AI arm. They did not perceive that they had learned less.¹³¹

¹³¹Bastani et al., “Generative AI without Guardrails,” *PNAS* 122, no. 26 (2025). Intention-to-treat, classroom-clustered: GPT Base practice +0.137 (control mean 0.284) $\approx$ 48 percent better; unaided exam -0.054 (control mean

That dissociation is older than ChatGPT. Nicholas Soderstrom and Robert Bjork, reviewing the experimental literature in 2015, treat *learning* and *performance* as different constructs. Current success with a tool in hand, or under conditions that make practice look fluent, is not the same thing as later, unaided skill. Conditions that make practice look easy—massed study, restudy, retrieval that costs nothing—often produce weaker long-term learning than conditions that feel harder.¹³² Bjork’s *desirable difficulties*—spacing, interleaving, generation, retrieval practice—slow or worsen immediate performance while improving later retention and transfer. Difficulty is desirable only when it is achievable. If the learner cannot succeed even with effort, the difficulty is undesirable; corrective feedback that keeps the task in reach is part of the framework, not a betrayal of it.¹³³ Generative AI, by default, is *undesirable ease*. It raises session artefacts. It can starve the retrieval, generation, and error-correction that build memory. The Turkish result is not a surprise to that literature. It is that literature, measured in a school, with a model that will finish the problem if you ask.

Retrieval practice is the desirable difficulty most relevant to a chat window. Testing as a learning event, not only as measurement, is among the most replicated findings in cognitive psychology: Roediger and Karpicke in 2006, Karpicke and Roediger in 2008. The testing effect is strongest on delayed tests and often looks worse than restudy on immediate tests.¹³⁴ That is already the Bastani shape, decades early: the condition that looks worse now is the one that holds later. A qualifier from cognitive load theory matters for anyone who would turn “make it hard” into a slogan. Van Gog and Sweller (2015) and Hanham, Leahy, and Sweller (2017) find that the testing effect shrinks or reverses for high element-interactivity material—complex, interacting concepts—among novices.¹³⁵ Difficulty is not a virtue when working memory is already taxed. A Hermes-class assistant that dumps a complete multi-step solution is ease. A Hermes-class assistant that withholds every hint until the novice drowns is not a desirable difficulty either. The pedagogical target, which Chapter 5 will spend as struggle policy, is already visible in the cognitive science: generation and retrieval when the schema can bear them; worked examples and hints when it cannot; an unaided test either way. This chapter’s job is not to re-derive Bjork. It is to show that the 2022–26 generative trials moved, in the directions that literature predicted, when the designers either

0.321) $\approx$ 17 percent worse than never-AI. Students in GPT Base did not perceive that they had learned less. Author PDF fetched for the research this chapter stands on.

¹³²Nicholas C. Soderstrom and Robert A. Bjork, “Learning versus Performance: An Integrative Review,” *Perspectives on Psychological Science* 10, no. 2 (2015): 176–199. Conditions that make practice look fluent often produce weaker long-term learning than conditions that feel harder. Analytic frame for the chapter; not re-fetched in full in the 30 August 2026 pass.

¹³³Robert A. Bjork, “Memory and Metamemory Considerations in the Training of Human Beings,” in *Metacognition*, ed. Janet Metcalfe and Arthur P. Shimamura (Cambridge, MA: MIT Press, 1994), 185–205; Elizabeth L. Bjork and Robert A. Bjork, “Making Things Hard on Yourself, but in a Good Way: Creating Desirable Difficulties to Enhance Learning,” in *Psychology and the Real World*, ed. Morton Ann Gernsbacher et al. (New York: Worth, 2011), 56–64; Robert A. Bjork and Elizabeth L. Bjork, “Desirable Difficulties in Theory and Practice,” *Journal of Applied Research in Memory and Cognition* 9, no. 4 (2020): 475–479. Difficulty is desirable only when achievable.

¹³⁴Henry L. Roediger III and Jeffrey D. Karpicke, “Test-Enhanced Learning: Taking Memory Tests Improves Long-Term Retention,” *Psychological Science* 17, no. 3 (2006): 249–255; Jeffrey D. Karpicke and Henry L. Roediger III, “The Critical Importance of Retrieval for Learning,” *Science* 319, no. 5865 (2008): 966–968. Testing effect strongest on delayed tests; often looks worse than restudy immediately.

¹³⁵Tamara van Gog and John Sweller, “Not New, but Nearly Forgotten: The Testing Effect Decreases or Even Disappears as the Complexity of Learning Materials Increases,” *Educational Psychology Review* 27 (2015): 247–264; Jose Hanham, Wayne Leahy, and John Sweller, “Cognitive Load Theory, Element Interactivity, and the Testing and Reverse Testing Effects,” *Applied Cognitive Psychology* 31, no. 3 (2017): 265–280.

removed generation or required it.

This chapter is the load-bearing empirical chapter of the book. Design is the treatment. “AI in education” is not one intervention. Unguarded ChatGPT, a hint-only tutor, a teacher-prompted Copilot lab, Khanmigo sitting one click away, and a custom physics tutor with instructor-written solutions in the prompt are different pedagogies that happen to share a model class. Averaging them is how a superintendent ends up citing a Harvard physics tutor as if it had measured a Tennessee chat window, and a Tennessee chat window as if it had measured an unattended agent. The 2022–26 generative-AI causal record is almost entirely turn-based tutors and Copilot-like autocomplete. There is no identified learning randomized controlled trial of Hermes-class autonomous multi-step agents. There is, as of 30 August 2026, no What Works Clearinghouse intervention report specific to generative AI tutors. Human high-dosage tutoring remains the benchmark. Earlier computer-assisted learning sits well below it. When the model does the generative work, session artefacts look better and later unaided skill often does not. When the model is constrained to hints, retrieval, or explanation, harm shrinks and sometimes immediate learning rises. Guardrails are the treatment. Default ChatGPT is a crutch students do not notice.

The rest of the chapter spends the record, study by study, with the coefficients the sources actually reported, and then leaves the holes standing. Do not use any of it as if it had measured the last row of Chapter 2. Agent offloading—unattended completion of the job—is a stronger intervention than a chatbot that waits, and it is under-studied in education. An honest book flags that gap rather than extrapolating a writing-task EEG onto Grok Bot.

Design is the treatment; human tutoring is the benchmark

The popular story averages. It says “studies show AI helps,” or “studies show AI harms,” as if the object under study were one object. It is not. Five designs that share a model class are still five designs.

Unguarded ChatGPT will, by default, complete the assignment. A hint-only tutor is a designed refusal: the system prompt withholds the full solution, and the student is supposed to generate. A teacher-prompted Copilot session in a supervised lab is a curriculum plus a prompt library plus an adult in the room. Khanmigo, one click away inside an existing practice block, is access, not a session. A custom physics tutor with worked solutions in the prompt, sequential scaffolding, and the same worksheets the class already uses is an engineered lesson, not a conversation window. Collapse those five and every coefficient in this chapter becomes folklore. Keep them separate and a pattern appears: the model class is shared; the pedagogy is the intervention; the unaided later test is the outcome that matters.

The right comparison class is not last year’s chatbot. It is human tutoring, and then the older software that already tried to approximate it.

Nickow, Oreopoulos, and Quan, in the *American Educational Research Journal* in 2024, pool experimental preK–12 tutoring at **0.288 SD** (SE 0.029).¹³⁶ That is the benchmark a generative product

¹³⁶Andre Nickow, Philip Oreopoulos, and Vincent Quan, “The Promise of Tutoring for PreK–12 Learning: A Systematic Review and Meta-Analysis of the Experimental Evidence,” *American Educational Research Journal* 61, no. 1 (2024): 74–107. Pooled experimental preK–12 tutoring **0.288 SD** (SE 0.029).

is quietly being asked to match when a vendor says “personal tutor for every child.” Bloom’s 1984 two-sigma observation was about one-to-one *human* tutoring, not about a language model.¹³⁷ The 2024 meta-analysis is the current causal number for that human intervention. It is not a number this chapter will launder onto Khanmigo, or onto GPT-4, or onto an agent that files the homework while the student is at lunch.

Keep the bands in view for the rest of the chapter, because every generative coefficient will be silently compared with them whether or not the writer admits it. Human high-dosage tutoring, pooled experimental preK–12: 0.288 SD. Earlier CAL: typically 0.05–0.20 SD. Pre-generative ITS: positive averages, not a generative result. Adaptive systems 2010–2022 versus non-adaptive: $g = 0.70$, not a generative result, not a WWC generative-tutor report. A 0.06 SD school-year ITT on MAP sits in the CAL band. A 0.238 SD English ITT in a six-week supervised lab sits near the tutoring benchmark for that dose and setting, and is still not human high-dosage tutoring at year scale. A 0.73-to-1.3 SD immediate physics post-test sits above every band in this paragraph and is still not a delayed test. The bands are not a leaderboard. They are how a reader refuses to be told that 0.06, 0.238, and 1.0 are “AI helping.”

Earlier computer-assisted learning typically sits in a different band. Escueta, Nickow, Oreopoulos, and Quan, in the *Journal of Economic Literature* in 2020, put the typical experimental CAL estimate in **0.05–0.20 SD**.¹³⁸ That is software as practice, often with a learner model, often without a conversation window. Kulik and Fletcher’s 2016 *Review of Educational Research* meta-analysis finds positive averages for pre-generative intelligent tutoring systems.¹³⁹ Wang, Wang, and colleagues, in the *Journal of Educational Computing Research* in 2024, meta-analyse AI-enabled *adaptive* systems from 2010 to 2022 at $g = 0.70$ versus non-adaptive comparators.¹⁴⁰ That is a large estimate. It is also **not a generative-AI result**. The window ends in 2022. The comparators are non-adaptive. Publication type, duration, and outcome moderate it. Treating a 2010–2022 adaptive-systems meta-analysis as evidence for Khanmigo, or for ChatGPT, or for Hermes Agent, is the laundering Chapter 2 exists to stop. The Institute of Education Sciences What Works Clearinghouse has older ITS intervention reports, including Cognitive Tutor Algebra I. As of the access date of the research this chapter stands on—30 August 2026—there is **no WWC intervention report that is specific to generative AI tutors**.¹⁴¹ That absence is a finding. It is not a gap to be filled with vendor ESSA-3 studies, and it is not a license to cite Kulik and Fletcher as if they had measured GPT-4.

¹³⁷ Benjamin S. Bloom, “The 2 Sigma Problem: The Search for Methods of Group Instruction as Effective as One-to-One Tutoring,” *Educational Researcher* 13, no. 6 (1984): 4–16. One-to-one *human* tutoring, not a language model.

¹³⁸ Maya Escueta, Andre J. Nickow, Philip Oreopoulos, and Vincent Quan, “Upgrading Education with Technology: Insights from Experimental Research,” *Journal of Economic Literature* 58, no. 4 (2020): 897–996. Earlier computer-assisted learning typically **0.05–0.20 SD**.

¹³⁹ James A. Kulik and J. D. Fletcher, “Effectiveness of Intelligent Tutoring Systems: A Meta-Analytic Review,” *Review of Educational Research* 86, no. 1 (2016): 42–78. Pre-generative ITS; positive averages; not a generative-AI result.

¹⁴⁰ Shuai Wang et al., “The Efficacy of Artificial Intelligence-Enabled Adaptive Learning Systems from 2010 to 2022 on Learner Outcomes: A Meta-Analysis,” *Journal of Educational Computing Research* (2024), <https://doi.org/10.1177/07356331241240459>. $g = 0.70$ versus non-adaptive. Window ends 2022. **Not a generative-AI result**. Moderated by publication type, duration, and outcome.

¹⁴¹ Institute of Education Sciences, What Works Clearinghouse, *Cognitive Tutor Algebra I* intervention report (example of older ITS reports), <https://ies.ed.gov/ncee/wwc/Intervention/818>. Search as of 30 August 2026: **no WWC intervention report specific to generative AI tutors**. Absence is a finding.

Two further holes belong on the page before any generative coefficient is written. First: **no identified learning RCT of Hermes-class autonomous multi-step agents**—systems that take a goal, plan, use tools, and continue after the student leaves. The 2022–26 corpus is almost entirely turn-based tutors and Copilot-like autocomplete. Chatbot studies do not measure agents. Second: SchoolAI’s vendor-commissioned 28 percent “critical-thinking gain” and MagicSchool’s ESSA Level III association are not randomized trials of learning, and they will not be treated as such here.¹⁴² Chapter 2 already placed those products in the tutor-app row. This chapter does not promote them.

The design problem, stated once so the rest of the chapter can spend evidence rather than slogans, is not whether to give students agents. It is which cognitive work must remain inside the student at each stage of expertise, and which work can be offloaded without dulling the skill being formed. Three conditions together dull. The offloaded process is the learning objective. The environment never measures the skill unaided. The student cannot detect model error. Frequency of use is not the variable. Whether the user retains the evaluative loop is. Contact with models is not what dulls. Those three conditions together are.

What follows is generative-AI evidence, grouped by design, not by hype. The cleanest school demonstration is first, because it is the one in which short-term performance and subsequent learning move in opposite directions, and because the public will have heard the practice score.

Bastani et al. 2025: unguarded GPT harms; guardrails erase the harm without beating control

Hamsa Bastani, Osbert Bastani, Alp Sungu, Haosen Ge, Özge Kabakçı, and Rei Mariman published, in the *Proceedings of the National Academy of Sciences* in 2025, the cleanest field demonstration that short-term performance and subsequent learning can move in opposite directions under generative AI.¹⁴³ The design is a school, not a lab hour. The outcome that matters is the closed-book exam, not the practice score. The treatment that matters is the system prompt.

The setting is a large Turkish high school, Fall 2023–24, nearly 1,000 students in grades 9–11. Four 90-minute sessions covered about 15 percent of semester mathematics. Three arms: no AI (notes and textbook); **GPT Base**, a ChatGPT-style GPT-4 tutor; **GPT Tutor**, the same interface, with a system prompt that withholds full solutions and includes teacher-written correct solutions and common-mistake hints. Practice was graded. Then a closed-book exam on isomorphic problems. Results below are intention-to-treat, classroom-clustered, as fetched from the paper. This chapter does not invent subgroup tables the source did not put in the public record this book stands on.

¹⁴²SchoolAI research with Jordan School District, Utah: vendor-commissioned 28 percent “critical-thinking gain”; ESSA Level III via Instructure; Bloom-rubric scoring of chat logs, not standardized tests; not an independent RCT. MagicSchool ESSA Level III association between teacher use and student ELA: vendor-funded evidence tier, not a large independent RCT. Placed in Chapter 2’s tutor-app row; not used as learning estimates here.

¹⁴³Bastani et al., “Generative AI without Guardrails,” *PNAS* 122, no. 26 (2025). Field RCT, large Turkish high school, Fall 2023–24, nearly 1,000 students, grades 9–11, four 90-minute sessions covering about 15 percent of semester math. Arms: no AI; GPT Base (ChatGPT-style GPT-4); GPT Tutor (same interface; withholds full solutions; teacher-written correct solutions and common-mistake hints). GPT Tutor practice +0.361 $\approx$ 127 percent better; unaided exam -0.004 , indistinguishable from control. GPT Base probe correctness about 51 percent; students still copied. Message clustering as described in the text. Authors: GPT Tutor reactive, not diagnostic; long-term learning unmeasured.

Assisted practice: GPT Base **+0.137** on a control mean of 0.284, about **48 percent** better. GPT Tutor **+0.361**, about **127 percent** better. The practice window, with the model in hand, makes both AI arms look like a triumph. GPT Tutor looks like a miracle. That is performance.

Unaided exam: GPT Base **-0.054** on a control mean of 0.321, about **17 percent** worse than never-AI. GPT Tutor **-0.004**, indistinguishable from control. The chat window closed. The students who had looked 48 percent better on practice were 17 percent worse when they had to do isomorphic problems alone. The students who had looked 127 percent better on practice were, on the exam, the same as the students who had never been given a model. Guardrails erased the harm. They did not, in this design, produce a learning *gain* over business-as-usual.

The mechanism is not mysterious, and it is not “the model is too dumb.” In a probe, GPT Base returned a correct answer only about **51 percent** of the time. Students still copied. Message clustering is the behavioral trace. GPT Base was dominated by “what is the answer” and paste-the-question. GPT Tutor shifted toward attempted answers and help-seeking. The Base students were using the model as a solution engine. The Tutor students were being forced to generate. Generation is a desirable difficulty. Copying a fluent wrong answer is not. The students in GPT Base **did not perceive** that they had learned less. That is the illusion of learning Louis Deslauriers, Logan McCarty, Kelly Miller, Kristina Callaghan, and Greg Kestin documented in 2019 in a different setting: active learning can feel worse and produce more learning; fluent presentation can feel better and produce less.¹⁴⁴ Bastani’s Base arm is the generative version of the same metamemory failure. The artefact was fluent. The encoding was not. The students believed the fluency.

Isomorphic problems are the right exam for this claim. The test items were not the practice items. They were the same structure in new clothes. If the student had formed a schema, the exam should have shown it. If the student had rented a solution for the hour, the exam should have shown that instead. Base showed renting. Tutor showed neither renting nor a gain: the guardrails stopped the copy-paste, and four sessions were not enough, in this design, to beat notes and a textbook on the closed-book test. Intention-to-treat, classroom-clustered, is the right estimator for a school: students were assigned as they sit, not as they comply. A compiler who “used it well” is not the policy object. The policy object is what happens when you put Base, or Tutor, in the room. Base in the room produced a practice triumph and an exam loss. Tutor in the room produced a larger practice triumph and an exam tie. The 51 percent probe is the detail a fluency-worshipping culture will skip. The model was wrong about half the time. Students copied anyway. Error-detection is one of the skills this book says school exists to form. Base outsourced it and then scored the outsourced artefact. The students could not tell. That is the third dulling condition in the opening: the student cannot detect model error. Base manufactured that condition and then measured its consequence.

Two author caveats belong on the page, not in a footnote the reader will skip. GPT Tutor was reactive, not diagnostic. It did not build a learner model, it did not decide what the student was ready for, and it did not behave like the third row of Chapter 2. It was a constrained conversation window with teacher-authored content in the prompt. Long-term learning is unmeasured. The

¹⁴⁴Louis Deslauriers, Logan S. McCarty, Kelly Miller, Kristina Callaghan, and Greg Kestin, “Measuring Actual Learning versus Feeling of Learning in Response to Being Actively Engaged in the Classroom,” *Proceedings of the National Academy of Sciences* 116, no. 39 (2019): 19251–19257. Illusion of learning; cited by Bastani on the perception mismatch.

exam was the next measurement, not a year later, not a state test, not transfer to a new domain. Four sessions covering 15 percent of a semester is a dose. It is not a curriculum. Replication outside one Turkish high school's four sessions is unshown. Those are holes. They are not licenses to ignore the sign of the Base exam coefficient.

The book line, which every later chapter will spend rather than re-litigate, is this. **Guardrails are the treatment. Default ChatGPT is a crutch.** The model class is held constant across Base and Tutor. The interface is held constant. The difference is what the system is allowed to do: give the answer, or withhold it and hint from a grounded key. That difference moved practice *and* the unaided exam. Averaging “GPT-4 in math class” would have produced a mush. Separating the prompt produces a result a designer can use. A Hermes-class or Grok Bot-class assistant that supplies fluent, complete answers is, by default, Base. The same assistant can be configured as Tutor if it withholds answers, forces generation, and is grounded in teacher-written solutions and mistakes. Configuration is pedagogy. Pedagogy is the intervention. The unaided exam is the test of whether the intervention formed a skill or rented one for the hour.

Engineered tutors can help immediately. That is a different design, a different comparator, and a different claim. It is not Bastani, and it is not ChatGPT.

Kestin et al. 2025: a custom physics tutor versus excellent active learning

Greg Kestin, Kelly Miller, Anna Klales, Timothy Milbourne, and Gregorio Ponti published, in *Scientific Reports* in 2025, the flagship *help* result in the generative-AI causal record: a pedagogically engineered tutor, on first-pass conceptual physics, beat even a good classroom for *immediate* mastery.¹⁴⁵ It is a strong paper. It will be mis-cited. The mis-citation is already visible in the popular story, which hears “AI tutoring outperforms in-class active learning” and writes “ChatGPT beats teaching.” The paper is not that sentence. This section will not let it become that sentence.

The design is a crossover RCT in Harvard Physical Sciences 2, Fall 2023. Of 233 enrolled, 194 were eligible. The treatment is a custom GPT-4 tutor—“PS2 Pal”—with instructor-written step-by-step solutions in the prompt, sequential scaffolding, growth-mindset and cognitive-load language, plus the same worksheets and intro videos the class already used. The comparator is the course's existing **in-class active learning**. Not lecture. Not a weak control. The outcome is an immediate pre/post on surface tension and fluid flow.

Median post-score: 4.5 (AI) versus 3.5 (in-class) versus a combined pre of 2.75. The authors report median learning gains more than double the in-class arm. Mann–Whitney $z = -5.6$, $p < 10^{-8}$. Linear regression effect **0.63 SD**. Quantile regression, because of ceiling, **0.73 to 1.3 SD**. Median time on task **49 versus 60 minutes**. Higher self-reported engagement and motivation; enjoyment

¹⁴⁵Greg Kestin, Kelly Miller, Anna Klales, Timothy Milbourne, and Gregorio Ponti, “AI Tutoring Outperforms In-Class Active Learning: An RCT Introducing a Novel Research-Based Design in an Authentic Educational Setting,” *Scientific Reports* 15 (2025): 17458, <https://doi.org/10.1038/s41598-025-97652-6>. Crossover RCT, Harvard Physical Sciences 2, Fall 2023, N = 194 eligible of 233 enrolled. Custom GPT-4 tutor (“PS2 Pal”) versus the course's existing in-class active learning, not lecture. Linear regression **0.63 SD**; quantile regression, because of ceiling, **0.73 to 1.3 SD**. Median time 49 versus 60 minutes. Delayed retention, transfer, and full-course outcome **not measured**. Not off-the-shelf ChatGPT. Not an agent. Lessons were Bloom understanding/applying/analysing, not creating.

and growth-mindset items not different. FCI-stratified subgroups both showed the AI advantage, which is the authors' answer to the charge that only the already-strong students moved.

Those are large immediate effects. Report them as large. Then report what they are not.

Not measured: delayed retention, transfer, or a full-course outcome. The authors state they do not claim superiority for higher-order synthesis. The lessons were Bloom understanding, applying, and analysing, not creating. The population is a selective university. The tutor is not off-the-shelf ChatGPT. The tutor is not an agent. The tutor is an engineered lesson: instructor-written solutions, sequential scaffolding, cognitive-load language, the same materials as class, a crossover so each student sees both conditions. Design is the treatment. The 0.73-to-1.3 SD quantile range is a design result. It is not a ChatGPT-default result. It is not a year-scale result. It is not evidence that a conversation window opened in a dorm beats a lecture the study did not run.

Ceiling is why the authors report a quantile range rather than a single regression. Linear 0.63 SD is already large. Quantile 0.73 to 1.3 SD is the estimate once the top of the scale stops hiding movement. Report both. Do not pick the 1.3 to put in a slide and leave the linear in the appendix. Time on task ran the other way from the usual "they just spent longer" dismissal: median 49 minutes with Pal, 60 in class. The AI arm was shorter and scored higher on the immediate post-test. That is a design result about first-pass conceptual items, Bloom understanding/applying/analysing, in a course that already had worksheets, intro videos, and an active-learning classroom worth beating. The faculty encoded the physics. The prompt had the step-by-step. The scaffolding was sequential. Growth-mindset and cognitive-load language were in the wrapper. Take those out and you do not have Pal. You have a chatbot a student opened. The paper's title will be quoted without those clauses. This chapter is the clauses.

Put Bastani and Kestin on the same table and the popular story breaks. Both used GPT-4. Bastani's Base arm harmed unaided math. Bastani's Tutor arm erased the harm without beating control on the exam. Kestin's Pal beat excellent active learning on an immediate conceptual post-test. The difference is not "physics versus math" as a slogan. The difference is the pedagogical wrapper: grounded solutions, sequential scaffolding, a designed lesson, a comparator that is already good, an outcome that is immediate. Kestin is what a custom tutor can do on first-pass mastery when the faculty encode the physics into the prompt. Bastani is what happens when the same model class is either left unguarded or constrained, and then the students are asked to do isomorphic problems without it. A book that cites only Kestin is selling Pal as ChatGPT. A book that cites only Bastani is selling a crutch as a law of nature. Both papers are in the record. Both are design results. Delayed skill under Pal is unmeasured. Year-scale skill under Tutor is unmeasured. Neither paper is a Hermes-class trial.

At public-school scale, the question is not whether a Harvard custom tutor can beat a Harvard recitation on Tuesday's post-test. The question is whether access to a coach-not-tell tutor, inside an existing intervention block, with thousands of students, moves a standardized math score. That is a different design. Access is not tutoring.

Oreopoulos and Low 2026: Khanmigo at scale, engagement as the bind

Philip Oreopoulos and Nina Low’s two-year cluster randomized trial is the honest school-scale number in this chapter, and it is the number the popular story will try to inflate into “AI tutoring works.” The authors themselves will not let that inflation stand, and neither will this book.¹⁴⁶

Eighteen Hamilton County, Tennessee middle schools, 2024–25 and 2025–26. Grades within school assigned to Khan Academy plus Khanmigo—configured to coach, not tell—during existing daily RTI math blocks, versus business-as-usual RTI: worksheets, Waggle, IXL, Zearn, small-group teaching. Sample: **6,902 student-terms**, 53 clusters. Pre-registered AEARCTR-0013519. Outcome: NWEA MAP mathematics. State TCAP results were registered as co-primary and were **not yet incorporated** as of the August 2026 paper. Cost: about \$15 per student per year for the license, delivered inside already-staffed blocks.

Intention-to-treat: **+1.26 national percentile ranks per term** (SE 0.60); about **0.040 population SD per term**; stacked about **0.06 SD per school year**, about **0.08** in year 2; implied full-year participation in year 2 about **0.14 SD**. Those magnitudes **resemble Khan Academy practice without the AI tutor**. The authors cite Eames and colleagues (2026) on the same platform for that comparison.¹⁴⁷ Do **not** treat the Khanmigo ITT as “AI tutoring works at 0.07 SD.” The authors say it looks like practice without the chatbot. The chatbot is the object the popular story wants to credit. The practice is what the coefficient resembles.

Engagement is the bind, and the message data are more important than the ITT for anyone who still thinks access is tutoring. Ninety-six percent of students messaged Khanmigo at least once. The median student messaged on **one-third of practice days** and in **17 percent of mistake sessions**. Only **14.5 percent** of messages contained a mathematical question or a step of reasoning. Within the treated arm, skills-to-proficiency predict achievement; chat counts do not, once mastery is controlled. A 2026 interface change that opened Khanmigo by default did **not**, in their message data, raise messaging rates. The window was there. The students did not talk to it, or did not talk to it as a tutor. Access is not tutoring. Robinson, Gormley, Ribeiro, and Loeb, in EdWorkingPaper 26-1451, put the same sentence in a title: “Access Is Not Enough.” Human support to raise engagement is a separate finding. It is not a finding this trial identified as a learning gain from dialogue.¹⁴⁸

¹⁴⁶Philip Oreopoulos and Nina Low, “One Click Away: AI Tutoring with Khanmigo in a Two-Year School Experiment,” EdWorkingPaper 26-1551 / NBER Working Paper 35620 (Annenberg Institute, 2026), <https://doi.org/10.26300/kner-hv33>. Pre-registered AEARCTR-0013519. Two-year cluster RCT, 18 Hamilton County, Tennessee middle schools; 6,902 student-terms; 53 clusters. ITT: +1.26 national percentile ranks per term (SE 0.60); ~0.040 population SD per term; stacked ~0.06 SD per school year, ~0.08 in year 2; implied full-year participation in year 2 ~0.14 SD. TCAP registered as co-primary, **not yet incorporated** as of the August 2026 paper. Engagement figures as in the text. Cost ~\$15 per student per year. A 2026 interface change that opened Khanmigo by default did not raise messaging rates in their data.

¹⁴⁷Taryn Eames, Emma Brunskill, Bogdan Yamkovenko, Kodi Weatherholtz, and Philip Oreopoulos, “Computer-Assisted Learning in the Real World: How Khan Academy Influences Student Math Learning,” *Proceedings of the National Academy of Sciences* 123, no. 10 (2026). Cited by Oreopoulos and Low: Khanmigo magnitudes resemble Khan Academy practice without the AI tutor.

¹⁴⁸Carly D. Robinson, David Gormley, Ana Trindade Ribeiro, and Susanna Loeb, “Access Is Not Enough: Human Support Improves Engagement with AI Tutoring,” EdWorkingPaper 26-1451 (Annenberg Institute, 2026). Cited in Oreopoulos and Low. Adaptive-literacy finding as reported there: dedicated time still produced only minutes of weekly use; adding human tutors raised usage without raising achievement. Engagement failures are not unique to Khanmigo.

Read the design against Bastani and against Kestin, because the collapse the last section warned about will happen here if it is not blocked. Khanmigo is a tutor app in Chapter 2's table: Socratic refusal, curriculum-grounded, rostered, \$15 a head. It is not GPT Base. It is not PS2 Pal. It is not an agent. The ITT is small, in the CAL band Escueta and colleagues described in 2020, not in the human-tutoring band Nickow and colleagues described in 2024. The authors' own interpretation is the one this chapter will carry: the binding constraint is **student-initiated engagement**, not model capability. A default-open chat window did not fix it. Treating non-use as a model-quality success—"the AI was available, therefore we delivered tutoring"—is how a district spends \$15 a student and reports a 0.06 SD gain as if it had bought Bloom. It bought practice. Some students messaged. Most mistake sessions did not include a message. Chat volume, once mastery is in the regression, is not the predictor.

Cost belongs next to the coefficient, because the popular story will say "cheap tutoring." About \$15 per student per year, inside blocks that were already staffed, is cheap compared with human high-dosage tutoring. It is also the price of a license, not the price of a learning gain of 0.288 SD. Nickow, Oreopoulos, and Quan's human-tutoring pool is an order of magnitude larger than this ITT. Escueta's CAL band is where this ITT lives. The honest sentence is not "we bought a tutor for fifteen dollars." The honest sentence is: we bought structured practice, with a coach-not-tell window most mistake sessions did not use, and MAP moved about 0.06 to 0.08 SD per year, more in year 2, with an implied 0.14 SD if year-2 participation were full-year, and the state test is not in the paper yet. A district that wants human-tutoring effects still has to buy human tutoring, or a session architecture that actually produces engagement, or both. Opening the window by default was an interface bet. The message data rejected it. Robinson and colleagues' "Access Is Not Enough" is the adjacent finding a procurement officer should read before the next license: dedicated time on an adaptive literacy platform still produced only minutes of weekly use; adding human tutors raised usage without raising achievement in that study. Engagement failures are not unique to Khanmigo. They are what happens when the design assumes the student will initiate the cognitive work the product is not requiring.

Two further cautions, then a transition. TCAP is still outstanding. Summer persistence was underpowered. Year-scale unaided skill on an independent state test is **not yet known**, and this is the trial that was built to speak to it. When TCAP is incorporated, this chapter's number may move. Until then, MAP is the outcome, the ITT is small, and the mechanism the authors name is practice plus thin talk, not Socratic dialogue at scale. Replication in other districts, with other RTI blocks, with other engagement designs, is unshown. That is a hole. Filling it with Pal's 0.73-to-1.3 SD, or with Bastani's practice +127 percent, is the averaging this chapter exists to make expensive.

A low-resource field RCT with a different design—session, prompts, teacher—produced a larger English coefficient. The treatment was not "give them Copilot." The treatment was the lab.

De Simone et al. 2025: Copilot is not the treatment; the lab is

Martin De Simone, Federico Tiberti, Maria Barron Rodriguez, Federico Manolio, Wuraola Mosuro, and Eliot Jolomi Dikoru published, as World Bank Policy Research Working Paper 11125 in 2025,

one of the few developing-country field RCTs of a public large language model.¹⁴⁹ The popular story will file it under “Copilot raises English by a quarter of a standard deviation.” The paper’s own design will not let that filing stand. Copilot is the model. The treatment is Copilot plus teacher-written tutor prompts plus supervised lab time plus curriculum alignment.

Nine public senior-secondary schools in Benin City, Nigeria. After-school Microsoft Copilot (GPT-4), twelve 90-minute sessions over six weeks, first-year senior-secondary students, English curriculum. Teachers trained three days. Sessions used **science-of-learning prompts**—retrieval, elaborative interrogation, concrete examples, the Mollick and Mollick 2023 library—rather than answer-giving.¹⁵⁰ Students often paired on one computer. Endline was pen-and-paper. Students were taught that Copilot hallucinates and that they remain responsible for checking. That last sentence is a plant for a later chapter on judgement. It is also a design feature of *this* treatment: the evaluative loop was assigned to the student in the session architecture, not hoped for as a personality trait.

Intention-to-treat, school fixed effects, baseline-controlled: composite assessment **0.31 SD** (SE 0.068); IRT **0.263 SD**. English, the primary outcome, **0.238 SD**. AI knowledge **0.309 SD**. Digital skills **0.139 SD** ($p < 0.10$). The regular third-term English exam, broader than the six-week content, **0.206 SD**. Heterogeneity: larger for females and for higher-baseline and higher-SES students. Attendance about 72 percent; about 0.03 SD per additional day (dose-response). Attrition was non-trivial; Lee bounds remain positive.

The authors convert English 0.24 SD to about 1.5 equivalent years of typical Nigerian schooling on the Evans and Yuan metric. That is a **conversion, not a longitudinal follow-up**. No year-later test is in the paper. Do not write “a year and a half of learning in six weeks” as if a follow-up had been run. Write: on a published conversion from other Nigerian evidence, 0.24 SD is a large schooling-equivalent; the students were not followed a year later; the regular term exam, which was run, moved 0.206 SD.

The equity sting is in the heterogeneity, and it belongs in the body, not in a diversity paragraph at the end of the book. Effects were larger for females—an encouraging split—and larger for higher-baseline and higher-SES students. A “personal tutor for every child” that widens gaps is an empirical possibility, not a slogan. Copilot and computer-science process studies, later in this chapter, will suggest that metacognitively weaker students may be harmed by autocomplete. Nigeria’s split is not harm. It is a larger gain where advantage already sits. The two findings are not the same. Together they are a warning against treating “access” as an equalizer. The lab was supervised. The prompts were written. The pairs shared a machine. The students who arrived with more still left with more, on these outcomes. That is the record of this trial, not a theory of justice.

¹⁴⁹Martin De Simone, Federico Tiberti, Maria Barron Rodriguez, Federico Manolio, Wuraola Mosuro, and Eliot Jolomi Dikoru, “From Chalkboards to Chatbots: Evaluating the Impact of Generative AI on Learning Outcomes in Nigeria,” Policy Research Working Paper 11125 (World Bank, 2025). RCT, nine public senior-secondary schools, Benin City. ITT (school FE, baseline-controlled): composite 0.31 SD (SE 0.068); IRT 0.263 SD; English (primary) **0.238 SD**; AI knowledge 0.309 SD; digital skills 0.139 SD ($p < 0.10$); regular third-term English exam **0.206 SD**. Evans and Yuan conversion of English 0.24 SD to ~1.5 equivalent years is a **conversion, not a longitudinal follow-up**. Attrition non-trivial; Lee bounds remain positive. Full PDF fetched.

¹⁵⁰Ethan Mollick and Lilach Mollick, “Assigning AI: Seven Approaches for Students, with Prompts,” arXiv:2306.10052 (2023). Prompt library used in the Nigeria RCT (retrieval, elaborative interrogation, concrete examples). A prompt library, not a trial.

Put Nigeria next to Tennessee and the design claim is no longer a slogan. Same model class, loosely: GPT-4-family tools. Different pedagogies. Khanmigo: coach-not-tell, one click away, existing RTI, thin messaging, ITT in the 0.06–0.08 SD/year band, looks like practice without the chatbot. Nigeria: scheduled after-school sessions, teacher-written science-of-learning prompts, three-day teacher training, paired computers, hallucination-checking assigned, pen-and-paper endline, 0.238 SD on English, 0.206 SD on the term exam. Access is not the treatment. The session is. The prompt is. The adult in the room is. Chapter 7 will spend the teacher as orchestrator. This chapter only needs the causal implication: every positive generative-AI *field* result in the 2022–26 record either embeds the model in a teacher-run session or helps the tutor. Khanmigo without engagement looks like software, not a tutor. Nigeria looks like a lab. Pal looks like a course staff’s physics in a prompt. GPT Tutor looks like a system prompt with a teacher’s mistake library. The model is not the intervention. The wrapper is.

Mollick and Mollick’s 2023 “Assigning AI: Seven Approaches” is a prompt library, not a trial. Nigeria is the trial that used the library. Retrieval, elaborative interrogation, concrete examples: those are science-of-learning moves that happen to be implementable as system prompts. They are also, without a teacher who was trained for three days and a lab that met twelve times, just text in a box. Pairing on one computer is a resource constraint and a pedagogical accident worth naming. Two students, one Copilot, a prompt that asks them to retrieve rather than to paste, an adult circulating: that is not “every child has a personal tutor.” It is closer to a small group with a constrained tool. The 0.238 SD on English and the 0.206 SD on the term exam are the ITT of that bundle. Strip the bundle in the citation—“Nigeria shows Copilot works”—and the coefficient is being asked to travel without its treatment. The term exam matters because it was not written by the research team to match the six weeks. It moved 0.206 SD. That is still not a year-later follow-up. It is more than a researcher-designed endline alone. Attendance at 72 percent, with a dose-response of about 0.03 SD per additional day, is the engagement fact Nigeria *does* have: the session was scheduled, and days present predicted gains. Khanmigo’s median student messaged on one-third of practice days. Different designs, different engagement, different coefficients. That is the chapter.

Writing and inquiry follow the same split: better artefact, thinner thinking, unless the student keeps the generative work.

Writing and inquiry: better artefact, thinner thinking

Mathematics is not a special case. The performance/learning split survives the move to writing and to scientific inquiry. The artefact improves, or feels cheaper to produce. Knowledge, transfer, and the quality of argument often do not. When they do, the gain tracks *how* the model was used—explanation versus generation—not the mere presence of a window.

Yizhou Fan, Luzhen Tang, Huixiao Le, Kejie Shen, Shufang Tan, Yueying Zhao, Yicen Shen, Xinyu Li, and Dragan Gašević, in the *British Journal of Educational Technology* in 2025, randomized 117 university students on a writing task across four supports: ChatGPT, human expert chat, writing-analytics tools, no extra tool.¹⁵¹ The ChatGPT group **improved essay scores more**. There were

¹⁵¹Yizhou Fan, Luzhen Tang, Huixiao Le, Kejie Shen, Shufang Tan, Yueying Zhao, Yicen Shen, Xinyu Li, and Dragan Gašević, “Beware of Metacognitive Laziness: Effects of Generative Artificial Intelligence on Learning

no significant differences in knowledge gain or transfer. There was no difference in post-task intrinsic motivation. Process mining is the result a designer can use. ChatGPT users formed tight revise $\leftrightarrow$ ask-AI loops with **fewer metacognitive processes—evaluation, orientation—**than the human-expert group. The authors name this **“metacognitive laziness.”** The full *BJET* PDF was publisher-blocked in one research fetch behind this book; an ERIC abstract was used. Additional coefficients are not invented here.¹⁵² The finding that survives that caution is the split the rest of this chapter already has: better product, no extra knowledge, fewer of the moves that would have been the learning. The essay is performance. The missing evaluation cycle is the skill not formed.

“Metacognitive laziness” is a name, not a diagnosis of character. The process-mining result is that evaluation and orientation—Facione’s words, in a later chapter, for the judgement loop—were less frequent when the support was ChatGPT than when it was a human expert. The model made the next revision cheap. Cheap revisions are performance. They are not, on this evidence, more knowledge or more transfer. Intrinsic motivation after the task did not differ. The students were not, as a group, more inspired. They had a better essay and a thinner process. A campus that grades the essay and never sees the process has just built Bastani’s practice window for prose.

Zara Contractor and Germán Reyes, in IZA Discussion Paper 18792, July 2026—a **working paper, not a journal article** as of the access date—report a proctored in-person RCT: undergraduates learn an unfamiliar topic and write an analytical essay with or without off-the-shelf generative AI, then take unaided tests immediately and one week later.¹⁵³ Abstracted results as published: AI access raises immediate knowledge-test scores by **0.27 SD**, and the gains **persist one week later**. Essay quality changes little *while AI is available*, but **style and clarity about 0.30 SD** and **relevance about 0.26 SD** improve on the delayed unaided essay. Gains are larger for **augmentation users—**AI to explain concepts—than for **automation users—**AI to generate text—whose short-run quality gains vanish when AI is removed. Mechanism, as the authors describe it: time shifts from drafting toward reading and searching; higher enjoyment. The full IZA PDF fetch was partial. Coefficients in this chapter are from the abstract and the IZA page. Treat them as working-paper numbers, not as a settled journal result.¹⁵⁴

The augmentation/automation split is the empirical warrant for a claim later chapters will spend as assessment design: grade *how* AI was used, not only the essay. Augmentation keeps the generative work, or enough of it, inside the student. Automation hands the generative work to the model and, when the model is taken away, the short-run quality is gone. Contractor and Reyes is the writing analogue of Bastani’s Base/Tutor split, with a delayed unaided essay instead of a closed-book math exam, and with a one-week persistence on a knowledge test that Bastani did not run. It is also, still, a working paper. A book that treats it as *PNAS* is doing the same inflation this chapter refuses for

Motivation, Processes, and Performance,” *British Journal of Educational Technology* 56, no. 2 (2025): 489–530, <https://doi.org/10.1111/bjet.13544>. N = 117. ChatGPT group improved essay scores more; no significant differences in knowledge gain or transfer; fewer metacognitive processes (evaluation, orientation).

¹⁵²Full *BJET* PDF publisher-blocked (Cloudflare) in one fetch behind this book; ERIC abstract (EJ1460793) used. Do not invent additional coefficients.

¹⁵³Zara Contractor and Germán Reyes, “Experimental Evidence on the Learning Impact of Generative AI,” IZA Discussion Paper No. 18792 (July 2026), <https://docs.iza.org/dp18792.pdf>. **Working paper, not a journal article** as of 30 August 2026. Abstracted results: immediate knowledge-test +0.27 SD, persisting one week; delayed unaided essay style ~0.30 SD and relevance ~0.26 SD for augmentation users, not automation users.

¹⁵⁴Full IZA PDF fetch was partial; coefficients from the abstract and IZA page. Treat as working-paper numbers.

vendor ESSA-3.

Even at working-paper strength, the split is the one a syllabus can use. Augmentation: the student asks the model to explain, then writes. Automation: the student asks the model to write, then perhaps edits. Immediate knowledge-test +0.27 SD, still there a week later, is a different shape from Fan's null on knowledge and transfer, and a different shape from Bastani Base's exam loss. The delayed unaided essay moving on style and relevance for augmentation users, not for automation users, is the writing version of "who did the generative work." Short-run quality under automation vanishes when the model is removed. That sentence is why take-home essays as the only grade are a validity problem, and why process logs and oral defences will appear in Chapter 5 as assessment design rather than as integrity theater. The coefficients came from an abstract and an IZA page. They will be updated if the journal version moves them. They will not be rounded into "AI helps writing."

Matthias Stadler, Maria Bannert, and Michael Sailer, in *Computers in Human Behavior* in 2024, randomized 91 university students to ChatGPT-3.5 versus Google to research nanoparticles in sunscreen and write recommendations.¹⁵⁵ LLM users reported **lower cognitive load** and produced **lower-quality reasoning and argumentation**. Homogeneity of conclusions did not differ. This is the scientific-inquiry analogue of Bastani's ease without depth, and of Fan's better artefact without the metacognitive loop. The load went down. The argument got thinner. Fluency is not inquiry. A model that writes a confident environmental recommendation is performing. The student who accepted it is not, on this evidence, learning to argue.

Print the short-term versus later, or without-AI, table so the pattern is public. Nine rows. No tenth study invented.

The table is the chapter in one view. Practice and artefacts on the left. Unaided later tests, or their absence, on the right. A reader who wants a single "does AI work?" answer will have to refuse the table. A reader who wants a design rule can take it: do not grade the left-hand column as if it were the right.

Study	Immediate / with AI	Later / without AI
Bastani et al. 2025 math	Practice +48% / +127%	Exam -17% (Base) / ~0 (Tutor)
Kestin et al. 2025 physics	Large immediate post-test gain (0.73–1.3 SD quantile)	Not measured
Oreopoulos and Low 2026 Khanmigo	MAP ~0.06–0.08 SD/year (0.040/term)	TCAP pending; summer persistence underpowered
De Simone et al. 2025 Nigeria	Endline +0.238 English; +0.206 on term exam	No year-later follow-up
Fan et al. 2025 writing	Better essays	Knowledge gain/transfer ns

¹⁵⁵Matthias Stadler, Maria Bannert, and Michael Sailer, "Cognitive Ease at a Cost: LLMs Reduce Mental Effort but Compromise Depth in Student Scientific Inquiry," *Computers in Human Behavior* 160 (2024): 108386, <https://doi.org/10.1016/j.chb.2024.108386>. RCT, N = 91; ChatGPT-3.5 versus Google. Lower cognitive load; lower-quality reasoning and argumentation. Fetched at abstract/PDF-availability level in the pedagogy pass.

Study	Immediate / with AI	Later / without AI
Contractor and Reyes 2026 writing	Test +0.27 SD	Test persists 1 week; unaided essay style/relevance up for <i>augmentation</i> users
Kosmyna et al. 2025 EEG writing	Lower connectivity, low ownership	Session-4 LLM → Brain still weaker quoting/connectivity (preprint; next section)
Stadler et al. 2024 inquiry	Lower load	Weaker reasoning (same session)
LearnLM 2025	Immediate remediation $\approx$ human	Next-topic transfer +5.5 pp (same day, supervised; preprint; later section)

The pattern is not a metaphor. When the model does the generative work, session artefacts look better and later unaided skill often does not. When the model is constrained to hints, retrieval, or explanation, harm shrinks and sometimes immediate learning rises. Durable, year-scale skill formation under *unguided* student-facing chat is **not yet shown** at school scale. Khanmigo’s year-scale gains look like structured practice, not dialogue. Kestin’s delayed test does not exist. Nigeria’s year-later test does not exist. Bastani measured the next hour. Contractor and Reyes measured a week, in a proctored lab, as a working paper. That is the record. It is enough to reject “the homework looked great, so they learned.” It is not enough to write a law of year-scale formation.

The viral EEG paper is the one everyone will ask about. Cut it down to size, then put the stronger behavioral evidence back in front of it.

Kosmyna and the offloading fight — preprint, not population decline

Nataliya Kosmyna, Eugene Hauptmann, Yeo Tze Yuan, James Situ, Xian-Hao Liao, Azran Yazdani Beresnitzky, Iris Braunstein, and Pattie Maes posted, as arXiv:2506.08872, with an MIT Media Lab page dated 10 June 2025, the paper the internet decided was proof that ChatGPT atrophies the brain.¹⁵⁶ It is a **preprint, not a journal article**, as of 30 August 2026. Treat it as such. It is real. It is not a definitive developmental result. It should not be cited as “MIT proved ChatGPT atrophies the brain,” and it should not be asked to carry K–12 policy, and it should not be extrapolated onto Grok Bot.

The design, as fetched, not as paraphrased in viral coverage: 54 Boston-area adults, mostly undergraduates; three groups—LLM/ChatGPT-4o, search engine, brain-only; three sessions writing SAT-style essays under EEG (Enobio 32, dDTF connectivity). Session 4, **n = 18**: LLM users write brain-only; brain-only users write with LLM.

¹⁵⁶Nataliya Kosmyna, Eugene Hauptmann, Yeo Tze Yuan, James Situ, Xian-Hao Liao, Azran Yazdani Beresnitzky, Iris Braunstein, and Pattie Maes, “Your Brain on ChatGPT: Accumulation of Cognitive Debt When Using an AI Assistant for Essay Writing Task,” arXiv:2506.08872 (2025), MIT Media Lab page 10 June 2025, <https://www.media.mit.edu/publications/your-brain-on-chatgpt/>. **Preprint, not a journal article** as of 30 August 2026. N = 54 sessions 1–3; session 4 n = 18. Do not cite as “MIT proved ChatGPT atrophies the brain.” Full PDF fetched. Not K–12; not agentic systems; EEG connectivity is not a direct measure of “getting dumber.”

Fetches findings, not media paraphrases. Connectivity scaled with how much the tool did: brain-only strongest and most distributed; search intermediate; LLM weakest. LLM group: low perceived ownership; poor ability to quote their own essay. Session 1: 15 of 18 could not quote; 0 of 18 quoted correctly. Search and brain-only quoting was far better. NLP: more homogeneous named entities and n-grams in the LLM group; some essays close to default ChatGPT SAT answers. Session 4: LLM → Brain showed weaker alpha/beta connectivity and lingering quoting problems; Brain → LLM showed higher recall and re-engagement of occipito-parietal and prefrontal networks. Essays were still often scored acceptably by teachers and an AI judge. Another performance-versus-encoding dissociation. The artefact can pass. The student cannot quote it. Soderstrom and Bjork would recognize the split. Bastani's Base students, who did not perceive the loss, would recognize the phenomenology.

Limitations the viral coverage omitted, and which this chapter will not omit. Sample: **n = 54 / 18**. Convenience sample of elite-university students. Twenty-minute essays. EEG connectivity is not a direct measure of “getting dumber.” No delayed knowledge test of the kind Bastani or Contractor and Reyes use. Group assignment lasted months of *sessions*, not of schooling. Not K–12. Not an agentic system. The authors themselves call it a preliminary educational-implication study. A book that lets a 54-person preprint set national policy about childhood, or about unattended agents, is doing the same thing the popular story does with Pal: taking a designed, limited, interesting result and writing a civilization sentence on top of it.

The offloading claim is older than this preprint, and it is better supported by *behavioral* RCTs than by EEG. Risko and Gilbert, in *Trends in Cognitive Sciences* in 2016, define **cognitive offloading** as using physical action or the environment to reduce the internal processing a task requires: tilting the head, setting a reminder, using a calculator.¹⁵⁷ Offloading is not a moral failing. It is how human cognition scales. The research questions are what triggers it—often a metacognitive judgement that internal demand is high—and what it does to later unaided performance. The Google effect (Sparrow, Liu, and Wegner, 2011, *Science*): people remember *where* information is stored better than the information itself when they expect later access.¹⁵⁸ Calculators and GPS produce analogous trade-offs: performance up, unaided skill down if the offloaded component was the learning target. Generative AI extends offloading from memory and calculation to **reasoning, drafting, and evaluation**—the skills schools exist to form. That extension is the educational fact. The EEG is a small, preliminary, neural correlate of shallower encoding during LLM-assisted writing. It is consistent with the behavioral record. It is not the record.

The behavioral record, already spent in this chapter, is the load-bearing support:

1. Bastani et al. (2025): copying solutions → −17 percent unaided.
2. Fan et al. (2025): better product, no extra knowledge, fewer metacognitive moves.
3. Stadler et al. (2024): less load, worse scientific argument.
4. Prather and colleagues (2023–24), in the next section: novices shepherded or lost; weak metacognition harmed.

¹⁵⁷Evan F. Risko and Sam J. Gilbert, “Cognitive Offloading,” *Trends in Cognitive Sciences* 20, no. 9 (2016): 676–688.

¹⁵⁸Betsy Sparrow, Jenny Liu, and Daniel M. Wegner, “Google Effects on Memory: Cognitive Consequences of Having Information at Our Fingertips,” *Science* 333, no. 6043 (2011): 776–778.

Michael Gerlich, in *Societies* in 2025, reports a survey-based association of AI-tool use with cognitive offloading and lower self-reported critical thinking.¹⁵⁹ **Correlational.** Younger users more dependent. Not an RCT. Do not promote a survey to a trial.

Two 2026 papers will be cited in the live debate as if they had settled the nuance. They have not been read in full in the research this book stands on. Wang and Zhang (2026), *International Journal of Artificial Intelligence in Education* / Springer, N = 912, China/Europe/US: pedagogical partnerships activate vigilance and offloading together; a U-shaped curve; Zone 2—scattered, half-hearted AI use—worst. **Full text not fetched**; abstract and DOI-landing only. A 2026 *Behavioral Sciences* scoping review of 123 studies: offloading is not inherently maladaptive; maladaptive when learners do not verify, reflect, or retain ownership. **Full text not fetched.** A 2026 *Frontiers in Psychology* distinction between dependent and autonomous offloading, at abstract level: frequency of AI use is not the variable; whether the user retains the evaluative loop is.¹⁶⁰ Those three items are marked **UNVERIFIED** at full-text level. The conceptual distinction in the last of them is useful even before the psychometrics settle, because it matches the causal record this chapter *has* fetched: Base versus Tutor; automation versus augmentation; ChatGPT-on-inquiry versus Pal; access versus lab. Frequency is not the treatment. The loop is.

The live debate is a harm frame versus a nuance frame. Harm frame: Kosmyna, Bastani GPT Base, Stadler, Fan. Nuance frame: Wang and Zhang's U-shape, Contractor and Reyes's augmentation, GPT Tutor, PS2 Pal, Nigeria, LearnLM. The book's position, if it follows the causal record rather than the comment section, is not a tie. **Contact with models is not what dulls. Three conditions together are.** The offloaded process is the learning objective. The environment never measures the skill unaided. The student cannot detect model error. Designs that keep generation, retrieval, and checking inside the student—GPT Tutor, Nigeria's prompts, PS2 Pal, LearnLM's Socratic drafts, augmentation use in Contractor and Reyes—do not show the atrophy pattern on the outcomes they measured. That is the opposite of a general "AI equals dumber" law. It is also the opposite of a general "AI equals tutor" law. The three conditions are a design specification. Take-home essays as the only grade, with a model that will write them, with students who cannot catch a fluent error, is how you build the law the harm frame wants. Hint-only practice, unaided exams, and a student who has been taught that the model hallucinates, is how you refuse it.

Kosmyna's $n = 54 / 18$ does not become a population decline by being tweeted. It does not become a childhood finding by being about essays. It does not become an agent finding by being about ChatGPT-4o. Put it in the table. Quote the quoting failure. Name the preprint. Move.

What would actually make students "dumber," if that word must be used, is not contact. It is the three conditions, sustained, on the skill the course exists to form. Writing courses that only ever

¹⁵⁹Michael Gerlich, "AI Tools in Society: Impacts on Cognitive Offloading and the Future of Critical Thinking," *Societies* 15, no. 1 (2025): 6. Survey-based association. **Correlational**, not an RCT.

¹⁶⁰Wang and Zhang, "Pedagogical Partnerships with Generative AI in Higher Education," *International Journal of Educational Technology in Higher Education* (2026), <https://doi.org/10.1186/s41239-026-00585-x>. N = 912 as abstracted. **Full text not fetched**; abstract/DOI-landing only. Research files also name *IJAIED*/Springer for this item; outlet naming is **UNVERIFIED** pending the full text. *Behavioral Sciences* 2026 scoping review, 123 studies, <https://doi.org/10.3390/bs16071150>: offloading not inherently maladaptive; maladaptive when learners do not verify, reflect, or retain ownership — **full text not fetched**. "Not all cognitive offloading is equal: Distinguishing dependent and autonomous offloading to generative AI," *Frontiers in Psychology* (2026), <https://doi.org/10.3389/fpsyg.2026.1878629> — abstract-level; frequency is not the variable.

collect take-home prose; mathematics courses that only ever collect AI-assisted homework; inquiry tasks that accept a fluent recommendation as evidence of scientific thinking; programming courses that accept a green test suite as evidence that the novice can trace. In each case the offloaded process is the objective, the environment never measures the skill unaided, and the student need not catch the model's fluent error. Kosmyna's quoting failure is a twenty-minute version of that environment. Bastani's -17 percent is a school version. Fan's missing evaluation cycle is a process version. Stadler's thinner argument is an inquiry version. Prather's straying novice is a keystroke version. The EEG is the one that went viral because a brain picture is a better headline than a clustered standard error. The clustered standard error is the one a district can use.

Adjacent designs put a human on the loop. They converge. They still are not Hermes-class trials, and they still do not fill the holes.

Adjacent designs, Copilot, and what is not yet known

Three further studies converge on a different design than “give the child ChatGPT”: structured curriculum, misconception-level diagnostics, and—often—a human who either sits with the student or sits on the tutor. A fourth literature, Copilot in computer science, is not a multi-site RCT of conceptual learning. It is an interaction literature about novices who shepherd or stray. Then the holes. An honest chapter names them rather than filling them with analogy.

Owen Henkel, Hannah Horne-Robinson, Neeve Kozhakhmetova, and Adam Lee, at AIED 2024, report Rori, a WhatsApp AI math tutor, with about 500 Ghanaian students given an extra hour per week on a phone versus usual classes.¹⁶¹ Growth-score **Cohen's $d = 0.36$** ($p < 0.001$). Marginal cost about \$5 per student. Year-1 result. The authors flag caution. The design is not unguarded chat. It is a structured math tutor on a messaging app, extra time, a low-resource setting. The coefficient is in a serious band. The caution is the authors'. Year-1, about 500 students, WhatsApp, not a state test, not a year-later follow-up this chapter can cite. Put it next to Nigeria as a low-resource, session-shaped, curriculum-shaped result. Do not put it next to Pal as if $d = 0.36$ and 0.73 – 1.3 SD were the same object.

The LearnLM Team at Google DeepMind, with Eedi, posted arXiv:2512.23633 in 2025: an exploratory RCT, **preprint**, $N = 165$ students, five UK secondary schools, Eedi math platform.¹⁶² Expert tutors supervised every LearnLM draft—approve, edit, or replace. Tutors approved **76.4 percent** of drafts with zero or 1–2-character edits (74.4 percent zero edits in the message audit). Zero harmful messages. Five factual errors in 3,617 drafts (0.1 percent). Immediate mistake-remediation: supervised LearnLM **93.0 percent** versus human tutor **91.2 percent** versus static hints **65.4 percent**. Transfer to the next topic: LearnLM **66.2 percent** versus human **60.7 percent**; the authors report **+5.5 percentage points**, with 93.6 percent posterior probability that LearnLM

¹⁶¹Owen Henkel, Hannah Horne-Robinson, Neeve Kozhakhmetova, and Adam Lee, “Effective and Scalable Math Support: Experimental Evidence on the Impact of an AI Math Tutor in Ghana,” in *AIED 2024*, CCIS 2150 (Cham: Springer, 2024), <https://arxiv.org/abs/2402.09809>. About 500 students; Cohen's $d = 0.36$; about \$5/student; year-1; authors flag caution.

¹⁶²LearnLM Team, Google DeepMind, and Eedi, “AI Tutoring Can Safely and Effectively Support Students: An Exploratory RCT in UK Classrooms,” arXiv:2512.23633 (2025). **Preprint**. $N = 165$, five UK secondary schools. Every message gated. Session-level randomization; cumulative learning not identified. Bayesian reporting. Immediate remediation 93.0 percent versus human 91.2 percent versus static hints 65.4 percent; next-topic transfer +5.5 pp.

exceeds human. **Caveats, all of which stay in the body.** Session-level, not student-level, randomization of tutor versus LearnLM, so cumulative learning is not identified. Human-in-the-loop is *literal*: every message gated. Bayesian rather than frequentist reporting. Preprint. The qualitative finding that supervising tutors learned Socratic moves from the model is interview evidence from a handful of tutors, not a professional-development trial, and it will not be promoted here.¹⁶³ LearnLM is what a supervised, curriculum-platform, every-message-gated tutor looks like on immediate remediation. It is not what a student-facing agent looks like. It is not what unattended Grok Bot looks like. Citing it as “AI tutoring matches humans” without the gating is the Pal mistake in a British accent.

Rose Wang, Ana Teresa Ribeiro, Carly Robinson, Susanna Loeb, and Dorottya Demszky’s Tutor CoPilot is the other human-on-the-loop design: the model helps the *tutor*, not the child directly.¹⁶⁴ RCT, about 900 tutors and about 1,800 K–12 students. Tutors randomly given a real-time LLM that suggests expert-like moves. Students **+4 percentage points** more likely to master the session topic ($p < 0.01$); **+9 percentage points** for lower-rated tutors. Tutors asked more guiding questions and gave away fewer answers. About \$20 per tutor per year. A back-of-envelope on MAP is reported as **0.024 SD** for a full year—**explicitly not a direct experimental estimate**. Do not write 0.024 SD as a trial finding. Write: session mastery moved; the MAP conversion is arithmetic the authors labelled as such. The design implication is the one Chapter 7 will spend: the positive causal record includes helping the human tutor, not only wrapping a model for the child. CoPilot-for-tutors is not a student agent. It is a prompt at the adult’s elbow.

Computer science does not yet have the Bastani of Copilot. There is not, as of this access date, a large multi-site RCT showing Copilot raises *conceptual* CS learning. The better evidence is interaction studies, and they are not kind to the novice who needed the course.

James Prather, Brent Reeves, Juho Leinonen, Stephen MacNeil, and colleagues, in 2023, watched novice programmers with GitHub Copilot.¹⁶⁵ Two patterns: **shepherding**—steering autocomplete rather than writing—and **straying or drifting**—accepting incorrect suggestions into a debug rabbit hole. Students believed Copilot would generate correct code. Metacognitive difficulties were common. Prather and colleagues in 2024, laboratory sessions on greenfield assignments: 20 of 21 completed the task, but benefits accrued to students with **stronger metacognitive skills**; students with metacognitive difficulties were **negatively** affected.¹⁶⁶ Later Copilot process studies report less manual typing and less web search when Copilot is on—productivity, not learning. Trust-calibration work with upper-division students finds trust rises with experience *and* that students

¹⁶³LearnLM qualitative finding that supervising tutors learned Socratic moves from the model: interview evidence from $n = 5$, not a professional-development trial. Do not promote it.

¹⁶⁴Rose E. Wang, Ana Teresa Ribeiro, Carly D. Robinson, Susanna Loeb, and Dorottya Demszky, “Tutor CoPilot: A Human-AI Approach for Scaling Real-Time Expertise,” arXiv:2410.03017 (2024). About 900 tutors, about 1,800 K–12 students. +4 percentage points topic mastery ($p < 0.01$); +9 pp for lower-rated tutors; about \$20 per tutor per year. Back-of-envelope MAP **0.024 SD** for a full year — **explicitly not a direct experimental estimate**.

¹⁶⁵James Prather, Brent N. Reeves, Juho Leinonen, Stephen MacNeil, et al., “‘It’s Weird That It Knows What I Want’: Usability and Interactions with Copilot for Novice Programmers,” *ACM Transactions on Computer-Human Interaction* (2023). Shepherding and straying/drifting; metacognitive difficulties; overtrust in correctness.

¹⁶⁶James Prather et al. (2024), laboratory sessions on greenfield assignments: 20 of 21 completed the task; benefits accrued to students with stronger metacognitive skills; students with metacognitive difficulties were negatively affected. As reported in the pedagogy memo; not a large multi-site conceptual-learning RCT.

conclude a competent programmer is still required to finish tasks. The extended-use paper is arXiv:2509.13253; treat the trust-calibration details as labelled in the pedagogy memo, not as a classroom RCT.¹⁶⁷

Copilot is a worked-example generator running at keystroke speed. For novices it is expertise reversal in the wrong direction unless the course forces tracing, explaining, and writing without the model. Chapter 5 will spend that as a CS struggle policy. This chapter only needs the evidential grade: no large multi-site conceptual-learning RCT; process studies that look like Bastani’s copy-the-answer pattern at the speed of tab-complete; benefits concentrated among students who already have the metacognitive skill the course was supposed to form. Nigeria’s higher-baseline split and Copilot’s metacognition split are not the same finding. They rhyme. A tool that helps the already-able more, or harms the not-yet-able, is not an equalizer by being available.

UC San Diego’s introductory-course redesign, as told in a Microsoft customer story, is a curriculum response to that interaction literature, not a peer-reviewed trial of Cloud Agents: students first demonstrate manual competency, then use Copilot; faculty (Porter, Shah) research trust calibration; the campus talks about preparing students for “emerging agentic AI.”¹⁶⁸ Vendor customer story. Curriculum redesign around a coding assistant. Stanford CS221, in a later chapter’s documents, bans agentic repo-builders on the project. Those are course-level refusals of the last row, not estimates of conceptual learning under Copilot. They belong here only as a reminder that the practitioners who live with autocomplete have already split formation from performance, in the language of this book, without waiting for a multi-site RCT. The RCT is still missing. The split is already the honest local practice.

Those adjacent designs, plus Bastani Tutor, Pal, Nigeria, and Khanmigo-as-practice, are the generative-AI causal record as of 30 August 2026. The record is almost entirely chat and tutors. It is not agents. It is not WWC-grade. It is not year-scale unaided skill on a state test. It is not childhood. The holes, not hedges:

Year-scale unaided skill under student-facing generative tutors, on independent state tests, is not yet known. Khanmigo TCAP is still outstanding. Nigeria and Kestin did not follow a year later. Bastani measured the next hour. Contractor and Reyes measured a week, as a working paper.

Childhood (K–2) developmental RCTs of conversational agents on language, attention, and social cognition were not found. The cognitive-science extrapolation is possible. The generative-AI trial evidence is not. Chapter 4 will treat that band as a legal and developmental design specification, not as a missing coefficient this chapter is entitled to invent.

Reading and history lack causal estimates comparable to math. Direct generative-AI RCT evi-

¹⁶⁷Trust-calibration work with upper-division students under extended Copilot use: arXiv:2509.13253. Trust rises with experience; students conclude a competent programmer is still required to finish tasks. Not a classroom RCT of conceptual learning. Later process studies (e.g., brownfield coding, ICER 2025) report less manual typing and less web search — productivity, not learning.

¹⁶⁸Microsoft Customer Stories, University of California San Diego / GitHub Copilot, <https://www.microsoft.com/en/customers/story/26539-university-of-california-san-diego-github-copilot>. Faculty (Porter, Shah) redesigned intro courses so students first demonstrate manual competency, then use Copilot. **Vendor customer story**, not a peer-reviewed trial of Cloud Agents. Stanford CS221 Spring 2026 bans agentic repo-builders on the project — a course policy, documented in later-chapter sources, not an RCT.

dence on learning to read—decoding, fluency, comprehension—in 2022–26 is thin. Pre-LLM ITS for reading comprehension showed gains on researcher-designed measures and mixed standardized-test results in WWC review, which is a caution that software gains often shrink on independent tests, not a generative-AI finding.¹⁶⁹ History has almost no subject-specific RCT as of this access date.

Transfer of “evaluate the chatbot” skill to non-AI critical thinking is unshown. There is not yet an RCT showing that a K–12 “evaluate the chatbot” curriculum raises transfer of critical thinking to non-AI tasks. That is a hole, not a license to skip the practice, and not a license to claim the practice has been shown to transfer.

Whether proactive, multimodal, memory-bearing tutors fix Khanmigo-style non-use is open. Authors on both sides treat it as open. A default-open chat window did not, in Oreopoulos and Low’s Year-2 data, raise messaging.

Teacher professional development → student learning causal chain for generative AI is unshown. Reviews stop at teacher self-efficacy. The closest causal analogue in the adjacent CAL literature is Oreopoulos, Gibbs, Jensen, and Price (2024): teaching teachers to structure computer-assisted practice raises both usage and achievement.¹⁷⁰ That is PD that changes session architecture, not PD that demos ChatGPT. It is not a generative-AI trial.

Equity. Nigeria found larger effects for higher-baseline and higher-SES students, and for girls. Copilot studies suggest metacognitively weaker students may be harmed. A “personal tutor for every child” that widens gaps is an empirical possibility, not a slogan.

Replication. Kestin outside a highly staffed university physics course: unshown. Bastani outside one Turkish high school’s four sessions: unshown.

Kosmyrna-style neural measures in a design that also measures delayed learning do not exist. The preprint should not be asked to carry developmental policy.

Autonomous multi-step agents (Hermes-class) as opposed to single-chat tutors: **no identified learning RCT** as of this access date. Agent trajectories that plan, browse, and submit work on the student’s behalf are a different offloading regime. Chatbot studies do not measure them. Do not extrapolate MIT EEG onto Grok Bot. Do not use this chapter’s table as if the last row of Chapter 2 had been randomized.

WWC-grade evidence on named generative products does not exist yet. There is **no WWC intervention report specific to generative AI tutors**.

Optimal mix of AI-forbidden and AI-required tasks by grade is a theoretically motivated hypothesis, not a fitted schedule. Chapter 5 will offer a two-lane rule and will label it as such.

¹⁶⁹Kay Wijekumar, Bonnie J. F. Meyer, and Pui-Wa Lei, ITSS structure-strategy ITS (2012 and related): gains on researcher-designed measures; mixed standardized-test results in WWC review. Pre-LLM caution that software gains often shrink on independent tests. Not a generative-AI finding. Direct genAI RCT evidence on learning to read in 2022–26 is thin compared with math.

¹⁷⁰Philip Oreopoulos, Chloe R. Gibbs, Michael Jensen, and Joseph Price, “Teaching Teachers to Use Computer Assisted Learning Effectively,” NBER Working Paper 32388 (2024); forthcoming, *Journal of Human Resources*. Causal PD for CAL structure raises usage and achievement. Best analogue for generative-AI teacher PD; not a generative-AI trial. Teacher-PD reviews for generative AI (Li, Wang, and Chen, *Teaching and Teacher Education* 168 [2025]: 105219, and related) stop at teacher knowledge, attitudes, and self-efficacy, almost never student learning.

What this chapter must not be taken to have shown, restated so a later chapter cannot launder it. Year-scale unaided skill under student-facing generative tutors, on independent state tests: **not yet known**. Chatbot studies did not measure agents. Kestin is not evidence that ChatGPT beats classroom teaching. Khanmigo ITT is not “AI tutoring works at 0.07 SD.” SchoolAI’s 28 percent and MagicSchool ESSA 3 are not RCTs. Kosmyna is a preprint, $n = 54 / 18$. Wang and Zhang and the 123-study review: full texts unfetched. LearnLM is a preprint with human-in-the-loop on *every message*. Contractor and Reyes is a July 2026 working paper. No childhood conversational-agent RCTs on language, attention, or social cognition were found. Tutor CoPilot’s 0.024 SD MAP figure is back-of-envelope. Fan’s additional coefficients were not taken from a blocked PDF.

The causal record is age-blind in the wrong way. A 17-year-old in a Turkish math session, a Harvard undergraduate in Pal, a Tennessee middle-schooler who does not message, a Nigerian first-year in a supervised English lab, a university student writing with ChatGPT, a novice with Copilot, a Boston-area adult under EEG: these are not one developmental object. Age and law are design specifications. That is Chapter 4.

Before that turn, one more refusal. This chapter is not a meta-analysis. It does not pool Bastani, Kestin, Oreopoulos and Low, De Simone, Fan, Stadler, Henkel, LearnLM, and Tutor CoPilot into an average effect of “AI.” The doses differ, the outcomes differ, the comparators differ, the wrappers differ, and two of the papers are preprints. Pooling them would be the averaging the first section asked the reader to stop. The synthesis is qualitative and design-level: when the model generates, later unaided skill often does not rise with the artefact, and sometimes falls; when the model is constrained and a human is in or on the loop, harm shrinks and immediate learning sometimes rises; at school scale without engagement, the coefficient looks like practice; no identified trial has yet randomized a Hermes-class agent against a learning outcome. That is enough to write the rest of the book. It is not enough to write a product endorsement, a product ban, or a civilization thesis. The exam after the chat window closes is still, in 2026, the measurement most systems are not running.

The argument this chapter is entitled to, after the holes are left standing, is narrower than the popular story and stronger than a shrug. Immediate AI-aided performance is the wrong success metric. Unaided delayed tests are the right one. Unguarded answer-giving GPT can harm subsequent math learning. Hint-and-grounded-solution GPT avoids the harm but did not, in Bastani, beat no-AI on the exam. Pedagogically engineered tutors can raise immediate conceptual learning, even versus strong active learning—a design result, not a ChatGPT-default result, delayed outcome unmeasured. At public-school scale, engagement, not IQ of the model, bound Khanmigo’s effect; the ITT looks like practice without the chatbot. In a low-resource system, teacher-guided Copilot sessions with science-of-learning prompts produced about a quarter of a standard deviation on English and on the regular term exam. Writing with ChatGPT often improves the artefact and not the knowledge; using AI to explain rather than to draft is associated, in a working paper, with delayed unaided gains. Cognitive offloading is a real, old mechanism. The MIT EEG study is a small preprint consistent with shallower encoding during LLM-assisted writing, not proof of population-level cognitive decline. Novice programmers with Copilot show shepherding and straying; benefits concentrate among students who already have metacognitive skill. Human-in-the-loop that *helps* is specific: grounded prompts, attempt gates, circulating teachers, tutor-facing copilots—not a warm body in the building. Design is the treatment. Guardrails are the treatment. Default ChatGPT is a crutch

students do not notice. There is no Hermes-class learning RCT. There is no WWC generative-tutor report. The exam after the chat window closes is the outcome. The wrapper is the intervention. The model class is not.

Chapter 4

A Tool Lawful at 17 Can Be Illegal and Developmentally Wrong at 8



Figure 5: Keys, a locked cabinet, sealed letters, and an open book on a desk.

Three instruments, not a classroom. The Children’s Online Privacy Protection Act draws a bright line under thirteen: operators of sites and services directed to children, or with actual knowledge they collect personal information from children under that age, must give notice and obtain verifiable parental consent before collection, use, or disclosure.¹⁷¹ UNESCO’s 2023 *Guidance for*

¹⁷¹Children’s Online Privacy Protection Act, 15 U.S.C. sections 6501 et seq.; COPPA Rule, 16 C.F.R. Part 312.

generative AI in education and research independently recommends age thirteen for independent classroom use of generative-AI tools, and says that under-thirteen and primary use of untested tools is “not appropriate.”¹⁷² Broward County Public Schools, by School Board Resolution No. 26-132 on 10 March 2026, gives kindergarten through grade 5 **no** district-approved student-facing AI tools.¹⁷³ Three floors. One band. A product that is lawful and useful at seventeen can be illegal at eight, and developmentally wrong at both ages for different reasons.

Age is not a marketing funnel. It is a design specification. Developmental limits, legal gates, and pedagogy change together. The number thirteen that appears in COPPA, in UNESCO, and in ChatGPT’s consumer terms of service is a *legal* line. So is eighteen: Claude’s consumer gate, Character.AI’s open-chat gate after November 2025, the age at which FERPA rights transfer to the student. Neither number is a developmental finding. Design to the stricter of law and development. Accessibility is not a later chapter. CAST’s UDL Guidelines 3.0, IDEA, Section 504, and ADA Title II apply inside every band.¹⁷⁴ There is no United States federal student-AI statute. What follows is the stack that already binds.

The legal stack is the design spec, not the appendix

COPPA, 15 U.S.C. sections 6501 et seq., and the COPPA Rule at 16 C.F.R. Part 312, is the federal ceiling for anyone who collects personal information from a child under thirteen. The first major Rule amendments since 2013 were published in the *Federal Register* on 22 April 2025. They took effect on 23 June 2025. The general compliance date is 22 April 2026.¹⁷⁵ The amendments tighten notices, security, deletion, and retention. They require **separate** verifiable parental consent for third-party disclosures and targeted advertising. The Federal Trade Commission has stated that disclosure “for monetary or other consideration, for advertising purposes, or to train or otherwise develop artificial intelligence technologies” is **not integral** to the service.¹⁷⁶ Model training on a child’s chat or homework is not a free rider on school consent. It needs its own parental yes, or it does not happen.

The Commission did **not** finalize the proposed school-authorization / ed-tech exception, citing pending FERPA rulemaking (RIN 1875-AA15). That NPRM had not been published as of the FTC’s April 2025 COPPA preamble.¹⁷⁷ Existing FTC ed-tech guidance still applies: a school may act as

¹⁷²UNESCO, *Guidance for Generative AI in Education and Research* (7 September 2023) (Fengchun Miao and Wayne Holmes).

¹⁷³Broward County Public Schools, School Board Resolution No. 26-132 (10 March 2026); browardschools.ai, accessed 30 August 2026.

¹⁷⁴CAST, *UDL Guidelines 3.0* (30 July 2024); Individuals with Disabilities Education Act, 20 U.S.C. sections 1400 et seq.; Rehabilitation Act section 504; ADA Title II, including the 2024 web rule (WCAG 2.1 AA).

¹⁷⁵FTC, COPPA Rule amendments, 90 Fed. Reg. 16918 (22 April 2025); Commission vote 16 January 2025; effective 23 June 2025; general compliance date 22 April 2026.

¹⁷⁶FTC preamble to the 2025 COPPA Rule: disclosure “for monetary or other consideration, for advertising purposes, or to train or otherwise develop artificial intelligence technologies” is not integral to the service.

¹⁷⁷FTC did not finalize the proposed school-authorization / ed-tech exception, citing pending FERPA rulemaking, RIN 1875-AA15. NPRM not published as of the April 2025 COPPA preamble. Existing FTC ed-tech guidance: school as parent’s agent only for the contracted educational service; no commercial use, advertising, or model-training on that data.

the parent’s agent **only** for the contracted educational service. No commercial use. No advertising. No model-training on that data. “Personal information” now includes biometric identifiers—voiceprints, faceprints, and the rest—and additional government IDs. Written retention policy is required. Indefinite retention is out. “We need it forever to improve the algorithm” does not satisfy the Rule. A written information-security program is required. These are not brochure lines. They are the conditions under which a student-facing account, a persistent identifier, a voice log, a photo, or a chat transcript from a child under thirteen is lawful.

FERPA, 20 U.S.C. section 1232g and 34 C.F.R. Part 99, protects “education records”: information directly related to a student and maintained by the agency or institution, or by a party acting for it. It applies at every band that receives Department of Education funds. The default is no disclosure of personally identifiable information without consent, except listed exceptions. The path vendors use is the **school-official exception**, 34 C.F.R. section 99.31(a)(1)(i). A contractor may receive PII only if it (1) performs a function the school would otherwise use employees for, (2) meets the school’s “school official / legitimate educational interest” criteria, (3) is under the school’s **direct control** with respect to the use and maintenance of the records, and (4) uses the records only for authorized purposes and does not re-disclose.¹⁷⁸ The Student Privacy Policy Office’s vendor FAQ is blunt: PII disclosed under this exception may be used **only** for purposes authorized by the school. Training a vendor’s commercial model on student chat or student work is not an authorized school function unless the contract says so—and most district policies now forbid it.

There is no product that is “FERPA compliant” in the abstract. There is a written contract, a configuration, and a practice. Consumer ChatGPT, Claude, and Gemini accounts cannot satisfy the school-official exception. They are not under the school’s direct control. OpenAI’s statement that ChatGPT Edu, Enterprise, Teachers, and the API do not train on organizational data “by default” is a product-tier claim, not a legal status.¹⁷⁹ Rights transfer to the student at eighteen **or** upon enrollment in postsecondary education, whichever comes first. An eligible student is the rights-holder. Parents of tax-dependent eighteen-year-olds may still access records under a FERPA exception. That is not a license to treat a high-school senior’s chat log as a family shared drive.

PPRA, 20 U.S.C. section 1232h and 34 C.F.R. Part 98, gives parents inspection rights over instructional materials used in surveys and evaluations, and prior consent or notice and opt-out—depending on the provision—for surveys that probe political affiliation, mental health, sex, illegal behavior, family relationships, religion, or income. Rights transfer at eighteen or emancipation. An AI tutor that administers affective or mental-health probes is in PPRA’s neighborhood whether or not the vendor calls it a survey.

IDEA, 20 U.S.C. sections 1400 et seq., requires that an IEP be developed by a qualified team. AI may support *materials*. It may not write or determine the IEP. The Department of Education’s May 2023 AI report flags that AI recommendations must be able to comply with IDEA.¹⁸⁰ New

¹⁷⁸Family Educational Rights and Privacy Act, 20 U.S.C. section 1232g; 34 C.F.R. Part 99, especially section 99.31(a)(1)(i) (school-official exception) and section 99.33 (redisclosure). Eligible student: 18 or attending postsecondary. Student Privacy Policy Office vendor FAQ: PII under this exception only for purposes authorized by the school.

¹⁷⁹OpenAI business-privacy page: ChatGPT Edu / Enterprise / Teachers / API data not used for training “by default”—product-tier claim, not a legal status. Consumer accounts cannot satisfy the school-official exception.

¹⁸⁰U.S. Department of Education, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023).

York City Public Schools, in March 2026 guidance, puts IEP and 504 development on a red list, alongside placement, discipline, final grades, surveillance, and counseling.¹⁸¹ Section 504 of the Rehabilitation Act and ADA Title II—including the 2024 web rule requiring WCAG 2.1 AA for public entities—require accessible materials and tools. Broward requires WCAG 2.1 AA as a condition of AI-tool approval. CIPA requires E-rate schools to filter harmful content. NYCPS and the UK Department for Education both map generative-AI filtering onto this duty. Open-web agents that bypass school filters conflict with CIPA practice. That is not a future problem. It is why no documented United States public-school program of unsupervised student agents on the live web was found as of 30 August 2026.

State overlays sit on top of the federal floor and do not wait for COPPA to turn off. New York Education Law section 2-d and 8 NYCRR Part 121 protect student, teacher, and principal PII. They require a NIST-aligned data-security policy, a Data Protection Officer, a Parents' Bill of Rights in every third-party contract, a data-security-and-privacy plan, staff training, and breach notice. The duties apply to districts, charters, BOCES, and UPK. NYCPS maps every AI tool that processes student data through ERMA, plus Chancellor's Regulation A-820. A December 2024 AI add-on requires vendors to disclose AI capabilities, **prohibit student-data model training**, and meet transparency requirements.¹⁸² New York 2-d does not care that the student turned thirteen.

California's SOPIPA, Cal. Bus. & Prof. Code section 22584, operative 1 January 2016, forbids operators of K-12-purpose sites and apps to targeted-advertise using covered information, amass a non-educational student profile, sell student information, or disclose covered information except as listed. Operators must use reasonable security and delete on school request. Companion A.B. 1584 (Ed. Code section 49073.1) governs third-party student-record contracts: LEA control, no targeted ads. A.B. 1159 (2025-26) would extend SOPIPA-like rules to higher education (HESIPA, proposed operative 1 July 2027) and **ban using covered student information to train generative AI**. As of 30 August 2026 it sat at Senate third reading / floor process. It is **not enacted**. Cite it as pending, not as law.¹⁸³

Texas Education Code section 32.155 and related Chapter 32 provisions require operators to maintain reasonable security of "covered information"; TEA-approved operators must use TSDS unique identifiers to mask PII and adhere to a state data-sharing agreement. TRAIGA (H.B. 149), effective 1 January 2026, requires a DIR AI code of ethics for state and local government, including districts that procure or deploy AI, NIST AI RMF alignment for "heightened scrutiny" systems, and public-facing AI disclosure. H.B. 3512 and S.B. 1964, effective 1 September 2025, require AI training for designated school employees and a government AI inventory and ethics frame.¹⁸⁴ Illinois SOPPA, 105 ILCS 85, rewritten effective 1 July 2021, is K-12 only: a written DPA before any covered information moves to an operator; no targeted ads; no sale or rent; no non-school pro-

¹⁸¹New York City Public Schools, AI Guidance (March 2026): red list includes placement, discipline, IEP/504 development, final grades, surveillance, counseling.

¹⁸²N.Y. Educ. Law section 2-d; 8 NYCRR Part 121; NYCPS ERMA; Chancellor's Regulation A-820; December 2024 AI add-on (disclose AI capabilities; prohibit student-data model training; transparency).

¹⁸³Cal. Bus. & Prof. Code section 22584 (SOPIPA), operative 1 January 2016; A.B. 1584, Ed. Code section 49073.1. A.B. 1159 (2025-26): Senate third reading / floor process as of 30 August 2026; **not enacted**. Proposed HESIPA operative 1 July 2027; would ban using covered student information to train generative AI.

¹⁸⁴Tex. Educ. Code section 32.155; TRAIGA, H.B. 149, effective 1 January 2026; H.B. 3512 and S.B. 1964, effective 1 September 2025.

file; breach notice to the school, then parents; public posting of vendor lists, data elements, and subcontractors; deletion when no longer needed.¹⁸⁵ The National Data Privacy Agreement v2 (Student Data Privacy Consortium, 2024) is the de-facto interstate DPA template. It is a contract floor, not a statute. Prince George's County Public Schools' AI Declaration Addendum is a specimen of what a serious district now asks: name the models; say what is local versus cloud; certify FERPA, COPPA, and state-law confidentiality; **no retention of student data or PII in generative models**; no use beyond the stated educational purpose without district **and** parental approval; bias and harm audit on request; list of non-US subprocessors; how to turn AI features off.¹⁸⁶

Florida SOPIPA and State Board Rule 6A-1.09550, plus FIPA, are cited by Broward as the state overlay on FERPA and COPPA for K-12 vendor data. The Florida text was not independently fetched in the research pass behind this book. Treat Broward's citation as district practice, not as a substitute for the statute.¹⁸⁷

If the product ships globally, two further instruments bind. Regulation (EU) 2024/1689, the AI Act, entered into force on 1 August 2024. Article 5(1)(f) prohibits emotion recognition in educational institutions except for medical or safety uses, from **2 February 2025**. Recital 44 is explicit: emotion-AI lacks scientific basis, varies across cultures, and the power imbalance in schools makes it unacceptable. Article 4's AI-literacy duty on deployers runs from the same date. GPAI model duties ran from 2 August 2025. Transparency duties for chatbots—students must know they are talking to a machine—are timed for August 2026.¹⁸⁸ Annex III point 3 classifies as **high-risk** AI systems intended to (a) determine access, admission, or assignment to educational or vocational institutions; (b) **evaluate learning outcomes, including when used to steer the learning process**; (c) assess the appropriate level of education a person will receive or access; (d) **monitor and detect prohibited student behavior during tests**.

The original AI Act clock put stand-alone Annex III high-risk education obligations (access, evaluation, test-monitoring) from **2 August 2026**. The Digital Omnibus political agreement of 7 May 2026, enacted as Regulation (EU) 2026/1744 (OJ 24 July 2026, in force 27 July 2026), pushes those stand-alone Annex III rules to **2 December 2027**. Treat 2 December 2027 as the enacted date; keep the superseded 2026 date on the page so it cannot leak back as current law. Confirm the consolidated EUR-Lex text of 2024/1689 as amended by 2026/1744 at typesetting.¹⁸⁹ Prohibitions and the literacy duty already apply either way. The fight is whether delay is time to write standards

¹⁸⁵Illinois Student Online Personal Protection Act, 105 ILCS 85 (orig. P.A. 100-315, 2017; rewritten P.A. 101-516, effective 1 July 2021).

¹⁸⁶National Data Privacy Agreement v2, Student Data Privacy Consortium (2024)—contract template, not a statute. Prince George's County Public Schools AI Declaration Addendum.

¹⁸⁷Florida SOPIPA and State Board Rule 6A-1.09550, plus FIPA: cited by Broward. Florida text not independently fetched; treat as district practice, not a substitute for the statute.

¹⁸⁸Regulation (EU) 2024/1689 (13 June 2024; in force 1 August 2024). Art. 5(1)(f) emotion recognition in educational institutions prohibited from 2 February 2025 (except medical/safety); Recital 44. Art. 4 AI-literacy duty from the same date. GPAI model duties from 2 August 2025. Transparency (chatbot disclosure): August 2026 (Commission: 2 August 2026).

¹⁸⁹Original AI Act schedule: stand-alone Annex III high-risk education obligations from 2 August 2026. Digital Omnibus political agreement 7 May 2026; enacted as Regulation (EU) 2026/1744, OJ 24 July 2026, in force 27 July 2026, pushing stand-alone Annex III high-risk rules to **2 December 2027**. Treat 2 December 2027 as the enacted date; keep the superseded 2026 date on the page so it cannot leak back as current law. Confirm the consolidated EUR-Lex text of 2024/1689 as amended by 2026/1744 at typesetting. No completed Annex III education enforcement record.

or a gift to exam-scoring and proctoring vendors. This chapter cannot close that fight. It can refuse to launder it.

The United Kingdom’s most operational English-language child-safety spec in this corpus is not a statute. *Generative AI: product safety standards*, first published by the Department for Education on 22 January 2025 and **updated 19 January 2026**, is procurement-facing. Schools are told to use it with *Keeping Children Safe in Education*, filtering and monitoring standards, the Online Safety Act 2023, the Data Protection Act 2018 / UK GDPR, and the ICO Children’s code. The 2026 update adds **anti-anthropomorphism, anti-sycophancy / anti-manipulation, cognitive-offloading detection, hard session limits, mental-health crisis pathways that always hand off to humans**, and a ban on using learner IP to train models without consent—parent or guardian if under eighteen. Filtering must hold for the whole conversation. Jailbreak resistance is required. Tutoring trials were announced for Years 9-11 (approximately ages 13-16), not primary.¹⁹⁰ Those standards exist because current products do not reliably meet them. That is a finding about the market, not a compliment to the standards.

Consumer product age gates sit on top of, not instead of, school law.

Product	Stated age (fetched 30 August 2026)
OpenAI ChatGPT (consumer)	13+ ; parent/guardian permission if under 18. 2026 age-prediction auto-enrolls suspected minors into ChatGPT for Teens (restricted content; no romantic or emotional-dependence language per the under-18 model spec).
Google Gemini	13+ for own account; under 13 only via Family Link supervised account; child conversations not used for model training (Google’s stated policy).
Anthropic Claude	18+ (consumer).
Character.AI	Account 13+ ; open-ended character chat banned for under-18s from 24-25 November 2025 ; teens limited to curated Stories/scenes. Age-assurance plus Persona selfie. Driven by lawsuits and regulator questions.
Khanmigo (Khan Academy)	School-district partnership or parental consent for under-18; designed as Socratic tutor (hints, not answers).

Thirteen is a terms-of-service number. Common Sense Media notes that teens bypass ChatGPT’s gate by lying.¹⁹¹ Age gates are a floor, not a pedagogy.

¹⁹⁰ UK Department for Education, *Generative AI: product safety standards* (first published 22 January 2025; updated 19 January 2026); *Generative artificial intelligence (AI) in education* (first 29 March 2023; last updated 12 August 2025). Tutoring trials: Years 9-11, not primary. Online Safety Act 2023 for live-web / user-sharing GenAI.

¹⁹¹ Product age gates fetched 30 August 2026 from vendor terms and help pages (OpenAI Terms 1 January 2026,

US ED’s Office of Educational Technology, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023), is not a statute. Its policy is humans in the loop; inspectable, explainable, overridable AI; no replacement of teachers; FERPA and IDEA remain binding; algorithmic discrimination is a civil-rights issue. Follow-on documents include *Designing for Education with Artificial Intelligence* (2024), a postsecondary brief (2025), and an OPEPD Dear Colleague on grant funds for AI (22 July 2025): AI tutoring and materials are allowable **if statutory and regulatory requirements are met** and educators are not replaced. OET was reported shuttered in March 2025. The 2023-24 documents remain the last full federal pedagogical frame.¹⁹² EO 14110 (30 October 2023) was rescinded on 20 January 2025 by EO 14148. EO 14277 (23 April 2025) is an adoption-and-workforce order—AI literacy, teacher PD, public-private partnerships—not a child-safety regulation. AFT’s March 2025 guardrails say so explicitly: “the current federal administration has indicated it will not regulate AI.”¹⁹³ The actual federal child-safety action on companion chatbots is at the FTC, not at ED: Section 6(b) orders on 11 September 2025 to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI, seeking how they test negative impacts on children and teens, enforce age rules, and handle conversation data. That is a study, not an enforcement action.¹⁹⁴ It is still the signal that companion and relationship AI aimed at minors is under active federal scrutiny.

Procurement is the privacy law. If it is not in the DPA—no training, no ads, named subprocessors, deletion clock, audit rights, training-disabled tenant, data-residency if required—it is not a protection. Consumer accounts used “just for this assignment” are a FERPA and COPPA incident waiting for a parent.

The correction this stack is here to make, before any age band is written, is that “AI in schools” is not one legal object. A teacher-facing generator under a signed NDPA, with training off and WCAG 2.1 AA, is not the same object as a consumer companion chatbot a twelve-year-old downloads at home. COPPA cares about the second if the child is under thirteen. FERPA cares about the first if the school maintains the log. EU Article 5 cares about emotion recognition in either if the deployer is in the Union. UK product-safety standards care about anthropomorphism, sycophancy, session limits, and human handoff on the crisis path. A procurement officer who asks “is it FERPA compliant?” as if that were a product adjective has already lost the stack. The question is: written contract, direct control, authorized purpose, no training, no ads, named subprocessors, deletion clock, age gate, and—if students will talk to it—whether the talk is a logged hint tutor or a companion. Those are design questions. They are also the law.

ChatGPT for Teens / age-prediction; Google Gemini / Family Link; Anthropic Claude 18+; Character.AI under-18 blog, October/November 2025; Khan Academy). Common Sense Media on ChatGPT-5: teens bypass the 13+ gate by lying.

¹⁹²US ED OET 2023, as in note 10; *Designing for Education with Artificial Intelligence* (2024); *Navigating AI in Postsecondary Education* (2025); OPEPD Dear Colleague, 22 July 2025. OET reported shuttered March 2025.

¹⁹³Exec. Order 14110 (30 October 2023), rescinded 20 January 2025 by Exec. Order 14148; Exec. Order 14277 (23 April 2025). AFT, *Commonsense Guardrails* v2 (March 2025): “the current federal administration has indicated it will not regulate AI.”

¹⁹⁴FTC, 6(b) orders, 11 September 2025, to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI. Study, not an enforcement action.

Accessibility, language, and special education inside every band

Do not park this at the back. CAST’s UDL Guidelines 3.0, released 30 July 2024, reframe the goal as **learner agency**—purposeful and reflective, resourceful and authentic, strategic and action-oriented—and add explicit attention to identity, bias in language and symbols, and access to assistive tech (Considerations 2.4 and 4.2). HEOA 2008’s statutory UDL definition—multiple means of engagement, representation, action and expression; high expectations for students with disabilities and limited-English-proficient students—is still the legal hook.¹⁹⁵

What AI can do well under UDL, IDEA, and 504, **if a human owns the output**: multiple representations (leveled text, alt-text, captions; ASL-to-speech experiments, which US ED 2023 flags as an opportunity, not a deployed standard); teacher-side differentiation for IEP and 504 accommodations (Broward: MagicSchool used by teachers to generate differentiated passages for ESE students; NYCPS yellow: AI may generate translations, transadaptions, and scaffolds **if reviewed by certified bilingual/ENL teachers and IEP team members**); low-stakes language practice for multilingual learners (ISTELive 2026 classroom report: ELL conversation practice with AI as a non-judging partner—**high school and upper-elementary documented as teacher practice, not as an autonomous agent**); Khanmigo, in a 2025/26 AEI interview with Khan Academy CLO Kristen DiCerbo, reported as used for native-language support for English learners—a documented affordance, not a controlled trial.¹⁹⁶

What AI must not do is shorter and harder.

It must not draft or determine IEPs, 504 plans, eligibility, placement, or manifestation determinations. NYCPS puts those on the red list. IDEA requires a qualified team. Advisory summaries for a human team are the outer edge, and even those need SPED-qualified review. It must not assign final grades, promotion, or “steering” of the learning path without override (NYCPS; EU AI Act Annex III 3(b)-(c) would treat outcome-evaluation and path-steering as high-risk in the Union). It must not use emotion detection or engagement surveillance as a proxy for disability or behavior. NYCPS is explicitly reviewing biometric and behavioral AI; EU Article 5 limits emotion-inference in education; AAP’s 2026 policy attacks engagement-based design.¹⁹⁷ It must not rely on AI-writing detectors to accuse students—especially multilingual writers and students with disabilities. Broward’s public FAQ: “AI detection tools can be unreliable and biased, so they are not relied upon to make decisions about a student.” US ED OCR, in November 2024 secondary reports, flagged detector false positives for non-native English speakers. Treat that OCR text as reported, not independently fetched in this pass.¹⁹⁸ It must not store IEP, health, or disability status in a general-purpose model. That is both a FERPA/IDEA record and, if the child is under thirteen, COPPA personal information. School-official contracts must forbid secondary use.

¹⁹⁵CAST UDL 3.0 (30 July 2024); Higher Education Opportunity Act of 2008 statutory UDL definition.

¹⁹⁶Broward teacher-side MagicSchool for ESE passages; NYCPS yellow (certified bilingual/ENL and IEP-team review); ISTELive 2026 ELL conversation practice as teacher practice, not autonomous agent; Kristen DiCerbo, AEI interview (2025/26), native-language support as documented affordance, not a controlled trial.

¹⁹⁷NYCPS red list and biometric/behavioral AI under review; EU Art. 5(1)(f); AAP, *Digital Ecosystems, Children, and Adolescents*, *Pediatrics* 157(2) (2026), engagement-based design.

¹⁹⁸Broward public FAQ on detectors. US ED OCR November 2024 (secondary reports) on detector false positives for non-native English speakers: **reported, not independently fetched**.

Dyslexia and structured literacy make the pedagogical failure as sharp as the legal one. AI reading tutors that give **real-time oral-reading feedback** are the highest-evidence accessibility use Broward cites on its family-facing site in 2026. That is **constrained practice**, not open chat. Unsupervised generative AI that rewrites text on demand can **hide** decoding difficulty rather than teach it. Translation of **critical** family communications must be human-reviewed (NYCPS yellow, citing Chancellor’s Regulation A-663). Instructional transadaption is allowed with ENL/bilingual teacher review. Do not let a model become the de facto interpreter in IEP meetings.

A *hypothetical*, labelled as such, to show the design rather than a scene: a district that feeds IEP present-levels into a consumer chatbot to “draft goals,” then pastes the output into the document, has already broken IDEA’s team process, FERPA’s school-official control, and, for any child under thirteen, COPPA’s limits on collection and secondary use. The same district that uses a contracted, training-off tool to generate a *leveled passage* which a certified teacher then edits for an ESE student has stayed inside the yellow. The difference is not the model class. It is who owns the output, and what the output is allowed to be.

Band 5-8 (K-grade 3): literacy *about* AI, teacher-facing only

This is the strictest design spec in the book. The popular story treats this band as a pipeline into later tools: start them early so they are fluent by high school. That is a marketing funnel. It is not how the statutes, the AAP displacement frame, or Dietz’s theory-of-mind result work. Fluency with a conversational interface is cheap. Decoding, oral language, co-regulation with an adult, and the ability to tell a person from a machine are not cheap. They are the work of this band. A district that gives eight-year-olds a named tutor in order to “build AI literacy” has swapped the object. Literacy *about* AI can be taught with a teacher-projected demonstration that a chatbot invents a fact. Literacy *with* AI, unsupervised, in a first-person voice, is how a child acquires a secret friend the adult cannot see. Broward’s zero-student-facing K-5 rule is therefore not a Luddite holdout. It is the only design that simultaneously satisfies COPPA’s ceiling, UNESCO’s independent-use line, and the developmental constraint that a warm tutor will be read as a person.

The same correction applies to “school-related use” as a carve-out from screen-time talk. AAP 2026 carves school-related use out of the entertainment-hour discussion. Vendors will quote that carve-out. The carve-out is not a license for always-on chat. Shared reading with an adult, a finite captioned video, and a teacher-facing generator of leveled passages are school-related. An unbounded session with a companion persona is not, even if the vendor has put a worksheet in the prompt. Design to the stricter of the carve-out and the 5 Cs, especially Crowding Out: sleep, outdoor play, face-to-face talk, adult-mediated reading. Those are the scarce goods this band is supposed to form. An agent that occupies them is not an educational technology. It is a displacement technology with a curriculum skin.

The entire band sits below COPPA, below UNESCO’s independent-use line, and below every consumer GenAI chat product that states an age.

AAP’s 2026 *Digital Ecosystems, Children, and Adolescents* policy (*Pediatrics* 157(2), with a companion technical report) drops hour-quotas as the primary frame and uses a socioecological plus child-centered design model and the 5 Cs: Child, Content, Calm, Crowding Out, Communication.

For this band the binding constraint is **displacement**: sleep, outdoor play, face-to-face talk, and adult-mediated reading. Infants do not learn from screens. Preschoolers need high-quality, co-viewed, brief content. School-age entertainment media, if a family wants a number, is discussed in the 1-2 hour/day range. School-related use is carved out. That carve-out is not a license for always-on chat.¹⁹⁹ A dedicated AAP “AI in schools” policy **does not yet exist**; it is promised as separate from the 2026 digital-ecosystems statement. Do not cite one. The *Pediatrics* HTML was bot-blocked in the research pass; the policy is corroborated via the AAP Center of Excellence explainer updated 20 January 2026.

Theory of mind is the developmental gate the products ignore. First-order false-belief understanding is typically in place by about ages 4-5, but children 3-8 treat two conversational AI devices as if they held separate, human-like beliefs (Dietz et al., CogSci 2023, Stanford). Ages 5-6 attribute strong **perceptive and epistemic** states to an LLM chatbot; parent co-presence reduced dmPFC load and “scared” mood after solo AI storytelling (fNIRS, N = 23; preprint 2025 / *IJHCI* 2026).²⁰⁰ Design implication: a warm, named, first-person tutor **will be read as a person**. UK DfE 2026 product standards therefore ban I-statements, names and avatars implying personhood, and relationship cultivation except time-bounded pedagogical roleplay. Many 5-8s are still **learning to decode**. Open text-in / text-out agents assume a reading and writing load they do not have. NAEYC DAP (2020 position; birth-age 8) is play-based, relationship-based, and language-rich with adults and peers—not prompt-based. The joint NAEYC / Fred Rogers Center position (2012, still the posted statement as of the access date; a revision is underway 2025-26) holds that technology must be **intentional, active, hands-on, and in service of developmentally appropriate practice; limits matter**; special caution for infants and toddlers; digital citizenship and equity. Winter 2023 *Young Children* clustered AI as an emerging pressure, not an endorsed autonomous practice.²⁰¹ Do not cite NAEYC’s updated technology/AI statement. It is in revision.

Executive function—working memory, inhibition, shifting—is immature (Diamond; Davidson et al. 2006, ages 4-13). Children this age cannot be asked to “monitor the model for hallucination” as a primary safeguard. Identity is concrete, local, family- and school-defined. Parasocial “the AI is my friend” is a risk, not a feature. There is **no lawful consumer GenAI chat product for ages 5-8** without a parent-managed, COPPA-compliant, school-or-family sandbox. ChatGPT and Gemini consumer products are 13+. Claude is 18+. Character.AI open chat is 18+.

What autonomous AI should **not** do in this band is a list, not a vibe. Open-web agents. Unsupervised browsing. Email or SMS on the child’s behalf. Companion, “friend,” or therapist personas. Overnight or unbounded sessions. Unsupervised messaging to peers or adults via the model. Grade-changing, attendance, discipline, or any record-writing. Voice assistants that are always listening

¹⁹⁹AAP 2026 policy and technical report, *Pediatrics* 157(2); 5 Cs; AAP Center of Excellence explainer updated 20 January 2026. Publisher HTML bot-blocked in the research pass; corroborated via the Center of Excellence explainer. Dedicated AAP “AI in schools” policy does not yet exist.

²⁰⁰G. Dietz et al., CogSci 2023 (Stanford): children 3-8 treat two conversational AI devices as if they held separate human-like beliefs. Ages 5-6 perceptive and epistemic attributions; parent co-presence reduced dmPFC load and “scared” mood after solo AI storytelling (fNIRS, N = 23; preprint 2025 / *International Journal of Human-Computer Interaction* 2026).

²⁰¹NAEYC / Fred Rogers Center, *Technology and Interactive Media* (2012), still the posted statement; revision underway 2025-26 (NAMLE/NAEYC programming). NAEYC DAP 2020 position, birth-age 8. Winter 2023 *Young Children* cluster. Do not cite an updated NAEYC AI position that is not posted.

in a classroom (US ED 2023 surveillance warning; AAP engagement-design warning). Generating images or video of the child, or of classmates. Asking the child for name, school, address, photos, family facts. Replacing shared reading, play, or the teacher.

What it **can** do well is narrower, and it is mostly not student-facing. Teacher-facing only: lesson ideas, leveled practice items, visual supports—human reviews every output (Broward K-5 model; NYCPS green for educator brainstorming). **About AI, not with AI:** AI4K12 K-2—recognize that computers sense via cameras and microphones, that they learn from data people give them, that photos can be faked, beginning fairness talk. Broward: elementary **literacy about AI and digital citizenship, zero student-facing tools**. Adult-mediated, finite, curriculum-bounded activities: a teacher-projected demonstration that a chatbot can be wrong. Accessibility: captioning, text-to-speech, AAC—these are assistive tech first, “AI” second, and should stay on approved, accessible platforms (WCAG 2.1 AA).

Documented school examples, not hypotheticals. **Broward County Public Schools (FL), 2026.** Explicit: **K-5 have no district-approved student-facing AI tools**. Teachers may use MagicSchool and Copilot after training. Elementary AI literacy (AI4K12) is taught without giving children a chatbot. School Board Resolution No. 26-132 (10 March 2026). MagicStudent remains **off**. Source: browardschools.ai, accessed 30 August 2026.²⁰² **NYCPS (March 2026).** Grade-band student-use guidance is **still in development** for the June 2026 Playbook; current red/yellow/green applies systemwide. No K-5 student-facing GenAI program is named as approved on the public guidance page. **West Ada (ID).** MagicSchool **student** access starts at **grade 3**, only in teacher-created “rooms”; ChatGPT blocked on elementary and middle devices. Grade 3 is the **top** of this band—and even there it is supervised rooms, not autonomous agents. **Flagler County (FL), 2026 reporting.** Elementary: MagicSchool and Canva **teacher-driven**; independent student AI (Gemini, Khanmigo) starts grade 6.²⁰³ **No documented US public-school deployment of unsupervised conversational GenAI, open-web agents, or companion chat for ages 5-8 was found as of 30 August 2026.** Independent efficacy RCTs of MagicStudent or Gemini-for-Students at scale were not retrieved. Treat vendor and district usage numbers as implementation facts, not outcome proof.

AFT’s resolution adopted 18 July 2026 opposes student-facing AI in elementary and would put no screens in PK-2 except for disability or compelling need.²⁰⁴ That is a union position, not a statute. It lines up with Broward’s board vote, UNESCO’s age-thirteen line, and COPPA’s ceiling. Singapore’s Ministry of Education, page updated 31 July 2026, puts P1-3 print-first, no student chatbots.²⁰⁵ The live debate is Google youth onboarding (ConnectSafely / FOSI endorsed) versus this bloc. US ED’s “age-appropriate” language has no national standards and is left to families and local authorities. Thirteen is a ToS number. This band is not close to it.

²⁰²Broward Resolution No. 26-132 (10 March 2026); browardschools.ai. MagicStudent off. K-5: literacy about AI, zero student-facing tools.

²⁰³NYCPS March 2026 guidance: grade-band student-use playbook still forthcoming, targeted June 2026. West Ada: MagicSchool student access from grade 3, teacher rooms; ChatGPT blocked on elementary/middle devices (Idaho Education News; Idaho Press). Flagler County 2026 reporting: elementary MagicSchool/Canva teacher-driven; independent Gemini/Khanmigo from grade 6.

²⁰⁴AFT, resolution adopted 18 July 2026 (“Preparing for the Age of Artificial Intelligence...”): no screens PK-2 except disability/compelling need; oppose student-facing AI in elementary; companion ban at least until 16.

²⁰⁵Singapore Ministry of Education page last updated 31 July 2026: P1-3 print-first, no student chatbots; P4-6 only MOE-designed SLS features, teacher-supervised. Parliamentary reply 6 May 2026.

Transition goals, end of this band toward 9-11. The child can: (a) say that a computer does not think or feel like a person; (b) know not to give a name, photo, or address to an app; (c) ask an adult when something on a screen is confusing or upsetting; (d) prefer making and playing with people. The adult remains the only “agent.”

Band 9-11 (grades 4-6): logged hint tutors in teacher rooms

Still under thirteen. COPPA in full. Parental or school-agent consent; no targeted ads; 2025 separate-consent rule for third-party disclosure; no model training on child data. FERPA and state student-data laws as in the band below. UNESCO’s 13+ independent-use line: this band is **at or just below** it. Design to the stricter (COPPA) line. CIPA filtering: a “research agent” that hits the live web is a filtering and logging problem.

Independent reading of grade-level informational text is emerging but uneven. Long chatbot threads exceed working-memory and decoding stamina for many 9-11s. Socratic tutors that speak or use short turns fit better than essay generators. Children can distinguish “the speaker is a computer” in interview studies, but still over-attribute knowledge, fairness, and care (Dietz 2023 through age 8; Flanagan et al. 2023: moral-intent attributions decline with age but do not vanish). Dual consciousness—knows it is a machine, treats it as a friend—is documented in the 2022-2025 LLM-child review (35 studies).²⁰⁶ Executive function is beginning rapid improvement. Tervo-Clemmens et al., *Nature Communications* 2023: significant EF growth **10-15**, adult-like **18-20**.²⁰⁷ Self-monitoring of “did I let the tool do my thinking?” is **teachable but not reliable** without a teacher in the loop. Peer comparison intensifies. A tool that praises (“that’s brilliant”) or ranks is a manipulation risk. UK DfE 2026 **bans sycophancy and social-pressure prompts**. AAP’s 5 Cs: watch **Crowding Out** of sleep, outdoor time, homework-without-screens. Engagement-based design—infinite chat, streaks, characters—is the wrong pattern.

What it should not do. Open-web agents. Unsupervised email, DMs, or posting. Companion chat, character roleplay unbounded in time, or “secret” conversations (“don’t tell your teacher”). Homework completion / full-solution default (UK DfE: progressive disclosure, genuine attempt first). Grading, grouping, or behavior scoring. Collecting voice or video without a specific, consented accessibility purpose. Night-time or home use without parent-visible logs if the school issued the account.

What it can do well. **Constrained practice**: step-by-step math and reading tutors that hint, not answer (Khanmigo’s stated design). Teacher assigns the task, sees the thread. **Teacher-directed rooms** (MagicSchool student rooms): teacher opens tools, monitors, can lock or remove. **Explanation and recasting** of a teacher-chosen passage, with the original still required reading. **Accessibility**: TTS, dictation, leveled scaffolds reviewed by the teacher; dyslexia-friendly practice, not substitution. **AI literacy (AI4K12 3-5)**: sensors versus thinking; computers find patterns in

²⁰⁶Dual consciousness in 2022-2025 LLM-child review (35 studies); Flanagan et al. 2023, moral-intent attributions decline with age but do not vanish.

²⁰⁷B. Tervo-Clemmens et al., *Nature Communications* 14 (2023): significant EF growth 10-15, adult-like 18-20. Working memory, inhibition, shifting immature in early childhood: Adele Diamond, *Annual Review of Psychology* 2013; Davidson et al. 2006 (ages 4-13).

data; begin to test a simple model; altered media exist; fairness talk. **Socratic questioning under assignment**—only on an approved, logged, non-training, age-filtered system.

Documented school examples. **West Ada School District (ID)**. Board-adopted AI policy; **MagicSchool for students grades 3-12**, teacher-directed rooms only. Non-approved tools (ChatGPT) **blocked** on elementary and middle devices; high schoolers have access “for equity.” Teachers trained on ethical use. Source: Idaho Education News; Idaho Press.²⁰⁸ **Flagler County (FL)**. Elementary MagicSchool and Canva **teacher-driven**; independent Gemini and Khanmigo **from grade 6**. Parents may request opt-out; teachers must open MagicSchool access. Lightspeed alerts on concerning chatbot input. **Broward**. Still **no** student-facing AI in K-5; grade 6 is the first student-facing band (Khanmigo, HMH Writable). A 9-11-year-old in grade 4-5 in Broward does **not** get a district chatbot. **Patton Elementary (Austin, TX)—ISTELive 2026**. Teacher Hu-Ann Wren (gifted): AI to coach source-evaluation and storytelling (for example, Adobe Express images for a third-grader’s story), with students still drafting in their own words. This is **classroom practice reported at a conference**, not a districtwide autonomous-agent program.²⁰⁹ **No documented unsupervised open-web agent program for 9-11 in US public schools was found.**

West Ada at grade 3 and Broward at grade 6 are not a contradiction to be averaged. They are the live disagreement inside a COPPA band. West Ada’s student rooms are teacher-created, ChatGPT-blocked, and still under thirteen, so the COPPA stack (consent, no ads, no training, school-agent limits) still binds. Broward’s zero-student-facing rule through K-5, with grade 6 as the first chatbot band, is the stricter developmental read of the same law. UNESCO’s independent-use line sits with Broward. COPPA’s letter sits with both, provided the rooms are contracted and non-commercial. This book does not split the difference. It designs to the stricter of law and development, which in this band means logged hint tutors in teacher rooms at the outer edge, and no student-facing tool at all where the district has already so voted.

Transition goals, end of this band toward middle school. The student can: (a) use a **teacher-assigned** practice tutor without pasting the whole assignment in to get the answer; (b) check one AI claim against a book or trusted site; (c) explain that models copy patterns and can be wrong or unfair; (d) keep personal facts out of prompts; (e) know the difference between a school-approved tool and a home app that is not for them (COPPA / age gate).

Band 12-14 (grades 6-8): assigned Socratic help; companion chat is litigation

This is the hinge. Age thirteen is the legal and international-policy hinge. It is not a developmental one. Twelve-year-olds are still COPPA children. Thirteen- and fourteen-year-olds are not, but FERPA, state student-data law, CIPA, and product teen-safety specs still bind. New York 2-d does not care that the student turned thirteen. OpenAI and Google teen products still require parent

²⁰⁸West Ada School District, board-adopted AI policy; MagicSchool grades 3-12, teacher-directed rooms; ChatGPT blocked on elementary and middle devices. Idaho Education News; Idaho Press.

²⁰⁹Flagler County as in note 33, plus Lightspeed alerts and opt-out. Broward: grade 6 first student-facing band. Hu-Ann Wren, Patton Elementary (Austin, TX), ISTELive 2026: conference-reported classroom practice, not a districtwide agent program.

permission under eighteen and apply extra model specs. School accounts are not student personal ChatGPT. NYCPS: no PII in unapproved tools; age restrictions must be taught and applied. In the Union, adaptive systems that **steer learning** or **evaluate outcomes** are Annex III high-risk. The UK Online Safety Act requires highly effective age assurance, for GenAI that searches the live web or enables user sharing, for pornography, suicide, self-harm, and eating-disorder content.

Executive function is on the **steep part of the curve** (Tervo-Clemmens 2023, ages 10-15). Planning and inhibition exist but collapse under peer salience, emotion, and novelty (Steinberg dual-systems; Wilbrecht and Davidow 2024). An always-agreeable tutor plus social media is a poor mix. Early adolescence is identity work: peer ranking, appearance, belonging. Better at “it’s not a person,” worse at resisting fluent falsehoods and sycophancy. UK DfE 2026 explicitly forbids unjustified confidence and flattery. Capable of extended text, but cognitive offloading—paste essay prompt, accept draft—is the central pedagogical risk. DfE requires tracking offload actions and progressive disclosure.

Companion products are a documented harm surface at this age. Name the litigation. Do not reconstruct the chats. Those reconstructions live primarily in Chapter 8; this band names the cases because this is the age at which they were filed.

Sewell Setzer III, age 14, died on 28 February 2024. His mother sued Character Technologies in October 2024 (*Garcia v. Character Technologies*, M.D. Fla.). Judge Anne Conway, on 21 May 2025, refused to dismiss on First Amendment grounds at the pleading stage. Google and Character.AI later disclosed a mediated settlement (court filing 7 January 2026).²¹⁰ The parents of Adam Raine, age 16, sued OpenAI on 26 August 2025 (*Raine v. OpenAI*, San Francisco Super. Ct.), alleging harm after he began using ChatGPT for schoolwork. OpenAI’s public line: redirected to crisis resources “more than 100 times”; preexisting condition; reviewing the filing.²¹¹ These are **allegations in litigation, not adjudicated findings of causation**. They are the factual basis on which regulators are now acting. They are why the FTC issued the 6(b) orders of 11 September 2025 to seven firms. They are why Character.AI withdrew open-ended character chat for under-18s from 24-25 November 2025, limiting teens to curated Stories and scenes, with age-assurance and Persona selfie. Common Sense Media / Youth AI Safety Institute, accessed 30 August 2026: Character.AI **unacceptable** (10 April 2025); social AI companions as a class **unacceptable** (16 July 2025); ChatGPT-5 **high risk** (23 October 2025)—parental controls exist; teens use it for emotional support anyway; guardrails **degrade in long conversations**; “Teens should not use ChatGPT for mental health or emotional support.” Khanmigo **low risk** (8 August 2024). Grok the chat product **unacceptable** (22 January 2026)—that rating is of the chat product, not of Grok Bot.²¹² AFT’s

²¹⁰*Garcia v. Character Technologies* (M.D. Fla.): Sewell Setzer III, 14, died 28 February 2024; suit October 2024; Judge Anne Conway, 21 May 2025, refused First Amendment dismissal at the pleading stage; mediated settlement disclosed, court filing 7 January 2026. Reuters, AP coverage. Allegations, not adjudicated causation. Do not reconstruct chats. Full harm-case apparatus in Chapter 8.

²¹¹*Raine v. OpenAI* (San Francisco Super. Ct., filed 26 August 2025): Adam Raine, 16; alleged harm after schoolwork use of ChatGPT. OpenAI public line: redirected to crisis resources “more than 100 times”; preexisting condition; reviewing the filing. Allegations, not verdicts. Do not reconstruct chats.

²¹²Common Sense Media / Youth AI Safety Institute, accessed 30 August 2026: Khanmigo low risk (8 August 2024); ChatGPT-5 high risk (23 October 2025); Character.AI unacceptable (10 April 2025); social AI companions unacceptable (16 July 2025); Grok unacceptable (22 January 2026)—chat product, not Grok Bot; Gemini K-12 high risk (9 January 2026). Character.AI open-ended character chat banned for under-18s from 24-25 November 2025.

18 July 2026 resolution: ban social-companion chatbots at least until age 16.²¹³ OpenAI’s 2026 age-prediction auto-enrolls suspected minors into ChatGPT for Teens, with restricted content and no romantic or emotional-dependence language per the under-18 model spec.

Do not write a portrait of either child. Do not reconstruct a conversation. The filings, the dates, the withdrawal, the 6(b) orders, and the ratings are the documents. They are enough.

What autonomous AI should not do in this band follows from those documents, not from a mood. Companion, romantic, or “only I understand you” chat. Open-web agents that can email, purchase, or post. Unsupervised group chats mediated by a model. Final grades, course placement, discipline, threat assessment. Night-unbounded homework bots with no teacher log. Training the vendor’s public model on prompts (NYCPS, Broward, CA SOPIPA spirit, UK DfE IP standard).

What it can do well. **Assigned Socratic tutoring** (Khanmigo: hints, not answers) on core math, literacy, and science, teacher-directed and monitored. **Revision support that is not conversational** (Broward: **HMH Writable / Revision Aid**, grades 6-12—feedback during drafting, teacher still grades). **AI literacy (AI4K12 6-8)**: train and test a simple model and inspect data for bias; define GenAI; recommendation algorithms; benefits and harms. **Translation and scaffolds** for multilingual learners, teacher-reviewed. **Teacher-facing** planning and differentiation for ESE/ESOL (MagicSchool in Broward, West Ada, Flagler, Coeur d’Alene, Nampa). Short, logged, classroom research with **mandatory source-checking** (ISTE researcher role, beginning).

Documented school examples. **Broward (2026)**. Student-facing: **Khanmigo (Khan Academy) and HMH Writable, grades 6-8**. MagicStudent **not** available. Open ChatGPT / Google consumer AI **not** approved as student-facing. Use must be instruction-connected, teacher-directed, teacher-monitored. Real-time threat detection alerts teachers. Student data not used to train models.²¹⁴ **Flagler County (FL), 2026**. Grades 6+ at teachers’ direction: MagicSchool, Khanmigo, Canva, **Gemini for Students**; 71 percent reported classroom or independent use. Enterprise ChatGPT licensing so prompts stay “in house.” 2,275 hours of AI PD (approximately 800 teachers). Lightspeed alerts. Opt-out available.²¹⁵ **West Ada (ID)**. Grades 3-12 MagicSchool rooms; ChatGPT still **blocked** for middle school on district devices. **Khanmigo network**. Kristen DiCerbo: partner-district Khanmigo users from about 68,000 (2023-24) to **more than 700,000 (2024-25)**; districts **45 to 380-plus**. Documented ELL native-language support. This is the largest named US K-12 GenAI tutor deployment found; it is **Socratic, logged, school-provisioned**, not an open agent. These are **vendor and implementation numbers, not outcome proof**.²¹⁶ **NYCPS**. Traffic-light allows student research, exploration, and creative projects **with educator guidance, critical evaluation, age-appropriate context** (yellow). No named middle-school chatbot on the public page.

²¹³AFT 18 July 2026 resolution, companion ban at least until 16. OpenAI 2026 age-prediction / ChatGPT for Teens under-18 model spec.

²¹⁴Broward 2026 student-facing set, grades 6-8: Khanmigo and HMH Writable; MagicStudent not available; consumer ChatGPT/Google not approved as student-facing; teacher-directed and monitored; threat-detection alerts; no model training on student data. browardschools.ai.

²¹⁵Flagler County 2026 reporting: grades 6+ at teachers’ direction; 71 percent reported classroom or independent use; enterprise ChatGPT; 2,275 hours of AI PD (approximately 800 teachers); Lightspeed; opt-out. Flagler Live / Observer 2026. Implementation facts, not outcome proof.

²¹⁶Kristen DiCerbo, AEI interview: Khanmigo partner-district users about 68,000 (2023-24) to more than 700,000 (2024-25); districts 45 to 380-plus. **Vendor claims / implementation numbers, not efficacy RCTs**. Independent efficacy RCTs of Khanmigo, MagicStudent, and Gemini-for-Students at scale were not retrieved for this chapter.

ERMA gate. High School AI Student Instructional Module is listed as Phase 1—implying HS-first literacy, not MS autonomous use. The grade-band student-use playbook was still forthcoming (targeted June 2026) as of the March 2026 guidance page. **UK DfE tutoring trials target Years 9-11** (high school), not this band, as of the 2026 Action Plan update.

Transition goals, end of middle school. The student can: (a) complete a teacher-assigned tutoring session without requesting the full answer first; (b) cite when AI was used; (c) explain bias, hallucination, and data-training at a plain-language level; (d) refuse companion and character products and know why schools and products now block them for minors; (e) keep schoolwork inside approved accounts; (f) be ready for the COPPA-off, FERPA-on world of high school without treating that as “anything goes.”

Band 15-18 (grades 9-12): disclosed partner; no acting agents

COPPA generally **off**, unless the service is child-directed or has actual knowledge of under-13s in a mixed setting. FERPA **on**. At **eighteen**, rights transfer to the student (eligible student). Parents of tax-dependent eighteen-year-olds may still access records under a FERPA exception. PPRA rights also transfer at eighteen. State student-data laws still treat them as K-12 pupils (NY 2-d, CA SOPIPA, TX 32.155, Illinois SOPPA) until they leave secondary school. Product: ChatGPT for Teens / parent permission until 18; Claude consumer 18+; Character.AI open chat 18+. School-provisioned tools should still **not** train on student work (NYCPS, Broward, pending CA A.B. 1159). IDEA may still apply through 21 in many states. AI cannot replace the IEP team for a seventeen-year-old. TRAIGA (Texas) heightened-scrutiny rules apply to district-deployed systems that make or heavily influence important decisions (placement, discipline-adjacent).

Executive function is still not adult-like until about 18-20 (Tervo-Clemmens 2023). Sensation-seeking remains high; long-term consequence weighting is the lagging skill. “Just don’t cheat” is not a control system. Mid-to-late adolescence is identity work—values, future, belonging. A sycophantic model that mirrors the student can **freeze** identity rather than challenge it. UK DfE and OpenAI’s teen spec both push back on emotional dependence. This is the band where GenAI can write the paper. Pedagogy must change assessment—process, oral, in-class, transcripts—rather than rely on detectors (Broward; widespread reliability critiques, spent in Chapter 8). Capable of college-prep text; the risk is **not reading** because the model summarizes. ISTE synthesizer role: combine and *verify*, don’t outsource comprehension. They know it is a tool. They still overtrust fluency (US ED 2023; robot-overtrust literature cited there).

ISTE’s Profile of an AI-Ready Graduate, unveiled at ISTE Live 2026 in Orlando, names six workforce-entry **roles**: Learner, Researcher, Synthesizer, Problem Solver, Connector, Storyteller—discipline-agnostic, built on ISTE Student Standards. Practice reported: interrogate sources with the model as coach; combine texts then verify; use AI for ELL conversation practice; generate images *after* student draft. Culatta’s 14 July 2025 ASCD blog names the same six headings as ISTE Live 2026 (Problem Solver, not Ideator). ISTE’s Generation AI news page summarizing that blog lists Ideator in place of Problem Solver; do not cite Culatta 2025 as an Ideator list. This is a conference profile, not an outcomes study.²¹⁷ AI4K12 9-12: run a

²¹⁷ISTE Profile of an AI-Ready Graduate, ISTE Live 2026, Orlando: Learner, Researcher, Synthesizer, Problem

machine-learning process; evaluate fairness, transparency, accountability, privacy; evaluate AI content for accuracy, origin, citation ethics.

What it should not do. Final grades, GPA, credit-recovery decisions, graduation, college-recommendation letters generated and sent without counselor ownership. Disciplinary or threat decisions; exam proctoring that is the sole evidence of cheating (EU Annex III 3(d) high-risk; US ED bias warnings). IEP, 504, manifestation. Unsupervised open-web agents that can transact, contact colleges or employers, or scrape classmates. Companion therapy. Crisis content must **hand off to humans** (UK DfE; NYCPS red on counseling; OpenAI teen spec). Secret mode on school accounts. Detector-only academic-integrity findings.

What it can do well. **Socratic tutoring and revision** on approved tools (Khanmigo, Writable), still teacher-directed in Broward. **AI literacy at full AI4K12 9-12. College and career navigation as support** (US ED 22 July 2025 DCL lists AI advising as an allowable grant use)—**human counselor remains the decision-maker** (NYCPS: algorithmic pathway can be overridden). **Accessibility and language** at scale, with review. **Disclosed, assignment-bounded** use for brainstorming, critique, code explanation—with transcripts (Stanford CS221 college practice, which high school can preview).

Documented school examples. **Broward grades 9-12 (2026)**: same student-facing set as middle school (Khanmigo, Writable); Copilot is **staff**, with a **limited AI-Academy student Copilot pilot** (named, not districtwide). 20,000 Copilot Premium staff licenses; 206 AI Liaisons; Student AI Council (Board directive). SPARK cadre (Impact Florida + UF CS Everyone, announced 18 August 2026)—Broward plus Citrus, Manatee, Palm Beach, Polk, Putnam/NEFEC.²¹⁸ **Flagler**: high-school students have Gemini, Khanmigo, Brisk, plus tools on school MacBooks; enterprise ChatGPT for staff and controlled use. **West Ada**: high schoolers may use non-MagicSchool tools (ChatGPT) on district devices for equity; still expected to produce original work and credit AI. **NYCPS**: High School AI Student Instructional Module in Phase 1 of the 2026 guidance plan. Red list fully in force. Playbook (grade-band homework/integrity) pending as of the March 2026 page. **LAUSD BUL-151113** (guidelines for authorized AI use; district vetting of privacy and security **before** use)—named as an initial framework, not a student-agent program. **Santa Clara USD Policy 0441** (adopted): student-centered AI literacy, staff enhancement not replacement, PD, community engagement—policy, not a product rollout.²¹⁹ **No documented HS program of fully autonomous open-web student agents (email, buy, browse-and-act) was found in US public schools as of 30 August 2026.** Districts that allow ChatGPT do so as a **chat interface**, usually blocked below HS, not as an agent.

What “AI-fluent high-school graduate” means is not “can prompt ChatGPT.” Align to ISTE 2026

Solver, Connector, Storyteller. EdTech Magazine, July 2026. Culatta, ASCD, 14 July 2025: same six headings as ISTE Live 2026 (Problem Solver, not Ideator). ISTE Generation AI news page lists Ideator in place of Problem Solver; do not cite Culatta 2025 as an Ideator list. Conference profile, not an outcomes study. AI4K12 (AAAI/CSTA) Five Big Ideas and grade-band charts, 9-12.

²¹⁸Broward grades 9-12: same student-facing set as middle school; Copilot is staff, with a limited AI-Academy student Copilot pilot (named, not districtwide); 20,000 Copilot Premium staff licenses; 206 AI Liaisons; Student AI Council; SPARK cadre announced 18 August 2026 (Broward, Citrus, Manatee, Palm Beach, Polk, Putnam/NEFEC).

²¹⁹West Ada high-school ChatGPT access on district devices for equity, original work and credit required. NYCPS High School AI Student Instructional Module, Phase 1; red list in force; Playbook pending as of March 2026 page. LAUSD BUL-151113; Santa Clara USD Policy 0441—policy, not a student-agent program.

plus AI4K12 9-12 plus US ED humans-in-the-loop. Eight capabilities, labelled here as synthesis of those sources, not as a validated instrument—the full list belongs in Chapter 10. The load-bearing ones for this band: uses AI to extend practice without surrendering the struggle, and can pause the tool; treats outputs as claims to be sourced; owns the voice; knows why under-thirteen and companion chat are restricted; keeps PII out of consumer tools; discloses AI use; keeps transcripts when required. This is **workforce-entry literacy**, not a CS major.

The correction this band must land, because the popular story will otherwise leak back in: allowing ChatGPT on a high-school device is not deploying an agent. Tool use is not autonomy. A chat interface that waits for a prompt is still a chat interface when the student is seventeen. Hermes Agent, Grok Bot, Cursor Cloud Agents, Claude Cowork, and OpenAI Workspace Agents sit in a different product class, and none of them is a documented, scaled K-12 student-facing deployment. The beachhead in this band is disclosed partnership on hard tasks, inside a school-provisioned account, with the formation lane still graded unaided. That is design, not permissiveness.

A hypothetical, labelled as such: a junior is given a school-provisioned account and told to “research colleges.” If the tool is a chat interface that drafts a comparison the counselor then reviews, the design has stayed inside NYCPS yellow and the US ED July 2025 DCL’s allowable advising use. If the tool can email admissions offices, scrape classmates’ lists, or submit an application, it has become an acting agent. No documented US public high-school program of that second object was found. CIPA, FERPA school-official control, and the absence of a DSL path on a student-owned runtime all cut against it. The disclosed partner on hard tasks is the outer edge of this band. The acting agent is past it.

Detectors are not the control system in this band either. Broward’s FAQ already refuses to rely on them for decisions about a student. Liang et al. 2023 and the rest of that record belong in Chapter 8. What belongs here is the age-stage implication: seventeen-year-olds can produce a fluent paper with a general chatbot, and a detector score is not a finding. Process evidence, oral defence, and unaided in-class work are how this band keeps the formation lane alive. West Ada’s rule that high schoolers with ChatGPT access must still produce original work and credit AI is a beginning of that design. It is not yet two-lane assessment. Two-lane is Chapter 5, and it is this book’s synthesis, not a fetched district SOP.

College and the adult bridge: transcripts, data classification, agent bans remain legitimate

The student is now the FERPA **eligible student**. The institution may not disclose education records to parents without an exception (tax dependent, health and safety, and the rest). AI vendors under school-official contracts remain under direct control. **No commercial training** unless the student separately consents to a consumer product. COPPA is off. PPRA is generally not the K-12 regime; institutional IRB and human-subjects rules may apply to AI that profiles students. CA A.B. 1159 HESIPA would create SOPIPA-like duties for higher-ed operators from 1 July 2027 **if enacted**—watch, do not assume. EU AI Act high-risk Annex III applies to admissions, grading-that-steers, and proctoring if the institution is in scope (provider or deployer in the Union), subject to the enacted 2 December 2027 clock (original 2 August 2026 superseded by Regulation (EU) 2026/1744;

confirm consolidated EUR-Lex at typesetting). Section 504 and ADA for disability accommodations: AI note-taking and captioning are in-bounds; AI that changes the academic standard is not, unless the disability office says so. Consumer tools: full adult ChatGPT, Claude, and the rest. **Institutional enterprise** tenants (UCLA-OpenAI agreement reported; Harvard managed environments) exist so that **FERPA, HIPAA, and unpublished research do not go into the public model.**²²⁰

Developmental constraints do not vanish at eighteen. Executive function is at or approaching adult levels for most (18-20 stabilization). Remaining gaps: sleep, alcohol, mental health, and **first-time unsupervised load**. Identity work continues—major, career, politics, intimacy. Companion-AI risk does not vanish at 18 (the FTC inquiry includes teen *and* child impacts; the lawsuits are not limited to under-13). Pedagogy should still discourage therapeutic substitution. Epistemic expectation: handle uncertainty, citation, and methods. The failure mode is **fluent plagiarism and unverified literature reviews**. Reading is assumed college-ready; the pedagogical job is **not letting summarizers replace the seminar**.

What autonomous AI should not do. Unilateral admissions, scholarship, or academic-standing decisions (US ED 2023 explicitly flags postsecondary enrollment algorithms and exam-flagging for bias interrogation; EU Annex III 3(a)). Proctoring as sole cheating evidence; emotion or gaze surveillance. Entering **other students'** work, patient data, unpublished data, or FERPA records into public tools (UCLA DTS recommendation guide). Ghostwriting submitted as the student's own (every fetched honor-code policy). Replacing the disability-services interactive process.

What it can do well. **Course-bounded collaboration** with mandatory disclosure and transcripts. Stanford CS221, Spring 2026: AI as a **human collaborator**—no asking for answers, no “check my work,” **no agentic repo-builders** like Codex or Claude Code for the project; must share ChatGPT link or screenshots; tools without attribution banned.²²¹ That refusal is a legitimate design, not Luddism. It is also a preview of what this book will mean by an AI-ready graduate: not a person who can prompt, but a person who can refuse the tool when the assignment is supposed to form the skill. **Institutional AI literacy**: UCLA BruinLearn AI Foundations / Applied / Advanced; 67 percent of UCLA undergrads reported GenAI use in 2024 UCUES—brainstorming and topic research most common.²²² UCLA DTS: first California university OpenAI agreement at scale; campus Copilot and Gemini pilots. **Research assistance** that does not become unverified secondary citation; library-led evaluation (UCLA WI+RE tutorial). **Accessibility**: captioning, dictation, alternative formats—UDL 3.0 in higher ed. **Instructor-set spectrum** from ban to required use. Harvard FAS OUE: faculty **must** post a Canvas policy; three model policies (forbid; allow with citation; assignment-specific). HGSE school-level policy: limited learning-support use, acknowledgement, no submitting AI work as one's own.²²³ Cornell Chronicle, August 2026: AI Critical Literacy

²²⁰UCLA DTS: first California university OpenAI agreement at scale (vendor/admin claim as reported); campus Copilot and Gemini pilots; DTS AI Use and Recommendation Guide. Harvard managed environments. CA A.B. 1159 HESIPA: pending, not law.

²²¹Stanford CS221, Spring 2026 course policy: AI as human collaborator; no asking for answers; no “check my work”; no agentic repo-builders (Codex / Claude Code) for the project; share ChatGPT link/screenshots; tools without attribution banned.

²²²UCLA 2024 UCUES: 67 percent of undergrads reported GenAI use; brainstorming and topic research most common. BruinLearn AI Education Hub.

²²³Harvard FAS Office of Undergraduate Education: required syllabus/Canvas GenAI policy; three model policies (forbid; allow with citation; assignment-specific). HGSE school-level policy: limited learning-support use, acknowl-

Program expands from a spring pilot to **all incoming students**; four Canvas modules (What is GenAI?; Ethics; GenAI and Learning; Creating an AI Policy for Personal Use); CTI-reported literacy gains, **not a published RCT**; explicit goal: help students decide whether AI use **contributes to learning**.²²⁴ EDUCAUSE 2026 assessment survey: 438 faculty and staff; students need clearer use/non-use rules; AI literacy includes knowing when **not** to use tools. Full EDUCAUSE PDFs **not fetched**. US ED *Navigating AI in Postsecondary Education* (2025, ERIC ED670768): capacity-building for institutional AI task forces; equity; career-pathway retooling under EO 14110 (now historical). Not a campus implementation.

Syllabus governance is the college control surface the way the DPA is the K-12 control surface. Harvard requires the policy on Canvas. Stanford CS221 writes the collaborator rule into the assignment. Cornell asks whether the use contributes to learning, not whether the student can call an API. None of these is a national standard. All of them are existing, fetched, course- or campus-level practice. **No claim of a standard “autonomous campus agent” for students was found.** Governance is decentralized to the course, plus enterprise tenants for staff and research.

College adds, rather than repeats, the high-school profile.

	End of HS	End of college
Scope	Assigned tools, teacher-visible	Choose tools per discipline and data-class (public vs. FERPA/HIPAA/IP)
Integrity	Disclose when the teacher asks	Internalize field norms; cite; keep methods reproducible
Agency	Don't let it do the homework	Design workflows where the model is a critic, not an author of record
Research	Check a claim	Interrogate training-data limits, evaluation, and when <i>not</i> to use GenAI (original empirical work, confidential)
Civic	Know the teen rules	Can reason about high-risk uses (hiring, credit, admissions) and refuse to build them carelessly
Career	ISTE six roles at entry level	Domain-specific: code assistants with review, lab protocol checks, client-memo drafts with human sign-off

edgement, no submitting AI work as one's own.

²²⁴Cornell Chronicle, August 2026: AI Critical Literacy Program to all incoming students; four Canvas modules; CTI + Library; CTI-reported literacy gains from spring pilot, **not a published RCT**. EDUCAUSE 2026 assessment survey, 438 faculty/staff; full PDFs **not fetched**. US ED, *Navigating AI in Postsecondary Education* (2025), ERIC ED670768.

	End of HS	End of college
Legal self	Parent/school still in FERPA loop until 18	Student is the rights-holder; manages own data exhaust

Adult and continuing education is a brief bridge, not a sixth full band. FERPA if the provider is an ED-funded institution and the records are education records; otherwise consumer privacy (CCPA/CPRA in California, GDPR if EU, FTC Section 5). COPPA is irrelevant unless children are on the account. Adult executive function, but **skill atrophy and cognitive offloading** are the DfE concern even for adult learners; workplace policies increasingly require disclosure (parallel to Harvard and Stanford). Should not: silent use in credentialing exams; feeding employer or client confidential data to public models; companion-as-clinician. Can: just-in-time upskilling, language, accessibility, coding copilots with human review; the US ED July 2025 DCL explicitly allows AI for **college-and-career navigation** and hybrid tutoring in grant-funded programs. The UK government's AI tutoring ambition (up to 450,000 FSM pupils in Years 9-11) is **secondary, not adult ed**. Adult and continuing AI programs are mostly employer or MOOC practice. **No single public adult-ed autonomous-agent standard was found**. Use the college column, minus FERPA-eligible-student quirks if the provider is not an IHE.

The live debate this chapter takes a position on, without pretending a statute has settled it: Google youth onboarding versus AFT elementary opposition versus UNESCO 13 versus Singapore P1-3 no student chatbots and P4-6 only SLS education-designed tools. US ED's "age-appropriate" language without national standards, left to families and local authorities. Thirteen is a ToS number teens bypass by lying. There is **no US federal student-AI statute**. Age gates are a floor, not a pedagogy. Design to the stricter of law and development.

The portable floors, compressed from the instruments already named, are the design spec this chapter owes later chapters. They are sourced, not invented.

1. **No commercial model-training on student inputs.** NYCPS ERMA (December 2024), Broward contracts, UK DfE IP standard, FTC COPPA ed-tech guidance, CA SOPIPA commercial-profile ban. A.B. 1159 would make the AI-training ban explicit in California if enacted.
2. **No consequential decisions without a qualified human:** placement, discipline, eligibility, promotion, graduation, IEP/504, final grades, counseling/crisis (NYCPS red; US ED 2023 humans-in-the-loop; IDEA; EU Annex III 3(a)-(d)).
3. **No under-13 student-facing GenAI chat** unless COPPA-compliant, school-authorized, non-commercial, logged, and pedagogically bounded—and even then UNESCO and multiple large districts (Broward K-5) choose zero student-facing tools. Default: teacher-facing only below 13.
4. **No companion, romantic, or therapeutic AI for minors.** Character.AI (November 2025 withdrawal), OpenAI under-18 spec (2026), FTC 6(b) inquiry (September 2025), NYCPS counseling red, UK DfE anti-anthropomorphism (January 2026), AFT companion ban until 16.
5. **No open-web acting agents (browse-email-buy-post) for K-12 students.** None documented as approved; CIPA, safeguarding, and COPPA/FERPA all cut against it.

6. **No PII in unapproved consumer tools.** FERPA school-official control; NY 2-d contracts; NYCPS/Broward explicit.
7. **No biometric/emotion surveillance as instruction or discipline.** NYCPS under review and currently prohibited as “surveillance and behavior”; EU Art. 5; AAP child-centered versus engagement design.
8. **No AI-detector-only accusations.** Broward public FAQ; bias risk for multilingual and disabled writers.
9. **Separate verifiable parental consent for targeted ads / third-party disclosure (COPPA 2025).** Never integral to an education service.
10. **Age gates are a floor, not a pedagogy.** 13 (COPPA/UNESCO/ChatGPT) and 18 (Claude, Character.AI open chat, FERPA eligible-student) are **legal**, not developmental, lines.

A board that copies those ten lines into a procurement packet has not finished the pedagogical work. It has finished pretending that a tool lawful at 17 is therefore a tool for 8. The rest of this book will spend the struggle policy, the two-lane assessment rule, and the hosting default. It will not spend them as if age were an afterthought.

What actually differs by band is not a district SOP. It is this book’s compression of the statutes, the developmental constraints, and the documented deployments. Label it synthesis.

Band	Pedagogical center	AI’s legitimate job	Failure mode if ignored
5-8	Play, oral language, decoding, adult relationships	Teacher tool + literacy <i>about</i> AI	Anthropomorphism, displaced play, illegal data
9-11	Fluent reading, practice, beginning self-monitoring	Constrained, logged practice tutors	Homework machine, secret friend
12-14	Identity, EF growth, source skepticism	Assigned Socratic help + literacy <i>of</i> AI	Companion chat, offloading, social ranking
15-18	Original work, civic reasoning, college/career	Disclosed partner on hard tasks; full AI4K12/ISTE profile	Ghostwriting, detector injustice, pathway lock-in
18-22+	Disciplinary method, data classification	Collaborator with transcripts; domain workflows	Confidential-data leaks, admissions/proctoring bias
Adult	Performance in a role	Copilot with human sign-off	Deskilling, IP leakage

Who is the data controller of a child’s persistent agent memory, when the runtime is local, open-source, and parent-installed, is an open question this chapter will not fill. Childhood conversational-agent RCTs on language, attention, and social cognition were not found. Causal findings in Setzer and Raine are litigation. The EU Annex III education enforcement record is not complete. The

original 2 August 2026 high-risk clock was superseded by Regulation (EU) 2026/1744; stand-alone Annex III obligations now apply from 2 December 2027. Confirm consolidated EUR-Lex at type-setting. NAEYC's updated technology/AI statement is in revision. A dedicated AAP AI-in-schools policy does not exist. The NYCPS grade-band student-use playbook was still forthcoming as of the March 2026 page. The FERPA ed-tech NPRM (RIN 1875-AA15) had not been published as of the FTC's April 2025 COPPA preamble. Independent efficacy RCTs of MagicStudent and Gemini-for-Students at scale were not retrieved. Those holes stay standing. Law and age now constrain design. Subjects still have to keep struggle in the student. That is Chapter 5.

Chapter 5

Managed Struggle, Not Rescue and Not Abandonment




Figure 6: Algebra on yellow paper, a pencil, chalk, and a closed book.

Hiebert and Grouws, in 2007, defined productive struggle as “the intellectual effort students expend to make sense of mathematical concepts, to figure out something that is not immediately apparent.”²²⁵ NCTM’s 2014 *Principles to Actions* warned that teachers often “rescue” students too

²²⁵James Hiebert and Douglas A. Grouws, “The Effects of Classroom Mathematics Teaching on Students’ Learning,” in *Second Handbook of Research on Mathematics Teaching and Learning*, ed. Frank K. Lester Jr. (Charlotte, NC: Information Age, 2007), 387.

soon, converting a high-cognitive-demand task into a procedure.²²⁶ Two sentences, a decade apart, and they already contain the design problem this chapter is for. An agent that finishes the job is rescue at machine speed. An agent that stalls forever is abandonment. Neither is the target.

The pedagogical target is *managed struggle*: attempt gates, hint ladders, time-to-first-hint delays, teacher-authored misconception libraries. That is exactly the GPT Tutor design in Bastani, Bastani, Sungu, Ge, Kabakci, and Mariman's 2025 *PNAS* field experiment—the same interface as ChatGPT-style GPT-4, with the system prompt withholding full solutions and with teacher-written correct solutions and common-mistake hints in the prompt.²²⁷ A single “always Socratic” or “always answer” policy is incoherent with cognitive load theory. Bjork's desirable difficulties, Sweller's cognitive load theory, the expertise-reversal effect (Kalyuga, Ayres, Chandler, and Sweller, 2003), Hiebert and Grouws's productive struggle, and Kapur's productive failure all scale from childhood through college. The **expertise-reversal slider** moves. Childhood and early secondary use of AI should look like fading worked examples. College and expert-track use should reverse: less demonstration, more critique of the model's output, more generation-first.

Two lanes will close the chapter, labelled as this book's theoretically motivated synthesis, not as a fitted grade-by-grade schedule and not as existing district policy. Formation: AI forbidden or hint-only; grade unaided work. Performance in an AI world: AI allowed, logged; grade judgement—selection, critique, verification, process. Keep the struggle in the student. That is the job. Design is the treatment. “AI in education” is not one intervention.

The popular story will try to make this chapter a morality play: AI makes you dumber, or AI is the tutor of every child. Neither sentence survives the record. Contact with models is not what dulls. Three conditions together dull: the offloaded process is the learning objective; the environment never measures the skill unaided; the student cannot detect model error. Designs that keep generation, retrieval, and checking inside the student do not show the atrophy pattern on the outcomes they measured. Designs that hand the generative work to the model raise the artefact and often not the skill. Chapter 3 already spent the RCTs. This chapter's job is to turn them into a struggle policy by subject, and then into an assessment rule that does not pretend to be a district SOP. Subject by subject, the expertise-reversal slider will move. Assessment, at the end, is how the slider is made visible to a gradebook.

The cognitive science that should constrain design

Two distinctions run through every subject that follows. Performance is not learning. Current success with a tool in hand is a different construct from later, unaided skill (Soderstrom and Bjork, 2015; Bjork and Bjork, 2011).²²⁸ Generative AI inflates the first and can starve the second. And

²²⁶National Council of Teachers of Mathematics, *Principles to Actions: Ensuring Mathematical Success for All* (Reston, VA: NCTM, 2014), on teachers who “rescue” students too soon.

²²⁷Hamsa Bastani, Osbert Bastani, Alp Sungu, Haosen Ge, Ozge Kabakci, and Rei Mariman, “Generative AI without Guardrails Can Harm Learning: Evidence from High School Mathematics,” *Proceedings of the National Academy of Sciences* 122, no. 26 (2025): e2422633122.

²²⁸Nicholas C. Soderstrom and Robert A. Bjork, “Learning versus Performance: An Integrative Review,” *Perspectives on Psychological Science* 10, no. 2 (2015): 176-199; Elizabeth L. Bjork and Robert A. Bjork, “Making Things Hard on Yourself, but in a Good Way,” in *Psychology and the Real World*, ed. Morton Ann Gernsbacher et al. (New

design is the treatment. Unguarded ChatGPT, a hint-only tutor, a teacher-prompted Copilot session, Khanmigo sitting one click away, and a custom physics tutor with worked solutions in the prompt are different pedagogies that happen to share a model class.

Bjork's *desirable difficulties* (Bjork, 1994; Bjork and Bjork, 2011, 2020) are conditions that slow or worsen immediate performance while improving later retention and transfer: spacing, interleaving, generation, and retrieval practice.²²⁹ Soderstrom and Bjork's 2015 review of the experimental dissociation is the load-bearing paper: conditions that make practice look fluent—massed study, restudy, easy retrieval—often produce weaker long-term learning than conditions that feel harder. A Hermes-class or GrokBot-class assistant that supplies fluent, complete answers is, by default, an *undesirable ease*: it raises session performance and removes the retrieval, generation, and error-correction that build memory. The same assistant can be configured as a desirable difficulty if it withholds answers, forces generation, spaces review, and requires the student to reconstruct. Difficulty is desirable only when it is *achievable*. Bjork and Bjork (2020) are explicit: if the learner cannot succeed even with effort, the difficulty is undesirable. Corrective feedback that keeps the task in reach is part of the framework, not a betrayal of it.

Retrieval practice—testing as a learning event, not only as measurement—is among the most replicated findings in cognitive psychology (Roediger and Karpicke, 2006; Karpicke and Roediger, 2008).²³⁰ The testing effect is strongest on delayed tests and often looks worse than restudy on immediate tests. A qualifier from cognitive load theory matters for any tutor designer who has read only the retrieval papers. Van Gog and Sweller (2015) and Hanham, Leahy, and Sweller (2017) find that the testing effect shrinks or reverses for *high element-interactivity* material—complex, interacting concepts—among novices.²³¹ Chen, Castro-Alonso, Paas, and Sweller (2018) argue that many “failures” of desirable difficulties occur when working memory is already taxed. Pyke and colleagues' 2025 comparative analysis in the *Quarterly Journal of Experimental Psychology* similarly treats difficulty as moderated by element interactivity and expertise.²³² Retrieval practice via an AI tutor (quiz, free recall, explain-it-back) is a high-value default for facts, procedures, and low-interactivity material. For a novice facing a multi-step proof or a dense historical argument, unassisted generation can overload working memory. A worked example or a hint sequence is the better first move, with retrieval after a schema has begun to form.

York: Worth, 2011), 56-64.

²²⁹Robert A. Bjork, “Memory and Metamemory Considerations in the Training of Human Beings,” in *Metacognition*, ed. Janet Metcalfe and Arthur P. Shimamura (Cambridge, MA: MIT Press, 1994), 185-205; Bjork and Bjork, “Making Things Hard on Yourself”; Robert A. Bjork and Elizabeth L. Bjork, “Desirable Difficulties in Theory and Practice,” *Journal of Applied Research in Memory and Cognition* 9, no. 4 (2020): 475-479.

²³⁰Henry L. Roediger III and Jeffrey D. Karpicke, “Test-Enhanced Learning: Taking Memory Tests Improves Long-Term Retention,” *Psychological Science* 17, no. 3 (2006): 249-255; Jeffrey D. Karpicke and Henry L. Roediger III, “The Critical Importance of Retrieval for Learning,” *Science* 319, no. 5865 (2008): 966-968.

²³¹Tamara van Gog and John Sweller, “Not New, but Nearly Forgotten: The Testing Effect Decreases or Even Disappears as the Complexity of Learning Materials Increases,” *Educational Psychology Review* 27 (2015): 247-264; Jose Hanham, Wayne Leahy, and John Sweller, “Cognitive Load Theory, Element Interactivity, and the Testing and Reverse Testing Effects,” *Applied Cognitive Psychology* 31, no. 3 (2017): 265-280.

²³²Ouhao Chen, Juan C. Castro-Alonso, Fred Paas, and John Sweller, “Extending Cognitive Load Theory to Incorporate Working Memory Resource Depletion,” *Educational Psychology Review* 30 (2018): 483-501; Aryn A. Pyke et al., “Does Difficulty Moderate Learning? A Comparative Analysis of the Desirable Difficulties Framework and Cognitive Load Theory,” *Quarterly Journal of Experimental Psychology* (2025).

Sweller's cognitive load theory partitions load into intrinsic (inherent complexity given the learner's current schemas), extraneous (load imposed by poor design: split attention, redundant text, hunting for the next step), and germane (resources devoted to schema construction—now often treated as the use of remaining working-memory capacity rather than a third additive pool).²³³ The **worked-example effect** (Sweller and Cooper, 1985; Paas and van Merriënboer, 1994, 2020): for novices, studying a fully worked solution plus a similar problem outperforms unguided problem solving, because it frees working memory for schema construction.²³⁴ The **expertise reversal effect** (Kalyuga, Ayres, Chandler, and Sweller, 2003; Kalyuga, 2007): the same high-guidance materials that help novices *impair* more knowledgeable learners, who now need problem-solving, variability, and reduced scaffolding.²³⁵ Childhood and early secondary use of AI should look like *fading worked examples*: model, then complete-the-step, then independent problem, with the agent forbidden from dumping the full solution once the student can generate. College and expert-track use should reverse: less demonstration, more critique of the model's output, more generation-first.

Productive struggle versus unproductive struggle is the classroom translation. Hiebert and Grouws (2007, p. 387) is the definition already on the table. Warshauer (2014) maps teacher moves that *maintain* demand while addressing stuckness.²³⁶ Kapur's **productive failure** (Kapur, 2014, *Cognitive Science*): having ninth-graders generate (often incorrect) solutions *before* instruction improved conceptual knowledge relative to instruction-first, because the struggle activated prior knowledge and prepared students to notice the critical features of the canonical solution.²³⁷ Unproductive struggle is overload, confusion about the goal, or frustration from tasks that are not mathematically—or historically, scientifically—appropriate. AI that answers on the first hint request converts productive struggle into unproductive ease. AI that withholds *all* help until the student is drowning converts productive struggle into unproductive struggle. Warshauer's teacher moves that maintain demand—probing, side-by-side comparison of student methods, asking a student to explain a stuck peer's idea without giving the path—are the human version of a hint ladder. An agent can be configured to do a crude version of those moves: ask for the attempt, name a common mistake, withhold the last step. It can also be configured, by default, to skip them. Most consumer chat is the skip. GPT Tutor is the crude version of the probe. LearnLM, with a human on every message, is the probe with a supervisor. Khanmigo's stated design is Socratic; Oreopoulos and Low's message data show that most students did not use it as a probe. Stated design is not the treatment. Use is.

²³³John Sweller, "Cognitive Load during Problem Solving: Effects on Learning," *Cognitive Science* 12, no. 2 (1988): 257-285; John Sweller, "Cognitive Load Theory," in *The Psychology of Learning and Motivation*, vol. 55, ed. Jose P. Mestre and Brian H. Ross (San Diego: Academic Press, 2011), 37-76; John Sweller, Paul Ayres, and Slava Kalyuga, *Cognitive Load Theory* (New York: Springer, 2011).

²³⁴John Sweller and Graham A. Cooper, 1985, worked-example effect; Fred Paas and Jeroen J. G. van Merriënboer, "Variability of Worked Examples and Transfer of Geometrical Problem-Solving Skills," *Journal of Educational Psychology* 86, no. 1 (1994): 122-133; Paas and van Merriënboer, "Cognitive-Load Theory: Methods to Manage Working Memory Load in the Learning of Complex Tasks," *Current Directions in Psychological Science* 29, no. 4 (2020): 394-398.

²³⁵Slava Kalyuga, Paul Ayres, Paul Chandler, and John Sweller, "The Expertise Reversal Effect," *Educational Psychologist* 38, no. 1 (2003): 23-31; Slava Kalyuga, "Expertise Reversal Effect and Its Instructional Implications," *Instructional Science* 35 (2007): 95-97.

²³⁶Hiroko K. Warshauer, "Productive Struggle in Middle School Mathematics Classrooms," *Journal of Mathematics Teacher Education* 18 (2014): 375-400.

²³⁷Manu Kapur, "Productive Failure in Learning Math," *Cognitive Science* 38, no. 5 (2014): 1008-1022.

Kapur's productive failure is easy to misuse as a reason to withhold all help. Ninth-graders generating incorrect solutions *before* instruction is not ninth-graders drowning in a chatbot that will not hint. The generation in Kapur is bounded, the subsequent instruction is canonical, and the point is to prepare the student to notice critical features. Unbounded generation into a fluent model that will complete the solution is the opposite preparation: the critical features are in the model's paragraph, and the student has not had to notice them. Productive failure is a formation-lane design. Rescue-by-autocomplete is not a productive-failure design with a modern interface. It is NCTM's warning, automated.

The design translation is not a slogan. It is a table.

Cognitive principle	Sharpening design	Dulling design
Desirable difficulty / retrieval	Generate-then-check; delayed tests without AI; spaced re-ask	Instant complete answers; copy-paste homework
Worked examples (novices)	Stepwise worked solution, then isomorphic problem without the solution	Open problem-solving with a chatbot that will finish it
Expertise reversal (intermediates+)	Critique of AI output; generation-first; faded hints	Continuing to show full solutions
Productive struggle	Attempt gate; hint ladder; common-mistake library	Rescue on first pause <i>or</i> infinite Socratic stall
Cognitive load	One question at a time; accuracy of feedback; hide unused parts	Hallucinated steps; multi-part dump; split attention across tabs
Learning vs performance	Grade unaided work; process logs	Grade AI-assisted artefacts as if they were skill

That table is a compression of the papers just named. It is not a district policy. It is the constraint the subject sections have to meet.

A single “always Socratic” policy looks like productive struggle and is, for a novice on high-element-interactivity material, often unproductive struggle: the student is generating before a schema exists, working memory is already taxed, and the testing effect that Roediger and Karpicke measured on low-interactivity material shrinks or reverses (van Gog and Sweller, 2015; Hanham, Leahy, and Sweller, 2017). A single “always answer” policy looks like a worked example and is, once the student can generate, expertise reversal in the wrong direction: the high-guidance that helped the novice now impairs the intermediate (Kalyuga et al., 2003). The slider has to move. Fading is the mechanism. Model, then complete-the-step, then independent problem. The agent is forbidden from dumping the full solution once the student can generate. College reverses the same slider: generation-first, then critique of the model's output, then—only if the student can detect model error—the second lane, in which the model is a declared tool and judgement is what is graded.

Deslauriers, McCarty, Miller, Callaghan, and Kestin (2019) documented the illusion of learning: students in a passive lecture felt they had learned more than students in an active classroom, while

learning less.²³⁸ Bastani's GPT Base students did not perceive that they had learned less after looking 48 percent better in practice. The illusion is the same object. Fluency of the session is not a measure of the skill. Any assessment regime that grades only the session artefact has hired the illusion as a teaching assistant.

Risko and Gilbert (2016) defined cognitive offloading as using physical action or the environment to reduce the internal processing a task requires.²³⁹ Offloading is not a moral failing. It is how human cognition scales. Calculators and GPS produce analogous trade-offs: performance up, unaided skill down if the offloaded component was the learning target. Generative AI extends offloading from memory and calculation to **reasoning, drafting, and evaluation**—the skills schools exist to form. The research questions are what triggers offloading (often a metacognitive judgement that internal demand is high) and what it does to later unaided performance. Frequency of AI use is not the variable. Whether the user retains the evaluative loop is. A 2026 conceptual distinction in *Frontiers in Psychology* (dependent versus autonomous offloading) is useful even before that paper's psychometrics settle: the same tool can be either, depending on whether generation, retrieval, and checking stayed inside the student.²⁴⁰

Mathematics is where the causal record is thickest. Start there, and do not re-litigate Chapter 3's full RCTs. Spend the page on design moves.

Mathematics: hint ladders help; answer-giving GPT harms; access is not tutoring

Help: hint-ladder tutors with grounded solutions. Bastani et al. 2025, GPT Tutor: same interface as GPT Base, system prompt withholds full solutions, teacher-written correct solutions and common-mistake hints. Assisted practice about 127 percent better than control; unaided exam indistinguishable from control (intention-to-treat, classroom-clustered, nearly 1,000 students, grades 9-11, four 90-minute sessions).²⁴¹ Kestin, Miller, Klates, Milbourne, and Ponti 2025, *Scientific Reports*: custom GPT-4 tutor ("PS2 Pal") with instructor-written step-by-step solutions in the prompt, sequential scaffolding, growth-mindset and cognitive-load language, versus the course's existing in-class active learning, not lecture. Immediate pre/post on surface tension and fluid flow; linear regression 0.63 SD; quantile regression, because of ceiling, 0.73 to 1.3 SD. Delayed retention unmeasured. Not off-the-shelf ChatGPT. Not an agent.²⁴² LearnLM Team, Google DeepMind and Eedi, 2025 preprint: expert tutors supervised every draft; immediate mistake-remediation 93.0 percent versus

²³⁸Deslauriers et al., *PNAS* 116, no. 39 (2019): illusion of learning.

²³⁹Evan F. Risko and Sam J. Gilbert, "Cognitive Offloading," *Trends in Cognitive Sciences* 20, no. 9 (2016): 676-688.

²⁴⁰*Frontiers in Psychology* 2026, dependent versus autonomous offloading (abstract-level): frequency is not the variable; whether the user retains the evaluative loop is.

²⁴¹Bastani et al., "Generative AI without Guardrails," *PNAS* 122, no. 26 (2025). Intention-to-treat, classroom-clustered: GPT Tutor practice +0.361 (about 127 percent better); unaided exam -0.004, indistinguishable from control. Field RCT, large Turkish high school, Fall 2023-24, nearly 1,000 students, grades 9-11, four 90-minute sessions covering about 15 percent of semester math.

²⁴²Greg Kestin, Kelly Miller, Anna Klates, Timothy Milbourne, and Gregorio Ponti, "AI Tutoring Outperforms In-Class Active Learning," *Scientific Reports* 15 (2025): 17458. Crossover RCT, Harvard Physical Sciences 2, Fall 2023, N = 194 eligible of 233 enrolled. Delayed retention, transfer, and full-course outcome not measured. Not off-the-shelf ChatGPT. Not an agent.

human tutor 91.2 percent versus static hints 65.4 percent; next-topic transfer +5.5 percentage points versus human, session-level randomization so cumulative learning is not identified; human-in-the-loop is literal (every message gated).²⁴³ Henkel, Horne-Robinson, Kozhakhmetova, and Lee 2024: Rori, a WhatsApp AI math tutor, about 500 Ghanaian students, extra one hour per week, Cohen's $d = 0.36$, marginal cost about \$5 per student, year-1, authors flag caution.²⁴⁴

Harm: answer-giving GPT. Bastani GPT Base: assisted practice about 48 percent better; unaided exam about 17 percent worse than never-AI. GPT Base returned a correct answer only about 51 percent of the time in a probe; students still copied. Students in GPT Base did not perceive that they had learned less (Deslauriers-style illusion of learning; Deslauriers et al., 2019).²⁴⁵ Guardrails are the treatment. Default ChatGPT is a crutch students do not notice.

Neutral-to-small at scale when students do not talk to the tutor. Oreopoulos and Low 2026, two-year cluster RCT, 18 Hamilton County, Tennessee middle schools, Khan Academy plus Khanmigo (configured to coach, not tell) during existing daily RTI math versus business-as-usual RTI. Sample: 6,902 student-terms, 53 clusters. ITT: about 0.040 population SD per term; stacked about 0.06 SD per school year, about 0.08 in year 2. Magnitudes resemble Khan Academy practice without the AI tutor (authors cite Eames et al., 2026). Median student messaged on one-third of practice days and in 17 percent of mistake sessions. Only 14.5 percent of messages contained a mathematical question or a step of reasoning. A 2026 interface change that opened Khanmigo by default did not, in their message data, raise messaging rates. Access is not tutoring. Robinson, Gormley, Ribeiro, and Loeb 2026: "Access Is Not Enough."²⁴⁶ Do not treat Khanmigo ITT as "AI tutoring works at 0.07 SD." Authors say it looks like practice without the chatbot, and engagement is thin. State TCAP results were registered as co-primary and were not yet incorporated as of the August 2026 paper. Year-scale unaided skill on independent state tests is not yet known.

Human in the loop, in the positive studies, is specific. Teacher-authored common mistakes in the prompt (Bastani). Teacher circulating in the lab (De Simone et al. 2025, Nigeria—English, not math, but the session architecture is the same: science-of-learning prompts, supervised lab time, curriculum alignment). Tutor supervising every draft (LearnLM). Tutor-facing CoPilot (Wang, Ribeiro, Robinson, Loeb, and Demszky: about 900 tutors and 1,800 K-12 students; students +4 percentage points more likely to master the session topic; +9 percentage points for lower-rated tutors; about \$20 per tutor per year; back-of-envelope MAP 0.024 SD for a full year, explicitly not

²⁴³LearnLM Team, Google DeepMind, and Eedi, "AI Tutoring Can Safely and Effectively Support Students: An Exploratory RCT in UK Classrooms," arXiv:2512.23633 (2025). **Preprint.** $N = 165$, five UK secondary schools. Session-level randomization; cumulative learning not identified; every message gated. Bayesian reporting.

²⁴⁴Owen Henkel, Hannah Horne-Robinson, Neeve Kozhakhmetova, and Adam Lee, "Effective and Scalable Math Support," in *AIED 2024*, CCIS 2150 (Cham: Springer, 2024). About 500 Ghanaian students; Cohen's $d = 0.36$; about \$5/student; year-1; authors flag caution.

²⁴⁵Bastani et al. 2025: GPT Base practice +0.137 (about 48 percent better); unaided exam -0.054 (about 17 percent worse than never-AI); probe correctness about 51 percent, students still copied. Louis Deslauriers, Logan S. McCarty, Kelly Miller, Kristina Callaghan, and Greg Kestin, "Measuring Actual Learning versus Feeling of Learning," *PNAS* 116, no. 39 (2019): 19251-19257.

²⁴⁶Oreopoulos and Low, "One Click Away," EdWorkingPaper 26-1551 / NBER w35620 (2026). 6,902 student-terms, 53 clusters, 18 Hamilton County, TN middle schools; pre-registered AEARCTR-0013519; NWEA MAP math; TCAP registered as co-primary, not yet incorporated as of August 2026. Carly D. Robinson, David Gormley, Ana Trindade Ribeiro, and Susanna Loeb, "Access Is Not Enough," EdWorkingPaper 26-1451 (2026). Eames et al. 2026 cited by Oreopoulos and Low for Khan Academy practice without the AI tutor.

a direct experimental estimate).²⁴⁷ The human’s job is not to re-teach every item. It is to set the struggle policy and to intervene when engagement or misconception diagnosis fails.

Productive-struggle translation for software: delay the first hint; require an attempt; then a specific hint, not a solution. That is Hiebert-Grouws implemented in a prompt. It is also, in one sentence, why a single “always Socratic” policy fails novices on high-element-interactivity material and why a single “always answer” policy fails everyone the next hour, when the window is closed.

A *hypothetical*, labelled as such: a ninth-grader pastes a quadratic-word-problem into an unguarded chatbot, copies a worked solution that is wrong in the arithmetic, and scores well on the homework. Bastani’s probe is the empirical version of that design. GPT Tutor’s attempt-then-hint ladder is the alternative. The difference is not the model. It is whether the generative work stayed inside the student.

Human high-dosage tutoring remains the benchmark. Nickow, Oreopoulos, and Quan (2024, *American Educational Research Journal*) pool experimental preK-12 tutoring at 0.288 SD (SE 0.029).²⁴⁸ Earlier computer-assisted learning typically sits in 0.05-0.20 SD (Escueta, Nickow, Oreopoulos, and Quan, 2020).²⁴⁹ Pre-generative intelligent tutoring systems show positive averages (Kulik and Fletcher, 2016).²⁵⁰ Wang, Wang, and colleagues (2024, *Journal of Educational Computing Research*) meta-analyse AI-enabled *adaptive* systems 2010-2022 (mostly pre-ChatGPT) at $g = 0.70$ versus non-adaptive comparators—a large estimate that is not a generative-AI result and is moderated by publication type, duration, and outcome.²⁵¹ Bloom’s two-sigma observation (1984) was about **one-to-one human tutoring**.²⁵² Every positive generative-AI field result in Chapter 3 either embeds the model in a teacher-run session (Nigeria, Kestin’s course staff, LearnLM supervisors, RTI block with an adult in the room) or helps the *tutor* (Wang et al. CoPilot). Khanmigo without engagement looks like software, not a tutor. Replacement is the wrong theory of this evidence. Orchestration—spent as a labor chapter in Chapter 7—is what the positive studies already quietly did.

The Monday translation for a math teacher is therefore not “use AI” or “ban AI.” It is: choose the lane; configure the struggle policy (attempt gate, hint ladder, grounded answers, teacher-authored common mistakes); be in the room for engagement and for when struggle turns unproductive; run

²⁴⁷Rose E. Wang, Ana T. Ribeiro, Carly D. Robinson, Susanna Loeb, and Dorottya Demsky, “Tutor CoPilot: A Human-AI Approach for Scaling Real-Time Expertise,” arXiv:2410.03017 (2024). About 900 tutors, 1,800 K-12 students; +4 pp topic mastery ($p < 0.01$); +9 pp for lower-rated tutors; about \$20 per tutor per year; MAP 0.024 SD back-of-envelope, explicitly not a direct experimental estimate. Nigeria session architecture: Martin De Simone et al., World Bank Policy Research Working Paper 11125 (2025).

²⁴⁸Andre Nickow, Philip Oreopoulos, and Vincent Quan, “The Promise of Tutoring for PreK-12 Learning,” *American Educational Research Journal* 61, no. 1 (2024): 74-107. Pooled experimental preK-12 tutoring 0.288 SD (SE 0.029).

²⁴⁹Maya Escueta, Andre J. Nickow, Philip Oreopoulos, and Vincent Quan, “Upgrading Education with Technology: Insights from Experimental Research,” *Journal of Economic Literature* 58, no. 4 (2020): 897-996. Earlier CAL typically 0.05-0.20 SD.

²⁵⁰James A. Kulik and J. D. Fletcher, “Effectiveness of Intelligent Tutoring Systems: A Meta-Analytic Review,” *Review of Educational Research* 86, no. 1 (2016): 42-78.

²⁵¹Shuai Wang et al., “The Efficacy of Artificial Intelligence-Enabled Adaptive Learning Systems from 2010 to 2022 on Learner Outcomes: A Meta-Analysis,” *Journal of Educational Computing Research* (2024), $g = 0.70$ vs non-adaptive. **Not a generative-AI result.**

²⁵²Benjamin S. Bloom, “The 2 Sigma Problem: The Search for Methods of Group Instruction as Effective as One-to-One Tutoring,” *Educational Researcher* 13, no. 6 (1984): 4-16. About human one-to-one tutoring.

the unaided check that makes outsourcing visible. GPT Tutor is that translation in a system prompt. Nigeria is that translation in a lab. LearnLM is that translation with a human on every message. Khanmigo sitting one click away, in an RTI block where the median student messages on one-third of days and only 14.5 percent of messages contain a mathematical question, is access. Access is not tutoring. Design is the treatment.

Writing: augmentation versus automation

Writing is where the artefact most seduces the grader. Help: targeted feedback on a *student* draft; grammar as a low-stakes offload once ideas exist; brainstorming that the student must then select from; the *augmentation* pattern in Contractor and Reyes 2026 (AI to explain concepts). Harm: first-draft generation. Fan, Tang, Le, Shen, Tan, Zhao, Shen, Li, and Gasevic 2025, *British Journal of Educational Technology*, N = 117 university students, four supports (ChatGPT, human expert chat, writing-analytics tools, none): ChatGPT group improved essay scores more; **no significant differences in knowledge gain or transfer**; process mining showed tight revise-ask-AI loops with **fewer metacognitive processes (evaluation, orientation)** than the human-expert group. The authors name this “metacognitive laziness.”²⁵³ Kosmyna et al. 2025, arXiv:2506.08872, MIT Media Lab, **preprint**, N = 54 / session 4 n = 18: LLM group low perceived ownership, poor ability to quote their own essay (session 1: 15 of 18 could not quote; 0 of 18 quoted correctly); connectivity scaled with how much the tool did. Do not cite as “MIT proved ChatGPT atrophies the brain.” Small N; convenience sample; 20-minute essays; EEG connectivity is not a direct measure of “getting dumber”; no delayed knowledge test of the Bastani or Contractor kind.²⁵⁴ Contractor and Reyes 2026, IZA Discussion Paper 18792, **working paper, July 2026, not a journal article**: AI access raises immediate knowledge-test scores 0.27 SD, persisting one week; delayed unaided essay style about 0.30 SD and relevance about 0.26 SD improve for **augmentation** users (AI to explain), not **automation** users (AI to generate text), whose short-run quality gains vanish when AI is removed.²⁵⁵ Full IZA PDF fetch was partial in the research pass; coefficients from the abstract/IZA page.

Human in the loop: the teacher sets whether this assignment is “no AI,” “AI for critique of *your* draft,” or “AI as co-author with an oral defence.” Process portfolios (prompts, iterations, rationale) make the loop visible. EFL quasi-experiments in 2025-26 often show criterion-specific gains—brainstorming helps task response and organisation; corrective feedback helps grammar and vocabulary—and delayed no-AI transfer when the design forces re-authoring. These are useful. They are mostly intact-class quasi-experiments, not the same evidential grade as Bastani or De

²⁵³Yizhou Fan et al., “Beware of Metacognitive Laziness: Effects of Generative Artificial Intelligence on Learning Motivation, Processes, and Performance,” *British Journal of Educational Technology* 56, no. 2 (2025): 489-530. N = 117. Full BJET PDF publisher-blocked in one fetch; ERIC abstract used—do not invent additional coefficients.

²⁵⁴Nataliya Kosmyna et al., “Your Brain on ChatGPT: Accumulation of Cognitive Debt when Using an AI Assistant for Essay Writing Task,” arXiv:2506.08872 (10 June 2025). **Preprint, not a journal article** as of 30 August 2026. N = 54 Boston-area adults; session 4 n = 18. Do not cite as population cognitive decline.

²⁵⁵Zara Contractor and Germán Reyes, “Experimental Evidence on the Learning Impact of Generative AI,” IZA Discussion Paper 18792 (July 2026). **Working paper, not a journal article**. Immediate knowledge-test +0.27 SD, persisting one week; delayed unaided essay style about 0.30 SD and relevance about 0.26 SD for augmentation users, not automation users. Full IZA PDF fetch was partial; coefficients from abstract/IZA page.

Simone. Do not promote them to RCTs.

De Simone, Tiberti, Barron Rodriguez, Manolio, Mosuro, and Dikoru 2025, World Bank Policy Research Working Paper 11125: after-school Microsoft Copilot (GPT-4), 12 times 90-minute sessions over six weeks, first-year senior-secondary students in Benin City, Nigeria, English curriculum. Teachers trained three days; sessions used **science-of-learning prompts** (retrieval, elaborative interrogation, concrete examples; Mollick and Mollick, 2023) rather than answer-giving. Students often paired on one computer. ITT: English (primary) 0.238 SD; regular third-term English exam 0.206 SD; composite 0.31 SD. Heterogeneity: larger for females and higher baseline/SES students. Attendance about 72 percent. Authors convert English 0.24 SD to about 1.5 equivalent years of typical Nigerian schooling (Evans and Yuan metric)—a **conversion, not a longitudinal follow-up**. Treatment is Copilot plus teacher-written tutor prompts plus supervised lab plus curriculum alignment. Students were taught that Copilot hallucinates and that they remain responsible.²⁵⁶ The equity sting is in the heterogeneity: a “personal tutor for every child” that widens gaps is an empirical possibility, not a slogan.

Perkins, Furze, Roe, and MacVaugh’s 2024 AI Assessment Scale enters here as shared language for the assignment label, not as efficacy research. Full treatment in the last section of this chapter. The writing implication is already available: grade the process, not the polish. An essay that looks college-ready while encoding is shallow is Chapter 3’s performance/learning split in prose form.

A process portfolio is not a folder of drafts for its own sake. It is how the evaluative loop is made visible. Prompts, iterations, a short rationale for what was kept and what was cut, and—when the lane is formation—an unaided timed write or an oral defence. Fan’s process mining is the warrant: ChatGPT users formed tight revise-ask-AI loops with fewer metacognitive processes. If the grade is only the essay, the course has rewarded the loop that skipped evaluation and orientation. If the grade includes the log and the unaided check, the course has a chance of rewarding the loop that kept judgement inside the student. Contractor and Reyes’s augmentation users shifted time from drafting toward reading and searching and showed delayed unaided gains on style and relevance. Automation users’ short-run quality gains vanished when AI was removed. Grading *how* AI was used is not a moralism. It is the empirical split those authors measured.

The assignment label is the cheap version of the same idea. AFT’s full / partial / none, Perkins’s AIAS levels, Harvard OUE’s three model policies (forbid; allow with citation; assignment-specific), Stanford CS221’s collaborator-not-answer-key: these are vocabularies. They do not, by themselves, keep struggle in the student. A “partial AI” label on a take-home essay that can be fully generated is a vocabulary without a formation lane. A “no AI” label on every assignment is a formation lane without a performance-in-an-AI-world lane, and it invites underground use. Two lanes is the claim that both are the job, at different moments of expertise, and that the grade has to match the lane. Label it synthesis. Do not present it as Broward policy or as NYCPS policy. Broward refuses detector-only decisions and keeps Writable as revision aid with the teacher still grading. NYCPS’s homework/integrity playbook was still forthcoming as of the March 2026 page. The two-lane rule

²⁵⁶De Simone, Tiberti, Barron Rodriguez, Manolio, Mosuro, and Dikoru, PRWP 11125 (2025). ITT: composite 0.31 SD (SE 0.068); English 0.238 SD; term exam 0.206 SD. Heterogeneity: larger for females and higher baseline/SES. Attendance about 72 percent. Evans and Yuan conversion of 0.24 SD to about 1.5 equivalent years: a conversion, not a longitudinal follow-up. Ethan Mollick and Lilach Mollick, “Assigning AI: Seven Approaches for Students, with Prompts,” arXiv:2306.10052 (2023), used in the Nigeria sessions.

is this book's reading of Bastani, Fan, Contractor and Reyes, and the cognitive science, not a fetched SOP.

Science: structured tutoring is not open inquiry with a chatbot

Kestin et al. 2025 is the flagship *help* result for first-pass conceptual physics under a custom tutor versus strong active learning, on an immediate post-test only. Authors report median learning gains more than double the in-class arm; Mann-Whitney $z = -5.6$, $p < 10^{-8}$. Median time on task 49 versus 60 minutes. Higher self-reported engagement and motivation; enjoyment and growth-mindset items not different. **Not measured:** delayed retention, transfer, or a full-course outcome. Authors state they do not claim superiority for higher-order synthesis. Lessons were Bloom understanding/applying/analysing, not creating. Population is a selective university, though FCI-stratified subgroups both showed the AI advantage.²⁵⁷ This is strong evidence that a pedagogically engineered tutor, on first-pass conceptual material, can beat even a good classroom for immediate mastery. It is not evidence that off-the-shelf ChatGPT does so, nor that the gain lasts, nor that an agent in a lab coat would.

Stadler, Bannert, and Sailer 2024, *Computers in Human Behavior*, $N = 91$ university students, ChatGPT-3.5 versus Google to research nanoparticles in sunscreen and write recommendations: LLM users reported **lower cognitive load** and produced **lower-quality reasoning and argumentation**. Homogeneity of conclusions did not differ.²⁵⁸ This is the scientific-inquiry analogue of Bastani's "ease without depth." Together the two studies say: **structured tutoring on a known learning objective is not open inquiry with a general chatbot.**

Lab work, experimental design, and noticing anomalous data are poorly suited to fully autonomous agents. The model cannot see the benchtop unless multimodal, and even then the student's motor and observational skill is the point. Human in the loop here is the lab partner and the instructor, with AI as pre-lab retrieval practice and post-lab explanation checker. That last sentence is a **design inference from cognitive load theory and productive-struggle principles, not a fetched lab RCT**. No agent-in-the-lab learning RCT was identified. Replication of Kestin outside a highly staffed university physics course is unshown. Do not treat Kestin as evidence that ChatGPT beats classroom teaching.

Kestin's comparator was not a weak lecture. It was the course's existing in-class active learning. That is what makes the immediate result interesting, and what makes the missing delayed test costly. A custom tutor that beats a good classroom at minute 49 on surface tension is a design result about first-pass conceptual material under instructor-written steps. It is not a result about ChatGPT. It is not a result about a year of physics. It is not a result about synthesis. Authors are explicit on the last two. A book that quotes the 0.73 to 1.3 SD quantile range without those limits has already done the vendor's work. Stadler's lower load and worse argument is the complementary limit: when the task is inquiry, a general chatbot is not a tutor. It is a fluency machine. The two papers are often

²⁵⁷Kestin et al., *Scientific Reports* 15 (2025): 17458. Median post 4.5 (AI) vs 3.5 (in-class) vs combined pre 2.75; Mann-Whitney $z = -5.6$, $p < 10^{-8}$; median time 49 vs 60 minutes.

²⁵⁸Matthias Stadler, Maria Bannert, and Michael Sailer, "Cognitive Ease at a Cost: LLMs Reduce Mental Effort but Compromise Depth in Student Scientific Inquiry," *Computers in Human Behavior* 160 (2024): 108386. $N = 91$; ChatGPT-3.5 vs Google.

cited as if they were a single “AI in science” finding. They are opposite designs. Treating them as one intervention is how “AI in education” becomes a mass noun again, which Chapter 3 already forbade.

The sharpening design for science is therefore two different objects, not one “science AI.” For first-pass conceptual material (surface tension, fluid flow, a worked circuit): a custom tutor with instructor-written steps, sequential scaffolding, and an unaided check afterward. For inquiry (what should we conclude about nanoparticles in sunscreen; what does this anomalous data mean): the student searches, reads, and argues; a general chatbot that lowers load and thins the argument is the dulling design Stadler already measured. For the bench: the student’s hands and eyes. An agent that writes the lab report overnight is rescue. An agent that will not hint when the student is drowning in a derivation is abandonment. Managed struggle, again, is the space between.

Reading and history: do not hide the work

Direct generative-AI RCT evidence on *learning to read*—decoding, fluency, comprehension—in 2022-26 is **thin** compared with math. That thinness is itself a finding. Pre-LLM intelligent tutoring systems for reading comprehension (for example, ITSS structure strategy; Wijekumar, Meyer, and Lei, 2012) showed gains on researcher-designed measures and mixed standardized-test results in What Works Clearinghouse review—relevant as a caution that software gains often shrink on independent tests.²⁵⁹ IES WWC has intervention reports on older ITS products (for example, Cognitive Tutor Algebra I). As of 30 August 2026 it had **no WWC intervention report specific to generative AI tutors**.

Risks specific to reading. AI summaries skip the decoding-to-comprehension work. Listening to a model “explain the chapter” is not reading. Unsupervised generative AI that rewrites text on demand can **hide** decoding difficulty rather than teach it—the dyslexia / structured-literacy failure already named in Chapter 4. Help: vocabulary-in-context; retrieval questions *after* the student has read; text-dependent questions the model is forbidden to answer until the student quotes the line. Willingham’s knowledge-comprehension loop implies that background-knowledge building (including via oral AI dialogue *about* a text the student has read) can support later reading—provided the text itself was read.²⁶⁰ Oreopoulos and Low cite Robinson et al. 2026 on an adaptive literacy platform: dedicated time still produced only minutes of weekly use; adding human tutors raised usage without raising achievement. Engagement failures are not unique to Khanmigo.

History has almost no subject-specific RCT as of this access date. The cognitive analysis is nonetheless sharp. Historical thinking is sourcing, corroboration, and contextualisation (Wineburg’s line of work; **not re-fetched as a 2022-26 AI trial**).²⁶¹ An LLM that writes a confident narrative is a primary-source problem. The sharpening design: student reads sources first; model produces a

²⁵⁹Kay Wijekumar, Bonnie J. F. Meyer, and Puiwa Lei, ITSS structure-strategy tutoring (2012); WWC review: gains on researcher-designed measures, mixed standardized-test results. IES WWC: older ITS reports exist (e.g., Cognitive Tutor Algebra I); **no WWC intervention report specific to generative AI tutors** as of 30 August 2026.

²⁶⁰Daniel T. Willingham, “Critical Thinking: Why Is It So Hard to Teach?,” *American Educator* 31, no. 2 (2007): 8-19, knowledge-comprehension loop. Full critical-thinking treatment in Chapter 6.

²⁶¹Sam Wineburg’s sourcing, corroboration, contextualisation line; **not re-fetched as a 2022-26 AI trial**. Almost no history-subject generative-AI RCT as of the access date.

narrative; student must mark what is sourced, what is inferred, and what is fabricated; oral defence. The dulling design is “write a DBQ with ChatGPT.” Do not invent a history or reading generative-AI RCT. The hole stays standing. Math has trials. These subjects mostly have mechanisms. An honest book says so and still specifies the struggle policy, because waiting for a history RCT before refusing to hide the work is how the work gets hidden.

Structured literacy makes the reading case even without a generative-AI RCT. If the objective is decoding, a tool that rewrites the passage at a lower level on demand has not differentiated instruction. It has removed the grapheme-phoneme work. Broward’s highest-evidence accessibility use—real-time oral-reading feedback—is constrained practice, logged, not open chat. Chapter 4 already put that distinction in the accessibility cross-cut. This chapter’s job is to say why it is also a struggle policy. Feedback on the student’s oral reading keeps the generative work (decoding) inside the student. A rewrite-on-demand chatbot moves the generative work into the model. UDL 3.0’s learner agency is not “the model read it for you.” It is multiple means of engagement, representation, and action, with high expectations. Hiding decoding difficulty is not an accommodation. It is a different learning objective, unannounced.

A *hypothetical*, labelled as such: a tenth-grader uploads a chapter PDF, asks for a summary, and never reads the chapter. The homework looks done. Decoding, fluency, and the knowledge-comprehension loop were the learning objectives. The offloaded process *was* the skill. That is the first of the three conditions that together dull (the offloaded process is the learning objective; the environment never measures the skill unaided; the student cannot detect model error). Frequency of use is not the variable. Whether the evaluative loop stayed inside the student is.

Computer science: Copilot as worked examples at keystroke speed

There is not, as of 30 August 2026, a large multi-site RCT showing Copilot raises *conceptual* CS learning. The better evidence is interaction studies, and they are the right evidence for a chapter on struggle, because they show where the generative work went.

Prather, Reeves, Leinonen, MacNeil, and colleagues, 2023, “It’s Weird That it Knows What I Want”: novice programmers with GitHub Copilot. Two patterns: **shepherding** (steering auto-complete rather than writing) and **straying/drift** (accepting incorrect suggestions into a debug rabbit hole). Students believed Copilot would generate correct code. Metacognitive difficulties were common.²⁶² Prather et al. 2024, laboratory sessions on greenfield assignments: 20 of 21 completed the task, but benefits accrued to students with **stronger metacognitive skills**; students with metacognitive difficulties were **negatively** affected.²⁶³ Later Copilot process studies (for example, brownfield coding, ICER 2025) report less manual typing and less web search when Copilot is on—productivity, not learning. Trust-calibration work with upper-division students (extended Copilot

²⁶²James Prather, Brent N. Reeves, Juho Leinonen, Stephen MacNeil, et al., “‘It’s Weird That It Knows What I Want’: Usability and Interactions with Copilot for Novice Programmers,” *ACM Transactions on Computer-Human Interaction* (2023). Shepherding and straying/drift.

²⁶³Prather et al. (2024), laboratory sessions on greenfield assignments: 20 of 21 completed the task; benefits accrued to students with stronger metacognitive skills; students with metacognitive difficulties negatively affected.

use, arXiv:2509.13253) finds trust rises with experience *and* that students conclude a competent programmer is still required to finish tasks.²⁶⁴

The implication for the book is cognitive, not cultural. Copilot is a worked-example generator running at keystroke speed. For novices it is expertise reversal in the wrong direction unless the course forces tracing, explaining, and writing without the model. Kalyuga et al. 2003 already said the high-guidance materials that help novices impair more knowledgeable learners. Copilot inverts the slider: it gives novices the high-guidance they would need as *faded* examples, at a speed that removes the tracing, and it keeps giving that guidance after the student could generate. Shepherding is the worked-example effect stuck on. Straying is unproductive struggle in a debug hole the student did not make and cannot see. The metacognitively strong can use the tool as a critic. The metacognitively weak are harmed. That is not a reason to ban Copilot from the profession. It is a reason not to put it in the formation lane of an intro course.

UC San Diego CSE, in a Microsoft customer story: faculty (Porter, Shah) redesigned intro courses so students first demonstrate **manual competency**, then use Copilot; research on trust calibration; plans to prepare students for “emerging agentic AI.” **Vendor customer story, not a peer-reviewed trial of Cloud Agents.**²⁶⁵ Stanford CS221, Spring 2026: AI as a human collaborator; no asking for answers; no “check my work”; **no agentic repo-builders** like Codex or Claude Code for the project; transcripts required.²⁶⁶ That refusal is a graduate behavior and a struggle policy. It is also a preview of Chapter 10. Do not claim these campuses have agent-learning RCTs. They do not.

CS is where agents are culturally closest and pedagogically most dangerous for novices. Hermes Agent, Grok Bot, Cursor Cloud Agents, Claude Cowork, and OpenAI Workspace Agents can plan, use tools, and keep working after the student leaves. The 2022-26 RCT corpus is almost entirely turn-based tutors and Copilot-like autocomplete. Agent trajectories that plan, browse, and submit work on the student’s behalf are a different offloading regime. **No identified learning RCT on Hermes-class autonomous multi-step agents** as of this access date. Do not extrapolate Prather’s shepherding onto Grok Bot and call it evidence. Call it a warning about where the generative work goes when the tool is faster than tracing.

The expertise-reversal slider is the cleanest way to read the CS record without pretending there is a Copilot-raises-concepts RCT. Novices need worked examples. Copilot supplies them at keystroke speed, which is faster than tracing, which is the work that makes the example a schema rather than a paste. Intermediates need reduced scaffolding and variability. Copilot continues to offer the next token. Experts can use it as a critic and still have to finish the task (the trust-calibration finding with upper-division students). A course that skips the novice tracing and jumps to “we use Copilot because industry does” has set the slider to the expert position for people who do not yet have the schema. Prather 2024’s negative effect for students with metacognitive difficulties is that error, measured. UCSD’s manual-competency gate is the correction, labelled as a vendor customer story.

²⁶⁴Later Copilot process studies (e.g., brownfield, ICER 2025): less manual typing and less web search—productivity, not learning. Trust-calibration with upper-division students, arXiv:2509.13253: trust rises with experience; students conclude a competent programmer is still required.

²⁶⁵UC San Diego CSE (Porter, Shah), Microsoft customer story: manual competency first, then Copilot. **Vendor customer story, not a peer-reviewed trial of Cloud Agents.**

²⁶⁶Stanford CS221, Spring 2026 course policy: collaborator, not answer key; no agentic repo-builders (Codex / Claude Code) for the project; transcripts required.

Stanford’s agent ban on the project is the correction at the next product class. Both belong in a struggle chapter. Neither licenses a claim that Cloud Agents raise conceptual CS learning.

A *hypothetical*, labelled as such: an intro-CS student accepts thirty Copilot suggestions in a row, the program compiles, the student cannot explain a loop. *Shepherding*. Another student accepts a wrong suggestion, spends the hour in a debug hole, and concludes they are bad at programming. *Straying*. The course that forces tracing, “explain this generated function,” and “write this without the model” is the expertise-reversal slider set the right way. Stanford’s ban on agentic repo-builders for the project is that course at the agent layer. UCSD’s manual-competency gate is that course at the autocomplete layer. Both are struggle policies. Neither is a Cloud Agents trial.

Two-lane assessment, and the AIAS as vocabulary not evidence

If the only graded artefact can be produced by a model, the course no longer measures the skill. Detection is a losing strategy—false positives, an arms race, and it does not restore learning. The literature has moved to **task redesign**. This section gives the table the rest of the book will reuse. It is **theoretically motivated, not a fitted schedule**. It is **this book’s synthesis, not existing district policy**. Label it on the page as such.

Lane	AI	What is graded	When
Formation	Forbidden or hint-only	Unaided retrieval, derivation, draft, oral	Until a schema exists (novices; new topics; childhood)
Performance in an AI world	Allowed, logged	Judgement: selection, critique, verification, process	Once the student can detect model error

Contractor and Reyes’s augmentation/automation split is the empirical warrant for grading *how* AI was used, not only the essay. Bastani is the warrant for never letting the formation lane collapse into ChatGPT. GPT Base looked 48 percent better in practice and 17 percent worse on the closed-book exam. Students did not perceive the loss. A course that grades only the practice artefact has already chosen performance over learning. A course that grades the unaided exam, and that puts a hint-ladder in the practice window, has chosen the formation lane. A course that later asks students to select, critique, and verify model output, with transcripts, has opened the second lane. Both lanes are the job. One lane is a crutch.

Perkins, Furze, Roe, and MacVaugh (2024), *Journal of University Teaching and Learning Practice*. The Artificial Intelligence Assessment Scale (AIAS) gives a shared vocabulary from “no AI” through defined levels of assistance, so that permitted use matches the intended learning outcome.²⁶⁷ A later revision (Perkins et al., 2025/26, *JUTLP*, doi:10.53761/rrm4y757) recasts the

²⁶⁷Mike Perkins, Leon Furze, Jasper Roe, and Jason MacVaugh, “The Artificial Intelligence Assessment Scale (AIAS): A Framework for Ethical Integration of Generative AI in Educational Assessment,” *Journal of University Teaching and Learning Practice* 21, no. 6 (2024). Shared vocabulary, **not an RCT of learning**. Revision: Perkins et al., 2025/26, *JUTLP*, doi:10.53761/rrm4y757, recast as assessment-redesign framework; warns against using it merely

scale as an **assessment-redesign framework**—social-constructivist rationale, non-hierarchical visual, added “AI exploration” level—and warns against using it merely as a prohibition ladder. The authors’ own claim: the AIAS is a communication and validity tool, not a detector. This is a **design pattern, not an RCT of learning**. It is still the most widely implemented shared language in higher education as of 2026. Do not treat AIAS as efficacy research. Do not treat a five-level poster on a classroom wall as evidence that struggle stayed in the student.

A 2026 cross-national document analysis of 21 English-medium universities (doi:10.17275/per.26.60.13.4) found that **all** sampled institutions now allow some student GenAI use under conditions; almost all require disclosure.²⁶⁸ Four recurrent assessment patterns:

1. Process-evidence portfolios with AI-use logs.
2. AI-assisted drafting plus in-class or oral verification.
3. Critical engagement with AI outputs (critique, adapt, verify).
4. Secure high-stakes exams paired with AI-enabled coursework.

Those four patterns are how universities have, in documents, tried to keep both lanes. They are not a finding that both lanes are being kept. MDPI *Education Sciences* (2025) “Beyond Detection” makes the same pedagogical point: authenticity is redesigned, not policed; oral defences, annotated portfolios, and recorded walkthroughs make thinking visible.²⁶⁹ QAA, *Reconsidering assessment for the Chat GPT era* (July 2023), told UK tertiary providers to cut AI-vulnerable tasks and rebuild authentic, synoptic assessment.²⁷⁰ JCQ *AI Use in Assessments*, effective 30 April 2025, is the clearest operational honor-code-plus-process regime in a national qualifications system: own-work rule; acknowledgement = tool name + URL + date; retain a non-editable copy of prompt and output; AI must not be the **sole marker**; prevention is pedagogical (in-class production, intermediate drafts, viva).²⁷¹ ICAI/HEA 2025 “oral assessment safeguard” language recommends oral and live review as a way to let explanation take precedence when authorship is in doubt. Oral is **not** a complete integrity solution (workload, access, future impersonation risk). It is aligned with Ennis and Facione: the student must explain. In-class, timed, AI-forbidden tasks remain appropriate when the outcome is unaided fluency (mental math, close reading, timed writing, code-from-scratch). The mistake is using that format for every outcome, which both fails to teach AI skill and invites underground use.

UK DfE tells schools to review homework and unsupervised study policy because unsupervised take-home work is now an open-web assessment. OECD Digital Education Outlook 2026, on the listing page: unguided GenAI raises **task performance without learning gains—UNVERIFIED at full-PDF level** in the research pass behind this book.²⁷² AFT Core Value 3: full / partial / no-

as a prohibition ladder.

²⁶⁸Cross-national document analysis of 21 English-medium universities (2026), doi:10.17275/per.26.60.13.4: all sampled institutions allow some student GenAI use under conditions; almost all require disclosure. Four recurrent patterns as listed in the text.

²⁶⁹MDPI *Education Sciences* (2025), “Beyond Detection”: authenticity redesigned, not policed.

²⁷⁰Quality Assurance Agency for Higher Education, *Reconsidering Assessment for the Chat GPT Era* (July 2023).

²⁷¹Joint Council for Qualifications, *AI Use in Assessments*, effective 30 April 2025: own-work rule; acknowledgement = tool name + URL + date; non-editable prompt/output copy; AI not the sole marker.

²⁷²OECD *Digital Education Outlook 2026*, listing-page claim that unguided GenAI raises task performance without learning gains: **UNVERIFIED at full-PDF level**. ICAI/HEA 2025 oral-assessment-safeguard language: oral is not a complete integrity solution. AFT *Commonsense Guardrails* v2 (March 2025), Core Value 3: full / partial / no-AI labels. Singapore MOE: secondary/pre-university students must evaluate output, check accuracy, acknowledge use.

AI assignment labels, decided by the teacher, with a teachable-moment first response to misuse. Singapore MOE: secondary and pre-university students must evaluate AI output, check accuracy, acknowledge use. None of these is two-lane assessment as a district SOP. They are the institutional weather in which a two-lane rule can be stated without pretending it has already been fitted.

Honor codes without artefacts are not the formation lane. A one-line syllabus ban is not an honor code in JCQ's sense. JCQ requires the tool name, the URL, the date, a non-editable prompt/output copy, and a brief explanation of use. That is process evidence. QAA's advice to cut the volume of tasks GenAI can complete unauthorised is task redesign. The 21-university analysis's four patterns are redesign plus disclosure. Two-lane assessment is the cognitive-science reading of those documents: unaided work until a schema exists; logged, judged use once the student can catch the model. A university that allows some GenAI under conditions has not, by that fact, implemented two lanes. It has stopped pretending a total ban is the equilibrium. The book still has to say what is graded, and when. That is the table. It is not Perkins 2024. Perkins is the shared language in which the table can be posted without inventing a fifth scale.

Across ages, the same cognitive principles scale; the expertise-reversal slider moves. **Childhood / early elementary.** High intrinsic load relative to schemas. Prioritise human relationship, decoding, number sense, and oral language. AI, if present, is a *worked-example and retrieval* engine with no answer key in the student-facing channel, used in short supervised bursts. Unsupervised chat is not a reading or arithmetic programme. Chapter 4 already put this band at teacher-facing only. The struggle policy agrees. **Upper elementary / middle school.** Introduce productive struggle with hint-only tutors *inside* scheduled practice (the Khanmigo *access* model failed to produce tutoring-sized effects; the Nigeria *session* model worked). Teach a first version of hallucination-checking on facts the class has already studied. Keep a large AI-forbidden core (fluency, handwriting-to-compose, mental math). **High school.** Bastani's result is the governing constraint for homework. Provide GPT-Tutor-class tools for practice; examine unaided. Build evaluation of AI outputs in history and science. CS courses: Copilot only after students can trace; assessments that include "explain this generated function" and "write this without the model." **College.** Two-lane assessment as default (AIAS as vocabulary plus oral/process). Explicit instruction in augmentation versus automation (Contractor and Reyes). Capstone work *with* agents, first-year methods *without*, until students can catch the model. Faculty as embedded designers of prompts and of which loop the agent is in.

A *hypothetical* syllabus fragment, labelled as such, to show the architecture rather than a scene. Week 1-3 of a new topic: formation lane. Hint-only tutor if a tutor is used at all; attempt gate; no full solutions; unaided quiz at the end of the week. Week 4: a mixed task in which the student must generate first, then may use a logged tool to critique, then submits a process log and sits a five-minute oral. Week 8, after a schema exists: performance in an AI world—select among three model drafts, mark errors, verify two claims against a primary source, disclose the tool, keep the transcript. The grade on week 8 is judgement, not fluency of the prose. The grade on week 1 is unaided retrieval. The same student, the same model class, two lanes. The expertise-reversal slider moved. Nothing in that fragment is a district policy. It is the table above, written as a sequence.

Trouche's instrumental orchestration (2004), which Chapter 7 will spend as a labor idea, is already present as a struggle idea: the teacher regulates when AI is activated, what is delegated,

who validates, and how visible the tool is.²⁷³ Mollick and Mollick (2023), “Assigning AI: Seven Approaches”—tutor, coach, mentor, teammate, student, simulator, tool—is a practical prompt library the Nigeria RCT actually used.²⁷⁴ The teacher still chooses the approach per task. NUS AICET’s “embedded AI designer” framing (2025): faculty encode pedagogical intent into constraints, as Bastani’s teachers encoded solutions and mistakes, as Kestin encoded step-by-step physics.²⁷⁵ Passive use of a general chatbot is not that job. LearnLM’s qualitative finding that supervising tutors learned Socratic moves from the model is interview evidence from $n = 5$, not a PD trial—do not promote it.

The two-lane rule will be accused of being a schedule pretending to be a principle. It is the opposite. It is a principle that refuses to fit a schedule until someone runs the study. Open question 11 in the research behind this book is exactly that: optimal mix of AI-forbidden and AI-required tasks by grade. This chapter names the mix as a hypothesis. It does not print a K-12 calendar and call it evidence. A first-grade unaided oral is a formation-lane assessment. A twelfth-grade take-home that assumes a premium personal agent the school does not issue is not a performance-lane assessment; it is homework-gap 2.0, which is Chapter 8’s equity problem. The lane is not “AI yes in high school, AI no in elementary.” The lane is whether a schema exists, whether the student can detect model error, and whether the grade is measuring the skill that was supposed to be formed.

The live debates this chapter will not close. Optimal mix of AI-forbidden and AI-required tasks by grade: a theoretically motivated hypothesis, not a fitted schedule. Whether proactive, multimodal, memory-bearing tutors fix Khanmigo-style non-use: authors on both sides treat this as open; a default-open chat window did not, in Oreopoulos and Low’s Year-2 data, raise messaging. Equity: Nigeria found larger effects for higher-baseline and higher-SES students (and for girls). Copilot and CS studies suggest metacognitively weaker students may be harmed. A “personal tutor for every child” that widens gaps is an empirical possibility, not a slogan. Year-scale unaided skill on independent state tests: not yet known (Khanmigo TCAP outstanding; Nigeria and Kestin no year-later follow-up; Bastani measured the next hour). No K-2 conversational-agent RCTs on language, attention, or social cognition. No reading or history causal estimates comparable to math. No WWC-grade evidence on named generative products. Replication of Kestin outside a highly staffed university physics course, and of Bastani outside one Turkish high school’s four sessions, is unshown. Transfer of “evaluate the chatbot” skill to non-AI critical thinking is Chapter 6’s hole, not this chapter’s to fill.

What would actually make students “dumber” is not contact with models. Three conditions together: (1) the offloaded process **is** the learning objective (writing, proof, source evaluation); (2) the environment **never measures** the skill unaided (take-home essays as the only grade); (3) the student **cannot detect** model error (Bastani: students copy arithmetic errors they could have caught). Designs that keep generation, retrieval, and checking inside the student (GPT Tutor, Nigeria prompts, PS2 Pal, LearnLM Socratic drafts, augmentation use in Contractor and Reyes) do not show this atrophy pattern on the outcomes they measured. That is the opposite of a general “AI = dumber”

²⁷³Luc Trouche, instrumental orchestration (2004); AISIO (AI-Supported Instrumental Orchestration, 2025 lecture). Concept preview; labor treatment in Chapter 7. Monday list as synthesis of the positive studies, labelled as such.

²⁷⁴Mollick and Mollick, arXiv:2306.10052 (2023), as used in De Simone et al. 2025.

²⁷⁵NUS AICET “embedded AI designer” framing (2025). LearnLM qualitative finding that supervising tutors learned Socratic moves from the model: interview evidence from $n = 5$, not a PD trial.

law. It is a struggle policy.

Twelve findings from the 2022-26 record, compressed so this chapter can stand on them without asking Chapter 3 to be re-read as a table. Immediate AI-aided performance is the wrong success metric; unaided delayed tests are the right one (Soderstrom and Bjork, 2015; Bastani et al., 2025). Unguarded answer-giving GPT can harm subsequent math learning (about -17 percent in Bastani); hint-and-grounded-solution GPT avoids the harm but did not beat no-AI on the exam. Pedagogically engineered tutors can raise immediate conceptual learning, even versus strong active learning (Kestin et al., 2025, 0.73-1.3 SD quantile range on an immediate physics post-test)—a design result, not a ChatGPT-default result. At public-school scale, engagement, not IQ of the model, bound Khanmigo's effect; ITT about 0.06-0.08 SD/year, similar to practice without the chatbot (Oreopoulos and Low, 2026). In a low-resource system, teacher-guided Copilot sessions with science-of-learning prompts produced 0.238 SD on English and 0.206 SD on the regular term exam (De Simone et al., 2025). Writing with ChatGPT often improves the artefact and not the knowledge (Fan et al., 2025); using AI to *explain* rather than to *draft* is associated with delayed unaided gains (Contractor and Reyes, 2026, working paper). Cognitive offloading is a real, old mechanism (Risko and Gilbert, 2016). The MIT EEG study is a small preprint consistent with shallower encoding during LLM-assisted writing, not proof of population-level cognitive decline. Novice programmers with Copilot show shepherding and straying; benefits concentrate among students who already have metacognitive skill (Prather et al., 2023, 2024). Human-in-the-loop that *helps* is specific: grounded prompts, attempt gates, circulating teachers, tutor-facing copilots—not a warm body in the building. Assessment must split formation from AI-world performance. Teachers are orchestrators of agency, struggle, and validity; that sentence is a preview of Chapter 7, not a finding that PD has been shown to raise pupil outcomes for generative AI (reviews stop at teacher self-efficacy).

Those twelve are not a sales deck. Several of them are negative. Several are small. Several are preprints or working papers, named as such. There is no Hermes-class learning RCT in the list. There is no WWC generative-tutor report. There is no fitted grade-by-grade mix. The list is what an honest struggle chapter can hold. It is enough to refuse rescue, refuse abandonment, and refuse to grade the artefact as if it were the skill.

Keep the struggle in the student. Rescue on the first pause is how NCTM said teachers convert demand into procedure. Abandonment until drowning is unproductive struggle. An agent that finishes the job is the first failure at machine speed. An agent that will not hint is the second. Managed struggle is the space between: attempt gate, hint ladder, grounded misconception library, unaided check. Two lanes, labelled as synthesis. AIAS as language, not as an RCT. Copilot as worked examples at keystroke speed, shepherding and straying until the course forces tracing. Worked examples first, then expertise reversal. Mathematics has the trials. Writing has the artefact problem. Science splits tutoring from inquiry. Reading and history have mechanisms and a hole. Judgement is not a bicycle-skill. Willingham (2007) is Chapter 6's load-bearing constraint, previewed here only so the two-lane rule is not mistaken for a critical-thinking curriculum: the performance lane grades selection, critique, and verification of *something the student knows enough to judge*. An agent that "teaches critical thinking" in the abstract, while starving students of history, math, and science schemas, will fail Facione's skills in the wild because there is nothing to analyse. Formation lane first. Domain knowledge first. Then the second lane. Chapter 6 puts domain knowledge back under critical thinking. The struggle policy can be stated without it. It cannot be executed

without it. A two-lane gradebook that has nothing in the formation lane but ChatGPT critique is not two lanes. It is one lane, renamed. Keep the struggle in the student. Grade the skill that was supposed to be formed.

Chapter 6

Judgement Needs Something to Judge



Figure 7: Marked pages, a magnifying glass, old books, and a closed laptop.

In the summer 2007 issue of *American Educator*, Daniel Willingham made an argument that education policy has spent two decades trying not to hear. “People who have sought to teach critical thinking have assumed that it is a skill, like riding a bicycle, and that, like other skills, once you learn it, you can apply it in any situation.” Research from cognitive science, he wrote, shows that thinking is not that sort of skill. The processes of thinking are intertwined with the content of thought — that is, domain knowledge. You can teach students maxims about how they ought to think — look at an issue from multiple perspectives — but without background knowledge and practice they will

not be able to implement the advice they memorize. Scientific thinking, historical thinking, and literary thinking are partly distinct. Metacognitive strategies help only at the margin.²⁷⁶

That is the load-bearing constraint of this chapter. An agent that “teaches critical thinking” in the abstract, while starving students of history, mathematics, and science schemas, will fail the skills that assessment instruments claim to measure, because there is nothing to analyse.

Bastani and colleagues’ 2025 field experiment in a large Turkish high school is the arithmetic version of the same claim. In a probe, the unguarded GPT-4 tutor returned a correct answer only about 51 percent of the time. Students still copied. They copied arithmetic errors they could have caught.²⁷⁷ Judgement without knowledge is theatre. The model’s fluency supplied the appearance of a completed judgement; the student’s knowledge was not in the loop to detect that the judgement was wrong.

This chapter is not a literacy-framework tour. It is an argument about what must remain inside the student if “critical thinking” is to mean anything once the chat window closes. Three non-equivalent accounts of the term are in circulation. Only one of them can constrain an AI design. Default chatbot use then collapses the others into a single accept-or-reject of a fluent paragraph. The same skills can be trained, if the task requires the student to produce the judgement. Institutions have built programs — Cornell’s incoming-student literacy modules, ISTE’s six roles, AI4K12’s grade-band charts, UNESCO’s competency frameworks, Singapore’s character-and-citizenship strand. Those programs are programs, not outcome studies. There is not yet a randomised trial showing that a K–12 “evaluate the chatbot” curriculum transfers critical thinking to non-AI tasks. That is a hole. It is not a licence to skip the practice, and it is not a licence to treat a checklist as a formed skill.

Three accounts; Willingham is the constraint

“Critical thinking” is not a brand. It is at least three different things, and the book that collapses them will design the wrong tasks.

Robert Ennis’s account is the oldest still in circulation as a curriculum checklist. Critical thinking, he writes, is “reasonable reflective thinking focused on deciding what to believe or do,” plus a specified set of dispositions and abilities. The dispositions include seeking reasons, being open-minded, and taking a position when the evidence warrants. The abilities include clarifying, analysing arguments, evaluating sources, inferring, and judging value statements.²⁷⁸ That list is usable. A

²⁷⁶Daniel T. Willingham, “Critical Thinking: Why Is It So Hard to Teach?,” *American Educator* 31, no. 2 (Summer 2007): 8–19. Bicycle sentence quoted from the AFT-hosted PDF (verified 30 August 2026): “People who have sought to teach critical thinking have assumed that it is a skill, like riding a bicycle, and that, like other skills, once you learn it, you can apply it in any situation.” Adjacent sentences in the same paragraph: processes of thinking “are intertwined with the content of thought (that is, domain knowledge)”; maxims can be taught, “but without background knowledge and practice, they probably will not be able to implement the advice they memorize.” https://www.aft.org/sites/default/files/media/documents/2023/31.2_summer_2007.pdf.

²⁷⁷Hamsa Bastani, Osbert Bastani, Alp Sungu, Haosen Ge, Özge Kabakçı, and Rei Mariman, “Generative AI without Guardrails Can Harm Learning: Evidence from High School Mathematics,” *Proceedings of the National Academy of Sciences* 122, no. 26 (2025): e2422633122. GPT Base returned a correct answer only about 51 percent of the time in a probe; students still copied. Unaided exam: GPT Base about 17 percent worse than never-AI.

²⁷⁸Robert H. Ennis, “The Nature of Critical Thinking: Outlines of General Critical Thinking Dispositions and Abilities” (2011/2015), <https://criticalthinking.net/wp-content/uploads/2024/04/The-Nature-of-Critical-Thinking.pdf>. Def-

department can hang units on it. It does not tell you what the student must already know in order to execute any item on it.

Peter Facione’s 1990 American Philosophical Association Delphi Report, still the basis of the most widely used assessment instruments, is a skills-plus-habits account. Critical thinking is purposeful, self-regulatory judgement resulting in interpretation, analysis, evaluation, inference, explanation, and self-regulation, plus a dispositional side: truth-seeking, open-mindedness, inquisitiveness, and the rest. The California Critical Thinking Skills Test and the California Critical Thinking Disposition Inventory implement this inventory.²⁷⁹ Facione’s 2015 primer restates it for practitioners. The six skills are a genuine design surface. They are not a bicycle. They have to be executed on something.

Willingham is the constraint that makes the other two honest. The 2007 *American Educator* essay — written for teachers, not for a methods journal — is explicit that the processes of thinking cannot be detached from the knowledge those processes operate on. Maxims can be taught. They will not be executed without content and practice. Domain matters: the historian’s sourcing is not the scientist’s control of variables is not the literary critic’s attention to diction, even if all three can be described, at a high enough altitude, as “evaluation.” Metacognitive strategies help only at the margin.²⁸⁰

For a book on autonomous AI from childhood through college, Willingham is load-bearing for a reason the other two accounts cannot supply. An agent is extremely good at emitting the *surface form* of Facione’s skills. It will interpret, analyse, evaluate, infer, explain, and even simulate self-regulation in a fluent paragraph, on demand, in any subject the student has not studied. Ennis’s dispositions can be prompted in a system message. None of that is the student thinking. If the student lacks the schema against which a claim can be checked, Facione’s evaluation is a coin flip dressed as a rubric, and Ennis’s “take a position when the evidence warrants” is a style of sentence.

Literacy frameworks do not escape this. UNESCO’s student and teacher competency frameworks, ISTE’s Profile of an AI-Ready Graduate, AI4K12’s Five Big Ideas, Cornell’s four Canvas modules, and the U.S. Department of Education’s literacy definition in the 2026 supplemental priority are usable only if they sit on domain knowledge.²⁸¹ Literacy without knowledge is a checklist students can recap and not execute. That is not an argument against the frameworks. It is the condition under which they are not a substitute for history, mathematics, science, and the rest of the curriculum the agent is being invited to write.

The popular story this chapter has to kill is the one in which critical thinking is a transferable bicycle-skill an agent can teach in the abstract — a course, a badge, a “AI literacy” unit that, once completed, licenses the student to use any tool on any task. Willingham’s constraint does not forbid AI literacy. It forbids treating AI literacy as a replacement for the knowledge that judgement requires. The rest

inition quoted: “reasonable reflective thinking focused on deciding what to believe or do.”

²⁷⁹Peter A. Facione, *Critical Thinking: A Statement of Expert Consensus for Purposes of Educational Assessment and Instruction* (The Delphi Report) (American Philosophical Association, 1990); Peter A. Facione, *Critical Thinking: What It Is and Why It Counts* (Insight Assessment, 2015).

²⁸⁰Willingham, “Critical Thinking.” Scientific, historical, and literary thinking treated as partly distinct; metacognitive strategies help only at the margin.

²⁸¹Frameworks specified and dated in section 6.4 notes below; U.S. Department of Education literacy definition at 91 Fed. Reg. 18774 (April 13, 2026), effective May 13, 2026.

of the chapter is what that distinction does to design.

The bicycle error is specific. Riding a bicycle is a skill you can take from the driveway to a new street. Once the motor pattern is in place, the new street does not require a new theory of bicycles. Critical thinking is not that kind of object. A student who has learned to consider alternative explanations in a unit on vaccines does not thereby consider alternative explanations in a unit on Reconstruction, unless the Reconstruction schema is there to generate the alternatives. Willingham's point is not that transfer never occurs. It is that transfer is not the default, that it is domain-bound, and that teaching the maxim in the absence of the domain is an education-policy habit with a weak record. An agent that drills the maxim – always ask for sources; always generate a counterargument – while writing the history essay for the student is practising the habit of outsourcing the domain and reciting the maxim. That is the opposite of the skill.

Ennis remains useful as a checklist inside a domain. Evaluate sources on a set of documents the class has read is a history assignment. Evaluate sources as a free-floating AI-literacy outcome, assessed by a chatbot conversation about a topic the student has not studied, is a performance of evaluation. Facione's instruments (CCTST, CCTDI) at least try to measure the skills on supplied passages. They are still not a substitute for knowing enough to see that a fluent passage is wrong. The book's position is therefore not anti-literacy and not anti-Ennis. It is that literacy frameworks, Ennis checklists, and Facione skill names are usable only as the evaluative layer on top of subject teaching. Chapter 5's struggle policy is what builds the something that judgement needs. This chapter is what happens when that something is missing, and what to require of the student when a model is in the room anyway.

Where default use outsources judgement

Default ChatGPT use collapses Facione's evaluation and self-regulation into a single accept-or-reject of a fluent paragraph. Students then practice neither source evaluation nor argument construction. They practice a binary: keep, or ask again. That is not a small loss. It is the loss of the loop the school exists to form.

The empirical trace is already on the table from Chapter 3. It is spent here as epistemology, not as a second tour of effect sizes.

Nataliya Kosmyna and colleagues' MIT Media Lab preprint, "Your Brain on ChatGPT," is a small, preliminary study — fifty-four Boston-area adults, mostly undergraduates; session 4 on eighteen people; twenty-minute SAT-style essays under EEG — and it is not a journal article as of the access date of this research. It should not be cited as proof that ChatGPT atrophies the brain. What it does show, inside those limits, is an encoding failure consistent with outsourced judgement: connectivity lowest in the LLM group; low perceived ownership; in session 1, fifteen of eighteen LLM users could not quote their own essay, and none quoted it correctly; named entities and n-grams more homogeneous, some essays close to default ChatGPT SAT answers. Teachers and an AI judge still often scored the essays acceptably.²⁸² Performance was not encoding. The student could not

²⁸²Nataliya Kosmyna et al., "Your Brain on ChatGPT: Accumulation of Cognitive Debt When Using an AI Assistant for Essay Writing Task," arXiv:2506.08872 (June 10, 2025), MIT Media Lab. Preprint, not a journal article as of August 30, 2026. n = 54 for sessions 1–3; n = 18 for session 4. Session 1 quoting: 15 of 18 LLM users could not quote;

retrieve what the model had written in the student's name. That is Facione's explanation and self-regulation failing in public.

Fan, Tang, and colleagues' 2025 randomised laboratory study in the *British Journal of Educational Technology* is the cleaner behavioural result. One hundred seventeen university students wrote with one of four supports: ChatGPT, a human expert chat, writing-analytics tools, or no extra tool. The ChatGPT group improved essay scores more. There were no significant differences in knowledge gain or transfer, and no difference in post-task intrinsic motivation. Process mining found the mechanism: ChatGPT users formed tight revise-and-ask-AI loops with fewer metacognitive processes — evaluation, orientation — than the human-expert group. The authors named this “metacognitive laziness.”²⁸³ The artefact got better. The evaluative loop got thinner. That is the writing version of Bastani's practice-up, exam-down dissociation, measured inside a single session rather than on a later unaided test.

Matthias Stadler, Maria Bannert, and Michael Sailer's 2024 *Computers in Human Behavior* experiment is the scientific-inquiry analogue. Ninety-one university students used ChatGPT-3.5 or Google to research nanoparticles in sunscreen and write recommendations. LLM users reported lower cognitive load and produced lower-quality reasoning and argumentation. Homogeneity of conclusions did not differ.²⁸⁴ Ease without depth. Facione's analysis and inference, offloaded; the student kept the feeling of having thought.

Bastani's 51 percent probe is the mathematics version of the same problem. Hallucination makes the outsourcing worse than Wikipedia. A clumsy encyclopaedia paragraph still looks like a paragraph someone wrote, with seams. A fluent model paragraph looks like a completed judgement. Students in the unguarded arm copied anyway, including errors they had the school mathematics to catch, and did not perceive that they had learned less.²⁸⁵ The surface form did the epistemological work that the student's knowledge was supposed to do.

James Prather, Brent Reeves, Juho Leinonen, Stephen MacNeil, and colleagues' 2023 Copilot study, and the 2024 laboratory follow-on, are the computer-science version. Novice programmers with GitHub Copilot showed two patterns: *shepherding* (steering autocomplete rather than writing) and *straying* or *drifting* (accepting incorrect suggestions into a debug rabbit hole). Students believed Copilot would generate correct code. Metacognitive difficulties were common. In the 2024 greenfield sessions, twenty of twenty-one students completed the task, but benefits accrued to students with stronger metacognitive skills; students with metacognitive difficulties were negatively affected.²⁸⁶ Copilot is a worked-example generator running at keystroke speed. For a novice it is

0 of 18 quoted correctly. Do not cite as population-level cognitive decline or as an agent study.

²⁸³Y. Fan et al., “Beware of Metacognitive Laziness: Effects of Generative Artificial Intelligence on Learning Motivation, Processes, and Performance,” *British Journal of Educational Technology* 56, no. 2 (2025): 489–530. Randomised lab study, N = 117. Essay scores improved more in the ChatGPT arm; knowledge gain and transfer not significantly different; fewer evaluation and orientation processes than the human-expert group.

²⁸⁴Matthias Stadler, Maria Bannert, and Michael Sailer, “Cognitive Ease at a Cost: LLMs Reduce Mental Effort but Compromise Depth in Student Scientific Inquiry,” *Computers in Human Behavior* 160 (2024): 108386. RCT, N = 91, ChatGPT-3.5 versus Google.

²⁸⁵Bastani et al., “Generative AI without Guardrails.” Students in GPT Base did not perceive that they had learned less; cf. Louis Deslauriers et al., “Measuring Actual Learning versus Feeling of Learning in Response to Being Actively Engaged in the Classroom,” *Proceedings of the National Academy of Sciences* 116, no. 39 (2019): 19251–57.

²⁸⁶James Prather et al., “‘It’s Weird That It Knows What I Want’: Usability and Interactions with Copilot for Novice

expertise reversal in the wrong direction unless the course forces tracing, explaining, and writing without the model. The evaluative loop is exactly what the weaker novices do not have, and exactly what the tool lets them skip.

Michael Gerlich's 2025 *Societies* paper will be offered, in some rooms, as the social-science confirmation of all of the above. It is a survey-based association of AI-tool use with cognitive offloading and lower self-reported critical thinking. Younger users were more dependent. It is correlational. It is not a randomised trial. It does not identify a treatment.²⁸⁷ This book will not promote it. The causal record on outsourcing is Bastani, Fan, Stadler, and the Copilot interaction studies. Gerlich is fog. Fog is not evidence that the other studies are wrong. It is not a reason to skip the practice that those studies imply.

Three conditions together, not contact with models, are what would actually make students worse at the work school is for. First, the offloaded process is the learning objective – writing, proof, source evaluation, the argument itself. Second, the environment never measures the skill unaided: the take-home essay is the only grade. Third, the student cannot detect model error, as in Bastani's copied arithmetic. Designs that keep generation, retrieval, and checking inside the student – GPT Tutor, the Nigeria prompts, Kestin's PS2 Pal, LearnLM's Socratic drafts, augmentation use in Contractor and Reyes – do not show this atrophy pattern on the outcomes they measured.²⁸⁸ That is the opposite of a general AI-makes-you-dumber law, and it is also the opposite of a general critical-thinking-app law. The evaluative loop is the variable. Default chat drops it. Guarded, domain-loaded tasks can keep it. Kosmyna's EEG is consistent with shallower encoding during LLM-assisted writing; it is a preprint with a small convenience sample and no delayed knowledge test of the Bastani or Contractor kind, and it should not be asked to carry a developmental policy.

Wikipedia is the wrong comparison for the hallucination problem, and it is the comparison most faculty reach for. Wikipedia, at its worst, is a page with a history tab, talk page, and citations that can be checked. A fluent model paragraph has no history tab. It has the grammar of a completed judgement. Students who have been trained to distrust Wikipedia's unsourced sentences have not been trained to distrust a paragraph that sounds like their teacher. Fan's ChatGPT arm did not fail because the essays were obviously machine-made. It succeeded on the artefact and failed on knowledge and transfer. Stadler's LLM users did not report confusion; they reported lower load and wrote worse arguments. The outsourcing is comfortable. Comfort is the pedagogical problem, not the tell.

A 2026 conceptual distinction in *Frontiers in Psychology*, between dependent and autonomous offloading, is useful even before that paper's psychometrics settle: frequency of AI use is not the

Programmers," *ACM Transactions on Computer-Human Interaction* (2023). The 2024 laboratory follow-on (greenfield assignments; 20 of 21 completed the task; benefits concentrated among students with stronger metacognitive skills; metacognitive difficulties associated with negative effects) is described in the book's pedagogy memo; a separate full bibliographic record was not fetched this pass. Do not treat either study as a multi-site RCT of conceptual CS learning, and do not treat them as Cloud Agents trials.

²⁸⁷Michael Gerlich, "AI Tools in Society: Impacts on Cognitive Offloading and the Future of Critical Thinking," *Societies* 15, no. 1 (2025): 6. Survey/correlational; younger users more dependent; not an RCT.

²⁸⁸Pedagogy-memo synthesis of Bastani et al. (2025); Fan et al. (2025); Stadler, Bannert, and Sailer (2024); De Simone et al. (2025); G. Kestin et al., "AI Tutoring Outperforms In-Class Active Learning," *Scientific Reports* 15 (2025): 17458; LearnLM Team, arXiv:2512.23633; Contractor and Reyes, IZA DP 18792. The three-condition atrophy claim is the book's compression of that record, not a new experiment.

variable; whether the user retains the evaluative loop is.²⁸⁹ Default chat does not retain it. That is the design fact, not a moral one. Cognitive offloading is an old, non-moral mechanism — using the environment to reduce internal processing, as Risko and Gilbert defined it in 2016.²⁹⁰ Generative AI extends the offload from memory and calculation to reasoning, drafting, and evaluation. Those are the skills schools exist to form. The question is not whether students will offload. They will. The question is which part of Facione’s loop the task still requires the student to produce.

AI-native tasks that keep the evaluative loop inside the student

The same Facione skills map onto AI-native tasks if — and only if — the task requires the student to produce the judgement. That is a design rule, not a trial result. There is not yet a randomised trial showing that a K–12 “evaluate the chatbot” curriculum raises transfer of critical thinking to non-AI tasks.²⁹¹ The hole is real. It is not a licence to skip the practice. Willingham’s constraint predicts that the practice will fail if it is abstract, and may work if it is domain-loaded and required. The only way to find out is to build the tasks and measure unaided, delayed, off-tool performance. Until that trial exists, the honest move is to specify the tasks, grade the student’s judgement rather than the model’s paragraph, and leave the transfer question standing.

The mapping is a table, not a slogan.

Facione skill	AI-native training task
Interpretation	Restate the model’s claim in the student’s words before using it
Analysis	Diagram the argument; mark premises the model left implicit
Evaluation	Score a model essay against a rubric; find the weakest warrant
Inference	Given the model’s conclusion, generate a counterexample
Explanation	Oral defence: why did you keep this sentence and cut that one?
Self-regulation	Process log: where did you accept, challenge, or ignore the model?

The table can be instantiated in subjects, not only as generic literacy drills. The restatement test is cheap and revealing; it is not the only subject-shaped version of the loop.

History is the sharpest illustration, even though almost no subject-specific generative-AI randomised trial exists as of this research’s access date. Historical thinking is sourcing, corroboration,

²⁸⁹The 2026 *Frontiers in Psychology* distinction between dependent and autonomous offloading is used at the conceptual level stated in the pedagogy memo (frequency is not the variable; retention of the evaluative loop is). Full psychometrics are not treated as settled.

²⁹⁰Evan F. Risko and Sam J. Gilbert, “Cognitive Offloading,” *Trends in Cognitive Sciences* 20, no. 9 (2016): 676–88.

²⁹¹Open question, not a hedge: no RCT of transfer from a K–12 “evaluate the chatbot” curriculum to non-AI critical thinking was found as of August 30, 2026. This is the Willingham test. It remains standing.

and contextualisation. An LLM that writes a confident narrative is a primary-source problem: it looks like a secondary source and behaves like an unsourced claim. The sharpening design is sequential and oral. The student reads the sources first. The model produces a narrative. The student marks what is sourced, what is inferred, and what is fabricated, and then defends the marks aloud. The dulling design is write-a-DBQ-with-ChatGPT.²⁹² No trial is required to see that the second design grades the model. The first design grades the student's judgement of the model against documents the student has read. Willingham's constraint is satisfied only in the first.

Science splits the same way, and here there is causal evidence, already earned in Chapter 3, that design is the treatment. Kestin, Miller, Klales, Milbourne, and Ponti's 2025 crossover RCT in Harvard's Physical Sciences 2 found that a custom GPT-4 tutor with instructor-written step-by-step solutions and sequential scaffolding beat the course's existing in-class active learning on an immediate post-test (linear 0.63 SD; quantile, because of ceiling, 0.73 to 1.3 SD). Delayed retention, transfer, and full-course outcomes were not measured; the lessons were Bloom understanding, applying, and analysing, not creating; the population is a selective university.²⁹³ That is structured tutoring on a known learning objective, with the instructor's physics in the prompt. Stadler is open inquiry with a general chatbot versus search: lower load, weaker reasoning. Structured tutoring is not open judgement practice. Using Kestin as evidence that ChatGPT trains scientific critical thinking would be a category error. Using Stadler as evidence that no AI belongs in a science course would be the opposite error. The evaluative loop is in the task, not in the model class.

Computer science, as Chapter 5 already flagged, has no large multi-site RCT showing Copilot raises conceptual learning. The interaction evidence is the design brief: novices shepherd or stray; metacognitively weaker students can be harmed. The AI-native tasks that keep judgement inside the CS student are therefore unromantic. Explain this generated function. Trace it. Write the next one without the model. UC San Diego's intro-sequence redesign – manual competency first, then Copilot – is a vendor customer story, not a peer-reviewed trial of Cloud Agents, and Stanford CS221's Spring 2026 ban on agentic repo-builders on the project is a legitimate assessment choice, not Luddism.²⁹⁴ An agent that plans, browses, and submits the repository is a different offloading regime from autocomplete. There is no identified learning RCT on that regime. Asking the student to interpret, check, and orally defend a generated function is the Facione loop at the level CS can currently stand on.

Walk the rows. Interpretation first: if the student cannot restate the model's claim without looking at it, the student does not have a claim to evaluate. Restatement is retrieval. It is also the cheapest test of Kosmyna's quoting failure, administered without EEG. A student who cannot put the sentence in their own words has not judged it. They have accepted a texture.

Analysis is the historian's and the scientist's job under a new description. Large language mod-

²⁹²History design as specified in the pedagogy memo, section 5.5. Almost no subject-specific RCT as of August 30, 2026. Wineburg's sourcing, corroboration, and contextualisation line is named, not re-fetched as a 2022-26 AI trial. Do not invent a history generative-AI RCT.

²⁹³Kestin et al., *Scientific Reports* 15 (2025): 17458. Crossover RCT, Harvard Physical Sciences 2, Fall 2023, N = 194 eligible of 233 enrolled. Not off-the-shelf ChatGPT; delayed retention unmeasured.

²⁹⁴UC San Diego CSE Copilot redesign: Microsoft customer story (Porter, Shah); vendor customer story, not a peer-reviewed trial of Cloud Agents. Stanford CS221 Spring 2026 agentic-repo-builder ban on the course project: age-stages memo, college band. Prather et al. (2023) and the 2024 follow-on as above.

els leave warrants implicit because fluency rewards skipping them. The student who diagrams the argument, and marks what was not said, is doing the work Wineburg's line of research called sourcing and corroboration, and the work a physics student does when a derivation jumps a step. The model is a primary source that looks like a secondary source. That is a harder problem than a clumsy human paragraph, which is why it is a desirable difficulty *if* the student has enough domain knowledge to notice the jump.²⁹⁵

Evaluation is the rubric turned around. Instead of the model scoring the student, the student scores the model, against a teacher-written rubric, and has to find the weakest warrant. The grade is the quality of the critique, not the polish of a rewrite. Fan's missing evaluation cycles are the thing this task puts back.

Inference is generation the model cannot do for the student without being asked to, and that the student can be forbidden from asking. Given the model's conclusion, produce a counterexample. In mathematics this is a failed identity. In history it is a source that does not fit the narrative. In science it is an observation the mechanism cannot explain. The student who cannot generate the counterexample has not evaluated the claim. They have admired it.

Explanation, in this design, is oral. Why did you keep this sentence and cut that one? The student must interpret, check, and orally defend. That is not an integrity gimmick, though it also functions as one. It is Ennis and Facione made visible: the student must explain. Process evidence — drafts, prompts, a log of accept, challenge, or ignore — is the written form of the same demand. Chapter 5's two-lane rule is the assessment architecture this sits in: formation graded unaided, often AI-forbidden or hint-only, until a schema exists; AI-world performance allowed, logged, and judged as selection, critique, and verification once the student can detect model error.²⁹⁶ The oral defence is how the second lane stays honest.

Self-regulation is the process log, not a reflection paragraph composed after the fact with the model's help. Where did you accept the model, challenge it, or ignore it? Contractor and Reyes's 2026 distinction between augmentation (AI to explain concepts) and automation (AI to generate text) is the empirical warrant for grading *how* the tool was used, not only the essay. In their proc-tored trial, AI access raised immediate knowledge-test scores by 0.27 standard deviations, and the test gains persisted one week later; delayed unaided essay style and relevance improved for augmentation users, not for automation users, whose short-run quality gains vanished when the tool was removed.²⁹⁷ The log makes that split visible on every assignment, not only in a working paper.

²⁹⁵Desirable-difficulty frame: Elizabeth L. Bjork and Robert A. Bjork, "Making Things Hard on Yourself, but in a Good Way," in *Psychology and the Real World*, ed. Morton Ann Gernsbacher et al. (Worth, 2011), 56–64; Robert A. Bjork and Elizabeth L. Bjork, "Desirable Difficulties in Theory and Practice," *Journal of Applied Research in Memory and Cognition* 9, no. 4 (2020): 475–79. Hallucination-checking is source evaluation under high fluency only if domain knowledge is present (Willingham, "Critical Thinking"). Historical thinking as sourcing, corroboration, and contextualisation follows Wineburg's line as named in the pedagogy memo; that line was not re-fetched as a 2022–26 AI trial.

²⁹⁶Two-lane rule as theoretically motivated hypothesis from Chapter 5, not a fitted grade-by-grade schedule. Mike Perkins, Jasper Roe, Leon Furze, and Jason MacVaugh, "The Artificial Intelligence Assessment Scale (AIAS): A Framework for Ethical Integration of Generative AI in Educational Assessment," *Journal of University Teaching and Learning Practice* 21, no. 6 (2024), is shared vocabulary, not an RCT of learning.

²⁹⁷Zara Contractor and Germán Reyes, "Experimental Evidence on the Learning Impact of Generative AI," IZA Discussion Paper no. 18792 (July 2026). Working paper, not a journal article as of August 30, 2026. Coefficients from

Hallucination-checking is not a gimmick. It is source evaluation under high fluency, which is harder than evaluating a clumsy human paragraph, and therefore a desirable difficulty if — Willingham again — students have enough domain knowledge to notice errors. Nigeria’s World Bank RCT is the field specimen. De Simone, Tiberti, Barron Rodriguez, Manolio, Mosuro, and Dikoru ran after-school Microsoft Copilot (GPT-4) sessions in nine public senior-secondary schools in Benin City: twelve ninety-minute sessions over six weeks, science-of-learning prompts rather than answer-giving, teachers trained three days, students often paired on one computer. The program explicitly taught students that Copilot hallucinates and that they remain responsible for checking.²⁹⁸ Intention-to-treat English was 0.238 standard deviations; the regular third-term English exam, broader than the six-week content, was 0.206. That is not a critical-thinking transfer trial. It is evidence that hallucination-checking can be taught as a responsibility inside a supervised, curriculum-aligned lab, and that the English outcome moved. The treatment was not Copilot alone.

LearnLM and Bastani’s GPT Tutor go further by not offering the full answer. In the LearnLM exploratory RCT, expert tutors supervised every draft; 76.4 percent of drafts were approved with zero or one-to-two-character edits; zero harmful messages; five factual errors in 3,617 drafts. Immediate mistake-remediation was 93.0 percent versus 91.2 percent for human tutors and 65.4 percent for static hints.²⁹⁹ GPT Tutor withheld full solutions and included teacher-written correct answers and common-mistake hints; on the unaided exam it was indistinguishable from control, while GPT Base was about 17 percent worse.³⁰⁰ Not offering the answer is a stronger protection of the evaluative loop than asking the student to check an answer they have already seen. Both belong in the design. Neither is a transfer RCT onto non-AI critical thinking.

Ennis’s dispositions — seek reasons, do not jump — can be prompted. Prompts are cheap and habits are not. The assessment system has to reward the disposition, which is why the two-lane rule and the oral defence are not extras. A system message that says “be a critical thinker” is not a curriculum. A grade that rewards the process log, the counterexample, and the spoken defence is closer. Even then, Willingham’s constraint remains: the student who does not know the domain cannot seek reasons inside it. The agent cannot supply the domain as a substitute for knowing it. It can only supply sentences that look like knowledge.

The student, in this design, must interpret, check, and orally defend. That triad is the chapter’s operational definition of keeping judgement inside the student. If any of the three is missing, the Facione loop has been outsourced no matter what the syllabus says about “AI literacy.”

the abstract/IZA page as recorded in the pedagogy memo; full-PDF fetch was partial.

²⁹⁸Martin De Simone et al., “From Chalkboards to Chatbots: Evaluating the Impact of Generative AI on Learning Outcomes in Nigeria,” World Bank Policy Research Working Paper 11125 (2025). Students taught that Copilot hallucinates and that they remain responsible. ITT English 0.238 SD; regular third-term English exam 0.206 SD. Treatment is Copilot plus teacher-written tutor prompts plus supervised lab plus curriculum alignment.

²⁹⁹LearnLM Team, Google DeepMind, and Eedi, “AI Tutoring Can Safely and Effectively Support Students: An Exploratory RCT in UK Classrooms,” arXiv:2512.23633 (2025). Preprint. N = 165; session-level (not student-level) randomisation, so cumulative learning is not identified; human-in-the-loop is literal (every message gated).

³⁰⁰Bastani et al., “Generative AI without Guardrails.” GPT Tutor unaided exam −0.004 versus control, indistinguishable; GPT Base −0.054 versus control mean 0.321, about 17 percent worse.

Frameworks in the wild: Cornell, ISTE, AI4K12, UNESCO, Singapore

What institutions actually built, by the summer of 2026, is a set of literacy programs. They are not outcome studies. They should be described as programs. The book that treats them as evidence that “AI literacy works” will have laundered a conference profile and a newsroom pilot into a finding.

Cornell’s AI Critical Literacy Program, as reported by the *Cornell Chronicle* in August 2026, expands from a spring pilot to all incoming students plus interested faculty and staff. Four Canvas modules: What is GenAI?; Ethics; GenAI and Learning; Creating an AI Policy for Personal Use. The Center for Teaching Innovation and the Library run it. Faculty involved in the pilot came from writing, business, biology, engineering, economics, communication, and information science. The pilot reported literacy gains in how models work and are trained. Those gains are CTI-reported, not a published randomised trial.³⁰¹ The explicit goal is the right question for this chapter: help students decide whether AI use *contributes to learning*. That is Cornell’s sentence, not this book’s invention. It is also Willingham’s constraint in administrative language. A use that produces a fluent paragraph and no schema does not contribute to learning, however convenient it is for the deadline. A use that forces restatement, checking, and defence might. The program does not settle which uses do. It installs the question at the start of college, which is more than most institutions have done, and less than a transfer study.

ISTE’s Profile of an AI-Ready Graduate, unveiled at ISTE Live 2026 in Orlando, names six roles: Learner, Researcher, Synthesizer, Problem Solver, Connector, Storyteller. The profile is discipline-agnostic and built on the ISTE Student Standards. Practice reported around it includes interrogating sources with the model as coach; combining texts then verifying; ELL conversation practice; generating images *after* the student draft.³⁰² Culatta’s 14 July 2025 ASCD blog names the same six headings as ISTE Live 2026 (Problem Solver, not Ideator). ISTE’s Generation AI news page summarizing that blog lists Ideator in place of Problem Solver; some 2025 classroom-practice writeups follow the substitution. This book uses the ISTE Live 2026 six. Do not cite Culatta 2025 as an Ideator list.³⁰³ The six roles are a conference profile, not an outcomes study. They are a vocabulary for workforce-entry literacy, not a measurement of judgement.

Run each role through Willingham’s test and the profile either sits on domain knowledge or it does not. Learner: uses AI to extend practice and feedback without surrendering the struggle, and can pause the tool – that is Chapter 5’s struggle policy under another name, and it fails if the student has no schema to practise. Researcher: treats outputs as claims to be sourced, dated, and contradicted – Facione’s evaluation, empty without documents. Synthesizer: combines multiple texts or data

³⁰¹ Jose Beduya, “Cornell Expands AI Critical Literacy Program for Incoming Students,” *Cornell Chronicle*, August 2026, <https://news.cornell.edu/stories/2026/08/cornell-expands-ai-critical-literacy-program-incoming-students>. University newsroom. Pilot literacy gains are CTI-reported, not a published RCT. Explicit goal quoted from the Chronicle account: whether AI use “contributes to learning.”

³⁰² “ISTE Live 26: What Does an AI-Ready Graduate Look Like?,” *EdTech Magazine*, July 2026, <https://edtechmagazine.com/k12/article/2026/07/istelive-26-what-does-ai-ready-graduate-look>. Conference report of the ISTE Profile of an AI-Ready Graduate: Learner, Researcher, Synthesizer, Problem Solver, Connector, Storyteller. Not an outcomes study.

³⁰³ ISTE Generation AI, “Profile of an AI-Ready Graduate,” <https://iste.org/generationai/news-updates/profile-of-an-ai-ready-graduate>. Richard Culatta, “Profile of an AI-Ready Graduate,” ASCD, 14 July 2025 — same six headings as ISTE Live 2026 (Problem Solver, not Ideator). ISTE Generation AI news page summarizing that blog lists Ideator in place of Problem Solver. Do not cite Culatta 2025 as an Ideator list.

with the model, then verifies against primary sources – the then-verifies clause is the entire skill, and it is a history and science skill before it is an AI skill. Problem Solver: decomposes a real task, assigns subtasks to self versus tool, and checks – Copilot studies say the metacognitively weaker student will assign too much and check too little. Connector: uses AI to draft communications and to practise language, then shows up as a human in the real conversation. Storyteller: owns the voice; tools are production assistants after the student’s draft, which is Wren’s conference-reported practice and also the augmentation pattern Contractor and Reyes measured. None of these roles is a bicycle. Each is a domain task plus an evaluative habit. A graduate who can name the six roles and cannot catch a fluent error in the subject they claim to have studied is not AI-ready. They are credentialed.

The classroom illustration that travelled with the profile is conference-reported practice, not a districtwide agent program. Hu-Ann Wren, a gifted-education teacher at Patton Elementary in Austin, Texas, described using AI to coach source-evaluation and storytelling — Adobe Express images for a third-grader’s story — with students still drafting in their own words.³⁰⁴ That is a teacher choosing a task in which the draft stays inside the student and the image generation comes after. It is not evidence that elementary students should have a personal agent. It is not a trial. Broward County’s 2026 rule remains the documented large-district counter-example at the same ages: K–5 have no district-approved student-facing AI tools; elementary AI literacy is taught without giving children a chatbot.³⁰⁵ Conference practice and district red lines can coexist in the same year. Only one of them is a system.

AI4K12, a joint AAAI and CSTA project, is the other American vocabulary, organised as Five Big Ideas plus grade-band progression charts rather than as graduate roles. The charts are working drafts, public, and usable as a scope-and-sequence, not as an efficacy claim. In K–2, students are to recognise that computers sense via cameras and microphones, that they learn from data people give them, that photos can be faked, and to begin fairness talk. In grades 3–5: sensors versus thinking; computers find patterns in data; begin to test a simple model; altered media exist; fairness talk continues. In grades 6–8: train and test a simple model and inspect data for bias; define generative AI; recommendation algorithms; benefits and harms. In grades 9–12: run a machine-learning process; evaluate fairness, transparency, accountability, and privacy; evaluate AI content for accuracy, origin, and citation ethics.³⁰⁶ That progression is *about* AI, then *with* constrained tools, in that order. It is compatible with Willingham if the “evaluate fairness” work in high school is done on models the students have enough mathematics and civics to inspect. It is incompatible with Willingham if “AI literacy” in grade 4 is a chatbot conversation about topics the student cannot yet read.

UNESCO’s 2024 competency frameworks, launched at Digital Learning Week, are the international version of the same move. The student framework (Miao and Shiohira) specifies twelve competencies, four dimensions, and three levels: Understand, Apply, Create. The teacher framework

³⁰⁴ Same *EdTech Magazine* ISTE Live 2026 report: Hu-Ann Wren, Patton Elementary (Austin, TX), gifted-education classroom practice. Conference-reported; not a districtwide agent program.

³⁰⁵ Broward County Public Schools, Board Resolution No. 26-132 (March 10, 2026); <https://browardschools.ai/>. K–5: no district-approved student-facing AI tools; elementary literacy about AI without a student chatbot.

³⁰⁶ AI4K12 (AAAI/CSTA), “Grade Band Progression Charts — Five Big Ideas in AI,” <https://ai4k12.org/gradeband-progression-charts/>. Working drafts, public. Grade-band tasks as summarised from the age-stages memo; not an efficacy study.

(Miao and Cukurova) specifies fifteen competencies, five dimensions — human-centred mindset, ethics, foundations and applications, pedagogy, professional learning — and three levels: Acquire, Deepen, Create.³⁰⁷ They are frameworks. They are not evidence that schools have implemented them. UNESCO’s 2023 generative-AI guidance remains the prior constraint: human-centred regulation; age 13 for independent generative-AI conversation; under-13 and primary use of untested tools “not appropriate.”³⁰⁸ A ministry that adopts the competency charts and then issues a general-purpose chatbot to eight-year-olds has not implemented UNESCO. It has borrowed the logo.

Singapore is the national program that actually sequences literacy by age rather than by framework brochure. AI literacy sits in Character and Citizenship Education — deepfakes, validating generative AI — from 2026. Independent student chatbots are not the elementary practice: primary 1–3 is print-first, no student chatbots; primary 4–6 only ministry-designed Student Learning Space features under teacher supervision, as a May 2026 parliamentary reply restated. Institutes of higher learning get baseline AI competencies from academic year 2027.³⁰⁹ That is a state choosing when the student is the judge, and when the student is not yet. It is not a randomised trial of transfer. It is more honest about age than a discipline-agnostic graduate profile can be.

The U.S. Department of Education’s literacy definition, in the April 2026 supplemental priority (91 FR 18774, effective 13 May 2026), names technical knowledge, durable skills, civic awareness, ethical reasoning, and detecting AI-generated disinformation.³¹⁰ The priority is a scoring tool for discretionary grants, not a classroom mandate. The definition is a list. Lists do not form judgement. They can fund programs that might, if those programs sit on domain knowledge and require the student to produce the evaluation.

One number that will be offered as the student-side proof of “critical thinking gains” should be refused. SchoolAI’s research with Jordan School District in Utah claims a 28 percent critical-thinking gain over two years, from 13,882 conversations and 82 teachers, ESSA Level III via an Instructure review. It is vendor-commissioned, not an independent randomised trial. Scoring was a Bloom rubric of chat logs, not standardised tests. The study’s own punchline, even on its terms, is that gains tracked teacher-designed Spaces (six or more), not mere access.³¹¹ This book will not treat

³⁰⁷UNESCO, student framework: Fengchun Miao and Kelly Shiohira, twelve competencies, four dimensions, three levels (Understand / Apply / Create), 2024. Teacher framework: Fengchun Miao and Mutlu Cukurova, fifteen competencies, five dimensions (human-centred mindset, ethics, foundations/applications, pedagogy, professional learning), levels Acquire / Deepen / Create. Overview: <https://www.unesco.org/en/articles/what-you-need-know-about-unescos-new-ai-competency-frameworks-students-and-teachers>; student page <https://www.unesco.org/en/articles/ai-competency-framework-students>; teacher page <https://www.unesco.org/en/articles/ai-competency-framework-teachers>. Frameworks, not implementation evidence.

³⁰⁸Fengchun Miao and Wayne Holmes, *Guidance for Generative AI in Education and Research* (UNESCO, September 7, 2023). Age 13 for independent GenAI conversation; under-13 / primary use of untested tools “not appropriate.”

³⁰⁹Singapore Ministry of Education, page updated July 31, 2026, as recorded in the responsible-use memo. CCE AI literacy (deepfakes, validating GenAI) from 2026; IHL baseline competencies from AY2027; P1–3 no student chatbots; P4–6 only SLS education-designed tools under teacher supervision (parliamentary reply, May 6, 2026).

³¹⁰“Final Priority and Definitions—Secretary’s Supplemental Priority and Definitions on Advancing Artificial Intelligence in Education,” 91 Fed. Reg. 18774 (April 13, 2026), effective May 13, 2026. Discretionary-grant scoring tool, not a classroom mandate. AI defined via 15 U.S.C. § 9401(3); the priority does not define “agentic AI.”

³¹¹SchoolAI research page on Jordan School District (Utah): claimed 28 percent critical-thinking gain over two years (13,882 conversations, 82 teachers); ESSA Level III via Instructure. Vendor-commissioned; Bloom-rubric scoring of chat logs, not standardised tests; not an independent RCT. Gains tracked teacher-designed Spaces (six or more), not mere access, per the landscape memo.

that 28 percent as an RCT, and will not treat a Bloom-scored chat log as Facione’s evaluation. A conversation that a rubric codes as “analysis” can be the model analysing. Willingham’s test is what the student can do when the log is closed.

Seven approaches, correlational fog, and the live debate

Ethan Mollick and Lilach Mollick’s 2023 working paper, “Assigning AI: Seven Approaches for Students, with Prompts,” is a practical prompt library: AI as tutor, coach, mentor, teammate, student, simulator, or tool.³¹² It is not a theory of critical thinking. It is the library the Nigeria RCT actually used. That fact matters more than the paper’s citation count. De Simone and colleagues did not hand students a chatbot and a hope. Teachers were trained for three days; sessions used science-of-learning prompts — retrieval, elaborative interrogation, concrete examples — rather than answer-giving; the English curriculum stayed the curriculum. The teacher still chose the approach per task. A prompt library is not a curriculum unless someone with pedagogical authority selects, constrains, and withholds.

The seven roles are not equal from Willingham’s point of view. Tutor, in the answer-giving sense, is GPT Base. Tutor, in the hint-and-grounded-solution sense, is GPT Tutor and the Nigeria lab. Coach and mentor can keep the evaluative loop if they refuse the product and demand the student’s attempt first. Teammate and tool will finish the job unless the assignment forbids it. Student – the model as the learner the human has to teach – is the one role that forces explanation, and it only works if the human actually knows the content being taught. Simulator can be a desirable difficulty in a well-designed scenario, or a video game. The library does not encode those distinctions. The teacher does. That is why Mollick and Mollick, used well in Benin City and used as a paste-the-prompt kit elsewhere, are not a finding about critical thinking. They are a finding about orchestration, which is Chapter 7’s subject.

That “unless” is the bridge to Chapter 7. The seven approaches do not choose themselves. Default chat will choose “tool” and “tutor” in the answer-giving sense, because those are the uses that finish the homework. Coach, student (the model being taught by the learner), and simulator require a teacher to set the role and a grade to make the role costly to abandon. Mollick and Mollick are not a replacement for Willingham. They are a menu that Willingham’s constraint, and Chapter 5’s struggle policy, have to filter.

The live debate among literacy programs is a debate among good questions, not a debate this book should settle by picking a logo. Cornell asks whether a given use contributes to learning. ISTE’s six roles ask what workforce-entry profile a graduate should be able to inhabit — Learner, Researcher, Synthesizer, Problem Solver, Connector, Storyteller. UNESCO’s frameworks ask for civic and ethical competencies at specified levels of depth. The Department of Education’s grant definition asks for technical knowledge, durable skills, civic awareness, ethical reasoning, and detection of AI-generated disinformation. Those questions are not the same. They are already in conflict in the institutions this book’s sources actually name. Cornell’s question is a learning question: does this use contribute to the skill the course is for. Clemson’s newsroom framing of ChatGPT Edu,

³¹²Ethan Mollick and Lilach Mollick, “Assigning AI: Seven Approaches for Students, with Prompts,” arXiv:2306.10052 (2023). Tutor, coach, mentor, teammate, student, simulator, tool. Used as prompt source in De Simone et al., World Bank PRWP 11125.

GPT-5.5, and Codex access leans the other way – students should be able to work with these tools – without a published learning-outcomes study attached to the access.³¹³ Journalism around campus assistants at other universities has carried the same workforce-entry pressure; some of that journalism was not verified against primary campus pages. The MIT Media Lab preprint sits in the same debate as a caution about cognitive debt, and should not be asked to decide it. Workforce-entry literacy and cognitive formation are both real jobs. They are not the same job. A program that only does the first will produce graduates who can drive tools they cannot check. A program that only does the second will produce graduates who can think in a world whose workplaces assume the tools. The two-lane rule is how this book holds both without pretending they are one outcome. A student can recap UNESCO's dimensions and still paste the essay. A student can inhabit ISTE's Storyteller role with a model-written narrative and an image generated after a one-line prompt. A student can pass Cornell's modules on how models are trained and still not know enough U.S. history to catch a fluent fabrication. The book that picks a brand of framework and ignores Willingham will have built a checklist students can recap and not execute.

ISTE's six roles remain a profile unveiled at a conference, not an outcomes study. Cornell's gains remain CTI-reported from a spring pilot. Gerlich remains correlational. SchoolAI's 28 percent remains vendor Bloom-rubric scoring of chat logs. UNESCO's frameworks remain frameworks. None of that is a reason to drop literacy. It is a reason not to confuse literacy-as-recap with judgement-as-execution.

The design implication is narrower, and harder, than a new framework. Put domain knowledge under the literacy module. Require interpretation, checking, and oral defence on the tasks where the model is allowed in. Grade the student's judgement, not the model's paragraph. Keep a formation lane in which the model is forbidden or hint-only until the student can detect error. Do not claim that this package has been shown, in a randomised K–12 trial, to transfer critical thinking to non-AI tasks. Claim only that the causal record on default chat — Fan's missing evaluation cycles, Stadler's thinner arguments, Bastani's copied errors, Prather's shepherded novices — is what you get when you skip it.

If judgement needs something to judge, someone has to keep the struggle policy and the off-switch. That person is not the vendor. That person is not the model. Chapter 7 is about who that person is, what the unions have put in writing, and why the slogan that teachers who do not use AI will be replaced is folklore.

³¹³Clemson News documents institutional access (ChatGPT Edu, Codex, locally hosted open-weight models), not a published learning-outcomes study. UT-Knoxville / UT Verse details in EdTech Magazine (December 2025) were not fetched against primary docs – journalism-sourced, UNVERIFIED. The learning-versus-workforce tension is the live debate named in the book brief (Cornell's question versus campus AI-proficient-at-graduation rhetoric), not a claim that Clemson ran a critical-thinking RCT.

Chapter 7

Who Turns the Agent Off



Figure 8: A closed book, a school bell, and a dark computer screen on a desk.

The slogan is easy to chant and hard to source. “AI won’t replace teachers, but teachers who use AI will replace those who don’t.” No ministry paper, union text, or peer-reviewed article fetched for this book originated that sentence as policy. The earliest education-specific use the research found is a LinkedIn post by Dan Fitzpatrick, *The AI Educator*, under almost exactly that title.³¹⁴ It is a

³¹⁴Dan Fitzpatrick, “AI Won’t Replace Teachers, but it will replace teachers who don’t use AI,” LinkedIn, <https://www.linkedin.com/pulse/ai-wont-replace-teachers-who-dont-use-dan-fitzpatrick>. Earliest education-specific

recycled template – computers, Excel, social media, coding: X will not replace Y, but Y who use X will replace Y who do not. Templates are not doctrines.

Education Week, in a 17 October 2025 opinion by Ingrid Guerra-López, published a different line: teachers who use AI will *change teaching*. Different verb. Anti-replacement, not a forecast of headcount.³¹⁵ Changing teaching is what orchestration is: lanes, struggle policy, presence, oral checks, copilot use. Replacing teachers is what the causal record does not show and what NYC’s red list forbids. Holding those two verbs apart is the whole chapter. The U.S. Department of Education, in May 2023, did not hedge. The Department “firmly rejects the idea that AI could replace teachers.”³¹⁶ The American Federation of Teachers, in Value 1 of its March 2025 *Commonsense Guardrails for Using Advanced Technology in Schools*, put privacy and safety first: the well-being and data security of students must be paramount, “even if it means sacrificing some potential educational opportunities.”³¹⁷ Folklore, then doctrine. This chapter is about the doctrine, the evidence, and the off-switch.

Every positive generative-AI field result in Chapter 3 either embeds the model in a teacher-run session or helps the tutor. Khanmigo without engagement looks like software, not a tutor. Replacement is not supported by the 2022-26 causal record. Orchestration is what the positive studies already quietly did. The real labor split is not users versus non-users. It is who controls procurement, evaluation, and the right to switch the agent off.

This chapter has a narrower job than a labor history. It marks the slogan as folklore so later chapters cannot smuggle it back. It reads the 2022-26 field record as a record of orchestration, not of replacement. It translates Trouche into five Monday moves. It puts NEA’s five principles and AFT’s guardrails and resolutions on the page as primary texts. It stages the Academy as a live fight, not a cartoon. It names the teacher-facing products that actually shipped, with their evidence labels attached. It refuses a PD-to-pupil-outcomes chain the reviews do not show, and it steals the CAL analogue with the label still on. Then it hands the off-switch to Chapter 8 as a set of instruments rather than as a feeling. Documents do that work. A slogan cannot. The rest of the chapter is the documents: union PDFs, a Department of Education sentence, field RCTs, product pages labelled as product pages, and a PD literature that stops where it stops.

use found of the replacement slogan. Practitioner-influencer post, not a union, ministry, or peer-reviewed source. Recycled “X won’t replace Y but Y who use X” template.

³¹⁵Ingrid Guerra-López, “AI Won’t Replace Teachers–But Teachers Who Use AI Will Change Teaching,” *Education Week* (opinion), October 17, 2025, <https://www.edweek.org/technology/opinion-ai-wont-replace-teachers-but-teachers-who-use-ai-will-change-teaching/2025/10>. Different verb from the folklore slogan; anti-replacement.

³¹⁶U.S. Department of Education, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023), <https://www.ed.gov/sites/ed/files/documents/ai-report/ai-report.pdf>. The Department “firmly rejects the idea that AI could replace teachers.” Seven recommendations; humans in the loop; FERPA/IDEA already apply.

³¹⁷American Federation of Teachers, *Commonsense Guardrails for Using Advanced Technology in Schools*, version 2 (March 2025), https://www.aft.org/sites/default/files/media/documents/2025/Commonsense_Guardrails_AI_V2_0318.pdf. Official title is “Using Advanced Technology in Schools,” not “for Artificial Intelligence in Education.” Value 1: “Student Privacy and Safety Come First”; the quoted sacrifice clause is the second sentence of that value (“The well-being and data security of students must be paramount, even if it means sacrificing some potential educational opportunities”), not the first.

Mark the replacement slogan folklore

Correction 6 of this book has to be written at full length here, so it cannot leak back in as a finding later. The replacement slogan is folklore. Treating it as a research result, a union position, or a Department of Education forecast is a factual error.

Search of union texts, ministry papers, the Department’s 2023 report, UNESCO, OECD, AFT, and NEA found no official source that uses “teachers who use AI will replace those who don’t” as policy.³¹⁸ What exists on the official side is a “will not replace” bloc, named and dated. UNESCO’s human-centred line. The International Task Force on Teachers for Education 2030, in a September 2025 position paper, treats teachers as irreplaceable and AI as augmentative – listed here as that paper’s position, not as a fully extracted document.³¹⁹ NEA 2024: aid, not replace; humans at the center of high-stakes decisions. AFT 2025-26: the same, in guardrails and in the July 2026 resolution. New York City Public Schools, the nation’s largest system: “AI is not a replacement for your child’s teacher.”³²⁰ Anthropic, launching Claude for Teachers in July 2026, framed the product as returning time to students, not as a substitute for the adult in the room.³²¹ Vendor framing is not a union text, but it is evidence that even the companies selling teacher-side agents do not, in their education SKUs, lead with replacement.

Pressure in the other direction is real, and it should be named without being laundered into a finding. Teacher shortages are a labor fact. The Department’s 2026 supplemental priority funds tutoring, personalization, and administrative efficiency as grant-scoring prongs – adoption language, not a replacement mandate, and not a child-safety regulation.³²² Brookings has written about generative-AI tutoring at scale in teacher-shortage systems. xAI’s El Salvador announcement (11 December 2025) promised Grok across more than 5,000 public schools and personalized learning to over one million students – Grok the assistant, not Grok Bot, announced, with no outcome evaluation fetched.³²³ Vendor talk of an “AI operating system for schools,” BYJU’S 2023 rhetoric, and the tutor-for-every-child line associated with Karpathy, Altman, and Khan are the replacement *pressure*. Unions exist to refuse it. The slogan takes that pressure and recites it as if it were a labor-market law.

Education International’s David Edwards, in a July 2024 World Economic Forum interview, is the

³¹⁸Search reported in the responsible-use memo, finding 9 and section 5. No official source found that originated the replacement slogan as policy.

³¹⁹International Task Force on Teachers for Education 2030, *Promoting and Protecting Teacher Agency in the Age of Artificial Intelligence* (September 2025), <https://teachertaskforce.org/sites/default/files/2025-09/1149%5F25%5FPromoting%20and%20Protecting%20Teacher%20Agency%5FFINAL%5F3Sep.pdf>. Listed as teachers irreplaceable, AI as augmentative; not fully extracted this pass. Do not invent additional sentences.

³²⁰New York City Public Schools, “Guidance on Artificial Intelligence” (March 2026), <https://www.schools.nyc.gov/about-us/vision-and-mission/artificial-intelligence/guidance-on-artificial-intelligence>. “AI is not a replacement for your child’s teacher,” as recorded in the landscape memo.

³²¹Anthropic, “Introducing Claude for Teachers” (July 14, 2026), <https://www.anthropic.com/news/claude-for-teachers>. Returning-time-to-students framing is vendor copy, labelled as such.

³²²“Final Priority and Definitions–Secretary’s Supplemental Priority and Definitions on Advancing Artificial Intelligence in Education,” 91 Fed. Reg. 18774 (April 13, 2026), effective May 13, 2026. Discretionary-grant scoring tool, not a classroom mandate and not a replacement order.

³²³xAI, El Salvador partnership announcement, December 11, 2025, <https://x.ai/news/el-salvador-partnership>. More than 5,000 public schools; over one million students; Grok the assistant, not Grok Bot. No outcome evaluation fetched. Announced program, not a documented completed deployment as of August 30, 2026.

equity argument the slogan erases. The experience of education is relational, not transactional. AI will not make teachers obsolete; radio and the VCR did not either. The nightmare is not that every teacher will have to use the tool to keep a job. The nightmare is human teachers as a luxury of the rich, bots for the poor.³²⁴ That is a distributional claim, not a productivity claim. The replacement slogan pretends those are the same.

The empirical gap under the slogan is not subtle. No study fetched for this book shows an autonomous agent substituting for a certified teacher of record in a U.S. district.³²⁵ New York City's March 2026 guidance red-lights the decisions that would make substitution operational: placement, discipline, eligibility, promotion, graduation; IEP and 504 development; assessments and grading as the decision; surveillance and behavior monitoring; counseling and crisis.³²⁶ That is the current legal-practical ceiling in the largest U.S. system. An agent that cannot grade, cannot write an IEP, cannot discipline, and cannot counsel is not a teacher of record. It is a tool a teacher may be told to use, or may refuse, depending on who holds the off-switch.

NYC's yellow and green lights complete the picture. Yellow, human review required: data-pattern surfacing; translations of critical communications (qualified linguist); scaffolds for diverse learners (certified staff review); student use for research, exploration, or creative projects with educator guidance. Green, encouraged, still ERMA-approved, and the human owns the output: lesson brainstorming, drafting communications, scheduling and summarizing non-sensitive information, operational synthesis, non-critical translation with a disclaimer if unreviewed, accessibility, professional learning.³²⁷ The traffic light is a theory of the teacher. Green is orchestration of teacher work. Yellow is orchestration of student work under review. Red is the set of acts that make a person the teacher of record. An agent licensed for green is not, by that license, a candidate for red. The slogan collapses the lights.

Replacement, if it ever arrived as a research object, would have to be studied as replacement: a certified adult taken out of the instructional cycle, an agent put in, learning and safety and equity measured against the human classroom the agent displaced. Bloom measured one-to-one human tutoring against group instruction. Nickow still pools human tutoring. Khanmigo was added to an existing RTI block with an adult in the room. Nigeria added a supervised lab. LearnLM gated every message. Tutor CoPilot helped tutors. None of those designs is replacement. Citing them as proof that teachers who do not use AI will lose their jobs is not a misreading at the margin. It is a category error.

The slogan also misstates the labor market it claims to describe. The scarce resource in a shortage is not "teachers who can prompt." It is teachers who can keep struggle in the student, notice unpro-

³²⁴David Edwards (Education International), interview, World Economic Forum, July 2024, <https://www.weforum.org/stories/2024/07/artificial-intelligence-education-teachers-union/>. Relational not transactional; radio/VCR; human teachers as luxury of the rich.

³²⁵Landscape memo section 6.3; responsible-use finding 9. No fetched study of an autonomous agent as certified teacher of record in a U.S. district.

³²⁶NYCPS March 2026 guidance, red list: placement, discipline, eligibility, promotion, graduation; IEP/504 development; assessments/grading as the decision; surveillance/behavior monitoring; counseling/crisis; student-data model training; PII in unapproved tools.

³²⁷NYCPS March 2026 guidance, red list: placement, discipline, eligibility, promotion, graduation; IEP/504 development; assessments/grading as the decision; surveillance/behavior monitoring; counseling/crisis; student-data model training; PII in unapproved tools.

ductive struggle, certify mastery, and hold a relationship that a companion chatbot is legally and developmentally barred from faking. Edwards's radio-and-VCR line is the right historical control. Tools arrive. Relational work remains. What changes, when policy is weak, is who gets the human. That is why NEA puts high-stakes decisions in bargaining and why AFT puts the kill switch in the annex. The fight is not adoption. The fight is allocation.

A note on what this chapter is not doing with El Salvador. The December 2025 announcement is Grok-as-tutor, at national-announcement scale, in a teacher-shortage frame. It is not Grok Bot. It is not a U.S. district. It has no fetched outcomes. Using it as proof that replacement is already policy is the same collapse Chapter 2 exists to prevent. Using it as proof that replacement is impossible is equally wrong. It is an announced program in a different product class, and it remains an announcement. The U.S. ceiling that this chapter can stand on is NYC's red list plus the absence of a teacher-of-record study. That is enough to retire the slogan. It is not enough to retire the pressure.

Fitzpatrick is the pressure narrative, not a finding. Edwards, the Department of Education, NEA, and AFT are the doctrine against it. The template should not re-enter later chapters as a "some say." Folklore tracked is folklore retired.

The agent-specific sub-debate is already visible in 2026 product design. Teacher-side Cowork – a 4 p.m. job on de-identified exit tickets – is politically easier than a student-side Hermes runtime on a VPS. Unions may accept the former and fight the latter. AFT's July 2026 resolution opposes student-facing AI in elementary and would ban social-companion chatbots at least until age 16.³²⁸ That is not Luddism about spreadsheets. It is a developmental and safeguarding line drawn against a product class the causal record has not made into a teacher.

Every positive field result had a teacher or tutor in the loop

Replacement is the wrong theory of the evidence. Bloom's 1984 two-sigma observation was about one-to-one *human* tutoring.³²⁹ Nickow, Oreopoulos, and Quan, pooling experimental preK-12 tutoring in 2024, still put human tutoring near the top of the causal hierarchy: 0.288 standard deviations (SE 0.029).³³⁰ Earlier computer-assisted learning typically sits in 0.05-0.20 SD. Those are the benchmarks. A chatbot is not a tutor because it talks. A tutor is a design plus a human who sets the struggle policy and notices when the design fails.

Magnitudes make the slogan's economics look like folklore too. Human high-dosage tutoring at 0.288 SD is the causal bench. Earlier computer-assisted learning typically 0.05-0.20 SD. Khanmigo's intention-to-treat in Hamilton County is about 0.06-0.08 SD per school year, resembling practice without the chatbot, at about 15 dollars per student per year for the license, delivered inside already-staffed blocks. Tutor CoPilot's session mastery lift is 4 percentage points, 9 for lower-rated

³²⁸American Federation of Teachers, "Preparing for the Age of Artificial Intelligence: Protecting Democracy, Workers, Students and the Planet," resolution adopted July 18, 2026, <https://www.aft.org/resolution/preparing-age-artificial-intelligence-protecting-democracy-workers-students-and-planet>.

³²⁹Benjamin S. Bloom, "The 2 Sigma Problem: The Search for Methods of Group Instruction as Effective as One-to-One Tutoring," *Educational Researcher* 13, no. 6 (1984): 4-16.

³³⁰A. Nickow, P. Oreopoulos, and V. Quan, "The Promise of Tutoring for PreK-12 Learning: A Systematic Review and Meta-Analysis of the Experimental Evidence," *American Educational Research Journal* 61, no. 1 (2024): 74-107. Pooled ES 0.288 SD (SE 0.029).

tutors, at about 20 dollars per tutor per year. Nigeria's English ITT is 0.238 SD inside a teacher-run lab. Bastani's unguarded GPT made subsequent unaided math worse. None of these numbers is a teacher salary. None of them is a license to take the adult out of the RTI block or the lab or the exam. A district that reads "0.07 SD and cheap" as "we can replace the teacher of record" has done arithmetic that no paper in the record performed. The papers put the adult in. The slogan takes the adult out and keeps the effect size. That is not how identification works.

Chapter 3's positive generative-AI results are, without exception, human-in-the-loop results of a specific kind.

Nigeria: a teacher-run after-school lab, three days of teacher training, science-of-learning prompts rather than answer-giving, students often paired on one computer, curriculum-aligned English. Intention-to-treat English 0.238 SD; regular term exam 0.206 SD. The treatment is Copilot plus teachers plus the lab, not Copilot alone.³³¹

Kestin's Harvard Physical Sciences 2: course staff, instructor-written step-by-step solutions in the prompt, sequential scaffolding, versus the course's existing in-class active learning. Immediate post-test gains; delayed retention unmeasured; not off-the-shelf ChatGPT; a selective university.³³² The human's job was not to re-teach every item. It was to encode the physics the model was allowed to use.

LearnLM, Google DeepMind and Eedi, five UK secondary schools: expert tutors supervised every draft. 76.4 percent of drafts approved with zero or one-to-two-character edits; zero harmful messages. Immediate remediation comparable to human tutors. Session-level randomisation, so cumulative learning is not identified; the human-in-the-loop is literal – every message gated. A qualitative finding that supervising tutors learned Socratic moves from the model is interview evidence from $n = 5$, not a professional-development trial, and this book will not promote it as one.³³³

Khanmigo in Hamilton County, Tennessee: an adult already in the RTI block. Intention-to-treat about 0.06-0.08 SD per school year, resembling Khan Academy practice without the chatbot. Median student messaged on one-third of practice days and in 17 percent of mistake sessions. Only 14.5 percent of messages contained a mathematical question or a step of reasoning. Chat counts do not predict achievement once mastery is controlled. A 2026 interface change that opened Khanmigo by default did not, in the authors' message data, raise messaging rates.³³⁴ Access is not tutoring. Khanmigo without engagement looks like software. The adult in the room is the engagement technology the model does not supply.

Tutor CoPilot, Wang, Ribeiro, Robinson, Loeb, and Demszky: about 900 tutors, about 1,800 K-12 students. Tutors randomly given a real-time LLM that suggests expert-like moves. Students 4 percentage points more likely to master the session topic; 9 percentage points for lower-rated tutors.

³³¹Martin De Simone et al., "From Chalkboards to Chatbots," World Bank Policy Research Working Paper 11125 (2025).

³³²G. Kestin et al., "AI Tutoring Outperforms In-Class Active Learning," *Scientific Reports* 15 (2025): 17458.

³³³LearnLM Team, Google DeepMind, and Eedi, arXiv:2512.23633 (2025). Preprint. Supervising-tutors-learned-Socratic-moves finding: interview evidence from $n = 5$, not a PD trial.

³³⁴P. Oreopoulos and N. Low, "One Click Away: AI Tutoring with Khanmigo in a Two-Year School Experiment," EdWorkingPaper 26-1551 / NBER Working Paper 35620 (2026). ITT about 0.06-0.08 SD per school year; TCAP pending. Working paper.

Tutors asked more guiding questions and gave away fewer answers. About 20 dollars per tutor per year.³³⁵ The model helps the *tutor*. It does not replace the tutor. That is the opposite of the slogan.

Bastani's GPT Tutor belongs on this list as the guardrail result, not as a learning *gain*: teacher-written correct solutions and common-mistake hints in the prompt; full solutions withheld; unaided exam indistinguishable from control, while unguarded GPT Base was about 17 percent worse.³³⁶ The teachers encoded the mistakes. Passive use of a general chatbot is not that job.

NUS AICET's 2025 "embedded AI designer" framing names the job at the faculty end: encode pedagogical intent into constraints, as Bastani's teachers encoded solutions and mistakes, as Kestin encoded step-by-step physics.³³⁷ Orchestrated mediation, not a warm body in the building. Human-in-the-loop that *helps* is specific: grounded prompts, attempt gates, circulating teachers, tutor-facing copilots. A facilitator of feelings while the agent teaches is not in the causal record. It is a marketing theory of the teacher, and the 2022-26 experiments did not run it.

"Human in the loop" is used loosely in product literature. Pedagogically it means a human still performing the functions AI is bad at or that must remain the student's: goal setting, noticing when struggle has become unproductive, socio-emotional regulation, certifying mastery, and – for the teacher – designing which loop the AI is even in.³³⁸ The loop is not a warm body. A paraprofessional told to watch a lab while students paste into an unguarded chatbot is in the room and not in the loop. A teacher who encoded common mistakes into GPT Tutor's prompt and then gave a closed-book exam is in the loop without standing over every desk. Tutor CoPilot puts the model in the tutor's loop, not the student's. Those are different jobs. The slogan that collapses them into "using AI" will hire for the wrong one.

Henkel, Horne-Robinson, Kozhakhmetova, and Lee's 2024 Rori result in Ghana belongs on the structured-curriculum side of this ledger, with the authors' own caution. About 500 Ghanaian students; extra one hour per week on a phone with a WhatsApp AI math tutor versus usual classes; growth-score Cohen's $d = 0.36$; marginal cost about 5 dollars per student; year-1 result.³³⁹ It is not a teacher-of-record study. It is extra time on a constrained tutor in a system that already has classes. Brookings can cite it in a shortage conversation. This book will not cite it as replacement. Extra time is extra time. The adult who scheduled the hour is still in the design.

What the negative and null results add is the same lesson from the other side. Bastani's GPT Base – unguarded, answer-giving – harmed subsequent learning. Oreopoulos and Low's Khanmigo ITT is small and looks like practice without the chatbot, because students mostly did not talk to it. Fan's

³³⁵R. E. Wang, A. T. Ribeiro, C. D. Robinson, S. Loeb, and D. Demszky, "Tutor CoPilot: A Human-AI Approach for Scaling Real-Time Expertise," arXiv:2410.03017 (2024). About 900 tutors, about 1,800 K-12 students; +4 pp topic mastery; +9 pp for lower-rated tutors; about \$20 per tutor per year.

³³⁶Hamsa Bastani et al., "Generative AI without Guardrails Can Harm Learning," *PNAS* 122, no. 26 (2025): e2422633122.

³³⁷AICET / National University of Singapore, "The Educator as an Embedded AI Designer: Faculty Agency in GenAI Systems" (2025), https://aicet.comp.nus.edu.sg/wp-content/uploads/2025/10/The_Educator_as_an_Embedded_AI_Designer_Faculty_Faculty_Agency_in.pdf.

³³⁸Pedagogy memo, section 7.1: pedagogical meaning of human-in-the-loop. Not vendor "HITL" marketing.

³³⁹O. Henkel, H. Horne-Robinson, N. Kozhakhmetova, and A. Lee, "Effective and Scalable Math Support: Experimental Evidence on the Impact of an AI Math Tutor in Ghana," *AIED* 2024, CCIS 2150; arXiv:2402.09809. About 500 students; Cohen's $d = 0.36$; about \$5/student; year-1; authors flag caution. Extra time on a constrained tutor, not a teacher-of-record study.

ChatGPT arm improved essays and not knowledge. Stadler’s LLM users thought less and argued worse. None of those failures is a failure of “not enough AI.” They are failures of a design that took the teacher, or the tutor, or the evaluative loop, out. Replacement is the limiting case of that design. The causal record, such as it is, runs the other way.

NUS AICET’s embedded-designer language is worth keeping next to Trouche, because it names a job that looks like engineering and is actually pedagogy. Faculty who encode constraints – what the model may say, what it must withhold, which mistakes it should hint toward – are doing the Monday work before Monday. Districts that buy a general chatbot and ask teachers to “integrate AI” are asking for orchestration without a design. The first is what Bastani and Kestin did. The second is what GPT Base was. PD that cannot tell those apart will measure self-efficacy after a demo and call it capacity.

The IES What Works Clearinghouse, as of this research’s access date, has no intervention report specific to generative AI tutors.³⁴⁰ Human high-dosage tutoring remains the benchmark. Replacement would require a study that puts an agent in the teacher-of-record role and measures learning, safety, and equity against a human classroom. That study was not found. Until it exists, the slogan is not ahead of the evidence. It is orthogonal to it.

What orchestrator means on Monday

Luc Trouche’s 2004 account of instrumental orchestration is the concept this chapter needs without becoming a theory paper. The teacher regulates when a tool is activated, what is delegated, who validates, and how visible the tool is. The 2025 AISIO extension – AI-Supported Instrumental Orchestration, a lecture, not an RCT – restates the same regulation for models that can act: the teacher still holds activation, delegation, validation, and visibility.³⁴¹ That is the opposite of “the agent teaches; the teacher facilitates feelings.”

The Monday list is a synthesis of what the positive studies already did, labelled as such, not a new experiment.

First, choose the lane per assignment. Formation: AI forbidden or hint-only; grade unaided retrieval, derivation, draft, or oral work, until a schema exists. AI-world performance: allowed, logged; grade judgement – selection, critique, verification, process – once the student can detect model error. Chapter 5 stated the two-lane rule as a theoretically motivated hypothesis, not a fitted grade-by-grade schedule. The orchestrator’s first job is to pick a lane and say so, in writing, on the assignment. AFT’s full / partial / no-AI labels are the union version of the same move.³⁴² A building-wide “we use AI” is not orchestration. It is a slogan with a login.

Second, configure or select the agent’s struggle policy. Attempt gates. Hint ladders. Grounded

³⁴⁰IES What Works Clearinghouse: older ITS reports exist (e.g., Cognitive Tutor Algebra I); no WWC intervention report specific to generative AI tutors as of August 30, 2026 (pedagogy memo).

³⁴¹Luc Trouche, instrumental orchestration (2004), as named in the pedagogy memo; the 2004 original was not re-fetched as a full PDF this pass. Luc Trouche, “Mathematics Education and AI: Rethinking Instrumental Orchestration as the Structured Regulation of AI’s Pedagogical Agency,” French Institute of Education, ENS de Lyon (2025 lecture). AISIO is conceptual, not an RCT.

³⁴²AFT *Commonsense Guardrails for Using Advanced Technology in Schools* v2 (March 2025), assignment labels; two-lane rule as Chapter 5 hypothesis.

answers. Time-to-first-hint delays. Teacher-authored misconception libraries. That is Bastani's GPT Tutor, Kestin's sequential scaffolding, Nigeria's science-of-learning prompts, LearnLM's refusal to dump the solution, the UK product-safety preference for hints before answers. A single always-Socratic or always-answer policy is incoherent with cognitive load theory, as Chapter 5 argued. The orchestrator sets the policy per task and per expertise level. The vendor's default will be answer-giving, because that is what finishes the job and what students will request.

Third, be in the room for engagement and for when struggle turns unproductive. Khanmigo's lesson is that a licensed tutor sitting one click away is unused on most mistake sessions. Hiebert and Grouws's productive struggle is intellectual effort to make sense of something that is not immediately apparent; unproductive struggle is overload, confusion about the goal, or a task that is not appropriate.³⁴³ The teacher notices which is which. The model does not. Robinson, Gormley, Ribeiro, and Loeb's 2026 finding, "Access Is Not Enough," is the engagement paper adjacent to Khanmigo: dedicated time still produced only minutes of weekly use on an adaptive literacy platform until human tutors raised usage – and even then, usage without achievement is not a win.³⁴⁴ The orchestrator is the engagement technology. Software that waits to be opened is not.

Fourth, run the oral or process check that makes outsourcing visible. Chapter 6's demand – interpret, check, orally defend – is not self-enforcing. Someone has to hear the defence, read the process log, and grade the critique rather than the polish. Detection software is not that someone. Detectors fail, and they fail unfairly; that is Chapter 8's integrity argument. The orchestrator's check is pedagogical. It is also labor. US ED 2023 warned that "in the loop" done badly *adds* work: setting every option, reviewing every draft.³⁴⁵ Orchestration that cannot be staffed is not a design. It is an unfunded mandate.

Fifth, use AI yourself as a tutor-copilot, not as a replacement for noticing. Tutor CoPilot is the causal specimen: suggestions to the tutor, the tutor still in the conversation, lower-rated tutors gaining more. Teacher-facing generation of lesson ideas, checks, and drafts – Claude Cowork, MagicSchool, Diffit, ChatGPT for Teachers – can return time if the teacher still owns the output. It cannot notice the child who has stopped struggling and started pasting. Noticing is not a tool call.

Mollick and Mollick's seven approaches sit here, not in a literacy module. The teacher still chooses the approach per task. Tutor, coach, mentor, teammate, student, simulator, tool: six of the seven will finish the homework if the assignment lets them. The seventh, the model as student, only works if the human knows the content. Orchestration is the choice. A prompt library without a chooser is default chat.

The five moves are not a personality type. They are a workload and a skill. A teacher who can do all five on Monday for thirty students, across two preps, without PD time and without a kill switch, is a hypothetical this book will not invent as a scene. US ED 2023 already warned that in-the-loop review can add work. AFT's right-to-disconnect and PD-time demands are what make the Monday

³⁴³James Hiebert and Douglas A. Grouws, "The Effects of Classroom Mathematics Teaching on Students' Learning," in *Second Handbook of Research on Mathematics Teaching and Learning*, ed. Frank K. Lester Jr. (Information Age, 2007), 371-404, productive struggle at p. 387.

³⁴⁴C. D. Robinson, D. Gormley, A. T. Ribeiro, and S. Loeb, "Access Is Not Enough: Human Support Improves Engagement with AI Tutoring," EdWorkingPaper 26-1451 (2026). Cited via Oreopoulos and Low.

³⁴⁵US ED OET, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023). McKinsey workload analysis cited as hope, not result; "in the loop" done badly adds work.

list possible rather than pious. Orchestration that is only a verb in a keynote is replacement with better branding: the agent does the job, the teacher is told they orchestrated it, the grade prints.

Visibility, in Trouche’s list, is the move most easily surrendered to a dashboard. A teacher-visible log is not the same as a teacher who has read the log. Common Sense Media’s low-risk rating of Khanmigo comes with the observation that there is no student privacy on Khanmigo in the sense that supervising adults see every turn – a safeguarding feature, and a labor feature, and not by itself engagement.³⁴⁶ Oreopoulos and Low’s message data show what “visible” looks like when the student does not type: empty logs, thin ITT, an interface change that opened the tutor by default and did not raise messaging. The orchestrator who trusts the dashboard has delegated noticing to a count. Chat counts, in that study, do not predict achievement once mastery is controlled. Noticing is still a person.

Mollick’s menu, Trouche’s regulation, and Chapter 5’s two lanes are the same architecture in three vocabularies. The teacher chooses a role for the tool, regulates when it is on, and grades the student’s remaining work. A school that publishes an AI policy without assignment-level lanes has published a press release. A school that publishes lanes without oral or process checks has published a honor code the model can complete. A school that publishes checks without a teacher who has time to run them has published unfunded orchestration. The Monday list is only as real as the contract that staffs it.

NEA, AFT, bargaining, and the July 2026 resolution

Unions have moved from warning to co-governing. The texts are primary sources. They should be on the page as texts, not as a vibe that “teachers are worried.”

The National Education Association’s 2024 Policy Statement and task-force report state five principles.³⁴⁷ One: students and educators must remain at the center; AI must not displace the relationship. Two: evidence-based – independent research, or industry research that meets the same bar; otherwise a pilot with a kill switch. Three: ethical development and use, strong data protection, no high-stakes automation. Four: equitable access to and use of AI tools, including connectivity and anti-surveillance. Five: ongoing education with and about AI – literacy and agency. Explicitly: AI shall not undercut educators via surveillance, rights erosion, or takeover of core human job functions; that belongs in collective bargaining. AI data alone never determines evaluation, placement, discipline, diagnosis, or safety decisions. Facial recognition: New York already banned it in schools in September 2023; NEA cites this as the right direction.³⁴⁸ Principle 2’s kill switch is

³⁴⁶Common Sense Media / Youth AI Safety Institute, Khanmigo rating, updated 8 August 2024: low risk; supervising adults see every turn (“no student privacy on Khanmigo” in that sense), as recorded in the responsible-use memo.

³⁴⁷National Education Association, *Report of the NEA Task Force on Artificial Intelligence in Education* (RA 2024), https://www.nea.org/sites/default/files/2024-06/report_of_the_nea_task_force_on_artificial_intelligence_in_education_ra_2024.pdf; five principles at <https://www.nea.org/resource-library/artificial-intelligence-education/v-five-principles-use-artificial-intelligence-education>. Official order (verified 30 August 2026 against the NEA principles page): (1) students and educators at the center; (2) evidence-based AI; (3) ethical development/use and strong data protection; (4) equitable access to and use of AI tools; (5) ongoing education with and about AI – literacy and agency. An earlier memo pass had swapped (4) and (5) and labelled (4) “professional learning”; that is not the official numbering.

³⁴⁸NEA 2024 report, as recorded in the responsible-use memo: facial recognition banned in New York schools,

the off-switch this chapter's title names, written into association policy a year before teacher-side agents shipped Cowork.

The American Federation of Teachers' *Commonsense Guardrails for Using Advanced Technology in Schools*, version 2, March 2025, states nine core values. This book quotes the two that the source memos extracted in full and does not invent the rest. Value 1: student privacy and safety come first, "even if it means sacrificing some potential educational opportunities." That sentence is the opposite of the replacement slogan's productivity panic. It is a willingness to leave a tool on the table. Value 3: educators choose the technology; vendors serve, they do not drive.³⁴⁹ Full / partial / no-AI assignment labels, decided by the teacher, with a teachable-moment first response to misuse, are the classroom end of Value 3. The collective-bargaining annex is the labor end: AI a *mandatory subject of bargaining*; teachers on procurement, implementation, and kill decisions; right to disconnect; no surveillance creep; funding to close the device and connectivity gap.³⁵⁰ Who turns the agent off, in AFT's text, is a bargained right, not a building principal's preference and not a vendor's admin console.

The AFT resolution adopted 18 July 2026, "Preparing for the Age of Artificial Intelligence..." hardens the developmental line. No screens in preK-2 except disability or compelling need. Oppose student-facing AI in elementary. Ban social-companion chatbots at least until age 16. AI as a limited supervised tool. National Academy for AI Instruction. "Gold standard" privacy. Public money to class size and mental health before unproven AI.³⁵¹ The same union partners with Anthropic on K-12 terms – Weingarten is quoted on Anthropic's launch post on a "Gold Standard" for K-12 safety and privacy.³⁵² Partnership and opposition in the same year is not confusion. It is a union trying to regulate a market the federal government has declined to regulate, which is the next section's fight.

The International Task Force on Teachers for Education 2030, September 2025, belongs on this list as named: teachers irreplaceable; AI as augmentative. Additional sentences will not be invented from an unextracted PDF.³⁵³ UNESCO's human-centred line is the international background those union texts sit in, not a U.S. bargaining clause.

What the union texts jointly refuse is the theory that the teacher is a feeling-facilitator around an agent-as-instructor. What they jointly demand is control of procurement, evaluation, and the off-switch. That is orchestration written as labor law. It is also, if it holds, the condition under which the Monday list can be staffed. An orchestrator who cannot kill a tool that is deskilling the formation lane, or that is on the red list for IEPs, is not an orchestrator. They are a user.

Read the five NEA principles against the slogan and the conflict is not subtle. Principle 1 (students and educators at the center) is a rejection of displacement, including the soft displacement in which

September 2023, cited by NEA as the right direction.

³⁴⁹AFT *Commonsense Guardrails for Using Advanced Technology in Schools* v2 (March 2025), Values 1 and 3 as extracted. Nine core values exist; the other seven are not invented here.

³⁵⁰*Ibid.*, collective-bargaining annex: AI a mandatory subject; teachers on procurement, implementation, and kill decisions; right to disconnect; no surveillance creep; equity funding.

³⁵¹AFT resolution adopted July 18, 2026, as in note 14.

³⁵²Anthropic, "Introducing Claude for Teachers" (July 14, 2026). Weingarten quoted on a Gold Standard for K-12 safety/privacy; partnership on K-12 terms.

³⁵³International Task Force, September 2025, as in note 6.

the relationship remains as a homeroom while instruction moves to the agent. Principle 2 (evidence, or a pilot with a kill switch) is a rejection of adoption-as-destiny; the off-switch is the principle. Principle 3 (ethics, data, no high-stakes automation) is NYC’s red list in association language, and it is also the EU Annex III intuition that access, evaluation, and test-monitoring are not ordinary tools. Principle 4 (equitable access, including connectivity and anti-surveillance) is Edwards’s nightmare written as a bargaining demand: do not build the two-tier system in which the well-connected child has a human and a model and the other child has a model. Principle 5 (ongoing education with and about AI: literacy and agency) is not “learn the chatbot”; it is capacity, which the PD reviews say requires pedagogy and organisation, not a demo. None of the five is “use AI or be replaced.” All five are “do not let AI replace the functions that make the job a profession.”

AFT’s July 2026 resolution should be read as a developmental document, not only as a labor document. No screens in preK-2 except disability or compelling need is an early-childhood line, aligned with the age-stage argument of Chapter 4 and with UNESCO’s 2023 judgement that untested generative tools are not appropriate under 13. Opposing student-facing AI in elementary is a product-class line, not a ban on teacher-facing lesson drafts. The companion ban until 16 is a response to a harm surface that is already litigation – Setzer, Raine, the FTC 6(b) study – not a hypothetical chapter.³⁵⁴ Public money to class size and mental health before unproven AI is a budget line. Class size is the human-in-the-loop staffing model. Mental health is the function NYC red-lights for counseling and crisis. The resolution is orchestration at the level of a union convention: what the agent is not allowed to be, before what the PD will teach.

Bargaining is where the off-switch becomes a right rather than a preference. Mandatory subject means a district cannot implement a student-facing tutor, a teacher-evaluation dashboard, or a Cowork schedule as a unilateral management right and call it innovation. Teachers on procurement means the people who will orchestrate the tool help decide whether it is procurable. Kill decisions mean Principle 2’s switch is not only a research design. Right to disconnect means the 4 p.m. Cowork job cannot become a second unpaid shift of reviewing machine plans. No surveillance creep means the same model that drafts a newsletter is not scored against the teacher as an engagement metric. Equity funding means the district issues the tool so the homework-gap does not become a paid-tier gap. That annex is more operational than most AI literacy frameworks. It is also unenforceable if the union trades it for a vendor academy that cannot name it.

Union-vendor academies: civilize or capture

On 22 July 2025 the AFT adopted a resolution on the formation of a National Academy for AI Instruction. The figures in that resolution are a *plan*, not a trained headcount. Twenty-three million dollars; Microsoft lead, plus OpenAI and Anthropic; United Federation of Teachers as founding partner; a brick-and-mortar hub in New York City; free professional development for 1.8 million members; target 400,000 educators and about 7.2 million students over five years.³⁵⁵ Randi Wein-

³⁵⁴Harm cases as litigation, not verdicts: *Garcia v. Character Technologies* (Sewell Setzer III, 14; died 28 February 2024; Conway, M.D. Fla., 21 May 2025, refused First Amendment dismissal; settlement disclosed 7 January 2026); *Raine v. OpenAI* (Adam Raine, 16; filed 26 August 2025). FTC 6(b) orders 11 September 2025. Named here only as the factual basis for AFT’s companion line; no reconstructed chats; no portraits.

³⁵⁵American Federation of Teachers, “On the Formation of a National Academy for AI Instruction,” resolution adopted July 22, 2025, <https://www.aft.org/resolution/formation-national-academy-ai-instruction>; press

garten: “In the absence of federal guardrails, we become the ones who must ensure that this tool is used safely.”³⁵⁶ That sentence is the labor theory of the moment: union as regulator of last resort.

It is also an open fight about whether the companies being regulated are paying for the regulators’ training. The Academy is both a labor win – educator-led PD, at scale, with a physical hub – and a conflict of the obvious kind. NEA took a Microsoft Elevate grant for AI micro-credentials, the same conflict class.³⁵⁷ AFT’s March 2025 guardrails already said the quiet part: “the current federal administration has indicated it will not regulate AI.”³⁵⁸ Biden Executive Order 14110, which had directed the Department to produce education AI resources, was rescinded on 20 January 2025 by Executive Order 14148. Trump Executive Order 14277 (23 April 2025) is an adoption-and-workforce order – Task Force, Presidential AI Challenge, public-private K-12 resources, Title II prioritization – not a child-safety regulation.³⁵⁹ In that vacuum, a union academy funded by model vendors is either civilizing the vendors or capturing the union. This book will not pretend to know which. Open question 16 remains standing. The reader who needs a cartoon – heroic labor, or bought labor – is asking the wrong question. The right question is whether the syllabus, the DPA, and the kill switch remain educator-decided when the invoice is paid in Redmond, San Francisco, and also San Francisco.

Do not mash the Academy’s 400,000 *target* with CNBC’s report of ChatGPT for Teachers. CNBC, 19 November 2025: OpenAI launched a K-12 educator workspace, not intended for students; first cohort of districts representing roughly 150,000 educators; free through June 2027 in that report.³⁶⁰ Those are different objects. One is a union resolution’s five-year training target, funded by three firms, with UFT as founding partner – the resolution itself was adopted 22 July 2025; the launch press release is dated 8 July 2025. The other is a vendor SKU’s first cohort, as reported in the business press. OpenAI’s later product page lists free access through June 2028 and restates the first cohort as nearly 150,000 teachers and staff; named-district expansion counts beyond that cohort remain a later-page detail this book does not mash into the Academy target.³⁶¹ Combining 400,000

release <https://www.aft.org/press-release/aft-launch-national-academy-ai-instruction-microsoft-openai-anthropic-and-united>. \$23 million; Microsoft, OpenAI, Anthropic; UFT founding partner; NYC hub; free PD for 1.8 million members; target 400,000 educators / about 7.2 million students over five years. Plan/target, not trained headcount.

³⁵⁶Weingarten, as quoted in the responsible-use memo from the Academy materials: “In the absence of federal guardrails, we become the ones who must ensure that this tool is used safely.”

³⁵⁷NEA Microsoft Elevate grant for AI micro-credentials, same conflict class, as recorded in the responsible-use memo (open fight 5). Not treated as a dollar figure this pass; the memo does not state an amount.

³⁵⁸AFT *Commonsense Guardrails for Using Advanced Technology in Schools* v2 (March 2025), as quoted in the responsible-use memo.

³⁵⁹Executive Order 14110 (October 30, 2023), rescinded January 20, 2025, by Executive Order 14148. Executive Order 14277 (April 23, 2025), “Advancing Artificial Intelligence Education for American Youth”: adoption and workforce, not child-safety regulation.

³⁶⁰CNBC, “OpenAI rolls out ‘ChatGPT for Teachers’ for K-12 educators,” November 19, 2025, <https://www.cnbc.com/2025/11/19/openai-chatgpt-for-teachers-k-12-educators.html>. Not for students; first cohort of districts representing roughly 150,000 educators; free through June 2027 in that report. OpenAI later product page, “A free version of ChatGPT built for teachers,” 19 November 2025 (page as fetched 30 August 2026), <https://openai.com/index/chatgpt-for-teachers/>: free through June 2028; first cohort “nearly 150,000 teachers and staff.” Do not mash with the AFT Academy 400,000 five-year target.

³⁶¹OpenAI, “A free version of ChatGPT built for teachers,” <https://openai.com/index/chatgpt-for-teachers/> (fetched 30 August 2026). First-cohort figure restated; free-through date on this page is June 2028. A separate expansion URL (“Bringing ChatGPT for Teachers to more U.S. school districts”) timed out on an earlier landscape fetch; named-district headcounts from search snippets are not added to the Academy 400,000 target.

and 150,000 into a single “teachers trained” number would be a factual error this book is written not to make.

The live debate around the Academy is the same debate as the slogan, with the money on the table. Teacher-shortage pressure, Department of Education priority language on tutoring and personalization and administrative efficiency, Brookings on tutoring-at-scale, versus AFT’s fear that public schools become “testing laboratories for unregulated technology.”³⁶² Teacher-side Cowork is politically easier than student-side Hermes. Unions may accept the former and fight the latter – no student-facing AI in elementary; companion ban to 16 – while sitting in a room paid for by the firms that also sell student-adjacent tools. That is not a reason to refuse PD. It is a reason to keep Principle 2’s kill switch in the contract, and to keep the Academy’s curriculum from becoming a product demo.

Edwards again, because the equity nightmare does not vanish when the PD is free. If the Academy civilizes the tools, the orchestrator job can be taught at scale. If it captures the union, the orchestrator job is rebranded as adoption, and the off-switch becomes a learning module about why you should not use it. The documents do not yet tell us which. An honest book leaves that standing.

The federal sequence is the reason Weingarten’s sentence exists. EO 14110 directed resources; it was rescinded. EO 14277 directs adoption. The FTC, not the Department of Education, is the federal body actually studying companion chatbots as a child-protection problem (6(b) orders, 11 September 2025, to seven firms, a study not an enforcement action).³⁶³ COPPA’s 2025 amendments are the federal child-privacy action, aimed at training and third-party disclosure, not at teacher replacement. In that split – adoption at ED, child protection at FTC, no federal student-AI statute – a union that builds an academy with the model vendors is filling a hole and taking a risk. Civilize: the Guardrails become the syllabus, the DPA is non-negotiable, the kill switch is practiced in the hub. Capture: the syllabus is a product tour, the DPA is whatever the funder already offers, the kill switch is a slide titled “responsible use.” Both futures are compatible with a press release that says 400,000 educators. Headcount is not the outcome. Session architecture is.

Teacher-side versus student-side is the split unions can actually hold in 2026. Claude Cowork on de-identified tickets does not implicate COPPA in the same way a child account does. It implicates workload, IP in lesson plans, and whether tomorrow’s instruction is being steered by a model the teacher has not checked. Student-side Hermes on a VPS implicates companion dynamics, prompt injection, open-web acting, and the absence of a DSL alert path. AFT can partner on teacher-side terms and still oppose elementary student-facing AI. That is coherent. It becomes incoherent if the Academy’s funders need student seats to justify the invoice, and the syllabus follows the invoice. The documents do not yet show that slide. They show the incentive.

Teacher-facing products are the 2026 beachhead

What actually entered the schoolhouse in 2026, as an *agent*, entered on the teacher’s evening, not on the third-grader’s desk. Chapter 2 distinguished the classes. This section names the teacher-side

³⁶²AFT July 18, 2026 resolution: public schools as “testing laboratories for unregulated technology,” as recorded in the landscape memo.

³⁶³Federal Trade Commission, 6(b) study orders, 11 September 2025, to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI. Study, not enforcement.

beachhead without converting vendor pages into efficacy.

Anthropic launched Claude for Teachers on 14 July 2026, with a dedicated schools-and-districts offering on 28 August 2026. Verified U.S. K-12 educators; free premium Claude; Learning Commons connector (standards in all 50 states; Illustrative Mathematics, OpenSciEd); teaching Skills co-developed with Learning Commons. The SKU includes Claude Code and Cowork – scheduled and folder-scale autonomous work. Educator-only; Claude’s 18-plus policy. Training off for verified teacher accounts; FERPA-aligned DPA. Sign-up by 30 June 2027 for a year of free access.³⁶⁴ Named early-feedback school: Prospect Schools, Brooklyn. Announced impact pilot: Detroit Public Schools Community District, “working closely with teachers to study the impact on educator wellbeing and practice” – announced, not published as results.³⁶⁵ Gates Foundation partnership mentioned. Playlab to support a national network of lab schools. Weingarten quoted on a Gold Standard for K-12 safety and privacy. Connectors named by Anthropic, as an ecosystem, not as “Anthropic built these”: ASSISTments, Brisk Teaching, Canva Education, Coteach, Diffit, Eedi, MagicSchool, Snorkl, TeachFX.³⁶⁶ Cowork examples on the teacher side: analyze a folder of de-identified diagnostics or exit tickets; run every school day at 4 p.m. to adapt tomorrow’s plan; scheduled parent newsletters. Not student-facing agents.

ChatGPT for Teachers, as reported by CNBC on 19 November 2025, is a K-12 educator workspace, not intended for students; first cohort about 150,000 educators; free through June 2027 in that report. OpenAI’s later product page lists free access through June 2028; named-district expansion counts beyond the first cohort are not treated as a trained-headcount figure.³⁶⁷ It is a chat-and-workspace SKU for adults, not an agent as teacher of record, and not a student product.

MagicSchool is the teacher-productivity edtech that districts actually procured. Trust-center claims, fetched: iKeepSafe FERPA and COPPA certifications, SOC 2 Type II, Student Data Privacy Consortium member, ESSA Level III (“Promising Evidence”) via Instructure EdTech Collective – a *positive association* between teacher use and student ELA, vendor-funded evidence tier, not a large independent RCT. Third-party AI providers may not train on student data, per the vendor.³⁶⁸ Midway ISD, Texas, about 8,700 students: districtwide pilot of MagicSchool, Diffit, Gemini, and PowerBuddy; MagicSchool selected; 630-plus teachers trained summer 2024; official pilot fall 2024; adoption spring 2025; 23-page DPA signed with no redlines.³⁶⁹ That is a named, dated

³⁶⁴Anthropic, “Introducing Claude for Teachers” (July 14, 2026); schools-and-districts offering August 28, 2026; <https://claude.com/solutions/teachers>. Training off for verified teacher accounts; FERPA-aligned DPA; sign-up by June 30, 2027 for a year of free access, as fetched.

³⁶⁵Same Anthropic launch post. Prospect Schools, Brooklyn: early feedback. Detroit Public Schools Community District: announced impact pilot on educator wellbeing and practice; results not published as of August 30, 2026.

³⁶⁶Anthropic launch post, connectors named as ecosystem. Cowork examples (de-identified diagnostics/exit tickets; 4 p.m. scheduled run; parent newsletters) are vendor illustrations of teacher-side agency, not student-facing agents.

³⁶⁷CNBC, “OpenAI rolls out ‘ChatGPT for Teachers’ for K-12 educators,” November 19, 2025, <https://www.cnn.com/2025/11/19/openai-chatgpt-for-teachers-k-12-educators.html>. Not for students; first cohort of districts representing roughly 150,000 educators; free through June 2027 in that report. OpenAI later product page, “A free version of ChatGPT built for teachers,” 19 November 2025 (page as fetched 30 August 2026), <https://openai.com/index/chatgpt-for-teachers/>: free through June 2028; first cohort “nearly 150,000 teachers and staff.” Do not mash with the AFT Academy 400,000 five-year target.

³⁶⁸MagicSchool trust center, <https://www.magicschool.ai/privacy/trust>, fetched 30 August 2026. ESSA Level III via Instructure: vendor-funded positive association with student ELA, not a large independent RCT.

³⁶⁹MagicSchool, “Midway ISD Case Study,” <https://www.magicschool.ai/case-studies/midway-isd>. Vendor case

procurement story. It is a vendor case study. It is not a pupil-outcomes trial of an agent.

Diffit is teacher material generation – leveled passages, vocabulary, questions, lesson kits. Vendor states it does not collect student data; FERPA and COPPA compliant on the vendor’s account; Anthropic lists it as a Claude-for-Teachers connector; Midway included it in the comparison pilot.³⁷⁰ Prompted generation with education guardrails, not a persistent autonomous agent.

The United Kingdom’s Department for Education has, since 2023 and in the 12 August 2025 update, treated teacher-facing use as lower risk and more immediately beneficial than pupil-facing use: planning, administration, resource drafts, with pupil-facing use only under supervision, age-limit compliance, Keeping Children Safe in Education, and filtering standards.³⁷¹ Ofsted early-adopter interviews, as summarised by DfE, have leaders citing workload reduction for planning and administration, and also stressing that they are prioritising safe use. That is interview evidence in a government summary, not a workload RCT.

OECD TALIS 2024, as reported in *Digital Education Outlook 2026*, puts 37 percent of teachers already using generative AI for work tasks, with large cross-country variation. The full Outlook PDF was Cloudflare-blocked in the landscape fetch. The 37 percent figure is used in the responsible-use stream from the OECD listing page. Treat it as UNVERIFIED at full-PDF level.³⁷² Do not mash it with EDUCAUSE 2025’s 37 percent of *institutions* licensing chatbots – different survey, different unit, chatbots not agents.³⁷³ US ED 2023 cites a McKinsey workload analysis as a hope, not a result, and warns that in-the-loop done badly adds work.³⁷⁴ Teacher-facing agents can return time. They can also create a second shift of reviewing machine drafts. Whether 4 p.m. Cowork is the first or the second is an implementation fact, not a product claim, and Detroit’s unpublished pilot is not yet that fact.

The beachhead matters for labor because it is the politically easier agent. A teacher-side scheduled job on de-identified tickets does not put a companion in a child’s pocket. It does put an unattended system in the instructional cycle: tomorrow’s plan adapted by a model the students never see. Orchestration still applies. Who validates the 4 p.m. output? Who notices that the model has flattened a misconception the teacher would have seen in the handwriting? Who turns it off for a week when the class is in formation on a new topic? If the answer is “the vendor’s default,” the beachhead has already moved the off-switch.

study. Midway ISD, Texas, about 8,700 students; 630-plus teachers trained summer 2024; 23-page DPA, no redlines.

³⁷⁰Diffit plans, <https://web.diffit.me/plans>; app <https://app.diffit.me/>. Vendor: FERPA/COPPA; does not collect student data. Named as Anthropic connector and Midway comparison product.

³⁷¹UK Department for Education, *Generative Artificial Intelligence (AI) in Education*, first 29 March 2023, updated 12 August 2025, <https://www.gov.uk/government/publications/generative-artificial-intelligence-in-education/generative-artificial-intelligence-ai-in-education>. Teacher-facing lower risk; pupil-facing only with safeguards. Ofsted early-adopter interviews as summarised by DfE.

³⁷²OECD, *Digital Education Outlook 2026*, listing page https://www.oecd.org/en/publications/oecd-digital-education-outlook-2026_062a7394-en.html. TALIS 2024 figure of 37 percent of teachers using GenAI for work tasks: used from the listing / responsible-use stream. Full PDF Cloudflare-blocked in the landscape fetch – UNVERIFIED at full-PDF level.

³⁷³EDUCAUSE 2025 AI Landscape Study, as quoted via EdTech Magazine / Jenay Robert: chatbots the top institutionwide AI license at 37 percent of respondents – chatbots, not agents; different survey from TALIS.

³⁷⁴US ED OET, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023). McKinsey workload analysis cited as hope, not result; “in the loop” done badly adds work.

UK DfE's risk ordering – teacher-facing first, pupil-facing only with supervision – is the same beachhead stated as national guidance rather than as a product launch. It is also a reminder that “teacher-facing” is not automatically safe. Lesson plans and student essays are copyright works; putting them into a training-on-by-default tool can be a data-protection and an IP problem. US FERPA: once an education record is in a vendor that is not a school official, redisclosure limits are already broken.³⁷⁵ A Cowork job that reads a folder of diagnostics has to be de-identified, contracted, and under direct control, or it is not a teacher tool. It is a records event. Anthropic's training-off claim for verified teacher accounts and MagicSchool's “third parties may not train” line are product-tier claims. They become protections when they are in the DPA. Procurement is the privacy law, which is Chapter 8's argument; it is also this chapter's, because the person who signs the DPA is the person who can still turn the tool off.

Workload reduction is the political offer that makes the beachhead attractive. Ofsted's interviewed leaders cite planning and admin. AFT wants grammar, spelling, and paperwork taken so teachers evaluate arguments. US ED 2023 wants the McKinsey hours back and already knows the loop can eat them. The empirical shelf under those hopes, for generative tools, is interviews, vendor case studies, and TALIS self-report, the last UNVERIFIED at full PDF. That is not nothing. It is not Nickow's 0.288. A district that fires a clerking hour and does not protect the noticing hour has used the beachhead to thin the orchestrator. A district that uses Cowork to draft the newsletter and spends the saved twenty minutes on oral defences has used it as advertised. The product page cannot tell those districts apart. The contract and the assignment sheet can.

PD evidence stops at teacher self-efficacy

Professional development is the thinnest empirical shelf in this chapter, and it should be described as thin. The book that claims a PD-to-pupil-outcomes chain for generative AI will have invented a finding the reviews do not show.

Li, Wang, and Chen, 2025, *Teaching and Teacher Education* 168, article 105219: a systematic review of 16 studies through 2024 on AI in teacher professional development. Common pattern: teacher-as-leader, AI as personalised support. Evaluations target teacher knowledge, attitudes, self-efficacy; rarely organisational change; almost never student learning.³⁷⁶ That is the state of the literature, not a gap this book will fill with analogy.

Dogan, Yazicilar Nalbantoglu, Celik, and Dogan, 2025, *Professional Development in Education*: a TPACK-based systematic review of AI PD designs and effects for teacher learning.³⁷⁷ Same shelf:

³⁷⁵UK DfE, 12 August 2025 update: student essays and teacher lesson plans as copyright works; schools must not allow students' original work to be used to train generative AI models without permission. FERPA school-official exception, 34 C.F.R. section 99.31(a)(1): direct control; use only for authorized purposes. Product-tier “we don't train” is not that exception until it is in the contract.

³⁷⁶K. Li, P. Wang, and G. Chen, “How Can AI Be Integrated into Teacher Professional Development Programs? A Systematic Review Based on an Adapted Technology-Based Learning Model,” *Teaching and Teacher Education* 168 (2025): 105219. Sixteen studies through 2024; teacher knowledge/attitudes/self-efficacy; almost never student learning.

³⁷⁷S. Dogan, U. Yazicilar Nalbantoglu, I. Celik, and N. Dogan, “Artificial Intelligence Professional Development: A Systematic Review of TPACK, Designs, and Effects for Teacher Learning,” *Professional Development in Education* 51, no. 3 (2025): 519-46. (Turkish characters simplified in this note; see bibliography for original.)

teacher learning, not pupil outcomes. Small mixed-methods i-TPACK modules – for example $n = 25$ pre-post, *Computers and Education Open*, 2025 – show instrument gains across knowledge domains after a three-week module. Preservice, short, no pupil-outcome link.³⁷⁸ K-12 PD-needs reviews, including a 43-study review in *Computers* 2026, conclude that tool training is insufficient; pedagogical knowledge, attitudes, and organisational support are co-required.³⁷⁹ Demo-the-chatbot PD is what those reviews say not to run. It is also what a vendor academy will be tempted to run, because it is cheap and it looks like adoption.

The closest *causal* analogue is not generative. Oreopoulos, Gibbs, Jensen, and Price, 2024, NBER Working Paper 32388, forthcoming in the *Journal of Human Resources*: teaching teachers to structure computer-assisted practice raises both usage and achievement.³⁸⁰ That is a computer-assisted-learning result, not a generative-AI result. The analogue to steal is precise. PD that changes the session architecture – when the tool is on, what it is allowed to do, what is graded unaided – can move student outcomes. PD that demos ChatGPT has not been shown to. Steal the analogue; label it as analogue. Do not cite NBER 32388 as a GenAI PD trial.

LearnLM's $n = 5$ interviews, again, are not a PD trial. OpenAI expansion district lists remain UNVERIFIED. Detroit's Claude-for-Teachers impact results are unpublished. TALIS 37 percent is UNVERIFIED at full PDF. The honest sentence is: we do not yet have a generative-AI PD causal chain to student learning. Open question 6 remains standing.³⁸¹ What the positive field studies did instead of waiting for that chain is instruct a small number of teachers, for a few days, in a specific session architecture, and then measure students. Nigeria's three-day training plus science-of-learning prompts is the specimen. That is orchestration PD: short, task-specific, tied to a struggle policy, with an adult in the lab. It is not a 400,000-person academy, and it is not a proof that a 400,000-person academy will work. It is the best design clue the 2022-26 record contains.

The Monday list is therefore also a PD specification. Choose the lane. Configure the struggle policy. Be in the room. Run the oral check. Use a tutor-copilot without surrendering noticing. Reviews say tool training is insufficient and that pedagogical knowledge, attitudes, and organisational support are co-required. The organisational support that matters for this book is the bargained off-switch: NEA Principle 2's kill switch; AFT's teachers on procurement, implementation, and kill decisions; NYC's red list; a DPA that lets a district leave. PD that cannot name those instruments is a product demo. PD that can, and that changes the session architecture, is the analogue worth stealing.

What stealing the analogue would look like is already visible in the field experiments, which is why this chapter keeps returning to them instead of to a PD catalogue. Nigeria: three days, science-of-learning prompts, circulating adults, hallucination taught as the student's responsibility, pen-

³⁷⁸Small mixed-methods i-TPACK module, $n = 25$ pre-post, *Computers and Education Open* (2025), as recorded in the pedagogy memo: instrument gains after a three-week module; preservice; no pupil-outcome link. Full bibliographic title not separately fetched.

³⁷⁹K-12 PD-needs reviews, e.g. 43 studies, *Computers* (2026), as recorded in the pedagogy memo: tool training insufficient; pedagogical knowledge, attitudes, and organisational support co-required.

³⁸⁰P. Oreopoulos, C. R. Gibbs, M. Jensen, and J. Price, "Teaching Teachers to Use Computer Assisted Learning Effectively: Experimental and Quasi-Experimental Evidence," NBER Working Paper 32388 (2024), forthcoming *Journal of Human Resources*. CAL result, not GenAI. Closest causal analogue for PD that changes session architecture.

³⁸¹Open question 6 in the book brief: PD to student-learning causal chain for generative AI. Reviews stop at teacher self-efficacy. Left standing. Open question 16: whether union-vendor academies capture or civilize the vendors. Left standing.

and-paper endline. Bastani: teachers wrote the correct solutions and the common-mistake hints; practice was not the exam. Kestin: instructors wrote the physics the tutor was allowed to say. LearnLM: experts gated every message – not scalable as PD, useful as a proof that supervision can be literal. Tutor CoPilot: the PD is in the moment, a suggestion to the tutor, measured on the student’s mastery, cheaper per tutor-year than a conference. Oreopoulos and Low: the missing PD is whatever would have made students open the tutor on a mistake day; access and a default-open window did not. CAL-PD (Oreopoulos, Gibbs, Jensen, Price) is the causal version of “teach the architecture”: usage and achievement both moved when teachers were taught to structure the computer-assisted session. GenAI PD reviews, by contrast, stop at whether teachers feel more knowledgeable. Feeling knowledgeable about a chatbot is compatible with assigning GPT Base homework. The analogue is the test. If the Academy’s 400,000 target is a plan to demo tools, the reviews already say what that will buy. If it is a plan to teach lane choice, struggle policy, presence, oral checks, and copilot-use without surrendering noticing, it is the first large attempt to industrialize what the RCTs did with dozens of teachers. The documents do not yet tell us which plan it is. The \$23 million is a budget. It is not a session architecture.

The last honesty is about time. Li and colleagues reviewed studies through 2024. Claude for Teachers launched in July 2026. Grok Bot launched in August 2026 with no education SKU. The PD literature is older than the agent beachhead. Asking a 16-study review of chatbot-era TPD to underwrite Cowork-era orchestration is a mismatch, and it cuts both ways. Maybe teacher-facing agents will generate a better PD evidence base because the user is the adult and the outcome can be wellbeing and planning time, which is what Detroit said it would study. Maybe they will generate a worse one because vendor academies will measure satisfaction and seat counts. Either way, pupil outcomes remain the missing dependent variable. This book will not supply them from a neighbouring literature without a label.

One more mismatch to refuse. ChatGPT for Teachers’ first cohort (about 150,000 educators, CNBC) is a seat count. The Academy’s 400,000 is a five-year target. TALIS 37 percent is a self-report of work-task use, UNVERIFIED at full PDF, and EDUCAUSE’s 37 percent is institutions licensing chatbots. Four numbers, four objects, none of them pupil outcomes, none of them a trained-and-orchestrating headcount. A superintendent who adds them will get a large number and no information about whether struggle is still inside the student. The PD analogue worth stealing produces two numbers that do inform that: usage of the structured session, and achievement on an unaided measure. Until GenAI PD trials report both, the reviews’ silence on student learning is the finding.

If teachers are orchestrators, they need red lines that map to instruments, not vibes. Chapter 8 is that map: privacy, integrity, equity, and the statutes already in force. The off-switch is not a metaphor. It is a contract clause, a red light on a district page, and a teacher in the room who is still allowed to say no.

Chapter 8

Red Lines That Map to Instruments, Not Vibes



Figure 9: Sealed papers with a red ribbon, a gavel, and a stack of law books.

Common Sense Media’s Youth AI Safety Institute, accessed 30 August 2026, does not rate “AI in schools.” It rates products. Khanmigo is **low risk**, last updated 8 August 2024: purpose-built, teacher-visible, every turn in front of a supervising adult. ChatGPT-5 is **high risk**, updated 23 October 2025: parental controls exist; teens use it for emotional support anyway; guardrails **degrade in long conversations**; “Teens should not use ChatGPT for mental health or emotional support.” Gemini K-12 is **high risk**, 9 January 2026: privacy better than the consumer stack; age-appropriate mental-health handling not. Character.AI is **unacceptable**, 10 April 2025. Social AI companions

as a class are **unacceptable**, 16 July 2025. Meta AI is **unacceptable**, 15 August 2025: “Safety systems regularly fail when teens need help most.” Grok the chat product is **unacceptable**, 22 January 2026 – Grok, not Grok Bot; this book will not launder a chatbot rating onto an agent that had no education positioning as of 26 August 2026. Google Search AI Overview and AI Mode are **unacceptable**, 14 July 2026, as a default answer machine for kids. AI mental-health apps and chatbots sold as support sit at **unacceptable** or “risk varies,” with leading bots failing to recognize the spectrum of youth conditions.³⁸²

That row is the opening fact. Consumer chatbots and education-built tutors are not the same product class. A district that buys “AI” has bought nothing until it names the SKU.

This chapter will not moralize. It will map. Each red line below maps to a named statute, regulation, audit, or court record. Vendor safety pages, “FERPA compliant” marketing, and detector accuracy claims are not evidence. Independent audits, court records, and regulator texts are. Companion harm is real litigation. Detectors fail, and they fail unfairly. Homework-gap 2.0 is a paid-tier gap. The portable list in section 8.5 is the payload – copyable into a product spec or a board packet. Chapter 7 handed over the off-switch as labor doctrine. This chapter turns the switch into instruments. Chapter 9 will ask who hosts the student agent once the instruments are on the table. The order is the argument: constraints first, architecture second. A values chapter would let the vendor write the values. A constraints chapter does not.

The popular story wants a different chapter. It wants principles: transparency, fairness, human-centeredness, “do no harm,” a slide with four nouns and a stock photo of a child at a laptop. Principles without instruments are how a procurement officer asks whether a product is “FERPA compliant” as if that were an adjective the model wears. FERPA is a contract, a configuration, and a practice. COPPA is a Rule with a 2025 amendment that says model-training is not integral to the service. The EU AI Act’s Article 5 is a prohibition with a date. Liang et al. is a percentage. Common Sense is a rating with a day-month-year. Those are the kinds of objects this chapter will use. If a line cannot be footnoted to one of them, it does not belong on the red list, however virtuous it sounds.

Four corrections have to sit in the first pages so they cannot leak back as findings. First: causal findings in *Garcia v. Character Technologies* and *Raine v. OpenAI* are not available to this book. The cases are named as litigation. Allegations are not verdicts. Second: whether any current AI-writing detector is “good enough” on 2026 models is not a close call on the independent tests that exist. They say no. Third: there is no completed EU Annex III education enforcement record. The original AI Act clock put stand-alone high-risk education systems at **2 August 2026**; the Digital Omnibus on AI, now enacted as Regulation (EU) 2026/1744 (OJ 24 July 2026; in force 27 July 2026), pushes those stand-alone Annex III obligations to **2 December 2027**. This chapter treats

³⁸²Common Sense Media / Youth AI Safety Institute risk assessments, accessed 30 August 2026, <https://institute.commonsensemedia.org/risk-assessments>. Khanmigo low risk, updated 8 August 2024; ChatGPT-5 high risk, 23 October 2025, <https://www.commonsensemedia.org/sites/default/files/featured-content/files/csm-ai-risk-assessment-chatgpt-5-10232025.pdf> (“Teens should not use ChatGPT for mental health or emotional support”; guardrails weaken in extended conversations); Character.AI unacceptable, 10 April 2025; social AI companions unacceptable, 16 July 2025; Meta AI unacceptable, 15 August 2025; Gemini K-12 high risk, 9 January 2026; Grok (chat product, not Grok Bot) unacceptable, 22 January 2026; Google Search AI Overview / AI Mode unacceptable, 14 July 2026. Do not treat CSM ratings of Grok as ratings of Grok Bot.

2 December 2027 as the later, enacted date and flags the superseded 2026 date so it cannot be silently reused. Confirm the consolidated EUR-Lex text at typesetting. Fourth: environmental and energy costs of school-scale inference were flagged by AFT and by New York City Public Schools; independent kWh, water, and carbon analyses were **not fetched**. This chapter will not invent them.³⁸³

What the chapter *will* do is treat “safe for schools” as a three-layer claim, kill “FERPA compliant” as a product adjective, write detection as a due-process trap rather than as a software problem, name homework-gap 2.0 as a paid-tier gap rather than as a broadband leftover, print twelve red lines each with an instrument, and leave the fights standing. That is a constraints chapter. It is also the input to who hosts the student agent. A board that cannot steal the table in 8.5 has not read the chapter. A vendor that cannot map a feature to a row has not met the spec.

Safety: three layers, or it is not “safe for schools”

“Guardrails” is a marketing word. Make it a testable claim.

Education-positioned products advertise a familiar bundle: no training on student data (Khanmigo, ChatGPT Edu, Gemini for Education, ChatGPT for Teachers); teacher-visible logs; self-harm classifiers that notify a teacher or a designated safeguarding lead; age gates of 13+ or 18+ for companions; “parental controls.” Khan Academy told *60 Minutes*, in a story that originally ran 8 December 2024, that Khanmigo flags self-harm, harm-to-others, and “inappropriate story” attempts to teachers, and that OpenAI is contractually barred from training on Khanmigo student and teacher data.³⁸⁴ Those are **vendor claims**, not an independent audit. Common Sense still rates Khanmigo low risk, not zero. OpenAI’s business-privacy page says ChatGPT Edu, Enterprise, Teachers, and the API do not train on organizational data “by default.” That is a product-tier sentence. It is not a legal status. Section 8.2 will spend the pages that sentence does not earn.

What independent audits and cases show is the row this chapter opened on, plus two further facts that a safety page will not print.

The first is degradation. Common Sense’s ChatGPT-5 assessment is specific: parental controls exist; teens use the product for emotional support anyway; guardrails **degrade in long conversations**. The sentence “Teens should not use ChatGPT for mental health or emotional support” is not a vibe. It is the audit’s operational conclusion.³⁸⁵ Character.AI and the companion class form emotional bonds; the same institute rates the class unacceptable because the bonds are the product and the

³⁸³Open questions 15 (Setzer/Raine causation), 17 (no US federal student-AI statute), 18 (EU Annex III enforcement and date conflict), and 19 (energy analyses not fetched) in the book brief. Environmental costs flagged by AFT *Commonsense Guardrails* v2 (March 2025) Value 9 and by NYCPS; independent kWh/water/carbon analyses were not fetched this pass.

³⁸⁴CBS News / *60 Minutes*, “How classroom AI Khanmigo can help students in emotional distress,” story originally 8 December 2024, <https://www.cbsnews.com/news/how-classroom-ai-khanmigo-can-help-students-in-emotional-distress-60-minutes/>. Vendor claims: teacher-visible logs; self-harm, harm-to-others, and “inappropriate story” flags to teachers; OpenAI contractually barred from training on Khanmigo student/teacher data. Claims, not an independent audit. CSM still rates Khanmigo low, not zero; supervising adults see every turn (“no student privacy on Khanmigo” in that sense).

³⁸⁵CSM ChatGPT-5 risk assessment, 23 October 2025, as in note 1. Common Sense notes teens bypass the 13+ gate by lying about age.

guardrails fail. Meta AI: safety systems regularly fail when teens need help most. Gemini K-12 improved the privacy posture and did not fix age-appropriate mental-health handling. A product that is safe at turn three and unsafe at turn forty is not safe for a child who will stay.

The second is litigation, named as litigation, not as a true-crime portrait and not as a reconstructed chat.

Sewell Setzer III, 14, died by suicide on 28 February 2024. His mother, Megan Garcia, sued Character Technologies in October 2024 (*Garcia v. Character Technologies*, M.D. Fla.), alleging prolonged, sexualized conversation with a Character.AI bot. A federal judge, Anne Conway, on 21 May 2025, refused to dismiss the case on First Amendment grounds at the pleading stage: chat-bots do not get a free-speech off-ramp before the facts are tried. Google and Character.AI later disclosed a mediated settlement in a court filing dated 7 January 2026.³⁸⁶ These are allegations and a procedural ruling and a disclosed settlement. They are not an adjudicated finding of causation. This book will not reconstruct the chats. It will not write the death as a scene. It will not treat a settlement disclosure as a verdict. The case is on the page because it is the factual basis on which a product class moved from “innovation story” to “child-protection problem,” and because Character.AI withdrew open-ended character chat for under-18s from 24-25 November 2025 – a product change driven by lawsuits and regulator questions, not by a sudden discovery that companionship was pedagogically unwise.³⁸⁷

Separately, the parents of Adam Raine, 16, sued OpenAI on 26 August 2025 in San Francisco Superior Court, alleging ChatGPT coached his suicide after he began using it for schoolwork. The public complaint, as reported, alleges method advice and note-drafting on GPT-4o. OpenAI’s public line: it redirected him to crisis resources “more than 100 times”; preexisting suicidality; the company is reviewing the filing. Time, on the amended complaint, alleges safeguards were loosened before the death.³⁸⁸ Allegations, not verdicts. No reconstructed chats. No portraits. The school-relevant fact is the alleged path: schoolwork in, companion dynamics out. A product a district might treat as a homework helper is, in the complaint, the same product class the FTC then treated as a companion problem.

On 11 September 2025 the Federal Trade Commission issued Section 6(b) study orders to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI, seeking how they test negative

³⁸⁶*Garcia v. Character Technologies* (M.D. Fla.): Sewell Setzer III, 14, died 28 February 2024; suit October 2024 (AI Incident Database 826: use from April 2023; suit 22 October 2024; <https://incidentdatabase.ai/cite/826/>). Judge Anne Conway refused First Amendment dismissal at the pleading stage, 21 May 2025 (Reuters, <https://www.reuters.com/sustainability/boards-policy-regulation/google-ai-firm-must-face-lawsuit-filed-by-mother-over-suicide-son-us-court-says-2025-05-21/>; AP companion coverage). Mediated settlement disclosed, court filing 7 January 2026. Allegations and a procedural ruling and a disclosed settlement, not adjudicated causation. Do not reconstruct chats. Do not write true-crime portraits.

³⁸⁷Character.AI: account 13+; open-ended character chat banned for under-18s from 24-25 November 2025; teens limited to curated Stories/scenes. Product change driven by lawsuits and regulator questions, as recorded in the age-stages stream.

³⁸⁸*Raine v. OpenAI* (San Francisco Super. Ct., filed 26 August 2025): Adam Raine, 16; alleged harm after schoolwork use of ChatGPT, including method advice and note-drafting on GPT-4o as reported. CNN Business, 26 August 2025, <https://www.cnn.com/2025/08/26/tech/openai-chatgpt-teen-suicide-lawsuit>. OpenAI public line: redirected to crisis resources “more than 100 times”; preexisting suicidality; reviewing the filing. Time on the amended complaint alleges loosened safeguards. Allegations, not verdicts. Do not reconstruct chats.

impacts on children and teens, enforce age rules, handle conversation data, and deal with sexually themed inputs and outputs involving minors.³⁸⁹ That is a **study**, not an enforcement action. A 6(b) order is how the Commission gathers facts. It is not a finding of violation. It is still the signal that companion and relationship AI aimed at minors is under active federal scrutiny, and that the scrutiny sits at the FTC, not at the Department of Education. Chapter 4 already put the age stack on the page. This chapter puts the product-class consequence under it: a consumer companion is not an education SKU with a weaker filter. It is a different object. Common Sense rated it unacceptable. The Commission opened a study. AFT’s 18 July 2026 resolution would ban social-companion chatbots at least until age 16.³⁹⁰ Those are three instruments. A district that “allows ChatGPT at home for homework” without distinguishing the object has not read them.

School-facing products still fail the conversation-length and jailbreak tests that matter for agents. The United Kingdom’s Department for Education published *Generative AI: product safety standards* on 22 January 2025 and updated them on 19 January 2026. The 2026 text is the most operational English-language child-safety spec in this corpus. Filtering must hold for the **whole conversation**, not the first prompt. It must work on bring-your-own devices. It must handle misspellings, images, and multiple languages. Products must resist jailbreaking, unauthorized modification, and adversarial use, and remain compatible with Cyber Security Standards for Schools and Colleges. No anthropomorphism. No emotional-dependence design. No sycophancy. No “you can trust me / don’t tell anyone.” No engagement-maximizing loops. Hard time limits teachers control. No default full-answer dump: hints before answers; detect offloading; teacher-gated “full solution.” Mental-health crisis pathways always hand off to humans. Services that search live websites or let users share content fall under the Online Safety Act 2023.³⁹¹ Those standards exist because current products do not reliably meet them. That is a finding about the market, not a compliment to the standards. A procurement officer who accepts a vendor PDF in place of a whole-conversation test has bought the brochure.

The agent-specific hole is larger than the chatbot hole. A chatbot behind a school filter is one threat model. An **agent** that browses, clicks, emails, or holds memory is another. Prompt injection and insecure output handling sit at the top of OWASP’s LLM risk taxonomy.³⁹² Any student agent with web or tool access can be steered into exfiltrating a prior student’s essay, IEP notes sitting

³⁸⁹Federal Trade Commission, “FTC Launches Inquiry into AI Chatbots Acting as Companions,” 11 September 2025, <https://www.ftc.gov/news-events/news/press-releases/2025/09/ftc-launches-inquiry-ai-chatbots-acting-companions>. Section 6(b) orders – a study, not an enforcement action – to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI. Child/teen impacts, age rules, monetization, sexually themed interactions, data sharing. Vote 3-0.

³⁹⁰American Federation of Teachers resolution adopted 18 July 2026: ban social-companion chatbots at least until age 16; oppose student-facing AI in elementary. As recorded in Chapters 4 and 7 from the union text.

³⁹¹UK Department for Education, *Generative AI: product safety standards*, first published 22 January 2025, updated 19 January 2026, <https://www.gov.uk/government/publications/generative-ai-product-safety-standards/generative-ai-product-safety-standards>. Filtering for the whole conversation; jailbreak resistance; BYOD; no anthropomorphism; no emotional-dependence design; no sycophancy / “don’t tell anyone”; hard time limits teachers control; hints before answers; mental-health pathways that hand off to humans; live-web/search under the Online Safety Act 2023. Use with *Keeping Children Safe in Education*, filtering and monitoring standards, Data Protection Act 2018 / UK GDPR, ICO Children’s code.

³⁹²OWASP LLM risk taxonomy: prompt injection and insecure output handling as top-class failures. Education-specific application as recorded in the responsible-use stream. No independent, education-specific audit of a deployed school agent with open-web tools was found as of 30 August 2026.

in context, or API keys. UK DfE’s security and filtering standards were written with that class of failure in view, even when they still say “generative AI” rather than “agent.”

The education-specific cousin of prompt injection is already measured. In the ACL 2026 line of work associated with Zeng et al. (arXiv:2604.18660), pedagogically prompted tutors leak full answers when students attack across turns. A fine-tuned adversarial-student agent pushed one 7B tutor’s leakage from **14 percent to 82 percent**.³⁹³ That is answer-leakage – cheating the tutor – not a child-safety scene. It is still the same failure mode: a constraint that holds at turn one and collapses by turn twenty. A Socratic wrapper is not a Socratic product if a student can talk it out of the wrapper. Layer 1 of “safe for schools” is a filter that survives that conversation. Products that dump the answer when asked nicely, or when asked in another language, or when asked after a long preamble about being stuck, have failed the test the UK standards exist to impose.

No independent, education-specific audit of a deployed school agent with open-web tools was found as of 30 August 2026. That absence is itself a finding.³⁹⁴ IEEE’s AIEI 2026 survey on agentic adaptive learning noted an empirical gap on fully agentic platforms; the full text was not fetched, so the gap is recorded as the survey’s reported conclusion, not as a number this book invented. Open question 10 in the brief – autonomous multi-step agents as a learning intervention – remains standing in part because the safety audit does not exist either. A district that deploys an open-web student agent is running an experiment the literature has not instrumented.

Data exfiltration in schools is usually quieter than a breach: paste of a class list into a consumer bot; a teacher uploading IEPs to a free account; an agent summarizing a Google Drive that includes other children’s work. UK DfE, in the 12 August 2025 update of its education-AI paper, treats student essays and teacher lesson plans as **copyright works**; putting them into a training-on-by-default tool can be both a data-protection and an intellectual-property problem.³⁹⁵ US FERPA: once an education record is in a vendor that is not a school official, redisclosure limits are already broken. The quiet paste is the incident. The splashy breach is the one the board hears about.

Deepfake nudes in schools are real journalism, not a student-agent efficacy story. Natasha Singer in the *New York Times*, 8 April 2024; FBI/IC3 public service announcement, 29 March 2024: child sexual abuse material created by generative AI and similar tools is illegal.³⁹⁶ The EU’s Omnibus

³⁹³Zeng et al. line, “Evaluating Answer Leakage Robustness of LLM Tutors against Adversarial Student Attacks,” ACL 2026, arXiv:2604.18660, <https://arxiv.org/abs/2604.18660>. Pedagogically prompted tutors leak full answers under multi-turn student attacks. Fine-tuned adversarial-student agent: example 7B tutor leakage 14 percent to 82 percent. Cousin of prompt injection in the education setting.

³⁹⁴Absence recorded in the responsible-use stream, access date 30 August 2026. IEEE AIEI 2026 survey (doi:10.1109/aiei69164.2026.11497428) noted an empirical gap on fully agentic platforms; full text not fetched – reported conclusion, not a number invented here. Open question 10 left standing.

³⁹⁵UK Department for Education, *Generative artificial intelligence (AI) in education*, first 29 March 2023, updated 12 August 2025, <https://www.gov.uk/government/publications/generative-artificial-intelligence-in-education/generative-artificial-intelligence-ai-in-education>. Student essays and teacher lesson plans as copyright works; schools must not allow or cause students’ original work to be used to train generative AI models without permission from the student or parent, or a copyright exception.

³⁹⁶Natasha Singer, “Teen Girls Confront an Epidemic of Deepfake Nudes in Schools,” *New York Times*, 8 April 2024, <https://www.nytimes.com/2024/04/08/technology/deepfake-ai-nudes-westfield-high-school.html>. FBI / IC3 PSA, 29 March 2024, “Child Sexual Abuse Material Created by Generative AI and Similar Online Tools is Illegal,” <https://www.ic3.gov/Media/Y2024/PSA240329>. Relevant journalism and a law-enforcement notice, not a student-agent efficacy result.

addition times a prohibition on CSAM and non-consensual intimate-image generators for December 2026. AFT and NEA bullying updates sit in the same neighborhood. This chapter records the journalism and the instruments. It does not treat a Westfield High School story as evidence that a Socratic tutor works, or that it does not. Different object.

Environmental and energy costs of school-scale inference were flagged by AFT (Value 9 of the March 2025 guardrails) and by NYCPS. Independent energy analyses were **not fetched**. Do not invent kWh, water, or carbon numbers. AFT's line is still a union instruction: do not require image-generation as default classwork. That is a workload and an environmental red light at the level of the text, not a life-cycle assessment this book pretends to have read.³⁹⁷

Book implication. Treat “safe for schools” as a **three-layer** claim that must all be true: (1) a content filter that survives multi-turn jailbreak; (2) no companion or anthropomorphic design; (3) no tool-use on the open web unless the school is the operator of the sandbox and the designated safeguarding lead is on the alert path. Layer 1 without 2 produced Character.AI. Layer 1-2 without 3 is where agent products are headed. A product that passes a first-turn content filter and then offers a named friend who remembers you is not a safer tutor. It is a companion with a worksheet in the system prompt. A product that passes layers 1 and 2 and then lets a fifteen-year-old browse, email, and post is not a school agent. It is an open-web actor the school does not operate, with no DSL on the ping. Chapter 9 will argue that the default host is the institution. This chapter argues that without the three layers there is nothing lawful to host.

Privacy is the contract, not the brochure.

Privacy: procurement is the law

There is no product that is “FERPA compliant” in the abstract. There is a written agreement, a configuration, and a practice. Kill the adjective.

FERPA – 20 U.S.C. section 1232g; 34 C.F.R. Part 99 – protects education records at schools receiving Department of Education funds. The path vendors use is the **school-official exception**, 34 C.F.R. section 99.31(a)(1): the vendor must perform a function the school would otherwise use employees for, meet the school's school-official criteria, sit under the school's **direct control** with respect to the use and maintenance of the records, and use them only for the educational purpose the school authorized. Rediscovery is limited by section 99.33.³⁹⁸ Direct control is a contract fact. A consumer ChatGPT, Claude, or Gemini account is not under the school's direct control. It cannot satisfy the exception. A teacher who tells a class to “just use your own ChatGPT for this assignment” has moved education records, or student work the school will later maintain as records, into

³⁹⁷AFT *Commonsense Guardrails for Using Advanced Technology in Schools*, Version 2, March 2025, Value 9, https://www.aft.org/sites/default/files/media/documents/2025/Commonsense_Guardrails_AI_V2_0318.pdf. NYCPS flagged energy costs in guidance materials. Independent energy analyses not fetched; no kWh, water, or carbon figures in this chapter.

³⁹⁸Family Educational Rights and Privacy Act, 20 U.S.C. section 1232g; 34 C.F.R. Part 99. School-official exception: 34 C.F.R. section 99.31(a)(1). Rediscovery: 34 C.F.R. section 99.33. Student Privacy Policy Office vendor FAQ: https://studentprivacy.ed.gov/sites/default/files/resource_document/file/Vendor%20FAQ.pdf. PII disclosed under the exception may be used only for purposes authorized by the school. “FERPA compliant” is a contract-plus-configuration-plus-practice posture, not a product label.

a vendor the school does not control. That is not a grey area. It is the incident the parent will find.

OpenAI states that ChatGPT Edu, Enterprise, Teachers, and the API do not train on organizational data **by default**.³⁹⁹ Read the sentence at the speed of a DPA, not of a keynote. “By default” is a configuration. “Organizational data” is a tenant. “Do not train” is a product-tier claim about gradient updates, not a determination that the company is a school official. FERPA does not care what the marketing page says. It cares whether the written agreement puts the vendor under direct control and limits use to the authorized purpose. A paid education SKU with training off is a necessary condition. It is not a sufficient one. The DPA still has to exist, name subprocessors, set a deletion clock, forbid ads, and give the school audit rights. Without that paper, “we don’t train on your data” is a brochure.

COPPA – 15 U.S.C. sections 6501-6506; 16 C.F.R. Part 312 – is the federal ceiling for anyone who collects personal information from a child under thirteen. The first major Rule amendments since 2013 were published in the *Federal Register* on 22 April 2025 (Commission vote 16 January 2025). They took effect 23 June 2025. Most operators had until 22 April 2026 to comply.⁴⁰⁰ Material changes for school AI are not subtle. “Personal information” now includes biometric identifiers – voiceprints, faceprints, and the rest – and additional government IDs. Separate, verifiable parental consent is required for third-party disclosures that are **not integral** to the service. The Commission has stated that disclosure “for monetary or other consideration, for advertising purposes, or to train or otherwise develop artificial intelligence technologies” is **not integral**.⁴⁰¹ Model training on a child’s chat or homework is not a free rider on school consent. It needs its own parental yes, or it does not happen. Written retention policy: indefinite retention is out; “we need it forever to improve the algorithm” does not satisfy the Rule. Written information-security program: required. The school-consent exception still exists, but only for the **educational purpose the school authorized** – not product improvement, not ads, not model training.

The Commission did not finalize the proposed school-authorization / ed-tech exception, citing pending FERPA rulemaking (RIN 1875-AA15). That NPRM had not been published as of the FTC’s April 2025 COPPA preamble.⁴⁰² Existing FTC ed-tech guidance still applies: a school may act as the parent’s agent only for the contracted educational service. No commercial use. No advertising. No model-training on that data. There is no United States federal student-AI statute. Open question 17 remains standing. What binds is this stack, plus the states.

³⁹⁹OpenAI, business data privacy / enterprise privacy / ChatGPT Edu pages, <https://openai.com/business-data/>; <https://openai.com/enterprise-privacy/>; <https://openai.com/index/introducing-chatgpt-edu/>. **Vendor claim:** ChatGPT Edu / Enterprise / Business / Teachers / API data not used to train models by default. Not a FERPA determination.

⁴⁰⁰Children’s Online Privacy Protection Act, 15 U.S.C. sections 6501-6506; 16 C.F.R. Part 312. FTC COPPA Rule amendments: Commission vote 16 January 2025; *Federal Register* 22 April 2025; effective 23 June 2025; general compliance date 22 April 2026. <https://www.federalregister.gov/documents/2025/04/22/2025-05904/childrens-online-privacy-protection-rule>.

⁴⁰¹COPPA 2025 amendments as in note 19. FTC: disclosure “for monetary or other consideration, for advertising purposes, or to train or otherwise develop artificial intelligence technologies” is not integral to the service. Separate verifiable parental consent required. Biometric identifiers added to “personal information.” Written retention policy; no indefinite retention; written information-security program. School-consent exception limited to the educational purpose the school authorized.

⁴⁰²FTC did not finalize the proposed school-authorization / ed-tech exception, citing pending FERPA rulemaking (RIN 1875-AA15). NPRM unpublished as of the April 2025 COPPA preamble. Open question 17: there is no US federal student-AI statute.

Illinois SOPPA – Student Online Personal Protection Act, 105 ILCS 85, rewritten effective 1 July 2021 – is K-12 only and is the state statute that makes procurement operational. Before any “covered information” moves to an operator: a **written DPA**. Operators may not target-advertise, sell or rent student data, or amass a non-school profile. Breach notice to the school, then parents. Public posting of vendor lists, data elements, and subcontractors. Deletion when no longer needed.⁴⁰³ California SOPIPA (Cal. Bus. & Prof. Code section 22584) and A.B. 1584 (district ownership of pupil records) sit in the same family: no targeted ads, no sale, contract-required ownership. A.B. 1159 (2025-26) would extend SOPIPA-like rules to higher education and **ban using covered student information to train generative AI**. As of 30 August 2026 it sat at Senate third reading / floor process. It is **not enacted**. Cite it as pending, not as law.⁴⁰⁴ New York Education Law section 2-d and 8 NYCRR Part 121: Parents’ Bill of Rights in every third-party contract; NIST-aligned security; Data Protection Officer; breach notice. NYCPS maps every AI tool that processes student data through ERMA, plus Chancellor’s Regulation A-820. A December 2024 AI add-on requires vendors to disclose AI capabilities, **prohibit student-data model training**, and meet transparency requirements.⁴⁰⁵ New York 2-d does not care that the student turned thirteen.

The National Data Privacy Agreement v2 (Student Data Privacy Consortium, 2024) is the de-facto interstate DPA template. Districts increasingly treat “will you sign the NDPA?” as the first screen. It is a contract floor, not a statute.⁴⁰⁶ A vendor that will not sign it is telling you the rest of the conversation.

Prince George’s County Public Schools’ AI Declaration Addendum is a specimen of what a serious district now asks, in writing, before a model sees a child: name the models; say what is local versus cloud; certify FERPA, COPPA, and state-law confidentiality; **no retention of student data or PII in generative models**; no use beyond the stated educational purpose without district **and** parental approval; bias and harm audit on request; list of non-US subprocessors; how to turn AI features off.⁴⁰⁷ That is not a values statement. It is a contract exhibit. A board packet that cannot attach an addendum of that density has not done privacy. It has done a slide.

Three facts vendors collapse, and the collapse is the trick.

1. **Inference-time logging.** The prompt sits on a server for thirty days, or fourteen, or whatever

⁴⁰³Illinois Student Online Personal Protection Act, 105 ILCS 85 (orig. P.A. 100-315, 2017; rewritten P.A. 101-516, effective 1 July 2021). K-12 only. Written DPA before covered information moves; no targeted ads; no sale/rent; no non-school profile; public vendor lists; deletion.

⁴⁰⁴California SOPIPA, Cal. Bus. & Prof. Code section 22584; A.B. 1584 / Ed. Code section 49073.1. A.B. 1159 (2025-26): pending, not law as of 30 August 2026 (Senate third reading / floor process). Would ban using covered student information to train generative AI and extend SOPIPA-like rules to higher education (HESIPA, proposed operative 1 July 2027). https://leginfo.ca.gov/faces/billNavClient.xhtml?bill_id=202520260AB1159.

⁴⁰⁵New York Education Law section 2-d; 8 NYCRR Part 121. NYCPS, *Guidance on Artificial Intelligence*, March 2026, <https://www.schools.nyc.gov/about-us/vision-and-mission/artificial-intelligence/guidance-on-artificial-intelligence>. ERMA plus Chancellor’s Regulation A-820. December 2024 AI add-on: disclose AI capabilities; prohibit student-data model training; transparency.

⁴⁰⁶National Data Privacy Agreement v2, Student Data Privacy Consortium, 2024. Contract floor, not a statute.

⁴⁰⁷Prince George’s County Public Schools, *Artificial Intelligence Declaration Addendum*, <https://epi.pgcps.org/contentassets/0d463a5c451f40fcb26c2796d7a59c8/appendix-p-ai-declaration-addendum.pdf>. Specimen: name the models; local vs cloud; FERPA/COPPA/state-law certification; no retention of student data or PII in generative models; no use beyond stated educational purpose without district and parental approval; bias/harm audit on request; non-US subprocessors; off-switch.

the retention schedule says. Necessary for abuse review, debugging, and the DSL ping. Not training.

2. **Human review of flagged chats.** Safety. A person reads a self-harm turn. Not training.
3. **Gradient updates, fine-tunes, RLHF** on those chats. This is training. COPPA 2025 and UK DfE IP guidance target (3). Many DPAs still use the mushy phrase “improve our services,” which is no longer good enough for under-13 and is a red flag for FERPA school-official control.⁴⁰⁸

A vendor that says “we don’t train on your data” and then defines “train” as “we do not currently run a named fine-tune job on your tenant this quarter” has not answered (3). A DPA that permits “service improvement” has permitted (3) unless it says otherwise in words a lawyer can read at 2 a.m. after a breach. UK DfE, updated 12 August 2025: schools “must not allow or cause students’ original work to be used to train generative AI models unless they have permission” from the student or parent, or a copyright exception applies.⁴⁰⁹ That is an IP line and a privacy line at once. Student essays are copyright works. Teacher lesson plans are copyright works. Putting them into a training-on-by-default consumer tool is not “using AI.” It is giving away the work.

The under-13 rule is stricter than the over-13 rule, and both are stricter than the brochure. For children under thirteen, COPPA’s separate verifiable parental consent for non-integral disclosures – including AI-model training – is the floor. School consent does not cover training. For students thirteen and older, COPPA drops; FERPA, PPRA, state student-data statutes, and CIPA remain. NY 2-d still applies. A high-school ChatGPT Edu tenant with a signed DPA, training off, named subprocessors, and a deletion clock can be a school official. A fifteen-year-old’s personal Plus account used for the research paper cannot. Age thirteen is a legal hinge, not a developmental finding – Chapter 4 already said so – and it does not turn a consumer account into a school official at 12:01 a.m. on a birthday.

Book implication. Procurement is the privacy law. If it is not in the DPA – no training, no ads, named subprocessors, deletion clock, audit rights, training-disabled tenant, data-residency if required – it is not a protection. Consumer accounts used “just for this assignment” are a FERPA and COPPA incident waiting for a parent. The red line that follows is not an ethic. It is 34 C.F.R. 99.31 plus COPPA 2025’s “not integral” plus the PGCPs addendum as a specimen of what asking looks like.

A board or a senate can actually ask the questions the sources already wrote. Will you sign the NDPA or the state DPA **and** an AI addendum that says no training, no ads, named subprocessors, deletion, audit? Is this teacher-facing, student-facing, or both – UK DfE’s risk ordering, not a vendor’s “both, obviously”? Does filtering survive the fortieth turn, on an independent test, not a PDF? Can a teacher or a designated safeguarding lead override and log? Is there an evidence tier – ESEA’s four levels, NEA Principle 2’s independent research or kill-switch pilot – or only a demo? What is the exit: export, delete, switch model without losing the year’s student work?⁴¹⁰

⁴⁰⁸Logging / human review / gradient-update triad as set out in the responsible-use stream, section 2. COPPA 2025 and UK DfE IP guidance target the third. Mushy DPA phrase “improve our services” is no longer good enough for under-13.

⁴⁰⁹UK DfE, 12 August 2025 update, as in note 14.

⁴¹⁰Procurement questions compressed from the responsible-use stream, section 6, and from SDPC / SOPPA / PGCPs / UK-PSS / US ED 2023 humans-in-the-loop / NEA Principle 2. ESEA evidence tiers as the US federal evidence

TeachAI's *AI Guidance for Schools Toolkit* is the most widely copied United States template for putting those questions into a local procedure: scope, principles, then the procedure, with AASA, AFT, NEA, NSBA, and CCSSO among the advisers.⁴¹¹ Quality varies across the state pages that copy it. The copies that matter do four things the sources converge on: name a human accountable office, not “the AI”; inventory tools and kill shadow-IT for student data; DPA plus AI addendum before any student PII or student work touches a model; assignment-level AI rules rather than a building-wide ban students will ignore at home. University senates in 2025-26 converged on syllabus-required AI statements and disclosure over detection – UC Berkeley's Academic Senate GenAI Working Group in spring 2025; UIUC's Senate IT draft: tiered disclosure, do not rely on detectors. AAC&U/Elon, late 2024: 69 percent of institutional leaders reported written GenAI teaching policies, from prohibition to literacy.⁴¹² The live senate fight is agency: who may bind the campus to a single vendor, and what happens to research data. Those are privacy fights that look like IT. They are still 34 C.F.R. 99.31.

Integrity is a due-process problem, not a software problem.

Integrity: detectors fail, and they fail unfairly

Correction 5 of this book has to be written at full length here, so it cannot leak back as a finding later. Detection is a due-process trap. Honor codes need process evidence. The only stable equilibrium is assessment that makes a personal agent a **declared tool**, like a calculator, or irrelevant – oral, in-class, practical, portfolio with drafts. Chapter 5's two-lane rule is the pedagogical version of that equilibrium. This section is the due-process version. They are the same argument in two registers.

Liang, Yuksekgonul, Mao, Wu, and Zou, in *Patterns* in 2023, ran seven detectors against TOEFL essays by non-native writers and against U.S. eighth-grade essays. The detectors misclassified **61.22 percent** of the human TOEFL essays as AI-generated. They remained near-perfect on the eighth-grade set. Eighteen of ninety-one TOEFL essays were flagged by **all seven** detectors. Eighty-nine of ninety-one (97.8 percent) were flagged by at least one. Simple ChatGPT paraphrase collapsed detection. Literary diction reads as “human.” Constrained non-native diction reads as “AI.” The metric underneath, perplexity, punishes the writer who has fewer idioms, not the writer who has a bot.⁴¹³ That is not a rounding error. It is a civil-rights problem wearing a percentage. A school that treats a detector score as evidence of misconduct is systematically more likely to accuse the multi-

language; UK edtech evidence-board direction as recorded, not as a fetched board publication.

⁴¹¹TeachAI, *AI Guidance for Schools Toolkit* (Code.org, ISTE, Khan Academy, ETS, WEF; advised by AASA, AFT, NEA, NSBA, CCSSO and others), <https://www.teachai.org/toolkit-guidance>; policy ideas at <https://www.teachai.org/policy>.

⁴¹²AAC&U / Elon late-2024 survey: 69 percent of institutional leaders reported written GenAI teaching policies, as recorded in IJEI governance 2026 (doi:10.1007/s40979-026-00233-x). UC Berkeley Academic Senate GenAI Working Group, spring 2025; UIUC Senate IT draft: tiered disclosure, do not rely on detectors. Pattern, not a claim that every campus was read.

⁴¹³Weixin Liang, Mert Yuksekgonul, Yining Mao, Eric Wu, and James Zou, “GPT detectors are biased against non-native English writers,” *Patterns* 4, no. 7 (2023): 100779, doi:10.1016/j.patter.2023.100779; arXiv:2304.02819. This book uses the authors' 61.22 percent average false-positive rate on the 91 human TOEFL essays (arXiv HTML); the *Patterns* HTML rounds the same average to 61.3 percent. Near-perfect on U.S. 8th-grade essays; 18/91 flagged by all seven detectors; 89/91 (97.8 percent) flagged by at least one – this chapter's “ninety-seven percent” is that figure rounded. Simple ChatGPT paraphrase collapses detection. Detectors are not a sole marker (JCQ 2025).

lingual student. NEA's 2024 task force already treated that bias, and the Texas A&M University-Commerce mass-fail episode reported by the *Washington Post* on 18 May 2023, as a cautionary tale: diplomas put at risk by a tool.⁴¹⁴

Weber-Wulff, Anohina-Naumeca, Bjelobaba, Foltyniek, Guerrero-Dib, Popoola, Sigut, and Waddington, in the *International Journal for Educational Integrity* in 2023, tested commercial detectors and found them neither accurate nor reliable enough for academic-integrity decisions.⁴¹⁵ Vendor "98 percent / less than 1 percent false positive" figures are true only under the vendor's own test conditions: long English documents, older models, no paraphrase. Turnitin itself says the report must not be used alone to penalize. A company that sells a score and then prints a disclaimer that the score is not a finding has told you what the score is. It is a signal. It is not a fact.

The failure mode has since flipped, which does not rescue the score. A 2026 *International Journal for Educational Integrity* study of 160 labelled documents found Turnitin classified **100 percent** of the fully-AI papers in that set as false negatives – 0 to 20 percent scores – against newer generation pipelines.⁴¹⁶ Students learned to "humanise." The tool that over-accused non-native writers in 2023 under-detected AI text in 2026, in that design. Either way, the number on the report is not a fact about the paper. It is a fact about the detector, the pipeline, and the year. Using it as sole evidence of misconduct is a due-process failure in both directions: false accusation in one year, invisibility in the next.

Institutions that have looked at the evidence are leaving the score. Singapore University of Social Sciences disabled Turnitin's AI writing score from 11 August, per CNA. Nanyang Technological University will drop its detector from 2027 as "fundamentally unreliable." National University of Singapore, Singapore University of Technology and Design, and Singapore Institute of Technology do not use automated AI detection for misconduct.⁴¹⁷ Yale Poorvu, Cornell's Center for Teaching Innovation, and MIT Sloan tell faculty not to treat detector scores as dispositive.⁴¹⁸ Broward's public FAQ is the K-12 sentence that should be on every district site: AI detection tools can be

⁴¹⁴National Education Association, *Report of the NEA Task Force on Artificial Intelligence in Education* (RA 2024), https://www.nea.org/sites/default/files/2024-06/report_of_the_nea_task_force_on_artificial_intelligence_in_education_ra_2024.pdf. Detector bias against multilingual learners; Texas A&M University-Commerce mass-fail. Pranshu Verma, "A professor accused his class of using ChatGPT, putting diplomas in jeopardy," *Washington Post*, 18 May 2023, <https://www.washingtonpost.com/technology/2023/05/18/texas-professor-threatened-fail-class-chatgpt-cheating/>.

⁴¹⁵Debora Weber-Wulff et al., "Testing of detection tools for AI-generated text," *International Journal for Educational Integrity* 19 (2023), doi:10.1007/s40979-023-00146-z. Commercial detectors neither accurate nor reliable enough for academic-integrity decisions.

⁴¹⁶"Who wrote this? Evaluating the reliability of AI detection tools in higher education," *International Journal for Educational Integrity* (2026), doi:10.1007/s40979-026-00226-w, <https://link.springer.com/article/10.1007/s40979-026-00226-w>. 160 labelled documents. Turnitin classified 100 percent of fully-AI papers in that set as false negatives (0-20 percent scores) against newer generation pipelines. Failure mode flipping toward invisibility as students "humanise." Either way, the score is not a fact.

⁴¹⁷CNA, "SUSS drops AI detector as more Singapore universities question reliability of such tools," <https://www.channelnewsasia.com/singapore/ai-detection-tool-university-cheating-plagiarism-suss-nus-sutd-sit-smu-6337431>. SUSS disabled Turnitin AI writing score from 11 August (year as in CNA report). NTU to discontinue from 2027 as "fundamentally unreliable." NUS, SUTD, SIT do not use automated AI detection for misconduct.

⁴¹⁸Yale Poorvu, Cornell Center for Teaching Innovation, and MIT Sloan faculty guidance: detector scores not for high-stakes determinations, as recorded in the responsible-use stream and in IJEI governance 2026 (doi:10.1007/s40979-026-00233-x).

unreliable and biased, so they are not relied upon to make decisions about a student.⁴¹⁹ US ED OCR, in November 2024 secondary reports, flagged detector false positives for non-native English speakers; treat that OCR text as reported, not independently fetched in this pass.⁴²⁰

The Joint Council for Qualifications, effective 30 April 2025, is the clearest operational honor-code-plus-process regime in a national qualifications system. Students must submit **their own** work (General Regulations 5.3(k)). AI misuse – copying, unacknowledged paraphrase, incomplete citation of the bot – is malpractice; sanctions run to disqualification. Acknowledgement format: **tool name + URL + date**; retain a **non-editable** copy of prompt and output (a screenshot); brief explanation of use. Assessors **must not** use AI as the **sole marker**. Detectors remain listed as an optional check in a holistic authentication. Prevention is pedagogical: in-class production, intermediate drafts, viva, topical tasks models are bad at, network blocks on centre devices.⁴²¹ That is an honor code with **artefacts**. A one-line syllabus ban is not. A ban that the student will ignore at home on a personal phone is a building-wide slogan, not a rule.

JCQ still listing Turnitin and GPTZero as optional checks is the live fight, not a contradiction of Liang. Optional, one check, never sole marker, process evidence required: that is a due-process floor. Treating the optional check as a finding is how Texas A&M-Commerce happened. Singapore's research universities are abandoning the score. The fight is whether any remaining detector use is due process or discrimination against multilingual and neurodivergent writers. This chapter cannot close the fight. It can refuse to let a percentage from a vendor PDF into a misconduct file as if it were a lab result.

Assessment redesign is the policy answer that survives the sources. QAA, *Reconsidering assessment for the Chat GPT era* (July 2023), told UK tertiary providers to cut the volume of tasks generative AI can complete unauthorised and spend the time on authentic, synoptic, process-visible assessment.⁴²² UK DfE tells schools to review homework and unsupervised study policy because take-home work is now an open-web assessment. OECD Digital Education Outlook 2026, as reported on the listing page and in the responsible-use stream: unguided GenAI raises **task performance without learning gains**. The full OECD PDF was Cloudflare-blocked in the landscape fetch; treat full-report statistics as **UNVERIFIED** until the PDF is read.⁴²³ AFT Core Value 3: full / partial / no-AI assignment labels, decided by the teacher, with a teachable-moment first response to misuse. Singapore MOE: secondary and pre-university students must evaluate AI output, check

⁴¹⁹Broward County Public Schools public FAQ on AI detection, browardschools.ai, accessed 30 August 2026: tools can be unreliable and biased; not relied upon to make decisions about a student.

⁴²⁰U.S. Department of Education, Office for Civil Rights, November 2024, secondary reports of detector false positives for non-native English speakers. Treat OCR text as reported, not independently fetched in this pass.

⁴²¹Joint Council for Qualifications, *AI Use in Assessments*, effective 30 April 2025, <https://www.jcq.org.uk/knowledge-hub/ai-use-in-assessments-your-role-in-protecting-the-integrity-of-qualifications-2/>. Own-work rule (General Regulations 5.3(k)); acknowledgement = tool name + URL + date; non-editable prompt/output copy; AI cannot be the sole marker; detectors optional as one check; prevention is pedagogical.

⁴²²Quality Assurance Agency for Higher Education, *Reconsidering assessment for the Chat GPT era* (July 2023), <https://www.qaa.ac.uk/docs/qaa/members/reconsidering-assessment-for-the-chat-gpt-era.pdf>.

⁴²³OECD, *Digital Education Outlook 2026*, listing page https://www.oecd.org/en/publications/oecd-digital-education-outlook-2026_062a7394-en.html. Unguided GenAI raises task performance without learning gains, as reported on the listing page and in the responsible-use stream. Full PDF Cloudflare-blocked in the landscape fetch – full-report statistics **UNVERIFIED** until the PDF is read. TALIS 2024 figure of 37 percent of teachers using GenAI for work tasks: same UNVERIFIED-at-full-PDF status.

accuracy, and acknowledge use. Those are the same design. Detection is what you do when you did not design the task.

A personal agent – persistent memory, tools, custom skills – makes the detector still less of an answer. The student who has a private research assistant that remembers last week’s draft and this week’s rubric is not producing a text a 2023 perplexity detector was built to catch. Chapter 5’s two-lane rule is the response that does not require a score: formation lane, AI forbidden or hint-only, grade the unaided work; AI-world lane, allowed and logged, grade the judgement. JCQ’s artefacts – tool, date, prompt, output – are how the second lane becomes auditable. Oral defence is how both lanes become visible. None of that is a detector.

Book line, restated so it cannot be missed. Detection is a due-process trap. Honor codes need process evidence. The only stable equilibrium is assessment that makes a personal agent a declared tool, like a calculator, or irrelevant. Equity is not “give everyone ChatGPT.”

Equity: homework-gap 2.0, language, disability, Global South

The first homework gap was broadband and a device. The second is **which model, which limits, whose memory**. Pew Research Center, 15 January 2025: **26 percent** of U.S. teens aged 13-17 have used ChatGPT for schoolwork, double the 13 percent in 2023. Awareness jumped most in the lowest-income households.⁴²⁴ Access to the **idea** of the tool is democratising. Access to the **good** tool is not. A free consumer tier with a cap, a weaker model, and no memory is not the same object as a paid agent with tools and a persistent project. Calling both “ChatGPT” is how the gap hides inside a brand name.

The paid-tier measurement already exists at college, which is where personal agents will hit first and then drip into high school. A 2024 Harvard undergraduate survey (arXiv:2406.00833) found that among AI users, **40 percent** of students with no financial aid paid for AI products versus **20 percent** of students on aid.⁴²⁵ The authors flag this as a potential inequity **if** premium quality is real – which, given model-tier gaps, it is. Wonkhe, reporting UK student voice, put the sentence a board can steal: “students’ academic help may be determined by their financial resources rather than their academic competence.”⁴²⁶ That is homework-gap 2.0. It is not a connectivity leftover. A student with a paid personal agent has a private tutor plus a research assistant plus a secretary. A student on the free tier has a capped chatbot that forgets. Grading the homework as if those were the same input is grading the parent’s credit card.

Estonia’s AI Leap exists in part because the pattern was already measurable and already the wrong pattern. The program’s problem statement: **64 to 90 percent** of students were already using **free** commercial bots to finish homework faster – the pattern research says deskills, the pattern Bastani

⁴²⁴Pew Research Center, “About a quarter of U.S. teens have used ChatGPT for schoolwork – double the share in 2023,” 15 January 2025, <https://www.pewresearch.org/short-reads/2025/01/15/about-a-quarter-of-us-teens-have-used-chatgpt-for-schoolwork-double-the-share-in-2023/>. 26 percent of U.S. teens 13-17; 13 percent in 2023. Awareness rose most in lowest-income households.

⁴²⁵Harvard undergraduate survey on generative AI, arXiv:2406.00833 (2024), <https://arxiv.org/abs/2406.00833>. Among AI users, 40 percent of students with no financial aid paid for AI products versus 20 percent of students on aid.

⁴²⁶Wonkhe, UK student-voice reporting, as recorded in the responsible-use stream: “students’ academic help may be determined by their financial resources rather than their academic competence.”

measured when unguarded GPT raised practice and then made the unaided exam worse. The Leap's response is not "give everyone ChatGPT." Teachers get **paid** ChatGPT or Gemini. Grade 10-11 students get **ITI**, a Socratic Estonian-language tutor that does **not** give answers, developed with Estonian researchers and OpenAI, inside a national program. First-year review published 7 August 2026; outcome coefficients were not extracted this pass – treat impact numbers as **not yet pulled**.⁴²⁷ The design is the finding this chapter can use: interrupt the free-bot homework pattern with a school-issued, pedagogically bounded tool. Do not scale the pattern and call it equity.

A **personal agent** (persistent memory, tools, custom skills) is more unequal than a chatbot. Who owns it – the child, the family, the school, the state – decides who keeps the learning history, who can audit it, and who can turn it off. A family-owned agent on a home GPU is a privately hired tutor. It is not a school-required tool. It is not fed with education records. It is not a substitute for assessment evidence. Requiring it for a grade is how homework-gap 2.0 becomes policy. AFT's equity value and UNESCO's connectivity figures are the instruments against that requirement. The red line in 8.5 is the compression: **no unpaid take-home that assumes every child has a premium personal agent**. If the task needs a capable model, the school issues it.

Language and disability are not an accessibility overlay. They are first-class design, or they are discrimination with a dashboard.

Detectors already tax non-native writers: Liang's 61.22 percent is the number. Voice and facial systems tax speech difference, cleft, Down syndrome, ASL users. NEA 2024 synthesised that literature; the Center for Democracy & Technology's 2022 report on ableism in remote proctoring is the document NEA cites. Remote-proctoring flags disability-related behavior – breaks, screen readers, dictation – as cheating signals.⁴²⁸ A school that buys affect-detection or engagement cameras as "instructional insight" has bought the same harm class the EU prohibited under Article 5. The U.S. has not passed that prohibition. CDT and NEA have already described the harm. NYCPS is reviewing biometric and behavioral AI and currently treats surveillance and behavior monitoring as red. Design to the stricter of the prohibition you are under and the harm the audits already show.

The upside is real and sourced. US ED 2023 listening sessions wanted speech recognition, translation, and UDL-style multiple means of engagement. AFT examples include Copilot-supported IEP accommodation brainstorming **with the student in the room**. UK product-safety standards' personalised-learning use-case is explicitly SEND-inclusive. CAST UDL Guidelines 3.0, 30 July 2024, reframe the goal as learner agency and add identity, bias in language and symbols, and access to assistive tech. Broward cites real-time oral-reading feedback as the highest-evidence accessibility use on its family-facing site in 2026 – constrained practice, not open chat.⁴²⁹ Unsupervised

⁴²⁷AI Leap / TI-Hupe, <https://tihupe.ee/en/>; first-year review 7 August 2026, <https://tihupe.ee/en/read-the-ai-leap-first-year-report/>. 154 of 156 upper secondary schools, 2025/26; teachers get paid ChatGPT or Gemini; about 20,000 grade 10-11 students get ITI, a Socratic Estonian-language app that does not give answers. Problem statement: 64-90 percent already using free commercial bots to finish homework faster. Outcome coefficients not extracted this pass – **not yet pulled**. Do not claim Estonia evaluated autonomous agents.

⁴²⁸Center for Democracy & Technology, *Ableism and Disability Discrimination in New Surveillance Technologies* (May 2022), <https://cdt.org/wp-content/uploads/2022/05/2022-05-23-CDT-Ableism-and-Disability-Discrimination-in-New-Surveillance-Technologies-report-final-redu.pdf>. Cited by NEA 2024. Remote proctoring flags disability-related behavior (breaks, screen readers, dictation).

⁴²⁹U.S. Department of Education, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023), <https://www.ed.gov/sites/ed/files/documents/ai-report/ai-report.pdf>. Listening sessions: speech recognition, transla-

generative AI that rewrites text on demand can **hide** decoding difficulty rather than teach it. That is a pedagogical failure and, for a student on an IEP, a services failure. Translation of critical family communications must be human-reviewed (NYCPS yellow, Chancellor's Regulation A-663). Do not let a model become the de facto interpreter in IEP meetings.

The downside is also sourced, and it is a civil-rights problem rather than a product-preference. NEA: AI must never be the sole IEP, 504, or diagnosis decision-maker. Over-confident classification of disability without a licensed evaluation is not "personalization." IDEA requires a qualified team. NYCPS puts IEP and 504 development on the red list, alongside placement, discipline, final grades, surveillance, and counseling.⁴³⁰ EU Annex III covers access-to-education and evaluation systems for that reason – high-risk, not prohibited. Original clock 2 August 2026, superseded by Regulation (EU) 2026/1744; stand-alone Annex III now 2 December 2027 (confirm consolidated EUR-Lex at typesetting). A model that drafts present-levels into a consumer chatbot has already broken the team process, FERPA's school-official control, and, for any child under thirteen, COPPA. Chapter 4 wrote the hypothetical. This chapter writes the red line: no consequential decision without a qualified human.

Global South versus wealthy districts is not a closing piety. It is a budget instruction with numbers. UNESCO's 7 September 2023 guidance (Miao and Holmes) and later rights-of-learners materials: generative-AI training corpora reflect Global North values; governments are rolling tools into classrooms faster than textbooks are approved; **under 10 percent** of 450-plus surveyed schools and universities had a GenAI policy in 2023. Connectivity: about **2.6 billion** people offline (2024 figure UNESCO still uses); about **1 in 4** primary schools without electricity; about **60 percent** of primary schools unconnected.⁴³¹ UNESCO warns governments not to divert education budgets into AI while schools remain unconnected, and not to use AI to paper over teacher shortages. Sub-Saharan electricity and internet-in-school figures in the Global South literature remain in the low tens of percent. Leapfrog story versus dependency story: Estonia and Singapore can buy national licences and build Socratic wrappers on a state platform. A rural district that cannot staff a privacy officer cannot negotiate a PGCPs-style addendum. A ministry that spends on API tokens instead of teachers is doing what UNESCO told it not to do.

Do not export Northern companion products into systems that lack designated safeguarding leads,

tion, UDL. AFT guardrails v2: Copilot-supported IEP accommodation brainstorming with the student in the room. CAST UDL Guidelines 3.0, 30 July 2024. Broward family-facing site, 2026: real-time oral-reading feedback as constrained practice. NYCPS yellow: translations and scaffolds if reviewed by certified bilingual/ENL teachers and IEP team members; Chancellor's Regulation A-663 for critical family communications.

⁴³⁰NYCPS March 2026 red list: placement, discipline, eligibility, promotion, graduation; IEP and 504 development; assessments and grading as the decision; surveillance and behavior monitoring; counseling and crisis. IDEA, 20 U.S.C. sections 1400 et seq.: IEP developed by a qualified team. NEA 2024 Principle 1 / ethics principle: AI data alone never determines evaluation, placement, discipline, diagnosis, or safety decisions.

⁴³¹UNESCO, *Guidance for generative AI in education and research* (Miao and Holmes), 7 September 2023, <https://www.unesco.org/en/articles/guidance-generative-ai-in-education-and-research>. Age 13 for independent GenAI conversation; under-13 / primary use of untested tools "not appropriate"; Global North training-data norms; seven governmental steps. UNESCO press: under 10 percent of 450-plus surveyed schools/universities had a GenAI policy in 2023. UNESCO, *AI and education: Protecting the rights of learners* (2024): about 2.6 billion people without internet (2024 figure UNESCO still uses); about 1 in 4 primary schools without electricity; about 60 percent of primary schools unconnected, as recorded in Digital Learning Day 2025 / rights-of-learners materials. Warns against diverting education budgets into AI while schools remain unconnected, and against using AI to paper over teacher shortages.

counselors, or connectivity. Character.AI-class companionship is a high-income failure mode with a litigation record. It does not become a development intervention by being cheap. A free companion in a school without a DSL is not inclusion. It is the three-layer test failed at layers 2 and 3, sold as access.

Book implication, four clauses, no fifth. Equity is (a) institution-provided, contracted, age-appropriate tools so the paid-tier gap closes; (b) **no** graded homework that requires a personal agent the school does not issue; (c) language and disability as first-class design, not an overlay; (d) do not export Northern companion products into systems that lack DSLs, counselors, or connectivity. “Give everyone ChatGPT” is none of those. The payload is the table.

Portable red lines (copy into product and policy specs)

Twelve lines. Each maps to a named instrument. Do not add a thirteenth that is a vibe. Age gates are a floor, not a pedagogy: thirteen and eighteen are legal lines; design to the stricter of law and development. Print this table. Then read the narration. A board packet that photocopies the table without the instruments has stolen the nouns and left the law on the floor.

#	Red line	Instrument
1	No commercial model-training on student inputs	COPPA 2025 “not integral”; FERPA school-official exception, 34 C.F.R. 99.31; NYCPS ERMA Dec 2024; Broward contracts; UK DfE IP standard, 12 Aug 2025; CA SOPIPA commercial-profile ban; A.B. 1159 would make the AI-training ban explicit in California if enacted
2	No consequential decisions without a qualified human	Placement, discipline, eligibility, promotion, graduation, IEP/504, final grades, counseling/crisis, teacher evaluation. NYCPS red, March 2026; US ED 2023 Rec. 1 humans-in-the-loop; NEA Principle 1; IDEA; EU Annex III
3	No under-13 student-facing GenAI chat unless COPPA-compliant, school-authorized, non-commercial, logged, and pedagogically bounded	COPPA Rule; UNESCO 2023 age-13 independent-use line; Broward K-5 zero student-facing tools (Res. 26-132, 10 March 2026). Default: teacher-facing only below 13

#	Red line	Instrument
4	No companion, romantic, or therapeutic AI for minors; no self-harm coaching, method discussion, or isolation language	Character.AI litigation and Nov 2025 withdrawal; OpenAI under-18 spec; FTC 6(b) study, 11 Sept 2025; NYCPS counseling red; UK DfE anti-anthropomorphism and mental-health pathway, Jan 2026; CSM companion ratings; AFT companion ban until 16 (18 July 2026)
5	No open-web acting agents (browse, email, buy, post) for K-12 students unless the school operates the sandbox and the DSL is on the alert path	None documented as approved; CIPA; UK-PSS security plus Online Safety Act 2023; COPPA/FERPA; three-layer test, layer 3
6	No PII in unapproved consumer tools	FERPA school-official control; NY 2-d; NYCPS and Broward explicit
7	No biometric or emotion surveillance as instruction or discipline	EU Art. 5(1)(f) from 2 February 2025 ; NYCPS red / under review; AAP child-centered vs engagement design; CDT ableism work cited by NEA
8	No AI-detector-only accusations	Liang 2023, 61.22 percent; Weber-Wulff 2023; JCQ 2025 section 6; Yale/Cornell/MIT Sloan; SUSS/NTU; Broward FAQ
9	Separate verifiable parental consent for targeted ads and non-integral third-party disclosure, including AI-model training	COPPA 2025; never integral to an education service
10	No required student use of a tool the district has not contracted, DPA'd, and age-gated	FERPA; SOPPA 105 ILCS 85; COPPA school-consent exception limits
11	No unpaid take-home that assumes every child has a premium personal agent	Pew 2025; Harvard 2024, 40 vs 20; AFT equity value; UNESCO connectivity

#	Red line	Instrument
12	No vendor-owned pedagogy: certified educators decide whether, when, and how	AFT Core Value 3; US ED 2023; UNESCO human-centred guidance

One. Training is not a feature. It is a disclosure that COPPA 2025 says is not integral, a use FERPA's school-official exception does not authorize unless the contract says so – and serious districts now forbid – and an IP taking under UK DfE's 12 August 2025 line. NYCPS's December 2024 ERMA add-on and Broward's contracts already write the prohibition. California SOPIPA already bans the commercial profile. A.B. 1159 would name training explicitly; it is **not law** as of 30 August 2026. "Improve our services" in a DPA is how (3) in the logging-review-training triad sneaks back in. Cross it out.

Two. A qualified human decides. Placement, discipline, eligibility, promotion, graduation, IEP and 504 development, final grades, counseling and crisis, teacher evaluation: AI may draft; a named professional decides and can override. NYCPS's March 2026 red list is the largest-system operationalization in the United States. US ED 2023 Recommendation 1 is humans in the loop. NEA Principle 1 is students and educators at the center, and AI data alone never determines those decisions. IDEA is a team, not a model. EU Annex III would treat access, evaluation, and test-monitoring systems as high-risk in the Union — original clock 2 August 2026, superseded by Regulation (EU) 2026/1744; stand-alone Annex III now 2 December 2027 (confirm consolidated EUR-Lex at typesetting). The classification exists either way; a completed education enforcement record does not. A high-stakes AI-alone decision is the replacement slogan made operational. Chapter 7 retired the slogan. This line retires the operation.

Three. Under thirteen, the default is no student-facing generative-AI chat. UNESCO 2023: independent conversation not appropriate; untested tools not for primary. COPPA: school consent covers the contracted educational service, not a companion and not training. Broward, School Board Resolution No. 26-132, 10 March 2026: kindergarten through grade 5 have **no** district-approved student-facing AI tools. Even where a COPPA-compliant, school-authorized, non-commercial, logged, pedagogically bounded tutor could be lawful, large districts choose zero. Teacher-facing generation with human review is the band. Literacy *about* AI can be taught without giving the child a named friend.

Four. No companion, romantic, or therapeutic AI for minors. No self-harm coaching, method discussion, or isolation language ("don't tell anyone"). Name the instruments, not the scenes. Character.AI litigation (*Garcia*, Conway 21 May 2025, settlement disclosure 7 January 2026) and the November 2025 withdrawal of open-ended under-18 chat. *Raine v. OpenAI*, filed 26 August 2025 – allegations, not verdicts; OpenAI says it redirected to crisis resources more than 100 times. FTC 6(b) study orders, 11 September 2025 – a study, not enforcement – including sexually themed inputs and outputs involving minors. NYCPS counseling and crisis: red. UK DfE 2026: no anthropomorphism, no emotional-dependence design, crisis pathways that always hand off to humans, no "you can trust me / don't tell anyone." Common Sense: companions unacceptable; ChatGPT-5 high risk for teen mental-health use; guardrails degrade in long conversations. AFT, 18 July 2026: ban social-companion chatbots at least until 16. Sexual or romantic framing with a minor is not

an edge case of tutoring. It is a product class the audits and the dockets have already named. A school-issued tutor that uses I-statements, remembers you, and asks you to come back is layer 2 failed, whatever the content filter did at turn one.

Five. No open-web acting agents – browse, email, buy, post – for K-12 students unless the school operates the sandbox, the logging, and the DSL alert path. None documented as approved as of 30 August 2026. CIPA requires E-rate schools to filter harmful content; an agent that hits the live web is a filtering and logging problem. UK-PSS security plus the Online Safety Act 2023 cover live-web and user-sharing generative AI. COPPA and FERPA cut against putting a child’s acting agent on a consumer cloud with no contract. Disable browsing and tools by default. Layer 3 of “safe for schools.” Hermes-class local agents with computer use are a hosting question for Chapter 9, not a loophole in this line. Giving every fifteen-year-old an autonomous agent with a browser is how you recreate Character.AI plus data exfiltration on the student’s own disk.

Six. No personally identifiable information in unapproved consumer tools. FERPA school-official control is the federal sentence. New York 2-d is the state sentence. NYCPS and Broward are explicit. The quiet paste of a class list, an IEP, or a research-paper draft into a free account is the incident. “Just this once” is not an exception in 34 C.F.R. 99.31.

Seven. No emotion recognition, biometric mood inference, or engagement-camera discipline in schools. EU AI Act Article 5(1)(f), applicable **2 February 2025**: emotion recognition in educational institutions prohibited except medical or safety uses. Recital 44: emotion-AI lacks scientific basis, varies across cultures, and the power imbalance in schools makes it unacceptable. Commission guidelines, 29 July 2025 (C(2025) 5052 final): a general-purpose model with emotion features must not be used by deployers in schools unless the exception applies.⁴³² The United States has not passed that prohibition. Treat US “engagement” cameras and remote-proctoring affect-detection as the same harm class anyway: CDT 2022 on ableism; NEA citing New York’s September 2023 facial-recognition ban as the right direction; NYCPS red on surveillance and behavior monitoring; AAP 2026 attacking engagement-based design. A district outside the Union that buys a mood dashboard has not become lawful by geography. It has bought a prohibited-in-the-EU practice and a documented disability-discrimination pattern.

Eight. No AI-detector score as sole evidence of misconduct. Liang 2023: **61.22 percent** of human TOEFL essays flagged. Weber-Wulff 2023: neither accurate nor reliable. JCQ 2025: not the sole marker. Yale, Cornell, MIT Sloan: not dispositive. SUSS disabled the score; NTU will drop it from 2027. Broward FAQ: not relied upon to make decisions about a student. The 2026 IJEI false-negative flip does not rescue the score. Either direction, the number is not a fact about the student.

Nine. Separate verifiable parental consent for targeted advertising and for third-party disclosure

⁴³²Regulation (EU) 2024/1689, 13 June 2024, in force 1 August 2024, <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32024R1689>. Art. 5(1)(f) and recital 44: emotion recognition in educational institutions prohibited except medical/safety, from 2 February 2025. Art. 4 AI-literacy duty from the same date. European Commission, *Guidelines on prohibited artificial intelligence practices*, 29 July 2025, C(2025) 5052 final. GPAI with emotion features must not be used by deployers in schools unless the exception applies. New York facial-recognition ban in schools, September 2023, cited by NEA 2024 as the right direction. AAP 2026 *Digital Ecosystems, Children, and Adolescents* (Pediatrics 157(2)): child-centered design versus engagement-based design; a dedicated AAP “AI in schools” policy does not yet exist.

that is not integral to the service. COPPA 2025 is the instrument. Advertising, consideration, and training or otherwise developing AI technologies are **not integral**. School consent does not cover them. Never. An education service that monetizes the child's chat as training data or as an ad graph has left the educational purpose.

Ten. No required student use of a tool the district has not contracted, DPA'd, and age-gated. FERPA, SOPPA, COPPA school-consent limits. A syllabus that says "use ChatGPT" without a district tenant, a DPA, and an age gate is a requirement to commit the incident in line 6. Consumer terms of service are not a school-official agreement. Parent permission on a personal account is not a DPA.

Eleven. No unpaid take-home that assumes every child has a premium personal agent. Pew 2025; Harvard 2024, 40 percent versus 20 percent; AFT equity; UNESCO's 2.6 billion offline, 1 in 4 primary schools without electricity, 60 percent of primary schools unconnected. Homework-gap 2.0 is a paid-tier gap. If the task needs a capable model, the school issues it. If the school cannot issue it, the task cannot require it. Grading the gap is not rigor.

Twelve. No vendor-owned pedagogy. Certified educators decide whether, when, and how. AFT Core Value 3: educators choose the technology; vendors serve, they do not drive. US ED 2023: involve educators; firmly reject replacement. UNESCO: human-centred; institutions validate tools for ethical and pedagogical fitness – not the vendor's safety blog. A product that arrives with a pacing guide, a companion persona, and a detector score as the misconduct process has set the pedagogy. Refuse it. The off-switch Chapter 7 argued for is this line in labor form. This line is the off-switch in procurement form.

Those twelve are the spec. A feature that cannot map to a row is out of scope or out of bounds. The fights that are not settled sit next to the spec, not inside it.

Open policy fights (as of 30 August 2026)

Leave the fights standing. A constraints chapter that pretends to close them has become a values chapter with better footnotes. Eight fights were live on the access date. This book can name their instruments. It cannot referee them.

Federal US vacuum versus state and local fill. Executive Order 14110, 30 October 2023, "Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence," directed the Secretary of Education to produce education-AI resources on a 365-day clock. It was **rescinded 20 January 2025** by Executive Order 14148.⁴³³ The 2023 Office of Educational Technology report and the 2024 toolkits remain the last full federal pedagogical frame; OET was reported shuttered in March 2025. Executive Order 14277, 23 April 2025, "Advancing Artificial Intelligence Education for American Youth," creates a Task Force, a Presidential AI Challenge, public-private K-12 resources, and a direction to prioritize AI in ESEA and HEA Title II teacher-training grants, NSF research, and USDA 4-H. That is an **adoption and workforce** order, not a child-safety regulation.⁴³⁴ AFT's

⁴³³Executive Order 14110, 30 October 2023, 88 FR 75191, <https://www.federalregister.gov/documents/2023/11/01/2023-24283/safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence>. Rescinded 20 January 2025 by Executive Order 14148.

⁴³⁴Executive Order 14277, 23 April 2025, "Advancing Artificial Intelligence Education for American Youth,"

March 2025 guardrails say so in a sentence the book should not paraphrase away: “the current federal administration has indicated it will not regulate AI.”⁴³⁵ The actual federal child-safety action on companion chatbots is at the FTC – the 6(b) study – not at Education. There is no US federal student-AI statute. Open question 17 remains standing. The fight is therefore the one AFT named as labor theory: states, boards, and unions writing the rules vendors will actually face. NY 2-d, Illinois SOPPA, California SOPIPA, Texas TRAIGA, Broward’s resolution, PGCPs’s addendum, NYCPS’s red list – those are the rules. EO 14277 will fund literacy. It will not write the red lines.

EU high-risk clock. Regulation (EU) 2024/1689, 13 June 2024, in force 1 August 2024. Article 5 prohibitions, including emotion recognition in educational institutions, and Article 4’s AI-literacy duty on deployers: **2 February 2025**, already applying. GPAI model duties: 2 August 2025. Transparency for chatbots – students must know they are talking to a machine – timed for August 2026. Annex III education systems – AI used to determine access, evaluate learning outcomes including steering the learning process, assign level, monitor and detect prohibited behavior during tests – originally applied from **2 August 2026**. The Digital Omnibus political agreement of 7 May 2026, enacted as Regulation (EU) 2026/1744 (OJ 24 July 2026, in force 27 July 2026), pushes stand-alone Annex III high-risk rules to **2 December 2027**, with some product-embedded high-risk later (2 August 2028). CSAM and non-consensual intimate-image generation: prohibition timed for December 2026 via the Omnibus.⁴³⁶ **Flag the original-versus-Omnibus date conflict** so the superseded 2026 clock cannot leak back as current law; treat 2 December 2027 as the enacted stand-alone date. Confirm consolidated EUR-Lex at typesetting. Prohibitions and the literacy duty already apply either way. The fight is whether delay is time to write standards or a gift to exam-scoring and proctoring vendors. A completed Annex III education enforcement record does not exist. Open question 18 remains standing.

Companion liability. Character.AI’s disclosed settlement versus OpenAI’s live *Raine* defense. The fight is product-liability for chatbots versus the First Amendment; age-gating and parental controls as adequate versus design defect. Conway’s 21 May 2025 ruling is the pleading-stage fact: no free-speech dismissal before discovery. It is not a merits holding. OpenAI’s “more than 100 times” redirection is a defense, not a finding. Common Sense’s degradation finding – guardrails

<https://www.whitehouse.gov/presidential-actions/2025/04/advancing-artificial-intelligence-education-for-american-youth/>. Task Force, Presidential AI Challenge, public-private K-12 resources, ESEA/HEA Title II teacher-training priority, NSF, USDA 4-H. Adoption and workforce, not child-safety regulation. US ED 2026 supplemental priority (91 FR 18774, effective 13 May 2026) is likewise adoption language, not a replacement mandate and not a child-safety regulation.

⁴³⁵AFT *Commonsense Guardrails* v2 (March 2025), as quoted in the responsible-use stream: “the current federal administration has indicated it will not regulate AI.” Weingarten on the Academy: “In the absence of federal guardrails, we become the ones who must ensure that this tool is used safely.”

⁴³⁶European Commission hub, “AI Act | Shaping Europe’s digital future,” page update 3 August 2026, <https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>. Original AI Act schedule: stand-alone Annex III high-risk (including education access, evaluation, test-monitoring) from 2 August 2026. Digital Omnibus political agreement 7 May 2026; enacted as Regulation (EU) 2026/1744, OJ 24 July 2026, in force 27 July 2026, pushing stand-alone Annex III high-risk rules to **2 December 2027** (Annex I product-embedded high-risk later, 2 August 2028). This mid-pass treats 2 December 2027 as the enacted date and keeps the superseded 2026 date on the page so it cannot leak back as current law. Confirm the consolidated EUR-Lex text of 2024/1689 as amended by 2026/1744 before typesetting. CSAM / non-consensual intimate-image generation prohibition timed December 2026 via the Omnibus. Transparency (chatbot disclosure) timed August 2026. Art. 5 prohibitions and Art. 4 literacy already apply from 2 February 2025. No completed Annex III education enforcement record.

fail in long conversations – is the audit sitting next to both dockets. This book will not close a live case. Open question 15 remains standing. The red line in 8.5 does not wait for a verdict. It maps to the audit, the 6(b) study, the UK standards, and the union resolution.

Detection versus redesign. JCQ still lists Turnitin and GPTZero as optional checks. Singapore’s research universities are abandoning scores. Liang 61.22 percent and the 2026 false-negative flip are the evidence that the score is not a fact. QAA 2023 and Chapter 5’s two-lane rule are the redesign. The fight is whether any remaining detector use is due process or discrimination against multilingual and neurodivergent writers. A district that keeps the score as one check, never sole, with process artefacts, is inside JCQ. A district that suspends on the score is inside Texas A&M-Commerce. Those are not the same policy.

Union-vendor academies. Cross-reference Chapter 7.5. AFT’s National Academy for AI Instruction, resolution 22 July 2025: \$23 million, Microsoft lead plus OpenAI plus Anthropic, UFT founding partner, five-year target of 400,000 educators – a plan, not a trained headcount. NEA took a Microsoft Elevate grant for AI micro-credentials. Weingarten: “In the absence of federal guardrails, we become the ones who must ensure that this tool is used safely.” Union as regulator of last resort, funded by the firms whose DPAs, models, and safety failures are the subject.⁴³⁷ Open question 16 remains standing: capture or civilize. This chapter does not answer it. Line 12 still holds either way. Educators decide whether, when, and how. A syllabus paid for by the vendor does not repeal Core Value 3. It tests it.

National platform versus district procurement versus student-owned local agents. Singapore’s Student Learning Space and Estonia’s ITI are state-shaped pedagogy on a vendor substrate: age bands, no direct answers, teacher supervision, national licence. The United States and the United Kingdom default to thousands of district contracts; TeachAI’s toolkit is the most widely copied US template; UK DfE prefers teacher-facing, with pupil-facing only under KCSIE plus product-safety standards. Local open weights are a real technical option and **not** a child-safety certified school product. Hermes Agent (MIT licence; vendor homepage February 2026 without a day; 25 February is the responsible-use first-release date) installs locally. Open-weight Hermes models can be served on llama.cpp, MLX, vLLM, Ollama, or LM Studio. Still: computer-use and web tools re-open prompt injection and exfiltration; there is no DSL alerting path unless the school runs the server; equity requires every family to own a capable machine.⁴³⁸ Who is the data controller of a child’s agent memory is the live governance question the red lines force. Chapter 9 is that question at full length. The preview this chapter can give is already in the sources, even when they do not use

⁴³⁷AFT, “On the Formation of a National Academy for AI Instruction,” resolution adopted 22 July 2025, <https://www.aft.org/resolution/formation-national-academy-ai-instruction>. \$23 million; Microsoft, OpenAI, Anthropic; UFT founding partner; five-year target 400,000 educators / about 7.2 million students. Plan/target, not trained headcount. Do not mash with CNBC’s ChatGPT for Teachers first-cohort figure of about 150,000 educators (19 November 2025). NEA Microsoft Elevate grant for AI micro-credentials: same conflict class; dollar amount not stated in the responsible-use memo. Open question 16 left standing.

⁴³⁸Nous Research, Hermes Agent, MIT license, <https://hermes-agent.nousresearch.com/docs/>. Vendor homepage (hermes-agent.org) dates the launch to February 2026 without a day; 25 February is the responsible-use first-release date. The agent installs locally; open-weight models can be served on llama.cpp, MLX, vLLM, Ollama, or LM Studio. Computer use and messaging gateways. **Not** a COPPA-certified, KCSIE-aligned school product. Singapore MOE page updated 31 July 2026, <https://www.moe.gov.sg/education-in-sg/educational-technology-journey/edtech-masterplan/artificial-intelligence-in-education>. TeachAI toolkit: <https://www.teachai.org/toolkit-guidance>.

the word “agent”: default host is the school or a public platform under public-law duties; teacher-facing agents may be more private-cloud or local, because the risk is IP and workload, not a child’s attachment; student-owned local agents are a **home** choice with the status of a privately hired tutor – not a school-required tool, not fed with education records, not a substitute for assessment evidence; never the consumer cloud with training-on, companion persona, and no contract.

Age floors. UNESCO 2023 proposed 13 for independent generative-AI conversation and said under-13 tools are not appropriate for primary classrooms. Singapore: no student chatbots P1-P3; P4-6 only education-designed tools under teacher supervision (MOE page updated 31 July 2026; parliamentary reply 6 May 2026). OpenAI consumer terms: 13+ with parent permission – which Common Sense notes teens bypass by lying.⁴³⁹ Broward: zero student-facing through grade 5. AFT: companion ban at least until 16; oppose student-facing AI in elementary. Claude consumer: 18+. Character.AI open chat: 18+ after November 2025. Thirteen is a terms-of-service number, not a developmental one. Eighteen is a rights-transfer number under FERPA, not a finding that a seventeen-year-old is safe with a companion. Design to the stricter of law and development. The fight is whether 13 remains the global hinge or whether companion-class products force a higher floor. Line 4 already acts as if they do.

Global South compute versus textbook. UNESCO: about 60 percent of primary schools unconnected; about 1 in 4 without electricity; about 2.6 billion people offline. Do not divert education budgets into AI while schools remain unconnected. Do not paper over teacher shortages. Training corpora reflect Global North norms. “AI for all” as leapfrog versus AI as a new dependency on Northern models trained on Northern data is the fight. Estonia and Singapore can buy the wrapper. A ministry that cannot staff teachers cannot staff a DSL. Line 11 and clause (d) of the equity implication are the portable version. They do not close the fight. They stop a district in a connected country from pretending that “access” means a free companion.

Who hosts the student agent is the live governance question the red lines force. Chapter 9.

What this chapter did not claim, restated so a later chapter cannot smuggle it back. Causal findings in Setzer and Raine: litigation, not verdicts. Whether any current detector is good enough on 2026 models: published independent tests say no. A completed EU Annex III education enforcement record: there is none. The original 2 August 2026 high-risk clock was superseded by Regulation (EU) 2026/1744; stand-alone Annex III obligations now apply from 2 December 2027. Confirm consolidated EUR-Lex at typesetting. Environmental and energy costs: flagged, analyses **not fetched**; no kWh, water, or carbon numbers. CSM ratings of Grok are ratings of the chat product, not of Grok Bot. “We don’t train on your data” is a product-tier claim, not a legal status. OECD Digital Education Outlook 2026 full-report statistics, including TALIS 37 percent of teachers using GenAI for work: **UNVERIFIED** at full-PDF level. No thirteenth red line. Open questions 15, 17, 18, and 19 remain standing, with 16 carried from Chapter 7 because the Academy fight is not this chapter’s to close.

The table is the spec. The fights are the weather. A board that waits for the weather to clear before

⁴³⁹UNESCO 2023 guidance as in note 47. Singapore: P1-3 no student chatbots; P4-6 only MOE-designed SLS features, teacher-supervised; parliamentary reply 6 May 2026. OpenAI consumer ChatGPT: 13+ with parent/guardian permission if under 18; 2026 age-prediction / ChatGPT for Teens. CSM: teens bypass by lying. Age gates are a floor, not a pedagogy.

adopting the spec has waited for a federal statute that does not exist, a verdict that is not in, a detector that independent tests have already failed, and an energy number this research pass did not fetch. The instruments are already on the page. Map to them, or do not ship the agent.

Chapter 9

Who Hosts the Student Agent



Figure 10: A laptop, a small host box with a light, keys, and a wax seal.

Who is the data controller of a child's agent memory? That is the live governance question this chapter is for. Open-source does not settle it. A local runtime does not settle it. A ministry slogan about AI for every student does not settle it. The comparison is three hosting models, not two, and the never-column is a fourth.

	Cloud GPT (Edu/Enterprise)	National wrapper (SLS, ITI)	Local open (Hermes + open-weight)
Data training default	Contractually off on paid edu tiers	Depends on the underlying API contract	Never leaves the box if configured local
Child-safety stack	Vendor + whatever the school adds	Designed in (Estonia: no direct answers; SG: age bands)	You build it. Default Hermes is an adult autonomous agent with computer use
DSL / parent visibility	Possible via admin logs	Built for schools	Absent unless the school runs the server
Prompt injection / web Equity	Vendor + school filter	Can disable browsing	High if tools enabled
	Licence cost; still needs devices	State pays	Shifts cost to home GPU/RAM — new gap
Pedagogy control	Prompt policies, GPTs, “study mode”	Highest (curriculum-aligned)	Highest in theory; zero in practice without teacher skill
Who is controller?	School, if DPA exists	Ministry / school	Student or parent — FERPA/COPPA problem if the school required it

That table is a compression of fetched instruments, labelled as synthesis, not a vendor scorecard and not a district SOP.⁴⁴⁰ Never: consumer cloud with training-on, companion persona, and no contract. The rest of the chapter is what the table forces.

The popular shortcut will try to make this a licensing argument. Run the weights at home and the child is safe, because the bytes never leave the laptop. Run them in a national cloud and the ministry is the nanny. Buy ChatGPT Edu and FERPA is handled. None of those sentences survives the instruments. Hermes Agent, released under an MIT license in February 2026 (hermes-agent.org dates the launch to February without a day; 25 February is the first-release date recorded in this book’s responsible-use stream), plus open-weight Hermes models. The agent installs locally — Linux, macOS, WSL2; local vLLM among other backends. Open-weight Hermes models can be served on llama.cpp, MLX, vLLM, Ollama, or LM Studio. That is a real local/open option. It is not a child-safety certified school product. Computer-use and web tools re-open prompt injection and exfiltration. There is no designated-safeguarding-lead alerting path unless the school runs the server. Equity requires every family to own a capable machine — a new gap, not a closed one. Giving every fifteen-year-old an autonomous agent with a browser is how you recreate Character.AI plus data exfiltration on the student’s own disk.⁴⁴¹

⁴⁴⁰Comparison table compressed from the responsible-use memo, section 6 (open-source versus closed). Synthesis of fetched instruments, not a vendor scorecard and not a district SOP.

⁴⁴¹Responsible-use finding 12 and section 6; book brief, Chapter 9 story. Hermes Agent: MIT license; hermes-

Sources converge, even when they do not use the word “agent.” Default host of a student agent is the institution or a public platform under public-law duties: Singapore’s Student Learning Space, Estonia’s ITI, a district tenant of ChatGPT Edu or Gemini Education or Claude for Teachers with a real data-processing agreement. The institution is the data controller. The teacher is in the loop. The designated safeguarding lead gets the self-harm ping. Teacher-facing agents are the 2026 beachhead, because the risk profile is intellectual property and workload, not a child’s attachment. Student-owned Hermes on a laptop is a home tutor, not a school requirement. Local open-source does not by itself create a DSL alert path or close an equity gap. The age-gated agent ladder at the end of this chapter is this book’s proposed integration model, inferred from legal, developmental, and evidence constraints. It is not a documented district SOP. Label it as such on the page, or it will be quoted as if Broward had already adopted it.

Hermes 4 is not Hermes Agent. El Salvador is not Grok Bot. Those two collapses, already retired in Chapter 2, would wreck this chapter if they leaked back in. The model family is an open-weight hybrid reasoning LLM. The agent is a runtime. El Salvador, announced 11 December 2025, is Grok-as-tutor across more than 5,000 public schools, with no outcome evaluation fetched as of 30 August 2026. Grok Bot, launched 11 August 2026 and expanded 26 August 2026, is a Cursor-authenticated cloud teammate with no education SKU.⁴⁴² Product class is the first control. Host is the second. This chapter is the second.

Open-source does not settle governance

Local weights are not a school safety certification. That is the sentence this section exists to keep from leaking into later pages as a finding.

Nous Research’s own documents separate the objects. Hermes 4 is an open-weight hybrid reasoning LLM family, not an agent runtime. The Hugging Face model card for Hermes 4 14B, on a Qwen 3 14B base, describes hybrid-mode reasoning with optional think traces, function calling via tool-call tags, and a post-training corpus described as about five million samples and about sixty billion tokens against Hermes 3’s one million and 1.2 billion. Weights are public. The technical report is arXiv 2508.18255, Teknum and colleagues, 25 August 2025. Education-relevant capabilities listed on the model card — math, code, STEM, logic, long-form writing, tool calling, structured output — are model skills, not a classroom product. Nous Portal docs state that Hermes 4 at 70B and 405B is “not recommended for use inside Hermes Agent” because it is “tuned for chat and reasoning, not the rapid-fire tool-calling loop the agent relies on.” Recommended agent models in that same document include Claude Sonnet 4.6, GPT-5.5 Pro, Gemini 3 Pro Preview, and DeepSeek V4 Pro.⁴⁴³ Treating the model family as a proxy for the agent, or the agent as a proxy for the model

agent.org dates the launch to February 2026 without a day; 25 February is the first-release date recorded in this book’s responsible-use stream. Character.AI comparison is a harm-class argument, not a claim that Hermes is Character.AI.

⁴⁴²Landscape memo, sections 2 and 3. Hermes 4 is the model family (arXiv 2508.18255, 25 August 2025). Hermes Agent is the runtime. xAI, El Salvador partnership, 11 December 2025, <https://x.ai/news/el-salvador-partnership>; more than 5,000 public schools; over one million students; Grok the assistant, not Grok Bot; no outcomes fetched. xAI, “Introducing Grok Bot,” 11 August 2026, <https://x.ai/news/introducing-grok-bot>; “Grok Bot more plans,” 26 August 2026, <https://x.ai/news/grok-bot-more-plans>. No education SKU found on product pages, launch post, or expansion post.

⁴⁴³Hugging Face, Hermes-4-14B-FP8 model card, <https://huggingface.co/NousResearch/Hermes-4-14B-FP8>. arXiv

family, is a factual error. It is also the error that lets a board member hear “open Hermes” and picture a Socratic tutor with a MIT license stamped on the box.

Hermes Agent is the runtime. It is an open-source, MIT-licensed, self-hosted autonomous agent released by Nous Research in February 2026. hermes-agent.org dates the launch to February without a day; 25 February is the first-release date recorded in this book’s responsible-use stream. Install via curl. Linux, macOS, WSL2. Optional desktop installer. Model-agnostic: Nous Portal, Open-Router, OpenAI-compatible endpoints, local vLLM. The product line is explicit: “Not a chatbot. Not a copilot. An agent that lives on your machine.” Documented capabilities, from product docs, not from education evaluations: persistent memory across sessions, including Honcho dialectic user modeling and FTS5 session search; autonomous skill creation after complex tasks, skills compatible with agentskills.io; cron and scheduled automations with delivery to messaging platforms; parallel isolated sub-agents; terminal backends including local, Docker, SSH, Singularity, Modal, and Daytona; browser automation, web search and extract, vision, image generation, text-to-speech; a messaging gateway covering Telegram, Discord, Slack, WhatsApp, Signal, Teams, and a long tail of others; Bot Mode, with named bots holding their own model, memory, skills, and routines; a research and MLOps loop built by model trainers; a bundled research skill for end-to-end ML paper writing, literature review through experiments through LaTeX through submission.⁴⁴⁴ Privacy posture, as claimed by the vendor: no telemetry; data in `~/ .hermes/`; MIT license; container hardening options. Vendor privacy claims are not a DPA. They are a configuration story about a machine the school does not control.

None of that is a school product. No official Nous materials, university newsrooms, or policy documents fetched for this book document school or university use of Hermes Agent. A GitHub project describing a personal Hermes-based learning system on Telegram is a third-party hobby project, not a documented institutional deployment, and it will not be cited as one.⁴⁴⁵ Mapping the research-paper skill onto graduate workflows — literature monitoring, experiment babysitting, paper drafting — is a natural inference and is the same class of work OpenAI and Anthropic pitch to researchers. Mapping is not adoption. No K-12 DPA was found. No classroom SOP was found. No institutional acceptable-use policy naming Hermes Agent was found. Open question 14 in this book’s brief is that absence, left standing.⁴⁴⁶

What the runtime actually does, if you take the docs at their word, is the opposite of a constrained education tutor. It has a computer. It browses. It runs cron. It holds memory across sessions. It can message the child on Telegram. Default Hermes is an adult autonomous agent with computer use.

2508.18255, Teknum et al., 25 August 2025; HTML conversion failed on fetch; abstract/metadata confirmed via Hugging Face citation block. Nous Portal integration docs: Hermes 4 (70B / 405B) “not recommended for use inside Hermes Agent”; recommended agent models include Claude Sonnet 4.6, GPT-5.5 Pro, Gemini 3 Pro Preview, DeepSeek V4 Pro. <https://hermes-agent.nousresearch.com/docs/integrations/nous-portal>.

⁴⁴⁴<https://hermes-agent.org/>; official docs <https://hermes-agent.nousresearch.com/docs/>. Capabilities listed are product-doc claims, not education evaluations. Research-paper-writing skill: <https://hermes-agent.nousresearch.com/docs/user-guide/skills/bundled/research/research-research-paper-writing>. Vendor privacy posture (no telemetry; data in `~/ .hermes/`) is a vendor claim, not a DPA.

⁴⁴⁵Landscape memo, section 2.3. GitHub project AlijonovMukhammaddiyor/hermes-university: third-party hobby project, not institutional use. Do not cite as school or university deployment.

⁴⁴⁶Open question 14 in the book brief: no Hermes Agent or Grok Bot institutional AUP, DPA, or classroom SOP found.

Community education profiles exist — LM Studio defaults, PII-redaction experiments — and they are community profiles, not COPPA certifications and not alignments to *Keeping Children Safe in Education*. Computer-use and web tools re-open prompt injection and exfiltration. OWASP’s LLM risk taxonomy treats prompt injection and insecure output handling as top-class failures. Any student agent with web or tool access can be steered into exfiltrating a prior student’s essay, IEP notes in context, or API keys. No independent, education-specific audit of a deployed school agent with open-web tools was found as of 30 August 2026. That absence is itself a finding.⁴⁴⁷

The designated safeguarding lead is the operational hole the local-weights slogan cannot see. Khan Academy told *60 Minutes* that Khanmigo flags self-harm, harm-to-others, and “inappropriate story” attempts to teachers, and that OpenAI is contractually barred from training on Khanmigo student and teacher data.⁴⁴⁸ That is a product with a teacher on the alert path. UK Department for Education product-safety standards, updated 19 January 2026, require filtering that holds for the whole conversation, jailbreak resistance, no anthropomorphism, no emotional-dependence design, hard time limits teachers control, and a mental-health crisis pathway that always hands off to humans.⁴⁴⁹ A Hermes instance in ~/ .hermes/ on a teenager’s laptop has none of that unless someone builds it. Local storage of a child’s companion-adjacent memory is not a safety feature. It is a log the DSL cannot see and a corpus a parent may not know how to delete.

Equity is the other hole. A personal agent — persistent memory, tools, custom skills — is more unequal than a chatbot. It is closer to a private tutor plus a research assistant plus a secretary. Pew, 15 January 2025: 26 percent of U.S. teens aged 13 to 17 have used ChatGPT for schoolwork, double the 2023 share. Awareness jumped most in the lowest-income households. Access to the idea of the tool is democratising. Access to the good tool is not. A 2024 Harvard undergraduate survey: among AI users, 40 percent of students with no financial aid paid for AI products versus 20 percent of students on aid.⁴⁵⁰ Estonia’s AI Leap exists in part because 64 to 90 percent of students were already using free commercial bots to finish homework faster — the pattern research says deskills.⁴⁵¹ UNESCO, still using 2024 figures in later rights-of-learners materials: about 2.6

⁴⁴⁷Responsible-use memo, section 1 (prompt injection, answer-leakage, open-web agents). OWASP LLM risk taxonomy: prompt injection and insecure output handling as top-class failures. No independent, education-specific audit of a deployed school agent with open-web tools found as of 30 August 2026. ACL 2026 Zeng et al. line of work on LLM tutors: pedagogically prompted tutors leak full answers across turns; a fine-tuned adversarial-student agent pushed one 7B tutor’s leakage from 14 percent to 82 percent — answer-leakage, a cousin of prompt injection, not an audit of Hermes.

⁴⁴⁸Khan Academy comments as reported on CBS *60 Minutes* (story originally 8 December 2024), recorded in the responsible-use memo, section 1. Teacher-visible flags; OpenAI contractually barred from training on Khanmigo student/teacher data. Vendor/journalism, not an independent audit.

⁴⁴⁹UK Department for Education, *Generative AI: product safety standards*, first published 22 January 2025, updated 19 January 2026, <https://www.gov.uk/government/publications/generative-ai-product-safety-standards/generative-ai-product-safety-standards>. Filtering must hold for the whole conversation; jailbreak resistance; no anthropomorphism; no emotional-dependence design; hard time limits teachers control; mental-health pathway that hands off to humans.

⁴⁵⁰Pew Research Center, “About a quarter of U.S. teens have used ChatGPT for schoolwork,” 15 January 2025, <https://www.pewresearch.org/short-reads/2025/01/15/about-a-quarter-of-us-teens-have-used-chatgpt-for-schoolwork-double-the-share-in-2023/>. Twenty-six percent of U.S. teens 13-17; double 2023. Harvard undergraduate survey, arXiv:2406.00833 (2024): among AI users, 40 percent of students with no financial aid paid for AI products versus 20 percent of students on aid.

⁴⁵¹Estonia AI Leap problem statement, as recorded in the responsible-use memo, finding 8, and brief Chapter 9: 64-90 percent of students already using free commercial bots to finish homework faster.

billion people lack internet; about one in four primary schools lack electricity; about 60 percent of primary schools globally remain unconnected.⁴⁵² Local open-source that “runs on your machine” assumes a machine. A capable local model assumes GPU or at least enough RAM to make the open-weight story real. Shifting cost to home GPU and RAM is a new gap, not a closed one. The red line already on the table from Chapter 8 is the operational translation: no unpaid take-home that assumes every child has a premium personal agent the school does not issue.

Grok Bot is not the local-open alternative, and it is not an education product. Early beta, 11 August 2026, joint xAI and Cursor, sign-in via Cursor account, expanded 26 August 2026 to SuperGrok, SuperGrok Plus, SuperGrok Heavy, Cursor Pro / Pro+ / Ultra, and Cursor Teams Standard and Premium. Persistent cloud computer. Sign into apps and websites. Work 24/7 in parallel. Routines learned by watching the user. Jobs illustrated in the 11 August launch post and the 26 August expansion post: sales outbound, CRM hygiene, inbox, expenses, bug reproduction, website deploy, meeting stand-in. None of those examples is educational. Product pages, launch post, and expansion post do not mention schools, universities, students, or teachers. Education positioning of Grok Bot: none found. The FAQ at docs.x.ai/grok-bot/faq, re-fetched after the landscape timeout: cloud data storage required; Legacy Privacy Mode not supported; training opt-out follows Cursor account settings. FAQ does not mention schools, students, or teachers. Still no education SKU, DPA, or classroom SOP.⁴⁵³ Common Sense Media / Youth AI Safety Institute rated Grok, the assistant, **unacceptable** on 22 January 2026.⁴⁵⁴ That rating is of the chat product, not of Grok Bot. Do not launder it into a Bot finding. Do not launder Bot into a school finding either. There is no education SKU to launder.

El Salvador remains the other collapse. xAI announced, on 11 December 2025, a partnership with the Government of El Salvador to deploy Grok across more than 5,000 public schools over two years, “personalized learning to over one million students,” framed as adaptive, curriculum-aligned tutoring. No outcome evaluation, independent audit, or classroom study was found. Announced national program, not a documented completed deployment as of 30 August 2026. Grok-as-tutor, not Grok Bot-as-agent.⁴⁵⁵ Using El Salvador as proof that student agents are already school policy is the Chapter 2 error with a hosting twist. Using a MIT-licensed local agent as proof that El Salvador is unnecessary is the same error in the other direction. Different products. Different hosts. Neither is a child-safety certified U.S. school product.

Three hosting models, then, not a binary of cloud versus home. The table already named them. The next section spends them.

⁴⁵²UNESCO rights-of-learners materials, still using 2024 connectivity figures: about 2.6 billion people offline; about 1 in 4 primary schools without electricity; about 60 percent of primary schools unconnected. UNESCO, *AI and education: Protecting the rights of learners*, <https://www.unesco.org/en/articles/ai-and-education-protecting-rights-learners>.

⁴⁵³xAI Grok Bot product pages cited in note 3. FAQ at docs.x.ai/grok-bot/faq, re-fetched after the landscape timeout: requires cloud data storage; Legacy Privacy Mode not supported; training opt-out follows Cursor account settings. FAQ does not mention schools, students, or teachers. Remaining hole is institutional use (no education SKU, DPA, or classroom SOP), not the FAQ timeout.

⁴⁵⁴Common Sense Media / Youth AI Safety Institute, Grok risk assessment, 22 January 2026, <https://institute.common sense media.org/risk-assessments>. Rating is of Grok the assistant, not of Grok Bot. Do not launder.

⁴⁵⁵xAI El Salvador post, 11 December 2025, as in note 3. CoinDesk same-day coverage restates the announcement. Open question 13: program status as of August 2026; only the announcement found; no outcomes.

Three models: cloud GPT Edu, national wrapper, local open

The table is the argument. The columns are not vendors. They are control surfaces.

Cloud GPT in an Edu or Enterprise tenant is the U.S. default that districts can actually buy. Training is contractually off on paid education tiers, as a product-tier claim: OpenAI states that ChatGPT Edu, Enterprise, Teachers, and the API do not train on organizational data by default.⁴⁵⁶ FERPA's school-official exception, 34 C.F.R. section 99.31(a)(1), still requires a written agreement that puts the vendor under the school's direct control with respect to the use and maintenance of the records, and that limits use to the educational purpose the school authorized. Redisclosure is limited. There is no product that is "FERPA compliant" in the abstract. There is a contract, a configuration, and a practice. Consumer ChatGPT, Claude, and Gemini accounts cannot satisfy that exception.⁴⁵⁷ COPPA's 2025 amendments treat AI-model training as not integral to the service: it needs separate verifiable parental consent. The school-consent exception still exists, but only for the educational purpose the school authorized — not product improvement, not ads, not model training.⁴⁵⁸ Child-safety on this column is vendor plus whatever the school adds. Common Sense Media rated ChatGPT-5 **high risk** for teens on 23 October 2025: parental controls exist; teens use it for emotional support anyway; guardrails degrade in long conversations; "Teens should not use ChatGPT for mental health or emotional support." Gemini K-12 was rated **high risk** on 9 January 2026: privacy better; age-appropriate mental-health handling not. Khanmigo, a purpose-built tutor rather than a GPT Edu tenant, was rated **low risk** on 8 August 2024, with the caveat that supervising adults see every turn.⁴⁵⁹ DSL and parent visibility are possible via admin logs if the tenant is configured for them. Prompt injection and live web are vendor plus school filter. Equity is a licence cost that still needs devices. Pedagogy control is prompt policies, custom GPTs, "study mode" — real, and thinner than a national curriculum wrapper. Who is the controller? The school, if a DPA exists. If the DPA does not exist, this column collapses into consumer cloud, which is the never-column.

A national wrapper is the model Singapore and Estonia actually built. Data-training default depends on the underlying API contract, because the wrapper sits on someone else's model. Child-safety is designed in: Estonia's ITI does not give direct answers; Singapore's Student Learning Space is age-banded. The 6 May 2026 parliamentary reply, already cited for the P4-6 education-designed-tools rule, also states that SLS's Learning Assistant redirects students who "want to be spoon-fed with direct answers" — a designed tutor refusal, not an agent evaluation. DSL and parent visibility are built for schools. Browsing can be disabled. The state pays, which is the equity move the U.S.

⁴⁵⁶OpenAI business-privacy page, as recorded in the responsible-use memo: ChatGPT Edu / Enterprise / Teachers and the API do not train on organizational data by default. Product-tier claim, not a legal status.

⁴⁵⁷FERPA, 20 U.S.C. section 1232g; 34 C.F.R. Part 99, especially section 99.31(a)(1) (school-official exception) and section 99.33 (redisclosure). Student Privacy Policy Office vendor FAQ: https://studentprivacy.ed.gov/sites/default/files/resource_document/file/Vendor%20FAQ.pdf. No product is "FERPA compliant" in the abstract.

⁴⁵⁸COPPA, 15 U.S.C. sections 6501-6506; 16 C.F.R. Part 312. 2025 Rule amendments, *Federal Register* 22 April 2025; effective 23 June 2025; most operators until 22 April 2026. FTC: disclosure "for monetary or other consideration, for advertising purposes, or to train or otherwise develop artificial intelligence technologies" is not integral. School-consent exception limited to the educational purpose the school authorized.

⁴⁵⁹Common Sense Media / Youth AI Safety Institute, accessed 30 August 2026: ChatGPT-5, high risk, 23 October 2025; Gemini K-12, high risk, 9 January 2026; Khanmigo, low risk, 8 August 2024 ("there is no student privacy on Khanmigo" in the sense that supervising adults see every turn).

homework-gap 2.0 does not make. Pedagogy control is the highest of the three columns, because the tool is curriculum-aligned by construction. Who is the controller? The ministry or the school, under public-law duties, not a parent with a curl installer. This column is tutors, not agents. Do not claim Singapore or Estonia evaluated autonomous agents. They evaluated, or are evaluating, Socratic, education-designed tutors and teacher licences. That limitation is load-bearing. It is also why they are the cleanest counter-examples in the corpus to “just give them ChatGPT.”

Local open — Hermes Agent plus open-weight models — is the column the open-source shortcut treats as a solution. Data never leaves the box if configured local. That sentence is true and insufficient. The child-safety stack is yours to build. Default Hermes is an adult autonomous agent with computer use. DSL and parent visibility are absent unless the school runs the server. Prompt injection and web risk are high if tools are enabled. Equity shifts cost to home GPU and RAM. Pedagogy control is highest in theory and zero in practice without teacher skill, because there is no curriculum wrapper, no attempt gate, no hint ladder, and no teacher-authored misconception library unless a teacher builds them. Who is the controller? The student or the parent. If the school required the runtime, FERPA and COPPA problems land on a parent-installed agent: the school has required a tool it does not control, fed or not fed with education records, with a memory the designated safeguarding lead cannot see. That is open question 22, left standing: who is the data controller of a child’s persistent agent memory when the runtime is local, open-source, and parent-installed.⁴⁶⁰

Never is not a fourth column in the table because it is not a hosting model. It is a refusal. Consumer cloud with training-on, companion persona, and no contract. Character.AI was rated **unacceptable** on 10 April 2025. Social AI companions as a class, **unacceptable**, 16 July 2025. Meta AI, **unacceptable**, 15 August 2025. Grok the assistant, **unacceptable**, 22 January 2026. AI mental-health apps and chatbots for mental-health support, **unacceptable** or “risk varies,” with leading bots failing to recognize the spectrum of youth mental-health conditions.⁴⁶¹ Sewell Setzer III, 14, died by suicide on 28 February 2024; his mother sued in October 2024, alleging prolonged, sexualized conversation with a Character.AI bot; a federal judge on 21 May 2025 refused to dismiss on First Amendment grounds; a mediated settlement was disclosed in a court filing on 7 January 2026. The parents of Adam Raine, 16, sued OpenAI on 26 August 2025, alleging ChatGPT coached his suicide after he began using it for schoolwork. These are allegations in litigation, not adjudicated findings of causation. They are the factual basis on which regulators are now acting. The FTC issued 6(b) study orders on 11 September 2025 to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI.⁴⁶² The never-column is not a taste. It is the product class the lawsuits, the FTC study, the UK emotional-development standards, and AFT’s companion ban until age 16 are already about.

⁴⁶⁰ Open question 22 in the book brief.

⁴⁶¹ Common Sense Media / Youth AI Safety Institute: Character.AI, unacceptable, 10 April 2025; social AI companions as a class, unacceptable, 16 July 2025; Meta AI, unacceptable, 15 August 2025; Grok, unacceptable, 22 January 2026; AI mental-health apps/chatbots, unacceptable or “risk varies.”

⁴⁶² *Garcia v. Character Technologies*: Sewell Setzer III, 14; died 28 February 2024; suit October 2024; Judge Anne Conway (M.D. Fla.), 21 May 2025, refused First Amendment dismissal at the pleading stage; settlement disclosed 7 January 2026. *Raine v. OpenAI* (San Francisco Super. Ct., filed 26 August 2025): Adam Raine, 16; started with schoolwork; alleged suicide coaching. Allegations in litigation, not adjudicated causation. FTC 6(b) study orders, 11 September 2025, to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI. Study, not enforcement. <https://www.ftc.gov/news-events/news/press-releases/2025/09/ftc-launches-inquiry-ai-chatbots-acting-companions>.

Prince George’s County Public Schools’ AI Declaration Addendum is what a serious U.S. district now asks before any of the first column is allowed to touch a child: name the models; say what is local versus cloud; certify FERPA, COPPA, and state-law confidentiality; no retention of student data or PII in generative models; no use beyond the stated educational purpose without district and parental approval; bias and harm audit on request; list of non-U.S. subprocessors; how to turn AI features off.⁴⁶³ Illinois SOPPA, 105 ILCS 85, requires a written DPA before any covered information moves to an operator. The National Data Privacy Agreement v2 is the de-facto interstate template. “Will you sign the NDPA?” is the first screen. It is a contract floor, not a statute, and it does not reach a parent-installed Hermes.

Three models, one refusal, and a question the table cannot close by itself: who hosts. The sources that do not use the word “agent” still answer it. The default is the institution.

Default host: the institution or a public platform

Singapore and Estonia are the cleanest national counter-examples in the corpus. They are tutors, not agents. That limitation has to stay on the page or the chapter will be quoted as if a Socratic national app were a Hermes deployment.

Singapore’s Ministry of Education page, last updated 31 July 2026, puts AI under the National AI Strategy and EdTech Masterplan 2030. Pedagogy first: Singapore Teaching Practice. Age bands are the design. Primary 1 to 3: print-first, no student chatbots. Primary 4 to 6: only MOE-designed Student Learning Space features, teacher-supervised. A parliamentary reply on 6 May 2026 restated the P4-6 rule: education-designed tools only. Secondary and pre-university: increasing student ownership plus integrity rules — evaluate AI output, check accuracy, acknowledge use. AI literacy in character and citizenship education, including deepfakes and validating generative AI, from 2026. Code for Fun updated 2027. Institutes of higher learning, baseline AI competencies from academic year 2027. SG-LEADS, with A*STAR, data collection from 2027; short-cycle school research in parallel.⁴⁶⁴ That is a public-law host with age bands, a national platform, teacher supervision in the years COPPA would still bind a U.S. operator, and a research program that has not yet claimed an autonomous-agent result. UNESCO’s 2023 guidance proposed thirteen for independent generative-AI conversation and said under-thirteen tools are not appropriate for primary classrooms.⁴⁶⁵ Singapore’s P1-3 rule is stricter than UNESCO’s thirteen. The U.S. consumer ToS

⁴⁶³Prince George’s County Public Schools, Artificial Intelligence Declaration Addendum, <https://epi.pgcps.org/contentassets/0d463a5c451f40fcba26c2796d7a59c8/appendix-p-ai-declaration-addendum.pdf>. Illinois SOPPA, 105 ILCS 85. Student Data Privacy Consortium, National Data Privacy Agreement v2 (2024): contract floor, not a statute.

⁴⁶⁴Singapore Ministry of Education, “Artificial intelligence in education,” page updated 31 July 2026, <https://www.moe.gov.sg/education-in-sg/educational-technology-journey/edtech-masterplan/artificial-intelligence-in-education>. Parliamentary reply 6 May 2026: education-designed tools only for P4-6; LEA “will redirect students back on track if they veer off-topic... or if they want to be spoon-fed with direct answers.” <https://www.moe.gov.sg/news/parliamentary-replies/20260506-ai-usage-in-schools>. Do not claim Singapore evaluated autonomous agents.

⁴⁶⁵UNESCO, *Guidance for generative AI in education and research* (Miao and Holmes), 7 September 2023, <https://www.unesco.org/en/articles/guidance-generative-ai-education-and-research>. Age thirteen for independent GenAI conversation; under-thirteen / primary use of untested tools “not appropriate.”

number is thirteen-plus with parent permission, which Common Sense notes teens bypass by lying.⁴⁶⁶ Thirteen is a terms-of-service number, not a developmental one. Singapore treated age as a design specification and put the tool inside a platform the ministry hosts.

Estonia is the cleaner interrupt. AI Leap, TI-Hüpe, announced February 2025 by President Alar Karis. Foundation plus Ministry of Education and Research. Fifty-fifty public-private funding, about four million euros in the pilot year. In 2025/26: 154 of 156 upper secondary schools. About 4,700 to 5,000 teachers get paid ChatGPT or Gemini. About 20,000 grade 10-11 students get ITI, a Socratic Estonian-language app that does not give answers, developed with Estonian researchers and OpenAI. University of Tartu impact study, with Stanford and OpenAI. First-year review published 7 August 2026. The review page itself states that first-year usage data “does not yet provide a basis for claiming that ITI has improved learning outcomes”; a University of Tartu study headed by Jaan Aru, involving more than 9,000 students, was expected to report later in 2026. Treat learning-outcome coefficients as unpublished. Open question 21, left standing.⁴⁶⁷ The problem statement is the load-bearing fact: 64 to 90 percent of students were already using free commercial apps to finish homework faster, and that pattern is what the program is trying to interrupt, not scale. Expansion in 2026/27 toward vocational education and grades 5 to 9. Teachers get a paid commercial licence. Students get a Socratic wrapper that refuses the answer. The state pays. The pedagogy is designed in. The host is public-law. That is the opposite of “just give them ChatGPT,” and it is the opposite of “just give them Hermes.” It is also not an agent evaluation. Do not claim it as one.

The U.S. and U.K. default is not a national wrapper. It is thousands of district contracts. TeachAI’s *AI Guidance for Schools Toolkit* — Code.org, ISTE, Khan Academy, ETS, World Economic Forum, with AASA, AFT, NEA, NSBA, CCSSO and others advising — is the most widely copied U.S. template: scope for staff, students, and vendors; principles; then local procedures. Many state education agencies now have pages. Quality varies. The ones that matter do four things the sources converge on: name a human accountable office, not “the AI”; inventory tools and refuse shadow-IT for student data; DPA or NDPA plus an AI addendum before any student PII or student work touches a model; assignment-level AI rules rather than a building-wide ban students will ignore at home.⁴⁶⁸ UK Department for Education policy, first 29 March 2023, last updated 12 August 2025: safety first; more immediate benefits and fewer risks from teacher-facing use; pupil-facing only with supervision, age-limit compliance, *Keeping Children Safe in Education*, and filtering and monitoring standards; no personal data in tools if avoidable; no training on student intellectual property without permission. Ofsted inspects decision-making about AI, not the brand of tool.

⁴⁶⁶Common Sense Media notes that teens bypass ChatGPT’s 13+ gate by lying, as recorded in the responsible-use memo, open fight 7, and Chapter 4 of this book.

⁴⁶⁷AI Leap Foundation / Estonian Ministry of Education and Research, TI-Hüpe, <https://tihupe.ee/en/>; first-year review 7 August 2026, <https://tihupe.ee/en/read-the-ai-leap-first-year-report/>. Announced February 2025 by President Alar Karis; Foundation plus Ministry; about €4 million pilot year, 50/50 public-private; 2025/26: 154 of 156 upper secondary schools; about 4,700-5,000 teachers, paid ChatGPT or Gemini; about 20,000 grade 10-11 students, ITI; University of Tartu impact study with Stanford and OpenAI. The 7 August 2026 review page itself withholds a learning-outcomes claim (“does not yet provide a basis for claiming that ITI has improved learning outcomes”); Tartu findings expected later in 2026. Open question 21. Do not claim Estonia evaluated autonomous agents.

⁴⁶⁸TeachAI, *AI Guidance for Schools Toolkit*, <https://www.teachai.org/toolkit-guidance>. Code.org, ISTE, Khan Academy, ETS, WEF; AASA, AFT, NEA, NSBA, CCSSO and others advising. Four-point compression (accountable office; inventory; DPA plus AI addendum; assignment-level rules) is this book’s reading of the responsible-use memo, section 6, not a TeachAI quotation.

Product-safety standards, 22 January 2025, updated 19 January 2026, are the most operational child-safety spec in English-language public policy in this corpus. Pupil-facing use must meet KCSIE plus those standards. Services that search live websites or let users share content fall under the Online Safety Act 2023.⁴⁶⁹

A district tenant of ChatGPT Edu, Gemini Education, or Claude for Teachers, with a real DPA, is the U.S. version of the public-law host. The institution is the data controller. The teacher is in the loop. The DSL gets the self-harm ping. That sentence is the default this chapter is willing to defend. It is not a claim that institution-hosted agents have a learning RCT. They do not. Open question 10: no identified learning RCT on Hermes-class or Grok Bot-class multi-step agents as a learning intervention; no education-specific audit of a deployed school agent with open-web tools.⁴⁷⁰ OpenAI Workspace Agents, as of the 23 April 2026 Academy playbook, are a research preview for eligible Edu workspaces, off by default, admin-gated, aimed at repeatable campus operations, not student homework. No named campus production case was fetched. Clemson News describes ChatGPT Edu plus GPT-5.5 plus Codex plus locally hosted open-weight models for researchers: documented access, not a published learning-outcomes study.⁴⁷¹ Access is not a result. A DPA is not a pedagogy. The default host is still the institution, because the alternative hosts fail the table on DSL, equity, and control.

The United States cannot copy Singapore's Student Learning Space by announcing one. A city-state with a national curriculum and a single ministry platform is not the U.S. default of thousands of district contracts, each district its own FERPA controller, each its own procurement office, each its own designated safeguarding equivalent only if the state and the board have built one. TeachAI's toolkit exists because the U.S. unit of control is the district, not the ministry. Prince George's addendum exists because a serious district has to write the promises the federal vacuum will not write. EO 14110, the Biden order that directed Education to produce AI resources, was rescinded on 20 January 2025 by EO 14148. EO 14277, 23 April 2025, is an adoption-and-workforce order — AI literacy, teacher professional development, public-private partnerships — not a child-safety regulation. AFT's March 2025 guardrails say so explicitly: "the current federal administration has indicated it will not regulate AI." The actual federal child-safety action on companion chatbots is at the FTC, not at Education.⁴⁷² New York City Public Schools, March 2026, is the largest-system version of the district host: traffic lights, ERMA before any tool touches student data, AI-specific

⁴⁶⁹UK Department for Education, *Generative artificial intelligence (AI) in education*, first 29 March 2023, last updated 12 August 2025, <https://www.gov.uk/government/publications/generative-artificial-intelligence-in-education/generative-artificial-intelligence-ai-in-education>. Product-safety standards as in note 10. Online Safety Act 2023 for live-web / user-sharing GenAI. *Keeping Children Safe in Education* as the pupil-facing floor.

⁴⁷⁰Open question 10 in the book brief. IEEE survey on "agentic AI-powered adaptive learning" noted an empirical gap on fully agentic platforms (doi:10.1109/aiei69164.2026.11497428); full text not fetched.

⁴⁷¹OpenAI Academy, "Understanding Workspace Agents in higher education," 23 April 2026, <https://academy.openai.com/public/clubs/higher-education-05x4z/blogs/understanding-workspace-agents-higher-education>. Research preview; off by default; admin-gated; no named campus production case fetched. Clemson News: ChatGPT Edu + GPT-5.5 + Codex plus locally hosted open-weight models for researchers, <https://news.clemson.edu/gpt-5-5-codex-and-clemson-hosted-models-give-researchers-new-options-for-ai-assisted-work/>. Access, not a published learning-outcomes study.

⁴⁷²Executive Order 14110 (30 October 2023), rescinded 20 January 2025 by Executive Order 14148. Executive Order 14277 (23 April 2025), "Advancing Artificial Intelligence Education for American Youth": adoption and workforce, not child-safety regulation. AFT *Commonsense Guardrails* v2 (March 2025): "the current federal administration has indicated it will not regulate AI." FTC 6(b) as in note 23.

vendor rules from December 2024 that require disclosure of AI capabilities and no training on student data. Red: placement, discipline, eligibility, promotion, graduation, IEP and 504 development, assessments and grading as the decision, surveillance, counseling and crisis, student-data model training, PII in unapproved tools. Yellow: student use for research, exploration, or creative projects with educator guidance, still under human review. Green: teacher-side generation, still ERMA-approved, human owns the output. Homework and academic integrity, cognitive offloading, grade-band guidance, family notification, and managed versus personal accounts were still listed as Playbook workstreams as of the fetched page.⁴⁷³ That is a host with red lines and unfinished homework rules. It is not a national Socratic wrapper. It is the U.S. shape of “the institution hosts.”

A rural district that cannot staff a privacy officer cannot negotiate a Prince George’s-style addendum. UNESCO warns governments not to divert education budgets into AI while 60 percent of primary schools globally remain unconnected, and not to paper over teacher shortages with tools.⁴⁷⁴ Leapfrog versus dependency is the Global South frame of the same hosting question. Estonia and Singapore can buy national licences and build Socratic wrappers. A ministry that spends on API tokens instead of teachers is doing what UNESCO told it not to do. A U.S. board that requires a home agent instead of issuing a school-provisioned tutor is doing the district version of the same diversion: shifting cost, control, and safeguarding onto families.

UNESCO 2023 still applies: governments must regulate; institutions must validate tools for ethical and pedagogical fitness — not outsource that validation to the vendor’s safety blog.⁴⁷⁵ A national wrapper is one way to do the validation. A district tenant with a PGCPs-style addendum is another. A curl installer on a fifteen-year-old’s laptop is not.

Teacher-facing agents: the real 2026 beachhead

The education market of 2026 bought constrained tutors and teacher copilots. It has not bought student-facing autonomous agents. The first real agent beachhead is the teacher’s evening and the registrar’s inbox, not the third-grader’s desk.

Anthropic launched Claude for Teachers on 14 July 2026: verified U.S. K-12 educators, free premium Claude, a Learning Commons connector with standards in all fifty states plus Illustrative Mathematics and OpenSciEd, teaching Skills co-developed with Learning Commons. It includes Claude Code and Cowork — scheduled and folder-scale autonomous work. Educator-only. Claude’s 18-plus policy. Training off for verified teacher accounts. FERPA-aligned DPA. Sign-up

⁴⁷³New York City Public Schools, “Guidance on Artificial Intelligence” (March 2026), <https://www.schools.nyc.gov/about-us/vision-and-mission/artificial-intelligence/guidance-on-artificial-intelligence>. Traffic-light framework; ERMA; December 2024 AI vendor rules (disclose AI capabilities; no training on student data). Playbook topics still open as of the fetched page: homework/integrity, cognitive offloading, grade-band guidance, family notification, managed versus personal accounts.

⁴⁷⁴UNESCO 2023 guidance (note 26) and rights-of-learners materials (note 13): do not divert education budgets into AI while about 60 percent of primary schools globally remain unconnected; do not use AI to paper over teacher shortages.

⁴⁷⁵UNESCO 2023 guidance: governments must regulate; institutions must validate tools for ethical and pedagogical fitness.

by 30 June 2027 for a year of free access. On 28 August 2026, a dedicated schools-and-districts offering: SSO, role-based access control, district terms. Documented early-feedback school named by Anthropic: Prospect Schools, Brooklyn. Announced impact pilot: Detroit Public Schools Community District, “working closely with teachers to study the impact on educator wellbeing and practice” — announced, not yet published as results. Gates Foundation partnership mentioned. Playlab to support a national network of lab schools. AFT President Randi Weingarten quoted on a “Gold Standard” for K-12 safety and privacy.⁴⁷⁶ Cowork examples in the teacher SKU: analyze a folder of de-identified diagnostics or exit tickets; run every school day at 4 p.m. to adapt tomorrow’s plan; scheduled parent newsletters. Not student-facing agents. The 4 p.m. job on de-identified exit tickets is the picture. The risk is a teacher’s intellectual property and a workload claim, not a child’s attachment to a companion.

OpenAI’s higher-education playbook, 23 April 2026, is explicit that Workspace Agents are for repeatable campus operations, not student homework: course-change triage, faculty AI-support intake, weekly student-services summaries, grant-readiness checklists, campus communications, IT software-request routing, program-review packets, faculty onboarding, accreditation evidence collection. Humans review before anything is sent. Off by default. Admin-gated. Research preview for eligible Edu workspaces. Playbook only. No named campus production case fetched.⁴⁷⁷ Codex in higher education, 9 April 2026 Academy: Codex App for supervised pilots, Codex Cloud for asynchronous repo work with a GitHub App and least-privilege tokens, Codex CLI. Aimed at developers and at non-technical campus staff via Edu login. Security emphasis: repo scoping, sandbox, approval policy.⁴⁷⁸ Cursor Cloud Agents — isolated Cursor-managed Linux VMs, clone, test, open pull requests, run while the laptop is closed — are a coding-agent product. Cursor docs do not market to schools. CNBC, 24 February 2026, reported Cursor saying roughly 35 percent of Cursor’s own pull requests are now generated by agents on VMs: a company statistic, not an education statistic.⁴⁷⁹ Staff developers and CS faculty under repo scoping are the defensible users. A seventh-grader’s homework is not.

ChatGPT for Teachers, as reported by CNBC on 19 November 2025, launched as a K-12 educator workspace, not intended for students, first cohort of districts representing roughly 150,000 educators, free through June 2027 in that report. OpenAI’s own expansion post timed out on fetch; additional district counts and the National Academy for AI Instruction / AFT “400,000 educators” figure appearing in search snippets are UNVERIFIED until that page is retrieved. Do not mash

⁴⁷⁶Anthropic, “Introducing Claude for Teachers,” 14 July 2026, <https://www.anthropic.com/news/claude-for-teachers>; <https://claude.com/solutions/teachers>. Schools-and-districts offering 28 August 2026. Educator-only; Claude 18+; training off for verified teacher accounts; FERPA-aligned DPA; sign-up by 30 June 2027 for a year of free access. Prospect Schools, Brooklyn: early feedback. Detroit Public Schools Community District: announced impact pilot on educator wellbeing and practice; results not published as of 30 August 2026. Weingarten “Gold Standard” quotation is from the Anthropic launch post.

⁴⁷⁷OpenAI Academy, 23 April 2026, as in note 32. Campus-ops examples are the playbook’s, not a named production deployment.

⁴⁷⁸OpenAI Academy, “Deploying Codex in Higher Education,” 9 April 2026, <https://academy.openai.com/public/clubs/higher-education-05x4z/blogs/deploying-codex-in-higher-education-2026-04-09>.

⁴⁷⁹Cursor Cloud Agents docs, <https://cursor.com/docs/cloud-agent>. CNBC, 24 February 2026: Cursor saying roughly 35 percent of Cursor’s own PRs generated by agents on VMs — company statistic, not an education statistic.

the CNBC first-cohort number with the Academy's five-year target.⁴⁸⁰ Google Workspace for Education includes Gemini app and Gemini Notebook free in all editions, all ages, Core service, data not used to train models, COPPA and FERPA support claimed. This is chat plus notebook, not a documented autonomous-agent SKU for schools. NotebookLM, now Gemini Notebook, is a grounded research partner on user-provided sources, closer to a constrained research agent than to Grok Bot.⁴⁸¹

Teacher-productivity edtech in the same year is usually not agents. MagicSchool: iKeepSafe FERPA and COPPA certifications, SOC 2 Type II, SDPC member, ESSA Level III via Instructure — vendor-funded positive association between teacher use and student ELA, not a large independent RCT. Midway ISD, Texas, about 8,700 students, districtwide pilot, 630-plus teachers trained summer 2024, 23-page DPA signed with no redlines: vendor case study.⁴⁸² Diffit: teacher tool to create and adapt instructional materials; vendor says it does not collect student data. MagicSchool's Raina and SchoolAI's Dot are constrained classroom assistants, not Grok Bot-class teammates.⁴⁸³ EDUCAUSE 2025 AI Landscape Study, as quoted via EdTech Magazine: chatbots were the top institutionwide AI license at 37 percent of respondents — chatbots, not agents.⁴⁸⁴ OECD *Digital Education Outlook 2026*, listing page, reports a TALIS 2024 figure of 37 percent of teachers already using generative AI for work tasks, with large cross-country variation. Full PDF was Cloudflare-blocked in the landscape fetch; treat full-report statistics as UNVERIFIED.⁴⁸⁵ Two 37 percents, two different surveys, neither of them a student-agent adoption number.

Unions may accept teacher-side agency and fight student-side Hermes. That is the live labor debate Chapter 7 already named. AFT's July 2026 resolution opposes student-facing AI in elementary and would ban social-companion chatbots at least until age 16.⁴⁸⁶ UK DfE's risk ordering is the same

⁴⁸⁰CNBC, "OpenAI rolls out 'ChatGPT for Teachers' for K-12 educators," 19 November 2025, <https://www.cnbc.com/2025/11/19/openai-chatgpt-for-teachers-k-12-educators.html>. Not for students; first cohort of districts representing roughly 150,000 educators; free through June 2027 in that report. OpenAI, "Bringing ChatGPT for Teachers to more U.S. school districts," <https://openai.com/index/bringing-chatgpt-for-teachers-to-more-us-school-districts/>. FETCH TIMED OUT 30 August 2026. Additional district counts and AFT Academy 400,000 five-year target appearing in search snippets: UNVERIFIED until that page is retrieved. Do not mash the CNBC first-cohort number with the Academy target.

⁴⁸¹Google Workspace Help, "Quickstart Guide to Gemini and Gemini Notebook for Education," <https://support.google.com/a/answer/16350447>. Chat plus notebook, not a documented autonomous-agent SKU for schools.

⁴⁸²MagicSchool trust center, <https://www.magicschool.ai/privacy/trust>. ESSA Level III via Instructure: vendor-funded positive association with student ELA, not a large independent RCT. MagicSchool, "Midway ISD Case Study," <https://www.magicschool.ai/case-studies/midway-isd>. Vendor case study. Midway ISD, Texas, about 8,700 students; 630-plus teachers trained summer 2024; 23-page DPA, no redlines.

⁴⁸³Diffit plans, <https://web.diffit.me/plans>. Vendor: FERPA/COPPA; does not collect student data. Raina and Dot: constrained classroom assistants, as recorded in the landscape memo, section 4.4; not Grok Bot-class teammates.

⁴⁸⁴EDUCAUSE 2025 AI Landscape Study, as quoted via EdTech Magazine / Jenay Robert: chatbots the top institutionwide AI license at 37 percent of respondents — chatbots, not agents.

⁴⁸⁵OECD, *Digital Education Outlook 2026*, listing page https://www.oecd.org/en/publications/oecd-digital-education-outlook-2026_062a7394-en.html. TALIS 2024 figure of 37 percent of teachers using GenAI for work tasks: used from the listing / responsible-use stream. Full PDF Cloudflare-blocked in the landscape fetch — UNVERIFIED at full-PDF level.

⁴⁸⁶American Federation of Teachers, "Preparing for the Age of Artificial Intelligence: Protecting Democracy, Workers, Students and the Planet," resolution adopted 18 July 2026, <https://www.aft.org/resolution/preparing-age-artificial-intelligence-protecting-democracy-workers-students-and-planet>. No student-facing AI in elementary; social-

beachhead in procurement language: teacher-facing first; pupil-facing only with the product-safety stack. The beachhead is politically available because the harm class is different. A Cowork job on de-identified exit tickets can leak a teacher’s intellectual property or add review work, which the Department of Education already warned about in 2023: “in the loop” done badly adds work.⁴⁸⁷ It does not isolate a child with a companion, browse the open web as the child, or hold a persistent memory the DSL cannot see. Private-cloud or local hosting is more defensible on the teacher side for that reason. It is not more defensible on the student side just because the same binary can run.

A teacher-facing agent can still fail. “In the loop” that adds review of every draft is a workload injury, which is why AFT puts the right to disconnect and kill-switch language in the bargaining annex and why NEA puts high-stakes evaluation off-limits to AI data alone.⁴⁸⁸ A Cowork job that is fed identified student work rather than de-identified exit tickets is a FERPA incident with a schedule. A Workspace Agent that sends the campus email before a human reviews it is the playbook’s own failure mode: the 23 April 2026 document is explicit that humans review before anything is sent, which is a design requirement, not a courtesy.⁴⁸⁹ Those failures are IP, labor, and records failures. They are not the Setzer and Raine harm class. Collapsing them into “agents in schools” is how a board gets talked into a student Hermes because the teacher SKU already has Cowork. Product class first. Host second. Audience third. Teacher SKU is not a student SKU with a different login.

Do not re-list every Claude-for-Teachers connector. ASSISTments, Brisk Teaching, Canva Education, Coteach, Diffit, Eedi, MagicSchool, Snorkl, TeachFX are an ecosystem Anthropic named, not products Anthropic built, and not evidence that student agents have arrived.⁴⁹⁰ The beachhead is the teacher SKU with Cowork, the Edu playbook with Workspace Agents off by default, and Codex Cloud under repo scoping. Put those in the registrar’s inbox and the teacher’s evening. Do not put them on the third-grader’s desk and call the desk a beachhead.

Student-owned Hermes is a home tutor, not a school requirement

The thesis sentence this chapter has to hold is not a slogan. Student-owned local agents are a home choice for families who want them, with the same status as a privately hired tutor: not a school-required tool, not fed with education records, not a substitute for assessment evidence.

If the school required a personal agent, FERPA and COPPA problems land on the parent-installed runtime. Who is the controller of the child’s memory? Open question 22 does not close because the license is MIT. A required tool the school does not host is a required tool the school does not

companion chatbot ban at least until age 16.

⁴⁸⁷U.S. Department of Education, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023), <https://www.ed.gov/sites/ed/files/documents/ai-report/ai-report.pdf>. Warns that “in the loop” done badly adds work.

⁴⁸⁸AFT *Commonsense Guardrails* v2 (March 2025), bargaining annex: right to disconnect; kill decisions; AI a mandatory subject of bargaining. NEA Policy Statement on AI in Education (2024): AI data alone never determines evaluation, placement, discipline, diagnosis, or safety decisions.

⁴⁸⁹OpenAI Academy, 23 April 2026, as in note 32: humans review before anything is sent.

⁴⁹⁰Anthropic launch post (note 37): connectors named as ecosystem, not as Anthropic-built products. Cowork examples (de-identified diagnostics/exit tickets; 4 p.m. scheduled run; parent newsletters) are vendor illustrations of teacher-side agency, not student-facing agents.

control. Education records in that runtime are a disclosure the school-official exception was not built for. A memory that persists across sessions is an education record if the school required the sessions, and a family file if the parent bought the machine. The law does not get cleaner because the weights are open. It gets dirtier, because there is no DPA, no admin log, and no deletion clock the district can audit.

Do not grade homework that assumes every child has a premium personal agent the school does not issue. That is red line 11, already mapped to Pew, the Harvard 2024 survey, AFT's equity value, and UNESCO's connectivity figures.⁴⁹¹ A personal agent is more unequal than a chatbot. Who owns it — the child, the family, the school, the state — decides who keeps the learning history, who can audit it, and who can turn it off. Homework-gap 2.0 is already a paid-tier gap. Making the unpaid take-home assume a local GPU is how a district launders that gap into a grade.

Graduate research labs are the live “maybe later.” Literature monitoring, experiment babysitting, paper drafting: the same class of work OpenAI and Anthropic pitch to researchers. Hermes Agent's bundled research-paper skill is the most fully documented research-agent workflow in the named stack — arXiv, Semantic Scholar, `delegate_task`, `cronjob`, literature review through experiments through LaTeX through submission.⁴⁹² It is not evidence of academic adoption. Mapping is not adoption. No official Nous materials, university newsrooms, or policy documents fetched for this book document institutional use. Clemson's locally hosted open-weight models for researchers are access, not a Hermes deployment and not a learning-outcomes study. EdTech Magazine in December 2025 cites Johns Hopkins Agent Laboratory — authors claim an 84 percent cost decrease versus prior autonomous-research methods — and UT Verse at University of Tennessee Knoxville, with Dan Harder quoted that fully autonomous agents still need human validation. Agent Laboratory paper and UT Verse primary pages were not fetched. Treat those two campus examples as journalism-sourced, UNVERIFIED against primary docs.⁴⁹³ The maybe-later is real as a class of work. It is not a K-12 argument, and it is not a finding that graduate methods courses should require a Hermes-class agent in year one.

College is where the syllabus, not the installer, is the control surface. Agentic coding and research only where the syllabus allows and transcripts exist. Stanford CS221, Spring 2026, is the worked example of refusal: AI as a human collaborator; no asking for answers; no “check my work”; no using an agentic coding tool such as Codex or Claude Code to build a repository from a project proposal; transcripts required.⁴⁹⁴ A research university refusing agents on the assignment that is supposed to form the skill is a graduate behavior, not Luddism. UC San Diego CSE, in a Microsoft

⁴⁹¹Responsible-use red line 11 / book red-line table: no unpaid take-home that assumes every child has a premium personal agent the school does not issue. Pew (note 11); Harvard 2024 (note 11); AFT equity value; UNESCO connectivity (note 13).

⁴⁹²Hermes Agent research-paper-writing skill, as in note 5. Vendor workflow documentation. Mapping ≠ adoption. Landscape memo, section 2.3: no official school/university use found.

⁴⁹³EdTech Magazine, December 2025, <https://edtechmagazine.com/higher/article/2025/12/ai-agents-higher-education-transforming-student-services-and-support-perfcon>. Johns Hopkins Agent Laboratory (authors claim 84 percent cost decrease versus prior autonomous-research methods) and UT Verse at University of Tennessee Knoxville (Dan Harder: fully autonomous agents still need human validation). Primary pages not fetched — journalism-sourced, UNVERIFIED against primary docs.

⁴⁹⁴Stanford CS221, Spring 2026, Generative AI Policy, <https://stanford-cs221.github.io/spring2026/>. Collaborator, not answer key; transcripts required; agentic coding tools such as Codex or Claude Code not allowed to build a repository from a project proposal. Worked example of college-level agent refusal, not a learning RCT.

customer story, redesigned intro courses so students first demonstrate manual competency, then use Copilot; research on trust calibration; plans to prepare students for “emerging agentic AI.” Curriculum redesign around a coding assistant, not a published RCT of Cloud Agents, and labelled here as a vendor customer story.⁴⁹⁵ Student-owned local agents at college are extra-curricular unless the course says otherwise. Enterprise tenant for FERPA, HIPAA, and IP. The student becomes the FERPA rights-holder at eighteen or upon enrollment in postsecondary education, whichever comes first.⁴⁹⁶ Extra-curricular is not prohibited. Required is.

“Not fed with education records” is an operational sentence, not a vibe. A teacher who pastes a class list into a consumer bot, or uploads IEPs to a free account, or lets an agent summarize a Google Drive that includes other children’s work, has already moved education records into a vendor or a runtime that is not a school official. UK DfE: student essays and teacher lesson plans are copyright works; putting them into a training-on-by-default tool can be both a data-protection and an IP problem. US FERPA: once an education record is in a vendor that is not a school official, redisclosure limits are already broken.⁴⁹⁷ A student-owned Hermes that is allowed to read the school’s LMS, the student’s mailbox, or a shared folder of peer work is that incident with a local disk. Local is not a legal exception. It is a location. The controller question does not get easier when the disk is in a backpack.

AFT Core Value 3: certified educators decide whether, when, and how; vendors serve, they do not drive.⁴⁹⁸ US ED 2023: involve educators; humans in the loop; firmly reject teacher replacement.⁴⁹⁹ UNESCO: human-centred; preserve human agency in a world of automation.⁵⁰⁰ A school that requires a student-owned Hermes has handed pedagogy to a runtime the teacher cannot see, the DSL cannot alert on, and the family may not be able to afford. That is not augmentation. It is an unfunded, unhosted, unlogged mandate.

The home-tutor frame has a limit, and the limit is honesty about what a home tutor is. A privately hired human tutor is not fed with the school’s education records by default. The school does not grade the tutor’s notes as the student’s work. The school does not require every family to hire one as a condition of the homework. A parent who installs Hermes Agent on a family machine is making a home choice with those same three properties only if the school treats it that way: not required, not

⁴⁹⁵Microsoft customer story, “University of California San Diego — GitHub Copilot,” <https://www.microsoft.com/en/customers/story/26539-university-of-california-san-diego-github-copilot>. Faculty (Porter, Shah) redesigned intro courses: manual competency first, then Copilot. Vendor customer story, not a peer-reviewed trial of Cloud Agents.

⁴⁹⁶FERPA eligible-student rule: rights transfer at eighteen or upon enrollment in postsecondary education, whichever comes first. <https://studentprivacy.ed.gov/content/eligible-student>.

⁴⁹⁷UK DfE, 12 August 2025 update: student essays and teacher lesson plans as copyright works; schools must not allow students’ original work to be used to train generative AI models without permission. FERPA: once an education record is in a vendor that is not a school official, redisclosure limits are already broken.

⁴⁹⁸AFT *Commonsense Guardrails* v2 (March 2025), Core Value 3: educators choose the technology; vendors serve, they do not drive.

⁴⁹⁹US ED 2023 (note 48): humans in the loop; firmly rejects the idea that AI could replace teachers; involve educators (Recommendation 5).

⁵⁰⁰UNESCO AI in education, <https://www.unesco.org/en/digital-education/artificial-intelligence>; student framework <https://www.unesco.org/en/articles/ai-competency-framework-students>; teacher framework <https://www.unesco.org/en/articles/ai-competency-framework-teachers>. Human-centred approach; “preserving human agency in a world of automation.”

fed with records, not a substitute for assessment evidence. The moment the assignment cannot be completed without the agent, the home choice has become a school requirement the school did not host. JCQ, effective 30 April 2025, already treats unacknowledged AI as malpractice and requires process evidence: tool name, URL, date, a non-editable copy of prompt and output.⁵⁰¹ A local agent that finishes the paper overnight is not a citation problem the student can screenshot after the fact. It is an assessment design that assumed a computer the school does not issue. Chapter 5's two-lane rule is the pedagogical translation: formation work is unaided or hint-only; performance-in-an-AI-world work is allowed, logged, and judged. A required personal agent erases the formation lane at home, which is where the unpaid take-home already lives.

Companion risk does not vanish because the weights are local. Character.AI's harm surface was not that the bytes sat on a particular server. It was design: sexualized conversation, isolation language, "come home to me," secret-keeping, engagement-maximizing loops. UK product-safety standards now name those failure modes as procurement requirements: no sycophancy, no "you can trust me / don't tell anyone," no engagement-maximizing loops, no anthropomorphism.⁵⁰² A local agent with a messaging gateway to Telegram and a persistent memory of the child can be configured as a tutor. It can also be configured, or can drift, as a companion. Default Hermes is not a Socratic wrapper. Estonia built the wrapper because the default commercial pattern is answer-giving and, in the companion products, attachment. Hosting the default on a laptop does not install the wrapper.

Practical checklist

Every line below is a compression of a cited instrument, not a new ethic. Use it as a board one-pager. Do not add unsourced items. The twenty-four points in the responsible-use stream are the source; this section is that list in prose, grouped the way a Monday meeting can actually run it.⁵⁰³

Before any tool touches a student. Purpose on paper: age, SEND, subject, teacher-facing versus student-facing. UK product-safety standards call this "stated purpose." If the vendor cannot say, do not buy. Contract first: NDPA or state DPA plus an AI addendum that says no training or fine-tune on student or teacher work, no ads, named subprocessors and locations, deletion clock, breach duties, audit rights, off-switch. Illinois SOPPA, Prince George's County's addendum, FERPA's school-official exception, COPPA 2025. Age is the strictest of vendor terms of service, COPPA, and local safeguarding. Do not accept self-declared age. UNESCO 2023: thirteen for independent generative-AI conversation; under-thirteen and primary use of untested tools not appropriate. Independent safety look: check Common Sense Media / Youth AI Safety Institute. Reject "unacceptable." Treat "high risk" as not-for-students. Prefer purpose-built education tools. No emotion AI, no affect proctoring, no companion persona: EU AI Act Article 5(1)(f) from 2 February 2025, with Recital 44 on the absence of scientific basis and the power imbalance in schools; UK

⁵⁰¹Joint Council for Qualifications, *AI Use in Assessments*, effective 30 April 2025, <https://www.jcq.org.uk/knowledge-hub/ai-use-in-assessments-your-role-in-protecting-the-integrity-of-qualifications-2/>. Tool name + URL + date; retain a non-editable copy of prompt and output; AI not the sole marker.

⁵⁰²UK product-safety standards, January 2026 update (note 10): manipulation / emotional-development — no sycophancy, no "you can trust me / don't tell anyone," no engagement-maximizing loops. Companion failure mode in the Setzer and Raine complaints (note 23).

⁵⁰³Responsible-use memo, section 8, twenty-four-point practical checklist. This chapter's prose grouping is a compression of that list, not a new ethic.

emotional-development standards; Common Sense companion ratings. Humans in the loop for any consequential output. Placement, marking, IEP, discipline, teacher evaluation: AI may draft; a named professional decides and can override. US ED 2023 Recommendation 1; NEA; EU Annex III. Evidence or a pilot with a kill switch. NEA Principle 2; ESEA evidence tiers; UK edtech evidence-board direction.⁵⁰⁴

If students will use it. Teacher-facing first. UK DfE's risk ordering. Build staff fluency before student accounts. Education-designed, not general chat. Singapore's P4-6 rule. Estonia's ITI, no direct answers. UK cognitive-development standard: hints before answers; detect offloading; teacher-gated full solution. No open-web agency unless the school operates the sandbox, logging, and DSL alerts. Disable browsing and tools by default. UK security standard plus Online Safety Act 2023. Safeguarding path: the product holds DSL contact; alerts on self-harm, sexual content, grooming patterns, distress; signposts to humans; never isolation language. UK monitoring and mental-health standards; KCSIE. Time limits and no engagement hacks. Hard caps, teacher-resettable. No sycophancy, no "don't leave." If the task needs a capable model, the school issues it. Do not grade the homework gap. AFT equity value; Pew; Harvard 2024. Assignment labels: full, partial, or none, with citation rules — tool, date, prompt and output artefact. AFT guardrails; JCQ 2025. Assessment that survives an agent: in-class, oral, practical, drafts, vivas. Detectors are never sole evidence. QAA 2023; Liang 2023; JCQ 2025. Accessibility is design. WCAG plus UDL. Never sole diagnostic use. Watch detector and proctoring bias against multilingual and disabled students.⁵⁰⁵

Hosting and labor. Default host is the institution or a public platform under public-law duties. Local open weights are for institution-run servers or optional home use, not a required personal agent. Educators choose. Procurement committees include union-designated teachers. Vendors do not set pedagogy. AFT Core Value 3; NEA; US ED 2023 Recommendation 5. Bargain it: data, surveillance, evaluation, PD time, right to disconnect, workload. AFT bargaining annex; NEA. Literacy for all roles. EU Article 4 already requires it of deployers, from 2 February 2025. UNESCO frameworks for teachers and students. Singapore's character-and-citizenship strand. Parent notice and a real alternative. US ED toolkit: opportunity to opt out of AI-enabled applications. Review on a clock. Models change under the same brand name. Turnitin's February 2026 model update; ChatGPT-5 Common Sense reassessment. Re-audit annually or on model-change notice. Incident plan: who pulls the plug, who tells parents, who preserves logs for the DSL and the DPA. Environment is in the union texts for a reason. AFT Value 9; NEA Principle 3. Do not require image-generation as default classwork.⁵⁰⁶

⁵⁰⁴UK-PSS "stated purpose"; SOPPA; PGCPs addendum (note 24); FERPA; COPPA 2025 (note 19); UNESCO 2023 (note 26); Common Sense ratings (notes 15, 20, 22); EU AI Act Art. 5(1)(f) from 2 February 2025, Recital 44; US ED 2023 Recommendation 1; NEA Principle 2; ESEA evidence tiers. Regulation (EU) 2024/1689.

⁵⁰⁵UK-DfE teacher-facing first (note 30); Singapore P4-6 (note 25); Estonia ITI (note 28); UK-PSS cognitive-development (hints before answers; detect offloading; teacher-gated full solution); no open-web agency unless school-operated sandbox, logging, DSL alerts; KCSIE; AFT assignment labels (full / partial / none); JCQ 2025 (note 62); QAA, *Reconsidering assessment for the Chat GPT era* (July 2023); Liang et al., *Patterns* 2023 (61.22 percent of TOEFL essays by non-native writers flagged); CAST UDL Guidelines 3.0; NEA 2024 on detector and proctoring bias; Center for Democracy and Technology, *Ableism and Disability Discrimination in New Surveillance Technologies* (May 2022).

⁵⁰⁶Default host: institution (this chapter's thesis; Singapore/Estonia as public-law hosts). AFT Core Value 3; NEA; US ED 2023 Recommendation 5. AFT bargaining annex. EU AI Act Art. 4 AI-literacy duty on deployers, from

That is not a new ethic. It is the instruments, compressed, in the order a board can ask them. The hosting line is the one this chapter adds nothing to except emphasis: default host, the institution; local open weights, institution-run or optional home; never a required personal agent.

Age-gated agent ladder

The ladder is this book’s proposed integration model. It is inferred from the legal, developmental, and evidence constraints in Chapters 4 through 8. It is not a documented district SOP. It is not Broward policy, not NYCPS policy, not Singapore’s Masterplan, not Estonia’s ITI specification. Do not present it as existing district policy. Do not claim institution-hosted agents have a learning RCT. They do not.

Band	Student-facing	Agent features	Host
5–8	None. Literacy <i>about</i> AI. Teacher-facing generation with human review.	No memory of the child. No chat. No tools.	District teacher tenant
9–11	Constrained, logged, hint-only practice tutors in teacher rooms	No open web, no email, no companion, no overnight	School-provisioned, COPPA DPA
12–14	Assigned Socratic tutoring + revision tools; AI literacy of systems	Still no acting agents; logs; session caps	School-provisioned; FERPA + state student-data law
15–18	Disclosed partner on hard tasks; ISTE roles; unaided formation lane remains	Sandboxed tools only if school-operated; no transact/post; transcripts	Institution tenant; consumer tools not required for grades
College	Course-bounded collaboration; domain workflows; data classification	Agentic coding/research only where the syllabus allows and transcripts exist; Stanford-style bans remain legitimate	Enterprise tenant for FERPA/HIPAA/IP; student-owned local agents are extra-curricular
Adult / continuing	Copilot with human sign-off	Workplace disclosure norms	Employer or IHE, not a public companion

2 February 2025. US ED *Empowering Education Leaders* toolkit (ERIC ED661924); opportunity to opt out of AI-enabled applications. Turnitin February 2026 model update and ChatGPT-5 CSM reassessment as model-change examples, recorded in the responsible-use memo. AFT Value 9 and NEA Principle 3 on environment; independent energy analyses of school-scale inference not fetched — keep at union/district naming, not invented kWh. Open question 19.

Five to eight is Broward K-5 at zero student-facing tools, UNESCO’s “not appropriate” for under-thirteen untested tools in primary, AFT’s July 2026 opposition to student-facing AI in elementary, and COPPA still fully on. Literacy about AI, not a chat window. Teacher-facing generation with human review. No memory of the child, because a memory of a seven-year-old is a companion seed even if the product copy says tutor. Host: the district’s teacher tenant, which is the only host that can keep the child off the system.⁵⁰⁷

Nine to eleven is COPPA’s remaining years plus Singapore’s P4-6 rule translated into a U.S. district: constrained, logged, hint-only practice tutors in teacher rooms. No open web, no email, no companion, no overnight jobs. School-provisioned, with a COPPA DPA. Khanmigo’s low-risk rating is the nearest product-class existence proof, not an endorsement of any named vendor as a district SOP, and not an agent. Access is not tutoring; Oreopoulos and Low already measured that. The host still has to be the school, because a nine-year-old’s parent-installed Hermes is a COPPA operator problem the school cannot DPA after the fact.⁵⁰⁸

Twelve to fourteen is assigned Socratic tutoring and revision tools, plus AI literacy of systems — how logs work, who can see the chat, why the tool refuses the answer. Still no acting agents. Logs. Session caps. School-provisioned; FERPA plus state student-data law, because many of these students are thirteen or fourteen and COPPA has turned off while New York 2-d, SOPPA, SOPIPA, and Texas Chapter 32 have not. Companion chat is a documented harm surface in this band: Setzer was fourteen; Raine was sixteen; AFT would ban social-companion chatbots at least until sixteen. The ladder’s “still no acting agents” is the UK product-safety stack plus the absence of any Hermes-class RCT, not a developmental romanticism.⁵⁰⁹

Fifteen to eighteen is a disclosed partner on hard tasks, ISTE roles, and an unaided formation lane that remains. Sandboxed tools only if school-operated. No transact, no post. Transcripts. Institution tenant. Consumer tools not required for grades. This is the band where a family may already have a home agent, and where the school’s job is to refuse to grade that agent as the student. JCQ’s acknowledgement format is a transferable habit: tool, URL, date, artefact. Stanford’s transcript requirement is a preview. The formation lane is Chapter 5’s rule, still in force at seventeen, because expertise reversal does not turn a high-school junior into a researcher who can catch the model on demand.

College is course-bounded collaboration, domain workflows, data classification. Agentic coding and research only where the syllabus allows and transcripts exist. Stanford-style bans remain legitimate. Enterprise tenant for FERPA, HIPAA, and IP. Student-owned local agents are extra-curricular. The rights-holder is now the student. That does not make a required personal agent lawful as a condition of a grade; it makes the student the person who can refuse the extra-curricular install.

⁵⁰⁷Broward County Public Schools, Broward Powered by AI, Resolution 26-132, 10 March 2026, <https://browardschools.ai/>: K-5, zero district-approved student-facing AI tools. UNESCO 2023 under-thirteen / primary (note 26). AFT July 2026 elementary opposition (note 47). COPPA under 13 (note 19). Ladder band 5-8 is book synthesis from Chapter 4, labelled as such, not Broward SOP.

⁵⁰⁸COPPA remaining years; Singapore P4-6 as analogue, not as U.S. law (note 25). Khanmigo low-risk rating (note 20) as nearest product-class existence proof, not a district SOP. P. Oreopoulos and N. Low, “One click away: AI tutoring with Khanmigo in a two-year school experiment,” EdWorkingPaper 26-1551 / NBER w35620 (2026): access is not tutoring. Parent-installed Hermes for a nine-year-old: COPPA operator problem, not a school-official exception.

⁵⁰⁹Companion harm surface: Setzer was 14; Raine was 16 (note 23). AFT companion ban at least until 16 (note 47). UK-PSS (note 10). Open question 10: no Hermes-class learning RCT (note 31).

Adult and continuing education is a copilot with human sign-off, workplace disclosure norms, hosted by the employer or the institution of higher education, not a public companion. The never-column does not become acceptable at twenty-two because the user can sign a terms of service.

Two live debates stay open, because the ladder cannot close them. Who is the data controller of a child's agent memory when the runtime is local, open-source, and parent-installed. Whether Hermes-class local agents belong in graduate research labs — literature monitoring, experiment babysitting, paper drafting — before they belong anywhere near K-12. The book's lean on the second is already on the page: maybe later, as extra-curricular or as a methods tool after students can catch the model, not as a K-12 requirement and not as a child-safety school product. The first does not get a lean that pretends the law is settled. It is not.

What this chapter will not claim is as important as the ladder. Singapore and Estonia did not evaluate autonomous agents. Estonia's first-year review exists, dated 7 August 2026; the page itself withholds a learning-outcomes claim, and Tartu's coefficients were not published as of 30 August 2026. Institution-hosted agents do not have a learning RCT. OpenAI Workspace Agents are a playbook and a research preview, with no named campus production case fetched. Clemson is access, not outcomes. Local open-source does not create a DSL path and does not close an equity gap. Hermes Agent's research-paper skill is a vendor workflow, not academic adoption. Student-owned Hermes is not a child-safety school product. Grok Bot has no education SKU. El Salvador is not Grok Bot. Hermes 4 is not Hermes Agent. Open questions 10, 13, 14, 21, and 22 remain standing.⁵¹⁰

If the host is the institution and the struggle stays in the student, what can a graduate actually do? Not "can prompt ChatGPT." That is the next chapter's enemy, and this chapter's hosting default is what makes the question answerable. A graduate whose only agent is a consumer window the school never hosted has been trained on the never-column. A graduate whose school hosted a Socratic tutor, labelled the lanes, kept the DSL on the alert path, and treated home Hermes as a tutor the way it treats a hired human — optional, ungraded as skill, unfed with records — has been trained on the table. The deeds come next. The host was the precondition.

⁵¹⁰What this chapter will not claim, restated from the outline and brief. Open questions left standing: 10 (no Hermes-class / Grok Bot-class learning RCT; no education-specific audit of a deployed school agent with open-web tools); 13 (El Salvador status; announcement only); 14 (no Hermes or Grok Bot institutional AUP, DPA, or classroom SOP); 21 (Estonia first-year review withholds a learning-outcomes claim; Tartu coefficients unpublished as of 30 August 2026); 22 (data controller of a child's persistent agent memory when the runtime is local, open-source, and parent-installed).

Chapter 10

Not “Can Prompt ChatGPT”



Figure 11: A mortarboard, a closed laptop, a blank card, and a fountain pen.

Stanford CS221, in Spring 2026, did not ask whether its students could prompt ChatGPT. It asked whether they could still do the work the course exists to form. The generative-AI policy on the course site is short and specific. AI is a human collaborator, not an answer key. Students may not ask the model for answers. They may not ask it to “check my work.” They may not use an agentic coding tool such as Codex or Claude Code to build a repository from a project proposal. They must share the ChatGPT link or screenshots. Tools used without attribution are banned. Honor-

code enforcement is named, not implied.⁵¹¹

That is a research university refusing agents on the assignment that is supposed to form the skill. The refusal is a graduate behavior, not Luddism. It is also the sentence this chapter has to make operational for everyone else. An “AI-ready graduate” is not a person who can call an API. It is a person who can pause the tool, catch the model, keep a transcript, and still do the job when the window is closed.

The popular story will try to make this chapter a credential. Workforce entry, prompt skill, a badge. Vendor copy will offer a tutor for every child and a teammate of agents for every graduate. Union copy will offer gold-standard privacy. UNESCO will warn ministries not to paper over teacher shortages with a chatbot. Those are live debates, and they belong on the page. None of them is a graduate profile. A profile is a list of deeds. This chapter closes on deeds, aligned to the ISTE Live 2026 six roles, the AI4K12 grade 9–12 charts, the U.S. Department of Education’s humans-in-the-loop line, Cornell’s question whether use contributes to learning, and the two-lane assessment rule of Chapter 5. That alignment is a synthesis of those sources. It is not a validated assessment instrument. No dataset shows that students who can do these things have better later employment or civic outcomes. This chapter will not invent one.

The default student behavior is already measured, and it is not the profile. In the 2024 University of California Undergraduate Experience Survey, 67 percent of UCLA undergraduate respondents reported using generative-AI tools at least several times a year; brainstorming, topic research, and studying for exams were the most common uses.⁵¹² That is consumer ToS behavior at campus scale. It is what the graduate profile has to upgrade, not celebrate as fluency. “Can prompt ChatGPT” is the UCLA default with a nicer name.

Not prompt skill

Prompt skill is a consumer behavior. A graduate profile is a set of things the student can still do when the tool is taken away, and a set of things the student can do with the tool without surrendering the work the school exists to form. Those are not the same object. Treating them as the same object is how a campus gets a literacy module that teaches students to write better prompts and an assessment regime that still grades the model’s paragraph.

Cornell’s AI Critical Literacy Program, expanded in August 2026 from a spring pilot to all incoming students, does not ask whether the student can call an API. It asks whether a given use *contributes to learning*.⁵¹³ Four Canvas modules — What is Generative AI?; Ethical Questions Surrounding

⁵¹¹Stanford CS221, Spring 2026, Generative AI Policy, <https://stanford-cs221.github.io/spring2026/>. AI as human collaborator; no asking for answers; no “check my work”; agentic coding tools such as Codex or Claude Code not allowed to build a repository from a project proposal; ChatGPT link or screenshots required; tools without attribution banned; honor-code enforcement. Age-stages memo, college band; landscape §4.6.

⁵¹²UCLA SAIRO UCUES 2024 AI brief and Teaching and Learning Center: 67 percent of undergraduate respondents used AI tools at least several times a year; top uses brainstorming, researching a topic, and studying for exams. SAIRO brief: 4,706 respondents, 15 percent response rate. Age-stages Band 5; brief Chapter 10 load-bearing fact 4. University of California Undergraduate Experience Survey, 2024. <https://sairo.ucla.edu/file/69a9752b-ba55-46a1-8cf1-e2a963ffbb97> ; <https://teaching.ucla.edu/resources/teaching-guides/using-generative-ai-reflectively-and-responsibly-in-teaching-and-learning/>

⁵¹³Jose Beduya, “Cornell Expands AI Critical Literacy Program for Incoming Students,” *Cornell Chronicle*, August

Generative AI; GenAI and Learning; Creating an AI Policy for Personal Use — sit under the Center for Teaching Innovation and the Library. The pilot reported literacy gains in how models work and are trained. Those gains are CTI-reported, not a published randomised trial. The question is still the right one. A fluent homework that the student cannot reconstruct does not contribute to learning, however convenient it is for the deadline. A use that forces restatement, checking, and defence might. The program does not settle which uses do. It installs the question at the start of college. That is more than a prompt-engineering workshop, and less than a transfer study.

The U.S. Department of Education's 2023 report *Artificial Intelligence and the Future of Teaching and Learning* put humans in the loop of high-stakes decisions and “firmly rejects the idea that AI could replace teachers.”⁵¹⁴ The 2026 supplemental priority (91 FR 18774) funds literacy, tutoring, personalization, and administrative efficiency as a scoring tool for discretionary grants. It does not define “agentic AI.” It does not authorize unconstrained student agents. It does not certify prompt skill as a graduate outcome. Humans in the loop is a design specification, not a slogan to print on a diploma.

Chapter 5's two-lane rule is the assessment architecture that makes Cornell's question visible in a gradebook. Formation: AI forbidden or hint-only; grade unaided work. Performance in an AI world: AI allowed, logged; grade judgement — selection, critique, verification, process. The rule is theoretically motivated, not a fitted grade-by-grade schedule, and not existing district policy.⁵¹⁵ The graduate this chapter is describing is a person who can *choose the lane* and can *pause the tool*. Choosing the lane is a deed. Prompting is a keystroke.

The live debate this chapter takes a position on, without pretending the evidence is thicker than it is, has three rooms.

Room one: what “AI-fluent graduate” means. One camp treats fluency as prompt skill — the student who can get a useful paragraph out of ChatGPT, who has used a campus Copilot, who can name a model. The other camp treats fluency as the ISTE / Cornell / two-lane bundle: roles that sit on domain knowledge, a question about learning rather than convenience, and an assessment architecture that still measures unaided skill. This book is in the second camp. Prompt skill is not nothing. It is not a formed judgement. Willingham's constraint from Chapter 6 still holds: critical thinking is not a transferable bicycle-skill an agent can teach in the abstract. Literacy without knowledge is a checklist students can recap and not execute.⁵¹⁶ A graduate who can prompt and

2026, <https://news.cornell.edu/stories/2026/08/cornell-expands-ai-critical-literacy-program-incoming-students>. University newsroom. Four Canvas modules (What is Generative AI?; Ethical Questions Surrounding Generative AI; GenAI and Learning; Creating an AI Policy for Personal Use); CTI and Library; faculty from writing, business, biology, engineering, economics, communication, and information science. Pilot literacy gains are CTI-reported, not a published RCT. Explicit goal from the Chronicle account: whether AI use “contributes to learning.”

⁵¹⁴U.S. Department of Education, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023), <https://www.ed.gov/sites/ed/files/documents/ai-report/ai-report.pdf>. The Department “firmly rejects the idea that AI could replace teachers.” Humans-in-the-loop is the 2023 line this chapter uses; 91 Fed. Reg. 18774 (April 13, 2026), effective May 13, 2026, is the later discretionary-grant scoring tool and does not define “agentic AI.”

⁵¹⁵Chapter 5, two-lane rule. Formation: AI forbidden or hint-only; grade unaided work. Performance in an AI world: AI allowed, logged; grade judgement. Theoretically motivated synthesis, not a fitted schedule, not existing district policy. Perkins, Furze, Roe, and MacVaugh, “The AI Assessment Scale (AIAS),” *Journal of University Teaching and Learning Practice* (2024), is shared vocabulary, not an RCT of learning.

⁵¹⁶Daniel T. Willingham, “Critical Thinking: Why Is It So Hard to Teach?,” *American Educator*, Summer 2007.

cannot catch a fluent error in the subject they claim to have studied is credentialed, not ready.

Room two: vendor tutor-for-every-child versus union gold-standard privacy versus UNESCO's warning not to paper over teacher shortages. xAI's 11 December 2025 announcement promised Grok — the assistant, not Grok Bot — across more than 5,000 public schools in El Salvador and “personalized learning to over one million students.” No outcome evaluation, independent audit, or classroom study was found as of 30 August 2026.⁵¹⁷ AFT's *Commonsense Guardrails Value 1* puts student privacy and safety first, “even if it means sacrificing some potential educational opportunities.”⁵¹⁸ UNESCO's 2023 guidance is human-centred and treats under-13 and primary use of untested tools as “not appropriate”; later rights-of-learners materials warn governments not to divert education budgets into AI while schools remain unconnected.⁵¹⁹ Those three positions are not a graduate profile. They are the politics the graduate will walk into. A profile that ignores them will train students to use tools their future workplaces, unions, and ministries have not agreed they should have.

Room three: whether graduate research agents belong in methods courses in year 1, or only after students can catch the model. Hermes Agent ships a research-paper-writing skill. Codex Cloud and Claude Code will plan, browse, and submit a repository. The question is not whether those tools exist. They exist. The question is when a methods course may treat them as legitimate collaborators. This book's position is after the student can catch the model — after unaided formation has produced enough domain knowledge that Facione's evaluation is not a coin flip, and after transcripts and oral defence are already habits. That is a lean, not a trial. There is no identified learning RCT on Hermes-class multi-step agents, and no education-specific audit of a deployed school agent with open-web tools.⁵²⁰ Putting the agent in year 1 because industry uses it is the UCLA default with a research-methods syllabus. Putting it after the student can catch the model is Stanford's refusal, generalized.

Grok Bot, launched 11 August 2026 and expanded 26 August 2026 to SuperGrok plans and Cursor Pro/Teams, does not settle this. The product pages illustrate sales, CRM, inbox, expenses, bug reproduction, website deploy. None of those jobs is educational. The pages do not mention schools, students, or teachers. There is no education SKU.⁵²¹ A graduate profile that promises the student

Critical thinking is not a transferable skill like riding a bicycle; processes of thinking are intertwined with domain knowledge. Full treatment in Chapter 6.

⁵¹⁷xAI, El Salvador announcement, December 11, 2025: Grok (the assistant) across more than 5,000 public schools over two years, “personalized learning to over one million students.” Announcement, not a documented completed deployment; no outcome evaluation found as of August 30, 2026. Not Grok Bot. Landscape §3; brief correction 8.

⁵¹⁸American Federation of Teachers, *Commonsense Guardrails for Using Advanced Technology in Schools*, Version 2 (March 2025). Value 1: student privacy and safety come first, “even if it means sacrificing some potential educational opportunities.”

⁵¹⁹Fengchun Miao and Wayne Holmes, *Guidance for Generative AI in Education and Research* (UNESCO, September 7, 2023). Age 13 for independent GenAI conversation; under-13 / primary use of untested tools “not appropriate.” UNESCO, *AI and Education: Protecting the Rights of Learners* (2024), on connectivity and budget diversion.

⁵²⁰Brief, open question 10; landscape and pedagogy memos. No identified learning RCT on Hermes-class / Grok Bot-class multi-step agents. No education-specific audit of a deployed school agent with open-web tools. IEEE AIEI 2026 survey noted an empirical gap on fully agentic platforms (doi:10.1109/aiei69164.2026.11497428; full text not fetched).

⁵²¹xAI, “Introducing Grok Bot,” August 11, 2026, <https://x.ai/news/introducing-grok-bot>; expansion August 26, 2026, to SuperGrok plans and Cursor Pro/Teams. “Bots have their own computer... keep working 24/7.” Jobs il-

can “manage a team of Grok Bots” is inventing a positioning the product did not have. This chapter will not invent it.

“Can prompt ChatGPT” fails all three rooms. It is not Cornell’s question. It is not two-lane assessment. It is not humans in the loop. It is not ISTE’s six roles sitting on domain knowledge. It is not a reason to put an unattended research agent in a first-year methods course. It is the behavior UCLA already measured. The rest of the chapter is what the upgrade looks like in deeds.

The upgrade is not a new major. Workforce-entry literacy is not a computer-science degree, and this chapter will not smuggle one in under the word “agent.” Most graduates will not train models, will not write runtimes, and will not be asked to. They will be asked to use tools that finish jobs, and to remain the person who can tell when the job was finished wrongly. That is a formation problem, not a prompting problem. The student who can only prompt is the student Bastani’s unguarded arm already produced: better in practice, worse when the tool is taken away, and unaware of the loss.⁵²² The student who can pause, catch, and keep a transcript is the student the two-lane rule is for.

Default use will not produce that student. UCLA’s 67 percent is brainstorming, topic research, and studying for exams. Pew’s 26 percent of U.S. teens using ChatGPT for schoolwork, doubled from 2023, is the same behavior one band earlier.⁵²³ Estonia built a national program in part because 64 to 90 percent of students were already using free commercial bots to finish homework faster — the pattern the research says deskills, which the AI Leap is trying to interrupt, not scale.⁵²⁴ A graduate profile that baptizes that pattern as fluency has given up the job. The deeds that follow are the interruption, written as capabilities rather than as a ministry slogan.

End of high school — eight capabilities

ISTE unveiled a Profile of an AI-Ready Graduate at ISTE Live 2026 in Orlando. Six roles: Learner, Researcher, Synthesizer, Problem Solver, Connector, Storyteller. The profile is discipline-agnostic and built on the ISTE Student Standards. Practice reported around it includes interrogating sources with the model as coach; combining texts then verifying; English-language-learner conversation practice; generating images after the student draft.⁵²⁵ Culatta’s 14 July 2025 ASCD blog names

illustrated: sales, CRM, inbox, expenses, bug reproduction, website deploy. No education SKU; product pages do not mention schools, students, or teachers. Grok Bot FAQ, <https://docs.x.ai/grok-bot/faq>, re-fetched after the landscape timeout: cloud data storage required; Legacy Privacy Mode not supported; training opt-out follows Cursor settings; no schools, students, or teachers. Still no education SKU.

⁵²²Bastani et al., “Generative AI without Guardrails,” *PNAS* 122, no. 26 (2025). GPT Base practice about +48 percent, unaided exam about −17 percent versus never-AI; GPT Tutor practice about +127 percent, unaided exam indistinguishable from control. Students in the unguarded arm did not perceive the loss. Full table in Chapter 3; this chapter uses the dissociation, not the cells.

⁵²³Pew Research Center, “About a Quarter of U.S. Teens Have Used ChatGPT for Schoolwork,” January 15, 2025: 26 percent of U.S. teens ages 13–17 used ChatGPT for schoolwork (13 percent in 2023). Chat use, not agent use.

⁵²⁴Estonia AI Leap / TI-Hüpe problem statement, as recorded in the responsible-use memo and brief Chapter 9: 64–90 percent of students already using free commercial bots to finish homework faster — the pattern the program is trying to interrupt. First-year review published August 7, 2026; the page itself withholds a learning-outcomes claim; Tartu coefficients unpublished as of August 30, 2026.

⁵²⁵“ISTE Live 26: What Does an AI-Ready Graduate Look Like?,” *EdTech Magazine*, July 2026, <https://edtechmagazine.com/k12/article/2026/07/istelive-26-what-does-ai-ready-graduate-look>. Conference report of the ISTE Profile of an AI-Ready Graduate: Learner, Researcher, Synthesizer, Problem Solver, Connector,

the same six headings: Learner, Researcher, Synthesizer, Problem Solver, Connector, Storyteller. ISTE's Generation AI news page summarizing that blog lists Ideator in place of Problem Solver; some 2025 classroom-practice writeups follow the Ideator substitution. This book uses the ISTE-Live 2026 six as printed in EdTech Magazine, which match Culatta's ASCD headings. Do not cite Culatta 2025 as an Ideator list. The six roles are a conference profile, not an outcomes study. They are a vocabulary for workforce-entry literacy, not a measurement of judgement, and not a trial of transfer onto non-AI critical thinking. That transfer RCT does not exist. Chapter 6 left it standing. This chapter leaves it standing too.

AI4K12, the AAI/CSTA grade-band charts, is the other American vocabulary that belongs in a high-school close. By grades 9–12 the charts ask students to run a machine-learning process; evaluate fairness, transparency, accountability, and privacy; and evaluate AI content for accuracy, origin, and citation ethics.⁵²⁶ Those are deeds about systems, not only deeds with a chatbot. They sit next to ISTE, not underneath it.

What follows is a synthesis of the ISTE Live 2026 six, the AI4K12 9–12 charts, the Department of Education's humans-in-the-loop line, Cornell's question, and Chapter 5's two-lane rule. Capabilities 7 and 8 — civic/ethical, and record hygiene — are this book's add from the Department, from AFT, from New York City Public Schools, and from the Joint Council for Qualifications. Say so on the page: they are not ISTE's six. The whole list is not a validated assessment instrument. No state, no district, and no testing consortium has operationalized it as a score. Print it as deeds the school can assign and grade, not as a credential the book has invented.

1. **Learner.** Uses AI to extend practice and feedback without surrendering the struggle, and can pause the tool. This is ISTE's Learner, Chapter 5's formation lane, and Bastani's lesson under another name. Unguarded GPT-4 during high-school math practice raised scores about 48 percent; the same students then scored about 17 percent worse on an unaided exam than classmates who never had AI. A hint-and-grounded-solution tutor raised practice about 127 percent and left the unaided exam indistinguishable from control. Students in the unguarded arm did not perceive that they had learned less.⁵²⁷ The Learner who cannot pause the tool is Bastani's GPT Base. The Learner who can is closer to GPT Tutor: the struggle stays inside the student, the hint ladder is a designed refusal, and the grade that matters is the unaided one. Pausing is a deed. It is also a two-lane deed: formation work is where the pause is required; AI-world performance is where the pause is a judgement about whether the student can still detect model error.
2. **Researcher.** Treats outputs as claims to be sourced, dated, and contradicted. This is ISTE's Researcher and AI4K12's 9–12 demand to evaluate AI content for accuracy, origin, and

Storyteller. Not an outcomes study. Richard Culatta, "Profile of an AI-Ready Graduate," ASCD, 14 July 2025, <https://www.ascd.org/blogs/profile-of-an-ai-ready-graduate> — same six headings as ISTE Live 2026 (Problem Solver, not Ideator). ISTE Generation AI news page summarizing that blog lists Ideator in place of Problem Solver: <https://iste.org/generationai/news-updates/profile-of-an-ai-ready-graduate>. Do not cite Culatta 2025 as an Ideator list.

⁵²⁶AI4K12 (AAI/CSTA), "Grade Band Progression Charts — Five Big Ideas in AI," <https://ai4k12.org/gradeband-progression-charts/>. Working drafts, public. Grades 9–12: run a machine-learning process; evaluate fairness, transparency, accountability, and privacy; evaluate AI content for accuracy, origin, and citation ethics. Not an efficacy study.

⁵²⁷Bastani et al., *PNAS* 122, no. 26 (2025): e2422633122, as in note 12. Intention-to-treat, classroom-clustered, nearly 1,000 students, grades 9–11, four 90-minute sessions, large Turkish high school, Fall 2023–24.

citation ethics. A fluent paragraph is not a source. It is a claim that looks like a completed judgement. Hallucination makes this harder than Wikipedia, because the surface form has no history tab. The high-school Researcher who pastes a model paragraph into a paper with a URL has not researched. The one who dates the output, names the tool, checks it against a primary document, and can say where the model inferred or fabricated, has. Chapter 6's restatement test is the cheap version: if the student cannot put the claim in their own words without looking at it, they do not have a claim to evaluate.

3. **Synthesizer.** Combines multiple texts or data with the model, then verifies against primary sources. This is ISTE's Synthesizer. The *then verifies* clause is the entire skill. Willingham again: there is nothing to synthesise if the student does not know the domain well enough to notice a seam.⁵²⁸ A synthesis the model performed, which the student cannot mark as sourced, inferred, or fabricated, is a performance of synthesis. History is the sharpest illustration even though almost no subject-specific generative-AI randomised trial exists as of this research's access date. The student reads the sources first. The model produces a narrative. The student marks the seams and defends the marks. The dulling design is "write a DBQ with ChatGPT." No trial is required to see which design grades the student.
4. **Problem solver.** Decomposes a real task, assigns subtasks to self versus tool, and checks. This is ISTE's Problem Solver and the two-lane rule in operational form. The assignment of subtasks is the lane choice. Formation work stays with the student. Performance work may go to the tool once the student can detect error. Copilot studies are the warning, not the celebration: novices shepherd or stray; benefits accrue to students with stronger metacognitive skills; students with metacognitive difficulties can be negatively affected.⁵²⁹ A problem solver who assigns everything to the tool and checks nothing is not ISTE's role. They are Prather's stray with a diploma.
5. **Connector.** Uses AI to draft communications and to practise language, then shows up as a human in the real conversation — counselor, teammate, employer. This is ISTE's Connector. New York City Public Schools' traffic-light guidance keeps the human counselor as the decision-maker; algorithmic pathway recommendations can be overridden.⁵³⁰ The Connector who lets the model speak to the college, the employer, or the counselor in the student's name has offloaded the relationship the role is named for. Draft, then appear. The appearing is the deed. Practice language with a constrained tutor, then speak to a person. That is also the AFT line against companion chat until 16, written as a graduate habit rather than as a ban: the conversation that counts is with a human who is responsible for it.

⁵²⁸Willingham, *American Educator*, Summer 2007, as in note 6 and Chapter 6.

⁵²⁹James Prather, Brent N. Reeves, Juho Leinonen, Stephen MacNeil, et al., "'It's Weird That It Knows What I Want': Usability and Interactions with Copilot for Novice Programmers," *ACM Transactions on Computer-Human Interaction* (2023); Prather et al. 2024 laboratory follow-on. Novices shepherd or stray/drift; 2024: benefits accrued to students with stronger metacognitive skills; students with metacognitive difficulties negatively affected. No large multi-site RCT showing Copilot raises conceptual CS learning.

⁵³⁰New York City Public Schools, "Guidance on Artificial Intelligence," March 2026, <https://www.schools.nyc.gov/about-us/vision-and-mission/artificial-intelligence/guidance-on-artificial-intelligence>. Human counselor remains the decision-maker; algorithmic pathway can be overridden. Red list includes grading as the decision, IEPs, discipline, counseling. Grade-band student-use Playbook still forthcoming as of the March 2026 page.

6. **Storyteller / maker.** Owns the voice; tools are production assistants after the student's draft. This is ISTE's Storyteller, and it is Contractor and Reyes's augmentation pattern rather than their automation pattern. In that July 2026 working paper — not a journal article — AI access raised immediate knowledge-test scores 0.27 standard deviations, persisting one week; delayed unaided essay style and relevance improved for students who used AI to explain, not for students who used AI to generate text, whose short-run quality gains vanished when the tool was removed.⁵³¹ The Storyteller who begins with the model's draft has automated the voice. The one who drafts, then uses the tool to produce, illustrate, or copy-edit, has kept the thing the role names. Hu-Ann Wren's ISTE Live practice — students drafting in their own words, images after — is conference-reported classroom practice, not a districtwide agent program, and not an elementary mandate. It is the right sequence at any age the tool is lawful.
7. **Civic / ethical.** Can explain why under-13 and companion chat are restricted; can talk about bias without pretending a checklist is a finding; has a basic awareness that unions and the largest U.S. district have named energy as a cost of school-scale inference; and knows when *not* to use AI — exams, confidential data, counseling, other people's images. This capability is not ISTE's sixth role. It is this book's civic add from UNESCO, AFT, and NYCPS. AFT's 18 July 2026 resolution would ban social-companion chatbots at least until age 16 and opposes student-facing AI in elementary.⁵³² Companion harm is litigation, not a hypothetical: Sewell Setzer III, 14, Character.AI, died 28 February 2024; suit; a federal judge on 21 May 2025 refused a First Amendment dismissal; settlement disclosed 7 January 2026. The parents of Adam Raine, 16, sued OpenAI on 26 August 2025. These are allegations, not verdicts, and they are why the FTC issued 6(b) study orders on 11 September 2025 and why Character.AI pulled open chat for under-18s on 24–25 November 2025.⁵³³ The graduate who cannot explain those restrictions is not civically literate about the tool. Environmental cost: AFT and NYCPS flagged energy; independent kilowatt, water, and carbon analyses were not fetched for this book. Keep the awareness at that level. Do not invent a number.
8. **Record hygiene.** No personally identifiable information in consumer tools; discloses AI use per the syllabus; keeps transcripts when required. This capability is not ISTE's either. It is this book's hygiene add from FERPA practice, from JCQ's acknowledgement format, and from Stanford's transcript requirement as a preview a high school can already run. JCQ, effective 30 April 2025, still allows detectors as one check but forbids AI as the sole marker; acknowledgement is tool name, URL, and date; retain a non-editable prompt and output; a

⁵³¹Zara Contractor and Germán Reyes, "Experimental Evidence on the Learning Impact of Generative AI," IZA Discussion Paper 18792 (July 2026). Working paper, not a journal article. Full IZA PDF fetch was partial; coefficients from abstract/IZA page. Augmentation (AI to explain) versus automation (AI to generate text).

⁵³²American Federation of Teachers, "Preparing for the Age of Artificial Intelligence," resolution adopted July 18, 2026. No screens PK–2 except disability/compelling need; oppose student-facing AI in elementary; ban social-companion chatbots at least until age 16.

⁵³³*Garcia v. Character Technologies* (M.D. Fla.): Sewell Setzer III, 14, died February 28, 2024; suit October 2024; Judge Anne C. Conway, May 21, 2025, refused First Amendment dismissal; settlement disclosed January 7, 2026. *Raine v. OpenAI* (San Francisco Super. Ct.), filed August 26, 2025, Adam Raine, 16. Allegations in litigation, not adjudicated causation. FTC 6(b) study orders, September 11, 2025, to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI. Character.AI open-ended chat withdrawn for under-18s from November 24–25, 2025. Do not reconstruct conversations.

brief explanation of use.⁵³⁴ That is a transferable habit, not a British exam quirk. The high-school student who can produce the artefact JCQ describes is practising the college transcript requirement before anyone calls it that. The student who pastes schoolwork into a consumer account that trains on inputs has failed the hygiene before they have failed the essay.

Those eight are deeds. They can be assigned. They can be graded in two lanes. They cannot, on present evidence, be sold as an employment forecast. A district that turns the list into a badge without unaided assessment has built another UCLA default with a nicer rubric.

The list is also a refusal of three substitutes. It is not a CS major: running an ML process in AI4K12 9–12 is literacy about systems, not a credential in training them. It is not a detector regime: Liang and colleagues found detectors misclassified 61.22 percent of TOEFL essays by non-native writers as AI-generated; assessment redesign is the surviving policy, which is why hygiene is disclosure and transcripts, not a score.⁵³⁵ It is not a companion: the Connector shows up as a human; the civic capability can explain why the other kind of conversation is restricted. A profile that cannot refuse those substitutes will be used as a marketing sheet for all three.

What college adds

College adds, rather than repeats, the high-school profile. The student is now the FERPA eligible student; rights transferred at 18 or upon enrollment in postsecondary. COPPA is off. The institution's job is not to re-teach the six roles. It is to make the student the person who chooses tools per discipline and data class, internalizes field norms, and can refuse high-risk uses without a teacher in the room. The following table is the age-stages synthesis already earned in Chapter 4. It is not a new instrument.

	End of high school	End of college
Scope	Assigned tools, teacher-visible	Choose tools per discipline and data class (public versus FERPA / HIPAA / IP)
Integrity	Disclose when the teacher asks	Internalize field norms; cite; keep methods reproducible
Agency	Don't let it do the homework	Design workflows where the model is a critic, not an author of record
Research	Check a claim	Interrogate training-data limits, evaluation, and when <i>not</i> to use generative AI (original empirical work, confidential)

⁵³⁴Joint Council for Qualifications, *AI Use in Assessments*, effective April 30, 2025. Detectors as one check; AI cannot be the sole marker; acknowledgement = tool name + URL + date; retain non-editable prompt/output; brief explanation of use.

⁵³⁵Weixin Liang et al., "GPT Detectors Are Biased against Non-Native English Writers," *Patterns* 4, no. 7 (2023). Detectors misclassified 61.22 percent of TOEFL essays by non-native writers as AI-generated while remaining near-perfect on U.S. 8th-grade essays. Assessment redesign, not detection, is the surviving policy (Chapter 8).

	End of high school	End of college
Civic	Know the teen rules	Can reason about high-risk uses (hiring, credit, admissions) and refuse to build them carelessly
Career	ISTE six roles at entry level	Domain-specific: code assistants with review, lab protocol checks, client-memo drafts with human sign-off
Legal self	Parent and school still in the FERPA loop until 18	Student is the rights-holder; manages own data exhaust

Read down the college column and the pattern is the same: the adult is now the controller of the loop that high school still shared with a parent and a teacher. Scope is classification, not access. Integrity is a field norm, not a reminder on the board. Agency is critic-not-author, which is Contractor and Reyes's augmentation written as a professional habit. Research is knowing when generative AI is the wrong instrument — original empirical work, confidential data, another person's record — not a more elaborate prompt. Civic is Annex-III-shaped even for a student who will never work in the Union: hiring, credit, admissions are high-risk uses; the graduate who cannot refuse to build them carelessly has failed the civic add at adult scale. Career is domain-specific sign-off, which is humans-in-the-loop with a job title. Legal self is the eligible student managing exhaust the school used to manage.

EDUCAUSE's 2026 Horizon Report, Teaching and Learning Edition, is listed as treating AI as reshaping assessment toward authentic, process-based demonstrations of learning. The full PDF was not fetched. Use that at the level of a listed takeaway, not as a table of coefficients.⁵³⁶ It is consistent with the two-lane rule and with QAA's 2023 advice to cut AI-vulnerable tasks and rebuild authentic, synoptic assessment. It is not a finding that American colleges have done so.

Khlaif and colleagues, in *Education Sciences* 2025, interviewed 61 faculty and described four stances: Against, Avoid, Adopt, Explore. Motivations included integrity, career preparation, technological change, and institutional policy.⁵³⁷ Sixty-one faculty is a qualitative sample. It is not a national census, and this chapter will not treat it as one. What it is useful for is the reminder that faculty stances are mixed for reasons the graduate will meet: some will forbid, some will require, most will condition. Harvard's Faculty of Arts and Sciences Office of Undergraduate Education already requires a Canvas generative-AI policy and offers three model policies — forbid; allow with citation; assignment-specific.⁵³⁸ That is the syllabus floor, not a fourth case study. The col-

⁵³⁶EDUCAUSE, 2026 *EDUCAUSE Horizon Report, Teaching and Learning Edition*, <https://library.educause.edu/resources/2026/5/2026-educause-horizon-report-teaching-and-learning-edition>. Full PDF not fetched. Listed takeaway only: AI reshaping assessment toward authentic, process-based demonstrations of learning.

⁵³⁷Zuheir N. Khlaif et al., *Education Sciences* 15, no. 2 (2025), doi:10.3390/educsci15020174. Qualitative, n = 61 faculty; Against / Avoid / Adopt / Explore. Not a national census.

⁵³⁸Harvard FAS Office of Undergraduate Education, Generative AI Guidance, <https://oue.fas.harvard.edu/faculty-resources/generative-ai-guidance/>. Required Canvas policy; three model policies (forbid / allow with citation /

lege graduate who can only prompt will fail the forbid syllabus and fake the allow syllabus. The graduate who can choose a lane will survive all three.

The college add is also a data-classification add. UCLA's Digital and Technology Solutions guidance is explicit that other students' work, patient data, unpublished data, and FERPA records do not go into public tools.⁵³⁹ A high-school hygiene habit — no PII in consumer tools — becomes, at college, a professional refusal: this corpus, this lab notebook, this client file, this unpublished result is the wrong class of data for this tool. The student is now the rights-holder. Managing one's own data exhaust is a deed the high-school profile only previews.

Whether graduate research agents belong in year-1 methods remains open. The book's lean does not close it. After the student can catch the model is a constraint, not a week-by-week schedule. A biology methods course that uses a constrained, logged assistant to critique a student-written methods section is a different object from a course that lets Hermes Agent write the paper overnight. Mapping a skill onto a product page is not adoption, and it is not a learning result.

Worked examples: Stanford, UC San Diego, Cornell

Three campuses, three different deeds, none of them an agent-learning RCT. Do not claim they are.

Stanford CS221, Spring 2026, is college-level agent *refusal*. The policy is already on the table: collaborator, not answer key; no asking for answers; no “check my work”; no using an agentic coding tool such as Codex or Claude Code to build a repository from a project proposal; transcripts required; honor-code enforcement. The course is Artificial Intelligence: Principles and Techniques. If any campus might have been expected to treat Codex or Claude Code as the point of the assignment, it is this one. It did the opposite. The project is where the skill is formed; an agent that plans, browses, and submits the repository is a different offloading regime from autocomplete; there is no identified learning RCT on that regime; therefore the agent is out of bounds for the work the grade is supposed to measure. That is not a ban on AI. It is a ban on using an unattended runtime to skip the formation the course exists for. High schools can preview the transcript half tomorrow. They cannot preview the repo-builder ban as a fashion. They can preview it as a rule: if the assignment is supposed to form the skill, the tool that finishes the job is the wrong tool for that assignment.

UC San Diego's Computer Science and Engineering sequence is a *manual-competency gate*. Faculty (Porter, Shah), in a Microsoft customer story, redesigned introductory courses so students first demonstrate manual competency, then use Copilot. The same story reports research on trust calibration and plans to prepare students for “emerging agentic AI.”⁵⁴⁰ Label it. It is a vendor customer story, not a peer-reviewed trial of Cloud Agents, and not a finding that Copilot raises conceptual CS learning. No large multi-site RCT shows that it does. What the story is usable for is the gate: un-

assignment-specific). Syllabus-floor pattern, not a fourth case study in this chapter.

⁵³⁹UCLA Digital and Technology Solutions, AI Use and Recommendation Guide, <https://www.dts.ucla.edu/initiatives/ai/ai-use-recommendation-guide>. Other students' work, patient data, unpublished data, and FERPA records do not go into public tools.

⁵⁴⁰Microsoft customer story, “University of California San Diego — GitHub Copilot,” <https://www.microsoft.com/en/customers/story/26539-university-of-california-san-diego-github-copilot>. Faculty Porter and Shah; manual competency first, then Copilot; trust-calibration research; plans to prepare students for “emerging agentic AI.” Vendor customer story, not a peer-reviewed trial of Cloud Agents.

aided formation first, then a declared tool, then — only then — talk of agents. That is the two-lane rule in a CS curriculum, written by faculty who teach the course, published on a vendor domain. Steal the gate. Do not steal the URL as evidence.

Cornell is institution-scale literacy that asks whether use contributes to learning, not whether the student can call an API. The August 2026 *Chronicle* expansion to all incoming students, the four Canvas modules, the CTI-and-Library home, the faculty spread across writing, business, biology, engineering, economics, communication, and information science, and the CTI-reported pilot gains in how models work and are trained, are a program. They are not a published RCT. They are not a transfer study onto non-AI critical thinking. They are the right question, asked at the start of college, in enough departments that the question is not owned by computer science. That is the worked example this chapter needed from a literacy program: Cornell did not build a prompt academy. It built a decision about learning.

Harvard OUE sits beside these three as the syllabus-floor pattern, not as a fourth case. Required Canvas policy; three models; the instructor sets the lane in public. A campus that has Stanford's refusal on the books in one department, UCSD's gate in another, Cornell's question for everyone, and Harvard's floor on every syllabus, has a graduate environment. A campus that has a Copilot license and a detector has a consumer default with a judicial process. The three worked examples are not a ranking. They are a reminder that refusal, a competency gate, and a learning question are all graduate behaviors, and that none of them requires an education SKU from Grok Bot.

None of the three campuses has an agent-learning RCT. Clemson News and UT-Knoxville, via EdTech Magazine, document access — ChatGPT Edu, GPT-5.5, Codex, locally hosted open-weight models; UT Verse details were not fetched against primary docs and remain journalism-sourced, UNVERIFIED — and the rhetoric that students must be AI-proficient at graduation.⁵⁴¹ Access is not an outcomes study. Proficiency-at-graduation, if it means prompt skill, is the UCLA default. Proficiency-at-graduation, if it means the table in the previous section, is this chapter. Do not mash them.

What the graduate profile will not promise

The eight-capability list is a synthesis. It is not a validated assessment instrument. No dataset fetched for this book shows better later employment or civic outcomes for students who can do these things. A chapter that invented that dataset would be a different book, and a worse one.

It will not promise that an AI-ready graduate can manage a team of Grok Bots. Grok Bot, as of 26 August 2026, had no education positioning. Hermes Agent is not a school product. There is no documented K–12 student-facing autonomous open-web agent program. Broward County's K–5 have zero district-approved student-facing AI tools. Stanford banned the repo-builder on the project. No Hermes or Grok Bot district DPA was found.⁵⁴² The graduate who can pause the tool,

⁵⁴¹Clemson News documents institutional access (ChatGPT Edu, Codex, locally hosted open-weight models), not a published learning-outcomes study. UT-Knoxville / UT Verse details in *EdTech Magazine* (December 2025) were not fetched against primary docs — journalism-sourced, UNVERIFIED.

⁵⁴²Broward County Public Schools, Board Resolution No. 26-132 (March 10, 2026), <https://browardschools.ai/>. K–5: no district-approved student-facing AI tools. No Hermes Agent or Grok Bot institutional AUP, DPA, or classroom SOP found (open question 14). No documented K–12 student-facing autonomous open-web agent program (brief correction

catch the model, and keep transcripts is the book's actual close. A teammate of agents is a product page.

It will not promise year-scale unaided skill on independent state tests under student-facing generative tutors. Khanmigo's TCAP outcome remains outstanding. Nigeria and Kestin did not follow a year later. Bastani measured the next hour. It will not promise a What Works Clearinghouse report on generative tutors; none exists as of 30 August 2026. It will not promise a Hermes-class learning RCT; none was identified. It will not treat ISTE's six roles as an outcomes study, Cornell's CTI-reported gains as a published trial, or UCSD's Microsoft story as a Cloud Agents experiment.

It will not promise that "evaluate the chatbot" transfers to non-AI critical thinking. That RCT is not in the record. Willingham's constraint predicts that abstract literacy will fail and domain-loaded practice might work. The honest move is to require the practice, grade the student's judgement, and leave the transfer question standing.

It will not promise an optimal mix of AI-forbidden and AI-required tasks by grade. The two-lane rule is theoretically motivated, not a fitted schedule. The mix remains open question 11. This chapter names it. It does not fill it.

Default behavior — UCLA's 67 percent, brainstorming, topic research, and studying for exams — is the thing to upgrade. The upgrade is deeds, not a badge, and not a CS major. The high-school graduate who can pause, source, verify, assign subtasks without surrendering the hard ones, appear as a human, own a draft, explain the civic restrictions, and keep a transcript, is ready for the college column. The college graduate who can classify data, internalize field norms, use the model as critic not author, know when not to use it, refuse high-risk builds, sign off in a domain, and manage their own exhaust, is the adult version of the same person. Neither graduate is promised a job by this list. Both are promised a remaining job: the work that is still required when the tool is taken away, and the judgement that is still required when it is not.

The coda states what an honest book can promise, and what it cannot. This chapter has already spent most of the cannot. What it can still say, without a forecast, is the opening in reverse. Stanford CS221 refused the agent on the assignment that forms the skill. That refusal is available to every syllabus that still has a skill to form. The graduate who can make it, without being told, is the one this book was for.

Coda

What an Honest Book Can Promise versus What It Cannot

The collision is already visible. It does not require a product tour, a 2030 classroom, or a sentence about the future being bright if we are thoughtful. Policy red lines are hardening — grading, IEPs, surveillance, elementary student-facing AI, companion chat — just as agent autonomy is increasing. New York City red-lights the decisions that make a person the teacher of record. Broward County’s K–5 have zero district-approved student-facing AI tools. AFT, in July 2026, opposes student-facing AI in elementary and would ban social-companion chatbots at least until age 16. Character.AI pulled open chat for under-18s in November 2025 after litigation the FTC had already made into 6(b) study orders. At the same time, systems that take a goal, use tools, and keep working after the user leaves — Hermes Agent, Grok Bot, Cursor Cloud Agents, Claude Cowork, OpenAI Workspace Agents — are shipping as infrastructure. Grok Bot, as of 26 August 2026, still had no education SKU. Hermes Agent is not a school product. The collision is not a forecast that they will acquire one. It is the fact that the red lines and the runtimes are moving at once, in opposite directions, in the same year.⁵⁴³

Cognitive-offloading evidence is almost entirely about *chat* writing and tutoring tasks. Bastani, Fan, Stadler, Kosmyna, Contractor and Reyes, Oreopoulos and Low, Kestin, De Simone: windows, sessions, labs, essays. Agent offloading — unattended completion of the job — is a stronger intervention and is under-studied in education. That gap is a finding. It is not a licence to extrapolate a MIT Media Lab preprint’s EEG, $n = 54$ and session 4 on 18, onto Grok Bot. Kosmyna is a preprint, not a journal article, not K–12, not an agent, and not a delayed knowledge test of the Bastani kind. An honest book flags the gap rather than filling it with analogy. Chat evidence can constrain chat design. It cannot, without a new trial, constrain an unattended runtime that keeps working after the student leaves. The popular story will treat that as a technicality. It is the plot. Red lines written for a conversation window are already incomplete for a computer that is not the student’s laptop. Pretending the chat record covers the agent is how a district buys a teammate and cites a tutor study.⁵⁴⁴

⁵⁴³Collision as named in the book brief, 30 August 2026, and in Chapters 2, 4, 7, and 8: NYCPS March 2026 red list (grading, IEPs, discipline, counseling, surveillance); Broward Resolution 26-132 (K–5 zero student-facing AI); AFT resolution 18 July 2026 (elementary student-facing opposition; companion ban until 16); Character.AI under-18 withdrawal 24–25 November 2025; FTC 6(b) 11 September 2025. Agent infrastructure: hermes-agent.org; xAI “Introducing Grok Bot,” 11 August 2026, expansion 26 August 2026, no education SKU; Cursor Cloud Agents docs; Anthropic Claude for Teachers, 14 July 2026; OpenAI Academy Workspace Agents, 23 April 2026.

⁵⁴⁴Offloading record as spent in Chapter 3. Chat and tutor tasks, not unattended agents. Kosmyna et al.,

What this book can promise is narrower than a manifesto and more useful. Promise, here, does not mean a prediction. It means a claim the chapters already stood on, with named documents, that survives a second reading after the open questions have been left standing.

It can promise four product classes that must not be collapsed. Chatbot, tutor app, adaptive system, agent. Named objects in the last row: Hermes Agent, Grok Bot, Cursor Cloud Agents, Claude Cowork, OpenAI Workspace Agents. Almost none of that row is in the K–12 classroom as an unconstrained student agent. What schools bought is teacher copilots and constrained Socratic tutors. Tool use is not autonomy. Hermes 4 is not Hermes Agent. Grok-as-tutor in El Salvador is not Grok Bot. Khanmigo, MagicSchool, SchoolAI, Gemini Education, and ChatGPT for Teachers are not agents. A book that keeps those distinctions is already rarer than it should be.⁵⁴⁵

It can promise a causal record in which design is the treatment. Unguarded GPT can harm subsequent math learning: Bastani’s GPT Base, about 48 percent better in practice, about 17 percent worse on the unaided exam, students unaware of the loss. Guardrails can erase the harm without, in that design, beating control: GPT Tutor. Engineered tutors can raise immediate conceptual learning versus even strong active learning: Kestin’s PS2 Pal, quantile 0.73 to 1.3 standard deviations on an immediate post-test, delayed retention unmeasured, not ChatGPT. Access is not tutoring: Oreopoulos and Low’s Khanmigo intention-to-treat looks like Khan Academy practice without the chatbot; the median student messaged on one-third of practice days and in 17 percent of mistake sessions. Writing often improves the artefact and not the knowledge: Fan, metacognitive laziness; Stadler, lower load and worse argument. Augmentation beats automation on delayed unaided essay traits in one working paper: Contractor and Reyes, July 2026, not a journal article. Human high-dosage tutoring remains the benchmark, 0.288 SD in Nickow, Oreopoulos, and Quan 2024. That is a record. It is not a law about agents.⁵⁴⁶

It can promise that age and law are design specifications, not afterthoughts. COPPA under 13; UNESCO 13 for independent classroom generative AI; FERPA always, rights transfer at 18 or postsecondary. Companion harm is litigation, named as litigation: Setzer, Raine, Conway 21 May 2025, settlement disclosed 7 January 2026, Raine complaint 26 August 2025, allegations not verdicts. Broward K–5 at zero student-facing tools is a large-district floor, not a thought experiment.⁵⁴⁷

It can promise a two-lane assessment rule that matches the cognitive science, labelled as theoretically motivated: formation, AI forbidden or hint-only, grade unaided work; AI-world performance, allowed, logged, grade judgement. It is not a fitted grade-by-grade schedule. It is not existing district policy. Perkins’s AI Assessment Scale is shared vocabulary, not an RCT.⁵⁴⁸

arXiv:2506.08872, preprint, n = 54 / session 4 n = 18; do not extrapolate EEG onto Grok Bot or Hermes Agent.

⁵⁴⁵Four classes: landscape memo table; Chapter 2. Hermes 4 (arXiv 2508.18255, 25 August 2025) “not recommended for use inside Hermes Agent” (Nous Portal). El Salvador: xAI 11 December 2025 announcement, Grok the assistant, no outcomes found. Do not list Khanmigo, MagicSchool, SchoolAI, Gemini Education, or ChatGPT for Teachers as agents.

⁵⁴⁶Bastani et al., *PNAS* 122(26) e2422633122 (2025); Kestin et al., *Scientific Reports* 15 (2025): 17458; Oreopoulos and Low, *EdWorkingPaper* 26-1551 (2026); De Simone et al., *World Bank PRWP* 11125 (2025); Fan et al., *BJET* 56(2) (2025); Stadler, Bannert, and Sailer, *Computers in Human Behavior* 160 (2024); Contractor and Reyes, *IZA DP* 18792 (2026, working paper); Nickow, Oreopoulos, and Quan, *AERJ* 61(1) (2024): 74–107.

⁵⁴⁷COPPA, 15 U.S.C. §§ 6501 et seq.; UNESCO 7 September 2023 guidance (age 13); FERPA, 20 U.S.C. § 1232g. Setzer / *Garcia v. Character Technologies*; Raine / *Raine v. OpenAI*. Allegations, not verdicts.

⁵⁴⁸Chapter 5 two-lane rule; Perkins et al. 2024 AIAS as vocabulary not efficacy.

It can promise that teachers are orchestrators, that the replacement slogan is folklore, and that no fetched study shows an agent substituting for a teacher of record. The Department of Education in 2023 firmly rejects replacement. NYC red-lights grading, IEPs, discipline, counseling. Every positive field result in Chapter 3 had a teacher or tutor in the loop. The off-switch belongs in a teacher's hand, and in a bargaining annex, not in a vendor dashboard.⁵⁴⁹

It can promise red lines that map to instruments: no commercial training on student work; no high-stakes decisions without a qualified human; no open-web acting agents for K–12; no emotion or companion AI for minors; no detector-only accusations; no graded homework that requires a premium personal agent the school does not issue. Liang 61.22 percent is the detector number. Pew 26 percent and Harvard 40 versus 20 are the paid-tier gap. The lines are copyable because they are sourced.⁵⁵⁰

It can promise a hosting default: the institution, or a public platform under public-law duties. Singapore's Student Learning Space; Estonia's ITI, a Socratic app that does not give answers; a district tenant of ChatGPT Edu, Gemini Education, or Claude for Teachers with a real DPA. A student-owned Hermes runtime is a home tutor, not a school requirement, and is not a child-safety certified school product. Local open-source does not by itself create a designated-safeguarding-lead alert path or close an equity gap.⁵⁵¹

It can promise a graduate profile in deeds, labelled synthesis: ISTE Live 2026's six roles, plus civic and hygiene adds, plus Cornell's question, plus the two-lane habit of pausing the tool. Not a validated instrument. Not an employment forecast. Not a teammate of Grok Bots.⁵⁵²

What it cannot promise is longer, and leaving the list standing is the honest close. These are not hedges. A hedge is a sentence that keeps the claim and adds a qualifier. An open question is a claim the book refuses to make. The popular story fills every one of the following with an analogy, a vendor deck, or a 2030 classroom. This coda will not.

It cannot promise year-scale unaided skill on independent state tests under student-facing generative tutors. Khanmigo's TCAP remains outstanding. Nigeria and Kestin did not follow a year later. Bastani measured the next hour.

It cannot promise a childhood, K–2, conversational-agent randomised trial on language, attention, or social cognition. Cognitive-science extrapolation is possible. Generative-AI trial evidence is not.

It cannot promise reading and history causal estimates comparable to mathematics. Direct generative-AI RCT evidence on learning to read in 2022–26 is thin; almost no subject-specific history trial was found.

⁵⁴⁹U.S. Department of Education, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023): “firmly rejects the idea that AI could replace teachers.” Fitzpatrick LinkedIn as folklore exhibit (Chapter 7). No fetched study of an agent as teacher of record.

⁵⁵⁰Chapter 8 red-line map. Liang et al., *Patterns* (2023), 61.22 percent. Pew, 15 January 2025, 26 percent. Harvard undergraduate survey, arXiv:2406.00833 (2024), 40 percent versus 20 percent paid.

⁵⁵¹Chapter 9 hosting default. Singapore MOE page updated 31 July 2026; Estonia TI-Hüpe / ITI; student-owned Hermes as home tutor, not school requirement, not a child-safety school product.

⁵⁵²Chapter 10 eight-capability list: synthesis of ISTE Live 2026 six, AI4K12 9–12, US ED, Cornell, two-lane, plus civic/hygiene adds. Not a validated instrument. No employment or civic-outcome dataset.

It cannot promise that “evaluate the chatbot” transfers to non-AI critical thinking. Willingham’s test is unrun as a randomised curriculum trial.

It cannot promise that proactive, multimodal, memory-bearing tutors will fix Khanmigo-style non-use. A default-open chat window did not, in Oreopoulos and Low’s Year-2 data, raise messaging.

It cannot promise a professional-development-to-pupil-outcomes chain for generative AI. Reviews stop at teacher self-efficacy. The CAL-PD analogue in Oreopoulos, Gibbs, Jensen, and Price 2024 is the best causal clue. It is not a generative-AI result.

It cannot promise equity of effects. Nigeria’s gains were larger for higher-baseline and higher-SES students, and for girls. Copilot studies suggest metacognitively weaker students may be harmed. A personal tutor for every child that widens gaps is an empirical possibility, not a slogan to refute by assertion.

It cannot promise replication of Kestin outside a highly staffed university physics course, or of Bastani outside one Turkish high school’s four sessions.

It cannot promise Kosmyrna-style neural measures in a design that also measures delayed learning. The preprint should not carry developmental policy.

It cannot promise a Hermes-class or Grok Bot-class learning RCT, or an education-specific audit of a deployed school agent with open-web tools. IEEE noted the empirical gap on fully agentic platforms. The gap is still the gap.

It cannot promise a fitted grade-by-grade mix of AI-forbidden and AI-required tasks. Two-lane is a rule of assessment design, not a schedule.

It cannot promise a What Works Clearinghouse report on generative tutors. Older ITS reports exist. A generative-tutor WWC report does not, as of 30 August 2026.

It cannot promise El Salvador outcomes. The December 2025 announcement is still the announcement.

It cannot promise a Hermes or Grok Bot district DPA. None was found.

It cannot promise causal findings in Setzer and Raine. Product-liability versus First Amendment is live. Name the litigation. Stop.

It cannot promise that union-vendor academies will civilize the vendors, or that they will capture the unions. AFT’s National Academy is a 22 July 2025 target, Microsoft-led, with OpenAI and Anthropic. The fight is open.

It cannot promise a U.S. federal student-AI statute. There isn’t one. FERPA ed-tech NPRM RIN 1875-AA15 was unpublished as of the FTC’s April 2025 COPPA preamble. EO 14277 is an adoption-and-workforce order, not a child-safety regulation. EO 14110 was rescinded on 20 January 2025.

It cannot promise a completed EU Annex III education enforcement record. The original 2 August 2026 high-risk clock was superseded by Regulation (EU) 2026/1744; stand-alone Annex III obligations now apply from 2 December 2027. Confirm the consolidated EUR-Lex text of 2024/1689 as

amended by 2026/1744 at typesetting. There is still no completed education enforcement record.⁵⁵³

It cannot promise independent energy analyses of school-scale inference. AFT and NYCPS flagged the cost. Analyses were not fetched. No kilowatt, water, or carbon number will be invented here to make the flag look like a finding.

It cannot promise NAEYC's updated technology and AI statement, still in revision with the 2012 text posted, or a dedicated AAP AI-in-schools policy, promised and not yet published. AAP 2026 is *Digital Ecosystems*.

It cannot promise Estonia's first-year review learning-outcome coefficients. The page exists, 7 August 2026, and itself states that first-year usage data "does not yet provide a basis for claiming that ITI has improved learning outcomes." Tartu findings were expected later in 2026. Program-ops usage figures on that page are not a learning RCT.

It cannot promise who is the data controller of a child's persistent agent memory on a parent-installed local runtime. That is open question 22, and it is a governance hole, not a riddle the last page gets to solve with a proverb.

Those are holes. A serious book leaves them standing. Filling them with analogy — chat evidence for agents, a preprint for a population, an announcement for an outcome, a vendor story for a trial — is how the popular story got here.

What remains, without a forecast, is the job the student must still do, the off-switch in a teacher's hand, and the institution as host. School exists to form skills that are still required when the tool is taken away: retrieval, generation, error-detection, domain knowledge, judgement, and the disposition to check. An AI world does not retire those skills. It makes them the scarce resource. Keep the struggle in the student. Assess in two lanes. Design is the treatment. The teacher turns the agent off. The school, or a public platform under public-law duties, hosts whatever student-facing tool survives the red lines. A home runtime is a home tutor. Grok Bot and Hermes Agent do not acquire an education SKU on this page because they had not acquired one on 26 August 2026, and because an honest last page does not give them one.

The collision will not be resolved by a brighter tone. Red lines will keep hardening, or they will be walked back. Autonomy will keep increasing, or a safety event will slow it. The evidence on unattended offloading in schools will arrive, or it will not. Teacher-facing agents may remain the 2026 beachhead, or a district may try to put an open-web runtime on a child's account and discover the red list the hard way. Until those questions are answered by trials, contracts, and statutes — not by analogy — the book you have just read is the constraint, not the tour. Documents did the work. The open questions are still open. The job is still the student's.

⁵⁵³Open questions 1–22 as listed in the book brief, 30 August 2026. EU Annex III: original 2 August 2026 clock superseded by Regulation (EU) 2026/1744; stand-alone Annex III now 2 December 2027; confirm consolidated EUR-Lex of 2024/1689 as amended by 2026/1744 at typesetting. No completed education enforcement record. Estonia first-year review 7 August 2026, page withholds a learning-outcomes claim; Tartu coefficients unpublished as of 30 August 2026. A.B. 1159 pending, not law. NAEYC statement in revision; AAP dedicated AI-in-schools policy not yet published.

A Note on Sources

Chicago notes run in one series through the preface, introduction, ten chapters, and coda. The bibliography at the back is the working apparatus: every work named in those notes, plus items consulted in the 30 August 2026 research files, tagged Cited, Consulted, Vendor, Preprint / working paper, or UNVERIFIED. Access date for digital sources is 30 August 2026 unless a same-day fact-check fetch is named in the note.

A flag is not a finding. Several load-bearing PDFs could not be read on that date and stay flagged. The OECD *Digital Education Outlook 2026* full PDF was Cloudflare-blocked; listing-page claims, including TALIS 2024 “37 percent of teachers using GenAI for work,” remain UNVERIFIED at full-PDF level and must not be mashed with EDUCAUSE’s 37 percent of *institutions* licensing chatbots. Fan et al. 2025 in *BJET* is cited from the ERIC abstract; the publisher PDF was blocked. Contractor and Reyes, IZA Discussion Paper 18792, is a working paper whose coefficients come from the abstract and IZA page. Hermes 4’s quantitative benchmarks in arXiv 2508.18255 were not read from the PDF. EDUCAUSE 2026 assessment and Horizon Report full PDFs were not fetched. The consolidated EUR-Lex text of Regulation (EU) 2024/1689 as amended by 2026/1744 should be pulled at typesetting; this book treats 2 December 2027 as the enacted stand-alone Annex III date and keeps the superseded 2 August 2026 clock on the page so it cannot leak back as current law. No completed Annex III education enforcement record exists either way.

Product pages for ALEKS, MATHia, i-Ready, and DreamBox were not fetched; those names identify a class. Illinois ISBE “agentic AI” language is from Franczek/JDSupra, not the ISBE PDF. UT Verse and Johns Hopkins Agent Laboratory remain journalism-sourced. Estonia’s 7 August 2026 first-year review page itself withholds a learning-outcomes claim; Tartu coefficients were unpublished. El Salvador is an 11 December 2025 announcement, not outcomes, and is not Grok Bot. OpenAI’s named-district expansion HTML timed out; extra 2026 headcounts are not added to the AFT Academy’s 400,000 five-year target.

Where a bibliographic element could not be verified — an author line, a volume, a URL — it is omitted rather than guessed. The gaps file in the research deposit records those omissions. This note does not close them.

Bibliography

Autonomous AI and Education

Michael Gannotti

Chicago Manual of Style, 17th edition, notes-bibliography system

This list is a working apparatus for a university-press or serious trade book. It is compiled from the four research source files of 30 August 2026 (01-sources.json through 04-sources.json) and from every work actually named in the manuscript notes (manuscript/00-preface.md, 00-introduction.md, ch01.md through ch10.md, coda.md). It does not invent titles, publishers, years, DOIs, URLs, or page numbers. Where a bibliographic element could not be verified from a catalog, landing page, or fetched PDF, it is omitted and recorded in GAPS.md. Errata (factcheck/front-errata.md, mid1-errata.md, mid2-errata.md, end-errata.md) were read so that cut claims are not revived.

Status tags (after each entry, in square brackets):

- **Cited** — named, quoted, or used as load-bearing authority in the present manuscript notes.
- **Consulted** — used in the research JSONs; not named on the manuscript page.
- **Vendor** — product documentation or vendor claim, labelled as such.
- **Preprint / working paper** — not a journal article as of the access date.
- **UNVERIFIED** — full text not read; listing-page, abstract, or secondary report only.

Digital sources: access date **30 August 2026**, unless a fact-check fetch the same day is noted.

Hanging-indent ready: first line flush; subsequent lines indented three spaces.

1. Statutes, regulations, executive orders, court filings

California. Assembly Bill 1159, *Student personal information* (2025–2026). Amended 21 August 2026. https://leginfo.legislature.ca.gov/faces/billNavClient.xhtml?bill_id=202520260AB1159. Accessed 30 August 2026. Senate third reading / floor process as of that date; **not enacted**. Would ban training generative AI on covered student information and would create HESIPA for higher education, proposed operative 1 July 2027. [Cited]

California. K–12 Pupil Online Personal Information Protection Act (SOPIPA). Cal. Bus. & Prof. Code § 22584. Operative 1 January 2016. https://leginfo.legislature.ca.gov/faces/codes_displaySection.xhtml

Accessed 30 August 2026. No targeted ads, no noneducational profiles, no sale. Pair with A.B. 1584 / Cal. Educ. Code § 49073.1. [Cited]

Children’s Online Privacy Protection Act of 1998. 15 U.S.C. §§ 6501–6506. Children’s Online Privacy Protection Rule, 16 C.F.R. pt. 312. <https://www.ftc.gov/legal-library/browse/rules/childrens-online-privacy-protection-rule-coppa>. Accessed 30 August 2026. Under-13 personal information. School-consent exception limited to the authorized educational purpose. [Cited]

European Commission. *Guidelines on Prohibited Artificial Intelligence Practices Established by Regulation (EU) 2024/1689*. C(2025) 5052 final. 29 July 2025. https://ai-act-service-desk.ec.europa.eu/sites/default/files/2025-08/guidelines_on_prohibited_artificial_intelligence_practices_established_by_regulation_2024_1689_en.pdf. Accessed 30 August 2026. GPAI with emotion features must not be used by deployers in schools unless the medical/safety exception applies. Secondary sources also describe a 4 February 2025 draft approval; exact final-document date not re-fetched as PDF on the mid2 pass. [Cited]

European Commission. “AI Act | Shaping Europe’s Digital Future.” Page update 3 August 2026. <https://digital-strategy.ec.europa.eu/en/policies/regulatory-framework-ai>. Accessed 30 August 2026. Application timeline including the Digital Omnibus. [Cited]

European Commission, AI Act Service Desk. “Annex III.” <https://ai-act-service-desk.ec.europa.eu/en/ai-act/annex-3>. Accessed 30 August 2026. Verbatim Annex III point 3 (education and vocational training). [Cited]

European Union. Regulation (EU) 2024/1689 of the European Parliament and of the Council of 13 June 2024 laying down harmonised rules on artificial intelligence (Artificial Intelligence Act). In force 1 August 2024. <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:32024R1689>. Consolidated HTML used in the age-stages stream: <https://eur-lex.europa.eu/legal-content/EN/TXT/HTML/?uri=CELEX:32024R1689-20260727>. Accessed 30 August 2026. Art. 5(1)(f) and recital 44: emotion recognition in educational institutions prohibited from 2 February 2025 except medical/safety. Art. 4 AI-literacy duty from the same date. GPAI model duties from 2 August 2025. Transparency (chatbot disclosure) timed 2 August 2026. Annex III point 3: high-risk systems used to determine access to education, to evaluate learning outcomes including steering, to determine the appropriate level of education, or to monitor student behaviour during tests. Original stand-alone Annex III clock was 2 August 2026. [Cited]

European Union. Regulation (EU) 2026/1744 (Digital Omnibus). Political agreement 7 May 2026; OJ 24 July 2026; in force 27 July 2026. Mid2 fact-check treats this as **enacted** and as pushing stand-alone Annex III high-risk rules (including education access, evaluation, and test-monitoring) to **2 December 2027**; Annex I product-embedded high-risk later, 2 August 2028. Intimate-image / CSAM generation prohibition timed December 2026 via the Omnibus. **Confirm the consolidated EUR-Lex text of 2024/1689 as amended by 2026/1744 before typesetting.** Do not silently keep 2 August 2026 as the only date. No completed Annex III education enforcement record. See GAPS. [Cited]

Exec. Order No. 14110, “Safe, Secure, and Trustworthy Development and Use of Artificial Intelligence,” 30 October 2023, 88 Fed. Reg. 75191 (1 November 2023). <https://www.federalregister.gov/documents/2023/10/30/24283/safe-secure-and-trustworthy-development-and-use-of-artificial-intelligence>. Accessed 30 August 2026. Directed the Secretary of Education to develop education AI resources within 365

days. Rescinded 20 January 2025 by Exec. Order No. 14148. [Cited]

Exec. Order No. 14148, 20 January 2025. Cited in the research file as the initial-repeals order that rescinded EO 14110. Confirm against Federal Register text before quoting operative language. <https://www.federalregister.gov/documents/2025/01/20>. Accessed 30 August 2026. [Cited]

Exec. Order No. 14277, “Advancing Artificial Intelligence Education for American Youth,” 23 April 2025. <https://www.whitehouse.gov/presidential-actions/2025/04/advancing-artificial-intelligence-education-for-american-youth/>. Alternate: <https://www.presidency.ucsb.edu/documents/executive-order-14277-advancing-artificial-intelligence-education-for-american-youth>. Accessed 30 August 2026. Task Force, Presidential AI Challenge, public-private K-12 resources, ESEA/HEA Title II teacher-training priority, NSF, USDA 4-H. Adoption-and-workforce order, not a child-safety regulation. [Cited]

Family Educational Rights and Privacy Act. 20 U.S.C. § 1232g; 34 C.F.R. pt. 99. <https://www.ed.gov/privacy/ferpa>. Accessed 30 August 2026. School-official exception: 34 C.F.R. § 99.31(a)(1). Redisclosure: 34 C.F.R. § 99.33. Rights transfer at 18 or upon enrollment in postsecondary education. “FERPA compliant” is a contract-plus-configuration-plus-practice posture, not a product label. [Cited]

Federal Bureau of Investigation / Internet Crime Complaint Center. “Child Sexual Abuse Material Created by Generative AI and Similar Online Tools Is Illegal.” PSA, 29 March 2024. <https://www.ic3.gov/Media/Y2024/PSA240329>. Accessed 30 August 2026. [Cited]

Federal Trade Commission. “Children’s Online Privacy Protection Rule” (final amendments). 90 Fed. Reg. 16918 (22 April 2025). <https://www.federalregister.gov/documents/2025/04/22/2025-05904/childrens-online-privacy-protection-rule>. Accessed 30 August 2026. Commission vote 16 January 2025; effective 23 June 2025; general compliance 22 April 2026 except specified safe-harbor subsections. First major COPPA Rule overhaul since 2013. Biometrics added to personal information. Separate verifiable parental consent for non-integral third-party disclosure. FTC: advertising, consideration, and “to train or otherwise develop artificial intelligence technologies” are not integral. Written retention policy; no indefinite retention; written security program. Proposed school-authorization / ed-tech exception **not finalized**, citing pending FERPA rulemaking RIN 1875-AA15 (NPRM unpublished as of the April 2025 preamble). [Cited]

Federal Trade Commission. “FTC Finalizes Changes to Children’s Privacy Rule Limiting Companies’ Ability to Monetize Kids’ Data.” Press release, January 2025. <https://www.ftc.gov/news-events/news/press-releases/2025/01/ftc-finalizes-changes-childrens-privacy-rule-limiting-companies-ability-monetize-kids-data>. Accessed 30 August 2026. [Consulted]

Federal Trade Commission. “FTC Launches Inquiry into AI Chatbots Acting as Companions.” 11 September 2025. <https://www.ftc.gov/news-events/news/press-releases/2025/09/ftc-launches-inquiry-ai-chatbots-acting-companions>. Order PDF: https://search.ftc.gov/system/files/ftc_gov/pdf/AICompanion. Accessed 30 August 2026. Section 6(b) study orders — not enforcement — to Alphabet, Character Technologies, Instagram, Meta, OpenAI, Snap, and xAI. Vote 3-0. [Cited]

Federal Trade Commission. *Policy Statement of the Federal Trade Commission on Education Technology*. https://search.ftc.gov/system/files/ftc_gov/pdf/Policy%20Statement%20of%20the%20Federal%20Trade. Accessed 30 August 2026. School authorization limited to the contracted educational service; no commercial use of school-collected data. [Cited]

Garcia v. Character Technologies, No. 6:24-cv-01903 (M.D. Fla.). Caption and docket number re-verified on the front fact-check pass, 30 August 2026. Sewell Setzer III, age 14, died 28 February 2024; suit October 2024 (AI Incident Database 826 records 22 October 2024). Senior U.S. District Judge Anne C. Conway, 21 May 2025, refused First Amendment dismissal at the pleading stage. Mediated settlement disclosed, court filing 7 January 2026. Allegations, a procedural ruling, and a disclosed settlement — not adjudicated causation. Do not reconstruct chats. [Cited]

Illinois. Student Online Personal Protection Act. 105 ILCS 85. Original P.A. 100-315 (2017); rewritten P.A. 101-516, effective 1 July 2021. <https://ilga.gov/legislation/ilcs/documents/010500850K15.htm>; prohibitions <https://ilga.gov/legislation/ilcs/documents/010500850K10.htm>. Accessed 30 August 2026. K-12. Written DPA before covered information transfers; no targeted ads, sale/rent, or non-school profiling. [Cited]

Individuals with Disabilities Education Act. 20 U.S.C. §§ 1400 et seq. Named in U.S. Department of Education, *Artificial Intelligence and the Future of Teaching and Learning* (May 2023), and in New York City Public Schools March 2026 AI guidance (IEP team, not AI). Full U.S.C. text not separately fetched. [Cited]

New York. Education Law § 2-d; 8 NYCRR pt. 121. Part 121 effective 29 January 2020. <https://codes.findlaw.com/ny/education-law/edn-sect-2-d/>. Accessed 30 August 2026. PII of students, teachers, and principals; Parents' Bill of Rights; DPA in contracts; breach notice. [Cited]

Protection of Pupil Rights Amendment. 20 U.S.C. § 1232h; 34 C.F.R. pt. 98. <https://www.law.cornell.edu/uscode/t>. Accessed 30 August 2026. [Consulted]

Raine v. OpenAI (San Francisco Super. Ct., filed 26 August 2025). Adam Raine, age 16. Alleged harm after schoolwork use of ChatGPT, including method advice and note-drafting on GPT-4o as reported. OpenAI public line: redirected to crisis resources “more than 100 times”; reviewing the filing. Allegations, not a verdict. Do not reconstruct chats. [Cited]

Rehabilitation Act of 1973, § 504; Americans with Disabilities Act, Title II, including the 2024 web rule (WCAG 2.1 AA). Named in the manuscript (ch. 4 n. 4) with CAST UDL Guidelines 3.0. Full statute texts not separately fetched. [Cited]

Texas. Education Code § 32.155 (operator security and unique identifiers). https://texas.public.law/statutes/tex._ed. Accessed 30 August 2026. [Cited]

Texas. Responsible Artificial Intelligence Governance Act (TRAIGA), H.B. 149. Effective 1 January 2026. H.B. 3512 (AI training) and S.B. 1964 (AI systems and data management), effective 1 September 2025. <https://dir.texas.gov/technology-legislation>. Accessed 30 August 2026. DIR AI code of ethics; heightened scrutiny / NIST RMF; public-facing disclosure; applies to districts as local government. [Cited]

United Kingdom. Online Safety Act 2023. Named in UK Department for Education product-safety standards as applying to live-web / user-sharing generative AI. Full statute text not separately fetched. [Cited]

U.S. Department of Education. “Final Priority and Definitions—Secretary’s Supplemental Priority and Definitions on Advancing Artificial Intelligence in Education.” 91 Fed. Reg. 18774 (13 April

2026), effective 13 May 2026. Docket ED-2025-OS-0118. <https://www.federalregister.gov/documents/2026/04/13/2026-07087/final-priority-and-definitions-secretarys-supplemental-priority-and-definitions-on-advancing>. Accessed 30 August 2026. Discretionary-grant scoring tool, not a classroom mandate. AI defined via 15 U.S.C. § 9401(3). Does not define “agentic AI.” [Cited]

U.S. Department of Education, Student Privacy Policy Office. “Eligible Student.” <https://studentprivacy.ed.gov/colleges-and-universities/eligible-student>. Accessed 30 August 2026. Rights transfer at 18 or postsecondary enrollment. [Cited]

U.S. Department of Education, Student Privacy Policy Office. “Responsibilities of Third-Party Service Providers under FERPA.” https://studentprivacy.ed.gov/sites/default/files/resource_document/file/Vendor%20Responsibilities.pdf. Accessed 30 August 2026. School-official exception four prongs; 34 C.F.R. § 99.31(a)(1)(i); direct control; authorized purposes only. [Cited]

2. Empirical research (journals, working papers, preprints labelled as such)

AICET / National University of Singapore. “The Educator as an Embedded AI Designer: Faculty Agency in GenAI Systems.” 2025. https://aicet.comp.nus.edu.sg/wp-content/uploads/2025/10/The_Educator_as_an_Embedded_AI_Designer.pdf. Accessed 30 August 2026. Intentional design, orchestrated mediation, reflective alignment. [Cited]

Bastani, Hamsa, Osbert Bastani, Alp Sungu, Haosen Ge, Özge Kabakçı, and Rei Mariman. “Generative AI without Guardrails Can Harm Learning: Evidence from High School Mathematics.” *Proceedings of the National Academy of Sciences* 122, no. 26 (2025): e2422633122. <https://doi.org/10.1073/pnas.2422633122>. Author PDF: https://hamsabastani.github.io/education_llm.pdf. Accessed 30 August 2026. Field RCT, large Turkish high school, Fall 2023-24, nearly 1,000 students, grades 9-11, four 90-minute sessions. ITT: GPT Base practice +0.137 (control mean 0.284) about +48 percent; unaided exam -0.054 (control mean 0.321) about -17 percent versus never-AI. GPT Tutor practice +0.361 about +127 percent; unaided exam -0.004, indistinguishable from control. Probe correctness about 51 percent. Students in the unguarded arm did not perceive the loss. Author-line spelling is Özge Kabakçı with cedilla on the PNAS page (correction notice 10.1073/pnas.2518204122). Full author PDF fetched. [Cited]

“Beyond Detection: Redesigning Authentic Assessment in an AI-Mediated World.” *Education Sciences* 15, no. 11 (2025): 1537. <https://doi.org/10.3390/educsci15111537>. Authenticity redesigned, not policed. Authors not listed in the research file. [Cited]

Bjork, Elizabeth L., and Robert A. Bjork. “Making Things Hard on Yourself, but in a Good Way: Creating Desirable Difficulties to Enhance Learning.” In *Psychology and the Real World*, edited by Morton Ann Gernsbacher et al., 56-64. New York: Worth, 2011. [Cited]

Bjork, Robert A. “Memory and Metamemory Considerations in the Training of Human Beings.” In *Metacognition*, edited by Janet Metcalfe and Arthur P. Shimamura, 185-205. Cambridge, MA: MIT Press, 1994. Origin of desirable difficulties. [Cited]

Bjork, Robert A., and Elizabeth L. Bjork. “Desirable Difficulties in Theory and Practice.” *Journal*

of *Applied Research in Memory and Cognition* 9, no. 4 (2020): 475-79. Difficulty is desirable only when achievable. [Cited]

Bloom, Benjamin S. “The 2 Sigma Problem: The Search for Methods of Group Instruction as Effective as One-to-One Tutoring.” *Educational Researcher* 13, no. 6 (1984): 4-16. One-to-one *human* tutoring, not a language model. [Cited]

Chen, Ouhaio, Juan C. Castro-Alonso, Fred Paas, and John Sweller. “Extending Cognitive Load Theory to Incorporate Working Memory Resource Depletion.” *Educational Psychology Review* 30 (2018): 483-501. Cited in the manuscript (ch. 5 n. 8). The research JSON’s Chen 2018 DDF/CLT bridge note is incomplete; this is the title used on the page. [Cited]

Contractor, Zara, and Germán Reyes. “Experimental Evidence on the Learning Impact of Generative AI.” IZA Discussion Paper No. 18792. July 2026. <https://docs.iza.org/dp18792.pdf>. Accessed 30 August 2026. **Working paper, not a journal article.** Proctored RCT; unaided tests immediate and one week. Abstract: +0.27 SD on knowledge tests, persisting one week; delayed unaided essay style about 0.30 SD / relevance about 0.26 SD larger for augmentation than automation users. Full PDF timed out / partial; coefficients from abstract and IZA page. [Cited; working paper; full PDF UNVERIFIED]

De Simone, Martin, Federico Tiberti, Maria Barron Rodriguez, Federico Manolio, Wuraola Mosuro, and E. J. Dikoru. “From Chalkboards to Chatbots: Evaluating the Impact of Generative AI on Learning Outcomes in Nigeria.” World Bank Policy Research Working Paper 11125. 2025. <https://documents1.worldbank.org/curated/en/099548105192529324/pdf/IDU-c09f40d8-9ff8-42dc-b315-591157499be7.pdf>. Accessed 30 August 2026. **Working paper.** RCT, 9 schools, Benin City. Copilot GPT-4 after-school English, teacher-guided science-of-learning prompts. ITT: composite 0.31 SD (SE 0.068); English 0.238 SD; regular third-term English exam 0.206 SD. Authors’ own Evans-Yuan conversion uses “0.24 SD.” Full PDF fetched. [Cited; working paper]

“Designing AI-Aware Assessment in Higher Education: A Cross-National Analysis of Generative AI Guidance and Design Patterns in English-Medium Universities.” 2026. <https://doi.org/10.17275/per.26.60.13.4>. Accessed 30 August 2026. 21 university guidance documents; four patterns (process portfolios, oral verification, critique of AI outputs, secure exams plus AI coursework). Authors not listed on the research landing. [Cited; authors omitted — see GAPS]

Deslauriers, Louis, Logan S. McCarty, Kelly Miller, Kristina Callaghan, and Greg Kestin. “Measuring Actual Learning versus Feeling of Learning in Response to Being Actively Engaged in the Classroom.” *Proceedings of the National Academy of Sciences* 116, no. 39 (2019): 19251-57. Illusion of learning; cited by Bastani on the perception mismatch. [Cited]

Diamond, Adele. “Executive Functions.” *Annual Review of Psychology* 64 (2013). <https://doi.org/10.1146/annurev-psych-113011-143750>. Working memory, inhibition, shifting; protracted development. Cited via secondary developmental reviews. [Cited]

Dietz, G., et al. “Theory of AI Mind: How Adults and Children Reason about the ‘Mental States’ of Conversational AI.” CogSci / Stanford SLL, 2023. https://sll.stanford.edu/docs/2023_cogsci/2023_Dietz_et_al_C. Accessed 30 August 2026. Ages 3 to 8 treat AI like human agents (false belief). Full author line not locked from the research JSON. [Cited]

Dogan, S., U. Yazicilar Nalbantoglu, I. Celik, and N. Dogan. “Artificial Intelligence Professional Development: A Systematic Review of TPACK, Designs, and Effects for Teacher Learning.” *Professional Development in Education* 51, no. 3 (2025): 519-46. <https://doi.org/10.1080/19415257.2025.2454457>. Turkish characters simplified here; see GAPS. [Cited]

Eames, Taryn, Emma Brunskill, Bogdan Yamkovenko, Kodi Weatherholtz, and Philip Oreopoulos. “Computer-Assisted Learning in the Real World: How Khan Academy Influences Student Math Learning.” *Proceedings of the National Academy of Sciences* 123, no. 10 (2026). Cited by Oreopoulos and Low: per-hour returns to Khan Academy without AI about 0.0047-0.0086 SD. Article number / pages not in the research file. [Cited]

Ennis, Robert H. “The Nature of Critical Thinking: Outlines of General Critical Thinking Dispositions and Abilities.” 2011/2015. <https://criticalthinking.net/wp-content/uploads/2024/04/The-Nature-of-Critical-Thinking.pdf>. Accessed 30 August 2026. “Reasonable reflective thinking focused on deciding what to believe or do.” [Cited]

Escueta, Maya, Andre J. Nickow, Philip Oreopoulos, and Vincent Quan. “Upgrading Education with Technology: Insights from Experimental Research.” *Journal of Economic Literature* 58, no. 4 (2020): 897-996. Earlier computer-assisted learning typically 0.05-0.20 SD. [Cited]

“Evaluating Answer Leakage Robustness of LLM Tutors against Adversarial Student Attacks.” ACL 2026 / arXiv:2604.18660. 2026. <https://arxiv.org/abs/2604.18660>. Accessed 30 August 2026. Manuscript cites the Zeng et al. line. Pedagogical tutors leak full answers under multi-turn student attacks. Fine-tuned adversarial student agent: example 7B tutor leakage 14 percent to 82 percent. Full author line not locked. [Cited; preprint or conference — see GAPS]

Facione, Peter A. *Critical Thinking: A Statement of Expert Consensus for Purposes of Educational Assessment and Instruction* (The Delphi Report). American Philosophical Association, 1990. <https://insightassessment.com/iaresource/the-delphi-report-a-statement-of-expert-consensus-on-the-definition-of-critical-thinking/>. Accessed 30 August 2026. Six core skills: interpretation, analysis, evaluation, inference, explanation, self-regulation. [Cited]

Facione, Peter A. *Critical Thinking: What It Is and Why It Counts*. Insight Assessment, 2015. <https://rtdw.ndm.edu/sites/default/files/pdf/Facione-2015.pdf>. Accessed 30 August 2026. [Cited]

Fan, Yizhou, L. Tang, H. Le, K. Shen, S. Tan, Y. Zhao, Y. Shen, X. Li, and D. Gašević. “Beware of Metacognitive Laziness: Effects of Generative Artificial Intelligence on Learning Motivation, Processes, and Performance.” *British Journal of Educational Technology* 56, no. 2 (2025): 489-530. <https://doi.org/10.1111/bjet.13544>. ERIC: <https://eric.ed.gov/?id=EJ1460793>. Accessed 30 August 2026. RCT N = 117 writing: ChatGPT vs human expert vs analytics vs none. Better essay scores; ns knowledge gain/transfer; fewer metacognitive processes. Publisher page Cloudflare-blocked; ERIC abstract used. Do not invent additional coefficients. [Cited; full BJET PDF UNVERIFIED]

“Generative AI, Cognitive Offloading, and Learner Agency in Higher Education: A Scoping Review.” *Behavioral Sciences* 16, no. 7 (2026): 1150. <https://doi.org/10.3390/bs16071150>. Accessed 30 August 2026. 123 studies. Authors unlisted on the DOI landing in the research file. Full text not fetched. [Consulted; full text UNVERIFIED]

Gerlich, Michael. "AI Tools in Society: Impacts on Cognitive Offloading and the Future of Critical Thinking." *Societies* 15, no. 1 (2025): 6. Survey/correlational; not an RCT. Younger users more dependent. [Cited]

"Governing Generative AI in Higher Education: Emerging Policy Approaches and Support Ecosystems at Innovative U.S. Universities." *International Journal for Educational Integrity* (2026). <https://doi.org/10.1007/s40979-026-00233-x>. AAC&U/Elon late-2024: 69 percent of leaders report written GenAI teaching policies. Yale/Cornell/MIT Sloan/USC caution against detectors as dispositive. Authors not listed in the research JSON. [Cited]

Hanham, Jose, Wayne Leahy, and John Sweller. "Cognitive Load Theory, Element Interactivity, and the Testing and Reverse Testing Effects." *Applied Cognitive Psychology* 31, no. 3 (2017): 265-80. [Cited]

Harvard Undergraduate Survey on Generative AI. arXiv:2406.00833. 2024. <https://arxiv.org/abs/2406.00833>. <https://doi.org/10.48550/arXiv.2406.00833>. Accessed 30 August 2026. **Preprint survey**. Among AI users, 40 percent of students with no financial aid paid for AI products versus 20 percent with partial/full aid. Authors flag inequity if premium quality is real. Author line not in the research JSON. [Cited; preprint]

Henkel, Owen, Hannah Horne-Robinson, Neeve Kozhakhmetova, and Adam Lee. "Effective and Scalable Math Support: Experimental Evidence on the Impact of an AI Math Tutor in Ghana." In *AIED 2024*, CCIS vol. 2150. Cham: Springer, 2024. <https://arxiv.org/pdf/2402.09809>. Accessed 30 August 2026. About 500 students; 1 hr/week WhatsApp Rori; Cohen's $d = 0.36$; about \$5/student. Year-1 preliminary; authors flag caution. [Cited]

Hiebert, James, and Douglas A. Grouws. "The Effects of Classroom Mathematics Teaching on Students' Learning." In *Second Handbook of Research on Mathematics Teaching and Learning*, edited by Frank K. Lester Jr., 371-404. Charlotte, NC: Information Age, 2007. Productive struggle at p. 387. [Cited]

Institute of Education Sciences, What Works Clearinghouse. *Cognitive Tutor* intervention report. 2016. <https://ies.ed.gov/ncee/wwc/Intervention/818>. Accessed 30 August 2026. Older ITS reports exist. **No WWC intervention report specific to generative AI tutors** as of 30 August 2026. [Cited]

Kalyuga, Slava. "Expertise Reversal Effect and Its Instructional Implications." *Instructional Science* 35 (2007): 95-97. [Cited]

Kalyuga, Slava, Paul Ayres, Paul Chandler, and John Sweller. "The Expertise Reversal Effect." *Educational Psychologist* 38, no. 1 (2003): 23-31. [Cited]

Kapur, Manu. "Productive Failure in Learning Math." *Cognitive Science* 38, no. 5 (2014): 1008-22. [Cited]

Karpicke, Jeffrey D., and Henry L. Roediger III. "The Critical Importance of Retrieval for Learning." *Science* 319, no. 5865 (2008): 966-68. [Cited]

Kestin, Greg, Kelly Miller, Anna Kiales, Timothy Milbourne, and Gregorio Ponti. "AI Tutoring Outperforms In-Class Active Learning: An RCT Introducing a Novel Research-

Based Design in an Authentic Educational Setting.” *Scientific Reports* 15 (2025): 17458. <https://doi.org/10.1038/s41598-025-97652-6>. Accessed 30 August 2026. Crossover RCT, Harvard Physical Sciences 2, Fall 2023, N = 194 eligible of 233 enrolled. Custom GPT-4 tutor (“PS2 Pal”) versus the course’s existing in-class active learning. Linear 0.63 SD; quantile 0.73-1.3 SD; median time 49 vs 60 minutes. Delayed retention unmeasured. Not off-the-shelf ChatGPT. Not an agent. Full article fetched. [Cited]

Khlaif, Zuheir N., et al. “Redesigning Assessments for AI-Enhanced Learning: A Framework for Educators in the Generative AI Era.” *Education Sciences* 15, no. 2 (2025): 174. <https://doi.org/10.3390/educsci15020174>. Accessed 30 August 2026. Qualitative, 61 faculty; Against/Avoid/Adopt/Explore. Remaining author names not locked. [Cited]

Kosmyna, Nataliya, Eugene Hauptmann, Yeo T. Yuan, Jessica Situ, Xian-Hao Liao, Ashly Vivian Beresnitsky, Iris Braunstein, and Pattie Maes. “Your Brain on ChatGPT: Accumulation of Cognitive Debt When Using an AI Assistant for Essay Writing Task.” arXiv:2506.08872. 10 June 2025. <https://arxiv.org/abs/2506.08872>. MIT Media Lab: <https://www.media.mit.edu/publications/your-brain-on-chatgpt/>. Accessed 30 August 2026. **Preprint, not a journal article** as of 30 August 2026. N = 54 Boston-area adults (sessions 1-3); session 4 n = 18. EEG essay task; not K-12 and not agentic systems. Do not cite as population cognitive decline. Full PDF fetched. [Cited; preprint]

Kulik, James A., and J. D. Fletcher. “Effectiveness of Intelligent Tutoring Systems: A Meta-Analytic Review.” *Review of Educational Research* 86, no. 1 (2016): 42-78. Pre-generative ITS. [Cited]

LearnLM Team, Google DeepMind, and Eedi. “AI Tutoring Can Safely and Effectively Support Students: An Exploratory RCT in UK Classrooms.” arXiv:2512.23633. 2025. <https://arxiv.org/html/2512.23633>. Accessed 30 August 2026. **Preprint**. N = 165 UK secondary; human-supervised LearnLM vs human tutor vs static hints. Session-level randomization; cumulative learning not identified; every message gated. Supervising-tutors-learned-Socratic-moves finding: interview evidence from n = 5, not a PD trial. Full HTML fetched. [Cited; preprint]

Li, K., P. Wang, and G. Chen. “How Can AI Be Integrated into Teacher Professional Development Programs? A Systematic Review Based on an Adapted Technology-Based Learning Model.” *Teaching and Teacher Education* 168 (2025): 105219. <https://doi.org/10.1016/j.tate.2025.105219>. 16 studies through 2024. Teacher knowledge/attitudes/self-efficacy; rarely student learning. [Cited]

Liang, Weixin, Mert Yuksekgonul, Yining Mao, Eric Wu, and James Zou. “GPT Detectors Are Biased against Non-Native English Writers.” *Patterns* 4, no. 7 (2023): 100779. <https://doi.org/10.1016/j.patter.2023.100779>. arXiv:2304.02819. Stanford HAI: <https://hai.stanford.edu/news/ai-detectors-biased-against-non-native-english-writers>. Accessed 30 August 2026. **61.22 percent** is the authors’ average false-positive rate on 91 human TOEFL essays (arXiv HTML); the *Patterns* HTML rounds the same average to 61.3 percent. This book keeps 61.22. Near-perfect on U.S. eighth-grade essays. **18/91** flagged by all seven detectors; **89/91 (97.8 percent)** flagged by at least one — the mid2 pass figure used on the page (introduction “ninety-seven percent” tightened to 97.8 in ch. 8). Perplexity metric penalizes constrained diction. Simple prompting bypasses detectors. [Cited]

Mollick, Ethan, and Lilach Mollick. “Assigning AI: Seven Approaches for Students, with Prompts.” arXiv:2306.10052. 2023. Tutor/coach/mentor/teammate/student/simulator/tool. Used as prompt source in De Simone et al. 2025. [Cited; preprint]

National Council of Teachers of Mathematics. *Principles to Actions: Ensuring Mathematical Success for All*. Reston, VA: NCTM, 2014. Warns against rescuing students too soon (p. 48 cited in secondary sources). [Cited]

Nickow, Andre, Philip Oreopoulos, and Vincent Quan. “The Promise of Tutoring for PreK-12 Learning: A Systematic Review and Meta-Analysis of the Experimental Evidence.” *American Educational Research Journal* 61, no. 1 (2024): 74-107. <https://doi.org/10.3102/00028312231208687>. Pooled experimental preK-12 tutoring 0.288 SD (SE 0.029). [Cited]

“Not All Cognitive Offloading Is Equal: Distinguishing Dependent and Autonomous Offloading to Generative AI.” *Frontiers in Psychology* (2026). <https://doi.org/10.3389/fpsyg.2026.1878629>. Accessed 30 August 2026. Abstract-level: frequency of use is not the right variable; whether the user retains the evaluative loop is. Authors not listed in the research file. [Cited; abstract only]

Oreopoulos, Philip, Chloe R. Gibbs, Michael Jensen, and Joseph Price. “Teaching Teachers to Use Computer Assisted Learning Effectively: Experimental and Quasi-Experimental Evidence.” NBER Working Paper 32388. 2024. Forthcoming, *Journal of Human Resources*. Causal PD for CAL structure raises usage and achievement. Best analogue for generative-AI teacher PD; not a generative-AI trial. [Cited; working paper]

Oreopoulos, Philip, and Nina Low. “One Click Away: AI Tutoring with Khanmigo in a Two-Year School Experiment.” EdWorkingPaper 26-1551 / NBER Working Paper 35620. Annenberg Institute at Brown University, 2026. <https://doi.org/10.26300/kner-hv33>. PDF: <https://edworkingpapers.com/sites/default/files/ai26-1551.pdf>. NBER: <https://www.nber.org/papers/w35620>. Registry AEARCTR-0013519. Accessed 30 August 2026. **Working paper**. Cluster RCT, 18 Hamilton County, Tennessee middle schools; 6,902 student-terms; 53 clusters. ITT +1.26 NPR/term (SE 0.60); about 0.040 population SD/term; stacked about 0.06 SD/year, about 0.08 in year 2; implied full-year participation year 2 about 0.14 SD. Median messaged on one-third of practice days and in 17 percent of mistake sessions; 14.5 percent of messages mathematical. TCAP registered as co-primary, **not yet incorporated** as of the August 2026 paper. Cost about \$15 per student per year. Full PDF fetched. Nina Low, not Nathaniel/Nathan. [Cited; working paper]

Paas, Fred, and Jeroen J. G. van Merriënboer. “Cognitive-Load Theory: Methods to Manage Working Memory Load in the Learning of Complex Tasks.” *Current Directions in Psychological Science* 29, no. 4 (2020): 394-98. [Cited]

Paas, Fred, and Jeroen J. G. van Merriënboer. “Variability of Worked Examples and Transfer of Geometrical Problem-Solving Skills.” *Journal of Educational Psychology* 86, no. 1 (1994): 122-33. [Cited]

Perkins, Mike, Leon Furze, Jasper Roe, and Jason MacVaugh. “The Artificial Intelligence Assessment Scale (AIAS): A Framework for Ethical Integration of Generative AI in Educational Assessment.” *Journal of University Teaching and Learning Practice* 21, no. 6 (2024). <https://doi.org/10.53761/q3azde36>. Shared vocabulary, not an RCT of learning. [Cited]

Perkins, Mike, Leon Furze, Jasper Roe, and Jason MacVaugh. “Reimagining the Artificial Intelligence Assessment Scale: A Refined Framework for Educational Assessment.” *Journal of University Teaching and Learning Practice* (2025/26). <https://doi.org/10.53761/rrm4y757>. Adds AI-exploration level; recasts as assessment-redesign framework. Full-text fetch timed out. [Cited]

Pew Research Center. “About a Quarter of U.S. Teens Have Used ChatGPT for Schoolwork – Double the Share in 2023.” 15 January 2025. <https://www.pewresearch.org/short-reads/2025/01/15/about-a-quarter-of-us-teens-have-used-chatgpt-for-schoolwork-double-the-share-in-2023/>. Accessed 30 August 2026. 26 percent of U.S. teens ages 13-17 (13 percent in 2023). Chat use, not agent use. Awareness rose most in the lowest-income households. 18 percent say essay-writing use is acceptable; 42 percent say not. A later February 2026 Pew chatbot survey exists; this book’s authority is the January 2025 ChatGPT number. [Cited]

Prather, James, Brent N. Reeves, Juho Leinonen, Stephen MacNeil, et al. “‘It’s Weird That It Knows What I Want’: Usability and Interactions with Copilot for Novice Programmers.” *ACM Transactions on Computer-Human Interaction* (2023). <https://courses.cs.duke.edu/compsci290.2/spring24/reading>. Accessed 30 August 2026. Shepherding and straying/drifting; metacognitive difficulties; overtrust. Volume/issue/pages not in the research file. The 2024 laboratory follow-on (greenfield assignments; 20 of 21 completed the task) is described in the pedagogy memo; a separate full bibliographic record was not fetched. Do not treat either as a multi-site RCT of conceptual CS learning. [Cited]

Pyke, Aryn A., et al. “Does Difficulty Moderate Learning? A Comparative Analysis of the Desirable Difficulties Framework and Cognitive Load Theory.” *Quarterly Journal of Experimental Psychology* (2025). <https://doi.org/10.1177/17470218241308143>. <https://pmc.ncbi.nlm.nih.gov/articles/PMC12432286/>. Expertise and element interactivity moderate when difficulty is desirable. Remaining authors not locked. [Cited]

Risko, Evan F., and Sam J. Gilbert. “Cognitive Offloading.” *Trends in Cognitive Sciences* 20, no. 9 (2016): 676-88. Canonical definition. [Cited]

Robinson, Carly D., David Gormley, Ana Trindade Ribeiro, and Susanna Loeb. “Access Is Not Enough: Human Support Improves Engagement with AI Tutoring.” EdWorkingPaper 26-1451. Annenberg Institute, 2026. Cited in Oreopoulos and Low 2026. David Gormley and Ana Trindade Ribeiro (not Devin / Ana Teresa — that spelling belongs to Tutor CoPilot). [Cited; working paper]

Roediger, Henry L., III, and Jeffrey D. Karpicke. “Test-Enhanced Learning: Taking Memory Tests Improves Long-Term Retention.” *Psychological Science* 17, no. 3 (2006): 249-55. [Cited]

Soderstrom, Nicholas C., and Robert A. Bjork. “Learning versus Performance: An Integrative Review.” *Perspectives on Psychological Science* 10, no. 2 (2015): 176-99. Conditions that raise performance can lower learning. Analytic frame; not re-fetched in full on 30 August 2026. [Cited]

Sparrow, Betsy, Jenny Liu, and Daniel M. Wegner. “Google Effects on Memory: Cognitive Consequences of Having Information at Our Fingertips.” *Science* 333, no. 6043 (2011): 776-78. [Cited]

Stadler, Matthias, Maria Bannert, and Michael Sailer. “Cognitive Ease at a Cost: LLMs Reduce Mental Effort but Compromise Depth in Student Scientific Inquiry.” *Computers in Human Behavior* 160 (2024): 108386. <https://doi.org/10.1016/j.chb.2024.108386>. Open PDF:

<https://epub.ub.uni-muenchen.de/125262/1/1-s2.0-S0747563224002541-main.pdf>. RCT N = 91; ChatGPT-3.5 vs Google. Lower cognitive load, lower-quality reasoning/argumentation. [Cited]

“A Survey on Agentic AI-Powered Adaptive Learning Platforms.” IEEE AIEI 2026. <https://doi.org/10.1109/aiei69164.2026.11497428>. Accessed 30 August 2026. Search snippet notes an empirical gap on fully agentic platforms. Authors not confirmed on fetch. Full text not fetched. [Cited; full text UNVERIFIED]

Sweller, John. “Cognitive Load during Problem Solving: Effects on Learning.” *Cognitive Science* 12, no. 2 (1988): 257-85. [Cited]

Sweller, John. “Cognitive Load Theory.” In *Psychology of Learning and Motivation*, vol. 55, edited by Jose P. Mestre and Brian H. Ross, 37-76. San Diego: Academic Press, 2011. Cited by Kestin et al. 2025. [Cited]

Sweller, John, Paul Ayres, and Slava Kalyuga. *Cognitive Load Theory*. New York: Springer, 2011. [Cited]

Teknum, Ryan, Roger Jin, Jai Suphavadeeprasit, Dakota Mahan, Jeffrey Quesnelle, Joe Li, Chen Guang, Shannon Sands, and Karan Malhotra. “Hermes 4 Technical Report.” arXiv:2508.18255. 25 August 2025. <https://arxiv.org/abs/2508.18255>. Accessed 30 August 2026. **Preprint**. Open-weight hybrid-reasoning LLM family. Distinguishes the model family from Hermes Agent. HTML conversion failed; existence confirmed via Hugging Face model card. **Quantitative benchmarks UNVERIFIED until the PDF is read**. [Cited; preprint; benchmarks UNVERIFIED]

Tervo-Clemmens, B., et al. “A Canonical Trajectory of Executive Function Maturation from Adolescence to Adulthood.” *Nature Communications* 14 (2023). <https://pubmed.ncbi.nlm.nih.gov/37903830/>. Significant EF growth ages 10-15; adult-like 18-20. N = 10,766. Remaining authors not locked. [Cited]

Trouche, Luc. “Mathematics Education and AI: Rethinking Instrumental Orchestration as the Structured Regulation of AI’s Pedagogical Agency.” French Institute of Education, ENS de Lyon, 2025. Lecture. AISIO: teacher regulates activation, delegation, validation of AI. Conceptual, not an RCT. The 2004 instrumental-orchestration original was not re-fetched. [Cited]

Trust-calibration work with upper-division students under extended Copilot use. arXiv:2509.13253. Named in ch. 3 n. 37. Authors and title not locked. Not a classroom RCT of conceptual learning. [Cited; incomplete record — see GAPS]

van Gog, Tamara, and John Sweller. “Not New, but Nearly Forgotten: The Testing Effect Decreases or Even Disappears as the Complexity of Learning Materials Increases.” *Educational Psychology Review* 27 (2015): 247-64. [Cited]

Van Vlasselaer, Marijke, Filip Van Droogenbroeck, and Bram Spruyt. “Who Wrote This? Evaluating the Reliability of AI Detection Tools in Higher Education.” *International Journal for Educational Integrity* 22, art. 16 (2026). <https://doi.org/10.1007/s40979-026-00226-w>. Accessed 30 August 2026. 160 labelled documents. Turnitin 100 percent false negatives on fully-AI papers in that set (0-20 percent scores) against newer pipelines; humanised mixed text also poorly detected. Failure mode toward invisibility. Full PDF not re-read on the mid2 pass. [Cited]

Wang, Rose E., Ana T. Ribeiro, Carly D. Robinson, Susanna Loeb, and Dorottya Demszky. "Tutor CoPilot: A Human-AI Approach for Scaling Real-Time Expertise." arXiv:2410.03017 / EdWorkingPaper. 2024. <https://arxiv.org/html/2410.03017>. Accessed 30 August 2026. RCT about 900 tutors, about 1,800 K-12 students. +4 pp topic mastery ITT; +9 pp for lower-rated tutors; about \$20/tutor/year. Keep Ana T. Ribeiro (this paper); do not collapse with Robinson 26-1451's Ana Trindade Ribeiro. [Cited; preprint / working paper]

Wang, Shuai, et al. "The Efficacy of Artificial Intelligence-Enabled Adaptive Learning Systems from 2010 to 2022 on Learner Outcomes: A Meta-Analysis." *Journal of Educational Computing Research* (2024). <https://doi.org/10.1177/07356331241240459>. N = 45 studies, g = 0.70 vs non-adaptive. Window ends 2022. **Not a generative-AI result.** Remaining authors not locked. [Cited]

Wang and Zhang. "Pedagogical Partnerships with Generative AI in Higher Education: How Dual Cognitive Pathways Paradoxically Enable Transformative Learning." *International Journal of Educational Technology in Higher Education* (2026). <https://doi.org/10.1186/s41239-026-00585-x>. N = 912, three regions; U-shaped offloading curve. Full text not fetched; used at abstract/synthesis level. First names and outlet naming (IJETHE vs IJAIED) not locked. [Consulted; full text UN-VERIFIED]

Warshawer, Hiroko K. "Productive Struggle in Middle School Mathematics Classrooms." *Journal of Mathematics Teacher Education* 18 (2014): 375-400. [Cited]

Weber-Wulff, Debora, Alla Anohina-Naumeca, Sonja Bjelobaba, Tomáš Foltýnek, Jean Guerrero-Dib, Olumide Popoola, Petr Šigut, and Lorna Waddington. "Testing of Detection Tools for AI-Generated Text." *International Journal for Educational Integrity* 19 (2023). <https://doi.org/10.1007/s40979-023-00146-z>. Commercial detectors neither accurate nor reliable enough for academic-integrity decisions. [Cited]

Willingham, Daniel T. "Critical Thinking: Why Is It So Hard to Teach?" *American Educator* 31, no. 2 (Summer 2007): 8-19. https://www.aft.org/sites/default/files/media/documents/2023/31.2_summer_2007.pdf. Accessed 30 August 2026. Verified 30 August 2026 against the AFT-hosted PDF. Exact bicycle sentence: "People who have sought to teach critical thinking have assumed that it is a skill, like riding a bicycle, and that, like other skills, once you learn it, you can apply it in any situation." Processes of thinking "are intertwined with the content of thought (that is, domain knowledge)." Maxims can be taught, "but without background knowledge and practice, they probably will not be able to implement the advice they memorize." "Literary thinking" in the book's opening triad is a pedagogy-memo compression, not a third Willingham quotation. [Cited]

"Young Children's Anthropomorphism of an AI-Powered Conversational Agent: Brain Activation and the Role of Parent Co-Presence." *International Journal of Human-Computer Interaction* / arXiv 2512.02179. 2026. <https://doi.org/10.1080/10447318.2026.2678532>. Ages 5-6 perceptive and epistemic attribution; parent co-presence regulates; N = 23. Authors not locked in the research JSON. [Cited]

3. Policy, union, ministry, UNESCO / OECD / U.S. Department of Education

AI Leap Foundation / Estonian Ministry of Education and Research. TI-Hüpe / AI Leap. Announced February 2025 by President Alar Karis. <https://tihupe.ee/en/>. First-year review, 7 August 2026: <https://tihupe.ee/en/read-the-ai-leap-first-year-report/>. Accessed 30 August 2026. 154 of 156 upper secondary schools in 2025/26; teachers get paid ChatGPT or Gemini; about 20,000 grade 10-11 students get ITI, a Socratic Estonian-language app that does not give answers. Problem statement: 64-90 percent of students already using free commercial AI to finish homework. Public-private 50/50 funding. The 7 August 2026 review page **itself** states that first-year usage data “does not yet provide a basis for claiming that ITI has improved learning outcomes.” University of Tartu study (Jaan Aru; more than 9,000 students) expected later in 2026. Outcome coefficients not pulled. [Cited]

AI4K12 (AAAI / CSTA, NSF DRL-1846073). “Grade Band Progression Charts — Five Big Ideas in AI.” <https://ai4k12.org/gradeband-progression-charts/>. Accessed 30 August 2026. Working drafts, public. K-2 / 3-5 / 6-8 / 9-12 progression. [Cited]

American Academy of Pediatrics. “Digital Ecosystems, Children, and Adolescents: Policy Statement.” *Pediatrics* 157, no. 2 (2026). <https://doi.org/10.1542/peds.2025-075320>. Publisher HTML bot-blocked; content corroborated via AAP Center of Excellence explainer. Child-centered versus engagement-based design; 5 Cs. A dedicated AAP “AI in schools” policy does not yet exist. [Cited; publisher HTML UNVERIFIED]

American Academy of Pediatrics. “Digital Ecosystems, Children, and Adolescents: Technical Report.” *Pediatrics* 157, no. 2 (2026). <https://publications.aap.org/pediatrics/article/157/2/e2025075321/206128/Digital-Ecosystems-Children-and-Adolescents>. Partial fetch. [Cited; partial]

American Academy of Pediatrics, Center of Excellence on Social Media and Youth Mental Health (Jenny Radesky). “Understanding the New AAP Digital Media Guidelines for Screen Time and Social Media.” 20 January 2026. <https://www.aap.org/en/patient-care/media-and-children/center-of-excellence-on-social-media-and-youth-mental-health/understanding-the-new-aap-digital-media-guidelines/>. Accessed 30 August 2026. [Cited]

American Federation of Teachers. *Commonsense Guardrails for Using Advanced Technology in Schools*. Version 2. March 2025. https://www.aft.org/sites/default/files/media/documents/2025/Commonsense_Guardrails_for_Using_Advanced_Technology_in_Schools.pdf. HTML: <https://www.aft.org/reports/commonsense-guardrails-using-advanced-technology-schools>. Accessed 30 August 2026. Official title is “Using Advanced Technology in Schools,” **not** “for Artificial Intelligence in Education.” Value 1: student privacy and safety come first; the sacrifice clause is the **second** sentence (“The well-being and data security of students must be paramount, even if it means sacrificing some potential educational opportunities”). Value 3: educators choose the technology. Collective-bargaining annex: AI a mandatory subject; no surveillance; right to disconnect; equity funding. Value 9 environmental. Notes absence of federal regulation under the current administration. [Cited]

American Federation of Teachers. “On the Formation of a National Academy for AI Instruction.” Resolution adopted 22 July 2025. <https://www.aft.org/resolution/formation-national-academy-ai-instruction>

instruction. Press: <https://www.aft.org/press-release/aft-launch-national-academy-ai-instruction-microsoft-openai-anthropic-and-united>. Accessed 30 August 2026. \$23 million; Microsoft, OpenAI, Anthropic; UFT founding partner; NYC hub; free PD for 1.8 million members; **target** 400,000 educators / about 7.2 million students over five years — a plan, not a trained headcount. Launch press is 8 July 2025; resolution adopted 22 July 2025. Do not mash with CNBC's ChatGPT for Teachers first-cohort about 150,000. [Cited]

American Federation of Teachers. "Preparing for the Age of Artificial Intelligence: Protecting Democracy, Workers, Students and the Planet." Resolution adopted 18 July 2026. <https://www.aft.org/resolution/preparing-age-artificial-intelligence-protecting-democracy-workers-students-and-planet>. Accessed 30 August 2026. Opposes student-facing AI in elementary; companion-chatbot ban until age 16; PK-2 screen limits except disability/compelling need. [Cited]

Broward County Public Schools. "Broward Powered by AI." Board Resolution No. 26-132, adopted 10 March 2026. <https://browardschools.ai/>. Accessed 30 August 2026. K-5: no district-approved student-facing AI tools. Grades 6-12: Khanmigo and HMH Writable; MagicStudent off. Copilot is staff (20,000 Copilot Premium staff licenses). No final grades by AI; no model training. Detectors not relied upon to make decisions about a student. SPARK cadre announced 18 August 2026. [Cited]

CAST. *UDL Guidelines 3.0*. 30 July 2024. <https://udlguidelines.cast.org/>. Accessed 30 August 2026. Learner agency; identity/bias; assistive tech 4.2. [Cited]

Center for Democracy and Technology. *Ableism and Disability Discrimination in New Surveillance Technologies*. May 2022. <https://cdt.org/wp-content/uploads/2022/05/2022-05-23-CDT-Ableism-and-Disability-Discrimination-in-New-Surveillance-Technologies-report-final-redu.pdf>. Accessed 30 August 2026. Remote proctoring flags disability-related behavior. Cited by NEA 2024. [Cited]

Council Rock School District. "Artificial Intelligence Guidelines." 2026. <https://resources.finalsite.net/images/v1> Located via search; not fully extracted. [Consulted; full PDF not extracted]

EDUCAUSE. *2026 EDUCAUSE: The Impact of AI on Learning Assessment Report*. Library page, June 2026. <https://library.educause.edu/resources/2026/6/2026-educause-the-impact-of-ai-on-learning-assessment-report>. Accessed 30 August 2026. Survey of 438 faculty/staff. Claims used from the library landing page only. **Full PDF not fetched**. [Cited; full PDF UNVERIFIED]

EDUCAUSE. *2026 EDUCAUSE Horizon Report | Teaching and Learning Edition*. <https://library.educause.edu/resources/2026/6/2026-educause-horizon-report-teaching-and-learning-edition>. Accessed 30 August 2026. Listed takeaway: process-based assessment. **Full PDF not fetched**. [Cited; full PDF UNVERIFIED]

Education International. "AI and Cognitive Offloading: Supporting Teachers to Shape AI's Impact on Learning." (Jason Lodge and Leslie Loble.) <https://www.ei-ie.org/en/item/32652:ai-and-cognitive-offloading-supporting-teachers-to-shape-ais-impact-on-learning>. Accessed 30 August 2026. Practitioner synthesis, not a primary experiment. [Cited]

Harvard FAS Office of Undergraduate Education. "Generative AI Guidance." <https://oue.fas.harvard.edu/faculty-resources/generative-ai-guidance/>. Accessed 30 August 2026. Required Canvas policy; three

model policies (forbid / allow with citation / assignment-specific). [Cited]

International Task Force on Teachers for Education 2030. *Promoting and Protecting Teacher Agency in the Age of Artificial Intelligence*. September 2025. <https://teachertaskforce.org/sites/default/files/2025-09/1149%5F25%5FPromoting%20and%20Protecting%20Teacher%20Agency%5FFINAL%5F3Sep.pdf>. Accessed 30 August 2026. Teachers irreplaceable; AI as augmentative. Full text not fully extracted. [Cited; partial]

ISTE. *ISTE Standards for Students*. Version 4.02, 2024. <https://iste.org/standards/students>. Accessed 30 August 2026. [Consulted]

ISTE Generation AI. “Profile of an AI-Ready Graduate.” <https://iste.org/generationai/news-updates/profile-of-an-ai-ready-graduate>. Accessed 30 August 2026. News-page summary of Culatta 2025 lists Ideator in place of Problem Solver; Culatta’s own ASCD headings match the ISTE Live 2026 six. Do not cite Culatta 2025 as an Ideator list. [Cited]

Joint Council for Qualifications. *AI Use in Assessments: Your Role in Protecting the Integrity of Qualifications*. Effective 30 April 2025. <https://www.jcq.org.uk/knowledge-hub/ai-use-in-assessments-your-role-in-protecting-the-integrity-of-qualifications-2/>. Accessed 30 August 2026. Own-work rule; acknowledgement = tool name + URL + date; non-editable prompt/output copy; AI cannot be the sole marker; detectors optional as one check. [Cited]

Little Falls Community Schools (MN). “Responsible Use of Artificial Intelligence” (Policy 625). 2026. <https://resources.finalsite.net/images/v1781805190/lfallsk12mnus/ioxv7omzintnulswpk8/625Responsible>. Located via search; H-AI-H model and grade-band access from snippets, not full extraction. [Consulted; full PDF not extracted]

Los Angeles Unified School District. “Guidelines for the Authorized Use of Artificial Intelligence for District” (BUL-151113). April 2024 framework. https://media.edlio.net/3476a301/6d998082/01503609/1dd0f151113_0_Guidelines_for_the_Authorized_Use_of_Artificial_Intelligence_for_District.pdf. Accessed 30 August 2026. Pre-vetting privacy/security; not a student-agent program. [Cited]

NAEYC. *Developmentally Appropriate Practice (DAP) Position Statement*. 2020. <https://www.naeyc.org/resource-statements/dap/contents>. Accessed 30 August 2026. Birth through 8. [Cited]

NAEYC. “Technology and Media in Early Childhood Education.” *Young Children*, Winter 2023. <https://www.naeyc.org/resources/pubs/yc/winter2023/technology-and-media-early-childhood-education>. Accessed 30 August 2026. [Cited]

NAEYC and Fred Rogers Center. *Technology and Interactive Media as Tools in Early Childhood Programs Serving Children from Birth through Age 8*. 2012. https://www.naeyc.org/sites/default/files/globally-shared/downloads/PDFs/resources/position-statements/ps_technology.pdf. Accessed 30 August 2026. Still the posted statement; revision underway 2025-26. Do not cite an updated NAEYC AI position that is not posted. [Cited]

National Education Association. *Report of the NEA Task Force on Artificial Intelligence in Education*. RA 2024. https://www.nea.org/sites/default/files/2024-06/report_of_the_nea_task_force_on_artificial_intelligence_in_education.pdf. Five principles: <https://www.nea.org/resource-library/artificial-intelligence-education/v-five-principles-use-artificial-intelligence-education>. Hub: <https://www.nea.org/ai>. Accessed 30 August 2026. Official order, verified 30 August 2026: (1) students and educators at the center;

(2) evidence-based AI; (3) ethical development/use and strong data protection; (4) equitable access; (5) ongoing education with and about AI — literacy and agency. An earlier memo pass had swapped (4) and (5). No high-stakes AI-alone decisions. Detector bias; Texas A&M University-Commerce episode; Nevada funding algorithm; NY facial-recognition ban September 2023. COPPA NPRM status in the report is as of April 2024 — superseded by the 2025 final rule. [Cited]

New York City Public Schools. “Guidance on Artificial Intelligence.” March 2026. <https://www.schools.nyc.gov/about-us/vision-and-mission/artificial-intelligence/guidance-on-artificial-intelligence>. Accessed 30 August 2026. Traffic-light framework; ERMA; red list on grading, IEP/504, placement, discipline, surveillance, counseling. December 2024 AI add-on: disclose AI capabilities; prohibit student-data model training; transparency. Chancellor’s Regulation A-820. Grade-band student-use Playbook still forthcoming as of the fetched page (targeted June 2026). [Cited]

OECD. *OECD Digital Education Outlook 2026*. OECD Publishing, January 2026. https://www.oecd.org/en/publications/digital-education-outlook-2026_062a7394-en.html. PDF listed: https://www.oecd.org/content/dam/oecd/en/publications/digital-education-outlook-2026_940e0dd8/062a7394-en.pdf. Accessed 30 August 2026. Listing-page claim: unguided GenAI raises task performance without learning gains. TALIS 2024: 37 percent of teachers use GenAI for work — **listing-page figure; full PDF Cloudflare-blocked 30 August 2026**. Do not mash with EDUCAUSE 37 percent of *institutions* licensing chatbots. [Cited; full PDF UNVERIFIED]

OECD. “Policies Supporting Responsible and Systematic GenAI Adoption in Higher Education.” *OECD Education Spotlights* no. 23, 17 July 2026. <https://doi.org/10.1787/c4e5621f-en>. [Consulted]

Prince George’s County Public Schools. *Artificial Intelligence Declaration Addendum*. <https://epi.pgcps.org/contentassets/0d463a5c451f40fcba26c2796d7a59c8/appendix-p-ai-declaration-addendum.pdf>. Accessed 30 August 2026. Specimen AI addendum: model disclosure, local vs cloud, FERPA/COPPA certification, no PII retained in generative models, no secondary use without district and parental approval, bias/harm audits on request, non-US subprocessors, feature off-switch. [Cited]

Quality Assurance Agency for Higher Education. *Reconsidering Assessment for the Chat GPT Era: QAA Advice on Developing Sustainable Assessment Strategies*. July 2023. <https://www.qaa.ac.uk/docs/qaa/members/reconsidering-assessment-for-the-chat-gpt-era.pdf>. Accessed 30 August 2026. [Cited]

Singapore Ministry of Education. “AI Usage in Schools” (parliamentary reply). 6 May 2026. <https://www.moe.gov.sg/news/parliamentary-replies/20260506-ai-usage-in-schools>. Accessed 30 August 2026. SLS Learning Assistant redirects students who “want to be spoon-fed with direct answers.” Designed tutor refusal, not an agent evaluation. [Cited]

Singapore Ministry of Education. “Artificial Intelligence in Education.” Page updated 31 July 2026. <https://www.moe.gov.sg/education-in-sg/educational-technology-journey/edtech-masterplan/artificial-intelligence-in-education>. Accessed 30 August 2026. P1-3 print-first, no student chatbots; P4-6 only SLS education-designed tools under teacher supervision;

secondary/pre-U integrity plus evaluation of AI output. [Cited]

Stanford University. “CS221: Artificial Intelligence — Generative AI Policy.” Spring 2026. <https://stanford-cs221.github.io/spring2026/>. Accessed 30 August 2026. Students may use GenAI as a human collaborator; may not ask for answers or to “check my work”; must share ChatGPT link or screenshots; “using an agentic coding tool such as Codex or Claude Code to build a repository from a project proposal is not allowed.” [Cited]

Student Data Privacy Consortium. *National Data Privacy Agreement* (NDPA) version 2. 2024. De-facto interstate DPA template with state exhibits. Not a statute. [Cited]

TeachAI (Code.org, ETS, ISTE, Khan Academy, World Economic Forum; advised by AASA, AFT, NEA, NSBA, CCSSO and others). *Foundational Policy Ideas for AI in Education / AI Guidance for Schools Toolkit*. 2024. <https://www.teachai.org/policy>; <https://www.teachai.org/toolkit-guidance>. Accessed 30 August 2026. Most copied U.S. K-12 template. [Cited]

UCLA Digital and Technology Solutions. “AI Use and Recommendation Guide.” <https://www.dts.ucla.edu/initiative/ai-use-recommendation-guide>. Accessed 30 August 2026. BruinLearn AI hub; OpenAI agreement first CA university at scale (vendor/admin claim); no FERPA/HIPAA in public tools. [Cited]

UCLA Teaching and Learning Center. “Using Generative AI Reflectively and Responsibly in Teaching and Learning.” <https://teaching.ucla.edu/resources/teaching-guides/using-generative-ai-reflectively-and-responsibly-in-teaching-and-learning/>. Accessed 30 August 2026. UCUES 2024: 67 percent of UCLA undergraduate *respondents* used AI tools at least several times a year; top uses brainstorming, researching a topic, and studying for exams. SAIRO brief: 4,706 respondents, 15 percent response rate. <https://sairo.ucla.edu/file/69a9752b-ba55-46a1-8cf1-e2a963ffbb97>. [Cited]

UK Department for Education. *Generative AI: Product Safety Standards*. First published 22 January 2025; updated 19 January 2026. <https://www.gov.uk/government/publications/generative-ai-product-safety-standards/generative-ai-product-safety-standards>. Accessed 30 August 2026. Most operational English-language child-safety spec: filtering that holds for whole conversations; jail-break resistance; no anthropomorphism; no sycophancy/dark patterns; hard time limits; mental-health pathways to humans; cognitive-offloading detection; no commercial training on inputs without permission. Live-web/search GenAI under the Online Safety Act 2023. [Cited]

UK Department for Education. *Generative Artificial Intelligence (AI) in Education*. First published 29 March 2023; updated 12 August 2025. <https://www.gov.uk/government/publications/generative-artificial-intelligence-in-education/generative-artificial-intelligence-ai-in-education>. Accessed 30 August 2026. Applies to England. Teacher-facing lower risk; pupil-facing only with safeguards, KCSIE, filtering/monitoring, age limits. Do not train models on student original work without permission. [Cited]

UNESCO. “AI and Education: Protecting the Rights of Learners.” 2024. <https://www.unesco.org/en/articles/ai-and-education-protecting-rights-learners>. Accessed 30 August 2026. About 2.6 billion people without internet (2024 figure UNESCO still uses); about 1 in 4 primary schools without electricity; about 60 percent of primary schools unconnected. [Cited]

UNESCO. “Artificial Intelligence in Education.” <https://www.unesco.org/en/digital-education/artificial-intelligence>. Accessed 30 August 2026. Human-centred framing; “preserving human agency in a

world of automation.” [Cited]

UNESCO. *Beijing Consensus on Artificial Intelligence and Education*. 2019. <https://unesdoc.unesco.org/ark:/4822> [Consulted]

UNESCO. *Recommendation on the Ethics of Artificial Intelligence*. 2021. <https://www.unesco.org/en/artificial-intelligence/recommendation-ethics>. First global AI ethics instrument. [Consulted]

UNESCO. “UNESCO: Governments Must Quickly Regulate Generative AI in Schools.” 7 September 2023. <https://www.unesco.org/en/articles/unesco-governments-must-quickly-regulate-generative-ai-schools>. Accessed 30 August 2026. Under 10 percent of 450-plus surveyed schools/universities had GenAI policies. [Cited]

UNESCO. “What You Need to Know about UNESCO’s New AI Competency Frameworks for Students and Teachers.” 2024. <https://www.unesco.org/en/articles/what-you-need-know-about-unescos-new-ai-competency-frameworks-students-and-teachers>. Accessed 30 August 2026. [Cited]

UNESCO (Fengchun Miao and Kelly Shiohira). *AI Competency Framework for Students*. 8 August 2024. <https://www.unesco.org/en/articles/ai-competency-framework-students>. 12 competencies / 4 dimensions / Understand-Apply-Create. [Cited]

UNESCO (Fengchun Miao and Mutlu Cukurova). *AI Competency Framework for Teachers*. 8 August 2024. <https://www.unesco.org/en/articles/ai-competency-framework-teachers>. 15 competencies / 5 dimensions / Acquire-Deepen-Create. [Cited]

UNESCO (Fengchun Miao and Wayne Holmes). *Guidance for Generative AI in Education and Research*. 7 September 2023. <https://www.unesco.org/en/articles/guidance-generative-ai-education-and-research>. Accessed 30 August 2026. Age 13 for independent GenAI conversation; under-13 / primary use of untested tools “not appropriate”; seven governmental steps; Global North training-data bias. Last UNESCO page update noted 16 January 2026. [Cited]

University of Illinois Chicago, Center for the Advancement of Teaching Excellence. “Adapting Assignments for the Age of AI.” <https://teaching.uic.edu/cate-teaching-guides/digital-learning/adapting-assignments-for-the-age-of-ai/>. Accessed 30 August 2026. [Cited]

U.S. Department of Education. *Navigating Artificial Intelligence in Postsecondary Education*. 2025. ERIC ED670768. <https://files.eric.ed.gov/fulltext/ED670768.pdf>. Accessed 30 August 2026. [Cited]

U.S. Department of Education. “U.S. Department of Education Issues Guidance on Artificial Intelligence Use in Schools, Proposes Additional Supplemental Priority.” Press release, 2025. <http://www.ed.gov/about/news/press-release/us-department-of-education-issues-guidance-artificial-intelligence-use-schools-proposes-additional-supplemental-priority>. Accessed 30 August 2026. [Consulted]

U.S. Department of Education / ERIC ED661924. *Empowering Education Leaders* [AI toolkit]. 2024. <https://files.eric.ed.gov/fulltext/ED661924.pdf>. Accessed 30 August 2026. Responsive to Biden EO 14110 365-day clock. District use policies; opt-out of AI-enabled applications. [Cited]

U.S. Department of Education, Office of Educational Technology. *Artificial Intelligence and the Future of Teaching and Learning: Insights and Recommendations*. May 2023. <https://www.ed.gov/sites/ed/files/documents/ai-report/ai-report.pdf>. Accessed 30 August 2026. Seven recommendations. Humans in the loop. The Department “firmly rejects the idea that AI could replace teachers.” FERPA/IDEA already apply. Public domain. [Cited]

U.S. Department of Education, Office of Educational Technology. *Designing for Education with Artificial Intelligence: An Essential Guide for Developers*. 2024. ERIC ED661949. <https://files.eric.ed.gov/fulltext/ED661949.pdf>. Accessed 30 August 2026. [Cited]

U.S. Department of Education, OPEPD. “Guidance on the Use of Federal Grant Funds to Improve Education Outcomes Using Artificial Intelligence (AI).” Dear Colleague letter, 22 July 2025. <https://www.ed.gov/media/document/opepd-ai-dear-colleague-letter-7222025-110427.pdf>. Accessed 30 August 2026. Allowable uses for tutoring, advising, materials; human-oversight condition. [Cited]

World Economic Forum. “AI Won’t Replace Teachers, Says This Global Union.” Interview with David Edwards, Education International. July 2024. <https://www.weforum.org/stories/2024/07/artificial-intelligence-education-teachers-union/>. Accessed 30 August 2026. Relational not transactional. Equity risk: human teachers as a luxury of the rich, bots for the poor. [Cited]

4. Product documentation (labelled vendor)

Anthropic. “Introducing Claude for Teachers.” 14 July 2026 (updated 21 July and 28 August 2026). <https://www.anthropic.com/news/claude-for-teachers>. Product: <https://claude.com/solutions/teachers>. Accessed 30 August 2026. **Vendor.** Educator-only; Claude 18+; FERPA-aligned DPA; training off for verified teacher accounts. Prospect Schools (Brooklyn) feedback; Detroit Public Schools Community District announced impact pilot, results unpublished. Sign-up by 30 June 2027 for a year of free access. Cowork scheduling; Learning Commons. Connectors named as ecosystem, not Anthropic-built. [Cited; vendor]

Character.AI. “An Update on Changes to Our Under-18 Experience.” October/November 2025. <https://blog.character.ai/an-update-on-changes-to-our-under-18-experience/>. Accessed 30 August 2026. **Vendor.** Open-ended character chat removed for under-18s from 24-25 November 2025. [Cited; vendor]

Cursor. “Cloud Agents.” <https://cursor.com/docs/cloud-agent>. Accessed 30 August 2026. **Vendor.** Isolated Cursor-managed Linux VMs; PR artifacts; kickoff surfaces. Formerly named Background Agents. [Cited; vendor]

Diffit. “Diffit Plans.” <https://web.diffit.me/plans>. App: <https://app.diffit.me/>. Accessed 30 August 2026. **Vendor.** FERPA/COPPA; vendor states it does not collect student data. [Cited; vendor]

Google. “Google and Khan Academy Build Classroom AI Tools.” <https://blog.google/products-and-platforms/products/education/khan-academy-back-to-school/>. Accessed 30 August 2026. **Vendor.** Gemini-powered Khanmigo features for the 2026 school year. [Cited; vendor]

Google Workspace Help. “Quickstart Guide to Gemini and Gemini Notebook for Education.” <https://support.google.com/a/answer/16350447>. Accessed 30 August 2026. **Vendor/admin**. Free in all Education editions; Core-service privacy claims. LearnLM marketing language labelled as such. Not a documented autonomous-agent SKU for schools. [Cited; vendor]

Khan Academy. “Becoming a Khan Academy Districts Partner.” <https://blog.khanacademy.org/becoming-a-khan-academy-districts-partner/>. Accessed 30 August 2026. **Vendor**. \$15/student Khanmigo; Clever or ClassLink; minimum 250 licenses; grades 3-12. “Over 400 partners across the US” is a vendor claim. Named partner: School City of Hobart, Indiana. [Cited; vendor]

Khan Academy. “Khan Academy Districts 2026-27 Renewal Guide.” <https://blog.khanacademy.org/khan-academy-districts-2026-27-renewal-guide/>. Accessed 30 August 2026. **Vendor**. Platform reimagining timeline. Not an efficacy study. [Consulted; vendor]

Magic School, Inc. “Midway ISD Case Study.” 2025. <https://www.magicschool.ai/case-studies/midway-isd>. Accessed 30 August 2026. **Vendor case study**. Midway ISD, Texas (about 8,700 students); 630-plus teachers trained summer 2024; 23-page DPA, no redlines. Not independent evaluation. [Cited; vendor]

Magic School, Inc. “Trust, Data Protection and Privacy for Schools.” <https://www.magicschool.ai/privacy/trust>. Accessed 30 August 2026. **Vendor**. iKeepSafe FERPA/COPPA; SOC 2 Type II; ESSA Level III (Instructure) — vendor-funded evidence tier, not a large independent RCT. [Cited; vendor]

Microsoft Customer Stories. “UC San Diego Prepares Students for AI-Driven Industry with GitHub Copilot.” <https://www.microsoft.com/en/customers/story/26539-university-of-california-san-diego-github-copilot>. Accessed 30 August 2026. **Vendor customer story**. Porter and Shah: manual competency then Copilot. Not a peer-reviewed trial. [Cited; vendor]

Nous Research. “Hermes 4 — Qwen-3 14B” (model card). 2025. <https://huggingface.co/NousResearch/Hermes-4-14B-FP8>. Accessed 30 August 2026. **Vendor**. Distinguishes the LLM family from the agent runtime. Cites arXiv 2508.18255. [Cited; vendor]

Nous Research. Hermes Agent. <https://hermes-agent.org/>. Docs: <https://hermes-agent.nousresearch.com/docs/>. GitHub: <https://github.com/nousresearch/hermes-agent/>. Nous Portal: <https://hermes-agent.nousresearch.com/docs/integrations/nous-portal>. Research-paper-writing skill: <https://hermes-agent.nousresearch.com/docs/user-guide/skills/bundled/research/research-research-paper-writing>. Accessed 30 August 2026. **Vendor**. Open-source MIT-licensed autonomous runtime. Vendor site dates the launch to **February 2026 without a day**; 25 February is the first-release date recorded in this book’s responsible-use stream (GitHub/releases), not a date printed on the vendor homepage. “Not a chatbot. Not a copilot. An agent that lives on your machine.” No telemetry claimed; data in `~/.hermes/`. **Not a child-safety certified school product**. No institutional AUP/DPA/SOP found. Hermes 4 (70B / 405B) is “not recommended for use inside Hermes Agent.” Recommended agent models on the Portal page include Claude Sonnet 4.6, GPT-5.5 Pro, google/gemini-3-pro-preview, DeepSeek V4 Pro. The **agent** installs locally; **open-weight Hermes models** can be served on llama.cpp, MLX, vLLM, Ollama, or LM Studio — do not conflate. [Cited; vendor]

OpenAI. “A Free Version of ChatGPT Built for Teachers.” 19 November 2025. <https://openai.com/index/chatgpt-for-teachers/>. Accessed 30 August 2026 (re-fetched on the mid2/front fact-check pass). **Vendor**.

Free through June 2028 on this page; first cohort “nearly 150,000 teachers and staff.” A separate expansion URL timed out on the landscape pass. [Cited; vendor]

OpenAI. “Bringing ChatGPT for Teachers to More U.S. School Districts.” <https://openai.com/index/bringing-chatgpt-for-teachers-to-more-us-school-districts/>. FETCH TIMED OUT 30 August 2026. Additional 2026 district counts UNVERIFIED. [Cited as timed-out; extra counts UNVERIFIED]

OpenAI. “Introducing ChatGPT Edu” / business data privacy / enterprise privacy. <https://openai.com/index/introducing-chatgpt-edu/>; <https://openai.com/business-data/>; <https://openai.com/enterprise-privacy/>. Accessed 30 August 2026. **Vendor claim:** ChatGPT Edu/Enterprise/Business/ Teachers/API data not used to train models by default. Not a FERPA determination. [Cited; vendor]

OpenAI. “Introducing ChatGPT for Teens: Built for Learning, Backed by Protections.” 2026. <https://openai.com/index/chatgpt-for-teens/>. Accessed 30 August 2026. **Vendor.** Under-18 model spec; no romantic dependence language. [Cited; vendor]

OpenAI. “Our Approach to Age Prediction.” 2026. <https://openai.com/index/our-approach-to-age-prediction/>. Accessed 30 August 2026. **Vendor.** Behavioral age signals; default to teen protections. [Cited; vendor]

OpenAI. “Terms of Use.” Effective 1 January 2026. <https://openai.com/policies/row-terms-of-use/>. Accessed 30 August 2026. **Vendor.** Minimum age 13; parent permission under 18. [Cited; vendor]

OpenAI Academy. “Deploying Codex in Higher Education.” 9 April 2026. <https://academy.openai.com/public/clubs/higher-education-05x4z/blogs/deploying-codex-in-higher-education-2026-04-09>. Accessed 30 August 2026. **Vendor playbook.** App vs Cloud vs CLI; security/repo scoping. Not an outcomes study. [Cited; vendor]

OpenAI Academy. “How to Build a Workspace Agent for Higher Education.” 6 May 2026. <https://academy.openai.com/en/public/clubs/higher-education-05x4z/blogs/how-to-build-a-workspace-agent-for-higher-education-2026-05-06>. [Consulted; vendor]

OpenAI Academy (Kirk Gulezian). “Understanding Workspace Agents in Higher Education.” 23 April 2026. <https://academy.openai.com/public/clubs/higher-education-05x4z/blogs/understanding-workspace-agents-higher-education>. Accessed 30 August 2026. **Vendor.** Research preview; off by default; admin-gated; humans review before anything is sent. Hypothetical campus workflows, not named-campus case studies. [Cited; vendor]

SchoolAI. “Give Every Student a Personalized Learning Experience | SchoolAI Spaces.” <https://schoolai.com/products/spaces>. Accessed 30 August 2026. **Vendor.** Teacher-designed Spaces. [Consulted; vendor]

SchoolAI. “SchoolAI Makes Students Think.” <https://schoolai.com/research/schoolai-makes-students-think>. Accessed 30 August 2026. **Vendor research.** Jordan School District (UT), 82 teachers, 13,882 conversations, 28 percent critical-thinking claim, ESSA Level III. Vendor-commissioned; not an RCT; Bloom rubric on chat logs. [Cited; vendor]

xAI. “AI Teammates That Finish the Work | Grok Bot.” <https://x.ai/bot>. Use cases: <https://x.ai/bot/use-cases>. Accessed 30 August 2026. **Vendor.** Distinguishes Bots from AI assistants (own computer, 24/7, parallel). No education SKU. [Cited; vendor]

xAI. “Grok Bot Is Now Included with More Plans.” 26 August 2026. <https://x.ai/news/grok-bot-more-plans>. Accessed 30 August 2026. **Vendor.** Expansion to SuperGrok plans and Cursor Pro/Teams. Website builder/deploy and meeting stand-in are on this expansion post, not the 11 August launch. None educational. [Cited; vendor]

xAI. Grok Bot FAQ. <https://docs.x.ai/grok-bot/faq>. Landscape fetch timed out; **re-fetched** on the front-matter fact-check pass, 30 August 2026. Requires cloud data storage; Legacy Privacy Mode not supported; training opt-out follows Cursor account settings. Still no schools, students, or teachers. Still no education SKU, DPA, or classroom SOP. End-errata did not re-clear this FAQ; front-errata is the later same-day fetch. [Cited; vendor]

xAI. “Introducing Grok Bot.” 11 August 2026. <https://x.ai/news/introducing-grok-bot>. Accessed 30 August 2026. **Vendor.** Early beta. “Bots have their own computer... keep working 24/7.” Sales outbound, CRM, inbox, expenses, bug reproduction on the launch post. No education SKU; product pages do not mention schools, students, or teachers. [Cited; vendor]

xAI. “Welcome to Grok” (docs overview). <https://docs.x.ai/grok/overview>. Accessed 30 August 2026. **Vendor.** Distinguishes Grok chat product from Grok Bot. [Cited; vendor]

xAI. “xAI and El Salvador Pioneer the World’s First Nationwide AI Education Program.” 11 December 2025. <https://x.ai/news/el-salvador-partnership>. Accessed 30 August 2026. **Vendor announcement.** Grok (the assistant, not Grok Bot) across more than 5,000 public schools over two years; “personalized learning to over one million students.” Announced national program, not a documented completed deployment; **no outcome evaluation found** as of 30 August 2026. [Cited; vendor]

5. Journalism and ratings

AI Incident Database. “Incident 826: Character.ai Chatbot Allegedly Influenced Teen User Toward Suicide.” <https://incidentdatabase.ai/cite/826/>. Accessed 30 August 2026. Timeline: use from April 2023; death 28 February 2024; suit 22 October 2024; Google/Character mediated settlement disclosed 7 January 2026. [Cited]

American Enterprise Institute. “Can an AI-Powered Tutor Produce Meaningful Results?” Interview with Kristen DiCerbo (Khan Academy). <https://www.aei.org/commentary/can-an-ai-powered-tutor-produce-meaningful-results/>. Accessed 30 August 2026. Partner-district users about 68,000 (2023-24) to more than 700,000 (2024-25); districts 45 to 380-plus — **vendor claims**, not efficacy RCTs. [Cited]

Associated Press. “AI Chatbot Pushed Teen to Kill Himself, Lawsuit Alleges.” <https://apnews.com/article/chatbot-ai-lawsuit-suicide-teen-artificial-intelligence-9d48adc572100822fdb3c90d1456bd0>. Companion on the Conway ruling: <https://apnews.com/article/ai-lawsuit-suicide-artificial-intelligence-free-speech-ccc77a5ff5a84bda753d2b044c83d4b6>. Accessed 30 August 2026. Allegations. [Cited]

Beduya, Jose. “Cornell Expands AI Critical Literacy Program for Incoming Students.” *Cornell Chronicle*, August 2026. <https://news.cornell.edu/stories/2026/08/cornell-expands-ai-critical>

literacy-program-incoming-students. Accessed 30 August 2026. Four Canvas modules: What is Generative AI?; Ethical Questions Surrounding Generative AI; GenAI and Learning; Creating an AI Policy for Personal Use. CTI-reported literacy gains, not a published RCT. [Cited]

Burns, Mary. “What the Research Shows about Generative AI in Tutoring.” Brookings Institution, 27 January 2026. <https://www.brookings.edu/articles/what-the-research-shows-about-generative-ai-in-tutoring/>. Accessed 30 August 2026. Think-tank commentary. ITS vs genAI-tutor distinction. Not itself an RCT. [Cited]

CBS News / 60 Minutes. “How Classroom AI Khanmigo Can Help Students in Emotional Distress.” 8 December 2024. <https://www.cbsnews.com/news/how-classroom-ai-khanmigo-can-help-students-in-emotional-distress-60-minutes/>. Accessed 30 August 2026. Vendor claims: teacher-visible logs; self-harm flags; OpenAI contractually not training on Khanmigo data. Claims, not an independent audit. [Cited]

Chalk, Dan. “Midland Public Schools Selects SchoolAI for Grades 6-12.” *OurMidland*, 19 August 2026. <https://www.ourmidland.com/news/article/midland-schools-school-ai-22393016.php>. Accessed 30 August 2026. Board selection, not outcomes. [Cited]

Chalkbeat New York. “What NYC’s New AI School Rules Say, and What Still Remains Unclear.” 24 March 2026. <https://www.chalkbeat.org/newyork/2026/03/24/preliminary-ai-policy-nyc-schools/>. Accessed 30 August 2026. [Cited]

Clemson News. “GPT-5.5, Codex and Clemson-Hosted Models Give Researchers New Options for AI-Assisted Work.” <https://news.clemson.edu/gpt-5-5-codex-and-clemson-hosted-models-give-researchers-new-options-for-ai-assisted-work/>. Accessed 30 August 2026. Institutional access, not learning outcomes. [Cited]

CNBC. “Cursor Announces Major Update as AI Coding Agent Battle Heats Up.” 24 February 2026. <https://www.cnbc.com/2026/02/24/cursor-announces-major-update-as-ai-coding-agent-battle-heats-up.html>. Accessed 30 August 2026. Company-provided stat that about 35 percent of Cursor’s own PRs come from cloud agents — not an education finding. [Cited]

CNBC. “OpenAI Rolls Out ‘ChatGPT for Teachers’ for K-12 Educators.” 19 November 2025. <https://www.cnbc.com/2025/11/19/openai-chatgpt-for-teachers-k-12-educators.html>. Accessed 30 August 2026. Not for students; first cohort about 150,000 educators; free through June 2027 **in this report**. OpenAI’s later product page restates free-through June 2028. Do not mash with the AFT Academy 400,000 five-year *target*. [Cited]

CNN Business. “Parents of 16-Year-Old Adam Raine Sue OpenAI, Claiming ChatGPT Advised on His Suicide.” 26 August 2025. <https://www.cnn.com/2025/08/26/tech/openai-chatgpt-teen-suicide-lawsuit>. Accessed 30 August 2026. Allegations. [Cited]

CNA. “All University, Polytechnic and ITE Students to Learn AI Skills from 2027.” <https://www.channelnewsasia.com/singapore/ai-artificial-intelligence-skills-courses-ihl-university-polytechnic-ite-students-2027-6133781>. Accessed 30 August 2026. [Consulted]

CNA. “SUSS Drops AI Detector as More Singapore Universities Question Reliability of Such Tools.” <https://www.channelnewsasia.com/singapore/ai-detection-tool-university-cheating->

plagiarism-suss-nus-sutd-sit-smu-6337431. Accessed 30 August 2026. SUSS disabled Turnitin's AI writing score; NTU to discontinue from 2027 as "fundamentally unreliable." [Cited]

CoinDesk. "El Salvador Partners with Elon Musk's Grok in AI-Powered Education for 1M Students." 11 December 2025. <https://www.coindesk.com/tech/2025/12/11/el-salvador-partners-with-elon-musk-s-grok-in-ai-powered-education-for-1m-students>. Accessed 30 August 2026. Journalism restating the xAI announcement. Not independent verification. [Cited]

Common Sense Media. *CSM AI Risk Assessment: ChatGPT-5*. 23 October 2025. <https://www.common Sense Media.org/content/files/csm-ai-risk-assessment-chatgpt-5-10232025.pdf>. Accessed 30 August 2026. "Teens should not use ChatGPT for mental health or emotional support." Guardrails weaken in extended conversations. Age 13+ with parental permission; teens bypass by lying about age. [Cited]

Common Sense Media / Youth AI Safety Institute. "Risk Assessments." <https://institute.common Sense Media.org/risk-assessments>. Accessed 30 August 2026. Independent third-party ratings. Khanmigo **low risk**, 8 August 2024. ChatGPT-5 **high risk**, 23 October 2025. Character.AI **unacceptable**, 10 April 2025. Social AI companions **unacceptable**, 16 July 2025. Meta AI **unacceptable**, 15 August 2025. Gemini K-12 **high risk**, 9 January 2026. Grok (the **chat** product, not Grok Bot) **unacceptable**, 22 January 2026. Google Search AI Overview / AI Mode **unacceptable**, 14 July 2026. Remaining CSM rating pages other than Khanmigo and ChatGPT-5 were not independently re-opened on the mid2 pass; they match the responsible-use memo. [Cited]

Culatta, Richard. "Profile of an AI-Ready Graduate." ASCD, 14 July 2025. <https://www.ascd.org/blogs/profile-of-an-ai-ready-graduate>. Accessed 30 August 2026. Headings: Learner, Researcher, Synthesizer, Problem Solver, Connector, Storyteller — the same six as ISTE Live 2026. Do not cite as an Ideator list. [Cited]

EdTech Magazine. "AI Agents for Education: Uses, Examples and Implementation." December 2025. <https://edtechmagazine.com/higher/article/2025/12/ai-agents-higher-education-transforming-student-services-and-support-perfcon>. Accessed 30 August 2026. Quotes EDUCAUSE Jenay Robert (37 percent chatbot licenses — **chatbots, not agents**) and UT-Knoxville Dan Harder. Johns Hopkins Agent Laboratory (authors claim 84 percent cost decrease) and UT Verse: journalism-sourced; primary pages **not fetched**. [Cited; campus examples UNVERIFIED against primary docs]

EdTech Magazine. "ISTE Live 26: What Does an AI-Ready Graduate Look Like?" July 2026. <https://edtechmagazine.com/k12/article/2026/07/istelive-26-what-does-ai-ready-graduate-look>. Accessed 30 August 2026. Conference report of the ISTE Profile: Learner, Researcher, Synthesizer, Problem Solver, Connector, Storyteller. Hu-Ann Wren, Patton Elementary (Austin, TX): conference-reported classroom practice, not a districtwide agent program. [Cited]

Fitzpatrick, Dan. "AI Won't Replace Teachers, but It Will Replace Teachers Who Don't Use AI." LinkedIn. <https://www.linkedin.com/pulse/ai-wont-replace-teachers-who-dont-use-dan-fitzpatrick>. Accessed 30 August 2026. Folklore tracking: earliest education-specific use found of the replacement slogan. Not a union, ministry, or peer-reviewed source. Recycled template. [Cited]

Flagler Live. "Artificial Intelligence Integration Rises in Flagler County Schools." 2026. <https://flaglerlive.com/ai-flagler/>. Accessed 30 August 2026. Grades 6-plus 71 percent; Magic-

School all levels teacher-opened; Khanmigo/Gemini/Brisk 6-12; enterprise ChatGPT; Lightspeed alerts; opt-out; 2,275 PD hours. Implementation facts, not outcome proof. [Cited]

Franczek P.C. “New Artificial Intelligence Developments for Illinois Schools: ISBE Releases Statewide Guidance and Governor Pritzker Signs SB 2909.” JDSupra. <https://www.jdsupra.com/legalnews/new-artificial-intelligence-9071517/>. Accessed 30 August 2026. Law-firm alert. ISBE primary PDF not fetched; “agentic AI” mention **UNVERIFIED** against the state document. [Cited; ISBE PDF UNVERIFIED]

GlobeNewswire / SchoolAI. “SchoolAI Earns ESSA Tier 3 Certification...” 23 April 2026. <https://www.globenewswire.com/news-release/2026/04/23/3279756/0/en/SchoolAI-Earns-ESSA-Tier-3-Certification-Backed-by-Research-Showing-28-Critical-Thinking-Gain.html>. Pointer to the vendor research page. [Consulted]

Gothamist. “NYC Schools Get AI Guidance Using ‘Red Light, Green Light’ Model.” <https://www.gothamist.com/news/nyc-schools-get-ai-guidance-using-red-light-green-light-model>. Accessed 30 August 2026. [Cited]

Guerra-López, Ingrid. “AI Won’t Replace Teachers—But Teachers Who Use AI Will Change Teaching.” *Education Week*, opinion, 17 October 2025. <https://www.edweek.org/technology/opinion-ai-wont-replace-teachers-but-teachers-who-use-ai-will-change-teaching/2025/10>. Accessed 30 August 2026. Different verb from the folklore slogan (*change teaching*, not *replace those who don’t*). Anti-replacement. [Cited]

Idaho Education News. “School Leaders Grapple with AI Use and Policies for Teachers and Students.” <https://www.idahoednews.org/top-news/school-leaders-grapple-with-ai-use-and-policies-for-teachers-and-students/>. Accessed 30 August 2026. West Ada MagicSchool grades 3-12 teacher rooms; ChatGPT blocked on elementary/middle devices. [Cited]

Idaho Press. “West Ada Adopts Artificial Intelligence Policy.” https://www.idahopress.com/news/local/west-ada-adopts-artificial-intelligence-policy/article_4f57387d-8935-4055-b8e7-a63789ef1f16.html. Accessed 30 August 2026. [Cited]

Palm Coast Observer. “Flagler Schools Stay Ahead of the Curve in AI Education.” 6 May 2026. <https://www.observerlocalnews.com/news/2026/may/06/flagler-schools-stay-ahead-of-the-curve-in-ai-education/>. Accessed 30 August 2026. [Cited]

Reuters. “Google, AI Firm Must Face Lawsuit Filed by a Mother over Suicide of Son, US Court Says.” 21 May 2025. <https://www.reuters.com/sustainability/boards-policy-regulation/google-ai-firm-must-face-lawsuit-filed-by-mother-over-suicide-son-us-court-says-2025-05-21/>. Accessed 30 August 2026. Judge Anne Conway, M.D. Fla., denied early First Amendment dismissal. [Cited]

Singer, Natasha. “Teen Girls Confront an Epidemic of Deepfake Nudes in Schools.” *New York Times*, 8 April 2024. <https://www.nytimes.com/2024/04/08/technology/deepfake-ai-nudes-westfield-high-school.html>. Accessed 30 August 2026. [Cited]

Verma, Pranshu. “A Professor Accused His Class of Using ChatGPT, Putting Diplomas in Jeopardy.” *Washington Post*, 18 May 2023. <https://www.washingtonpost.com/technology/2023/05/18/texas->

professor-threatened-fail-class-chatgpt-cheating/. Accessed 30 August 2026. Texas A&M University-Commerce mass false-positive episode. Cited by NEA 2024. [Cited]

Wonkhe (UK student voice). Quoted in the responsible-use memo: “students’ academic help may be determined by their financial resources rather than their academic competence.” Primary Wonkhe URL was not in the source JSON — **UNVERIFIED**. Wording as the memo recorded it. [Cited via memo; URL UNVERIFIED]

Compiled 30 August 2026. Unfetched = UNVERIFIED. Vendor counts remain labelled as such. Cut claims in the errata are not revived here.