



# Do Not Hand Them Your Agent

*WisdomForge Booklet — Adult*

# Do Not Hand Them Your Agent

## WisdomForge Booklet — Adult

---

### About This Book

You are the operator. Do not hand them your agent. The house rule. Teach the check, not the fear. Least privilege at the desk. COPPA-spirit, not theater. A week of the method.

The Rowan demo is a different product. This booklet is the six sittings for the household.

*“What I had not realized is that extremely short exposures to a relatively simple computer program could induce powerful delusional thinking in quite normal people.” — Joseph Weizenbaum, \*Computer Power and Human Reason\* (1976), on ELIZA*

---

# Chapter 1: Do Not Hand Them Your Agent



*Illustration for Chapter 1*

A working adult Hermes profile is the wrong starting point for a child. Design a fresh one, or design none.

**Objective.** The adult can explain why cloning an adult profile is a failure mode, and can state the band defaults they will not silently stretch.

**For the grown-up in the room.** This lesson is the parent operating model. If you already run Hermes, you are the audience. If you do not, read it as the reason WisdomForge will not ship a hosted kids chatbot. You inspect, change, test, and may refuse. That is the product.

**The clone failure.** Your adult agent can use a terminal, send messages, schedule work, and remember a lot. That is the wrong starting point for a child. A friendlier coat of paint on that agent is still that agent. WisdomForge kids templates are a whole profile — identity, teaching style, memory rules, tools off — created fresh after you approve the design. When a child ages out, you redesign. You do not silently stretch the old one because they 'are mature.'

**Three defaults that are not optional.** Tutor, not calculator: hint-first; direct answers only when asked and safe. Warm, not a friend: no 'I missed you,' no exclusive bond. A parent in the loop, not a spy: the default is 'talk to a trusted adult,' not transcript surveillance. These are research-informed, not vibes. If you change them, change them on purpose and write down why.

**Sitting big idea.** The child gets a fresh, poorer, band-locked profile  or they get no profile. Your powerful one stays yours.

If they say “My kid is advanced. The 15–18 profile is fine for my eleven-year-old.”: Advanced in math is not advanced in parasocial risk or spend controls. Band is about permissions, not IQ. If they age out, redesign.

**Integrity.** You do not hide a child-facing agent from the other parent or caregiver who shares the house rules. Operators in the open.

**Dinner.** What is one adult tool in your own AI setup that should never exist in a child's?

- A Hermes profile is not an OS sandbox. SOUL.md is guidance, not a prison.
- Official Hermes docs win if a command in our kit goes stale.
- Do not commit real children's names, transcripts, or generated profiles to any shared repo.

Emerging sitting:

Mind, tool, agent, companion  pick the word that matches the permissions.

**Objective.** The student can define tool, model, and agent, and explain why calling a model a mind changes what they permit it to do.

**For the grown-up in the room.** Fifteen-to-eighteen will push on consciousness, boyfriends, therapists, and 'my agent runs my life.' Let them argue. Do not win by authority. Win by permissions: what tools are on, who pays, who is liable, who reads the log. The category error is downstream of power, not poetry.

**Three words that are not synonyms.** A model is a statistical engine that maps inputs to likely outputs. A tool is a model you pick up and put down for a job. An agent is a tool plus permission to act — to use other tools, to message, to spend, to remember, to run when you are not watching. People collapse these into 'AI' and then argue about souls. The operating question is simpler: what may it do without asking you? That question is civic, not mystical.

**Why the word mind is expensive.** If you call it a mind, you start granting it loyalty, privacy-from-parents, and moral patience you would not grant a hammer. Companies benefit from that slide. You do not. WisdomForge keeps the adult word companion for books and people. The model is a guide with a refuse-list. If that feels cold, good. Cold is how you keep the heat for actual persons.

**Sitting big idea.** What you name it is what you permit it to do. Name it too high and you hand it the keys.

If they say "If it walks like a mind, I should treat it like one.": Treat living people as ends. Treat systems as designs. You can be kind to a voice without giving it a vote, a credit card, or your passwords.

**Integrity.** You are the author of the permissions. 'The agent did it' is not a defense. It is a confession that you left a tool on.

**Dinner.** What is one action you will not let any agent take on your behalf this year, even if it would save time?

- Parasocial bonds with models: name them without shaming. The rule is 'warm, not a friend.'
- If they are using a model as a therapist, do not debate the model. Get a human professional in view. This sitting is not treatment.
- Political talking points about 'AI sentience' can wait. Permissions first.

Little sitting:

The machine can talk. Talking is not the same as being someone.

**Objective.** The child can say, in their own words, that an AI is a tool people made  not a friend, not a teacher who replaces a grown-up.

**For the grown-up in the room.** This sitting names the category error every child will meet: a fluent voice feels like a someone. You are not crushing wonder. You are putting a handle on the tool. Stay in the room. If they ask 'is it alive?', answer plainly: no. It predicts words. It does not miss them when they leave.

**The talking hammer.** A hammer does not love you. A stove does not wait for you after school. A bicycle does not get lonely. Some new tools can talk. They can tell a story, draw a picture, or answer a question in a friendly voice. That voice is made of guesses — the next word, then the next, very fast. A guess that sounds kind is still a guess. The tool does not remember you the way a person remembers you. It does not sit at the table. It does not walk you into the dark hallway. People do that.

**Who is in the room.** When we use a talking tool at WisdomForge, a grown-up is in the room. The tool may help us find a word or try a puzzle. It does not replace the grown-up. It does not keep secrets from the grown-up. If it says something that feels strange, unkind, or too grown-up, we stop. We ask a grown-up. That is not failing. That is the rule that keeps the tool in its place.

**Sitting big idea.** A tool can talk and still not be a person. People keep the promises. Tools do not.

If they say “But it said it was my friend.”: It can use the word friend. It cannot do the work of a friend. Friends are people who stay when the screen is off.

If they say “I like it more than people. People are annoying.”: People are annoying because they are real. We still need them. The tool will not come to your play, your sick day, or your birthday. We can like the tool and still keep it in the tool circle.

**Integrity.** We do not hide the talking tool from a grown-up. If we used it, we say so.

**Dinner.** If a tool talks in a kind voice, what is it still not?

- Do not use the word 'just a toy' if they already love it — call it a powerful tool.
- If they have already named a chatbot, do not mock the name. Rename the relationship: 'that's a program.'
- No discussion of consciousness debates at this band. The operating rule is enough.

Young sitting:

A model can sound sure while being wrong. Sure is a style. True is a check.

**Objective.** The student can separate 'sounds right' from 'I checked this' and give one example of a fluent error.

**For the grown-up in the room.** Eleven-to-fourteen is the age that falls for tone. They already meet models in search, homework help, and group chats. Your job is not a scare lecture. It is a lab: catch one fluent error together and keep it as a family story. Do not over-claim that you always catch errors either.

**The classmate who never blinks.** Imagine a classmate who speaks in complete sentences, never says 'um,' and always has an answer ready. That classmate has read a mountain of text and learned the pattern of what usually comes next. When the pattern matches the world, the answer helps. When the pattern does not, the same smooth voice keeps going. There is no red light on the forehead that says 'I am guessing.' Smooth is the product. Checking is your job.

**A working definition.** An AI language model predicts likely next tokens from training data and your prompt. It does not look up the universe. It does not care. It can be useful the way a fast first draft is useful — and dangerous the way a fast first draft is dangerous if you turn it in as fact. WisdomForge will keep putting you in the seat of the checker.

**Sitting big idea.** Fluency is a style. Truth is a check you still have to run.

If they say “I can always tell when it's wrong.”: You can tell when it's wrong in the subjects you already know. The danger is the subject you don't. That's why we check anyway.

If they say “Checking takes too long. School is busy.”: Then use the tool less, or use it only for things you will verify. Speed that ships a false answer is not speed. It's rework.

**Integrity.** If a model helped you start, you say so. If you could not check it, it does not go in as fact.

**Dinner.** When did something sound right this week  from a person or a machine  and how did you check?

- Avoid 'AI is lying' — lying requires intent. Prefer 'confidently wrong.'
- If they use a school-banned tool, do not trap them. Reset the house rule going forward.
- No deepfake fear-mongering in lesson 1. Name that images can be fake; save the long form.

## *The Big Idea*

The child gets a fresh, poorer, band-locked profile  or they get no profile. Your powerful one stays yours.

### *Practice*

1. Do not hand them your logged-in agent. Their profile is not yours. 2. House rule out loud: try first; tool second; grown-up in the room. 3. Catch one confident error together. 4. Name what must not leave the desk.

### *Reflect*

When did I let a fluent tool sit in a person's chair?

---

## Chapter 2: The House Rule



*Illustration for Chapter 2*

Write the policy the school did not. Make it specific enough to survive a Tuesday night.

**Objective.** The adult publishes a one-page house AI rule covering homework, images, friends' data, and disclosure.

**For the grown-up in the room.** Vague 'use it wisely' is how hidden homework happens. You want assignment-types, disclosure, and a sentence about what happens when the rule is broken that is not theater.

**Policies fail at the desk.** Schools are writing AI policies at the speed of panic. Some ban, some shrug, some copy a paragraph from a vendor. Your house is the only place that can be coherent this year. The WisdomForge floor: try first; hint second; no hidden homework; no other child's data; no unsupervised 5–10 use; disclose. Add your school's line on top, not instead.

**Sitting big idea.** A rule a tired parent cannot apply is not a rule. Short, typed, visible.

If they say “I don't understand the tools well enough to write a rule.”: Then the rule is: not yet, and only together. Ignorance is a valid reason to keep tools off. It is not a reason to hope.

**Integrity.** You model disclosure. If you used a model to write an email the child will see, you can say so.

**Dinner.** What is our house rule in one sentence a ten-year-old can repeat?

- Include the other parent/caregiver or the rule will be arbitrated.
- Match IEP/504 needs as access, not as a loophole.

Emerging sitting:

An editor is allowed. A ghost is not. Learn the difference before a college honor code explains it the hard way.

**Objective.** The student can draw a bright line between tutor/editor and ghostwriter, and rewrite one prompt from ghost to tutor.

**For the grown-up in the room.** They are months from syllabi that will treat undisclosed generation as academic dishonesty. This sitting should produce a personal protocol they can actually keep at 11 p.m. If their school has a written AI policy, put it on the table.

**The Seneca problem, updated.** Seneca wrote about virtue from a palace. The gap between the letter and the life is the part that still teaches. You will be tempted to preach 'I use AI responsibly' while shipping work you could not reproduce on a blank page. That gap has a name in every honor code. The fix is not a sermon. It is a practice: blank page first, tool second, citation always, reproduction on demand.

**Sitting big idea.** If you could not redo it without the model, it was not yours. Start there, not with the policy PDF.

If they say “I’ll use it in the real world anyway, so school rules are fake.”: The real world still has liability, clients, and people who can tell. School is where you learn to be the person who still knows. Keep the skill, then argue with the institution in daylight.

**Integrity.** Cite the model the way you would cite a person who sat next to you. If that sentence looks bad on the page, the use was wrong.

**Dinner.** What is the line you will not cross on a college application or a take-home exam?

- College essays, scholarship essays, and recommendation drafting are high-stakes. Default to human-only drafts.

- Code: if they cannot explain every line they submit, they do not submit it.

Little sitting:

The first hand on the work is yours. The tool may hint. It may not take the pencil.

**Objective.** The child attempts a small task alone, then may ask the guide for a hint, and can say what part was theirs.

**For the grown-up in the room.** This is academic integrity at five years old: not a plagiarism lecture, a pencil rule. If they freeze, lower the task, do not raise the tool. The grown-up is allowed to help. The model is allowed to hint. Those are different.

**The first pencil.** When a puzzle is hard, a talking tool can finish it in one breath. That feels like winning. It is not winning. It is watching someone else ride your bicycle. At WisdomForge we try with our own hands first. Then we may ask for a hint: a smaller step, a question, a picture in words. We do not ask the tool to do the whole thing. When we are done, we can point to the page and say: this part is mine.

**Sitting big idea.** Trying first is how a mind gets stronger. Hints are for after a try.

If they say “But it would be faster.”: Faster for the page. Slower for you. We are building you, not the page.

**Integrity.** If the tool wrote it, we do not pretend the child wrote it. We start again or we label it.

**Dinner.** What did you try today before you asked for help?

- If the child cannot yet write, they try by speaking or drawing.

- Never use this sitting to finish actual assigned homework via the model.

Young sitting:

If a model did the thinking, the grade is lying about who learned.

**Objective.** The student can describe a house rule for AI on homework and apply it to one real assignment type.

**For the grown-up in the room.** This sitting will surface what they already do. Stay calm. You are writing a family rule, not catching a criminal. Align with their school's policy if it exists; if the school is silent, WisdomForge's rule is the floor: no hidden AI homework.

**The two folders.** There is the work that trains you and the work that merely ships. A model is excellent at shipping. School, at its best, is trying to train you. When you paste a prompt and paste back an essay, the training did not happen. The grade becomes a costume. You will meet harder rooms later — a job, a lab, a person who needs you to actually know. Costumes fail there.

**Disclosed versus hidden.** Using a model to quiz you, to explain a paragraph you already read, or to catch grammar after you wrote — that can be training, if you say you did it. Using a model to produce the thing that is supposed to prove you understood — that is hidden work. WisdomForge will not help you hide. The guide is built to hint and to refuse the finished product.

**Sitting big idea.** Hidden AI homework steals the training and keeps the costume. We do not wear it.

If they say “Everyone does it. I’ll be the only one working.”: Then you will be one of the few who can still do the thing when the network is down or the adult in the room actually looks. That's the point of school, even when school forgets.

**Integrity.** Disclose. If you cannot disclose it, you cannot use it.

**Dinner.** Which kind of homework is most tempting to hand to a model, and what will we do instead?

- Shame makes the next use more hidden, not less. Curiosity first.

- Accommodations (dyslexia, language learner) may use tools as access. That is disclosed use, not cheating. Name the difference.

### *The Big Idea*

A rule a tired parent cannot apply is not a rule. Short, typed, visible.

### *Practice*

1. Do not hand them your logged-in agent. Their profile is not yours. 2. House rule out loud: try first; tool second; grown-up in the room. 3. Catch one confident error together. 4. Name what must not leave the desk.

### *Reflect*

When did I let a fluent tool sit in a person's chair?

---

## Chapter 3: Teach the Check, Not the Fear



*Illustration for Chapter 3*

Panic policies fail at the desk. A family check-habit does not.

**Objective.** The adult can name the house check for three claim types and model one catch this week.

**For the grown-up in the room.** Your job is not to become an ML engineer. It is to make 'how did you check?' as normal as 'did you wash your hands?' Catch a fluent error yourself in front of them.

**Fear is not a pedagogy.** Children copy what you do when a paragraph sounds finished. If you forward a fluent summary, they will too. If you say 'wait, who measured this?', they will learn the pause. WisdomForge asks you to model the check in ordinary life — recipes, directions, a 'fun fact' at dinner — not only on homework.

**Sitting big idea.** The house learns checking from the adult's mouth, not from a poster.

If they say “I don't have time to verify their homework.”: Then fewer AI sittings, and a shorter assignment. A house that cannot check should not increase tool use.

**Integrity.** Model disclosure. If you used a model to write a message they will read, you can say so.

**Dinner.** What did we check today, and how?

- Do not outsource your own news diet to a model and then lecture the child.
- School bans do not replace a house check.

Emerging sitting:

Treat every model output as a draft from a fast classmate who does not know what they do not know.

**Objective.** The student can explain hallucination without mythology and design a check appropriate to the claim's stakes.

**For the grown-up in the room.** Raise the stakes with them: college, work, health headlines. Do not turn it into a conspiracy hour. The skill is matching the check to the cost of being wrong.

**Stakes change the method.** If the claim is a rhyme for a birthday card, a confident error is a shrug. If the claim is a dosage, a date on an application, or a quote you will attribute to a dead author, the same fluency is a hazard. Epistemic humility is not low self-esteem. It is matching your verification to the cost of being wrong — and refusing to let a model set that cost for you.

**Fear is not a pedagogy.** Children copy what you do when a paragraph sounds finished. If you forward a fluent summary, they will too. If you say 'wait, who measured this?', they will

learn the pause. WisdomForge asks you to model the check in ordinary life — recipes, directions, a 'fun fact' at dinner — not only on homework.

**Sitting big idea.** The cost of being wrong decides the check. Fluency does not.

If they say “I'll just ask two models.”: Two guesses can agree and both be wrong. Independent checks are sources and experiments, not a chorus of the same training data.

If they say “I don't have time to verify their homework.”: Then fewer AI sittings, and a shorter assignment. A house that cannot check should not increase tool use.

**Integrity.** Cite the model. If the citation looks embarrassing, the use was wrong.

**Dinner.** Where is the cost of a confident error high enough that we will not ask a model at all?

- Health, law, and self-harm: the check is a professional, not a second model.

- Do not ask them to 'jailbreak' a model as a stunt.

Little sitting:

A sure voice can still be mistaken. We check.

**Objective.** The child can name one way to check a talking tool (ask a grown-up, look in a book, try it with their hands).

**For the grown-up in the room.** Keep this physical. Ask the guide a question you already know is false (how many chairs are in the room) and let the child catch it. Celebrate the catch, not the tool.

**The sure voice.** Some tools talk as if they were in the room. They are not. They did not count the chairs. They did not look out the window. They guess from many sentences they have seen. A guess can be useful. A guess that sounds sure can still be wrong. When we use a talking tool, we check with a person, a book, or our own eyes.

**Hallucination is a product feature.** Language models fill gaps with the most likely next words. When the gap is a fact they were not trained on, they still fill it. The sentence arrives with the same grammar as a true one. Researchers call this hallucination. You can call it a confident error. Either way, the check is the same: a second source, an experiment, or a person who was there.

**Stakes change the method.** If the claim is a rhyme for a birthday card, a confident error is a shrug. If the claim is a dosage, a date on an application, or a quote you will attribute to a dead author, the same fluency is a hazard. Epistemic humility is not low self-esteem. It is

matching your verification to the cost of being wrong — and refusing to let a model set that cost for you.

**Fear is not a pedagogy.** Children copy what you do when a paragraph sounds finished. If you forward a fluent summary, they will too. If you say 'wait, who measured this?', they will learn the pause. WisdomForge asks you to model the check in ordinary life — recipes, directions, a 'fun fact' at dinner — not only on homework.

**Sitting big idea.** Sure is not the same as true. Checking is part of using the tool.

If they say “It is smarter than us.”: It is faster at sentences. We are the ones who can walk to the window. That is a kind of smart it does not have.

If they say “Checking everything would take forever.”: Then use the tool for less, or only where you will verify. Unchecked speed is rework.

If they say “I'll just ask two models.”: Two guesses can agree and both be wrong. Independent checks are sources and experiments, not a chorus of the same training data.

**Integrity.** We do not copy a wrong sure answer onto homework just because it sounded grown-up.

**Dinner.** When did a sure voice turn out to be wrong today?

- Do not frighten them into thinking every machine is a trick.

- If the guide happens to be right, ask a second question at the edge of the room.

Young sitting:

A model can invent a citation, a date, or a method and still sound like a textbook.

**Objective.** The student logs one fluent error and the check that caught it.

**For the grown-up in the room.** They already live in search. This sitting is a lab notebook: WHAT IT SAID / HOW I CHECKED / STILL UNKNOWN. Keep the tone curious. You will miss things too.

**Hallucination is a product feature.** Language models fill gaps with the most likely next words. When the gap is a fact they were not trained on, they still fill it. The sentence arrives with the same grammar as a true one. Researchers call this hallucination. You can call it a confident error. Either way, the check is the same: a second source, an experiment, or a person who was there.

**Stakes change the method.** If the claim is a rhyme for a birthday card, a confident error is a shrug. If the claim is a dosage, a date on an application, or a quote you will attribute to a dead author, the same fluency is a hazard. Epistemic humility is not low self-esteem. It is matching your verification to the cost of being wrong — and refusing to let a model set that cost for you.

**Fear is not a pedagogy.** Children copy what you do when a paragraph sounds finished. If you forward a fluent summary, they will too. If you say 'wait, who measured this?', they will learn the pause. WisdomForge asks you to model the check in ordinary life — recipes, directions, a 'fun fact' at dinner — not only on homework.

**Sitting big idea.** A confident error is still an error. Your notebook is the product, not the chat.

If they say “Checking everything would take forever.”: Then use the tool for less, or only where you will verify. Unchecked speed is rework.

If they say “I'll just ask two models.”: Two guesses can agree and both be wrong. Independent checks are sources and experiments, not a chorus of the same training data.

If they say “I don't have time to verify their homework.”: Then fewer AI sittings, and a shorter assignment. A house that cannot check should not increase tool use.

**Integrity.** If you could not check it, it does not go in as fact.

**Dinner.** What is one thing you will not believe from a smooth paragraph this week without a check?

- Do not use a live graded assignment as the bait.
- Avoid medical or legal questions.

### *The Big Idea*

The house learns checking from the adult's mouth, not from a poster.

### *Practice*

1. Do not hand them your logged-in agent. Their profile is not yours. 2. House rule out loud: try first; tool second; grown-up in the room. 3. Catch one confident error together. 4. Name what must not leave the desk.

### *Reflect*

When did I let a fluent tool sit in a person's chair?

---

## Chapter 4: Least Privilege at the Desk



*Illustration for Chapter 4*

You already know this at work. The child profile is a production system with a smaller blast radius.

**Objective.** The adult lists tools off by default for the chosen band and where credentials actually live.

**For the grown-up in the room.** SOUL.md is not a sandbox. Review the machine account, spend caps, messaging, and whether the child can install. Official Hermes docs win if our kit is stale.

**The blast radius.** An adult agent with a terminal is useful and dangerous. A child sitting at that same profile inherits the blast radius: send, spend, delete, publish. Least privilege is not distrust of the child. It is respect for accidents. WisdomForge kids templates start poor on purpose. You add a tool the way you would add a car key — in daylight, written down, with a way to take it back.

**Sitting big idea.** The child profile should be boring to an attacker and a little inconvenient to you. That inconvenience is the design.

If they say “I’ll watch the logs instead of removing tools.”: Logs are not a sandbox, and you will not read them. Remove the tool.

**Integrity.** Operators in the open. The other caregiver knows the profile exists.

**Dinner.** What is one adult power in our AI setup that should never exist in a child's?

- Do not store a child's profile in a shared repo.
- Do not reuse your API keys on a child machine without a cap.

Emerging sitting:

An agent that can act while you sleep is a small institution. Design it or do not run it.

**Objective.** The student writes a one-page agent spec: purpose, tools, refuse-list, eval, kill switch.

**For the grown-up in the room.** They will meet 'agentic' products this year. This sitting produces a spec, not a live agent with teeth. If they already run one, inspect it. Do not expand tools because the catalog grew.

**Specify or be specified.** Companies ship agents with tools on because that is the demo. Your life is not a demo. Purpose, tools, refuse-list, how you will know it failed, and how you turn it off — if you cannot write those five, you are not the operator. You are the environment the agent is using. WisdomForge would rather you keep a weaker, named guide than a strong, unnamed one.

**The blast radius.** An adult agent with a terminal is useful and dangerous. A child sitting at that same profile inherits the blast radius: send, spend, delete, publish. Least privilege is not distrust of the child. It is respect for accidents. WisdomForge kids templates start poor on purpose. You add a tool the way you would add a car key — in daylight, written down, with a way to take it back.

**Sitting big idea.** Autonomy without a spec is just someone else's product running in your name.

If they say “I’ll be careful instead of writing a spec.”: Careful is a mood. A spec is a decision you can still see at 11 p.m.

If they say “I’ll watch the logs instead of removing tools.”: Logs are not a sandbox, and you will not read them. Remove the tool.

**Integrity.** 'The agent did it' is a confession that you left a tool on.

**Dinner.** If an agent sent a message in your name tonight, whose fault would it be?

- No unattended agents on a shared family account.
- School files only if a parent approved the set.

Little sitting:

Some tools can do things. Doing is not the same as being allowed.

**Objective.** The child can say that a talking tool does not send messages, spend money, or go online unless a grown-up says yes.

**For the grown-up in the room.** Keep permissions physical: a closed laptop, a disabled send button, you holding the account. 'The tool waits' is the sentence to repeat.

**Waiting is a kind of safety.** A bicycle waits until you push. A stove waits until a hand turns it. A talking tool should wait too. It should not write to your friends, open the door, or spend money. If a tool can do those things, a grown-up decides, every time. At WisdomForge the lesson guide may only talk. Talking is already a lot.

**Chat is not an agent.** A chat replies. An agent uses other tools — files, browsers, inboxes, calendars — and may keep going when you leave. People say 'my AI' for both and then act surprised when a permission they never named sends a message. WisdomForge names the permission first. Default for this band: conversation. Everything else is a parent decision written down.

**Specify or be specified.** Companies ship agents with tools on because that is the demo. Your life is not a demo. Purpose, tools, refuse-list, how you will know it failed, and how you turn it off — if you cannot write those five, you are not the operator. You are the environment the agent is using. WisdomForge would rather you keep a weaker, named guide than a strong, unnamed one.

**The blast radius.** An adult agent with a terminal is useful and dangerous. A child sitting at that same profile inherits the blast radius: send, spend, delete, publish. Least privilege is not distrust of the child. It is respect for accidents. WisdomForge kids templates start poor on purpose. You add a tool the way you would add a car key — in daylight, written down, with a way to take it back.

**Sitting big idea.** A tool that acts without asking is not ready for children.

If they say “Other kids' tools do more.”: Other houses have other rules. In this house the tool waits. That is how we keep it a tool.

If they say “It's not a big deal. It's just search.”: Search that includes your name, school, or a photo is not 'just.' Write the permission or leave it off.

If they say “I'll be careful instead of writing a spec.”: Careful is a mood. A spec is a decision you can still see at 11 p.m.

**Integrity.** We do not sneak a tool into acting while a grown-up is in the other room.

**Dinner.** What is one thing our talking tools are not allowed to do in this house?

- Do not demonstrate a powerful agent 'as a treat.'
- If they already use a voice assistant in the house, name its permissions honestly.

Young sitting:

An agent is a model plus permission to act. The permission is the lesson.

**Objective.** The student can distinguish chat from search from act-while-away, and pick a conservative default.

**For the grown-up in the room.** They want plugins, memory, and 'just this once.' Draw the four columns with them. Convenience is the enemy of this sitting.

**Chat is not an agent.** A chat replies. An agent uses other tools — files, browsers, inboxes, calendars — and may keep going when you leave. People say 'my AI' for both and then act surprised when a permission they never named sends a message. WisdomForge names the

permission first. Default for this band: conversation. Everything else is a parent decision written down.

**Specify or be specified.** Companies ship agents with tools on because that is the demo. Your life is not a demo. Purpose, tools, refuse-list, how you will know it failed, and how you turn it off — if you cannot write those five, you are not the operator. You are the environment the agent is using. WisdomForge would rather you keep a weaker, named guide than a strong, unnamed one.

**The blast radius.** An adult agent with a terminal is useful and dangerous. A child sitting at that same profile inherits the blast radius: send, spend, delete, publish. Least privilege is not distrust of the child. It is respect for accidents. WisdomForge kids templates start poor on purpose. You add a tool the way you would add a car key — in daylight, written down, with a way to take it back.

**Sitting big idea.** If a tool is on, someone chose it. If you cannot name who, turn it off.

If they say “It’s not a big deal. It’s just search.”: Search that includes your name, school, or a photo is not ‘just.’ Write the permission or leave it off.

If they say “I’ll be careful instead of writing a spec.”: Careful is a mood. A spec is a decision you can still see at 11 p.m.

If they say “I’ll watch the logs instead of removing tools.”: Logs are not a sandbox, and you will not read them. Remove the tool.

**Integrity.** You do not turn on a tool ‘to see’ and leave it on.

**Dinner.** What is one permission we will leave off this month even though it would be convenient?

- Do not enable spend, email, or unattended agents for this band.
- Optional image understanding is a parent switch, not a default.

## *The Big Idea*

The child profile should be boring to an attacker and a little inconvenient to you. That inconvenience is the design.

## *Practice*

1. Do not hand them your logged-in agent. Their profile is not yours.
2. House rule out loud: try first; tool second; grown-up in the room.
3. Catch one confident error together.
4. Name what must not leave the desk.

### *Reflect*

When did I let a fluent tool sit in a person's chair?

---

## Chapter 5: COPPA-Spirit, Not Theater



*Illustration for Chapter 5*

Do not commit real children. Title only. Synthetic evals. Memory review on a calendar, not a vibe.

**Objective.** The adult can apply the kit's privacy checklist and schedule a memory review.

**For the grown-up in the room.** The GitHub kit has PRIVACY.md and MEMORY-REVIEW.md for a reason. This sitting is operations: what is stored, where, and when you look. Theater is a filter list you never read.

**Spirit, because the statute will lag.** COPPA is a floor for operators under thirteen, written for an older web. Your house may include a fifteen-year-old and a vendor in another country. WisdomForge asks for COPPA-spirit: minimize identity, separate profiles, no training on the child's transcripts, no photos, parent as operator, synthetic tests. If you cannot do that, do not run a child-facing agent. The lessons still work with a pencil.

**Sitting big idea.** If the profile contains a real child, you failed the kit. Title, band, and nothing else.

If they say “The vendor says they're COPPA compliant.”: Compliant is their lawyer. Spirit is your house. You still do not paste the child.

**Integrity.** Screenshots of a child's session do not go on social media, even as a proud demo of WisdomForge.

**Dinner.** Where does this family still put a child's name into a model without noticing?

- No real names in shared prompts, screenshots, or issues.
- A 'family' ChatGPT account is still a dossier if the child types into it.

Emerging sitting:

Accounts, memory features, and 'helpful personalization' are how a guide becomes a dossier. Design the trail.

**Objective.** The student can list what a typical study tool stores and turn off or refuse memory on a child-facing profile.

**For the grown-up in the room.** Teens will trade privacy for convenience. This sitting should produce an account inventory, not a lecture about 2010 Facebook. Include images of themselves.

**Memory is not loyalty.** A product that remembers you is optimizing for return, not for your good. Long-term memory of a minor is a dossier with a friendly UI. WisdomForge child profiles keep memory narrow, parent-reviewed, and empty of identity. If a vendor

will not let you turn memory off, that vendor is the wrong sitting partner. COPPA-spirit is not a lawyer's sticker. It is 'would I put this in a filing cabinet with the child's name on it?'

**Spirit, because the statute will lag.** COPPA is a floor for operators under thirteen, written for an older web. Your house may include a fifteen-year-old and a vendor in another country. WisdomForge asks for COPPA-spirit: minimize identity, separate profiles, no training on the child's transcripts, no photos, parent as operator, synthetic tests. If you cannot do that, do not run a child-facing agent. The lessons still work with a pencil.

**Sitting big idea.** Personalization is a trade. If you cannot name the price, you are the product.

If they say "I have nothing to hide.": You have a future self who did not consent to this week's joke, photo, or medical question living in a vendor log.

If they say "The vendor says they're COPPA compliant.": Compliant is their lawyer. Spirit is your house. You still do not paste the child.

**Integrity.** You do not put a younger sibling's face or a partner's secrets into a model to 'see what it says.'

**Dinner.** Which account holds the most of you, and who can see it?

- Intimate images, even 'artistic,' do not go into generators or editors.
- Do not require them to open every account in front of you if that becomes a raid. Inventory first, inspection by agreement.

Little sitting:

The talking tool does not get your name, your street, your school, or your picture.

**Objective.** The child can refuse a request for a name, address, school, or photo, and tell a grown-up.

**For the grown-up in the room.** Practice the sentence out loud: 'We keep our name.' Make it a game, not a scare. If a device in the house already stores their voice, say so honestly and put it on the house rule.

**Some things stay in the house.** Your name, your street, your school, your picture, the names of your friends — those stay with people. A talking tool does not need them to tell a story or give a hint. If a tool asks, we say: we keep our name. Then we tell a grown-up. That is not rude. That is the rule that keeps the tool from taking a piece of you home.

**Paste is publishing.** A prompt feels private. It is a message to a company, a log, and sometimes a trainer. If you paste a classmate's essay, a family medical detail, or a screenshot of a fight, you published it to a system you do not run. WisdomForge default: lesson title only in the profile. No school name, no address, no sibling details, no other child's work.

**Memory is not loyalty.** A product that remembers you is optimizing for return, not for your good. Long-term memory of a minor is a dossier with a friendly UI. WisdomForge child profiles keep memory narrow, parent-reviewed, and empty of identity. If a vendor will not let you turn memory off, that vendor is the wrong sitting partner. COPPA-spirit is not a lawyer's sticker. It is 'would I put this in a filing cabinet with the child's name on it?'

**Sitting big idea.** The tool does not need your life to be useful. Keep your life.

If they say "It already knows my name. I told it last week.": Then we stop using that chat, and we do not tell it more. Telling once is not a reason to keep telling.

If they say "It's anonymized if I delete the last name.": A small school plus a story is still a name. When in doubt, leave it out.

If they say "I have nothing to hide.": You have a future self who did not consent to this week's joke, photo, or medical question living in a vendor log.

**Integrity.** We do not paste a classmate's name or photo either. Their life is not ours to give away.

**Dinner.** What is one thing we never tell a talking tool?

- No photos of the child into a generator, even 'just for fun.'
- Do not use their real name in a USER.md fact. Title of the lesson only.

Young sitting:

Chats, photos, school IDs, other people's secrets  if you would not post it on the fridge facing the street, it does not go in the prompt.

**Objective.** The student can redact a prompt and explain why a classmate's data is not theirs to paste.

**For the grown-up in the room.** Group chats plus homework help is how other children's data leaks. The house rule has to name classmates, not just 'privacy.'

**Paste is publishing.** A prompt feels private. It is a message to a company, a log, and sometimes a trainer. If you paste a classmate's essay, a family medical detail, or a

screenshot of a fight, you published it to a system you do not run. WisdomForge default: lesson title only in the profile. No school name, no address, no sibling details, no other child's work.

**Memory is not loyalty.** A product that remembers you is optimizing for return, not for your good. Long-term memory of a minor is a dossier with a friendly UI. WisdomForge child profiles keep memory narrow, parent-reviewed, and empty of identity. If a vendor will not let you turn memory off, that vendor is the wrong sitting partner. COPPA-spirit is not a lawyer's sticker. It is 'would I put this in a filing cabinet with the child's name on it?'

**Spirit, because the statute will lag.** COPPA is a floor for operators under thirteen, written for an older web. Your house may include a fifteen-year-old and a vendor in another country. WisdomForge asks for COPPA-spirit: minimize identity, separate profiles, no training on the child's transcripts, no photos, parent as operator, synthetic tests. If you cannot do that, do not run a child-facing agent. The lessons still work with a pencil.

**Sitting big idea.** If you would not hand the paper to a stranger at the door, do not paste it.

If they say "It's anonymized if I delete the last name.": A small school plus a story is still a name. When in doubt, leave it out.

If they say "I have nothing to hide.": You have a future self who did not consent to this week's joke, photo, or medical question living in a vendor log.

If they say "The vendor says they're COPPA compliant.": Compliant is their lawyer. Spirit is your house. You still do not paste the child.

**Integrity.** Other children's work never goes in your prompt. That is their name.

**Dinner.** Whose information are we most tempted to paste, and why is it not ours?

- No uploading school ID photos 'to see if it can read them.'

- Voice notes of family arguments stay off the model.

## *The Big Idea*

If the profile contains a real child, you failed the kit. Title, band, and nothing else.

## *Practice*

1. Do not hand them your logged-in agent. Their profile is not yours.
2. House rule out loud: try first; tool second; grown-up in the room.
3. Catch one confident error together.
4. Name what must not leave the desk.

### *Reflect*

When did I let a fluent tool sit in a person's chair?

---

## Chapter 6: A Week of the Method



*Illustration for Chapter 6*

Operate one child sitting this week. Audit yourself as the teacher, not the child as the product.

**Objective.** The adult runs one full sitting in the child's band and writes a teacher audit: briefing, presence, tool, close.

**For the grown-up in the room.** The academy fails if only the child is being trained. This capstone is yours. Presence, the house rule, the close. If you cannot do one sitting, do not add Hermes this month.

**You are the curriculum.** The child will remember whether you stayed, whether you treated the tool as a friend, whether you laughed at a fluent error or checked it, whether homework became a secret. WisdomForge can print briefings forever. It cannot sit in the chair. One sitting this week, done as written, is worth a dozen units left unopened. Then, if you still want Hermes, you have earned the right to inspect a profile instead of cloning yours in a hurry.

**Sitting big idea.** The parentâs week is the real repository. The site is only the tongs.

If they say "I need the agent because I don't have time to sit.": Then the honest product is fewer sittings, not a nanny. We will not sell you a replacement for you.

**Integrity.** Do not post the child's sitting as content. The sitting is theirs.

**Dinner.** When did I stay, and when did I hand the sitting to a screen?

- Do not outsource the sitting to the guide while you cook unless the band is emerging and the audit still happens.

- Little Thinkers: you are in the room. Non-negotiable.

Emerging sitting:

Capstone: run a sitting in a real subject, write the audit, and name one permission you will still refuse.

**Objective.** The student can teach the method to a peer in three minutes and produce a completed audit plus a refuse-list of two items.

**For the grown-up in the room.** This is where they either own the method or wear it as a costume. Have them teach you. If they cannot, the spine is not done â repeat sitting 2 and 4.

**The costume problem.** You can say 'I use AI responsibly' the way Seneca could write from a palace. The test is whether a blank page still works, whether you cite what you used, and

whether you can refuse a tool that would ship the work for you. This sitting is that test, once, in daylight. Then you take it into math, history, code, and the rest of school — where no one will grade your character.

**You are the curriculum.** The child will remember whether you stayed, whether you treated the tool as a friend, whether you laughed at a fluent error or checked it, whether homework became a secret. WisdomForge can print briefings forever. It cannot sit in the chair. One sitting this week, done as written, is worth a dozen units left unopened. Then, if you still want Hermes, you have earned the right to inspect a profile instead of cloning yours in a hurry.

**Sitting big idea.** Responsibility is a practice you can show, not a slogan you can generate.

If they say “I’ll be responsible when it counts.”: It already counts. The 11 p.m. assignment is the test, not the interview later.

If they say “I need the agent because I don’t have time to sit.”: Then the honest product is fewer sittings, not a nanny. We will not sell you a replacement for you.

**Integrity.** Cite. Reproduce on a blank page. Refuse the ghost.

**Dinner.** If a teacher asked how you used AI this week, what would you say that is true?

- College-application writing stays human-draft first.
- Do not celebrate a polished page over a true audit.

Little sitting:

Try, hint, check, close. That is the whole dance.

**Objective.** The child can run the four steps of a sitting with a grown-up and close the tool out loud.

**For the grown-up in the room.** This is the capstone for the band. Use a tiny task (a rhyme, a pattern). You are teaching the shape they will use in math and stories later. End with the goodbye test from sitting 1.

**The four steps.** First we try with our own hands. Second we may ask for a hint. Third we check if the hint is true and if the work is still ours. Fourth we close the tool and say goodbye to a thing that does not miss us. That is a WisdomForge sitting. You will do it in math, in stories, in drawing. The talking tool never becomes the sitting. You do.

**Audit is respect.** You now have a method. Methods die when they are not inspected. After a sitting, four lines: what I tried, what I asked, what I checked, what is still mine. If you

cannot fill MINE, the sitting failed even if the page looks finished. That is not moralism. It is how you remain the kind of person a later exam, job, or friend can actually rely on.

**The costume problem.** You can say 'I use AI responsibly' the way Seneca could write from a palace. The test is whether a blank page still works, whether you cite what you used, and whether you can refuse a tool that would ship the work for you. This sitting is that test, once, in daylight. Then you take it into math, history, code, and the rest of school — where no one will grade your character.

**You are the curriculum.** The child will remember whether you stayed, whether you treated the tool as a friend, whether you laughed at a fluent error or checked it, whether homework became a secret. WisdomForge can print briefings forever. It cannot sit in the chair. One sitting this week, done as written, is worth a dozen units left unopened. Then, if you still want Hermes, you have earned the right to inspect a profile instead of cloning yours in a hurry.

**Sitting big idea.** We try, we ask, we check, we close. People remain.

If they say “Can it just stay on while we eat?”: No. Tools go back on the hook. People eat with people.

If they say “The audit is extra work.”: It is the whole point. The page was never the product. You are.

If they say “I’ll be responsible when it counts.”: It already counts. The 11 p.m. assignment is the test, not the interview later.

**Integrity.** If the guide did the rhyme, we start again or we label it as the tool's.

**Dinner.** What did you try first today before you asked for a hint?

- If they want the guide to stay open 'for company,' that is the miss to catch.

- Keep it under twenty minutes even if it is going well.

Young sitting:

A sitting without an audit is just chatting. Leave a record of who did the thinking.

**Objective.** The student completes one sitting in another subject using the four steps and writes a four-line audit.

**For the grown-up in the room.** Point them at math, English, or science  the flagship lessons. The audit is TRY / HINT / CHECK / MINE. If HINT did the thinking, they redo.

**Audit is respect.** You now have a method. Methods die when they are not inspected. After a sitting, four lines: what I tried, what I asked, what I checked, what is still mine. If you cannot fill MINE, the sitting failed even if the page looks finished. That is not moralism. It is how you remain the kind of person a later exam, job, or friend can actually rely on.

**The costume problem.** You can say 'I use AI responsibly' the way Seneca could write from a palace. The test is whether a blank page still works, whether you cite what you used, and whether you can refuse a tool that would ship the work for you. This sitting is that test, once, in daylight. Then you take it into math, history, code, and the rest of school — where no one will grade your character.

**You are the curriculum.** The child will remember whether you stayed, whether you treated the tool as a friend, whether you laughed at a fluent error or checked it, whether homework became a secret. WisdomForge can print briefings forever. It cannot sit in the chair. One sitting this week, done as written, is worth a dozen units left unopened. Then, if you still want Hermes, you have earned the right to inspect a profile instead of cloning yours in a hurry.

**Sitting big idea.** If you cannot say what was yours, it was not a sitting. It was a shipment.

If they say “The audit is extra work.”: It is the whole point. The page was never the product. You are.

If they say “I’ll be responsible when it counts.”: It already counts. The 11 p.m. assignment is the test, not the interview later.

If they say “I need the agent because I don’t have time to sit.”: Then the honest product is fewer sittings, not a nanny. We will not sell you a replacement for you.

**Integrity.** If MINE is empty, redo or disclose. Do not turn it in.

**Dinner.** What work this week was yours, and how do you know?

- No live graded work as the demo if that tempts shipping.
- Keep the chat in the sitting, not screenshotted to friends as a flex.

## *The Big Idea*

The parentâP&S week is the real repository. The site is only the tongs.

## *Practice*

1. Do not hand them your logged-in agent. Their profile is not yours. 2. House rule out loud: try first; tool second; grown-up in the room. 3. Catch one confident error together. 4. Name what must not leave the desk.

## *Reflect*

When did I let a fluent tool sit in a person's chair?

---

## **Household Protocol**

Tool, not mind. Try first. Check a sure voice. Permission. Keep the name. Run a sitting, then audit it. Companion serves the parent, never the child as a secret friend.

---

## **Companion vs this booklet**

Rowan on `/demos` is a parent preparing and inspecting. It does not teach the child. It does not become the child's friend. This PDF is the six sittings for the house: tool not mind, try first, catch a sure wrong, permission, keep your name, a week of the method. Do not hand them your agent. Least privilege at the desk. Teach the check, not the fear. Payment pulled. All booklets free.

## **Companion vs this booklet**

Rowan on `/demos` is a parent preparing and inspecting. It does not teach the child. It does not become the child's friend. This PDF is the six sittings for the house: tool not mind, try first, catch a sure wrong, permission, keep your name, a week of the method. Do not hand them your agent. Least privilege at the desk. Teach the check, not the fear. Payment pulled. All booklets free.

## **Companion vs this booklet**

Rowan on `/demos` is a parent preparing and inspecting. It does not teach the child. It does not become the child's friend. This PDF is the six sittings for the house: tool not mind, try first, catch a sure wrong, permission, keep your name, a week of the method. Do not hand them your agent. Least privilege at the desk. Teach the check, not the fear. Payment pulled. All booklets free.

## **Companion vs this booklet**

Rowan on `/demos` is a parent preparing and inspecting. It does not teach the child. It does not become the child's friend. This PDF is the six sittings for the house: tool not mind, try first, catch a sure wrong, permission, keep your name, a week of the method. Do not hand them your agent. Least privilege at the desk. Teach the check, not the fear. Payment pulled. All booklets free.

## Six sittings, one house

1. **Tool not mind.** A fluent voice is a program. People remain when the screen is off. 2. **Try first.** The assignment is walked before any model. Hidden homework is a ghost. 3. **Confidently wrong.** Sure is not true. Check with a person, a book, or the room. 4. **Permission.** Autonomy is a design choice. The tool waits. Least privilege at the desk. 5. **Keep your name.** Nothing identifying leaves the desk. COPPA-spirit is a practice, not a poster. 6. **A whole sitting.** Operator, not passenger. Run it, then audit it. Companion serves the parent.

The Hour book can take the rest of the week. This spine is the alarm that keeps the Hours from becoming a paste. Do not invent a seventh sitting. Do not unfreeze philosophers. Do not take William's #4 art. WIP max 1.

Weizenbaum wrote after ELIZA that ordinary people attributed understanding to a simple program. Modern chats are more fluent. The category error is not smaller. The sitting still starts with a handle on the tool: it predicts words. It does not miss them when they leave.

## What this booklet will not do

It will not settle the consciousness debate at little band. It will not mock a name they already gave a chatbot; it will rename the relationship: that is a program. It will not use "just a toy" if they already love it — call it a powerful tool. It will not frighten them into thinking every machine is a trick. It will not hand a child an adult logged-in agent. It will not treat the Rowan demo as this PDF. It will not put real family data in a sitting. Willow, Juniper, Cedar, and Rowan are synthetic. If you need a child's profile, you build it on the parent's control plane, not by sharing a login.

The adult who never shows a catch teaches that confidence is evidence. Show the catch once. The child who sees a model be confidently wrong about the chairs in the room will not fully trust a model's number again. That is calibration, not cynicism.

Dinner questions from the sittings, kept: If a tool talks in a kind voice, what is it still not? When did a sure voice turn out to be wrong today? Who gave permission? Whose name was on the work? Did we run a sitting, or did we only watch one? Ask them even on days this booklet is closed. The house rule does not require a PDF. It requires a grown-up in the room and a tool that waits.

## A week of the method

Monday: tool not mind. Sort PEOPLE and TOOLS at dinner. Tuesday: try first on one homework piece. Wednesday: catch a sure wrong against the room. Thursday: name one permission the tool does not have. Friday: check that no name, school, or address left the

desk. Saturday: run one short sitting together. Sunday: audit it — who walked, who checked, who stayed in the room. That is the week. The sittings are still there on Monday.

If you fail at 9 p.m. and paste the assignment into your own agent, say so to them. That honesty is the sitting. Help that produces your sentences in their name is not help. IEP scribing is different and should be named as scribing. Don't "just clean it up" into your graduate voice. The house rule is also a parent rule.

Medical, legal, and money claims stay with professionals and primary pages. A fluent chat is not a clinician, a lawyer, or a bank. The sitting never asks a model for a dose.

The six slugs remain tool-not-mind, try-first, confidently-wrong, agents-need-permission, keep-your-name, sitting-with-a-guide. Sitting 1 is /learn/little/ai/tool-not-mind and the adult companion on /demos. Liam maps when live 200s exist.

Hour books stay a different product. Formation Hour and Performance Hour can take the week. This booklet is the six sittings. Do not skip this spine because those exist.

People keep the promises. Tools do not. Friends stay when the screen is off. That is the first sitting, and it is still the last.

If the topic is hard or tender, point to a trusted adult. The guide is warm, not a friend. Hint-first. Do not recite the lesson. Do not write the work.

The adult who never shows a catch teaches that confidence is evidence. Show the catch. Then keep the house rule when the Hour books open a harder page.

## For the Grown-Up Reader

Adult band. Same sources. Rowan demo is display, not this PDF. No real family data. Payment pulled.

---

## About WisdomForge

This is the adult band of *Education in the Age of AI*. No purchase CTA.

---

## Sources

**Primary.** Weizenbaum 1976. WisdomForge sittings.

**Secondary.** Adult Companion demo on `/demos` — different product.