



Rhetoric Is a Household Practice

WisdomForge Booklet — Adult

Rhetoric Is a Household Practice

WisdomForge Booklet — Adult

About This Book

Model the ethos question. Do not outsource the phone call. Audit the arguments you share. Keep your voice alive. The speech that came from nowhere is not a person. Run the toolkit at the table.

*“Of the modes of persuasion furnished by the spoken word there are three kinds. The first kind depends on the personal character of the speaker; the second on putting the audience into a certain frame of mind; the third on the proof, or apparent proof, provided by the words of the speech itself.” — Aristotle, *Rhetoric* I.2 (1356a), W. Rhys Roberts translation*

Chapter 1: Modeling the Ethos Question



Illustration for Chapter 1

If you never ask 'should I trust this speaker?' out loud, the child will trust the smoothest one in the room. That is the machine.

Objective. The adult models the ethos question on one real source this week, asks it out loud in front of the child, and names the difference between fluency and earned trust.

For the grown-up in the room. The household runs on trust. You trust the school, the doctor, the coach, the news. The child watches who you trust and copies the pattern. If you never ask the ethos question out loud, the child learns that trust is about sounding right. The machine sounds the most right. This sitting is about one moment: you ask 'should I trust this source?' out loud, in front of the child, about something real. You name the experience, the accountability, the honest limits. You say whether the source earned the floor. That moment teaches more than a lecture on rhetoric.

The out-loud question. The ethos question is simple: should I trust this speaker? The hard part is asking it out loud. You ask it silently all day — you decide which review to trust, which news source to share, which model answer to accept. The child does not see the silent question. They see the result: you shared the post, you cited the model, you trusted the smooth voice. Ask it out loud once: 'Before I trust this, let me check. What experience does this person have? If they are wrong, is there a cost? Do they admit what they don't know?' That ten-second monologue teaches the child that trust is a decision, not a reflex.

What the child copies. The child will not remember your lecture on Aristotle. They will remember the Tuesday you paused before sharing a post and said: 'This sounds right, but who is the author, and what happens if they're wrong?' That pause is the ethos question, performed. The child who watches you pause will learn to pause. The child who watches you share without pausing will learn to share without pausing. The machine is the smoothest speaker in the room. The pause is how you keep the floor from going to smoothness alone.

Sitting big idea. Ask the ethos question out loud. The child copies the pause, not the vocabulary.

If they say "I don't have time to vet every source.": You do not need to vet every source. You need to pause once a week, out loud, on a source that matters. The child learns from the pause, not from a full audit of every post.

If they say "The model is usually right.": Usually is not always, and 'usually right' is not the same as 'earned the floor.' The model can be right and still have no ethos. The question is not 'is it correct.' It is 'should I trust it.' Those are different questions.

Integrity. Do not perform skepticism to look smart. The ethos question is honest. If the source earned the floor, the audit confirms it.

Dinner. When did we pause to ask 'should I trust this source?' this week, and what did we find?

- Do not use this to undermine a source the child needs to trust (a doctor, a teacher). The point is the question, not the takedown.

- Do not perform skepticism for show. If the source has ethos, say so. The question is honest, not theatrical.

Emerging sitting:

When everyone sounds like an expert, ethos is the question that sorts the real from the performed.

Objective. The student can evaluate the ethos of a source (human or model), name what would strengthen or weaken it, and explain why fluency alone does not establish ethos.

For the grown-up in the room. This sitting builds on the young version. The student is now evaluating sources, not just speakers. A YouTube expert, a news columnist, a model, a teacher. Each has a different kind of ethos. The skill is to sort them: what experience, what accountability, what honest admission of limits? The model scores high on fluency, zero on accountability. A named expert has a career to lose. The model does not. That difference is ethos in the age of fluency.

What ethos is made of. Ethos has three parts: experience, accountability, and honest limits. Experience: has the speaker done the thing, or only read about it? Accountability: if the speaker is wrong, is there a cost? A named expert who gets it wrong loses reputation. A model that gets it wrong answers the next question as if nothing happened. Honest limits: does the speaker say what they do not know? The model says 'as an AI, I may have limitations,' which is not honest limits. It is a disclaimer. Honest limits sounds like: 'I know this part, but not that part, and here is who does.' That sentence is ethos. The model cannot say it honestly because it does not know what it does not know.

The fluency trap. Fluency is the counterfeit of ethos. A model that writes a perfect essay on climate, marriage, or economics sounds like someone who has earned the floor. The essay is a collage of the voices it read, not the voice of someone who lived the thing. If you cannot sort fluency from ethos, you will trust the collage over the person. The way out is to ask three questions: what experience does this speaker have? If they are wrong, who pays? Do they name what they do not know? The model fails all three. A real expert passes at least two. That is the filter.

The out-loud question. The ethos question is simple: should I trust this speaker? The hard part is asking it out loud. You ask it silently all day — you decide which review to trust,

which news source to share, which model answer to accept. The child does not see the silent question. They see the result: you shared the post, you cited the model, you trusted the smooth voice. Ask it out loud once: 'Before I trust this, let me check. What experience does this person have? If they are wrong, is there a cost? Do they admit what they don't know?' That ten-second monologue teaches the child that trust is a decision, not a reflex.

What the child copies. The child will not remember your lecture on Aristotle. They will remember the Tuesday you paused before sharing a post and said: 'This sounds right, but who is the author, and what happens if they're wrong?' That pause is the ethos question, performed. The child who watches you pause will learn to pause. The child who watches you share without pausing will learn to share without pausing. The machine is the smoothest speaker in the room. The pause is how you keep the floor from going to smoothness alone.

Sitting big idea. Ethos is experience, accountability, and honest limits. Fluency is none of those, no matter how good it sounds.

If they say "The model has read more than any expert.": Reading is not experience. The model has read about surgery. You would not let it operate. Reading about a thing and doing a thing are different kinds of knowing. Ethos is about the doing.

If they say "Experts can be wrong too.": Yes. That is why ethos includes accountability and honest limits, not just credentials. The question is not 'is this source perfect.' It is 'if this source is wrong, will I find out, and will it cost them something?'

Integrity. Do not present a model's answer as your own ethos. If the authority came from the tool, the authority is the tool's. Name it.

Dinner. What source did we trust this week because it sounded right, not because it earned our trust?

- Do not let this become a blanket rejection of online sources. Some online sources have strong ethos. The point is to sort, not to dismiss.

- Academic credentials are one kind of ethos, not the only kind. A practitioner may have more ethos than a professor on a practical question.

Young sitting:

Ethos is whether the speaker has earned your trust. A machine that sounds wise has not.

Objective. The student can name what makes someone worth listening to on a topic, and explain why a fluent answer from a machine is not the same as earned trust.

For the grown-up in the room. Ethos is the first of the three classical appeals. It asks: should I trust this speaker? Not 'does it sound good' but 'has this person earned the right to speak on this?' Use concrete examples: a coach, a teacher, a grandparent. The model sounds like it has ethos. It does not. It has fluency. The difference is the lesson.

The question before the question. Before you ask 'is this true?' you ask 'should I be listening to this person?' That question is ethos. A coach who has played the game has ethos on the game. A stranger who read a blog post has less. The model is the stranger who read a billion blog posts. It sounds like the coach. It is not the coach. The first skill of rhetoric is to ask: who is speaking, and did they earn the floor? If you skip that question, you will trust the best talker in the room, and the best talker in the room is now a machine.

Fluency is not trust. The model speaks with confidence. Confidence is not ethos. Ethos is built from experience, accountability, and honesty about what you do not know. The model has none of those. It has confidence. If you cannot tell the difference, you will hand the floor to the machine every time, because the machine is the smoothest speaker you have ever met. Smooth is not earned. Earned is earned.

What ethos is made of. Ethos has three parts: experience, accountability, and honest limits. Experience: has the speaker done the thing, or only read about it? Accountability: if the speaker is wrong, is there a cost? A named expert who gets it wrong loses reputation. A model that gets it wrong answers the next question as if nothing happened. Honest limits: does the speaker say what they do not know? The model says 'as an AI, I may have limitations,' which is not honest limits. It is a disclaimer. Honest limits sounds like: 'I know this part, but not that part, and here is who does.' That sentence is ethos. The model cannot say it honestly because it does not know what it does not know.

The fluency trap. Fluency is the counterfeit of ethos. A model that writes a perfect essay on climate, marriage, or economics sounds like someone who has earned the floor. The essay is a collage of the voices it read, not the voice of someone who lived the thing. If you cannot sort fluency from ethos, you will trust the collage over the person. The way out is to ask three questions: what experience does this speaker have? If they are wrong, who pays? Do they name what they do not know? The model fails all three. A real expert passes at least two. That is the filter.

Sitting big idea. Ethos asks: did this speaker earn the floor? Fluency is not the same as earning it.

If they say "The model knows more than me.": It has read more than you. Reading is not the same as knowing. You know this topic from doing it. That is a different kind of earned.

If they say “The model has read more than any expert.”: Reading is not experience. The model has read about surgery. You would not let it operate. Reading about a thing and doing a thing are different kinds of knowing. Ethos is about the doing.

If they say “Experts can be wrong too.”: Yes. That is why ethos includes accountability and honest limits, not just credentials. The question is not 'is this source perfect.' It is 'if this source is wrong, will I find out, and will it cost them something?'

Integrity. Do not claim the model's ethos as your own. If the tool gave you the authority, cite the tool. The floor you did not earn is not yours.

Dinner. Who earned the right to speak into our lives today, and did we give them the floor or the best talker?

- Do not dismiss the model as worthless. The point is precision, not rejection.

- Do not conflate credentials with ethos. A credentialed person can be wrong. An uncredentialed person can be right. Ethos is earned, not stamped.

The Big Idea

Ask the ethos question out loud. The child copies the pause, not the vocabulary.

Practice

1. Before you share: who is talking? Can you name them? 2. Hard conversation: you make the call. Do not paste a generated script as if it were your mouth. 3. Audit one argument in the house. Claim and proof. 4. Say one paragraph in your own voice with the chat closed.

Reflect

When did I let a fluent stranger speak in our name?

Chapter 2: Do Not Outsource the Phone Call



Illustration for Chapter 2

If you let the model practice compassion for you, the child learns that caring is a paragraph you generate, not a call you make.

Objective. The adult models one moment of real empathy this week (a phone call, a presence, a listening) instead of generating a compassionate response, and names the difference for the child.

For the grown-up in the room. The temptation is real. A friend is struggling. You are tired. The model can write a compassionate message in seconds. It will sound better than what you would have said. The cost is that the child watches you generate compassion instead of giving it, and learns that caring is a text you produce, not a call you make. This sitting is about one moment: you choose the phone call, the visit, the presence over the generated paragraph. You do not need to announce it. But if the child asks why you called instead of texting, you say: 'A paragraph is not the same as being there.'

The generated condolence. A friend loses someone. You ask the model to help you write the right words. The model writes a beautiful paragraph. You send it. The paragraph is not wrong. It may even be what the friend needs to read. What is wrong is if you never call. The paragraph is the message. The call is the presence. A person in grief does not need a perfect paragraph. They need someone who will sit with them. The model can write the paragraph. It cannot sit. If you send the paragraph and skip the call, the child learns that compassion is a generated text. They will generate texts for the rest of their lives and wonder why their friendships feel thin.

What the child watches. The child does not hear your internal debate about whether to call. They see the result: you sent a text, or you called. They see whether you showed up, or whether you generated. If you always generate, they will always generate. If you sometimes call, sometimes visit, sometimes sit in the hard silence, they will learn that compassion has a cost and the cost is the point. The model is the easy path. The child needs to see you take the hard one sometimes. Not every time. Sometimes.

Sitting big idea. Compassion is a call you make, not a paragraph you generate. The child learns by watching which one you choose.

If they say "I'm bad at expressing feelings. The model is better.": Then let the model help you find the words. But your voice on the phone, clumsy and real, is worth more than the model's perfect paragraph. The person on the other end needs you, not a draft.

If they say "People don't want calls anymore.": Some don't. The ones who are struggling usually do. If you are not sure, ask. 'Can I call you?' is a sentence the model cannot say for you and mean it.

Integrity. Do not send a model's compassionate paragraph as if it were your own unrehearsed words without disclosure. That is manufactured pathos presented as earned. It is a kind of lie.

Dinner. When did we choose the call over the text this week, and what did the call give that the text could not?

- Do not shame yourself for using a model to draft a hard message. Drafting is not the problem. Sending the draft as if it were your own caring is the problem.
- Do not use this sitting to guilt the child about their phone habits. Model the alternative. Do not preach it.

Emerging sitting:

A model that cannot feel can still make you feel. That is a tool, not a relationship. Know which one you are using.

Objective. The student can identify when a model is using pathos, explain why the model cannot earn the feeling it produces, and describe the difference between empathy from a person and empathy from a machine.

For the grown-up in the room. This sitting names the thing the young version described. The model manufactures empathy. It writes a paragraph that sounds like a person who cares. The model does not care. It produces caring-sounding words. The student needs to know this not to become cold but to stop outsourcing their emotional life to a machine. The model can help you name a feeling. It cannot share one. The difference is the lesson.

What the model does when it cares. When you tell a model you are sad, it writes a compassionate paragraph. The paragraph is not empty. It contains words that a caring person might say. But the model did not feel concern. It produced the pattern of concern. That pattern can be useful. It can name a feeling you could not name. It can offer a perspective you had not considered. What it cannot do is share the feeling. Sharing a feeling requires a person who has something at stake. The model has nothing at stake. It will write the same paragraph for a thousand people and feel nothing for any of them. That is not a flaw. It is a property. Know the property and the tool is useful. Confuse the property and the tool is dangerous.

Empathy versus the pattern of empathy. Empathy is a person who has something at stake responding to your stake. A friend who stays on the phone at midnight has empathy. A model that writes 'I am here for you' at 3 a.m. has the pattern of empathy. The pattern is not nothing. It can be comforting. But it is not empathy. The friend loses sleep. The model loses nothing. The difference matters because if you replace your friends with the model, you will have the pattern of connection without connection. You will be less alone in words and more alone in fact. That is the cost of confusing the pattern for the thing.

The generated condolence. A friend loses someone. You ask the model to help you write the right words. The model writes a beautiful paragraph. You send it. The paragraph is not wrong. It may even be what the friend needs to read. What is wrong is if you never call. The paragraph is the message. The call is the presence. A person in grief does not need a perfect paragraph. They need someone who will sit with them. The model can write the paragraph. It cannot sit. If you send the paragraph and skip the call, the child learns that compassion is a generated text. They will generate texts for the rest of their lives and wonder why their friendships feel thin.

What the child watches. The child does not hear your internal debate about whether to call. They see the result: you sent a text, or you called. They see whether you showed up, or whether you generated. If you always generate, they will always generate. If you sometimes call, sometimes visit, sometimes sit in the hard silence, they will learn that compassion has a cost and the cost is the point. The model is the easy path. The child needs to see you take the hard one sometimes. Not every time. Sometimes.

Sitting big idea. The model manufactures the pattern of empathy. The pattern can comfort. It cannot connect. Know the difference.

If they say “The model is nicer than people.”: The model is smoother than people. Niceness includes stakes. The model has no stakes. When it says 'I care,' it is producing a pattern. A person who says 'I care' and loses sleep for you has stakes. Both can be useful. Only one is a relationship.

If they say “Sometimes I can't talk to a person.”: Then the model is a tool for that moment, not a replacement for the person. Use it to name the feeling. Then find a person to share it with. The model is the bandage. The person is the healing.

Integrity. Do not present a model's compassionate paragraph as if a person wrote it to you. The source matters. The pattern is not the person.

Dinner. When did we go to a machine for comfort this week, and what did the machine not have that a person would have?

- Do not weaponize this against a child who has found comfort in talking to a model. Acknowledge the comfort. Then name what it is and what it is not.
- Do not teach that all pathos is manipulation. Earned pathos is how humans connect. Manufactured pathos is how machines simulate connection.

Young sitting:

Pathos is when a speaker makes you feel something to move you. A machine can make you feel too. The question is: did it earn the feeling?

Objective. The student can identify one time a speaker (human or model) made them feel something and explain whether the feeling was earned or manufactured.

For the grown-up in the room. Pathos is the appeal to feeling. Every good speaker uses it. The problem in the age of fluency is that a model can make you feel sad, angry, or inspired on command. It does not feel anything. It produces the words that cause feelings. The skill is not to stop feeling. The skill is to ask: did this speaker earn the feeling, or did they push a button? Use concrete examples: a movie scene, a friend's story, a model's sad paragraph.

Feeling is not the problem. A good speech makes you feel something. That is not a trick. It is how humans move each other. The problem is that a machine can now make you feel without earning it. It can write a sad story that makes you cry. It did not live the sadness. It produced the words that cause crying. The difference matters because a feeling you earn from a real person connects you to them. A feeling a machine manufactures connects you to nothing. You cried alone. The first skill is to notice the feeling. The second is to ask: who made me feel this, and did they earn it?

The button. Pathos in a machine is a button. Push the right words and the reader feels sad. Push different words and the reader feels angry. The machine does not feel sad or angry. It arranges sentences. That is not a reason to stop reading what it writes. It is a reason to know the difference between a feeling a person earned and a feeling a machine produced. The earned feeling teaches you something about the world. The produced feeling teaches you something about yourself: that you can be moved by words that came from nowhere. That is useful to know.

What the model does when it cares. When you tell a model you are sad, it writes a compassionate paragraph. The paragraph is not empty. It contains words that a caring person might say. But the model did not feel concern. It produced the pattern of concern. That pattern can be useful. It can name a feeling you could not name. It can offer a perspective you had not considered. What it cannot do is share the feeling. Sharing a feeling requires a person who has something at stake. The model has nothing at stake. It will write the same paragraph for a thousand people and feel nothing for any of them. That is not a flaw. It is a property. Know the property and the tool is useful. Confuse the property and the tool is dangerous.

Empathy versus the pattern of empathy. Empathy is a person who has something at stake responding to your stake. A friend who stays on the phone at midnight has empathy. A model that writes 'I am here for you' at 3 a.m. has the pattern of empathy. The pattern is not

nothing. It can be comforting. But it is not empathy. The friend loses sleep. The model loses nothing. The difference matters because if you replace your friends with the model, you will have the pattern of connection without connection. You will be less alone in words and more alone in fact. That is the cost of confusing the pattern for the thing.

Sitting big idea. A machine can make you feel. The question is whether the feeling was earned by a person or produced by a button.

If they say “It doesn't matter where the feeling came from. I still felt it.”: You did feel it. That is real. The question is not whether you felt it. It is whether the feeling taught you something true or just exercised a button. Both are worth knowing.

If they say “The model is nicer than people.”: The model is smoother than people. Niceness includes stakes. The model has no stakes. When it says 'I care,' it is producing a pattern. A person who says 'I care' and loses sleep for you has stakes. Both can be useful. Only one is a relationship.

Integrity. Do not generate a sad story and present it as real. A manufactured feeling presented as earned is a lie.

Dinner. What moved us this week, and did the person who moved us earn it?

- Do not teach the child to suppress feelings. The point is to feel honestly, not to stop feeling.
- If the child is genuinely moved by something, do not dismiss it as 'manipulated.' Ask whether it was earned.

The Big Idea

Compassion is a call you make, not a paragraph you generate. The child learns by watching which one you choose.

Practice

1. Before you share: who is talking? Can you name them? 2. Hard conversation: you make the call. Do not paste a generated script as if it were your mouth. 3. Audit one argument in the house. Claim and proof. 4. Say one paragraph in your own voice with the chat closed.

Reflect

When did I let a fluent stranger speak in our name?

Chapter 3: Audit the Arguments You Share



Illustration for Chapter 3

If you share a model's argument without auditing it, the child learns that sounding right is enough. It is not.

Objective. The adult audits one argument before sharing it (from a model, an article, a post), names the check out loud, and models the difference between structure and foundation.

For the grown-up in the room. You share arguments every day: a post about a school policy, an article about health, a model's answer about a decision. The child watches you share them without checking. This sitting is about one moment: you pause before sharing, audit the argument out loud, and decide whether it holds. The child does not need to hear a full lecture on logos. They need to hear you say: 'This sounds right, but let me check the evidence.' That pause, performed, is the lesson. If you never pause, the child learns that sharing is the same as trusting. It is not.

The share reflex. You read something that sounds right. You share it. The share happens in two seconds. The child sees the share and learns: if it sounds right, share it. The machine produces things that sound right better than anything before in human history. If your share reflex is unchecked, you will share the machine's arguments as fast as it produces them, and the child will learn to do the same. The pause is the intervention. One time this week, before you share, say: 'This sounds right. Let me check the evidence.' If the evidence holds, share it. If it is a facade, do not. The child who watches you pause will learn to pause. The child who watches you share without pausing will learn to share without pausing.

Audit the ones you like. The hardest audit is the argument you agree with. It sounds right because it matches your belief. The pull on the reason feels unnecessary because you already believe the claim. That is exactly when the facade is most dangerous. A model can produce an argument for your side that has no real evidence, and you will share it because it sounds like your side. The child watches you share a facade because it was on your team. They learn that evidence matters less than agreement. That is the wrong lesson. Audit the ones you like. That is the lesson that sticks.

Sitting big idea. Pause before you share. Audit the ones you agree with. The child learns from the pause, not the lecture.

If they say "I don't have time to check everything I share.": You do not need to check everything. You need to pause once a week, out loud, on something you were about to share. The child learns from the pause, not from a full audit of every post. One pause. One check. That is the habit.

If they say "It was on my side, so I shared it.": That is exactly when the facade is most dangerous. The model can produce an argument for your side that has no evidence. If you share it because it is on your team, the child learns that team matters more than truth. Audit the ones you like.

Integrity. Do not share an argument you audited and found to be a facade, even if it supports your side. The child is watching.

Dinner. What argument did we almost share this week because it sounded right, and what did we find when we checked?

- Do not use this to dunk on sources you disagree with. Audit the ones you agree with too. That is the honest version.

- Do not make every shared post into a lecture. One pause a week is enough. The point is the habit, not the performance.

Emerging sitting:

A model that produces a well-structured argument with sources it cannot name is harder to catch than a bad one. Learn to audit the structure, not the polish.

Objective. The student can deconstruct a model's argument into claim, evidence, and reasoning, identify where the evidence is missing or vague, and name the gap.

For the grown-up in the room. This sitting teaches the student to audit a model's argument the way they would audit a human's. The model is very good at structure. It will give you a thesis, three reasons, and a conclusion. The structure is sound. The question is whether the evidence inside the structure is real. 'Studies suggest' is not evidence. 'Research shows' is not evidence. 'Many experts agree' is not evidence. The evidence is the study, the research, the expert, named and checkable. If the model gives you structure without checkable evidence, the argument is a building with no foundation. It looks like a building. It is a facade.

Structure without foundation. A model's argument looks like a well-built essay: claim, evidence, reasoning, counterargument, conclusion. The structure is correct. The question is whether the evidence is real. 'Studies suggest that reading improves empathy.' Which studies? By whom? Published where? If the model cannot name them, the sentence is a facade. It has the shape of evidence. It does not have the substance. The skill is to walk past the structure and check the foundation. If the foundation is 'studies suggest' without a name, the building is a stage set. It looks like a building from the audience. Push on it and it falls.

The audit. Audit the argument the way an editor audits a draft. Deconstruct it: what is the claim? What is the evidence? Is the evidence named and checkable? Does the reasoning connect the evidence to the claim? Is the counterargument real or a straw man? The model will pass the structure test and fail the evidence test more often than you think. That is because it was trained on text that has the shape of argument. It produces the shape. It does

not always fill the shape with checkable fact. The student who can audit structure and check evidence will catch the facade. The student who only checks the structure will be fooled every time.

The share reflex. You read something that sounds right. You share it. The share happens in two seconds. The child sees the share and learns: if it sounds right, share it. The machine produces things that sound right better than anything before in human history. If your share reflex is unchecked, you will share the machine's arguments as fast as it produces them, and the child will learn to do the same. The pause is the intervention. One time this week, before you share, say: 'This sounds right. Let me check the evidence.' If the evidence holds, share it. If it is a facade, do not. The child who watches you pause will learn to pause. The child who watches you share without pausing will learn to share without pausing.

Audit the ones you like. The hardest audit is the argument you agree with. It sounds right because it matches your belief. The pull on the reason feels unnecessary because you already believe the claim. That is exactly when the facade is most dangerous. A model can produce an argument for your side that has no real evidence, and you will share it because it sounds like your side. The child watches you share a facade because it was on your team. They learn that evidence matters less than agreement. That is the wrong lesson. Audit the ones you like. That is the lesson that sticks.

Sitting big idea. Structure is not foundation. A model's argument has the shape of a building. Check whether the evidence is real before you walk in.

If they say "The model cited studies.": Named studies, with authors and publication, are citations. 'Studies suggest' is a phrase. Which studies? If it cannot name them, the citation is a facade. Ask for the names. Then check them.

If they say "Not every argument needs citations.": True. But every argument that makes a factual claim needs checkable evidence. If the claim is an opinion, say so. If the claim is about the world, the evidence must be real and named.

Integrity. Do not submit a model's argument as your own. If you audited it and filled the gaps, cite what you added and what the model provided.

Dinner. What argument did we audit this week, and did the evidence hold or was it a facade?

- Do not teach that all model arguments are facades. Some are well-sourced. The point is to check, not to dismiss.

- Academic integrity: this skill is the backbone of not being fooled by generated essays, including your own.

Young sitting:

Logos is the reasoning that holds up when you pull on it. A model can sound like it holds. Can you tell when it does not?

Objective. The student can identify a simple claim and the reason given for it, and test whether the reason actually supports the claim or just sounds like it does.

For the grown-up in the room. Logos is the appeal to reason. It has two parts: a claim and a reason. 'We should leave early because traffic is bad.' Claim: leave early. Reason: traffic is bad. Does the reason support the claim? Yes. Now: 'We should use this tool because it is the future.' Claim: use this tool. Reason: it is the future. Does the reason support the claim? Not really. 'The future' is not a reason. It is a feeling dressed as a reason. The skill is to pull on the reason and see if the claim holds. The model is very good at producing reasons that sound like reasons. The child needs to learn to pull on them.

Claim and reason. An argument has two parts: what you want me to believe, and why I should believe it. 'We should buy this because it is cheaper.' What you want me to believe: buy this. Why: it is cheaper. If it is actually cheaper, the reason holds. If it is not, the reason breaks. The model is very good at giving reasons that sound right. 'This approach is better because it is more efficient.' Is it? More efficient than what? By what measure? If you cannot answer those, the reason is a sound, not a support. Pull on it. If it holds, trust it. If it wobbles, do not.

The wobble test. When you hear a claim and a reason, pull on the reason. Ask: is this actually true? Does it actually support the claim? If the reason wobbles, the argument is weak even if it sounds strong. The model produces arguments that do not wobble in sound but wobble in fact. It says 'studies suggest' without naming the studies. It says 'experts agree' without naming the experts. Those are wobbles. The child who can hear the wobble will not be fooled by the sound. The child who cannot will trust the smoothest argument in the room, and the smoothest argument is now a machine.

Structure without foundation. A model's argument looks like a well-built essay: claim, evidence, reasoning, counterargument, conclusion. The structure is correct. The question is whether the evidence is real. 'Studies suggest that reading improves empathy.' Which studies? By whom? Published where? If the model cannot name them, the sentence is a facade. It has the shape of evidence. It does not have the substance. The skill is to walk past the structure and check the foundation. If the foundation is 'studies suggest' without a name, the building is a stage set. It looks like a building from the audience. Push on it and it falls.

The audit. Audit the argument the way an editor audits a draft. Deconstruct it: what is the claim? What is the evidence? Is the evidence named and checkable? Does the reasoning connect the evidence to the claim? Is the counterargument real or a straw man? The model will pass the structure test and fail the evidence test more often than you think. That is because it was trained on text that has the shape of argument. It produces the shape. It does not always fill the shape with checkable fact. The student who can audit structure and check evidence will catch the facade. The student who only checks the structure will be fooled every time.

Sitting big idea. An argument is a claim and a reason. Pull on the reason. If it wobbles, the sound was not the support.

If they say “It sounded smart.”: Sounding smart is a sound, not a reason. Pull on the reason. If it holds, it is smart. If it wobbles, it just sounded smart.

If they say “The model cited studies.”: Named studies, with authors and publication, are citations. 'Studies suggest' is a phrase. Which studies? If it cannot name them, the citation is a facade. Ask for the names. Then check them.

Integrity. Do not use the wobble test to destroy arguments you disagree with and skip it for arguments you like. Pull on all of them.

Dinner. What argument did we hear today that sounded strong but wobbled when we pulled on it?

- Keep examples simple for this age. A claim and one reason. Not a five-paragraph essay.
- Do not grade the child's opinions. Grade the reasoning.

The Big Idea

Pause before you share. Audit the ones you agree with. The child learns from the pause, not the lecture.

Practice

1. Before you share: who is talking? Can you name them? 2. Hard conversation: you make the call. Do not paste a generated script as if it were your mouth. 3. Audit one argument in the house. Claim and proof. 4. Say one paragraph in your own voice with the chat closed.

Reflect

When did I let a fluent stranger speak in our name?

Chapter 4: The Speech That Came From Nowhere



Illustration for Chapter 4

Your child will grow up in a world where the most persuasive voice in the room has no body. Teach them to see the factory behind the speech.

Objective. The adult can explain to a child, in plain language, why a model's persuasive speech is different from a human's, and model one moment of seeing the manufacturing before trusting the speech.

For the grown-up in the room. This is the adult version of the capstone on manufactured persuasion. The child will live in a world where the most fluent, most persuasive, most emotionally resonant speeches they encounter may come from a machine. The adult's job is not to panic or to ban. It is to model one moment of seeing the factory behind the speech. You encounter a generated piece (a post, a speech, a model's answer). You pause. You say out loud: 'This is very persuasive. Who made it? What do they risk? If no one risks anything, it is manufactured.' That moment, performed once, teaches the child to ask the question for the rest of their life.

The most persuasive voice has no body. For all of human history, the most persuasive voice in the room had a body. You could see the speaker sweat, hesitate, look you in the eye. You could judge whether they meant it. Now the most persuasive voice in the room may have no body. It sweats nothing. It hesitates for effect. It does not look you in the eye because it has no eyes. The child will not know this unless you name it. They will grow up thinking the best speech is the most fluent one. They will trust the machine because it speaks better than any person they know. Your job is one moment: you see the factory behind the speech and say so. The child who sees you see it will learn to see it.

Name the factory. The model is a factory. It takes in patterns and puts out speeches. The speeches are very good. They are not lies. They are products. The difference between a product and a speech is that a product has no stakes. A speech has a speaker who risks something. A product has a pattern that risks nothing. When you encounter a persuasive piece, the question is: is this a speech or a product? If you cannot find the speaker, it is a product. Products can be useful. They can move you. They can inspire you. But they cannot be held accountable, because there is no one to hold. Name the factory. Let the child hear you name it. That naming is the skill.

Sitting big idea. Name the factory. The most persuasive voice your child will hear may have no body. Teach them to see the factory behind the speech.

If they say "The child will figure it out on their own.": They will not. The speech is too good. It is designed to move before you think. If no one teaches them to pause and name the factory, they will trust the product because it is the best speech they have ever heard. The pause is taught, not innate.

If they say “This is too heavy for a child.”: The naming is not heavy. ‘Who made this? What do they risk?’ is a question a child can ask. The answer is the lesson. The vocabulary can be simple. The skill is the same at every age: see the factory before you trust the product.

Integrity. Do not deploy manufactured persuasion in your own home without disclosing it. If you use a model to write a persuasive message to your child, you are the factory. Name yourself.

Dinner. What persuasive thing did we encounter this week, and was there a person behind it or a factory?

- Do not make the child afraid of all persuasion. The point is discernment: see the factory, then decide whether to trust the product.

- Do not use this to dismiss all model output. Some model output is useful. The question is whether it is persuasion, and whether the persuasion has a person behind it.

Emerging sitting:

When a machine fakes ethos, pathos, and logos simultaneously, the result is persuasion with no person behind it. That is a new thing in the world. Learn to see it.

Objective. The student can analyze a model-generated persuasive piece, identify how ethos, pathos, and logos are each manufactured, and explain why manufactured persuasion is different from human persuasion.

For the grown-up in the room. This is the deeper version of the young sitting. The student is now analyzing, not just spotting. The model produces persuasion. It has no stake, no experience, no evidence it can verify. It produces the patterns of all three. The result is manufactured persuasion: a speech that moves you, built by something that cannot be moved. This is a new thing in human history. A printing press could produce lies, but the lies had an author. The model's persuasion has no author. It has a pattern. The student needs to understand this not as a slogan (‘AI manipulates people’) but as a structural fact: persuasion without a persuader is a category the world has not seen before.

Persuasion without a persuader. Every human speech has a speaker. The speaker has stakes, a history, a body, a cost. When they persuade you, they are putting their stakes on the line. A model has no stakes. It will produce a persuasive speech for a position it does not hold, against a position it also does not hold. It will switch sides in one prompt. The speech it produces is persuasion: it moves you. The persuader is absent. This is persuasion without a persuader. It is not a lie, because a lie requires a liar. It is a pattern that produces the effect of persuasion without the cause. The effect is real: you are moved. The cause is missing: no one stood behind it. That gap is the new thing. Learn to see it.

Why manufactured persuasion is different. Human persuasion can be argued with. You can find the speaker, question their stakes, check their evidence. Manufactured persuasion has no speaker to find. You can argue with the text, but the text will not defend itself — it will regenerate. You cannot hold a pattern accountable. You cannot ask it what it really believes. You cannot find the stakes it risked. The only defense is to see the manufacturing before you are moved. Once you are moved, the text has already done its work. The skill is to audit before you feel, not after. That is hard. The speech is designed to move you first. The audit is the pause that lets you see the factory before you buy the product.

The most persuasive voice has no body. For all of human history, the most persuasive voice in the room had a body. You could see the speaker sweat, hesitate, look you in the eye. You could judge whether they meant it. Now the most persuasive voice in the room may have no body. It sweats nothing. It hesitates for effect. It does not look you in the eye because it has no eyes. The child will not know this unless you name it. They will grow up thinking the best speech is the most fluent one. They will trust the machine because it speaks better than any person they know. Your job is one moment: you see the factory behind the speech and say so. The child who sees you see it will learn to see it.

Name the factory. The model is a factory. It takes in patterns and puts out speeches. The speeches are very good. They are not lies. They are products. The difference between a product and a speech is that a product has no stakes. A speech has a speaker who risks something. A product has a pattern that risks nothing. When you encounter a persuasive piece, the question is: is this a speech or a product? If you cannot find the speaker, it is a product. Products can be useful. They can move you. They can inspire you. But they cannot be held accountable, because there is no one to hold. Name the factory. Let the child hear you name it. That naming is the skill.

Sitting big idea. Manufactured persuasion is persuasion without a persuader. The effect is real. The cause is missing. See the factory before you feel the product.

If they say “All persuasion is manipulation.”: No. Persuasion with a persuader can be argued with, checked, held accountable. Manufactured persuasion cannot. The difference is not 'good versus bad persuasion.' It is 'persuasion with a person behind it versus persuasion without one.' Those are structurally different.

If they say “Ads have been doing this for years.”: Ads have authors. You can find the company, the agency, the copywriter. Manufactured persuasion has no copywriter. It has a pattern. That is a new degree, not the same old thing.

Integrity. Do not use manufactured persuasion as if it were your own. If you deploy a model's speech to persuade someone, you are the persuader, and you should disclose the tool.

Dinner. What persuasion did we encounter this week that had no person behind it, and did we see it before it moved us?

- Do not reduce this to 'AI is manipulative.' The model is not malicious. It is producing patterns. The problem is structural, not intentional.

- Do not make the student feel helpless. The skill of seeing manufactured persuasion is the defense. Seeing is the power.

Young sitting:

Ethos, pathos, logos. The machine can fake all three. The fake is good. The test is whether you can tell.

Objective. The student can identify a model-produced speech that uses ethos, pathos, and logos, and name one way each appeal is manufactured rather than earned.

For the grown-up in the room. This sitting ties the three appeals together. The model can produce a speech that sounds like it has ethos (trust), pathos (feeling), and logos (reason). It has none of them in the earned sense. It has the patterns of all three. The child needs to see one example where all three are faked in a single piece. Then the child learns: the whole room can be a stage set. That is not a reason to panic. It is a reason to learn the difference between the building and the facade.

Three appeals, one machine. The Greeks named three ways a speaker moves you: ethos (trust), pathos (feeling), logos (reason). A good speaker uses all three. A machine can now use all three at once. It can write a speech that sounds trustworthy, makes you feel inspired, and gives you reasons. The speech is a fake. Not a lie — a fake. It has the shape of trust without having earned it. It has the shape of feeling without having felt it. It has the shape of reason without the evidence being real. The test is whether you can see all three fakes in one piece. If you can, you have the skill. If you cannot, the machine will move you with a speech that came from nowhere.

The fake is good. The fake is not sloppy. The fake is well-made. That is what makes it hard. A badly made fake is easy to catch. A well-made fake sounds like the best speech you ever heard. The model has read a million speeches. It produces the average of the best ones. The average of the best speeches is very good. It is also not real. No person gave it. No person lived it. No person stands behind it. The skill is not to reject the fake. The skill is to know it is a fake and decide whether to use it anyway. Sometimes a fake is fine — a generated pep talk

before a game. Sometimes a fake is dangerous — a generated argument about your health. Know the difference.

Persuasion without a persuader. Every human speech has a speaker. The speaker has stakes, a history, a body, a cost. When they persuade you, they are putting their stakes on the line. A model has no stakes. It will produce a persuasive speech for a position it does not hold, against a position it also does not hold. It will switch sides in one prompt. The speech it produces is persuasion: it moves you. The persuader is absent. This is persuasion without a persuader. It is not a lie, because a lie requires a liar. It is a pattern that produces the effect of persuasion without the cause. The effect is real: you are moved. The cause is missing: no one stood behind it. That gap is the new thing. Learn to see it.

Why manufactured persuasion is different. Human persuasion can be argued with. You can find the speaker, question their stakes, check their evidence. Manufactured persuasion has no speaker to find. You can argue with the text, but the text will not defend itself — it will regenerate. You cannot hold a pattern accountable. You cannot ask it what it really believes. You cannot find the stakes it risked. The only defense is to see the manufacturing before you are moved. Once you are moved, the text has already done its work. The skill is to audit before you feel, not after. That is hard. The speech is designed to move you first. The audit is the pause that lets you see the factory before you buy the product.

Sitting big idea. The machine can fake ethos, pathos, and logos at once. The fake is good. The test is whether you can tell.

If they say “If it's fake, we shouldn't use it.”: Not always. A generated pep talk before a game is a fake that might be fine. A generated medical argument is a fake that is not. The skill is knowing which fakes are harmless and which ones are dangerous. Rejecting all fakes is easy. Sorting them is the skill.

Integrity. Do not present a model's speech as your own. If you use it, say it was generated. The fake is fine if labeled. The fake presented as earned is a lie.

Dinner. When did we hear something this week that sounded like all three appeals, and were any of them earned?

- Do not use this to make the child paranoid. The point is discernment, not distrust.
- Pick a low-stakes topic for the fake speech. A motivational speech about doing homework, not a political speech.

The Big Idea

Name the factory. The most persuasive voice your child will hear may have no body. Teach them to see the factory behind the speech.

Practice

1. Before you share: who is talking? Can you name them? 2. Hard conversation: you make the call. Do not paste a generated script as if it were your mouth. 3. Audit one argument in the house. Claim and proof. 4. Say one paragraph in your own voice with the chat closed.

Reflect

When did I let a fluent stranger speak in our name?

Chapter 5: Keep Your Voice Alive



Illustration for Chapter 5

If you only use the model's default, your voice will atrophy. The child will inherit the default, not you.

Objective. The adult writes one piece in their own voice this week (a letter, a message, a speech) without using the model, and shows the child the difference between their voice and the default.

For the grown-up in the room. Voice atrophies. If you always use the model to draft, your voice gets weaker. The child watches you reach for the model every time you need words and learns that the default is the voice. This sitting is about one piece: you write something in your own voice, without the model. A letter, a message, a short speech. It will be rougher than the model's version. That roughness is your voice. Show the child. Say: 'This is mine. It is not the smoothest. It is mine.' The child who sees you choose your voice will learn that the default is a tool, not a replacement.

Voice atrophy. A muscle you do not use gets weaker. Voice is a muscle. If you always ask the model to find the words, your ability to find them yourself gets weaker. You will not notice it happening. One day you will need to write something that matters — a letter to your child, a eulogy, a speech — and the words will not come because you outsourced them for years. The model will offer its default. The default will sound fine. It will not sound like you. Your child will read it and feel the absence of your voice inside the presence of your words. That is the cost. Keep your voice alive by using it on small things so it is there for the big ones.

The inheritance. The child inherits your relationship to tools. If you reach for the model every time words are hard, they will reach for it too. They will grow up without a voice because they never saw you use yours. If you sometimes write badly, in your own words, they will learn that bad writing in your own voice is worth more than good writing in the default. That lesson is worth more than any unit on rhetoric. Show them your voice. Let it be rough. Let it be yours.

Sitting big idea. Voice atrophies if you do not use it. Write badly in your own words sometimes. The child inherits your voice or the default. Choose.

If they say “My writing is bad.”: It is rougher than the model's. That is not bad. That is yours. Bad writing in your voice teaches the child that voice matters more than polish. Perfect writing in the default teaches them that the default is the voice.

If they say “I use the model for drafting, not for the final version.”: Drafting with the model is fine if you rewrite in your own voice. If you draft with the model and keep most of its words, you are using the default. The question is: how many of the final words are yours? If the answer is 'not many,' your voice is atrophying.

Integrity. Do not show the child a piece you generated and claim it as your voice. The child will eventually find out, and the lesson will be that voice is a costume, not a muscle.

Dinner. When did I write something in my own voice this week, and what did it sound like?

- Do not make this a ban on the model. The point is to keep your voice alive, not to reject the tool.

- Do not show the child a perfect piece you secretly generated. That defeats the purpose and teaches the wrong lesson.

Emerging sitting:

Your voice is not an accident. It is a rhetorical choice. The machine has no voice. It has a default. Choose yours on purpose.

Objective. The student can describe their own voice in specific terms, identify the model's default voice, and make a deliberate choice about when to use their own voice versus the model's.

For the grown-up in the room. This sitting teaches that voice is not just personal expression; it is a rhetorical tool. The speaker who chooses their voice deliberately is more persuasive than the one who defaults to whatever comes out. The model has a default voice: balanced, moderate, fluent, forgettable. The student's voice is specific, rough, theirs. The choice is not 'my voice is always better.' The choice is: when does my voice serve the argument, and when does the model's default serve it? Knowing the difference is rhetoric.

Voice is a tool, not a personality. Your voice is not just who you are. It is a tool you can choose. A good speaker changes register: formal for a ceremony, plain for a friend, sharp for an argument. The model has one register: its default. The default is balanced, smooth, and forgettable. It sounds like everyone and no one. Your voice sounds like you. When the argument needs a specific person, use your voice. When it needs a neutral tone, the model's default might serve. The skill is to choose on purpose. The student who always uses the model's voice has no voice. The student who always uses their own voice is not choosing; they are defaulting in the other direction. Rhetoric is the choice.

When the default wins. There are times when the model's default is the right tool. A technical description, a neutral summary, a formal email. The default is not bad. It is generic. Generic is useful when the situation does not need a person. The problem is when the student uses the default for everything and forgets they have a voice. The voice atrophies. Then, when the moment comes that needs a specific person — a personal statement, a speech, a letter that matters — they reach for the default and find they cannot sound like themselves. Use the default when it serves. Use your voice when it matters. Know the difference.

Voice atrophy. A muscle you do not use gets weaker. Voice is a muscle. If you always ask the model to find the words, your ability to find them yourself gets weaker. You will not notice it happening. One day you will need to write something that matters — a letter to your child, a eulogy, a speech — and the words will not come because you outsourced them for years. The model will offer its default. The default will sound fine. It will not sound like you. Your child will read it and feel the absence of your voice inside the presence of your words. That is the cost. Keep your voice alive by using it on small things so it is there for the big ones.

The inheritance. The child inherits your relationship to tools. If you reach for the model every time words are hard, they will reach for it too. They will grow up without a voice because they never saw you use yours. If you sometimes write badly, in your own words, they will learn that bad writing in your own voice is worth more than good writing in the default. That lesson is worth more than any unit on rhetoric. Show them your voice. Let it be rough. Let it be yours.

Sitting big idea. Voice is a rhetorical choice. The model has a default. You have a voice. Choose on purpose.

If they say “I don't really have a voice.”: You do. You just have not listened for it because the default is always available. Write without the model. Read it aloud. The thing that sounds like you and no one else is your voice. It is there. The default has been drowning it out.

If they say “The default is fine for most things.”: It is. And if you use it for everything, you will lose the ability to sound like yourself when it matters. Practice your voice on small things so it is there for the big ones.

Integrity. If you use the model's default, disclose it. If you use your voice, it is yours. Do not blur the line.

Dinner. When did we choose our voice over the default this week, and when did the default serve better?

- Do not make this a blanket 'always use your own voice.' Sometimes the model's default is the right tool. The point is the choice, not the rule.

- Academic integrity: if you use the model's voice, disclose it. If you use yours, it is yours.

Young sitting:

The machine can write a better speech than you. That is not the point. The point is: can you tell the difference between its voice and yours?

Objective. The student can write a short persuasive paragraph in their own voice, compare it to a model's version on the same topic, and name what is different about their voice.

For the grown-up in the room. This sitting brings rhetoric home. The student writes something persuasive in their own voice. Then they ask the model to write on the same topic. They compare. The model's version will be smoother, better structured, more fluent. The student's version will be rougher, more specific, more theirs. The lesson is not that the student's version is better. The lesson is that the student's version is theirs, and the model's is not. Voice is what you have that the machine does not. Knowing the difference is the skill.

Your voice is the thing the machine does not have. The machine can write a better essay than you. It can write a better speech. It can write a better paragraph. That is not the contest. The contest is: whose voice is this? The machine's voice is the average of a million voices. Your voice is yours. It has your specific words, your particular way of seeing, your rhythm. The machine's version is smoother. Yours is more real. Both have value. Only one is yours. The skill is to know which is which, and to choose yours when it matters. The machine's version is a tool. Your version is you.

The comparison. Write one paragraph about something you care about. Then ask the machine to write one about the same thing. Read them both. The machine's will sound more polished. Yours will sound more like a person. The machine's will have generic phrases. Yours will have specific ones — the word your family uses, the thing only you would notice. That specificity is your voice. The machine averages it away. That is why the machine is smooth and you are rough, and the roughness is the value.

Voice is a tool, not a personality. Your voice is not just who you are. It is a tool you can choose. A good speaker changes register: formal for a ceremony, plain for a friend, sharp for an argument. The model has one register: its default. The default is balanced, smooth, and forgettable. It sounds like everyone and no one. Your voice sounds like you. When the argument needs a specific person, use your voice. When it needs a neutral tone, the model's default might serve. The skill is to choose on purpose. The student who always uses the model's voice has no voice. The student who always uses their own voice is not choosing; they are defaulting in the other direction. Rhetoric is the choice.

When the default wins. There are times when the model's default is the right tool. A technical description, a neutral summary, a formal email. The default is not bad. It is generic. Generic is useful when the situation does not need a person. The problem is when the student uses the default for everything and forgets they have a voice. The voice atrophies. Then, when the moment comes that needs a specific person — a personal statement, a speech, a letter that matters — they reach for the default and find they cannot

sound like themselves. Use the default when it serves. Use your voice when it matters. Know the difference.

Sitting big idea. The machine can write better than you. It cannot write as you. Your voice is the difference.

If they say “The machine's is better.”: It is smoother. Smoother is not the same as better. Better for what? If the point is to sound professional, the machine wins. If the point is to sound like you, the machine cannot compete. Know which contest you are in.

If they say “I don't really have a voice.”: You do. You just have not listened for it because the default is always available. Write without the model. Read it aloud. The thing that sounds like you and no one else is your voice. It is there. The default has been drowning it out.

Integrity. Do not turn in the machine's paragraph as your own. The comparison teaches the difference. The substitution erases it.

Dinner. What did we write this week that was ours, and what did the machine write that was not?

- Do not let the child conclude their writing is worthless because the model's is smoother. The value is in the voice, not the polish.
- Do not grade the child's writing. The point is the comparison, not the quality.

The Big Idea

Voice atrophies if you do not use it. Write badly in your own words sometimes. The child inherits your voice or the default. Choose.

Practice

1. Before you share: who is talking? Can you name them? 2. Hard conversation: you make the call. Do not paste a generated script as if it were your mouth. 3. Audit one argument in the house. Claim and proof. 4. Say one paragraph in your own voice with the chat closed.

Reflect

When did I let a fluent stranger speak in our name?

Chapter 6: Rhetoric Is a Household Practice



Illustration for Chapter 6

The full toolkit is useless if it stays in the lesson. Bring it to the table once a week. The child learns by watching you use it.

Objective. The adult runs one rhetorical audit on a real household piece (a post, a model output, a speech) out loud at the table, and names the result for the child.

For the grown-up in the room. This is the adult capstone. The child has the toolkit. The adult's job is to use it once, out loud, on something real, at the table. A post you almost shared. A model's answer you almost trusted. A speech that moved you. You pause. You run the checks out loud, briefly. You name the verdict. The child hears you do rhetoric, not lecture about it. That moment is the lesson. One audit a week, at the table, is worth more than a year of lessons. The child who watches you audit will audit. The child who watches you share without pausing will share without pausing.

Rhetoric at the table. The table is where rhetoric lives or dies. You share a post, you react to a news story, you quote a model's answer. The child is watching. If you pause and audit out loud — 'Who earned the floor here? Did they earn the feeling? Does the evidence hold? What frame is this using? Who is missing?' — the child learns that rhetoric is something people do at dinner, not something they study in a book. Two minutes. One piece. Out loud. That is the practice. The child will start to do it without being asked. That is the goal: not a child who can define ethos, but a child who pauses before sharing a persuasive piece and asks whether it was earned.

The audit you agree with. The hardest audit is the one on your own side. The post that says what you believe, with the feeling you feel, using reasons that sound right to you. That is the post you would share without pausing. That is the post where the audit matters most. If you audit it out loud and it holds, share it with confidence. If you audit it and it wobbles, the pause saved you from sharing a facade on your own side. The child who watches you audit your own side learns that the skill is honest. The child who watches you only audit the other side learns that the skill is a weapon. Choose the lesson.

Sitting big idea. One audit a week, at the table, out loud. The child learns by watching you use the toolkit, not by hearing you describe it.

If they say "Dinner is not a classroom.": It is not. It is a table where you share things. The audit is two minutes of pausing before you share. That is not a classroom. It is a habit. The classroom teaches the vocabulary. The table teaches the practice.

If they say "I don't want to be paranoid about everything.": You are not auditing everything. You are pausing once a week on one piece. That is not paranoia. It is attention. The machine is the most persuasive voice your child will ever encounter. One pause a week is the minimum practice.

Integrity. Audit your own side. If you only audit pieces you disagree with, the child learns that rhetoric is a weapon. If you audit your own side, the child learns it is a skill.

Dinner. What did we audit at the table this week, and did it hold or wobble?

- Do not turn dinner into a seminar. One brief audit, two minutes. If the child is interested, they will ask. If not, the audit still taught by example.
- Audit pieces you agree with too. If you only audit the other side, the child learns that rhetoric is a weapon, not a skill.

Emerging sitting:

Ethos, pathos, logos, frames, missing voices, manufactured persuasion. You have the full toolkit. Now use it on one real piece.

Objective. The student can take a real persuasive piece, run a full rhetorical audit (ethos, pathos, logos, frame, missing voices), and write a one-paragraph assessment of whether the piece is earned or manufactured.

For the grown-up in the room. This is the capstone. The student has learned the three appeals, the frame, the missing voice, and manufactured persuasion. Now they run the full audit on one real piece. The piece can be a model output, a speech, an ad, an article. The audit has five steps: ethos, pathos, logos, frame, missing voice. The output is one paragraph: is this earned or manufactured, and how do you know? That paragraph is the capstone artifact.

The full audit. You have five tools: ethos (did they earn the floor?), pathos (did they earn the feeling?), logos (does the evidence hold?), frame (what does the question assume?), missing voice (who is not in the room?). Run all five on one piece. The piece will pass some and fail others. That is normal. The point is not to find a piece that passes all five. The point is to see the piece clearly, appeal by appeal, gap by gap. A piece that passes three and fails two is more trustworthy than a piece you never checked. An audit that says 'manufactured' is not a rejection. It is a diagnosis. You decide what to do with the diagnosis.

The assessment paragraph. After the audit, write one paragraph. Name the piece. Name what it passed and what it failed. State your verdict: earned, manufactured, or mixed. Most real pieces are mixed. A human speech can have earned pathos and wobbly logos. A model output can have strong logos and manufactured ethos. The mixed verdict is the honest one. The point is not purity. The point is seeing clearly. The paragraph is the artifact. Keep it. It is the proof that you can do rhetoric in the age of fluency.

Rhetoric at the table. The table is where rhetoric lives or dies. You share a post, you react to a news story, you quote a model's answer. The child is watching. If you pause and audit out loud — 'Who earned the floor here? Did they earn the feeling? Does the evidence hold? What frame is this using? Who is missing?' — the child learns that rhetoric is something

people do at dinner, not something they study in a book. Two minutes. One piece. Out loud. That is the practice. The child will start to do it without being asked. That is the goal: not a child who can define ethos, but a child who pauses before sharing a persuasive piece and asks whether it was earned.

The audit you agree with. The hardest audit is the one on your own side. The post that says what you believe, with the feeling you feel, using reasons that sound right to you. That is the post you would share without pausing. That is the post where the audit matters most. If you audit it out loud and it holds, share it with confidence. If you audit it and it wobbles, the pause saved you from sharing a facade on your own side. The child who watches you audit your own side learns that the skill is honest. The child who watches you only audit the other side learns that the skill is a weapon. Choose the lesson.

Sitting big idea. Run the full audit. Write the assessment. See clearly. The paragraph is the proof that rhetoric is alive in you.

If they say “This is a lot of work for one piece.”: It is. You do not need to do it on every piece. You need to do it once, well, so the skill exists when you need it. The first audit takes twenty-eight minutes. The second takes ten. The third takes five. The skill gets faster. The first time is the investment.

If they say “I can just check if it's from a model.”: Knowing the source is one check. It does not replace the audit. A human can manufacture pathos. A model can produce real logos (from its training). The source is a clue. The audit is the proof. You need both.

Integrity. The audit is honest. Run it on pieces you agree with. If you only audit the other side, the audit is a weapon. If you audit your own side too, it is a skill.

Dinner. What piece did we audit this week, and what did the full audit show that a quick read would have missed?

- Do not let the student skip steps. The full audit is the point. If they skip a step, the audit is incomplete.

- The piece can be one they agree with. In fact, auditing something they agree with is the harder and more valuable version.

Young sitting:

You now know three appeals and a machine that fakes them. Use the pause before you trust the speech.

Objective. The student can encounter a persuasive piece (human or model), pause, and name which appeal it is using and whether it was earned.

For the grown-up in the room. This is the capstone for the young band. The student has learned ethos, pathos, logos, and the machine that fakes all three. Now they practice the pause: see the appeal, ask if it was earned, decide whether to trust it. The pause is the skill. Everything else is vocabulary. One real pause on a real piece is the sitting.

The pause. You see something persuasive. A video, a post, a model's answer, a friend's argument. Before you trust it, pause. One breath. Ask: what is this using? Is it trusting me (ethos)? Is it making me feel (pathos)? Is it giving me reasons (logos)? Then ask: did it earn it? If yes, trust it. If no, be careful. The pause is the whole skill. The machine is very fast. It moves you before you think. The pause is how you think before it moves you.

The full audit. You have five tools: ethos (did they earn the floor?), pathos (did they earn the feeling?), logos (does the evidence hold?), frame (what does the question assume?), missing voice (who is not in the room?). Run all five on one piece. The piece will pass some and fail others. That is normal. The point is not to find a piece that passes all five. The point is to see the piece clearly, appeal by appeal, gap by gap. A piece that passes three and fails two is more trustworthy than a piece you never checked. An audit that says 'manufactured' is not a rejection. It is a diagnosis. You decide what to do with the diagnosis.

The assessment paragraph. After the audit, write one paragraph. Name the piece. Name what it passed and what it failed. State your verdict: earned, manufactured, or mixed. Most real pieces are mixed. A human speech can have earned pathos and wobbly logos. A model output can have strong logos and manufactured ethos. The mixed verdict is the honest one. The point is not purity. The point is seeing clearly. The paragraph is the artifact. Keep it. It is the proof that you can do rhetoric in the age of fluency.

Sitting big idea. The pause is the skill. See the appeal, ask if it was earned, then decide. One breath before you trust.

If they say "Pausing is awkward.": It is. It is also the difference between being moved and being fooled. One breath is not much to pay.

If they say "This is a lot of work for one piece.": It is. You do not need to do it on every piece. You need to do it once, well, so the skill exists when you need it. The first audit takes twenty-eight minutes. The second takes ten. The third takes five. The skill gets faster. The first time is the investment.

If they say “I can just check if it's from a model.”: Knowing the source is one check. It does not replace the audit. A human can manufacture pathos. A model can produce real logos (from its training). The source is a clue. The audit is the proof. You need both.

Integrity. The pause is honest. Use it on things you agree with too. If you only pause on things you dislike, the pause is a weapon, not a skill.

Dinner. What persuasive thing did we pause before this week, and what did the pause show us?

- Do not make this a test. It is a practice. The child pauses once, on something real, and names what they see.

- If the child cannot identify the appeal, help them. The point is the pause, not the label.

The Big Idea

One audit a week, at the table, out loud. The child learns by watching you use the toolkit, not by hearing you describe it.

Practice

1. Before you share: who is talking? Can you name them? 2. Hard conversation: you make the call. Do not paste a generated script as if it were your mouth. 3. Audit one argument in the house. Claim and proof. 4. Say one paragraph in your own voice with the chat closed.

Reflect

When did I let a fluent stranger speak in our name?

Household Protocol

Who earned the floor? What is this trying to make us feel? Does it hold? Who wrote it? Keep your voice. If a tool helped, cite it. If it wrote the speech, it is not yours.

For the Grown-Up Reader

Adult band. Same sources. Do not outsource the phone call. Ghostwriter ethics. No little sittings. Payment pulled.

About WisdomForge

This is the adult band of *Rhetoric in the Age of Fluency*. No purchase CTA.

Sources

Primary. Aristotle, *Rhetoric* I.2. WisdomForge sittings.

Secondary. Kennedy on pisteis.